

1

Natural Language Processing

2

Jacob Eisenstein

3

July 1, 2018

4 Contents

5	Contents	1
6	1 Introduction	13
7	1.1 Natural language processing and its neighbors	13
8	1.2 Three themes in natural language processing	17
9	1.2.1 Learning and knowledge	17
10	1.2.2 Search and learning	19
11	1.2.3 Relational, compositional, and distributional perspectives	20
12	1.3 Learning to do natural language processing	22
13	1.3.1 Background	22
14	1.3.2 How to use this book	23
15	I Learning	27
16	2 Linear text classification	29
17	2.1 Naïve Bayes	32
18	2.1.1 Types and tokens	34
19	2.1.2 Prediction	35
20	2.1.3 Estimation	36
21	2.1.4 Smoothing and MAP estimation	38
22	2.1.5 Setting hyperparameters	38
23	2.2 Discriminative learning	39
24	2.2.1 Perceptron	40
25	2.2.2 Averaged perceptron	42
26	2.3 Loss functions and large-margin classification	43
27	2.3.1 Large margin classification	46
28	2.3.2 Support vector machines	47
29	2.3.3 Slack variables	48
30	2.4 Logistic regression	50
31	2.4.1 Regularization	51

32	2.4.2 Gradients	52
33	2.5 Optimization	52
34	2.5.1 Batch optimization	53
35	2.5.2 Online optimization	54
36	2.6 *Additional topics in classification	56
37	2.6.1 Feature selection by regularization	56
38	2.6.2 Other views of logistic regression	56
39	2.7 Summary of learning algorithms	58
40	3 Nonlinear classification	61
41	3.1 Feedforward neural networks	62
42	3.2 Designing neural networks	64
43	3.2.1 Activation functions	64
44	3.2.2 Network structure	65
45	3.2.3 Outputs and loss functions	66
46	3.2.4 Inputs and lookup layers	67
47	3.3 Learning neural networks	67
48	3.3.1 Backpropagation	69
49	3.3.2 Regularization and dropout	71
50	3.3.3 *Learning theory	72
51	3.3.4 Tricks	73
52	3.4 Convolutional neural networks	75
53	4 Linguistic applications of classification	81
54	4.1 Sentiment and opinion analysis	81
55	4.1.1 Related problems	83
56	4.1.2 Alternative approaches to sentiment analysis	84
57	4.2 Word sense disambiguation	85
58	4.2.1 How many word senses?	86
59	4.2.2 Word sense disambiguation as classification	87
60	4.3 Design decisions for text classification	88
61	4.3.1 What is a word?	88
62	4.3.2 How many words?	91
63	4.3.3 Count or binary?	92
64	4.4 Evaluating classifiers	92
65	4.4.1 Precision, recall, and <i>F</i> -MEASURE	93
66	4.4.2 Threshold-free metrics	95
67	4.4.3 Classifier comparison and statistical significance	96
68	4.4.4 *Multiple comparisons	99
69	4.5 Building datasets	99
70	4.5.1 Metadata as labels	100

71	4.5.2 Labeling data	100
72	5 Learning without supervision	107
73	5.1 Unsupervised learning	107
74	5.1.1 K -means clustering	108
75	5.1.2 Expectation Maximization (EM)	110
76	5.1.3 EM as an optimization algorithm	114
77	5.1.4 How many clusters?	115
78	5.2 Applications of expectation-maximization	116
79	5.2.1 Word sense induction	116
80	5.2.2 Semi-supervised learning	117
81	5.2.3 Multi-component modeling	118
82	5.3 Semi-supervised learning	119
83	5.3.1 Multi-view learning	120
84	5.3.2 Graph-based algorithms	121
85	5.4 Domain adaptation	122
86	5.4.1 Supervised domain adaptation	123
87	5.4.2 Unsupervised domain adaptation	124
88	5.5 *Other approaches to learning with latent variables	126
89	5.5.1 Sampling	126
90	5.5.2 Spectral learning	128
91	II Sequences and trees	135
92	6 Language models	137
93	6.1 N -gram language models	138
94	6.2 Smoothing and discounting	141
95	6.2.1 Smoothing	141
96	6.2.2 Discounting and backoff	142
97	6.2.3 *Interpolation	143
98	6.2.4 *Kneser-Ney smoothing	145
99	6.3 Recurrent neural network language models	146
100	6.3.1 Backpropagation through time	148
101	6.3.2 Hyperparameters	149
102	6.3.3 Gated recurrent neural networks	149
103	6.4 Evaluating language models	151
104	6.4.1 Held-out likelihood	151
105	6.4.2 Perplexity	152
106	6.5 Out-of-vocabulary words	153

107	7 Sequence labeling	155
108	7.1 Sequence labeling as classification	155
109	7.2 Sequence labeling as structure prediction	157
110	7.3 The Viterbi algorithm	159
111	7.3.1 Example	162
112	7.3.2 Higher-order features	163
113	7.4 Hidden Markov Models	163
114	7.4.1 Estimation	165
115	7.4.2 Inference	165
116	7.5 Discriminative sequence labeling with features	167
117	7.5.1 Structured perceptron	170
118	7.5.2 Structured support vector machines	170
119	7.5.3 Conditional random fields	172
120	7.6 Neural sequence labeling	177
121	7.6.1 Recurrent neural networks	177
122	7.6.2 Character-level models	179
123	7.6.3 Convolutional Neural Networks for Sequence Labeling	180
124	7.7 *Unsupervised sequence labeling	180
125	7.7.1 Linear dynamical systems	182
126	7.7.2 Alternative unsupervised learning methods	182
127	7.7.3 Semiring Notation and the Generalized Viterbi Algorithm	182
128	8 Applications of sequence labeling	185
129	8.1 Part-of-speech tagging	185
130	8.1.1 Parts-of-Speech	186
131	8.1.2 Accurate part-of-speech tagging	190
132	8.2 Morphosyntactic Attributes	192
133	8.3 Named Entity Recognition	193
134	8.4 Tokenization	195
135	8.5 Code switching	196
136	8.6 Dialogue acts	197
137	9 Formal language theory	199
138	9.1 Regular languages	200
139	9.1.1 Finite state acceptors	201
140	9.1.2 Morphology as a regular language	202
141	9.1.3 Weighted finite state acceptors	204
142	9.1.4 Finite state transducers	209
143	9.1.5 *Learning weighted finite state automata	214
144	9.2 Context-free languages	215
145	9.2.1 Context-free grammars	216

146	9.2.2 Natural language syntax as a context-free language	219
147	9.2.3 A phrase-structure grammar for English	221
148	9.2.4 Grammatical ambiguity	226
149	9.3 *Mildly context-sensitive languages	226
150	9.3.1 Context-sensitive phenomena in natural language	227
151	9.3.2 Combinatory categorial grammar	228
152	10 Context-free parsing	233
153	10.1 Deterministic bottom-up parsing	234
154	10.1.1 Recovering the parse tree	236
155	10.1.2 Non-binary productions	236
156	10.1.3 Complexity	237
157	10.2 Ambiguity	237
158	10.2.1 Parser evaluation	238
159	10.2.2 Local solutions	239
160	10.3 Weighted Context-Free Grammars	240
161	10.3.1 Parsing with weighted context-free grammars	241
162	10.3.2 Probabilistic context-free grammars	243
163	10.3.3 *Semiring weighted context-free grammars	245
164	10.4 Learning weighted context-free grammars	245
165	10.4.1 Probabilistic context-free grammars	246
166	10.4.2 Feature-based parsing	246
167	10.4.3 *Conditional random field parsing	247
168	10.4.4 Neural context-free grammars	249
169	10.5 Grammar refinement	250
170	10.5.1 Parent annotations and other tree transformations	251
171	10.5.2 Lexicalized context-free grammars	252
172	10.5.3 *Refinement grammars	256
173	10.6 Beyond context-free parsing	257
174	10.6.1 Reranking	257
175	10.6.2 Transition-based parsing	258
176	11 Dependency parsing	261
177	11.1 Dependency grammar	261
178	11.1.1 Heads and dependents	262
179	11.1.2 Labeled dependencies	263
180	11.1.3 Dependency subtrees and constituents	264
181	11.2 Graph-based dependency parsing	266
182	11.2.1 Graph-based parsing algorithms	268
183	11.2.2 Computing scores for dependency arcs	269
184	11.2.3 Learning	271

185	11.3 Transition-based dependency parsing	272
186	11.3.1 Transition systems for dependency parsing	273
187	11.3.2 Scoring functions for transition-based parsers	277
188	11.3.3 Learning to parse	278
189	11.4 Applications	281
190	III Meaning	285
191	12 Logical semantics	287
192	12.1 Meaning and denotation	288
193	12.2 Logical representations of meaning	289
194	12.2.1 Propositional logic	289
195	12.2.2 First-order logic	290
196	12.3 Semantic parsing and the lambda calculus	294
197	12.3.1 The lambda calculus	295
198	12.3.2 Quantification	297
199	12.4 Learning semantic parsers	299
200	12.4.1 Learning from derivations	300
201	12.4.2 Learning from logical forms	302
202	12.4.3 Learning from denotations	303
203	13 Predicate-argument semantics	309
204	13.1 Semantic roles	311
205	13.1.1 VerbNet	312
206	13.1.2 Proto-roles and PropBank	313
207	13.1.3 FrameNet	314
208	13.2 Semantic role labeling	316
209	13.2.1 Semantic role labeling as classification	316
210	13.2.2 Semantic role labeling as constrained optimization	319
211	13.2.3 Neural semantic role labeling	321
212	13.3 Abstract Meaning Representation	322
213	13.3.1 AMR Parsing	325
214	13.4 Applications of Predicate-Argument Semantics	326
215	14 Distributional and distributed semantics	333
216	14.1 The distributional hypothesis	333
217	14.2 Design decisions for word representations	335
218	14.2.1 Representation	335
219	14.2.2 Context	336
220	14.2.3 Estimation	337

221	14.3 Latent semantic analysis	338
222	14.4 Brown clusters	339
223	14.5 Neural word embeddings	343
224	14.5.1 Continuous bag-of-words (CBOW)	343
225	14.5.2 Skipgrams	344
226	14.5.3 Computational complexity	344
227	14.5.4 Word embeddings as matrix factorization	346
228	14.6 Evaluating word embeddings	347
229	14.6.1 Intrinsic evaluations	347
230	14.6.2 Extrinsic evaluations	348
231	14.7 Distributed representations beyond distributional statistics	349
232	14.7.1 Word-internal structure	350
233	14.7.2 Lexical semantic resources	352
234	14.8 Distributed representations of multiword units	352
235	14.8.1 Purely distributional methods	352
236	14.8.2 Distributional-compositional hybrids	353
237	14.8.3 Supervised compositional methods	354
238	14.8.4 Hybrid distributed-symbolic representations	355
239	15 Reference Resolution	359
240	15.1 Forms of referring expressions	360
241	15.1.1 Pronouns	360
242	15.1.2 Proper Nouns	365
243	15.1.3 Nominals	366
244	15.2 Algorithms for coreference resolution	366
245	15.2.1 Mention-pair models	367
246	15.2.2 Mention-ranking models	368
247	15.2.3 Transitive closure in mention-based models	369
248	15.2.4 Entity-based models	370
249	15.3 Representations for coreference resolution	375
250	15.3.1 Features	376
251	15.3.2 Distributed representations of mentions and entities	378
252	15.4 Evaluating coreference resolution	381
253	16 Discourse	385
254	16.1 Segments	385
255	16.1.1 Topic segmentation	386
256	16.1.2 Functional segmentation	387
257	16.2 Entities and reference	387
258	16.2.1 Centering theory	388
259	16.2.2 The entity grid	389

260	16.2.3 *Formal semantics beyond the sentence level	390
261	16.3 Relations	390
262	16.3.1 Shallow discourse relations	391
263	16.3.2 Hierarchical discourse relations	394
264	16.3.3 Argumentation	398
265	16.3.4 Applications of discourse relations	399
266	IV Applications	405
267	17 Information extraction	407
268	17.1 Entities	409
269	17.1.1 Entity linking by learning to rank	410
270	17.1.2 Collective entity linking	412
271	17.1.3 *Pairwise ranking loss functions	413
272	17.2 Relations	415
273	17.2.1 Pattern-based relation extraction	416
274	17.2.2 Relation extraction as a classification task	417
275	17.2.3 Knowledge base population	420
276	17.2.4 Open information extraction	424
277	17.3 Events	425
278	17.4 Hedges, denials, and hypotheticals	426
279	17.5 Question answering and machine reading	428
280	17.5.1 Formal semantics	428
281	17.5.2 Machine reading	429
282	18 Machine translation	435
283	18.1 Machine translation as a task	435
284	18.1.1 Evaluating translations	437
285	18.1.2 Data	439
286	18.2 Statistical machine translation	440
287	18.2.1 Statistical translation modeling	441
288	18.2.2 Estimation	443
289	18.2.3 Phrase-based translation	444
290	18.2.4 *Syntax-based translation	445
291	18.3 Neural machine translation	446
292	18.3.1 Neural attention	448
293	18.3.2 *Neural machine translation without recurrence	450
294	18.3.3 Out-of-vocabulary words	451
295	18.4 Decoding	453
296	18.5 Training towards the evaluation metric	455

297	19 Text generation	459
298	19.1 Data-to-text generation	459
299	19.1.1 Latent data-to-text alignment	461
300	19.1.2 Neural data-to-text generation	462
301	19.2 Text-to-text generation	466
302	19.2.1 Neural abstractive summarization	466
303	19.2.2 Sentence fusion for multi-document summarization	468
304	19.3 Dialogue	469
305	19.3.1 Finite-state and agenda-based dialogue systems	469
306	19.3.2 Markov decision processes	470
307	19.3.3 Neural chatbots	472
308	A Probability	475
309	A.1 Probabilities of event combinations	475
310	A.1.1 Probabilities of disjoint events	476
311	A.1.2 Law of total probability	477
312	A.2 Conditional probability and Bayes' rule	477
313	A.3 Independence	479
314	A.4 Random variables	480
315	A.5 Expectations	481
316	A.6 Modeling and estimation	482
317	B Numerical optimization	485
318	B.1 Gradient descent	486
319	B.2 Constrained optimization	486
320	B.3 Example: Passive-aggressive online learning	487
321	Bibliography	489

322 Notation

323 As a general rule, words, word counts, and other types of observations are indicated with
324 Roman letters (a, b, c); parameters are indicated with Greek letters (α, β, θ). Vectors are
325 indicated with bold script for both random variables \mathbf{x} and parameters $\boldsymbol{\theta}$. Other useful
326 notations are indicated in the table below.

Basics

$\exp x$	the base-2 exponent, 2^x
$\log x$	the base-2 logarithm, $\log_2 x$
$\{x_n\}_{n=1}^N$	the set $\{x_1, x_2, \dots, x_N\}$
x_i^j	x_i raised to the power j
$x_i^{(j)}$	indexing by both i and j

Linear algebra

$\mathbf{x}^{(i)}$	a column vector of feature counts for instance i , often word counts
$\mathbf{x}_{j:k}$	elements j through k (inclusive) of a vector \mathbf{x}
$[\mathbf{x}; \mathbf{y}]$	vertical concatenation of two column vectors
$[\mathbf{x}, \mathbf{y}]$	horizontal concatenation of two column vectors
\mathbf{e}_n	a “one-hot” vector with a value of 1 at position n , and zero everywhere else
$\boldsymbol{\theta}^\top$	the transpose of a column vector $\boldsymbol{\theta}$
$\boldsymbol{\theta} \cdot \mathbf{x}^{(i)}$	the dot product $\sum_{j=1}^N \theta_j \times x_j^{(i)}$
\mathbf{X}	a matrix
$x_{i,j}$	row i , column j of matrix \mathbf{X}
$\text{Diag}(\mathbf{x})$	a matrix with \mathbf{x} on the diagonal, e.g., $\begin{pmatrix} x_1 & 0 & 0 \\ 0 & x_2 & 0 \\ 0 & 0 & x_3 \end{pmatrix}$
\mathbf{X}^{-1}	the inverse of matrix \mathbf{X}

Text datasets

w_m	word token at position m
N	number of training instances
M	length of a sequence (of words or tags)
V	number of words in vocabulary
$y^{(i)}$	the true label for instance i
\hat{y}	a predicted label
\mathcal{Y}	the set of all possible labels
K	number of possible labels $K = \mathcal{Y} $
\square	the start token
\blacksquare	the stop token
$\mathbf{y}^{(i)}$	a structured label for instance i , such as a tag sequence
$\mathcal{Y}(\mathbf{w})$	the set of possible labelings for the word sequence \mathbf{w}
\diamond	the start tag
\blacklozenge	the stop tag

Probabilities

$\Pr(A)$	probability of event A
$\Pr(A B)$	probability of event A , conditioned on event B
$p_B(b)$	the marginal probability of random variable B taking value b ; written $p(b)$ when the choice of random variable is clear from context
$p_{B A}(b a)$	the probability of random variable B taking value b , conditioned on A taking value a ; written $p(b a)$ when clear from context
$A \sim p$	the random variable A is distributed according to distribution p . For example, $X \sim \mathcal{N}(0, 1)$ states that the random variable X is drawn from a normal distribution with zero mean and unit variance.
$A B \sim p$	conditioned on the random variable B , A is distributed according to p . ¹

Machine learning

$\Psi(\mathbf{x}^{(i)}, y)$	the score for assigning label y to instance i
$\mathbf{f}(\mathbf{x}^{(i)}, y)$	the feature vector for instance i with label y
θ	a (column) vector of weights
$\ell^{(i)}$	loss on an individual instance i
L	objective function for an entire dataset
\mathcal{L}	log-likelihood of a dataset
λ	the amount of regularization

327 **Chapter 1**

328 **Introduction**

329 Natural language processing is the set of methods for making human language accessible
330 to computers. In the past decade, natural language processing has become embedded
331 in our daily lives: automatic machine translation is ubiquitous on the web and in social
332 media; text classification keeps emails from collapsing under a deluge of spam; search
333 engines have moved beyond string matching and network analysis to a high degree of
334 linguistic sophistication; dialog systems provide an increasingly common and effective
335 way to get and share information.

336 These diverse applications are based on a common set of ideas, drawing on algo-
337 rithms, linguistics, logic, statistics, and more. The goal of this text is to provide a survey
338 of these foundations. The technical fun starts in the next chapter; the rest of this current
339 chapter situates natural language processing with respect to other intellectual disciplines,
340 identifies some high-level themes in contemporary natural language processing, and ad-
341 vises the reader on how best to approach the subject.

342 **1.1 Natural language processing and its neighbors**

343 One of the great pleasures of working in this field is the opportunity to draw on many
344 other intellectual traditions, from formal linguistics to statistical physics. This section
345 briefly situates natural language processing with respect to some of its closest neighbors.

346 **Computational Linguistics** Most of the meetings and journals that host natural lan-
347 guage processing research bear the name “computational linguistics”, and the terms may
348 be thought of as essentially synonymous. But while there is substantial overlap, there is
349 an important difference in focus. In linguistics, language is the object of study. Compu-
350 tational methods may be brought to bear, just as in scientific disciplines like computational
351 biology and computational astronomy, but they play only a supporting role. In contrast,

352 natural language processing is focused on the design and analysis of computational al-
 353 gorithms and representations for processing natural human language. The goal of natu-
 354 ral language processing is to provide new computational capabilities around human lan-
 355 guage: for example, extracting information from texts, translating between languages, an-
 356 swering questions, holding a conversation, taking instructions, and so on. Fundamental
 357 linguistic insights may be crucial for accomplishing these tasks, but success is ultimately
 358 measured by whether and how well the job gets done.

359 **Machine Learning** Contemporary approaches to natural language processing rely heav-
 360 ily on machine learning, which makes it possible to build complex computer programs
 361 from examples. Machine learning provides an array of general techniques for tasks like
 362 converting a sequence of discrete tokens in one vocabulary to a sequence of discrete to-
 363 kens in another vocabulary — a generalization of what normal people might call “transla-
 364 tion.” Much of today’s natural language processing research can be thought of as applied
 365 machine learning. However, natural language processing has characteristics that distin-
 366 guish it from many of machine learning’s other application domains.

- 367 • Unlike images or audio, text data is fundamentally discrete, with meaning created
 368 by combinatorial arrangements of symbolic units. This is particularly consequential
 369 for applications in which text is the output, such as translation and summarization,
 370 because it is not possible to gradually approach an optimal solution.
- 371 • Although the set of words is discrete, new words are always being created. Further-
 372 more, the distribution over words (and other linguistic elements) resembles that of a
 373 **power law** (Zipf, 1949): there will be a few words that are very frequent, and a long
 374 tail of words that are rare. A consequence is that natural language processing algo-
 375 rithms must be especially robust to observations that do not occur in the training
 376 data.
- 377 • Language is **recursive**: units such as words can combine to create phrases, which
 378 can combine by the very same principles to create larger phrases. For example, a
 379 **noun phrase** can be created by combining a smaller noun phrase with a **preposi-**
 380 **tional phrase**, as in *the whiteness of the whale*. The prepositional phrase is created by
 381 combining a preposition (in this case, *of*) with another noun phrase (*the whale*). In
 382 this way, it is possible to create arbitrarily long phrases, such as,

383 (1.1) ...huge globular pieces of the whale of the bigness of a human head.¹

384 The meaning of such a phrase must be analyzed in accord with the underlying hier-
 385 archical structure. In this case, *huge globular pieces of the whale* acts as a single noun
 386 phrase, which is conjoined with the prepositional phrase of *the bigness of a human*

¹Throughout the text, this notation will be used to introduce linguistic examples.

387 *head*. The interpretation would be different if instead, *huge globular pieces* were con-
388 joined with the prepositional phrase *of the whale of the bigness of a human head* —
389 implying a disappointingly small whale. Even though text appears as a sequence,
390 machine learning methods must account for its implicit recursive structure.

391 **Artificial Intelligence** The goal of artificial intelligence is to build software and robots
392 with the same range of abilities as humans (Russell and Norvig, 2009). Natural language
393 processing is relevant to this goal in several ways. The capacity for language is one of the
394 central features of human intelligence, and no artificial intelligence program could be said
395 to be complete without the ability to communicate in words.²

396 Much of artificial intelligence research is dedicated to the development of systems
397 that can reason from premises to a conclusion, but such algorithms are only as good as
398 what they know (Dreyfus, 1992). Natural language processing is a potential solution to
399 the “knowledge bottleneck”, by acquiring knowledge from natural language texts, and
400 perhaps also from conversations; This idea goes all the way back to Turing’s 1949 pa-
401 per *Computing Machinery and Intelligence*, which proposed the **Turing test** and helped to
402 launch the field of artificial intelligence (Turing, 2009).

403 Conversely, reasoning is sometimes essential for basic tasks of language processing,
404 such as determining who a pronoun refers to. **Winograd schemas** are examples in which
405 a single word changes the likely referent of a pronoun, in a way that seems to require
406 knowledge and reasoning to decode (Levesque et al., 2011). For example,

407 (1.2) The trophy doesn’t fit into the brown suitcase because **it** is too [small/large].
408 When the final word is *small*, then the pronoun *it* refers to the suitcase; when the final
409 word is *large*, then *it* refers to the trophy. Solving this example requires spatial reasoning;
410 other schemas require reasoning about actions and their effects, emotions and intentions,
411 and social conventions.

412 The Winograd schemas demonstrate that natural language understanding cannot be
413 achieved in isolation from knowledge and reasoning. Yet the history of artificial intelli-
414 gence has been one of increasing specialization: with the growing volume of research in
415 subdisciplines such as natural language processing, machine learning, and computer vi-

²This view seems to be shared by some, but not all, prominent researchers in artificial intelligence. Michael Jordan, a specialist in machine learning, has said that if he had a billion dollars to spend on any large research project, he would spend it on natural language processing (https://www.reddit.com/r/MachineLearning/comments/2fxi6v/ama_michael_i_jordan/). On the other hand, in a public discussion about the future of artificial intelligence in February 2018, computer vision researcher Yann LeCun argued that language was perhaps the “50th most important” thing to work on, and that it would be a great achievement if AI could attain the capabilities of an orangutan, which presumably do not include language (<http://www.abigailsee.com/2018/02/21/deep-learning-structure-and-innate-priors.html>).

416 sion, it is difficult for anyone to maintain expertise across the entire field. Still, recent work
417 has demonstrated interesting connections between natural language processing and other
418 areas of AI, including computer vision (e.g., Antol et al., 2015) and game playing (e.g.,
419 Branavan et al., 2009). The dominance of machine learning throughout artificial intel-
420 ligence has led to a broad consensus on representations such as graphical models and
421 knowledge graphs, and on algorithms such as backpropagation and combinatorial opti-
422 mization. Many of the algorithms and representations covered in this text are part of this
423 consensus.

424 **Computer Science** The discrete and recursive nature of natural language invites the ap-
425 plication of theoretical ideas from computer science. Linguists such as Chomsky and
426 Montague have shown how formal language theory can help to explain the syntax and
427 semantics of natural language. Theoretical models such as finite-state and pushdown au-
428 tomata are the basis for many practical natural language processing systems. Algorithms
429 for searching the combinatorial space of analyses of natural language utterances can be
430 analyzed in terms of their computational complexity, and theoretically motivated approx-
431 imations can sometimes be applied.

432 The study of computer systems is also relevant to natural language processing. Large
433 datasets of unlabeled text are a natural application for parallelization techniques like
434 MapReduce (Dean and Ghemawat, 2008; Lin and Dyer, 2010); high-volume data sources
435 such as social media are a natural application for approximate streaming and sketching
436 techniques (Goyal et al., 2009). When deep neural networks are implemented in pro-
437 duction systems, it is possible to eke out speed gains using techniques such as reduced-
438 precision arithmetic (Wu et al., 2016). Many classical natural language processing algo-
439 rithms are not naturally suited to graphics processing unit (GPU) parallelization, suggest-
440 ing directions for further research at the intersection of natural language processing and
441 computing hardware (Yi et al., 2011).

442 **Speech Processing** Natural language is often communicated in spoken form, and speech
443 recognition is the task of converting an audio signal to text. From one perspective, this is
444 a signal processing problem, which might be viewed as a preprocessing step before nat-
445 ural language processing can be applied. However, context plays a critical role in speech
446 recognition by human listeners: knowledge of the surrounding words influences percep-
447 tion and helps to correct for noise (Miller et al., 1951). For this reason, speech recognition
448 is often integrated with text analysis, particularly with statistical **language model**, which
449 quantify the probability of a sequence of text (see chapter 6). Beyond speech recognition,
450 the broader field of speech processing includes the study of speech-based dialogue sys-
451 tems, which are briefly discussed in chapter 19. Historically, speech processing has often
452 been pursued in electrical engineering departments, while natural language processing

453 has been the purview of computer scientists. For this reason, the extent of interaction
454 between these two disciplines is less than it might otherwise be.

455 **Others** Natural language processing plays a significant role in emerging interdisciplinary
456 fields like **computational social science** and the **digital humanities**. Text classification
457 (chapter 4), clustering (chapter 5), and information extraction (chapter 17) are particularly
458 useful tools; another is probabilistic **topic models** (Blei, 2012), which are not covered in
459 this text. **Information retrieval** (Manning et al., 2008) makes use of similar tools, and
460 conversely, techniques such as latent semantic analysis (§ 14.3) have roots in information
461 retrieval. **Text mining** is sometimes used to refer to the application of data mining tech-
462 niques, especially classification and clustering, to text. While there is no clear distinction
463 between text mining and natural language processing (nor between data mining and ma-
464 chine learning), text mining is typically less concerned with linguistic structure, and more
465 interested in fast, scalable algorithms.

466 1.2 Three themes in natural language processing

467 Natural language processing covers a diverse range of tasks, methods, and linguistic phe-
468 nomena. But despite the apparent incommensurability between, say, the summarization
469 of scientific articles (§ 16.3.4.1) and the identification of suffix patterns in Spanish verbs
470 (§ 9.1.4.3), some general themes emerge. Each of these themes can be expressed as an
471 opposition between two extreme viewpoints on how to process natural language, and in
472 each case, existing approaches can be placed on a continuum between these two extremes.

473 1.2.1 Learning and knowledge

474 A recurring topic of debate is the relative importance of machine learning and linguistic
475 knowledge. On one extreme, advocates of “natural language processing from scratch” (Col-
476 lobert et al., 2011) propose to use machine learning to train end-to-end systems that trans-
477 mute raw text into any desired output structure: e.g., a summary, database, or transla-
478 tion. On the other extreme, the core work of natural language processing is sometimes
479 taken to be transforming text into a stack of general-purpose linguistic structures: from
480 subword units called **morphemes**, to word-level **parts-of-speech**, to tree-structured repre-
481 sentations of grammar, and beyond, to logic-based representations of meaning. In theory,
482 these general-purpose structures should then be able to support any desired application.

483 The end-to-end learning approach has been buoyed by recent results in computer vi-
484 sion and speech recognition, in which advances in machine learning have swept away
485 expert-engineered representations based on the fundamentals of optics and phonology (Krizhevsky
486 et al., 2012; Graves and Jaitly, 2014). But while some amount of machine learning is an el-
487 ement of nearly every contemporary approach to natural language processing, linguistic

488 representations such as syntax trees have not yet gone the way of the visual edge detector
 489 or the auditory triphone. Linguists have argued for the existence of a “language faculty”
 490 in all human beings, which encodes a set of abstractions specially designed to facilitate
 491 the understanding and production of language. The argument for the existence of such
 492 a language faculty is based on the observation that children learn language faster and
 493 from fewer examples than would be reasonably possible, if language was learned from
 494 experience alone.³ Regardless of the cognitive validity of these arguments, it seems that
 495 linguistic structures are particularly important in scenarios where training data is limited.

496 Moving away from the extreme ends of the continuum, there are a number of ways in
 497 which knowledge and learning can be combined in natural language processing. Many
 498 supervised learning systems make use of carefully engineered **features**, which transform
 499 the data into a representation that can facilitate learning. For example, in a task like doc-
 500 ument classification, it may be useful to identify each word’s **stem**, so that a learning
 501 system can more easily generalize across related terms such as *whale*, *whales*, *whalers*, and
 502 *whaling*. This is particularly important in the many languages that exceed English in the
 503 complexity of the system of affixes that can attach to words. Such features could be ob-
 504 tained from a hand-crafted resource, like a dictionary that maps each word to a single
 505 root form. Alternatively, features can be obtained from the output of a general-purpose
 506 language processing system, such as a parser or part-of-speech tagger, which may itself
 507 be built on supervised machine learning.

508 Another synthesis of learning and knowledge is in model structure: building machine
 509 learning models whose architectures are inspired by linguistic theories. For example, the
 510 organization of sentences is often described as **compositional**, with meaning of larger
 511 units gradually constructed from the meaning of their smaller constituents. This idea
 512 can be built into the architecture of a deep neural network, which is then trained using
 513 contemporary deep learning techniques (Dyer et al., 2016).

514 The debate about the relative importance of machine learning and linguistic knowl-
 515 edge sometimes becomes heated. No machine learning specialist likes to be told that their
 516 engineering methodology is unscientific alchemy;⁴ nor does a linguist want to hear that
 517 the search for general linguistic principles and structures has been made irrelevant by big
 518 data. Yet there is clearly room for both types of research: we need to know how far we
 519 can go with end-to-end learning alone, while at the same time, we continue the search for
 520 linguistic representations that generalize across applications, scenarios, and languages.
 521 For more on the history of this debate, see Church (2011); for an optimistic view of the
 522 potential symbiosis between computational linguistics and deep learning, see Manning

³The *Language Instinct* (Pinker, 2003) articulates these arguments in an engaging and popular style. For arguments against the innateness of language, see Elman et al. (1998).

⁴Ali Rahimi argued that much of deep learning research was similar to “alchemy” in a presentation at the 2017 conference on Neural Information Processing Systems. He was advocating for more learning theory, not more linguistics.

523 (2015).

524 **1.2.2 Search and learning**

525 Many natural language processing problems can be written mathematically in the form
 526 of optimization,⁵

$$\hat{\mathbf{y}} = \operatorname{argmax}_{\mathbf{y} \in \mathcal{Y}(\mathbf{x})} \Psi(\mathbf{x}, \mathbf{y}; \boldsymbol{\theta}), \quad [1.1]$$

527 where,

- 528 • \mathbf{x} is the input, which is an element of a set \mathcal{X} ;
- 529 • \mathbf{y} is the output, which is an element of a set $\mathcal{Y}(\mathbf{x})$;
- 530 • Ψ is a scoring function (also called the **model**), which maps from the set $\mathcal{X} \times \mathcal{Y}$ to
 531 the real numbers;
- 532 • $\boldsymbol{\theta}$ is a vector of parameters for Ψ ;
- 533 • $\hat{\mathbf{y}}$ is the predicted output, which is chosen to maximize the scoring function.

534 This basic structure can be used across a huge range of problems. For example, the
 535 input \mathbf{x} might be a social media post, and the output \mathbf{y} might be a labeling of the emotional
 536 sentiment expressed by the author (chapter 4); or \mathbf{x} could be a sentence in French, and the
 537 output \mathbf{y} could be a sentence in Tamil (chapter 18); or \mathbf{x} might be a sentence in English,
 538 and \mathbf{y} might be a representation of the syntactic structure of the sentence (chapter 10); or
 539 \mathbf{x} might be a news article and \mathbf{y} might be a structured record of the events that the article
 540 describes (chapter 17).

541 By adopting this formulation, we make an implicit decision that language processing
 542 algorithms will have two distinct modules:

543 **Search.** The search module is responsible for computing the argmax of the function Ψ . In
 544 other words, it finds the output $\hat{\mathbf{y}}$ that gets the best score with respect to the input
 545 \mathbf{x} . This is easy when the search space $\mathcal{Y}(\mathbf{x})$ is small enough to enumerate, or when
 546 the scoring function Ψ has a convenient decomposition into parts. In many cases,
 547 we will want to work with scoring functions that do not have these properties, moti-
 548 vating the use of more sophisticated search algorithms. Because the outputs are
 549 usually discrete in language processing problems, search often relies on the machin-
 550 ery of **combinatorial optimization**.

⁵Throughout this text, equations will be numbered by square brackets, and linguistic examples will be numbered by parentheses.

551 **Learning.** The learning module is responsible for finding the parameters θ . This is typ-
 552 ically (but not always) done by processing a large dataset of labeled examples,
 553 $\{(\mathbf{x}^{(i)}, \mathbf{y}^{(i)})\}_{i=1}^N$. Like search, learning is also approached through the framework
 554 of optimization, as we will see in chapter 2. Because the parameters are usually
 555 continuous, learning algorithms generally rely on **numerical optimization**, search-
 556 ing over vectors of real numbers for parameters that optimize some function of the
 557 model and the labeled data. Some basic principles of numerical optimization are
 558 reviewed in Appendix B.

559 The division of natural language processing into separate modules for search and
 560 learning makes it possible to reuse generic algorithms across a range of different tasks
 561 and models. This means that the work of natural language processing can be focused on
 562 the design of the model Ψ , while reaping the benefits of decades of progress in search,
 563 optimization, and learning. Much of this textbook will focus on specific classes of scoring
 564 functions, and on the algorithms that make it possible to search and learn efficiently with
 565 them.

566 When a model is capable of making subtle linguistic distinctions, it is said to be *expres-*
 567 *sive*. Expressiveness is often traded off against the efficiency of search and learning. For
 568 example, a word-to-word translation model makes search and learning easy, but it is not
 569 expressive enough to distinguish good translations from bad ones. Unfortunately many
 570 of the most important problems in natural language processing seem to require expres-
 571 sive models, in which the complexity of search grows exponentially with the size of the
 572 input. In these models, exact search is usually impossible. Intractability threatens the neat
 573 modular decomposition between search and learning: if search requires a set of heuristic
 574 approximations, then it may be advantageous to learn a model that performs well under
 575 these specific heuristics. This has motivated some researchers to take a more integrated
 576 approach to search and learning, as briefly mentioned in chapters 11 and 15.

577 1.2.3 Relational, compositional, and distributional perspectives

578 Any element of language — a word, a phrase, a sentence, or even a sound — can be
 579 described from at least three perspectives. Consider the word *journalist*. A *journalist* is a
 580 subcategory of a *profession*, and an *anchorwoman* is a subcategory of *journalist*; furthermore,
 581 a *journalist* performs *journalism*, which is often, but not always, a subcategory of *writing*.
 582 This relational perspective on meaning is the basis for semantic **ontologies** such as **Word-**
 583 **Net** (Fellbaum, 2010), which enumerate the relations that hold between words and other
 584 elementary semantic units. The power of the relational perspective is illustrated by the
 585 following example:

586 (1.3) Umashanthi interviewed Ana. She works for the college newspaper.

587 Who works for the college newspaper? The word *journalist*, while not stated in the ex-
588 ample, implicitly links the *interview* to the *newspaper*, making *Umashanthi* the most likely
589 referent for the pronoun. (A general discussion of how to resolve pronouns is found in
590 chapter 15.)

591 Yet despite the inferential power of the relational perspective, it is not easy to formalize
592 computationally. Exactly which elements are to be related? Are *journalists* and *reporters*
593 distinct, or should we group them into a single unit? Is the kind of *interview* performed by
594 a journalist the same as the kind that one undergoes when applying for a job? Ontology
595 designers face many such thorny questions, and the project of ontology design hearkens
596 back to Borges' (1993) *Celestial Emporium of Benevolent Knowledge*, which divides animals
597 into:

598 (a) belonging to the emperor; (b) embalmed; (c) tame; (d) suckling pigs; (e)
599 sirens; (f) fabulous; (g) stray dogs; (h) included in the present classification;
600 (i) frenzied; (j) innumerable; (k) drawn with a very fine camelhair brush; (l) et
601 cetera; (m) having just broken the water pitcher; (n) that from a long way off
602 resemble flies.

603 Difficulties in ontology construction have led some linguists to argue that there is no task-
604 independent way to partition up word meanings (Kilgarriff, 1997).

605 Some problems are easier. Each member in a group of *journalists* is a *journalist*: the -s
606 suffix distinguishes the plural meaning from the singular in most of the nouns in English.
607 Similarly, a *journalist* can be thought of, perhaps colloquially, as someone who produces or
608 works on a *journal*. (Taking this approach even further, the word *journal* derives from the
609 French *jour+nal*, or *day+ly* = *daily*.) In this way, the meaning of a word is constructed from
610 the constituent parts — the principle of **compositionality**. This principle can be applied
611 to larger units: phrases, sentences, and beyond. Indeed, one of the great strengths of the
612 compositional view of meaning is that it provides a roadmap for understanding entire
613 texts and dialogues through a single analytic lens, grounding out in the smallest parts of
614 individual words.

615 But alongside *journalists* and *anti-parliamentarians*, there are many words that seem to
616 be linguistic atoms: think, for example, of *whale*, *blubber*, and *Nantucket*. Furthermore,
617 idiomatic phrases like *kick the bucket* and *shoot the breeze* have meanings that are quite
618 different from the sum of their parts (Sag et al., 2002). Composition is of little help for such
619 words and expressions, but their meanings can be ascertained — or at least approximated
620 — from the contexts in which they appear. Take, for example, *blubber*, which appears in
621 such contexts as:

- 622 (1.4) The blubber served them as fuel.
623 (1.5) ... extracting it from the blubber of the large fish ...

624 (1.6) Amongst oily substances, blubber has been employed as a manure.

625 These contexts form the **distributional properties** of the word *blubber*, and they link it to
 626 words which can appear in similar constructions: *fat*, *pelts*, and *barnacles*. This distribu-
 627 tional perspective makes it possible to learn about meaning from unlabeled data alone;
 628 unlike relational and compositional semantics, no manual annotation or expert knowl-
 629 edge is required. Distributional semantics is thus capable of covering a huge range of
 630 linguistic phenomena. However, it lacks precision: *blubber* is similar to *fat* in one sense, to
 631 *pelts* in another sense, and to *barnacles* in still another. The question of *why* all these words
 632 tend to appear in the same contexts is left unanswered.

633 The relational, compositional, and distributional perspectives all contribute to our un-
 634 derstanding of linguistic meaning, and all three appear to be critical to natural language
 635 processing. Yet they are uneasy collaborators, requiring seemingly incompatible repre-
 636 sentations and algorithmic approaches. This text presents some of the best known and
 637 most successful methods for working with each of these representations, but it is hoped
 638 that future research will reveal new ways to combine them.

639 1.3 Learning to do natural language processing

640 This text began with the notes that I use for teaching Georgia Tech’s undergraduate and
 641 graduate courses on natural language processing, CS 4650 and 7650. There are several
 642 other good resources (e.g., Manning and Schütze, 1999; Jurafsky and Martin, 2009; Smith,
 643 2011; Collins, 2013), but the goal of this text is focus on a core subset of the field, uni-
 644 fied by the concepts of learning and search. A remarkable thing about natural language
 645 processing is that so many problems can be solved by a compact set of methods:

646 **Search.** Viterbi, CKY, minimum spanning tree, shift-reduce, integer linear programming,
 647 beam search.

648 **Learning.** Naïve Bayes, logistic regression, perceptron, expectation-maximization, matrix
 649 factorization, backpropagation, recurrent neural networks.

650 This text explains how these methods work, and how they can be applied to problems
 651 that arise in the computer processing of natural language: document classification, word
 652 sense disambiguation, sequence labeling (part-of-speech tagging and named entity recog-
 653 nition), parsing, coreference resolution, relation extraction, discourse analysis, language
 654 modeling, and machine translation.

655 1.3.1 Background

656 Because natural language processing draws on many different intellectual traditions, al-
 657 most everyone who approaches it feels underprepared in one way or another. Here is a

658 summary of what is expected, and where you can learn more:

659 **Mathematics and machine learning.** The text assumes a background in multivariate cal-
660 culus and linear algebra: vectors, matrices, derivatives, and partial derivatives. You
661 should also be familiar with probability and statistics. A review of basic proba-
662 bility is found in Appendix A, and a minimal review of numerical optimization is
663 found in Appendix B. For linear algebra, the online course and textbook from Strang
664 (2016) are excellent sources of review material. Deisenroth et al. (2018) are currently
665 preparing a textbook on *Mathematics for Machine Learning*, and several chapters can
666 be found online.⁶ For an introduction to probabilistic modeling and estimation, see
667 James et al. (2013); for a more advanced and comprehensive discussion of the same
668 material, the classic reference is Hastie et al. (2009).

669 **Linguistics.** This book assumes no formal training in linguistics, aside from elementary
670 concepts like nouns and verbs, which you have probably encountered in the study
671 of English grammar. Ideas from linguistics are introduced throughout the text as
672 needed, including discussions of morphology and syntax (chapter 9), semantics
673 (chapters 12 and 13), and discourse (chapter 16). Linguistic issues also arise in the
674 application-focused chapters 4, 8, and 18. A short guide to linguistics for students
675 of natural language processing is offered by Bender (2013); you are encouraged to
676 start there, and then pick up a more comprehensive introductory textbook (e.g., Ak-
677 majian et al., 2010; Fromkin et al., 2013).

678 **Computer science.** The book is targeted at computer scientists, who are assumed to have
679 taken introductory courses on the analysis of algorithms and complexity theory. In
680 particular, you should be familiar with asymptotic analysis of the time and memory
681 costs of algorithms, and should have seen dynamic programming. The classic text
682 on algorithms is offered by Cormen et al. (2009); for an introduction to the theory of
683 computation, see Arora and Barak (2009) and Sipser (2012).

684 1.3.2 How to use this book

685 The textbook is organized into four main units:

686 **Learning.** This section builds up a set of machine learning tools that will be used through-
687 out the rest of the textbook. Because the focus is on machine learning, the text
688 representations and linguistic phenomena are mostly simple: “bag-of-words” text
689 classification is treated as a model example. Chapter 4 describes some of the more
690 linguistically interesting applications of word-based text analysis.

⁶<https://mml-book.github.io/>

691 **Sequences and trees.** This section introduces the treatment of language as a structured
 692 phenomena. It describes sequence and tree representations and the algorithms that
 693 they facilitate, as well as the limitations that these representations impose. Chapter
 694 9 introduces finite state automata and briefly overviews a context-free account of
 695 English syntax.

696 **Meaning.** This section takes a broad view of efforts to represent and compute meaning
 697 from text, ranging from formal logic to neural word embeddings. It also includes
 698 two topics that are closely related to semantics: resolution of ambiguous references,
 699 and analysis of multi-sentence discourse structure.

700 **Applications.** The final section offers chapter-length treatments on three of the most prominent
 701 applications of natural language processing: information extraction, machine
 702 translation, and text generation. Each of these applications merits a textbook length
 703 treatment of its own (Koehn, 2009; Grishman, 2012; Reiter and Dale, 2000); the chapters
 704 here explain some of the most well known systems using the formalisms and
 705 methods built up earlier in the book, while introducing methods such as neural attention.
 706

707 Each chapter contains some advanced material, which is marked with an asterisk.
 708 This material can be safely omitted without causing misunderstandings later on. But
 709 even without these advanced sections, the text is too long for a single semester course, so
 710 instructors will have to pick and choose among the chapters.

711 Chapters 2 and 3 provide building blocks that will be used throughout the book, and
 712 chapter 4 describes some critical aspects of the practice of language technology. Lan-
 713 guage models (chapter 6), sequence labeling (chapter 7), and parsing (chapter 10 and 11)
 714 are canonical topics in natural language processing, and distributed word embeddings
 715 (chapter 14) are so ubiquitous that students will complain if you leave them out. Of the
 716 applications, machine translation (chapter 18) is the best choice: it is more cohesive than
 717 information extraction, and more mature than text generation. In my experience, nearly
 718 all students benefit from the review of probability in Appendix A.

- 719 • A course focusing on machine learning should add the chapter on unsupervised
 720 learning (chapter 5). The chapters on predicate-argument semantics (chapter 13),
 721 reference resolution (chapter 15), and text generation (chapter 19) are particularly
 722 influenced by recent machine learning innovations, including deep neural networks
 723 and learning to search.
- 724 • A course with a more linguistic orientation should add the chapters on applica-
 725 tions of sequence labeling (chapter 8), formal language theory (chapter 9), semantics
 726 (chapter 12 and 13), and discourse (chapter 16).

- 727 • For a course with a more applied focus — for example, a course targeting under-
728 graduates — I recommend the chapters on applications of sequence labeling (chap-
729 ter 8), predicate-argument semantics (chapter 13), information extraction (chapter 17),
730 and text generation (chapter 19).

731 **Acknowledgments**

732 Several of my colleagues and students read early drafts of chapters in their areas of exper-
733 tise, including Yoav Artzi, Kevin Duh, Heng Ji, Jessy Li, Brendan O'Connor, Yuval Pinter,
734 Nathan Schneider, Pamela Shapiro, Noah A. Smith, Sandeep Soni, and Luke Zettlemoyer.
735 I would also like to thank the following people for helpful discussions: Kevin Murphy,
736 Shawn Ling Ramirez, William Yang Wang, and Bonnie Webber. Several students, col-
737 leagues, friends, and family found mistakes in early drafts: Parminder Bhatia, Kimberly
738 Caras, Justin Chen, Murtaza Dhuliawala, Barbara Eisenstein, Luiz C. F. Ribeiro, Chris Gu,
739 Joshua Killingsworth, Jonathan May, Taha Merghani, Gus Monod, Raghavendra Murali,
740 Nidish Nair, Brendan O'Connor, Yuval Pinter, Nathan Schneider, Zhewei Sun, Rubin Tsui,
741 Ashwin Cunnappakkam Vinjimir, Clay Washington, Ishan Waykul, and Yuyu Zhang. Clay
742 Washington pilot tested several of the programming exercise. Special thanks to the many
743 students in Georgia Tech's CS 4650 and 7650 who suffered through early versions of the
744 text.

745

Part I

746

Learning

747

Chapter 2

748

Linear text classification

749 We'll start with the problem of **text classification**: given a text document, assign it a dis-
750 crete label $y \in \mathcal{Y}$, where \mathcal{Y} is the set of possible labels. This problem has many appli-
751 cations, from spam filtering to analysis of electronic health records. Text classification is
752 also a building block that is used throughout more complex natural language processing
753 tasks.

754 To perform this task, the first question is how to represent each document. A common
755 approach is to use a vector of word counts, e.g., $\mathbf{x} = [0, 1, 1, 0, 0, 2, 0, 1, 13, 0 \dots]^T$, where
756 x_j is the count of word j . The length of \mathbf{x} is $V \triangleq |\mathcal{V}|$, where \mathcal{V} is the set of possible words
757 in the vocabulary.

758 The object \mathbf{x} is a vector, but colloquially we call it a **bag of words**, because it includes
759 only information about the count of each word, and not the order in which the words
760 appear. We have thrown out grammar, sentence boundaries, paragraphs — everything
761 but the words. Yet the bag of words model is surprisingly effective for text classification.
762 If you see the word *freeee* in an email, is it a spam email? What if you see the word
763 *Bayesian*? For many labeling problems, individual words can be strong predictors.

764 To predict a label from a bag-of-words, we can assign a score to each word in the
765 vocabulary, measuring the compatibility with the label. In the spam filtering case, we
766 might assign a positive score to the word *freeee* for the label SPAM, and a negative score
767 to the word *Bayesian*. These scores are called **weights**, and they are arranged in a column
768 vector θ .

769 Suppose that you want a multiclass classifier, where $K \triangleq |\mathcal{Y}| > 2$. For example, we
770 might want to classify news stories about sports, celebrities, music, and business. The goal
771 is to predict a label \hat{y} , given the bag of words \mathbf{x} , using the weights θ . For each label $y \in \mathcal{Y}$,
772 we compute a score $\Psi(\mathbf{x}, y)$, which is a scalar measure of the compatibility between the
773 bag-of-words \mathbf{x} and the label y . In a linear bag-of-words classifier, this score is the vector

774 inner product between the weights θ and the output of a **feature function** $f(\mathbf{x}, y)$,

$$\Psi(\mathbf{x}, y) = \theta \cdot f(\mathbf{x}, y). \quad [2.1]$$

775 As the notation suggests, f is a function of two arguments, the word counts \mathbf{x} and the
 776 label y , and it returns a vector output. For example, given arguments \mathbf{x} and y , element j
 777 of this feature vector might be,

$$f_j(\mathbf{x}, y) = \begin{cases} x_{freeee}, & \text{if } y = \text{SPAM} \\ 0, & \text{otherwise} \end{cases} \quad [2.2]$$

778 This function returns the count of the word *freeee* if the label is SPAM, and it returns zero
 779 otherwise. The corresponding weight θ_j then scores the compatibility of the word *freeee*
 780 with the label SPAM. A positive score means that this word makes the label more likely.

To formalize this feature function, we define $f(\mathbf{x}, y)$ as a column vector,

$$f(\mathbf{x}, y = 1) = [\mathbf{x}; \underbrace{0; 0; \dots; 0}_{(K-1) \times V}] \quad [2.3]$$

$$f(\mathbf{x}, y = 2) = [\underbrace{0; 0; \dots; 0}_V; \mathbf{x}; \underbrace{0; 0; \dots; 0}_{(K-2) \times V}] \quad [2.4]$$

$$f(\mathbf{x}, y = K) = [\underbrace{0; 0; \dots; 0}_{(K-1) \times V}; \mathbf{x}], \quad [2.5]$$

781 where $\underbrace{[0; 0; \dots; 0]}_{(K-1) \times V}$ is a column vector of $(K - 1) \times V$ zeros, and the semicolon indicates
 782 vertical concatenation. This arrangement is shown in Figure 2.1; the notation may seem
 783 awkward at first, but it generalizes to an impressive range of learning settings.

Given a vector of weights, $\theta \in \mathbb{R}^{V \times K}$, we can now compute the score $\Psi(\mathbf{x}, y)$. This
 inner product gives a scalar measure of the compatibility of the observation \mathbf{x} with label
 y .¹ For any document \mathbf{x} , we predict the label \hat{y} ,

$$\hat{y} = \operatorname{argmax}_{y \in \mathcal{Y}} \Psi(\mathbf{x}, y) \quad [2.6]$$

$$\Psi(\mathbf{x}, y) = \theta \cdot f(\mathbf{x}, y). \quad [2.7]$$

784 This inner product notation gives a clean separation between the *data* (\mathbf{x} and y) and the
 785 *parameters* (θ). This notation also generalizes nicely to **structured prediction**, in which

¹Only $V \times (K - 1)$ features and weights are necessary. By stipulating that $\Psi(\mathbf{x}, y = K) = 0$ regardless of \mathbf{x} , it is possible to implement any classification rule that can be achieved with $V \times K$ features and weights. This is the approach taken in binary classification rules like $y = \text{Sign}(\beta \cdot \mathbf{x} + a)$, where β is a vector of weights, a is an offset, and the label set is $\mathcal{Y} = \{-1, 1\}$. However, for multiclass classification, it is more concise to write $\theta \cdot f(\mathbf{x}, y)$ for all $y \in \mathcal{Y}$.

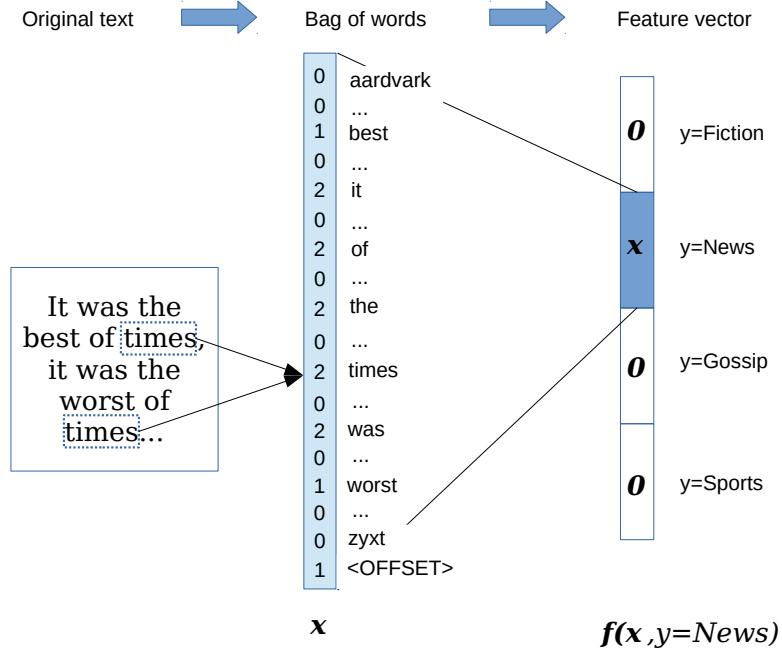


Figure 2.1: The bag-of-words and feature vector representations, for a hypothetical text classification task.

786 the space of labels \mathcal{Y} is very large, and we want to model shared substructures between
 787 labels.

788 It is common to add an **offset feature** at the end of the vector of word counts \mathbf{x} , which
 789 is always 1. We then have to also add an extra zero to each of the zero vectors, to make the
 790 vector lengths match. This gives the entire feature vector $\mathbf{f}(\mathbf{x}, y)$ a length of $(V + 1) \times K$.
 791 The weight associated with this offset feature can be thought of as a bias for or against
 792 each label. For example, if we expect most documents to be spam, then the weight for
 793 the offset feature for $y = \text{SPAM}$ should be larger than the weight for the offset feature for
 794 $y = \text{HAM}$.

Returning to the weights θ , where do they come from? One possibility is to set them by hand. If we wanted to distinguish, say, English from Spanish, we can use English and Spanish dictionaries, and set the weight to one for each word that appears in the

associated dictionary. For example,²

$$\begin{array}{ll} \theta_{(E,bicycle)} = 1 & \theta_{(S,bicycle)} = 0 \\ \theta_{(E,bicicleta)} = 0 & \theta_{(S,bicicleta)} = 1 \\ \theta_{(E,con)} = 1 & \theta_{(S,con)} = 1 \\ \theta_{(E,ordinateur)} = 0 & \theta_{(S,ordinateur)} = 0. \end{array}$$

Similarly, if we want to distinguish positive and negative sentiment, we could use positive and negative **sentiment lexicons** (see § 4.1.2), which are defined by social psychologists (Tausczik and Pennebaker, 2010).

But it is usually not easy to set classification weights by hand, due to the large number of words and the difficulty of selecting exact numerical weights. Instead, we will learn the weights from data. Email users manually label messages as SPAM; newspapers label their own articles as BUSINESS or STYLE. Using such **instance labels**, we can automatically acquire weights using **supervised machine learning**. This chapter will discuss several machine learning approaches for classification. The first is based on probability. For a review of probability, consult Appendix A.

2.1 Naïve Bayes

The **joint probability** of a bag of words \mathbf{x} and its true label y is written $p(\mathbf{x}, y)$. Suppose we have a dataset of N labeled instances, $\{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^N$, which we assume are **independent and identically distributed (IID)** (see § A.3). Then the joint probability of the entire dataset, written $p(\mathbf{x}^{(1:N)}, y^{(1:N)})$, is equal to $\prod_{i=1}^N p_{X,Y}(\mathbf{x}^{(i)}, y^{(i)})$.³

What does this have to do with classification? One approach to classification is to set the weights $\boldsymbol{\theta}$ so as to maximize the joint probability of a **training set** of labeled documents. This is known as **maximum likelihood estimation**:

$$\hat{\boldsymbol{\theta}} = \operatorname{argmax}_{\boldsymbol{\theta}} p(\mathbf{x}^{(1:N)}, y^{(1:N)}; \boldsymbol{\theta}) \quad [2.8]$$

$$= \operatorname{argmax}_{\boldsymbol{\theta}} \prod_{i=1}^N p(\mathbf{x}^{(i)}, y^{(i)}; \boldsymbol{\theta}) \quad [2.9]$$

$$= \operatorname{argmax}_{\boldsymbol{\theta}} \sum_{i=1}^N \log p(\mathbf{x}^{(i)}, y^{(i)}; \boldsymbol{\theta}). \quad [2.10]$$

²In this notation, each tuple (language, word) indexes an element in $\boldsymbol{\theta}$, which remains a vector.

³The notation $p_{X,Y}(\mathbf{x}^{(i)}, y^{(i)})$ indicates the joint probability that random variables X and Y take the specific values $\mathbf{x}^{(i)}$ and $y^{(i)}$ respectively. The subscript will often be omitted when it is clear from context. For a review of random variables, see Appendix A.

Algorithm 1 Generative process for the Naïve Bayes classifier

for Document $i \in \{1, 2, \dots, N\}$ **do**:

Draw the label $y^{(i)} \sim \text{Categorical}(\boldsymbol{\mu})$;

Draw the word counts $\mathbf{x}^{(i)} | y^{(i)} \sim \text{Multinomial}(\boldsymbol{\phi}_{y^{(i)}})$.

810 The notation $p(\mathbf{x}^{(i)}, y^{(i)}; \boldsymbol{\theta})$ indicates that $\boldsymbol{\theta}$ is a *parameter* of the probability function. The
 811 product of probabilities can be replaced by a sum of log-probabilities because the log func-
 812 tion is monotonically increasing over positive arguments, and so the same $\boldsymbol{\theta}$ will maxi-
 813 mize both the probability and its logarithm. Working with logarithms is desirable because
 814 of numerical stability: on a large dataset, multiplying many probabilities can **underflow**
 815 to zero.⁴

816 The probability $p(\mathbf{x}^{(i)}, y^{(i)}; \boldsymbol{\theta})$ is defined through a **generative model** — an idealized
 817 random process that has generated the observed data.⁵ Algorithm 1 describes the gener-
 818 ative model describes the **Naïve Bayes** classifier, with parameters $\boldsymbol{\theta} = \{\boldsymbol{\mu}, \boldsymbol{\phi}\}$.

- 819 • The first line of this generative model encodes the assumption that the instances are
 820 mutually independent: neither the label nor the text of document i affects the label
 821 or text of document j .⁶ Furthermore, the instances are identically distributed: the
 822 distributions over the label $y^{(i)}$ and the text $\mathbf{x}^{(i)}$ (conditioned on $y^{(i)}$) are the same
 823 for all instances i .
- 824 • The second line of the generative model states that the random variable $y^{(i)}$ is drawn
 825 from a categorical distribution with parameter $\boldsymbol{\mu}$. Categorical distributions are like
 826 weighted dice: the vector $\boldsymbol{\mu} = [\mu_1, \mu_2, \dots, \mu_K]^\top$ gives the probabilities of each la-
 827 bel, so that the probability of drawing label y is equal to μ_y . For example, if $\mathcal{Y} =$
 828 $\{\text{POSITIVE}, \text{NEGATIVE}, \text{NEUTRAL}\}$, we might have $\boldsymbol{\mu} = [0.1, 0.7, 0.2]^\top$. We require
 829 $\sum_{y \in \mathcal{Y}} \mu_y = 1$ and $\mu_y \geq 0, \forall y \in \mathcal{Y}$.⁷
- 830 • The third line describes how the bag-of-words counts $\mathbf{x}^{(i)}$ are generated. By writing
 831 $\mathbf{x}^{(i)} | y^{(i)}$, this line indicates that the word counts are conditioned on the label, so

⁴Throughout this text, you may assume all logarithms and exponents are base 2, unless otherwise indicated. Any reasonable base will yield an identical classifier, and base 2 is most convenient for working out examples by hand.

⁵Generative models will be used throughout this text. They explicitly define the assumptions underlying the form of a probability distribution over observed and latent variables. For a readable introduction to generative models in statistics, see Blei (2014).

⁶Can you think of any cases in which this assumption is too strong?

⁷Formally, we require $\boldsymbol{\mu} \in \Delta^{K-1}$, where Δ^{K-1} is the $K - 1$ **probability simplex**, the set of all vectors of K nonnegative numbers that sum to one. Because of the sum-to-one constraint, there are $K - 1$ degrees of freedom for a vector of size K .

832 that the joint probability is factored using the chain rule,

$$p_{X,Y}(x^{(i)}, y^{(i)}) = p_{X|Y}(x^{(i)} | y^{(i)}) \times p_Y(y^{(i)}). \quad [2.11]$$

The specific distribution $p_{X|Y}$ is the **multinomial**, which is a probability distribution over vectors of non-negative counts. The probability mass function for this distribution is:

$$p_{\text{mult}}(x; \phi) = B(x) \prod_{j=1}^V \phi_j^{x_j} \quad [2.12]$$

$$B(x) = \frac{(\sum_{j=1}^V x_j)!}{\prod_{j=1}^V (x_j)!} \quad [2.13]$$

833 As in the categorical distribution, the parameter ϕ_j can be interpreted as a proba-
 834 bility: specifically, the probability that any given token in the document is the word
 835 j . The multinomial distribution involves a product over words, with each term in
 836 the product equal to the probability ϕ_j , exponentiated by the count x_j . Words that
 837 have zero count play no role in this product, because $\phi_j^0 = 1$. The term $B(x)$ doesn't
 838 depend on ϕ , and can usually be ignored. Can you see why we need this term at
 839 all?⁸

840 The notation $p(x | y; \phi)$ indicates the conditional probability of word counts x given
 841 label y , with parameter ϕ , which is equal to $p_{\text{mult}}(x; \phi_y)$. By specifying the multino-
 842 mial distribution, we describe the **multinomial naïve Bayes** classifier. Why “naïve”?
 843 Because the multinomial distribution treats each word token independently: the
 844 probability mass function factorizes across the counts.⁹

845 2.1.1 Types and tokens

846 A slight modification to the generative model of Naïve Bayes is shown in Algorithm 2.
 847 Instead of generating a vector of counts of **types**, x , this model generates a *sequence of*
 848 **tokens**, $w = (w_1, w_2, \dots, w_M)$. The distinction between types and tokens is critical: $x_j \in$
 849 $\{0, 1, 2, \dots, M\}$ is the count of word type j in the vocabulary, e.g., the number of times
 850 the word *cannibal* appears; $w_m \in \mathcal{V}$ is the identity of token m in the document, e.g. $w_m =$
 851 *cannibal*.

⁸Technically, a multinomial distribution requires a second parameter, the total number of word counts in x . In the bag-of-words representation is equal to the number of words in the document. However, this parameter is irrelevant for classification.

⁹You can plug in any probability distribution to the generative story and it will still be Naïve Bayes, as long as you are making the “naïve” assumption that the features are conditionally independent, given the label. For example, a multivariate Gaussian with diagonal covariance is naïve in exactly the same sense.

Algorithm 2 Alternative generative process for the Naïve Bayes classifier

```

for Document  $i \in \{1, 2, \dots, N\}$  do:
    Draw the label  $y^{(i)} \sim \text{Categorical}(\boldsymbol{\mu})$ ;
    for Token  $m \in \{1, 2, \dots, M_i\}$  do:
        Draw the token  $w_m^{(i)} | y^{(i)} \sim \text{Categorical}(\boldsymbol{\phi}_{y^{(i)}})$ .

```

852 The probability of the sequence \mathbf{w} is a product of categorical probabilities. Algo-
 853 rithm 2 makes a conditional independence assumption: each token $w_m^{(i)}$ is independent
 854 of all other tokens $w_{n \neq m}^{(i)}$, conditioned on the label $y^{(i)}$. This is identical to the “naïve”
 855 independence assumption implied by the multinomial distribution, and as a result, the
 856 optimal parameters for this model are identical to those in multinomial Naïve Bayes. For
 857 any instance, the probability assigned by this model is proportional to the probability un-
 858 der multinomial Naïve Bayes. The constant of proportionality is the factor $B(\mathbf{x})$, which
 859 appears in the multinomial distribution. Because $B(\mathbf{x}) \geq 1$, the probability for a vector
 860 of counts \mathbf{x} is at least as large as the probability for a list of words \mathbf{w} that induces the
 861 same counts: there can be many word sequences that correspond to a single vector of
 862 counts. For example, *man bites dog* and *dog bites man* correspond to an identical count vec-
 863 tor, $\{bites : 1, dog : 1, man : 1\}$, and $B(\mathbf{x})$ is equal to the total number of possible word
 864 orderings for count vector \mathbf{x} .

865 Sometimes it is useful to think of instances as counts of types, \mathbf{x} ; other times, it is
 866 better to think of them as sequences of tokens, \mathbf{w} . If the tokens are generated from a
 867 model that assumes conditional independence, then these two views lead to probability
 868 models that are identical, except for a scaling factor that does not depend on the label or
 869 the parameters.

870 **2.1.2 Prediction**

The Naïve Bayes prediction rule is to choose the label y which maximizes $\log p(\mathbf{x}, y; \boldsymbol{\mu}, \boldsymbol{\phi})$:

$$\hat{y} = \underset{y}{\operatorname{argmax}} \log p(\mathbf{x}, y; \boldsymbol{\mu}, \boldsymbol{\phi}) \quad [2.14]$$

$$= \underset{y}{\operatorname{argmax}} \log p(\mathbf{x} | y; \boldsymbol{\phi}) + \log p(y; \boldsymbol{\mu}) \quad [2.15]$$

Now we can plug in the probability distributions from the generative story.

$$\log p(\mathbf{x} | y; \boldsymbol{\phi}) + \log p(y; \boldsymbol{\mu}) = \log \left[B(\mathbf{x}) \prod_{j=1}^V \phi_{y,j}^{x_j} \right] + \log \mu_y \quad [2.16]$$

$$= \log B(\mathbf{x}) + \sum_{j=1}^V x_j \log \phi_{y,j} + \log \mu_y \quad [2.17]$$

$$= \log B(\mathbf{x}) + \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}, y), \quad [2.18]$$

where

$$\boldsymbol{\theta} = [\boldsymbol{\theta}^{(1)}; \boldsymbol{\theta}^{(2)}; \dots; \boldsymbol{\theta}^{(K)}] \quad [2.19]$$

$$\boldsymbol{\theta}^{(y)} = [\log \phi_{y,1}; \log \phi_{y,2}; \dots; \log \phi_{y,V}; \log \mu_y] \quad [2.20]$$

871 The feature function $\mathbf{f}(\mathbf{x}, y)$ is a vector of V word counts and an offset, padded by
 872 zeros for the labels not equal to y (see Equations 2.3-2.5, and Figure 2.1). This construction
 873 ensures that the inner product $\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}, y)$ only activates the features whose weights are
 874 in $\boldsymbol{\theta}^{(y)}$. These features and weights are all we need to compute the joint log-probability
 875 $\log p(\mathbf{x}, y)$ for each y . This is a key point: through this notation, we have converted the
 876 problem of computing the log-likelihood for a document-label pair (\mathbf{x}, y) into the compu-
 877 tation of a vector inner product.

878 2.1.3 Estimation

879 The parameters of the categorical and multinomial distributions have a simple interpre-
 880 tation: they are vectors of expected frequencies for each possible event. Based on this
 881 interpretation, it is tempting to set the parameters empirically,

$$\phi_{y,j} = \frac{\text{count}(y, j)}{\sum_{j'=1}^V \text{count}(y, j')} = \frac{\sum_{i:y^{(i)}=y} x_j^{(i)}}{\sum_{j'=1}^V \sum_{i:y^{(i)}=y} x_{j'}^{(i)}}, \quad [2.21]$$

882 where $\text{count}(y, j)$ refers to the count of word j in documents with label y .

883 Equation 2.21 defines the **relative frequency estimate** for ϕ . It can be justified as a
 884 **maximum likelihood estimate**: the estimate that maximizes the probability $p(\mathbf{x}^{(1:N)}, y^{(1:N)}; \boldsymbol{\theta})$.
 885 Based on the generative model in Algorithm 1, the log-likelihood is,

$$\mathcal{L}(\boldsymbol{\phi}, \boldsymbol{\mu}) = \sum_{i=1}^N \log p_{\text{mult}}(\mathbf{x}^{(i)}; \boldsymbol{\phi}_{y^{(i)}}) + \log p_{\text{cat}}(y^{(i)}; \boldsymbol{\mu}), \quad [2.22]$$

which is now written as a function \mathcal{L} of the parameters ϕ and μ . Let's continue to focus on the parameters ϕ . Since $p(y)$ is constant with respect to ϕ , we can drop it:

$$\mathcal{L}(\phi) = \sum_{i=1}^N \log p_{\text{mult}}(\mathbf{x}^{(i)}; \phi_{y^{(i)}}) = \sum_{i=1}^N \log B(\mathbf{x}^{(i)}) + \sum_{j=1}^V x_j^{(i)} \log \phi_{y^{(i)}, j}, \quad [2.23]$$

where $B(\mathbf{x}^{(i)})$ is constant with respect to ϕ .

We would now like to optimize the log-likelihood \mathcal{L} , by taking derivatives with respect to ϕ . But before we can do that, we have to deal with a set of constraints:

$$\sum_{j=1}^V \phi_{y,j} = 1 \quad \forall y \quad [2.24]$$

These constraints can be incorporated by adding a set of Lagrange multipliers (see Appendix B for more details). Solving separately for each label y , we obtain the Lagrangian,

$$\ell(\phi_y) = \sum_{i:y^{(i)}=y} \sum_{j=1}^V x_j^{(i)} \log \phi_{y,j} - \lambda \left(\sum_{j=1}^V \phi_{y,j} - 1 \right). \quad [2.25]$$

It is now possible to differentiate the Lagrangian with respect to the parameter of interest,

$$\frac{\partial \ell(\phi_y)}{\partial \phi_{y,j}} = \sum_{i:y^{(i)}=y} x_j^{(i)} / \phi_{y,j} - \lambda \quad [2.26]$$

The solution is obtained by setting each element in this vector of derivatives equal to zero,

$$\lambda \phi_{y,j} = \sum_{i:y^{(i)}=y} x_j^{(i)} \quad [2.27]$$

$$\phi_{y,j} \propto \sum_{i:y^{(i)}=y} x_j^{(i)} = \sum_{i=1}^N \delta(y^{(i)} = y) x_j^{(i)} = \text{count}(y, j), \quad [2.28]$$

where $\delta(y^{(i)} = y)$ is a **delta function**, also sometimes called an **indicator function**, which returns one if $y^{(i)} = y$, and zero otherwise. Equation 2.28 shows three different notations for the same thing: a sum over the word counts for all documents i such that the label $y^{(i)} = y$. This gives a solution for each ϕ_y up to a constant of proportionality. Now recall the constraint $\sum_{j=1}^V \phi_{y,j} = 1$, which arises because ϕ_y represents a vector of probabilities for each word in the vocabulary. This constraint leads to an exact solution,

$$\phi_{y,j} = \frac{\text{count}(y, j)}{\sum_{j'=1}^V \text{count}(y, j')}. \quad [2.29]$$

891 This is equal to the relative frequency estimator from Equation 2.21. A similar derivation
 892 gives $\mu_y \propto \sum_{i=1}^N \delta(y^{(i)} = y)$.

893 2.1.4 Smoothing and MAP estimation

894 With text data, there are likely to be pairs of labels and words that never appear in the
 895 training set, leaving $\phi_{y,j} = 0$. For example, the word *Bayesian* may have never yet ap-
 896 peared in a spam email. But choosing a value of $\phi_{\text{SPAM}, \text{Bayesian}} = 0$ would allow this single
 897 feature to completely veto a label, since $p(\text{SPAM} | x) = 0$ if $x_{\text{Bayesian}} > 0$.

898 This is undesirable, because it imposes high **variance**: depending on what data hap-
 899 pens to be in the training set, we could get vastly different classification rules. One so-
 900 lution is to **smooth** the probabilities, by adding a “pseudocount” of α to each count, and
 901 then normalizing.

$$\phi_{y,j} = \frac{\alpha + \text{count}(y, j)}{V\alpha + \sum_{j'=1}^V \text{count}(y, j')} \quad [2.30]$$

902 This is called **Laplace smoothing**.¹⁰ The pseudocount α is a **hyperparameter**, because it
 903 controls the form of the log-likelihood function, which in turn drives the estimation of ϕ .

904 Smoothing reduces variance, but it takes us away from the maximum likelihood esti-
 905 mate: it imposes a **bias**. In this case, the bias points towards uniform probabilities. Ma-
 906 chine learning theory shows that errors on heldout data can be attributed to the sum of
 907 bias and variance (Mohri et al., 2012). Techniques for reducing variance typically increase
 908 the bias, leading to a **bias-variance tradeoff**.

- 909 • Unbiased classifiers may **overfit** the training data, yielding poor performance on
 910 unseen data.
- 911 • But if the smoothing is too large, the resulting classifier can **underfit** instead. In the
 912 limit of $\alpha \rightarrow \infty$, there is zero variance: you get the same classifier, regardless of the
 913 data. However, the bias is likely to be large.

914 2.1.5 Setting hyperparameters

915 How should we choose the best value of hyperparameters like α ? Maximum likelihood
 916 will not work: the maximum likelihood estimate of α on the training set will always be
 917 $\alpha = 0$. In many cases, what we really want is **accuracy**: the number of correct predictions,
 918 divided by the total number of predictions. (Other measures of classification performance
 919 are discussed in § 4.4.) As we will see, it is hard to optimize for accuracy directly. But for
 920 scalar hyperparameters like α can be tuned by a simple heuristic called **grid search**: try a

¹⁰Laplace smoothing has a Bayesian justification, in which the generative model is extended to include ϕ as a random variable. The resulting estimate is called **maximum a posteriori**, or MAP.

921 set of values (e.g., $\alpha \in \{0.001, 0.01, 0.1, 1, 10\}$), compute the accuracy for each value, and
922 choose the setting that maximizes the accuracy.

923 The goal is to tune α so that the classifier performs well on *unseen* data. For this reason,
924 the data used for hyperparameter tuning should not overlap the training set, where very
925 small values of α will be preferred. Instead, we hold out a **development set** (also called
926 a **tuning set**) for hyperparameter selection. This development set may consist of a small
927 fraction of the labeled data, such as 10%.

928 We also want to predict the performance of our classifier on unseen data. To do this,
929 we must hold out a separate subset of data, called the **test set**. It is critical that the test set
930 not overlap with either the training or development sets, or else we will overestimate the
931 performance that the classifier will achieve on unlabeled data in the future. The test set
932 should also not be used when making modeling decisions, such as the form of the feature
933 function, the size of the vocabulary, and so on (these decisions are reviewed in chapter 4.)
934 The ideal practice is to use the test set only once — otherwise, the test set is used to guide
935 the classifier design, and test set accuracy will diverge from accuracy on truly unseen
936 data. Because annotated data is expensive, this ideal can be hard to follow in practice,
937 and many test sets have been used for decades. But in some high-impact applications like
938 machine translation and information extraction, new test sets are released every year.

939 When only a small amount of labeled data is available, the test set accuracy can be
940 unreliable. *K*-fold **cross-validation** is one way to cope with this scenario: the labeled
941 data is divided into *K* folds, and each fold acts as the test set, while training on the other
942 folds. The test set accuracies are then aggregated. In the extreme, each fold is a single data
943 point; this is called **leave-one-out** cross-validation. To perform hyperparameter tuning in
944 the context of cross-validation, another fold can be used for grid search. It is important
945 not to repeatedly evaluate the cross-validated accuracy while making design decisions
946 about the classifier, or you will overstate the accuracy on truly unseen data.

947 2.2 Discriminative learning

948 Naïve Bayes is easy to work with: the weights can be estimated in closed form, and the
949 probabilistic interpretation makes it relatively easy to extend. However, the assumption
950 that features are independent can seriously limit its accuracy. Thus far, we have defined
951 the **feature function** $f(\mathbf{x}, y)$ so that it corresponds to bag-of-words features: one feature
952 per word in the vocabulary. In natural language, bag-of-words features violate the as-
953 sumption of conditional independence — for example, the probability that a document
954 will contain the word *naïve* is surely higher given that it also contains the word *Bayes* —
955 but this violation is relatively mild.

956 However, good performance on text classification often requires features that are richer
957 than the bag-of-words:

- To better handle out-of-vocabulary terms, we want features that apply to multiple words, such as prefixes and suffixes (e.g., *anti-*, *un-*, *-ing*) and capitalization.
- We also want *n-gram* features that apply to multi-word units: **bigrams** (e.g., *not good*, *not bad*), **trigrams** (e.g., *not so bad*, *lacking any decency*, *never before imagined*), and beyond.

These features flagrantly violate the Naïve Bayes independence assumption. Consider what happens if we add a prefix feature. Under the Naïve Bayes assumption, we make the following approximation:¹¹

$$\Pr(\text{word} = \text{unfit}, \text{prefix} = \text{un-} \mid y) \approx \Pr(\text{prefix} = \text{un-} \mid y) \times \Pr(\text{word} = \text{unfit} \mid y).$$

To test the quality of the approximation, we can manipulate the left-hand side by applying the chain rule,

$$\Pr(\text{word} = \text{unfit}, \text{prefix} = \text{un-} \mid y) = \Pr(\text{prefix} = \text{un-} \mid \text{word} = \text{unfit}, y) \quad [2.31]$$

$$\times \Pr(\text{word} = \text{unfit} \mid y) \quad [2.32]$$

But $\Pr(\text{prefix} = \text{un-} \mid \text{word} = \text{unfit}, y) = 1$, since *un-* is guaranteed to be the prefix for the word *unfit*. Therefore,

$$\Pr(\text{word} = \text{unfit}, \text{prefix} = \text{un-} \mid y) = 1 \times \Pr(\text{word} = \text{unfit} \mid y) \quad [2.33]$$

$$\gg \Pr(\text{prefix} = \text{un-} \mid y) \times \Pr(\text{word} = \text{unfit} \mid y), \quad [2.34]$$

because the probability of any given word starting with the prefix *un-* is much less than one. Naïve Bayes will systematically underestimate the true probabilities of conjunctions of positively correlated features. To use such features, we need learning algorithms that do not rely on an independence assumption.

The origin of the Naïve Bayes independence assumption is the learning objective, $p(\mathbf{x}^{(1:N)}, y^{(1:N)})$, which requires modeling the probability of the observed text. In classification problems, we are always given \mathbf{x} , and are only interested in predicting the label y , so it seems unnecessary to model the probability of \mathbf{x} . **Discriminative learning** algorithms focus on the problem of predicting y , and do not attempt to model the probability of the text \mathbf{x} .

2.2.1 Perceptron

In Naïve Bayes, the weights can be interpreted as parameters of a probabilistic model. But this model requires an independence assumption that usually does not hold, and limits

¹¹The notation $\Pr(\cdot)$ refers to the probability of an event, and $p(\cdot)$ refers to the probability density or mass for a random variable (see Appendix A).

Algorithm 3 Perceptron learning algorithm

```

1: procedure PERCEPTRON( $\mathbf{x}^{(1:N)}, y^{(1:N)}$ )
2:    $t \leftarrow 0$ 
3:    $\boldsymbol{\theta}^{(0)} \leftarrow \mathbf{0}$ 
4:   repeat
5:      $t \leftarrow t + 1$ 
6:     Select an instance  $i$ 
7:      $\hat{y} \leftarrow \operatorname{argmax}_y \boldsymbol{\theta}^{(t-1)} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y)$ 
8:     if  $\hat{y} \neq y^{(i)}$  then
9:        $\boldsymbol{\theta}^{(t)} \leftarrow \boldsymbol{\theta}^{(t-1)} + \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) - \mathbf{f}(\mathbf{x}^{(i)}, \hat{y})$ 
10:    else
11:       $\boldsymbol{\theta}^{(t)} \leftarrow \boldsymbol{\theta}^{(t-1)}$ 
12:    until tired
13:   return  $\boldsymbol{\theta}^{(t)}$ 

```

976 our choice of features. Why not forget about probability and learn the weights in an error-
977 driven way? The **perceptron** algorithm, shown in Algorithm 3, is one way to do this.

978 Here's what the algorithm says: if you make a mistake, increase the weights for fea-
979 tures that are active with the correct label $y^{(i)}$, and decrease the weights for features that
980 are active with the guessed label \hat{y} . This is an **online learning** algorithm, since the clas-
981 sifier weights change after every example. This is different from Naïve Bayes, which
982 computes corpus statistics and then sets the weights in a single operation — Naïve Bayes
983 is a **batch learning** algorithm. Algorithm 3 is vague about when this online learning pro-
984 cedure terminates. We will return to this issue shortly.

985 The perceptron algorithm may seem like a cheap heuristic: Naïve Bayes has a solid
986 foundation in probability, but the perceptron is just adding and subtracting constants from
987 the weights every time there is a mistake. Will this really work? In fact, there is some nice
988 theory for the perceptron, based on the concept of **linear separability**:

989 **Definition 1** (Linear separability). *The dataset $\mathcal{D} = \{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^N$ is linearly separable iff
990 (if and only if) there exists some weight vector $\boldsymbol{\theta}$ and some margin ρ such that for every instance
991 $(\mathbf{x}^{(i)}, y^{(i)})$, the inner product of $\boldsymbol{\theta}$ and the feature function for the true label, $\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)})$, is
992 at least ρ greater than inner product of $\boldsymbol{\theta}$ and the feature function for every other possible label,
993 $\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y')$.*

$$\exists \boldsymbol{\theta}, \rho > 0 : \forall (\mathbf{x}^{(i)}, y^{(i)}) \in \mathcal{D}, \quad \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) \geq \rho + \max_{y' \neq y^{(i)}} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y'). \quad [2.35]$$

994 Linear separability is important because of the following guarantee: if your data is

995 linearly separable, then the perceptron algorithm will find a separator (Novikoff, 1962).¹²
 996 So while the perceptron may seem heuristic, it is guaranteed to succeed, if the learning
 997 problem is easy enough.

998 How useful is this proof? Minsky and Papert (1969) famously proved that the simple
 999 logical function of *exclusive-or* is not separable, and that a perceptron is therefore inca-
 1000 pable of learning this function. But this is not just an issue for the perceptron: any linear
 1001 classification algorithm, including Naïve Bayes, will fail on this task. In natural language
 1002 classification problems usually involve high dimensional feature spaces, with thousands
 1003 or millions of features. For these problems, it is very likely that the training data is indeed
 1004 separable. And even if the data is not separable, it is still possible to place an upper bound
 1005 on the number of errors that the perceptron algorithm will make (Freund and Schapire,
 1006 1999).

1007 2.2.2 Averaged perceptron

1008 The perceptron iterates over the data repeatedly — until “tired”, as described in Algo-
 1009 rithm 3. If the data is linearly separable, the perceptron will eventually find a separator,
 1010 and we can stop once all training instances are classified correctly. But if the data is not
 1011 linearly separable, the perceptron can *thrash* between two or more weight settings, never
 1012 converging. In this case, how do we know that we can stop training, and how should
 1013 we choose the final weights? An effective practical solution is to *average* the perceptron
 1014 weights across all iterations.

1015 This procedure is shown in Algorithm 4. The learning algorithm is nearly identical,
 1016 but we also maintain a vector of the sum of the weights, \mathbf{m} . At the end of the learning
 1017 procedure, we divide this sum by the total number of updates t , to compute the average
 1018 weights, $\bar{\theta}$. These average weights are then used for prediction. In the algorithm sketch,
 1019 the average is computed from a running sum, $\mathbf{m} \leftarrow \mathbf{m} + \theta$. However, this is inefficient,
 1020 because it requires $|\theta|$ operations to update the running sum. When $f(\mathbf{x}, y)$ is sparse,
 1021 $|\theta| \gg |f(\mathbf{x}, y)|$ for any individual (\mathbf{x}, y) . This means that computing the running sum will
 1022 be much more expensive than computing of the update to θ itself, which requires only
 1023 $2 \times |f(\mathbf{x}, y)|$ operations. One of the exercises is to sketch a more efficient algorithm for
 1024 computing the averaged weights.

1025 Even if the data is not separable, the averaged weights will eventually converge. One
 1026 possible stopping criterion is to check the difference between the average weight vectors
 1027 after each pass through the data: if the norm of the difference falls below some predefined
 1028 threshold, we can stop training. Another stopping criterion is to hold out some data,
 1029 and to measure the predictive accuracy on this heldout data. When the accuracy on the
 1030 heldout data starts to decrease, the learning algorithm has begun to **overfit** the training

¹²It is also possible to prove an upper bound on the number of training iterations required to find the separator. Proofs like this are part of the field of **statistical learning theory** (Mohri et al., 2012).

Algorithm 4 Averaged perceptron learning algorithm

```

1: procedure AVG-PERCEPTRON( $\mathbf{x}^{(1:N)}, \mathbf{y}^{(1:N)}$ )
2:    $t \leftarrow 0$ 
3:    $\boldsymbol{\theta}^{(0)} \leftarrow 0$ 
4:   repeat
5:      $t \leftarrow t + 1$ 
6:     Select an instance  $i$ 
7:      $\hat{y} \leftarrow \operatorname{argmax}_y \boldsymbol{\theta}^{(t-1)} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y)$ 
8:     if  $\hat{y} \neq y^{(i)}$  then
9:        $\boldsymbol{\theta}^{(t)} \leftarrow \boldsymbol{\theta}^{(t-1)} + \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) - \mathbf{f}(\mathbf{x}^{(i)}, \hat{y})$ 
10:    else
11:       $\boldsymbol{\theta}^{(t)} \leftarrow \boldsymbol{\theta}^{(t-1)}$ 
12:     $\mathbf{m} \leftarrow \mathbf{m} + \boldsymbol{\theta}^{(t)}$ 
13:   until tired
14:    $\bar{\boldsymbol{\theta}} \leftarrow \frac{1}{t} \mathbf{m}$ 
15:   return  $\bar{\boldsymbol{\theta}}$ 

```

1031 set. At this point, it is probably best to stop; this stopping criterion is known as **early**
 1032 **stopping**.

1033 **Generalization** is the ability to make good predictions on instances that are not in
 1034 the training data. Averaging can be proven to improve generalization, by computing an
 1035 upper bound on the generalization error (Freund and Schapire, 1999; Collins, 2002).

1036 **2.3 Loss functions and large-margin classification**

1037 Naïve Bayes chooses the weights $\boldsymbol{\theta}$ by maximizing the joint log-likelihood $\log p(\mathbf{x}^{(1:N)}, \mathbf{y}^{(1:N)})$.
 1038 By convention, optimization problems are generally formulated as minimization of a **loss**
 1039 **function**. The input to a loss function is the vector of weights $\boldsymbol{\theta}$, and the output is a non-
 1040 negative scalar, measuring the performance of the classifier on a training instance. The
 1041 loss $\ell(\boldsymbol{\theta}; \mathbf{x}^{(i)}, y^{(i)})$ is then a measure of the performance of the weights $\boldsymbol{\theta}$ on the instance
 1042 $(\mathbf{x}^{(i)}, y^{(i)})$. The goal of learning is to minimize the sum of the losses across all instances in
 1043 the training set.

We can trivially reformulate maximum likelihood as a loss function, by defining the

loss function to be the *negative log-likelihood*:

$$\log p(\mathbf{x}^{(1:N)}, y^{(1:N)}; \boldsymbol{\theta}) = \sum_{i=1}^N \log p(\mathbf{x}^{(i)}, y^{(i)}; \boldsymbol{\theta}) \quad [2.36]$$

$$\ell_{\text{NB}}(\boldsymbol{\theta}; \mathbf{x}^{(i)}, y^{(i)}) = -\log p(\mathbf{x}^{(i)}, y^{(i)}; \boldsymbol{\theta}) \quad [2.37]$$

$$\hat{\boldsymbol{\theta}} = \operatorname{argmin}_{\boldsymbol{\theta}} \sum_{i=1}^N \ell_{\text{NB}}(\boldsymbol{\theta}; \mathbf{x}^{(i)}, y^{(i)}) \quad [2.38]$$

$$= \operatorname{argmax}_{\boldsymbol{\theta}} \sum_{i=1}^N \log p(\mathbf{x}^{(i)}, y^{(i)}; \boldsymbol{\theta}). \quad [2.39]$$

1044 The problem of minimizing ℓ_{NB} is thus identical to the problem of maximum-likelihood
1045 estimation.

1046 Loss functions provide a general framework for comparing machine learning objectives.
1047 For example, an alternative loss function is the **zero-one loss**,

$$\ell_{0-1}(\boldsymbol{\theta}; \mathbf{x}^{(i)}, y^{(i)}) = \begin{cases} 0, & y^{(i)} = \operatorname{argmax}_y \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y) \\ 1, & \text{otherwise} \end{cases} \quad [2.40]$$

1048 The zero-one loss is zero if the instance is correctly classified, and one otherwise. The
1049 sum of zero-one losses is proportional to the error rate of the classifier on the training
1050 data. Since a low error rate is often the ultimate goal of classification, this may seem
1051 ideal. But the zero-one loss has several problems. One is that it is **non-convex**,¹³ which
1052 means that there is no guarantee that gradient-based optimization will be effective. A
1053 more serious problem is that the derivatives are useless: the partial derivative with respect
1054 to any parameter is zero everywhere, except at the points where $\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y) = \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, \hat{y})$
1055 for some \hat{y} . At those points, the loss is discontinuous, and the derivative is undefined.

1056 The perceptron optimizes the following loss function:

$$\ell_{\text{PERCEPTRON}}(\boldsymbol{\theta}; \mathbf{x}^{(i)}, y^{(i)}) = \max_{y \in \mathcal{Y}} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y) - \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}), \quad [2.41]$$

1057 When $\hat{y} = y^{(i)}$, the loss is zero; otherwise, it increases linearly with the gap between the
1058 score for the predicted label \hat{y} and the score for the true label $y^{(i)}$. Plotting this loss against
1059 the input $\max_{y \in \mathcal{Y}} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y) - \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)})$ gives a hinge shape, motivating the name
1060 **hinge loss**.

¹³A function f is **convex** iff $\alpha f(x_i) + (1-\alpha)f(x_j) \geq f(\alpha x_i + (1-\alpha)x_j)$, for all $\alpha \in [0, 1]$ and for all x_i and x_j on the domain of the function. In words, any weighted average of the output of f applied to any two points is larger than the output of f when applied to the weighted average of the same two points. Convexity implies that any local minimum is also a global minimum, and there are many effective techniques for optimizing convex functions (Boyd and Vandenberghe, 2004). See Appendix B for a brief review.

1061 To see why this is the loss function optimized by the perceptron, take the derivative
 1062 with respect to θ ,

$$\frac{\partial}{\partial \theta} \ell_{\text{PERCEPTRON}}(\theta; \mathbf{x}^{(i)}, y^{(i)}) = \mathbf{f}(\mathbf{x}^{(i)}, \hat{y}) - \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}). \quad [2.42]$$

1063 At each instance perceptron algorithm takes a step of magnitude one in the opposite direction
 1064 of this **gradient**, $\nabla_{\theta} \ell_{\text{PERCEPTRON}} = \frac{\partial}{\partial \theta} \ell_{\text{PERCEPTRON}}(\theta; \mathbf{x}^{(i)}, y^{(i)})$. As we will see in § 2.5,
 1065 this is an example of the optimization algorithm **stochastic gradient descent**, applied to
 1066 the objective in Equation 2.41.

1067 **Breaking ties with subgradient descent** Careful readers will notice the tacit assumption
 1068 that there is a unique \hat{y} that maximizes $\theta \cdot \mathbf{f}(\mathbf{x}^{(i)}, y)$. What if there are two or more labels
 1069 that maximize this function? Consider binary classification: if the maximizer is $y^{(i)}$, then
 1070 the gradient is zero, and so is the perceptron update; if the maximizer is $\hat{y} \neq y^{(i)}$, then the
 1071 update is the difference $\mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) - \mathbf{f}(\mathbf{x}^{(i)}, \hat{y})$. The underlying issue is that the perceptron
 1072 loss is not **smooth**, because the first derivative has a discontinuity at the hinge point,
 1073 where the score for the true label $y^{(i)}$ is equal to the score for some other label \hat{y} . At this
 1074 point, there is no unique gradient; rather, there is a set of **subgradients**. A vector v is a
 1075 subgradient of the function g at u_0 iff $g(u) - g(u_0) \geq v \cdot (u - u_0)$ for all u . Graphically,
 1076 this defines the set of hyperplanes that include $g(u_0)$ and do not intersect g at any other
 1077 point. As we approach the hinge point from the left, the gradient is $\mathbf{f}(\mathbf{x}, \hat{y}) - \mathbf{f}(\mathbf{x}, y)$; as we
 1078 approach from the right, the gradient is 0. At the hinge point, the subgradients include all
 1079 vectors that are bounded by these two extremes. In subgradient descent, *any* subgradient
 1080 can be used (Bertsekas, 2012). Since both 0 and $\mathbf{f}(\mathbf{x}, \hat{y}) - \mathbf{f}(\mathbf{x}, y)$ are subgradients at the
 1081 hinge point, either one can be used in the perceptron update.

1082 **Perceptron versus Naïve Bayes** The perceptron loss function has some pros and cons
 1083 with respect to the negative log-likelihood loss implied by Naïve Bayes.

- 1084 • Both ℓ_{NB} and $\ell_{\text{PERCEPTRON}}$ are convex, making them relatively easy to optimize. How-
 1085 ever, ℓ_{NB} can be optimized in closed form, while $\ell_{\text{PERCEPTRON}}$ requires iterating over
 1086 the dataset multiple times.
- 1087 • ℓ_{NB} can suffer **infinite** loss on a single example, since the logarithm of zero probabil-
 1088 ity is negative infinity. Naïve Bayes will therefore overemphasize some examples,
 1089 and underemphasize others.
- 1090 • $\ell_{\text{PERCEPTRON}}$ treats all correct answers equally. Even if θ only gives the correct answer
 1091 by a tiny margin, the loss is still zero.

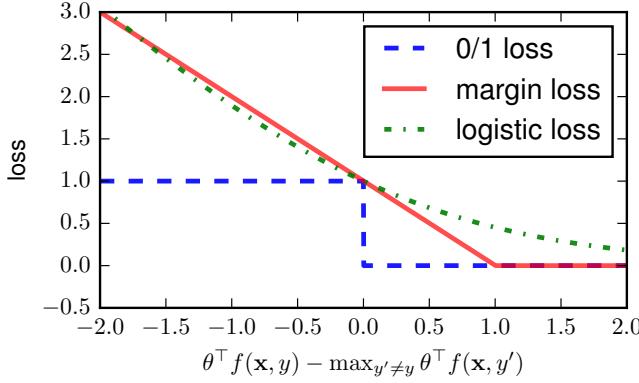


Figure 2.2: Margin, zero-one, and logistic loss functions.

1092 **2.3.1 Large margin classification**

1093 This last comment suggests a potential problem with the perceptron. Suppose a test ex-
 1094 ample is very close to a training example, but not identical. If the classifier only gets the
 1095 correct answer on the training example by a small margin, then it may get the test instance
 1096 wrong. To formalize this intuition, define the **margin** as,

$$\gamma(\boldsymbol{\theta}; \mathbf{x}^{(i)}, y^{(i)}) = \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) - \max_{y \neq y^{(i)}} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y). \quad [2.43]$$

The margin represents the difference between the score for the correct label $y^{(i)}$, and the score for the highest-scoring label. The intuition behind **large margin classification** is that it is not enough just to label the training data correctly — the correct label should be separated from other labels by a comfortable margin. This idea can be encoded into a loss function,

$$\ell_{\text{MARGIN}}(\boldsymbol{\theta}; \mathbf{x}^{(i)}, y^{(i)}) = \begin{cases} 0, & \gamma(\boldsymbol{\theta}; \mathbf{x}^{(i)}, y^{(i)}) \geq 1, \\ 1 - \gamma(\boldsymbol{\theta}; \mathbf{x}^{(i)}, y^{(i)}), & \text{otherwise} \end{cases} \quad [2.44]$$

$$= (1 - \gamma(\boldsymbol{\theta}; \mathbf{x}^{(i)}, y^{(i)}))_+, \quad [2.45]$$

1097 where $(x)_+ = \max(0, x)$. The loss is zero if there is a margin of at least 1 between the
 1098 score for the true label and the best-scoring alternative \hat{y} . This is almost identical to the
 1099 perceptron loss, but the hinge point is shifted to the right, as shown in Figure 2.2. The
 1100 margin loss is a convex upper bound on the zero-one loss.

1101 **2.3.2 Support vector machines**

If a dataset is linearly separable, then there is some hyperplane θ that correctly classifies all training instances with margin ρ (by Definition 1). This margin can be increased to any desired value by multiplying the weights by a constant. Now, for any datapoint $(x^{(i)}, y^{(i)})$, the geometric distance to the separating hyperplane is given by $\frac{\gamma(\theta; x^{(i)}, y^{(i)})}{\|\theta\|_2}$,

where the denominator is the norm of the weights, $\|\theta\|_2 = \sqrt{\sum_j \theta_j^2}$. The geometric distance is sometimes called the **geometric margin**, in contrast to the **functional margin** $\gamma(\theta; x^{(i)}, y^{(i)})$. Both are shown in Figure 2.3. The geometric margin is a good measure of the robustness of the separator: if the functional margin is large, but the norm $\|\theta\|_2$ is also large, then a small change in $x^{(i)}$ could cause it to be misclassified. We therefore seek to maximize the minimum geometric margin, subject to the constraint that the functional margin is at least one:

$$\begin{aligned} \max_{\theta} . & \quad \min_i . & & \frac{\gamma(\theta; x^{(i)}, y^{(i)})}{\|\theta\|_2} \\ \text{s.t.} & \quad \gamma(\theta; x^{(i)}, y^{(i)}) \geq 1, \quad \forall i. & & [2.46] \end{aligned}$$

1102 This is a **constrained optimization** problem, where the second line describes constraints
 1103 on the space of possible solutions θ . In this case, the constraint is that the functional
 1104 margin always be at least one, and the objective is that the minimum geometric margin
 1105 be as large as possible.

Any scaling factor on θ will cancel in the numerator and denominator of the geometric margin. This means that if the data is linearly separable at ρ , we can increase this margin to 1 by rescaling θ . We therefore need only minimize the denominator $\|\theta\|_2$, subject to the constraint on the functional margin. The minimizer of $\|\theta\|_2$ is also the minimizer of $\frac{1}{2}\|\theta\|_2^2 = \frac{1}{2}\sum_{j=1}^V \theta_j^2$, which is easier to work with. This gives the optimization problem,

$$\begin{aligned} \min_{\theta} . & \quad \frac{1}{2}\|\theta\|_2^2 \\ \text{s.t.} & \quad \gamma(\theta; x^{(i)}, y^{(i)}) \geq 1, \quad \forall i. & & [2.47] \end{aligned}$$

1106 This optimization problem is a **quadratic program**: the objective is a quadratic function
 1107 of the parameters, and the constraints are all linear inequalities. The resulting classifier
 1108 is better known as the **support vector machine**. The name derives from one of the
 1109 solutions, which is to incorporate the constraints through Lagrange multipliers $\alpha_i \geq 0, i =$
 1110 $1, 2, \dots, N$. The instances for which $\alpha_i > 0$ are the **support vectors**; other instances are
 1111 irrelevant to the classification boundary.



Figure 2.3: Functional and geometric margins for a binary classification problem. All separators that satisfy the margin constraint are shown. The separator with the largest geometric margin is shown in bold.

1112 2.3.3 Slack variables

If a dataset is not linearly separable, then there is no θ that satisfies the margin constraint. To add more flexibility, we introduce a set of **slack variables** $\xi_i \geq 0$. Instead of requiring that the functional margin be greater than or equal to one, we require that it be greater than or equal to $1 - \xi_i$. Ideally there would not be any slack, so the slack variables are penalized in the objective function:

$$\begin{aligned} \min_{\theta, \xi} \quad & \frac{1}{2} \|\theta\|_2^2 + C \sum_{i=1}^N \xi_i \\ \text{s.t.} \quad & \gamma(\theta; \mathbf{x}^{(i)}, y^{(i)}) + \xi_i \geq 1, \quad \forall i \\ & \xi_i \geq 0, \quad \forall i. \end{aligned} \quad [2.48]$$

1113 The hyperparameter C controls the tradeoff between violations of the margin con-
 1114 straint and the preference for a low norm of θ . As $C \rightarrow \infty$, slack is infinitely expensive,
 1115 and there is only a solution if the data is separable. As $C \rightarrow 0$, slack becomes free, and
 1116 there is a trivial solution at $\theta = 0$. Thus, C plays a similar role to the smoothing parame-
 1117 ter in Naïve Bayes (§ 2.1.4), trading off between a close fit to the training data and better
 1118 generalization. Like the smoothing parameter of Naïve Bayes, C must be set by the user,
 1119 typically by maximizing performance on a heldout development set.

1120 To solve the constrained optimization problem defined in Equation 2.48, we can first

1121 solve for the slack variables,

$$\xi_i \geq (1 - \gamma(\boldsymbol{\theta}; \mathbf{x}^{(i)}, y^{(i)}))_+. \quad [2.49]$$

The inequality is tight, because the slack variables are penalized in the objective, and there is no advantage to increasing them beyond the minimum value (Ratliff et al., 2007; Smith, 2011). The problem can therefore be transformed into the unconstrained optimization,

$$\min_{\boldsymbol{\theta}} \frac{\lambda}{2} \|\boldsymbol{\theta}\|_2^2 + \sum_{i=1}^N (1 - \gamma(\boldsymbol{\theta}; \mathbf{x}^{(i)}, y^{(i)}))_+, \quad [2.50]$$

1122 where each ξ_i has been substituted by the right-hand side of Equation 2.49, and the factor
 1123 of C on the slack variables has been replaced by an equivalent factor of $\lambda = \frac{1}{C}$ on the
 1124 norm of the weights.

1125 Now define the **cost** of a classification error as,¹⁴

$$c(y^{(i)}, \hat{y}) = \begin{cases} 1, & y^{(i)} \neq \hat{y} \\ 0, & \text{otherwise.} \end{cases} \quad [2.51]$$

Equation 2.50 can be rewritten using this cost function,

$$\min_{\boldsymbol{\theta}} \frac{\lambda}{2} \|\boldsymbol{\theta}\|_2^2 + \sum_{i=1}^N \left(\max_{y \in \mathcal{Y}} (\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y) + c(y^{(i)}, y)) - \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) \right)_+. \quad [2.52]$$

1126 This objective maximizes over all $y \in \mathcal{Y}$, in search of labels that are both *strong*, as mea-
 1127 sured by $\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y)$, and *wrong*, as measured by $c(y^{(i)}, y)$. This maximization is known
 1128 as **cost-augmented decoding**, because it augments the maximization objective to favor
 1129 high-cost predictions. If the highest-scoring label is $y = y^{(i)}$, then the margin constraint is
 1130 satisfied, and the loss for this instance is zero. Cost-augmentation is only for learning: it
 1131 is not applied when making predictions on unseen data.

Differentiating Equation 2.52 with respect to the weights gives,

$$\nabla_{\boldsymbol{\theta}} L_{\text{SVM}} = \lambda \boldsymbol{\theta} + \sum_{i=1}^N \mathbf{f}(\mathbf{x}^{(i)}, \hat{y}) - \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) \quad [2.53]$$

$$\hat{y} = \operatorname{argmax}_{y \in \mathcal{Y}} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y) + c(y^{(i)}, y), \quad [2.54]$$

1132 where L_{SVM} refers to minimization objective in Equation 2.52. This gradient is very similar
 1133 to the perceptron update. One difference is the additional term $\lambda \boldsymbol{\theta}$, which **regularizes** the

¹⁴We can also define specialized cost functions that heavily penalize especially undesirable errors (Tsochantaridis et al., 2004). This idea is revisited in chapter 7.

weights towards 0. The other difference is the cost $c(y^{(i)}, y)$, which is added to $\theta \cdot \mathbf{f}(\mathbf{x}, y)$ when choosing \hat{y} during training. This term derives from the margin constraint: large margin classifiers learn not only from instances that are incorrectly classified, but also from instances for which the correct classification decision was not sufficiently confident.

2.4 Logistic regression

Thus far, we have seen two broad classes of learning algorithms. Naïve Bayes is a probabilistic method, where learning is equivalent to estimating a joint probability distribution. The perceptron and support vector machine are discriminative, error-driven algorithms: the learning objective is closely related to the number of errors on the training data. Probabilistic and error-driven approaches each have advantages: probability makes it possible to quantify uncertainty about the predicted labels, but the probability model of Naïve Bayes makes unrealistic independence assumptions that limit the features that can be used.

Logistic regression combines advantages of discriminative and probabilistic classifiers. Unlike Naïve Bayes, which starts from the **joint probability** $p_{X,Y}$, logistic regression defines the desired **conditional probability** $p_{Y|X}$ directly. Think of $\theta \cdot \mathbf{f}(\mathbf{x}, y)$ as a scoring function for the compatibility of the base features \mathbf{x} and the label y . To convert this score into a probability, we first exponentiate, obtaining $\exp(\theta \cdot \mathbf{f}(\mathbf{x}, y))$, which is guaranteed to be non-negative. Next, we normalize, dividing over all possible labels $y' \in \mathcal{Y}$. The resulting conditional probability is defined as,

$$p(y | \mathbf{x}; \theta) = \frac{\exp(\theta \cdot \mathbf{f}(\mathbf{x}, y))}{\sum_{y' \in \mathcal{Y}} \exp(\theta \cdot \mathbf{f}(\mathbf{x}, y'))}. \quad [2.55]$$

Given a dataset $\mathcal{D} = \{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^N$, the weights θ are estimated by **maximum conditional likelihood**,

$$\log p(\mathbf{y}^{(1:N)} | \mathbf{x}^{(1:N)}; \theta) = \sum_{i=1}^N \log p(y^{(i)} | \mathbf{x}^{(i)}; \theta) \quad [2.56]$$

$$= \sum_{i=1}^N \theta \cdot \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) - \log \sum_{y' \in \mathcal{Y}} \exp(\theta \cdot \mathbf{f}(\mathbf{x}^{(i)}, y')). \quad [2.57]$$

The final line is obtained by plugging in Equation 2.55 and taking the logarithm.¹⁵ Inside

¹⁵The log-sum-exp term is a common pattern in machine learning. It is numerically unstable, because it will underflow if the inner product is small, and overflow if the inner product is large. Scientific computing libraries usually contain special functions for computing `logsumexp`, but with some thought, you should be able to see how to create an implementation that is numerically stable.

1148 the sum, we have the (additive inverse of the) **logistic loss**,

$$\ell_{\text{LOGREG}}(\boldsymbol{\theta}; \mathbf{x}^{(i)}, y^{(i)}) = -\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) + \log \sum_{y' \in \mathcal{Y}} \exp(\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y')) \quad [2.58]$$

1149 The logistic loss is shown in Figure 2.2. A key difference from the zero-one and hinge
 1150 losses is that logistic loss is never zero. This means that the objective function can always
 1151 be improved by assigning higher confidence to the correct label.

1152 2.4.1 Regularization

1153 As with the support vector machine, better generalization can be obtained by penalizing
 1154 the norm of $\boldsymbol{\theta}$. This is done by adding a term of $\frac{\lambda}{2} \|\boldsymbol{\theta}\|_2^2$ to the minimization objective.
 1155 This is called L_2 regularization, because $\|\boldsymbol{\theta}\|_2^2$ is the squared L_2 norm of the vector $\boldsymbol{\theta}$.
 1156 Regularization forces the estimator to trade off performance on the training data against
 1157 the norm of the weights, and this can help to prevent overfitting. Consider what would
 1158 happen to the unregularized weight for a base feature j that is active in only one instance
 1159 $\mathbf{x}^{(i)}$: the conditional log-likelihood could always be improved by increasing the weight
 1160 for this feature, so that $\boldsymbol{\theta}_{(j,y^{(i)})} \rightarrow \infty$ and $\boldsymbol{\theta}_{(j,\tilde{y} \neq y^{(i)})} \rightarrow -\infty$, where (j, y) is the index of
 1161 feature associated with $x_j^{(i)}$ and label y in $\mathbf{f}(\mathbf{x}^{(i)}, y)$.

In § 2.1.4, we saw that smoothing the probabilities of a Naïve Bayes classifier can be justified in a hierarchical probabilistic model, in which the parameters of the classifier are themselves random variables, drawn from a prior distribution. The same justification applies to L_2 regularization. In this case, the prior is a zero-mean Gaussian on each term of $\boldsymbol{\theta}$. The log-likelihood under a zero-mean Gaussian is,

$$\log N(\theta_j; 0, \sigma^2) \propto -\frac{1}{2\sigma^2} \theta_j^2, \quad [2.59]$$

1162 so that the regularization weight λ is equal to the inverse variance of the prior, $\lambda = \frac{1}{\sigma^2}$.

1163 **2.4.2 Gradients**

Logistic loss is minimized by optimization along the gradient. Here is the gradient with respect to the logistic loss on a single example,

$$\ell_{\text{LOGREG}} = -\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) + \log \sum_{y' \in \mathcal{Y}} \exp(\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y')) \quad [2.60]$$

$$\frac{\partial \ell}{\partial \boldsymbol{\theta}} = -\mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) + \frac{1}{\sum_{y'' \in \mathcal{Y}} \exp(\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y''))} \times \sum_{y' \in \mathcal{Y}} \exp(\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y')) \times \mathbf{f}(\mathbf{x}^{(i)}, y') \quad [2.61]$$

$$= -\mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) + \sum_{y' \in \mathcal{Y}} \frac{\exp(\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y'))}{\sum_{y'' \in \mathcal{Y}} \exp(\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y''))} \times \mathbf{f}(\mathbf{x}^{(i)}, y') \quad [2.62]$$

$$= -\mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) + \sum_{y' \in \mathcal{Y}} p(y' | \mathbf{x}^{(i)}; \boldsymbol{\theta}) \times \mathbf{f}(\mathbf{x}^{(i)}, y') \quad [2.63]$$

$$= -\mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) + E_{Y|X}[\mathbf{f}(\mathbf{x}^{(i)}, y)]. \quad [2.64]$$

1164 The final step employs the definition of a conditional expectation (§ A.5). The gradient of
 1165 the logistic loss is equal to the difference between the expected counts under the current
 1166 model, $E_{Y|X}[\mathbf{f}(\mathbf{x}^{(i)}, y)]$, and the observed feature counts $\mathbf{f}(\mathbf{x}^{(i)}, y^{(i)})$. When these two
 1167 vectors are equal for a single instance, there is nothing more to learn from it; when they
 1168 are equal in sum over the entire dataset, there is nothing more to learn from the dataset as
 1169 a whole. The gradient of the hinge loss is nearly identical, but it involves the features of
 1170 the predicted label under the current model, $\mathbf{f}(\mathbf{x}^{(i)}, \hat{y})$, rather than the expected features
 1171 $E_{Y|X}[\mathbf{f}(\mathbf{x}^{(i)}, y)]$ under the conditional distribution $p(y | \mathbf{x}; \boldsymbol{\theta})$.

The regularizer contributes $\lambda \boldsymbol{\theta}$ to the overall gradient:

$$L_{\text{LOGREG}} = \frac{\lambda}{2} \|\boldsymbol{\theta}\|_2^2 - \sum_{i=1}^N \left(\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) - \log \sum_{y' \in \mathcal{Y}} \exp \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y') \right) \quad [2.65]$$

$$\nabla_{\boldsymbol{\theta}} L_{\text{LOGREG}} = \lambda \boldsymbol{\theta} - \sum_{i=1}^N \left(\mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) - E_{y|\mathbf{x}}[\mathbf{f}(\mathbf{x}^{(i)}, y)] \right). \quad [2.66]$$

1172 **2.5 Optimization**

1173 Each of the classification algorithms in this chapter can be viewed as an optimization
 1174 problem:

- 1175 • In Naïve Bayes, the objective is the joint likelihood $\log p(\mathbf{x}^{(1:N)}, \mathbf{y}^{(1:N)})$. Maximum
 1176 likelihood estimation yields a closed-form solution for $\boldsymbol{\theta}$.

- 1177 • In the support vector machine, the objective is the regularized margin loss,

$$L_{\text{SVM}} = \frac{\lambda}{2} \|\boldsymbol{\theta}\|_2^2 + \sum_{i=1}^N (\max_{y \in \mathcal{Y}} (\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y) + c(y^{(i)}, y)) - \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}))_+, \quad [2.67]$$

1178 There is no closed-form solution, but the objective is convex. The perceptron algo-
1179 rithm minimizes a similar objective.

- 1180 • In logistic regression, the objective is the regularized negative log-likelihood,

$$L_{\text{LOGREG}} = \frac{\lambda}{2} \|\boldsymbol{\theta}\|_2^2 - \sum_{i=1}^N \left(\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) - \log \sum_{y \in \mathcal{Y}} \exp(\boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y)) \right) \quad [2.68]$$

1181 Again, there is no closed-form solution, but the objective is convex.

1182 These learning algorithms are distinguished by *what* is being optimized, rather than
1183 *how* the optimal weights are found. This decomposition is an essential feature of con-
1184 temporary machine learning. The domain expert's job is to design an objective function
1185 — or more generally, a **model** of the problem. If the model has certain characteristics,
1186 then generic optimization algorithms can be used to find the solution. In particular, if an
1187 objective function is differentiable, then gradient-based optimization can be employed;
1188 if it is also convex, then gradient-based optimization is guaranteed to find the globally
1189 optimal solution. The support vector machine and logistic regression have both of these
1190 properties, and so are amenable to generic **convex optimization** techniques (Boyd and
1191 Vandenberghe, 2004).

1192 2.5.1 Batch optimization

In **batch optimization**, each update to the weights is based on a computation involving the entire dataset. One such algorithm is **gradient descent**, which iteratively updates the weights,

$$\boldsymbol{\theta}^{(t+1)} \leftarrow \boldsymbol{\theta}^{(t)} - \eta^{(t)} \nabla_{\boldsymbol{\theta}} L, \quad [2.69]$$

1193 where $\nabla_{\boldsymbol{\theta}} L$ is the gradient computed over the entire training set, and $\eta^{(t)}$ is the **step size**
1194 at iteration t . If the objective L is a convex function of $\boldsymbol{\theta}$, then this procedure is guaranteed
1195 to terminate at the global optimum, for appropriate schedule of learning rates, $\eta^{(t)}$.¹⁶

¹⁶Specifically, the learning rate must have the following properties (Bottou et al., 2016):

$$\sum_{t=1}^{\infty} \eta^{(t)} = \infty \quad [2.70]$$

$$\sum_{t=1}^{\infty} (\eta^{(t)})^2 < \infty. \quad [2.71]$$

1196 In practice, gradient descent can be slow to converge, as the gradient can become
 1197 infinitesimally small. Faster convergence can be obtained by second-order Newton opti-
 1198 mization, which incorporates the inverse of the **Hessian matrix**,

$$H_{i,j} = \frac{\partial^2 L}{\partial \theta_i \partial \theta_j} \quad [2.72]$$

1199 The size of the Hessian matrix is quadratic in the number of features. In the bag-of-words
 1200 representation, this is usually too big to store, let alone invert. **Quasi-Network optimiza-**
 1201 **tion** techniques maintain a low-rank approximation to the inverse of the Hessian matrix.
 1202 Such techniques usually converge more quickly than gradient descent, while remaining
 1203 computationally tractable even for large feature sets. A popular quasi-Newton algorithm
 1204 is **L-BFGS** (Liu and Nocedal, 1989), which is implemented in many scientific computing
 1205 environments, such as `scipy` and `Matlab`.

1206 For any gradient-based technique, the user must set the learning rates $\eta^{(t)}$. While con-
 1207 vergence proofs usually employ a decreasing learning rate, in practice, it is common to fix
 1208 $\eta^{(t)}$ to a small constant, like 10^{-3} . The specific constant can be chosen by experimentation,
 1209 although there is research on determining the learning rate automatically (Schaul et al.,
 1210 2013; Wu et al., 2018).

1211 2.5.2 Online optimization

1212 Batch optimization computes the objective on the entire training set before making an up-
 1213 date. This may be inefficient, because at early stages of training, a small number of train-
 1214 ing examples could point the learner in the correct direction. **Online learning** algorithms
 1215 make updates to the weights while iterating through the training data. The theoretical
 1216 basis for this approach is a stochastic approximation to the true objective function,

$$\sum_{i=1}^N \ell(\boldsymbol{\theta}; \mathbf{x}^{(i)}, y^{(i)}) \approx N \times \ell(\boldsymbol{\theta}; \mathbf{x}^{(j)}, y^{(j)}), \quad (\mathbf{x}^{(j)}, y^{(j)}) \sim \{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^N, \quad [2.73]$$

1217 where the instance $(\mathbf{x}^{(j)}, y^{(j)})$ is sampled at random from the full dataset.

1218 In **stochastic gradient descent**, the approximate gradient is computed by randomly
 1219 sampling a single instance, and an update is made immediately. This is similar to the
 1220 perceptron algorithm, which also updates the weights one instance at a time. In **mini-**
 1221 **batch** stochastic gradient descent, the gradient is computed over a small set of instances.
 1222 A typical approach is to set the minibatch size so that the entire batch fits in memory on a
 1223 graphics processing unit (GPU; Neubig et al., 2017). It is then possible to speed up learn-
 1224 ing by parallelizing the computation of the gradient over each instance in the minibatch.

These properties can be obtained by the learning rate schedule $\eta^{(t)} = \eta^{(0)} t^{-\alpha}$ for $\alpha \in [1, 2]$.

Algorithm 5 Generalized gradient descent. The function BATCHER partitions the training set into B batches such that each instance appears in exactly one batch. In gradient descent, $B = 1$; in stochastic gradient descent, $B = N$; in minibatch stochastic gradient descent, $1 < B < N$.

```

1: procedure GRADIENT-DESCENT( $\mathbf{x}^{(1:N)}, \mathbf{y}^{(1:N)}, L, \eta^{(1:\infty)}$ , BATCHER,  $T_{\max}$ )
2:    $\boldsymbol{\theta} \leftarrow \mathbf{0}$ 
3:    $t \leftarrow 0$ 
4:   repeat
5:      $(\mathbf{b}^{(1)}, \mathbf{b}^{(2)}, \dots, \mathbf{b}^{(B)}) \leftarrow \text{BATCHER}(N)$ 
6:     for  $n \in \{1, 2, \dots, B\}$  do
7:        $t \leftarrow t + 1$ 
8:        $\boldsymbol{\theta}^{(t)} \leftarrow \boldsymbol{\theta}^{(t-1)} - \eta^{(t)} \nabla_{\boldsymbol{\theta}} L(\boldsymbol{\theta}^{(t-1)}; \mathbf{x}^{(b_1^{(n)}, b_2^{(n)}, \dots)}, \mathbf{y}^{(b_1^{(n)}, b_2^{(n)}, \dots)})$ 
9:       if Converged( $\boldsymbol{\theta}^{(1, 2, \dots, t)}$ ) then
10:        return  $\boldsymbol{\theta}^{(t)}$ 
11:   until  $t \geq T_{\max}$ 
12:   return  $\boldsymbol{\theta}^{(t)}$ 

```

1225 Algorithm 5 offers a generalized view of gradient descent. In standard gradient de-
 1226 scendent, the batcher returns a single batch with all the instances. In stochastic gradient de-
 1227 scent, it returns N batches with one instance each. In mini-batch settings, the batcher
 1228 returns B minibatches, $1 < B < N$.

There are many other techniques for online learning, and the field is currently quite active (Bottou et al., 2016). Some algorithms use an adaptive step size, which can be different for every feature (Duchi et al., 2011). Features that occur frequently are likely to be updated frequently, so it is best to use a small step size; rare features will be updated infrequently, so it is better to take larger steps. The **AdaGrad** (adaptive gradient) algorithm achieves this behavior by storing the sum of the squares of the gradients for each feature, and rescaling the learning rate by its inverse:

$$\mathbf{g}_t = \nabla_{\boldsymbol{\theta}} L(\boldsymbol{\theta}^{(t)}; \mathbf{x}^{(i)}, y^{(i)}) \quad [2.74]$$

$$\boldsymbol{\theta}_j^{(t+1)} \leftarrow \boldsymbol{\theta}_j^{(t)} - \frac{\eta^{(t)}}{\sqrt{\sum_{t'=1}^t g_{t,j}^2}} g_{t,j}, \quad [2.75]$$

1229 where j iterates over features in $\mathbf{f}(\mathbf{x}, y)$.

1230 In most cases, the number of active features for any instance is much smaller than the
 1231 number of weights. If so, the computation cost of online optimization will be dominated
 1232 by the update from the regularization term, $\lambda \boldsymbol{\theta}$. The solution is to be “lazy”, updating
 1233 each $\boldsymbol{\theta}_j$ only as it is used. To implement lazy updating, store an additional parameter τ_j ,
 1234 which is the iteration at which $\boldsymbol{\theta}_j$ was last updated. If $\boldsymbol{\theta}_j$ is needed at time t , the $t - \tau$

1235 regularization updates can be performed all at once. This strategy is described in detail
 1236 by Kummerfeld et al. (2015).

1237 2.6 *Additional topics in classification

1238 Throughout this text, advanced topics will be marked with an asterisk.

1239 2.6.1 Feature selection by regularization

1240 In logistic regression and large-margin classification, generalization can be improved by
 1241 regularizing the weights towards 0, using the L_2 norm. But rather than encouraging
 1242 weights to be small, it might be better for the model to be **sparse**: it should assign weights
 1243 of exactly zero to most features, and only assign non-zero weights to features that are
 1244 clearly necessary. This idea can be formalized by the L_0 norm, $L_0 = \|\theta\|_0 = \sum_j \delta(\theta_j \neq 0)$,
 1245 which applies a constant penalty for each non-zero weight. This norm can be thought
 1246 of as a form of **feature selection**: optimizing the L_0 -regularized conditional likelihood is
 1247 equivalent to trading off the log-likelihood against the number of active features. Reduc-
 1248 ing the number of active features is desirable because the resulting model will be fast,
 1249 low-memory, and should generalize well, since irrelevant features will be pruned away.
 1250 Unfortunately, the L_0 norm is non-convex and non-differentiable. Optimization under L_0
 1251 regularization is **NP-hard**, meaning that it can be solved efficiently only if P=NP (Ge et al.,
 1252 2011).

1253 A useful alternative is the L_1 norm, which is equal to the sum of the absolute values
 1254 of the weights, $\|\theta\|_1 = \sum_j |\theta_j|$. The L_1 norm is convex, and can be used as an approxima-
 1255 tion to L_0 (Tibshirani, 1996). Conveniently, the L_1 norm also performs feature selection,
 1256 by driving many of the coefficients to zero; it is therefore known as a **sparsity inducing**
 1257 **regularizer**. The L_1 norm does not have a gradient at $\theta_j = 0$, so we must instead optimize
 1258 the L_1 -regularized objective using **subgradient** methods. The associated stochastic sub-
 1259 gradient descent algorithms are only somewhat more complex than conventional SGD;
 1260 Sra et al. (2012) survey approaches for estimation under L_1 and other regularizers.

1261 Gao et al. (2007) compare L_1 and L_2 regularization on a suite of NLP problems, finding
 1262 that L_1 regularization generally gives similar accuracy to L_2 regularization, but that L_1
 1263 regularization produces models that are between ten and fifty times smaller, because more
 1264 than 90% of the feature weights are set to zero.

1265 2.6.2 Other views of logistic regression

In binary classification, we can dispense with the feature function, and choose y based on
 the inner product of $\theta \cdot x$. The conditional probability $p_{Y|X}$ is obtained by passing this

inner product through a **logistic function**,

$$\sigma(a) \triangleq \frac{\exp(a)}{1 + \exp(a)} = (1 + \exp(-a))^{-1} \quad [2.76]$$

$$p(y | \mathbf{x}; \boldsymbol{\theta}) = \sigma(\boldsymbol{\theta} \cdot \mathbf{x}). \quad [2.77]$$

1266 This is the origin of the name **logistic regression**. Logistic regression can be viewed as
 1267 part of a larger family of **generalized linear models** (GLMs), in which various other “link
 1268 functions” convert between the inner product $\boldsymbol{\theta} \cdot \mathbf{x}$ and the parameter of a conditional
 1269 probability distribution.

1270 In the early NLP literature, logistic regression is frequently called **maximum entropy**
 1271 classification (Berger et al., 1996). This name refers to an alternative formulation, in
 1272 which the goal is to find the maximum entropy probability function that satisfies **moment-**
 1273 **matching** constraints. These constraints specify that the empirical counts of each feature
 1274 should match the expected counts under the induced probability distribution $p_{Y|X;\boldsymbol{\theta}}$.

$$\sum_{i=1}^N f_j(\mathbf{x}^{(i)}, y^{(i)}) = \sum_{i=1}^N \sum_{y \in \mathcal{Y}} p(y | \mathbf{x}^{(i)}; \boldsymbol{\theta}) f_j(\mathbf{x}^{(i)}, y), \quad \forall j \quad [2.78]$$

1275 The moment-matching constraint is satisfied exactly when the derivative of the condi-
 1276 tional log-likelihood function (Equation 2.64) is equal to zero. However, the constraint
 1277 can be met by many values of $\boldsymbol{\theta}$, so which should we choose?

1278 The **entropy** of the conditional probability distribution $p_{Y|X}$ is,

$$H(p_{Y|X}) = - \sum_{\mathbf{x} \in \mathcal{X}} p_X(\mathbf{x}) \sum_{y \in \mathcal{Y}} p_{Y|X}(y | \mathbf{x}) \log p_{Y|X}(y | \mathbf{x}), \quad [2.79]$$

1279 where \mathcal{X} is the set of all possible feature vectors, and $p_X(\mathbf{x})$ is the probability of observing
 1280 the base features \mathbf{x} . The distribution p_X is unknown, but it can be estimated by summing
 1281 over all the instances in the training set,

$$\tilde{H}(p_{Y|X}) = - \frac{1}{N} \sum_{i=1}^N \sum_{y \in \mathcal{Y}} p_{Y|X}(y | \mathbf{x}^{(i)}) \log p_{Y|X}(y | \mathbf{x}^{(i)}). \quad [2.80]$$

1282 If the entropy is large, the likelihood function is smooth across possible values of y ;
 1283 if it is small, the likelihood function is sharply peaked at some preferred value; in the
 1284 limiting case, the entropy is zero if $p(y | x) = 1$ for some y . The maximum-entropy cri-
 1285 terion chooses to make the weakest commitments possible, while satisfying the moment-
 1286 matching constraints from Equation 2.78. The solution to this constrained optimization
 1287 problem is identical to the maximum conditional likelihood (logistic-loss) formulation
 1288 that was presented in § 2.4.

1289 2.7 Summary of learning algorithms

1290 It is natural to ask which learning algorithm is best, but the answer depends on what
 1291 characteristics are important to the problem you are trying to solve.

1292 **Naïve Bayes** *Pros:* easy to implement; estimation is fast, requiring only a single pass over
 1293 the data; assigns probabilities to predicted labels; controls overfitting with smoothing
 1294 parameter. *Cons:* often has poor accuracy, especially with correlated features.

1295 **Perceptron** *Pros:* easy to implement; online; error-driven learning means that accuracy
 1296 is typically high, especially after averaging. *Cons:* not probabilistic; hard to know
 1297 when to stop learning; lack of margin can lead to overfitting.

1298 **Support vector machine** *Pros:* optimizes an error-based metric, usually resulting in high
 1299 accuracy; overfitting is controlled by a regularization parameter. *Cons:* not proba-
 1300 bilistic.

1301 **Logistic regression** *Pros:* error-driven and probabilistic; overfitting is controlled by a reg-
 1302 ularization parameter. *Cons:* batch learning requires black-box optimization; logistic
 1303 loss can “overtrain” on correctly labeled examples.

1304 One of the main distinctions is whether the learning algorithm offers a probability
 1305 over labels. This is useful in modular architectures, where the output of one classifier
 1306 is the input for some other system. In cases where probability is not necessary, the sup-
 1307 port vector machine is usually the right choice, since it is no more difficult to implement
 1308 than the perceptron, and is often more accurate. When probability is necessary, logistic
 1309 regression is usually more accurate than Naïve Bayes.

1310 Additional resources

1311 For more on classification, you can consult a textbook on machine learning (e.g., Mur-
 1312 phy, 2012), although the notation will differ slightly from what is typical in natural lan-
 1313 guage processing. Probabilistic methods are surveyed by Hastie et al. (2009), and Mohri
 1314 et al. (2012) emphasize theoretical considerations. Online learning is a rapidly moving
 1315 subfield of machine learning, and Bottou et al. (2016) describes progress through 2016.
 1316 Kummerfeld et al. (2015) empirically review several optimization algorithms for large-
 1317 margin learning. The python toolkit `scikit-learn` includes implementations of all of
 1318 the algorithms described in this chapter (Pedregosa et al., 2011).

1319 **Exercises**

- 1320 1. Let \mathbf{x} be a bag-of-words vector such that $\sum_{j=1}^V x_j = 1$. Verify that the multinomial
1321 probability $p_{\text{mult}}(\mathbf{x}; \phi)$, as defined in Equation 2.12, is identical to the probability of
1322 the same document under a categorical distribution, $p_{\text{cat}}(\mathbf{w}; \phi)$.
- 1323 2. Derive the maximum-likelihood estimate for the parameter μ in Naïve Bayes.
- 1324 3. As noted in the discussion of averaged perceptron in § 2.2.2, the computation of the
1325 running sum $\mathbf{m} \leftarrow \mathbf{m} + \boldsymbol{\theta}$ is unnecessarily expensive, requiring $K \times V$ operations.
1326 Give an alternative way to compute the averaged weights $\bar{\boldsymbol{\theta}}$, with complexity that is
1327 independent of V and linear in the sum of feature sizes $\sum_{i=1}^N |\mathbf{f}(\mathbf{x}^{(i)}, y^{(i)})|$.
- 1328 4. Consider a dataset that is comprised of two identical instances $\mathbf{x}^{(1)} = \mathbf{x}^{(2)}$ with
1329 distinct labels $y^{(1)} \neq y^{(2)}$. Assume all features are binary $x_j \in \{0, 1\}$ for all j .

1330 Now suppose that the averaged perceptron always chooses $i = 1$ when t is even,
1331 and $i = 2$ when t is odd, and that it will terminate under the following condition:

$$\epsilon \geq \max_j \left| \frac{1}{t} \sum_t \theta_j^{(t)} - \frac{1}{t-1} \sum_t \theta_j^{(t-1)} \right|. \quad [2.81]$$

1332 In words, the algorithm stops when the largest change in the averaged weights is
1333 less than or equal to ϵ . Compute the number of iterations before the averaged per-
1334 ceptron terminates.

- 1335 5. Suppose you have two labeled datasets D_1 and D_2 , with the same features and la-
1336 bels.
 - 1337 • Let $\boldsymbol{\theta}^{(1)}$ be the unregularized logistic regression (LR) coefficients from training
1338 on dataset D_1 .
 - 1339 • Let $\boldsymbol{\theta}^{(2)}$ be the unregularized LR coefficients (same model) from training on
1340 dataset D_2 .
 - 1341 • Let $\boldsymbol{\theta}^*$ be the unregularized LR coefficients from training on the combined
1342 dataset $D_1 \cup D_2$.

Under these conditions, prove that for any feature j ,

$$\begin{aligned} \theta_j^* &\geq \min(\theta_j^{(1)}, \theta_j^{(2)}) \\ \theta_j^* &\leq \max(\theta_j^{(1)}, \theta_j^{(2)}). \end{aligned}$$

1343

1344 **Chapter 3**

1345 **Nonlinear classification**

1346 Linear classification may seem like all we need for natural language processing. The bag-
1347 of-words representation is inherently high dimensional, and the number of features is
1348 often larger than the number of training instances. This means that it is usually possible
1349 to find a linear classifier that perfectly fits the training data. Moving to nonlinear classifi-
1350 cation may therefore only increase the risk of overfitting. For many tasks, **lexical features**
1351 (words) are meaningful in isolation, and can offer independent evidence about the in-
1352 stance label — unlike computer vision, where individual pixels are rarely informative,
1353 and must be evaluated holistically to make sense of an image. For these reasons, natu-
1354 ral language processing has historically focused on linear classification to a greater extent
1355 than other machine learning application domains.

1356 But in recent years, nonlinear classifiers have swept through natural language pro-
1357 cessing, and are now the default approach for many tasks (Manning, 2016). There are at
1358 least three reasons for this change.

- 1359 • There have been rapid advances in **deep learning**, a family of nonlinear meth-
1360 ods that learn complex functions of the input through multiple layers of computa-
1361 tion (Goodfellow et al., 2016).
- 1362 • Deep learning facilitates the incorporation of **word embeddings**, which are dense
1363 vector representations of words. Word embeddings can be learned from large amounts
1364 of unlabeled data, and enable generalization to words that do not appear in the an-
1365notated training data (word embeddings are discussed in detail in chapter 14).
- 1366 • A third reason for the rise of deep nonlinear learning algorithms is hardware. Many
1367 deep learning models can be implemented efficiently on graphics processing units
1368 (GPUs), offering substantial performance improvements over CPU-based comput-
1369 ing.

1370 This chapter focuses on **neural networks**, which are the dominant approach for non-

1371 linear classification in natural language processing today.¹ Historically, a few other non-
 1372 linear learning methods have been applied to language data:

- 1373 • **Kernel methods** are generalizations of the **nearest-neighbor** classification rule, which
 1374 classifies each instance by the label of the most similar example in the training
 1375 set (Hastie et al., 2009). The application of the **kernel support vector machine** to
 1376 information extraction is described in chapter 17.
- 1377 • **Decision trees** classify instances by checking a set of conditions. Scaling decision
 1378 trees to bag-of-words inputs is difficult, but decision trees have been successful in
 1379 problems such as coreference resolution (chapter 15), where more compact feature
 1380 sets can be constructed (Soon et al., 2001).
- 1381 • **Boosting** and related **ensemble methods** work by combining the predictions of sev-
 1382 eral “weak” classifiers, each of which may consider only a small subset of features.
 1383 Boosting has been successfully applied to text classification (Schapire and Singer,
 1384 2000) and syntactic analysis (Abney et al., 1999), and remains one of the most suc-
 1385 cessful methods on machine learning competition sites such as Kaggle (Chen and
 1386 Guestrin, 2016).

1387 3.1 Feedforward neural networks

1388 Consider the problem of building a classifier for movie reviews. The goal is to predict
 1389 a label $y \in \{\text{GOOD}, \text{BAD}, \text{OKAY}\}$ from a representation of the text of each document, x .
 1390 But what makes a good movie? The story, acting, cinematography, soundtrack, and so
 1391 on. Now suppose the training set contains labels for each of these additional features,
 1392 $z = [z_1, z_2, \dots, z_{K_z}]^\top$. With such information, we could build a two-step classifier:

- 1393 1. **Use the text x to predict the features z .** Specifically, train a logistic regression clas-
 1394 sifier to compute $p(z_k | x)$, for each $k \in \{1, 2, \dots, K_z\}$.
- 1395 2. **Use the features z to predict the label y .** Again, train a logistic regression classifier
 1396 to compute $p(y | z)$. On test data, z is unknown, so we use the probabilities $p(z | x)$
 1397 from the first layer as the features.

1398 This setup is shown in Figure 3.1, which describes the proposed classifier in a **compu-
 1399 tation graph**: the text features x are connected to the middle layer z , which in turn is
 1400 connected to the label y .

1401 Since each $z_k \in \{0, 1\}$, we can treat $p(z_k | x)$ as a binary classification problem, using
 1402 binary logistic regression:

$$\Pr(z_k = 1 | x; \Theta^{(x \rightarrow z)}) = \sigma(\boldsymbol{\theta}_k^{(x \rightarrow z)} \cdot x) = (1 + \exp(-\boldsymbol{\theta}_k^{(x \rightarrow z)} \cdot x))^{-1}, \quad [3.1]$$

¹I will use “deep learning” and “neural networks” interchangeably.

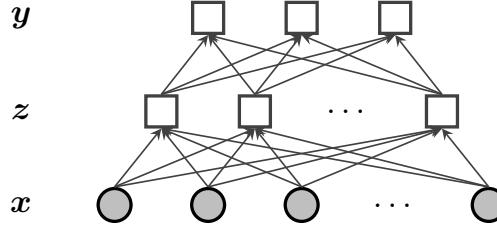


Figure 3.1: A feedforward neural network. Shaded circles indicate observed features, usually words; squares indicate nodes in the computation graph, which are computed from the information carried over the incoming arrows.

1403 where $\sigma(\cdot)$ is the **sigmoid** function (shown in Figure 3.2), and the matrix $\Theta^{(x \rightarrow z)} \in \mathbb{R}^{K_z \times V}$
 1404 is constructed by stacking the weight vectors for each z_k ,

$$\Theta^{(x \rightarrow z)} = [\theta_1^{(x \rightarrow z)}, \theta_2^{(x \rightarrow z)}, \dots, \theta_{K_z}^{(x \rightarrow z)}]^\top. \quad [3.2]$$

1405 We will assume that x contains a term with a constant value of 1, so that a corresponding
 1406 offset parameter is included in each $\theta_k^{(x \rightarrow z)}$.

1407 The output layer is computed by the multi-class logistic regression probability,

$$\Pr(y = j \mid z; \Theta^{(z \rightarrow y)}, b) = \frac{\exp(\theta_j^{(z \rightarrow y)} \cdot z + b_j)}{\sum_{j' \in \mathcal{Y}} \exp(\theta_{j'}^{(z \rightarrow y)} \cdot z + b_{j'})}, \quad [3.3]$$

1408 where b_j is an offset for label j , and the output weight matrix $\Theta^{(z \rightarrow y)} \in \mathbb{R}^{K_y \times K_z}$ is again
 1409 constructed by concatenation,

$$\Theta^{(z \rightarrow y)} = [\theta_1^{(z \rightarrow y)}, \theta_2^{(z \rightarrow y)}, \dots, \theta_{K_y}^{(z \rightarrow y)}]^\top. \quad [3.4]$$

1410 The vector of probabilities over each possible value of y is denoted,

$$p(y \mid z; \Theta^{(z \rightarrow y)}, b) = \text{SoftMax}(\Theta^{(z \rightarrow y)} z + b), \quad [3.5]$$

1411 where element j in the output of the **SoftMax** function is computed as in Equation 3.3.

We have now defined a multilayer classifier, which can be summarized as,

$$p(z \mid x; \Theta^{(x \rightarrow z)}) = \sigma(\Theta^{(x \rightarrow z)} x) \quad [3.6]$$

$$p(y \mid z; \Theta^{(z \rightarrow y)}, b) = \text{SoftMax}(\Theta^{(z \rightarrow y)} z + b), \quad [3.7]$$

1412 where $\sigma(\cdot)$ is now applied **elementwise** to the vector of inner products,

$$\sigma(\Theta^{(x \rightarrow z)} x) = [\sigma(\theta_1^{(x \rightarrow z)} \cdot x), \sigma(\theta_2^{(x \rightarrow z)} \cdot x), \dots, \sigma(\theta_{K_z}^{(x \rightarrow z)} \cdot x)]^\top. \quad [3.8]$$

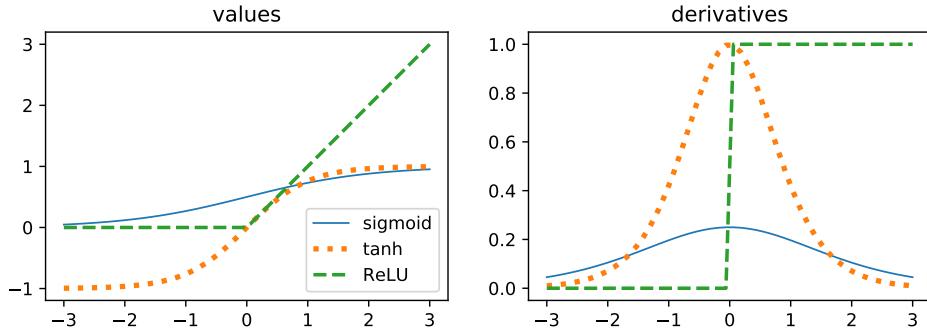


Figure 3.2: The sigmoid, tanh, and ReLU activation functions

Now suppose that the hidden features z are never observed, even in the training data. We can still construct the architecture in Figure 3.1. Instead of predicting y from a discrete vector of predicted values z , we use the probabilities $\sigma(\theta_k \cdot x)$. The resulting classifier is barely changed:

$$z = \sigma(\Theta^{(x \rightarrow z)} x) \quad [3.9]$$

$$p(y | x; \Theta^{(z \rightarrow y)}, b) = \text{SoftMax}(\Theta^{(z \rightarrow y)} z + b). \quad [3.10]$$

1413 This defines a classification model that predicts the label $y \in \mathcal{Y}$ from the base features x ,
1414 through a “hidden layer” z . This is a **feedforward neural network**.²

1415 3.2 Designing neural networks

1416 This feedforward neural network can be generalized in a number of ways.

1417 3.2.1 Activation functions

1418 If the hidden layer is viewed as a set of latent features, then the sigmoid function repre-
1419 presents the extent to which each of these features is “activated” by a given input. However,
1420 the hidden layer can be regarded more generally as a nonlinear transformation of the in-
1421 put. This opens the door to many other activation functions, some of which are shown in
1422 Figure 3.2. At the moment, the choice of activation functions is more art than science, but
1423 a few points can be made about the most popular varieties:

- 1424 • The range of the sigmoid function is $(0, 1)$. The bounded range ensures that a cas-
1425 cade of sigmoid functions will not “blow up” to a huge output, and this is impor-

²The architecture is sometimes called a **multilayer perceptron**, but this is misleading, because each layer is not a perceptron as defined in Algorithm 3.

tant for deep networks with several hidden layers. The derivative of the sigmoid is $\frac{\partial}{\partial a} \sigma(a) = \sigma(a)(1 - \sigma(a))$. This derivative becomes small at the extremes, which can make learning slow; this is called the **vanishing gradient** problem.

- The range of the **tanh activation function** is $(-1, 1)$: like the sigmoid, the range is bounded, but unlike the sigmoid, it includes negative values. The derivative is $\frac{\partial}{\partial a} \tanh(a) = 1 - \tanh(a)^2$, which is steeper than the logistic function near the origin (LeCun et al., 1998). The tanh function can also suffer from vanishing gradients at extreme values.
- The **rectified linear unit (ReLU)** is zero for negative inputs, and linear for positive inputs (Glorot et al., 2011),

$$\text{ReLU}(a) = \begin{cases} a, & a \geq 0 \\ 0, & \text{otherwise.} \end{cases} \quad [3.11]$$

The derivative is a step function, which is 1 if the input is positive, and zero otherwise. Once the activation is zero, the gradient is also zero. This can lead to the problem of **dead neurons**, where some ReLU nodes are zero for all inputs, throughout learning. A solution is the **leaky ReLU**, which has a small positive slope for negative inputs (Maas et al., 2013),

$$\text{Leaky-ReLU}(a) = \begin{cases} a, & a \geq 0 \\ .0001a, & \text{otherwise.} \end{cases} \quad [3.12]$$

Sigmoid and tanh are sometimes described as **squashing functions**, because they squash an unbounded input into a bounded range. Glorot and Bengio (2010) recommend against the use of the sigmoid activation in deep networks, because its mean value of $\frac{1}{2}$ can cause the next layer of the network to be saturated, with very small gradients on their own parameters. Several other activation functions are reviewed by Goodfellow et al. (2016), who recommend ReLU as the “default option.”

3.2.2 Network structure

Deep networks stack up several hidden layers, with each $z^{(d)}$ acting as the input to the next layer, $z^{(d+1)}$. As the total number of nodes in the network increases, so does its capacity to learn complex functions of the input. For a fixed number of nodes, an architectural decision is whether to emphasize width (large K_z at each layer) or depth (many layers). At present, this tradeoff is not well understood.³

³With even a single hidden layer, a neural network can approximate any continuous function on a closed and bounded subset of \mathbb{R}^N to an arbitrarily small non-zero error; see section 6.4.1 of Goodfellow et al. (2016) for a survey of these theoretical results. However, depending on the function to be approximated, the width

1453 It is also possible to “short circuit” a hidden layer, by propagating information directly
 1454 from the input to the next higher level of the network. This is the idea behind **residual net-**
 1455 **works**, which propagate information directly from the input to the subsequent layer (He
 1456 et al., 2016),

$$z = f(\Theta^{(x \rightarrow z)} \mathbf{x}) + \mathbf{x}, \quad [3.13]$$

where f is any nonlinearity, such as sigmoid or ReLU. A more complex architecture is the **highway network** (Srivastava et al., 2015; Kim et al., 2016), in which an addition **gate** controls an interpolation between $f(\Theta^{(x \rightarrow z)} \mathbf{x})$ and \mathbf{x} :

$$t = \sigma(\Theta^{(t)} \mathbf{x} + \mathbf{b}^{(t)}) \quad [3.14]$$

$$z = t \odot f(\Theta^{(x \rightarrow z)} \mathbf{x}) + (1 - t) \odot \mathbf{x}, \quad [3.15]$$

1457 where \odot refers to an elementwise vector product, and $\mathbf{1}$ is a column vector of ones. The
 1458 sigmoid function is applied elementwise to its input; recall that the output of this function
 1459 is restricted to the range $[0, 1]$. Gating is also used in the **long short-term memory (LSTM)**,
 1460 which is discussed in chapter 6. Residual and highway connections address a problem
 1461 with deep architectures: repeated application of a nonlinear activation function can make
 1462 it difficult to learn the parameters of the lower levels of the network, which are too distant
 1463 from the supervision signal.

1464 3.2.3 Outputs and loss functions

In the multi-class classification example, a softmax output produces probabilities over each possible label. This aligns with a negative **conditional log-likelihood**,

$$-\mathcal{L} = -\sum_{i=1}^N \log p(y^{(i)} | \mathbf{x}^{(i)}; \Theta). \quad [3.16]$$

1465 where $\Theta = \{\Theta^{(x \rightarrow z)}, \Theta^{(z \rightarrow y)}, \mathbf{b}\}$ is the entire set of parameters.

This loss can be written alternatively as follows:

$$\tilde{y}_j \triangleq \Pr(y = j | \mathbf{x}^{(i)}; \Theta) \quad [3.17]$$

$$-\mathcal{L} = -\sum_{i=1}^N e_{y^{(i)}} \cdot \log \tilde{y} \quad [3.18]$$

1466 where $e_{y^{(i)}}$ is a **one-hot vector** of zeros with a value of 1 at position $y^{(i)}$. The inner product
 1467 between $e_{y^{(i)}}$ and $\log \tilde{y}$ is also called the multinomial **cross-entropy**, and this terminology
 1468 is preferred in many neural networks papers and software packages.

of the hidden layer may need to be arbitrarily large. Furthermore, the fact that a network has the capacity to approximate any given function does not say anything about whether it is possible to *learn* the function using gradient-based optimization.

It is also possible to train neural networks from other objectives, such as a margin loss. In this case, it is not necessary to use softmax at the output layer: an affine transformation of the hidden layer is enough:

$$\Psi(y; \mathbf{x}^{(i)}, \Theta) = \theta_y^{(z \rightarrow y)} \cdot \mathbf{z} + b_y \quad [3.19]$$

$$\ell_{\text{MARGIN}}(\Theta; \mathbf{x}^{(i)}, y^{(i)}) = \max_{y \neq y^{(i)}} \left(1 + \Psi(y; \mathbf{x}^{(i)}, \Theta) - \Psi(y^{(i)}; \mathbf{x}^{(i)}, \Theta) \right)_+ \quad [3.20]$$

- 1469 In regression problems, the output is a scalar or vector (see § 4.1.2). For these problems, a
 1470 typical loss function is the squared error $(y - \hat{y})^2$ or squared norm $\|\mathbf{y} - \hat{\mathbf{y}}\|_2^2$.

1471 3.2.4 Inputs and lookup layers

1472 In text classification, the input layer \mathbf{x} can refer to a bag-of-words vector, where x_j is
 1473 the count of word j . The input to the hidden unit z_k is then $\sum_{j=1}^V \theta_{j,k}^{(x \rightarrow z)} x_j$, and word j is
 1474 represented by the vector $\theta_j^{(x \rightarrow z)}$. This vector is sometimes described as the **embedding** of
 1475 word j , and can be learned from unlabeled data, using techniques discussed in chapter 14.
 1476 The columns of $\Theta^{(x \rightarrow z)}$ are each K_z -dimensional word embeddings.

1477 Chapter 2 presented an alternative view of text documents, as a sequence of word
 1478 tokens, w_1, w_2, \dots, w_M . In a neural network, each word token w_m is represented with
 1479 a one-hot vector, $e_{w_m} \in \mathbb{R}^V$. The matrix-vector product $\Theta^{(x \rightarrow z)} e_{w_m}$ returns the embed-
 1480 ding of word w_m . The complete document can be represented by horizontally concatenating
 1481 these one-hot vectors, $\mathbf{W} = [e_{w_1}, e_{w_2}, \dots, e_{w_M}]$, and the bag-of-words representation can
 1482 be recovered from the matrix-vector product $\mathbf{W} \mathbf{1}$, which simply sums each row over the
 1483 tokens $m = \{1, 2, \dots, M\}$. The matrix product $\Theta^{(x \rightarrow z)} \mathbf{W}$ contains the horizontally con-
 1484 catenated embeddings of each word in the document, which will be useful as the starting
 1485 point for **convolutional neural networks** (see § 3.4). This is sometimes called a **lookup**
 1486 **layer**, because the first step is to lookup the embeddings for each word in the input text.

1487 3.3 Learning neural networks

The feedforward network in Figure 3.1 can now be written in a more general form,

$$\mathbf{z} \leftarrow f(\Theta^{(x \rightarrow z)} \mathbf{x}^{(i)}) \quad [3.21]$$

$$\tilde{\mathbf{y}} \leftarrow \text{SoftMax} \left(\Theta^{(z \rightarrow y)} \mathbf{z} + \mathbf{b} \right) \quad [3.22]$$

$$\ell^{(i)} \leftarrow -e_{y^{(i)}} \cdot \log \tilde{y}, \quad [3.23]$$

- 1488 where f is an elementwise activation function, such as σ or ReLU.

Let us now consider how to estimate the parameters $\Theta^{(x \rightarrow z)}$, $\Theta^{(z \rightarrow y)}$ and \mathbf{b} , using online gradient-based optimization. The simplest such algorithm is stochastic gradient descent (Algorithm 5). The relevant updates are,

$$\mathbf{b} \leftarrow \mathbf{b} - \eta^{(t)} \nabla_{\mathbf{b}} \ell^{(i)} \quad [3.24]$$

$$\boldsymbol{\theta}_k^{(z \rightarrow y)} \leftarrow \boldsymbol{\theta}_k^{(z \rightarrow y)} - \eta^{(t)} \nabla_{\boldsymbol{\theta}_k^{(z \rightarrow y)}} \ell^{(i)} \quad [3.25]$$

$$\boldsymbol{\theta}_k^{(x \rightarrow z)} \leftarrow \boldsymbol{\theta}_k^{(x \rightarrow z)} - \eta^{(t)} \nabla_{\boldsymbol{\theta}_k^{(x \rightarrow z)}} \ell^{(i)}, \quad [3.26]$$

where $\eta^{(t)}$ is the learning rate on iteration t , $\ell^{(i)}$ is the loss at instance (or minibatch) i , and $\boldsymbol{\theta}_k^{(x \rightarrow z)}$ is column k of the matrix $\Theta^{(x \rightarrow z)}$, and $\boldsymbol{\theta}_k^{(z \rightarrow y)}$ is column k of $\Theta^{(z \rightarrow y)}$.

The gradients of the negative log-likelihood on \mathbf{b} and $\boldsymbol{\theta}_k^{(z \rightarrow y)}$ are very similar to the gradients in logistic regression,

$$\nabla_{\boldsymbol{\theta}_k^{(z \rightarrow y)}} \ell^{(i)} = \left[\frac{\partial \ell^{(i)}}{\partial \theta_{k,1}^{(z \rightarrow y)}}, \frac{\partial \ell^{(i)}}{\partial \theta_{k,2}^{(z \rightarrow y)}}, \dots, \frac{\partial \ell^{(i)}}{\partial \theta_{k,K_y}^{(z \rightarrow y)}} \right]^\top \quad [3.27]$$

$$\frac{\partial \ell^{(i)}}{\partial \theta_{k,j}^{(z \rightarrow y)}} = - \frac{\partial}{\partial \theta_{k,j}^{(z \rightarrow y)}} \left(\boldsymbol{\theta}_{y^{(i)}}^{(z \rightarrow y)} \cdot \mathbf{z} - \log \sum_{y \in \mathcal{Y}} \exp \boldsymbol{\theta}_y^{(z \rightarrow y)} \cdot \mathbf{z} \right) \quad [3.28]$$

$$= \left(\Pr(y = j \mid \mathbf{z}; \Theta^{(z \rightarrow y)}, \mathbf{b}) - \delta(j = y^{(i)}) \right) z_k, \quad [3.29]$$

where $\delta(j = y^{(i)})$ is a function that returns one when $j = y^{(i)}$, and zero otherwise. The gradient $\nabla_{\mathbf{b}} \ell^{(i)}$ is similar to Equation 3.29.

The gradients on the input layer weights $\Theta^{(x \rightarrow z)}$ can be obtained by applying the chain rule of differentiation:

$$\frac{\partial \ell^{(i)}}{\partial \theta_{n,k}^{(x \rightarrow z)}} = \frac{\partial \ell^{(i)}}{\partial z_k} \frac{\partial z_k}{\partial \theta_{n,k}^{(x \rightarrow z)}} \quad [3.30]$$

$$= \frac{\partial \ell^{(i)}}{\partial z_k} \frac{\partial f(\boldsymbol{\theta}_k^{(x \rightarrow z)} \cdot \mathbf{x})}{\partial \theta_{n,k}^{(x \rightarrow z)}} \quad [3.31]$$

$$= \frac{\partial \ell^{(i)}}{\partial z_k} \times f'(\boldsymbol{\theta}_k^{(x \rightarrow z)} \cdot \mathbf{x}) \times x_n, \quad [3.32]$$

where $f'(\boldsymbol{\theta}_k^{(x \rightarrow z)} \cdot \mathbf{x})$ is the derivative of the activation function f , applied at the input

$\theta_k^{(x \rightarrow z)} \cdot \mathbf{x}$. For example, if f is the sigmoid function, then the derivative is,

$$\frac{\partial \ell^{(i)}}{\partial \theta_{n,k}^{(x \rightarrow z)}} = \frac{\partial \ell^{(i)}}{\partial z_k} \times \sigma(\theta_k^{(x \rightarrow z)} \cdot \mathbf{x}) \times (1 - \sigma(\theta_k^{(x \rightarrow z)} \cdot \mathbf{x})) \times x_n \quad [3.33]$$

$$= \frac{\partial \ell^{(i)}}{\partial z_k} \times z_k \times (1 - z_k) \times x_n. \quad [3.34]$$

1493 For intuition, consider each of the terms in the product.

- 1494 • If the negative log-likelihood $\ell^{(i)}$ does not depend much on z_k , $\frac{\partial \ell^{(i)}}{\partial z_k} \rightarrow 0$, then it
1495 doesn't matter how z_k is computed, and so $\frac{\partial \ell^{(i)}}{\partial \theta_{n,k}^{(x \rightarrow z)}} \rightarrow 0$.
- 1496 • If z_k is near 1 or 0, then the curve of the sigmoid function (Figure 3.2) is nearly flat,
1497 and changing the inputs will make little local difference. The term $z_k \times (1 - z_k)$ is
1498 maximized at $z_k = \frac{1}{2}$, where the slope of the sigmoid function is steepest.
- 1499 • If $x_n = 0$, then it does not matter how we set the weights $\theta_{n,k}^{(x \rightarrow z)}$, so $\frac{\partial \ell^{(i)}}{\partial \theta_{n,k}^{(x \rightarrow z)}} = 0$.

1500 3.3.1 Backpropagation

1501 In the equations above, the value $\frac{\partial \ell^{(i)}}{\partial z_k}$ is reused in the derivatives with respect to each
1502 $\theta_{n,k}^{(x \rightarrow z)}$. It should therefore be computed once, and then cached. Furthermore, we should
1503 only compute any derivative once we have already computed all of the necessary "inputs"
1504 demanded by the chain rule of differentiation. This combination of sequencing, caching,
1505 and differentiation is known as **backpropagation**. It can be generalized to any directed
1506 acyclic **computation graph**.

1507 A computation graph is a declarative representation of a computational process. At
1508 each node t , compute a value v_t by applying a function f_t to a (possibly empty) list of
1509 parent nodes, π_t . For example, in a feedforward network with one hidden layer, there are
1510 nodes for the input $\mathbf{x}^{(i)}$, the hidden layer \mathbf{z} , the predicted output $\tilde{\mathbf{y}}$, and the parameters
1511 $\{\Theta^{(x \rightarrow z)}, \Theta^{(z \rightarrow y)}, \mathbf{b}\}$. During training, there is also a node for the observed label $y^{(i)}$ and
1512 the loss $\ell^{(i)}$. Computation graphs have three main types of nodes:

1513 **Variables.** The variables include the *inputs* \mathbf{x} , the *hidden nodes* \mathbf{z} , the outputs \mathbf{y} , and the
1514 loss function. Inputs are variables that do not have parents. Backpropagation com-
1515putes the gradients with respect to all variables except the inputs, but does not up-
1516 date the variables during learning.

1517 **Parameters.** In a feedforward network, the parameters include the weights and offsets.
1518 Parameter nodes do not have parents, and they are updated during learning.

Algorithm 6 General backpropagation algorithm. In the computation graph G , every node contains a function f_t and a set of parent nodes π_t ; the inputs to the graph are $x^{(i)}$.

1519 **Objective.** The *objective* node is not the parent of any other node. Backpropagation begins
1520 by computing the gradient with respect to this node.

If the computation graph is a directed acyclic graph, then it is possible to order the nodes with a topological sort, so that if node t is a parent of node t' , then $t < t'$. This means that the values $\{v_t\}_{t=1}^T$ can be computed in a single forward pass. The topological sort is reversed when computing gradients: each gradient g_t is computed from the gradients of the children of t , implementing the chain rule of differentiation. The general backpropagation algorithm for computation graphs is shown in Algorithm 6, and illustrated in Figure 3.3.

While the gradients with respect to each parameter may be complex, they are composed of products of simple parts. For many networks, all gradients can be computed through **automatic differentiation**. This means that end users need only specify the feed-forward computation, and the gradients necessary for learning can be obtained automatically. There are many software libraries that perform automatic differentiation on computation graphs, such as Torch (Collobert et al., 2011), TensorFlow (Abadi et al., 2016), and DyNet (Neubig et al., 2017). One important distinction between these libraries is whether they support **dynamic computation graphs**, in which the structure of the computation graph varies across instances. Static computation graphs are compiled in advance, and can be applied to fixed-dimensional data, such as bag-of-words vectors. In many natural language processing problems, each input has a distinct structure, requiring a unique computation graph.

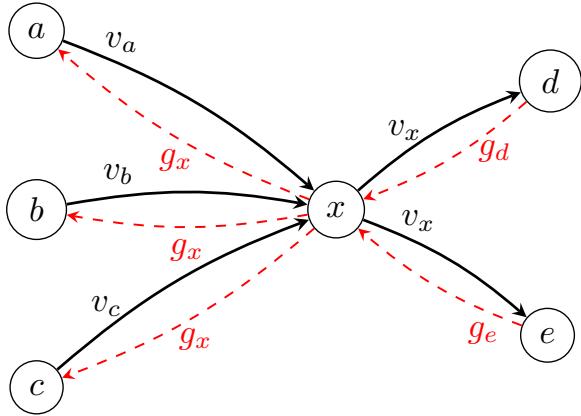


Figure 3.3: Backpropagation at a single node x in the computation graph. The values of the predecessors v_a, v_b, v_c are the inputs to x , which computes v_x , and passes it on to the successors d and e . The gradients at the successors g_d and g_e are passed back to x , where they are incorporated into the gradient g_x , which is then passed back to the predecessors a, b , and c .

1540 3.3.2 Regularization and dropout

1541 In linear classification, overfitting was addressed by augmenting the objective with a reg-
 1542 ularization term, $\lambda \|\theta\|_2^2$. This same approach can be applied to feedforward neural net-
 1543 works, penalizing each matrix of weights:

$$L = \sum_{i=1}^N \ell^{(i)} + \lambda_{z \rightarrow y} \|\Theta^{(z \rightarrow y)}\|_F^2 + \lambda_{x \rightarrow z} \|\Theta^{(x \rightarrow z)}\|_F^2, \quad [3.35]$$

1544 where $\|\Theta\|_F^2 = \sum_{i,j} \theta_{i,j}^2$ is the squared **Frobenius norm**, which generalizes the L_2 norm
 1545 to matrices. The bias parameters b are not regularized, as they do not contribute to the
 1546 sensitivity of the classifier to the inputs. In gradient-based optimization, the practical
 1547 effect of Frobenius norm regularization is that the weights “decay” towards zero at each
 1548 update, motivating the alternative name **weight decay**.

1549 Another approach to controlling model complexity is **dropout**, which involves ran-
 1550 domly setting some computation nodes to zero during training (Srivastava et al., 2014).
 1551 For example, in the feedforward network, on each training instance, with probability ρ we
 1552 set each input x_n and each hidden layer node z_k to zero. Srivastava et al. (2014) recom-
 1553 mend $\rho = 0.5$ for hidden units, and $\rho = 0.2$ for input units. Dropout is also incorporated
 1554 in the gradient computation, so if node z_k is dropped, then none of the weights $\theta_k^{(x \rightarrow z)}$ will
 1555 be updated for this instance. Dropout prevents the network from learning to depend too
 1556 much on any one feature or hidden node, and prevents **feature co-adaptation**, in which a

hidden unit is only useful in combination with one or more other hidden units. Dropout is a special case of **feature noising**, which can also involve adding Gaussian noise to inputs or hidden units (Holmstrom and Koistinen, 1992). Wager et al. (2013) show that dropout is approximately equivalent to “adaptive” L_2 regularization, with a separate regularization penalty for each feature.

3.3.3 *Learning theory

Chapter 2 emphasized the importance of **convexity** for learning: for convex objectives, the global optimum can be found efficiently. The negative log-likelihood and hinge loss are convex functions of the parameters of the output layer. However, the output of a feed-forward network is generally not a convex function of the parameters of the input layer, $\Theta^{(x \rightarrow z)}$. Feedforward networks can be viewed as function composition, where each layer is a function that is applied to the output of the previous layer. Convexity is generally not preserved in the composition of two convex functions — and furthermore, “squashing” activation functions like tanh and sigmoid are not convex.

The non-convexity of hidden layer neural networks can also be seen by permuting the elements of the hidden layer, from $z = [z_1, z_2, \dots, z_{K_z}]$ to $\tilde{z} = [z_{\pi(1)}, z_{\pi(2)}, \dots, z_{\pi(K_z)}]$. This corresponds to applying π to the rows of $\Theta^{(x \rightarrow z)}$ and the columns of $\Theta^{(z \rightarrow y)}$, resulting in permuted parameter matrices $\Theta_\pi^{(x \rightarrow z)}$ and $\Theta_\pi^{(z \rightarrow y)}$. As long as this permutation is applied consistently, the loss will be identical, $L(\Theta) = L(\Theta_\pi)$: it is *invariant* to this permutation. However, the loss of the linear combination $L(\alpha\Theta + (1 - \alpha)\Theta_\pi)$ will generally not be identical to the loss under Θ or its permutations. If $L(\Theta)$ is better than the loss at any points in the immediate vicinity, and if $L(\Theta) = L(\Theta_\pi)$, then the loss function does not satisfy the definition of convexity (see § 2.3). One of the exercises asks you to prove this more rigorously.

In practice, the existence of multiple optima is not necessarily problematic, if all such optima are permutations of the sort described in the previous paragraph. In contrast, “bad” local optima are better than their neighbors, but much worse than the global optimum. Fortunately, in large feedforward neural networks, most local optima are nearly as good as the global optimum (Choromanska et al., 2015), which helps to explain why backpropagation works in practice. More generally, a **critical point** is one at which the gradient is zero. Critical points may be local optima, but they may also be **saddle points**, which are local minima in some directions, but local *maxima* in other directions. For example, the equation $x_1^2 - x_2^2$ has a saddle point at $x = (0, 0)$.⁴ In large networks, the overwhelming majority of critical points are saddle points, rather than local minima or maxima (Dauphin et al., 2014). Saddle points can pose problems for gradient-based optimization, since learning will slow to a crawl as the gradient goes to zero. However, the noise introduced by

⁴Thanks to Rong Ge’s blogpost for this example, <http://www.offconvex.org/2016/03/22/saddlepoints/>

1593 stochastic gradient descent, and by feature noising techniques such as dropout, can help
 1594 online optimization to escape saddle points and find high-quality optima (Ge et al., 2015).
 1595 Other techniques address saddle points directly, using local reconstructions of the Hessian
 1596 matrix (Dauphin et al., 2014) or higher-order derivatives (Anandkumar and Ge, 2016).

1597 **3.3.4 Tricks**

1598 Getting neural networks to work effectively sometimes requires heuristic “tricks” (Bottou,
 1599 2012; Goodfellow et al., 2016; Goldberg, 2017b). This section presents some tricks that are
 1600 especially important.

Initialization Initialization is not especially important for linear classifiers, since convexity ensures that the global optimum can usually be found quickly. But for multilayer neural networks, it is helpful to have a good starting point. One reason is that if the magnitude of the initial weights is too large, a sigmoid or tanh nonlinearity will be saturated, leading to a small gradient, and slow learning. Large gradients are also problematic. Initialization can help avoid these problems, by ensuring that the variance over the initial gradients is constant and bounded throughout the network. For networks with tanh activation functions, this can be achieved by sampling the initial weights from the following uniform distribution (Glorot and Bengio, 2010),

$$\theta_{i,j} \sim U \left[-\frac{\sqrt{6}}{\sqrt{d_{\text{in}}(n) + d_{\text{out}}(n)}}, \frac{\sqrt{6}}{\sqrt{d_{\text{in}}(n) + d_{\text{out}}(n)}} \right], \quad [3.36]$$

[3.37]

1601 For the weights leading to a ReLU activation function, He et al. (2015) use similar argu-
 1602 mentation to justify sampling from a zero-mean Gaussian distribution,

$$\theta_{i,j} \sim N(0, \sqrt{2/d_{\text{in}}(n)}) \quad [3.38]$$

Rather than initializing the weights independently, it can be beneficial to initialize each layer jointly as an **orthonormal matrix**, ensuring that $\Theta^\top \Theta = \mathbb{I}$ (Saxe et al., 2014). Orthonormal matrices preserve the norm of the input, so that $\|\Theta x\| = \|x\|$, which prevents the gradients from exploding or vanishing. Orthogonality ensures that the hidden units are uncorrelated, so that they correspond to different features of the input. Orthonormal initialization can be performed by applying **singular value decomposition** to a matrix of

values sampled from a standard normal distribution:

$$a_{i,j} \sim N(0, 1) \quad [3.39]$$

$$\mathbf{A} = \{a_{i,j}\}_{i=1,j=1}^{d_{\text{in}}(j), d_{\text{out}}(j)} \quad [3.40]$$

$$\mathbf{U}, \mathbf{S}, \mathbf{V}^\top = \text{SVD}(\mathbf{A}) \quad [3.41]$$

$$\Theta^{(j)} \leftarrow \mathbf{U}. \quad [3.42]$$

1603 The matrix \mathbf{U} contains the **singular vectors** of \mathbf{A} , and is guaranteed to be orthonormal.
 1604 For more on singular value decomposition, see chapter 14.

1605 Even with careful initialization, there can still be significant variance in the final re-
 1606 sults. It can be useful to make multiple training runs, and select the one with the best
 1607 performance on a heldout development set.

1608 **Clipping and normalizing the gradients** As already discussed, the magnitude of the
 1609 gradient can pose problems for learning: too large, and learning can diverge, with suc-
 1610 cessive updates thrashing between increasingly extreme values; too small, and learning can
 1611 grind to a halt. Several heuristics have been proposed to address this issue.

1612 • In **gradient clipping** (Pascanu et al., 2013), an upper limit is placed on the norm of
 1613 the gradient, and the gradient is rescaled when this limit is exceeded,

$$\text{CLIP}(\hat{\mathbf{g}}) = \begin{cases} \mathbf{g} & \|\hat{\mathbf{g}}\| < \tau \\ \frac{\tau}{\|\mathbf{g}\|} \mathbf{g} & \text{otherwise.} \end{cases} \quad [3.43]$$

1614 • In **batch normalization** (Ioffe and Szegedy, 2015), the inputs to each computation
 1615 node are recentered by their mean and variance across all of the instances in the
 minibatch \mathcal{B} (see § 2.5.2). For example, in a feedforward network with one hidden
 layer, batch normalization would transform the inputs to the hidden layer as follows:

$$\boldsymbol{\mu}^{(\mathcal{B})} = \frac{1}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \mathbf{x}^{(i)} \quad [3.44]$$

$$\mathbf{s}^{(\mathcal{B})} = \frac{1}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} (\mathbf{x}^{(i)} - \boldsymbol{\mu}^{(\mathcal{B})})^2 \quad [3.45]$$

$$\bar{\mathbf{x}}^{(i)} = (\mathbf{x}^{(i)} - \boldsymbol{\mu}^{(\mathcal{B})}) / \sqrt{\mathbf{s}^{(\mathcal{B})}}. \quad [3.46]$$

1614 Empirically, this speeds convergence of deep architectures. One explanation is that
 1615 it helps to correct for changes in the distribution of activations during training.

- In **layer normalization** (Ba et al., 2016), the inputs to each nonlinear activation function are recentered across the layer:

$$\mathbf{a} = \Theta^{(x \rightarrow z)} \mathbf{x} \quad [3.47]$$

$$\mu = \frac{1}{K_z} \sum_{k=1}^{K_z} a_k \quad [3.48]$$

$$s = \frac{1}{K_z} \sum_{k=1}^{K_z} (a_k - \mu)^2 \quad [3.49]$$

$$z = (\mathbf{a} - \mu) / \sqrt{s}. \quad [3.50]$$

1616 Layer normalization has similar motivations to batch normalization, but it can be
 1617 applied across a wider range of architectures and training conditions.

Online optimization The trend towards deep learning has spawned a cottage industry of **online optimization** algorithms, which attempt to improve on stochastic gradient descent. **AdaGrad** was reviewed in § 2.5.2; its main innovation is to set adaptive learning rates for each parameter by storing the sum of squared gradients. Rather than using the sum over the entire training history, we can keep a running estimate,

$$v_j^{(t)} = \beta v_j^{(t-1)} + (1 - \beta) g_{t,j}^2, \quad [3.51]$$

1618 where $g_{t,j}$ is the gradient with respect to parameter j at time t , and $\beta \in [0, 1]$. This term
 1619 places more emphasis on recent gradients, and is employed in the **AdaDelta** (Zeiler, 2012)
 1620 and **Adam** (Kingma and Ba, 2014) optimizers. Online optimization and its theoretical
 1621 background are reviewed by Bottou et al. (2016). **Early stopping**, mentioned in § 2.2.2,
 1622 can help to avoid overfitting, by terminating training after reaching a plateau in the per-
 1623 formance on a heldout validation set.

1624 3.4 Convolutional neural networks

1625 A basic weakness of the bag-of-words model is its inability to account for the ways in
 1626 which words combine to create meaning, including even simple reversals such as *not*
 1627 *pleasant, hardly a generous offer*, and *I wouldn't mind missing the flight*. Similarly, computer
 1628 vision faces the challenge of identifying the semantics of images from pixel features that
 1629 are uninformative in isolation. An earlier generation of computer vision research fo-
 1630 cused on designing *filters* to aggregate local pixel-level features into more meaningful
 1631 representations, such as edges and corners (e.g., Canny, 1987). Similarly, earlier NLP re-
 1632 search attempted to capture multiword linguistic phenomena by hand-designed lexical
 1633 patterns (Hobbs et al., 1997). In both cases, the output of the filters and patterns could

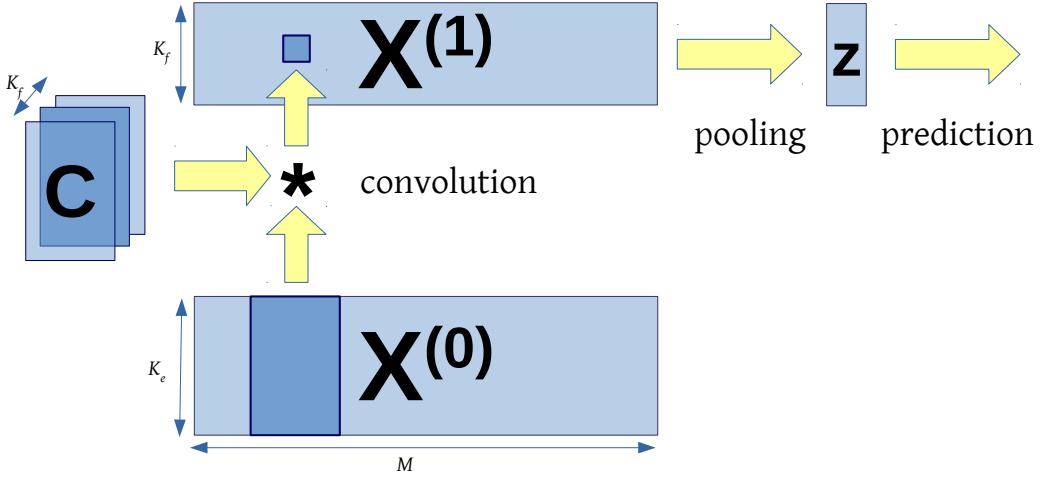


Figure 3.4: A convolutional neural network for text classification

then act as base features in a linear classifier. But rather than designing these feature extractors by hand, a better approach is to learn them, using the magic of backpropagation. This is the idea behind **convolutional neural networks**.

Following § 3.2.4, define the base layer of a neural network as,

$$\mathbf{X}^{(0)} = \Theta^{(x \rightarrow z)}[\mathbf{e}_{w_1}, \mathbf{e}_{w_2}, \dots, \mathbf{e}_{w_M}], \quad [3.52]$$

where \mathbf{e}_{w_m} is a column vector of zeros, with a 1 at position w_m . The base layer has dimension $\mathbf{X}^{(0)} \in \mathbb{R}^{K_e \times M}$, where K_e is the size of the word embeddings. To merge information across adjacent words, we *convolve* $\mathbf{X}^{(0)}$ with a set of filter matrices $\mathbf{C}^{(k)} \in \mathbb{R}^{K_e \times h}$. Convolution is indicated by the symbol $*$, and is defined,

$$\mathbf{X}^{(1)} = f(\mathbf{b} + \mathbf{C} * \mathbf{X}^{(0)}) \implies x_{k,m}^{(1)} = f \left(b_k + \sum_{k'=1}^{K_e} \sum_{n=1}^h c_{k',n}^{(k)} \times x_{k',m+n-1}^{(0)} \right), \quad [3.53]$$

where f is an activation function such as tanh or ReLU, and \mathbf{b} is a vector of offsets. The convolution operation slides the matrix $\mathbf{C}^{(k)}$ across the columns of $\mathbf{X}^{(0)}$; at each position m , compute the elementwise product $\mathbf{C}^{(k)} \odot \mathbf{X}_{m:m+h-1}^{(0)}$, and take the sum.

A simple filter might compute a weighted average over nearby words,

$$\mathbf{C}^{(k)} = \begin{bmatrix} 0.5 & 1 & 0.5 \\ 0.5 & 1 & 0.5 \\ \dots & \dots & \dots \\ 0.5 & 1 & 0.5 \end{bmatrix}, \quad [3.54]$$

1642 thereby representing trigram units like *not so unpleasant*. In **one-dimensional convolution**,
 1643 each filter matrix $\mathbf{C}^{(k)}$ is constrained to have non-zero values only at row k (Kalchbrenner et al., 2014).

1645 To deal with the beginning and end of the input, the base matrix $\mathbf{X}^{(0)}$ may be padded
 1646 with h column vectors of zeros at the beginning and end; this is known as **wide convolution**. If padding is not applied, then the output from each layer will be $h - 1$ units smaller
 1647 than the input; this is known as **narrow convolution**. The filter matrices need not have
 1648 identical filter widths, so more generally we could write h_k to indicate width of filter
 1649 $\mathbf{C}^{(k)}$. As suggested by the notation $\mathbf{X}^{(0)}$, multiple layers of convolution may be applied,
 1650 so that $\mathbf{X}^{(d)}$ is the input to $\mathbf{X}^{(d+1)}$.

After D convolutional layers, we obtain a matrix representation of the document $\mathbf{X}^{(D)} \in \mathbb{R}^{K_z \times M}$. If the instances have variable lengths, it is necessary to aggregate over all M word positions to obtain a fixed-length representation. This can be done by a **pooling** operation, such as max-pooling (Collobert et al., 2011) or average-pooling,

$$\mathbf{z} = \text{MaxPool}(\mathbf{X}^{(D)}) \implies z_k = \max(x_{k,1}^{(D)}, x_{k,2}^{(D)}, \dots, x_{k,M}^{(D)}) \quad [3.55]$$

$$\mathbf{z} = \text{AvgPool}(\mathbf{X}^{(D)}) \implies z_k = \frac{1}{M} \sum_{m=1}^M x_{k,m}^{(D)}. \quad [3.56]$$

1652 The vector \mathbf{z} can now act as a layer in a feedforward network, culminating in a prediction
 1653 \hat{y} and a loss $\ell^{(i)}$. The setup is shown in Figure 3.4.

Just as in feedforward networks, the parameters $(\mathbf{C}^{(k)}, \mathbf{b}, \Theta)$ can be learned by backpropagating from the classification loss. This requires backpropagating through the max-pooling operation, which is a discontinuous function of the input. But because we need only a local gradient, backpropagation flows only through the argmax m :

$$\frac{\partial z_k}{\partial x_{k,m}^{(D)}} = \begin{cases} 1, & x_{k,m}^{(D)} = \max(x_{k,1}^{(D)}, x_{k,2}^{(D)}, \dots, x_{k,M}^{(D)}) \\ 0, & \text{otherwise.} \end{cases} \quad [3.57]$$

1654 The computer vision literature has produced a huge variety of convolutional architectures,
 1655 and many of these bells and whistles can be applied to text data. One avenue for
 1656 improvement is more complex pooling operations, such as k -max pooling (Kalchbrenner
 1657 et al., 2014), which returns a matrix of the k largest values for each filter. Another innovation
 1658 is the use of **dilated convolution** to build multiscale representations (Yu and Koltun,
 1659 2016). At each layer, the convolutional operator applied in *strides*, skipping ahead by s
 1660 steps after each feature. As we move up the hierarchy, each layer is s times smaller than
 1661 the layer below it, effectively summarizing the input. This idea is shown in Figure 3.5.
 1662 Multi-layer convolutional networks can also be augmented with “shortcut” connections,
 1663 as in the ResNet model from § 3.2.2 (Johnson and Zhang, 2017).

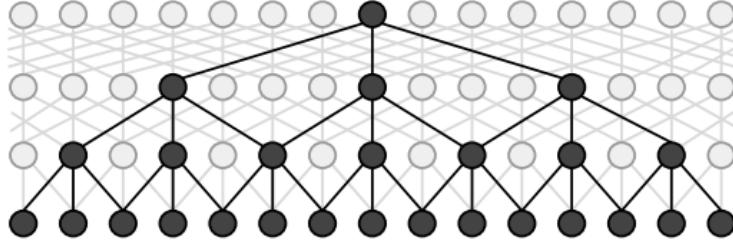


Figure 3.5: A dilated convolutional neural network captures progressively larger context through recursive application of the convolutional operator (Strubell et al., 2017) [todo: permission]

1664 Additional resources

1665 The deep learning textbook by Goodfellow et al. (2016) covers many of the topics in this
 1666 chapter in more detail. For a comprehensive review of neural networks in natural lan-
 1667 guage processing, see (Goldberg, 2017b). A seminal work on deep learning in natural
 1668 language processing is the aggressively titled “Natural Language Processing (Almost)
 1669 from Scratch”, which uses convolutional neural networks to perform a range of language
 1670 processing tasks (Collobert et al., 2011). This chapter focuses on feedforward and con-
 1671 volutional neural networks, but recurrent neural networks are one of the most important
 1672 deep learning architectures for natural language processing. They are covered extensively
 1673 in chapters 6 and 7.

1674 The role of deep learning in natural language processing research has caused angst
 1675 in some parts of the natural language processing research community (e.g., Goldberg,
 1676 2017a), especially as some of the more zealous deep learning advocates have argued that
 1677 end-to-end learning from “raw” text can eliminate the need for linguistic constructs such
 1678 as sentences, phrases, and even words (Zhang et al., 2015, originally titled *Text understand-
 1679 ing from scratch*). These developments were surveyed by Manning (2016).

1680 Exercises

- 1681 1. Prove that the softmax and sigmoid functions are equivalent when the number of
 1682 possible labels is two. Specifically, for any $\Theta^{(z \rightarrow y)}$ (omitting the offset b for sim-
 1683 plicity), show how to construct a vector of weights θ such that,

$$\text{SoftMax}(\Theta^{(z \rightarrow y)} z)[0] = \sigma(\theta \cdot z). \quad [3.58]$$

- 1684 2. Design a feedforward network to compute the XOR function:

$$f(x_1, x_2) = \begin{cases} -1, & x_1 = 1, x_2 = 1 \\ 1, & x_1 = 1, x_2 = 0 \\ 1, & x_1 = 0, x_2 = 1 \\ -1, & x_1 = 0, x_2 = 0 \end{cases}. \quad [3.59]$$

1685 Your network should have a single output node which uses the Sign activation func-
 1686 tion. Use a single hidden layer, with ReLU activation functions. Describe all weights
 1687 and offsets.

- 1688 3. Consider the same network as above (with ReLU activations for the hidden layer),
 1689 with an arbitrary differentiable loss function $\ell(y^{(i)}, \tilde{y})$, where \tilde{y} is the activation of
 1690 the output node. Suppose all weights and offsets are initialized to zero. Prove that
 1691 gradient-based optimization cannot learn the desired function from this initializa-
 1692 tion.
- 1693 4. The simplest solution to the previous problem relies on the use of the ReLU activa-
 1694 tion function at the hidden layer. Now consider a network with arbitrary activations
 1695 on the hidden layer. Show that if the initial weights are any uniform constant, then
 1696 it is not possible to learn the desired function.
- 1697 5. Consider a network in which: the base features are all binary, $\mathbf{x} \in \{0, 1\}^M$; the
 1698 hidden layer activation function is sigmoid, $z_k = \sigma(\boldsymbol{\theta}_k \cdot \mathbf{x})$; and the initial weights
 1699 are sampled independently from a standard normal distribution, $\theta_{j,k} \sim N(0, 1)$.

- 1700 • Show how the probability of a small initial gradient on any weight, $\frac{\partial z_k}{\partial \theta_{j,k}} < \alpha$,
 1701 depends on the size of the input M . **Hint:** use the lower bound,

$$\Pr(\sigma(\boldsymbol{\theta}_k \cdot \mathbf{x}) \times (1 - \sigma(\boldsymbol{\theta}_k \cdot \mathbf{x})) < \alpha) \geq 2 \Pr(\sigma(\boldsymbol{\theta}_k \cdot \mathbf{x}) < \alpha), \quad [3.60]$$

1702 and relate this probability to the variance $V[\boldsymbol{\theta}_k \cdot \mathbf{x}]$.

- 1703 • Design an alternative initialization that removes this dependence.

- 1704 6. Suppose that the parameters $\Theta = \{\Theta^{(x \rightarrow z)}, \Theta(z \rightarrow y), \mathbf{b}\}$ are a local optimum of a
 1705 feedforward network in the following sense: there exists some $\epsilon > 0$ such that,

$$\begin{aligned} & \left(\|\tilde{\Theta}^{(x \rightarrow z)} - \Theta^{(x \rightarrow z)}\|_F^2 + \|\tilde{\Theta}^{(z \rightarrow y)} - \Theta^{(z \rightarrow y)}\|_F^2 + \|\tilde{\mathbf{b}} - \mathbf{b}\|_2^2 < \epsilon \right) \\ & \Rightarrow \left(L(\tilde{\Theta}) > L(\Theta) \right) \end{aligned} \quad [3.61]$$

1706 Define the function π as a permutation on the hidden units, as described in § 3.3.3,
 1707 so that for any Θ , $L(\Theta) = L(\Theta_\pi)$. Prove that if a feedforward network has a local
 optimum in the sense of Equation 3.61, then its loss is not a convex function of the
 parameters Θ , using the definition of convexity from § 2.3

- 1708 7. Consider a network with a single hidden layer, and a single output,

$$y = \theta^{(z \rightarrow y)} \cdot g(\Theta^{(x \rightarrow z)} x). \quad [3.62]$$

1709 Assume that g is the ReLU function. Prove that for any matrix of weights $\Theta^{(x \rightarrow z)}$, it
1710 is permissible to rescale each row to have a norm of one, because an identical output
1711 can be obtained by finding a corresponding rescaling of $\theta^{(z \rightarrow y)}$.

¹⁷¹² Chapter 4

¹⁷¹³ Linguistic applications of ¹⁷¹⁴ classification

¹⁷¹⁵ Having learned some techniques for classification, this chapter shifts the focus from mathematics to linguistic applications. Later in the chapter, we will consider the design decisions involved in text classification, as well as evaluation practices.

¹⁷¹⁸ 4.1 Sentiment and opinion analysis

¹⁷¹⁹ A popular application of text classification is to automatically determine the **sentiment**
¹⁷²⁰ or **opinion polarity** of documents such as product reviews and social media posts. For
¹⁷²¹ example, marketers are interested to know how people respond to advertisements, ser-
¹⁷²² vices, and products (Hu and Liu, 2004); social scientists are interested in how emotions
¹⁷²³ are affected by phenomena such as the weather (Hannak et al., 2012), and how both opin-
¹⁷²⁴ ions and emotions spread over social networks (Coviello et al., 2014; Miller et al., 2011).
¹⁷²⁵ In the field of **digital humanities**, literary scholars track plot structures through the flow
¹⁷²⁶ of sentiment across a novel (Jockers, 2015).¹

¹⁷²⁷ Sentiment analysis can be framed as a direct application of document classification,
¹⁷²⁸ assuming reliable labels can be obtained. In the simplest case, sentiment analysis is a
¹⁷²⁹ two or three-class problem, with sentiments of POSITIVE, NEGATIVE, and possibly NEU-
¹⁷³⁰ TRAL. Such annotations could be annotated by hand, or obtained automatically through
¹⁷³¹ a variety of means:

- ¹⁷³² • Tweets containing happy emoticons can be marked as positive, sad emoticons as
¹⁷³³ negative (Read, 2005; Pak and Paroubek, 2010).

¹Comprehensive surveys on sentiment analysis and related problems are offered by Pang and Lee (2008) and Liu (2015).

- 1734 • Reviews with four or more stars can be marked as positive, two or fewer stars as
1735 negative (Pang et al., 2002).
- 1736 • Statements from politicians who are voting for a given bill are marked as positive
1737 (towards that bill); statements from politicians voting against the bill are marked as
1738 negative (Thomas et al., 2006).

1739 The bag-of-words model is a good fit for sentiment analysis at the document level: if
1740 the document is long enough, we would expect the words associated with its true senti-
1741 ment to overwhelm the others. Indeed, **lexicon-based sentiment analysis** avoids machine
1742 learning altogether, and classifies documents by counting words against positive and neg-
1743 ative sentiment word lists (Taboada et al., 2011).

1744 Lexicon-based classification is less effective for short documents, such as single-sentence
1745 reviews or social media posts. In these documents, linguistic issues like **negation** and **ir-**
1746 **realis** (Polanyi and Zaenen, 2006) — events that are hypothetical or otherwise non-factual
1747 — can make bag-of-words classification ineffective. Consider the following examples:

- 1748 (4.1) That's not bad for the first day.
- 1749 (4.2) This is not the worst thing that can happen.
- 1750 (4.3) It would be nice if you acted like you understood.
- 1751 (4.4) There is no reason at all to believe that the polluters are suddenly going to be-
1752 come reasonable. (Wilson et al., 2005)
- 1753 (4.5) This film should be brilliant. The actors are first grade. Stallone plays a happy,
1754 wonderful man. His sweet wife is beautiful and adores him. He has a fascinat-
1755 ing gift for living life fully. It sounds like a great plot, **however**, the film is a
1756 failure. (Pang et al., 2002)

1757 A minimal solution is to move from a bag-of-words model to a bag-of-**bigrams** model,
1758 where each base feature is a pair of adjacent words, e.g.,

$$(that's, not), (not, bad), (bad, for), \dots \quad [4.1]$$

1759 Bigrams can handle relatively straightforward cases, such as when an adjective is immedi-
1760 ately negated; trigrams would be required to extend to larger contexts (e.g., *not the worst*).
1761 But this approach will not scale to more complex examples like (4.4) and (4.5). More
1762 sophisticated solutions try to account for the syntactic structure of the sentence (Wilson
1763 et al., 2005; Socher et al., 2013), or apply more complex classifiers such as **convolutional**
1764 **neural networks** (Kim, 2014), which are described in chapter 3.

1765 **4.1.1 Related problems**

1766 **Subjectivity** Closely related to sentiment analysis is **subjectivity detection**, which re-
1767 quires identifying the parts of a text that express subjective opinions, as well as other non-
1768 factual content such as speculation and hypotheticals (Riloff and Wiebe, 2003). This can be
1769 done by treating each sentence as a separate document, and then applying a bag-of-words
1770 classifier: indeed, Pang and Lee (2004) do exactly this, using a training set consisting of
1771 (mostly) subjective sentences gathered from movie reviews, and (mostly) objective sen-
1772 tences gathered from plot descriptions. They augment this bag-of-words model with a
1773 graph-based algorithm that encourages nearby sentences to have the same subjectivity
1774 label.

1775 **Stance classification** In debates, each participant takes a side: for example, advocating
1776 for or against proposals like adopting a vegetarian lifestyle or mandating free college ed-
1777 ucation. The problem of stance classification is to identify the author’s position from the
1778 text of the argument. In some cases, there is training data available for each position,
1779 so that standard document classification techniques can be employed. In other cases, it
1780 suffices to classify each document as whether it is in support or opposition of the argu-
1781 ment advanced by a previous document (Anand et al., 2011). In the most challenging
1782 case, there is no labeled data for any of the stances, so the only possibility is group docu-
1783 ments that advocate the same position (Somasundaran and Wiebe, 2009). This is a form
1784 of **unsupervised learning**, discussed in chapter 5.

1785 **Targeted sentiment analysis** The expression of sentiment is often more nuanced than a
1786 simple binary label. Consider the following examples:

1787 (4.6) The vodka was good, but the meat was rotten.

1788 (4.7) Go to Heaven for the climate, Hell for the company. —Mark Twain

1789 These statements display a mixed overall sentiment: positive towards some entities (e.g.,
1790 *the vodka*), negative towards others (e.g., *the meat*). **Targeted sentiment analysis** seeks to
1791 identify the writer’s sentiment towards specific entities (Jiang et al., 2011). This requires
1792 identifying the entities in the text and linking them to specific sentiment words — much
1793 more than we can do with the classification-based approaches discussed thus far. For
1794 example, Kim and Hovy (2006) analyze sentence-internal structure to determine the topic
1795 of each sentiment expression.

1796 **Aspect-based opinion mining** seeks to identify the sentiment of the author of a review
1797 towards predefined aspects such as PRICE and SERVICE, or, in the case of (4.7), CLIMATE
1798 and COMPANY (Hu and Liu, 2004). If the aspects are not defined in advance, it may again
1799 be necessary to employ **unsupervised learning** methods to identify them (e.g., Branavan
1800 et al., 2009).

1801 **Emotion classification** While sentiment analysis is framed in terms of positive and neg-
 1802 ative categories, psychologists generally regard **emotion** as more multifaceted. For ex-
 1803 ample, Ekman (1992) argues that there are six basic emotions — happiness, surprise, fear,
 1804 sadness, anger, and contempt — and that they are universal across human cultures. Alm
 1805 et al. (2005) build a linear classifier for recognizing the emotions expressed in children’s
 1806 stories. The ultimate goal of this work was to improve text-to-speech synthesis, so that
 1807 stories could be read with intonation that reflected the emotional content. They used bag-
 1808 of-words features, as well as features capturing the story type (e.g., jokes, folktales), and
 1809 structural features that reflect the position of each sentence in the story. The task is diffi-
 1810 cult: even human annotators frequently disagreed with each other, and the best classifiers
 1811 achieved accuracy between 60-70%.

1812 4.1.2 Alternative approaches to sentiment analysis

1813 **Regression** A more challenging version of sentiment analysis is to determine not just
 1814 the class of a document, but its rating on a numerical scale (Pang and Lee, 2005). If the
 1815 scale is continuous, it is most natural to apply **regression**, identifying a set of weights θ
 1816 that minimize the squared error of a predictor $\hat{y} = \theta \cdot x + b$, where b is an offset. This
 1817 approach is called **linear regression**, and sometimes **least squares**, because the regression
 1818 coefficients θ are determined by minimizing the squared error, $(y - \hat{y})^2$. If the weights are
 1819 regularized using a penalty $\lambda \|\theta\|_2^2$, then it is **ridge regression**. Unlike logistic regression,
 1820 both linear regression and ridge regression can be solved in closed form as a system of
 1821 linear equations.

1822 **Ordinal ranking** In many problems, the labels are ordered but discrete: for example,
 1823 product reviews are often integers on a scale of 1 – 5, and grades are on a scale of A – F.
 1824 Such problems can be solved by discretizing the score $\theta \cdot x$ into “ranks”,

$$\hat{y} = \underset{r: \theta \cdot x \geq b_r}{\operatorname{argmin}} r, \quad [4.2]$$

1825 where $\mathbf{b} = [b_1 = -\infty, b_2, b_3, \dots, b_K]$ is a vector of boundaries. It is possible to learn the
 1826 weights and boundaries simultaneously, using a perceptron-like algorithm (Crammer and
 1827 Singer, 2001).

1828 **Lexicon-based classification** Sentiment analysis is one of the only NLP tasks where
 1829 hand-crafted feature weights are still widely employed. In **lexicon-based classification** (Taboada
 1830 et al., 2011), the user creates a list of words for each label, and then classifies each docu-
 1831 ment based on how many of the words from each list are present. In our linear classifica-
 1832 tion framework, this is equivalent to choosing the following weights:

$$\theta_{y,j} = \begin{cases} 1, & j \in \mathcal{L}_y \\ 0, & \text{otherwise,} \end{cases} \quad [4.3]$$

1833 where \mathcal{L}_y is the lexicon for label y . Compared to the machine learning classifiers discussed
 1834 in the previous chapters, lexicon-based classification may seem primitive. However, su-
 1835 pervised machine learning relies on large annotated datasets, which are time-consuming
 1836 and expensive to produce. If the goal is to distinguish two or more categories in a new
 1837 domain, it may be simpler to start by writing down a list of words for each category.

1838 An early lexicon was the *General Inquirer* (Stone, 1966). Today, popular sentiment lex-
 1839 cons include sentiwordnet (Esuli and Sebastiani, 2006) and an evolving set of lexicons
 1840 from Liu (2015). For emotions and more fine-grained analysis, *Linguistic Inquiry and Word*
 1841 *Count* (LIWC) provides a set of lexicons (Tausczik and Pennebaker, 2010). The MPQA lex-
 1842 icon indicates the polarity (positive or negative) of 8221 terms, as well as whether they are
 1843 strongly or weakly subjective (Wiebe et al., 2005). A comprehensive comparison of senti-
 1844 ment lexicons is offered by Ribeiro et al. (2016). Given an initial **seed lexicon**, it is possible
 1845 to automatically expand the lexicon by looking for words that frequently co-occur with
 1846 words in the seed set (Hatzivassiloglou and McKeown, 1997; Qiu et al., 2011).

1847 4.2 Word sense disambiguation

1848 Consider the the following headlines:

- 1849 (4.8) Iraqi head seeks arms
- 1850 (4.9) Prostitutes appeal to Pope
- 1851 (4.10) Drunk gets nine years in violin case²

1852 These headlines are ambiguous because they contain words that have multiple mean-
 1853 ings, or **senses**. Word sense disambiguation is the problem of identifying the intended
 1854 sense of each word token in a document. Word sense disambiguation is part of a larger
 1855 field of research called **lexical semantics**, which is concerned with meanings of the words.

1856 At a basic level, the problem of word sense disambiguation is to identify the correct
 1857 sense for each word token in a document. Part-of-speech ambiguity (e.g., noun versus
 1858 verb) is usually considered to be a different problem, to be solved at an earlier stage.
 1859 From a linguistic perspective, senses are not properties of words, but of **lemmas**, which
 1860 are canonical forms that stand in for a set of inflected words. For example, *arm*/N is a
 1861 lemma that includes the inflected form *arms*/N — the /N indicates that it we are refer-
 1862 ring to the noun, and not its **homonym** *arm*/V, which is another lemma that includes
 1863 the inflected verbs (*arm*/V, *arms*/V, *armed*/V, *arming*/V). Therefore, word sense disam-
 1864 biguation requires first identifying the correct part-of-speech and lemma for each token,

²These examples, and many more, can be found at <http://www.ling.upenn.edu/~beatrice/humor/headlines.html>

1865 and then choosing the correct sense from the inventory associated with the corresponding
 1866 lemma.³ (Part-of-speech tagging is discussed in § 8.1.)

1867 **4.2.1 How many word senses?**

1868 Words sometimes have many more than two senses, as exemplified by the word *serve*:

- 1869 • [FUNCTION]: *The tree stump served as a table*
- 1870 • [CONTRIBUTE TO]: *His evasive replies only served to heighten suspicion*
- 1871 • [PROVIDE]: *We serve only the rawest fish*
- 1872 • [ENLIST]: *She served in an elite combat unit*
- 1873 • [JAIL]: *He served six years for a crime he didn't commit*
- 1874 • [LEGAL]: *They were served with subpoenas*⁴

1875 These sense distinctions are annotated in **WordNet** (<http://wordnet.princeton.edu>), a lexical semantic database for English. WordNet consists of roughly 100,000 **synsets**,
 1876 which are groups of lemmas (or phrases) that are synonymous. An example synset is
 1877 $\{chump^1, fool^2, sucker^1, mark^9\}$, where the superscripts index the sense of each lemma that
 1878 is included in the synset: for example, there are at least eight other senses of *mark* that
 1879 have different meanings, and are not part of this synset. A lemma is **polysemous** if it
 1880 participates in multiple synsets.

1882 WordNet defines the scope of the word sense disambiguation problem, and, more
 1883 generally, formalizes lexical semantic knowledge of English. (WordNets have been cre-
 1884 ated for a few dozen other languages, at varying levels of detail.) Some have argued
 1885 that WordNet's sense granularity is too fine (Ide and Wilks, 2006); more fundamentally,
 1886 the premise that word senses can be differentiated in a task-neutral way has been criti-
 1887 cized as linguistically naïve (Kilgarriff, 1997). One way of testing this question is to ask
 1888 whether people tend to agree on the appropriate sense for example sentences: accord-
 1889 ing to Mihalcea et al. (2004), people agree on roughly 70% of examples using WordNet
 1890 senses; far better than chance, but less than agreement on other tasks, such as sentiment
 1891 annotation (Wilson et al., 2005).

1892 ***Other lexical semantic relations** Besides **synonymy**, WordNet also describes many
 1893 other lexical semantic relationships, including:

- 1894 • **antonymy**: *x* means the opposite of *y*, e.g. FRIEND-ENEMY;

³Navigli (2009) provides a survey of approaches for word-sense disambiguation.

⁴Several of the examples are adapted from WordNet (Fellbaum, 2010).

- 1895 • **hyponymy:** x is a special case of y , e.g. RED-COLOR; the inverse relationship is
1896 **hyperonymy**;
- 1897 • **meronymy:** x is a part of y , e.g., WHEEL-BICYCLE; the inverse relationship is **holonymy**.

1898 Classification of these relations can be performed by searching for characteristic pat-
1899 terns between pairs of words, e.g., X , *such as* Y , which signals hyponymy (Hearst, 1992),
1900 or X *but* Y , which signals antonymy (Hatzivassiloglou and McKeown, 1997). Another ap-
1901 proach is to analyze each term's **distributional statistics** (the frequency of its neighboring
1902 words). Such approaches are described in detail in chapter 14.

1903 4.2.2 Word sense disambiguation as classification

1904 How can we tell living *plants* from manufacturing *plants*? The context is often critical:

- 1905 (4.11) Town officials are hoping to attract new manufacturing plants through weakened
1906 environmental regulations.
- 1907 (4.12) The endangered plants play an important role in the local ecosystem.

It is possible to build a feature vector using the bag-of-words representation, by treat-
ing each context as a pseudo-document. The feature function is then,

$$f((\text{plant}, \text{The endangered plants play an ...}), y) = \\ \{(the, y) : 1, (\text{endangered}, y) : 1, (\text{play}, y) : 1, (\text{an}, y) : 1, \dots\}$$

1908 As in document classification, many of these features are irrelevant, but a few are very
1909 strong predictors. In this example, the context word *endangered* is a strong signal that
1910 the intended sense is biology rather than manufacturing. We would therefore expect a
1911 learning algorithm to assign high weight to (*endangered*, BIOLOGY), and low weight to
1912 (*endangered*, MANUFACTURING).⁵

It may also be helpful to go beyond the bag-of-words: for example, one might encode
the position of each context word with respect to the target, e.g.,

$$f((\text{bank}, I \text{ went to the bank to deposit my paycheck}), y) = \\ \{(i - 3, \text{went}, y) : 1, (i + 2, \text{deposit}, y) : 1, (i + 4, \text{paycheck}, y) : 1\}$$

1913 These are called **collocation features**, and they give more information about the specific
1914 role played by each context word. This idea can be taken further by incorporating addi-
1915 tional syntactic information about the grammatical role played by each context feature,
1916 such as the **dependency path** (see chapter 11).

⁵The context bag-of-words can be also used to perform word-sense disambiguation without machine learning: the Lesk (1986) algorithm selects the word sense whose dictionary definition best overlaps the local context.

Using such features, a classifier can be trained from labeled data. A **semantic concordance** is a corpus in which each open-class word (nouns, verbs, adjectives, and adverbs) is tagged with its word sense from the target dictionary or thesaurus. SemCor is a semantic concordance built from 234K tokens of the Brown corpus (Francis and Kucera, 1982), annotated as part of the WordNet project (Fellbaum, 2010). SemCor annotations look like this:

(4.13) As of Sunday¹_N night¹_N there was⁴_V no word²_N ...,

with the superscripts indicating the annotated sense of each polysemous word, and the subscripts indicating the part-of-speech.

As always, supervised classification is only possible if enough labeled examples can be accumulated. This is difficult in word sense disambiguation, because each polysemous lemma requires its own training set: having a good classifier for the senses of *serve* is no help towards disambiguating *plant*. For this reason, **unsupervised** and **semisupervised** methods are particularly important for word sense disambiguation (e.g., Yarowsky, 1995). These methods will be discussed in chapter 5. Unsupervised methods typically lean on the heuristic of “one sense per discourse”, which means that a lemma will usually have a single, consistent sense throughout any given document (Gale et al., 1992). Based on this heuristic, we can propagate information from high-confidence instances to lower-confidence instances in the same document (Yarowsky, 1995).

4.3 Design decisions for text classification

Text classification involves a number of design decisions. In some cases, the design decision is clear from the mathematics: if you are using regularization, then a regularization weight λ must be chosen. Other decisions are more subtle, arising only in the low level “plumbing” code that ingests and processes the raw data. Such decision can be surprisingly consequential for classification accuracy.

4.3.1 What is a word?

The bag-of-words representation presupposes that extracting a vector of word counts from text is unambiguous. But text documents are generally represented as sequences of characters (in an encoding such as ascii or unicode), and the conversion to bag-of-words presupposes a definition of the “words” that are to be counted.

4.3.1.1 Tokenization

The first subtask for constructing a bag-of-words vector is **tokenization**: converting the text from a sequence of characters to a sequence of **word tokens**. A simple approach is

Whitespace	Isn't Ahab, Ahab? ;)
Treebank	Is n't Ahab , Ahab ? ;)
Tweet	Isn't Ahab , Ahab ? ;)
TokTok (Dehdari, 2014)	Isn ' t Ahab , Ahab ? ;)

Figure 4.1: The output of four `nltk` tokenizers, applied to the string *Isn't Ahab, Ahab? ;)*

1950 to define a subset of characters as whitespace, and then split the text on these tokens.
 1951 However, whitespace-based tokenization is not ideal: we may want to split conjunctions
 1952 like *isn't* and hyphenated phrases like *prize-winning* and *half-asleep*, and we likely want
 1953 to separate words from commas and periods that immediately follow them. At the same
 1954 time, it would be better not to split abbreviations like *U.S.* and *Ph.D.* In languages with
 1955 Roman scripts, tokenization is typically performed using regular expressions, with mod-
 1956 ules designed to handle each of these cases. For example, the `nltk` package includes a
 1957 number of tokenizers (Loper and Bird, 2002); the outputs of four of the better-known tok-
 1958 enizers are shown in Figure 4.1. Social media researchers have found that emoticons and
 1959 other forms of orthographic variation pose new challenges for tokenization, leading to the
 1960 development of special purpose tokenizers to handle these phenomena (O'Connor et al.,
 1961 2010).

1962 Tokenization is a language-specific problem, and each language poses unique chal-
 1963 lenges. For example, Chinese does not include spaces between words, nor any other
 1964 consistent orthographic markers of word boundaries. A “greedy” approach is to scan the
 1965 input for character substrings that are in a predefined lexicon. However, Xue et al. (2003)
 1966 notes that this can be ambiguous, since many character sequences could be segmented in
 1967 multiple ways. Instead, he trains a classifier to determine whether each Chinese character,
 1968 or *hanzi*, is a word boundary. More advanced sequence labeling methods for word seg-
 1969 mentation are discussed in § 8.4. Similar problems can occur in languages with alphabetic
 1970 scripts, such as German, which does not include whitespace in compound nouns, yield-
 1971 ing examples such as *Freundschaftsbezeugungen* (demonstration of friendship) and *Dilett-*
1972 tantenaufdringlichkeiten (the importunities of dilettantes). As Twain (1997) argues, “*These*
1973 things are not words, they are alphabetic processions.” Social media raises similar problems
 1974 for English and other languages, with hashtags such as *#TrueLoveInFourWords* requiring
 1975 decomposition for analysis (Brun and Roux, 2014).

1976 4.3.1.2 Normalization

1977 After splitting the text into tokens, the next question is which tokens are really distinct.
 1978 Is it necessary to distinguish *great*, *Great*, and *GREAT*? Sentence-initial capitalization may
 1979 be irrelevant to the classification task. Going further, the complete elimination of case
 1980 distinctions will result in a smaller vocabulary, and thus smaller feature vectors. However,

Original	The	Williams	sisters	are	leaving	this	tennis	centre
Porter stemmer	the	william	sister	are	leav	thi	tenni	centr
Lancaster stemmer	the	william	sist	ar	leav	thi	ten	cent
WordNet lemmatizer	The	Williams	sister	are	leaving	this	tennis	centre

Figure 4.2: Sample outputs of the Porter (1980) and Lancaster (Paice, 1990) stemmers, and the WordNet lemmatizer

1981 case distinctions might be relevant in some situations: for example, *apple* is a delicious
 1982 pie filling, while *Apple* is a company that specializes in proprietary dongles and power
 1983 adapters.

1984 For Roman script, case conversion can be performed using unicode string libraries.
 1985 Many scripts do not have case distinctions (e.g., the Devanagari script used for South
 1986 Asian languages, the Thai alphabet, and Japanese kana), and case conversion for all scripts
 1987 may not be available in every programming environment. (Unicode support is an im-
 1988 portant distinction between Python’s versions 2 and 3, and is a good reason for mi-
 1989 grating to Python 3 if you have not already done so. Compare the output of the code
 1990 "\à l\'hôtel".upper() in the two language versions.)⁶

1991 Case conversion is a type of **normalization**, which refers to string transformations that
 1992 remove distinctions that are irrelevant to downstream applications (Sproat et al., 2001).
 1993 Other normalizations include the standardization of numbers (e.g., 1,000 to 1000) and
 1994 dates (e.g., August 11, 2015 to 2015/11/08). Depending on the application, it may even be
 1995 worthwhile to convert all numbers and dates to special tokens, !NUM and !DATE. In social
 1996 media, there are additional orthographic phenomena that may be normalized, such as ex-
 1997 pressive lengthening, e.g., *coooooool* (Aw et al., 2006; Yang and Eisenstein, 2013). Similarly,
 1998 historical texts feature spelling variations that may need to be normalized to a contempo-
 1999 rary standard form (Baron and Rayson, 2008).

2000 A more extreme form of normalization is to eliminate **inflectional affixes**, such as the
 2001 -ed and -s suffixes in English. On this view, *bike*, *bikes*, *biking*, and *biked* all refer to the
 2002 same underlying concept, so they should be grouped into a single feature. A **stemmer** is
 2003 a program for eliminating affixes, usually by applying a series of regular expression sub-
 2004 stitutions. Character-based stemming algorithms are necessarily approximate, as shown
 2005 in Figure 4.2: the Lancaster stemmer incorrectly identifies -ers as an inflectional suffix of
 2006 *sisters* (by analogy to *fix/fixer*s), and both stemmers incorrectly identify -s as a suffix of *this*
 2007 and *Williams*. Fortunately, even inaccurate stemming can improve bag-of-words classifi-
 2008 cation models, by merging related strings and thereby reducing the vocabulary size.

2009 Accurately handling irregular orthography requires word-specific rules. **Lemmatizers**

⁶[todo: I want to make this a footnote, but can't figure out how.]

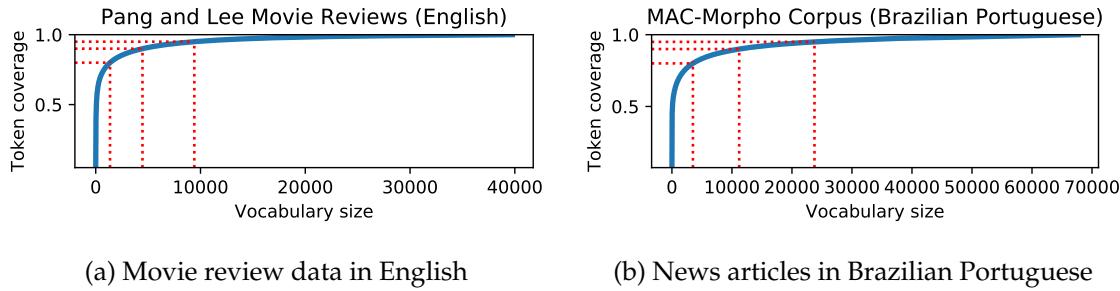


Figure 4.3: Tradeoff between token coverage (y-axis) and vocabulary size, on the `nltk` movie review dataset, after sorting the vocabulary by decreasing frequency. The red dashed lines indicate 80%, 90%, and 95% coverage.

2010 are systems that identify the underlying lemma of a given wordform. They must avoid the
 2011 over-generalization errors of the stemmers in Figure 4.2, and also handle more complex
 2012 transformations, such as *geese*→*goose*. The output of the WordNet lemmatizer is shown in
 2013 the final line of Figure 4.2. Both stemming and lemmatization are language-specific: an
 2014 English stemmer or lemmatizer is of little use on a text written in another language. The
 2015 discipline of **morphology** relates to the study of word-internal structure, and is described
 2016 in more detail in § 9.1.2.

2017 The value of normalization depends on the data and the task. Normalization re-
 2018 duces the size of the feature space, which can help in generalization. However, there
 2019 is always the risk of merging away linguistically meaningful distinctions. In supervised
 2020 machine learning, regularization and smoothing can play a similar role to normalization
 2021 — preventing the learner from overfitting to rare features — while avoiding the language-
 2022 specific engineering required for accurate normalization. In unsupervised scenarios, such
 2023 as content-based information retrieval (Manning et al., 2008) and topic modeling (Blei
 2024 et al., 2003), normalization is more critical.

2025 4.3.2 How many words?

2026 Limiting the size of the feature vector reduces the memory footprint of the resulting mod-
 2027 els, and increases the speed of prediction. Normalization can help to play this role, but
 2028 a more direct approach is simply to limit the vocabulary to the N most frequent words
 2029 in the dataset. For example, in the movie-reviews dataset provided with `nltk` (orig-
 2030 inally from Pang et al., 2002), there are 39,768 word types, and 1.58M tokens. As shown
 2031 in Figure 4.3a, the most frequent 4000 word types cover 90% of all tokens, offering an
 2032 order-of-magnitude reduction in the model size. Such ratios are language-specific: in for
 2033 example, in the Brazilian Portuguese Mac-Morpho corpus (Aluísio et al., 2003), attain-
 2034 ing 90% coverage requires more than 10000 word types (Figure 4.3b). This reflects the

2035 morphological complexity of Portuguese, which includes many more inflectional suffixes
 2036 than English.

2037 Eliminating rare words is not always advantageous for classification performance: for
 2038 example, names, which are typically rare, play a large role in distinguishing topics of news
 2039 articles. Another way to reduce the size of the feature space is to eliminate **stopwords** such
 2040 as *the*, *to*, and *and*, which may seem to play little role in expressing the topic, sentiment,
 2041 or stance. This is typically done by creating a **stoplist** (e.g., `nltk.corpus.stopwords`),
 2042 and then ignoring all terms that match the list. However, corpus linguists and social psy-
 2043 chologists have shown that seemingly inconsequential words can offer surprising insights
 2044 about the author or nature of the text (Biber, 1991; Chung and Pennebaker, 2007). Further-
 2045 more, high-frequency words are unlikely to cause overfitting in discriminative classifiers.
 2046 As with normalization, stopword filtering is more important for unsupervised problems,
 2047 such as term-based document retrieval.

2048 Another alternative for controlling model size is **feature hashing** (Weinberger et al.,
 2049 2009). Each feature is assigned an index using a hash function. If a hash function that
 2050 permits collisions is chosen (typically by taking the hash output modulo some integer),
 2051 then the model can be made arbitrarily small, as multiple features share a single weight.
 2052 Because most features are rare, accuracy is surprisingly robust to such collisions (Ganchev
 2053 and Dredze, 2008).

2054 4.3.3 Count or binary?

2055 Finally, we may consider whether we want our feature vector to include the *count* of each
 2056 word, or its *presence*. This gets at a subtle limitation of linear classification: it worse to
 2057 have two *failures* than one, but is it really twice as bad? Motivated by this intuition, Pang
 2058 et al. (2002) use binary indicators of presence or absence in the feature vector: $f_j(x, y) \in$
 2059 $\{0, 1\}$. They find that classifiers trained on these binary vectors tend to outperform feature
 2060 vectors based on word counts. One explanation is that words tend to appear in clumps:
 2061 if a word has appeared once in a document, it is likely to appear again (Church, 2000).
 2062 These subsequent appearances can be attributed to this tendency towards repetition, and
 2063 thus provide little additional information about the class label of the document.

2064 4.4 Evaluating classifiers

2065 In any supervised machine learning application, it is critical to reserve a held-out test set.
 2066 This data should be used for only one purpose: to evaluate the overall accuracy of a single
 2067 classifier. Using this data more than once would cause the estimated accuracy to be overly
 2068 optimistic, because the classifier would be customized to this data, and would not perform
 2069 as well as on unseen data in the future. It is usually necessary to set hyperparameters or

2070 perform feature selection, so you may need to construct a **tuning** or **development set** for
 2071 this purpose, as discussed in § 2.1.5.

2072 There are a number of ways to evaluate classifier performance. The simplest is **accuracy**:
 2073 the number of correct predictions, divided by the total number of instances,

$$\text{acc}(\mathbf{y}, \hat{\mathbf{y}}) = \frac{1}{N} \sum_i^N \delta(y^{(i)} = \hat{y}). \quad [4.4]$$

2074 Exams are usually graded by accuracy. Why are other metrics necessary? The main
 2075 reason is **class imbalance**. Suppose you are building a classifier to detect whether an
 2076 electronic health record (EHR) describes symptoms of a rare disease, which appears in
 2077 only 1% of all documents in the dataset. A classifier that reports $\hat{y} = \text{NEGATIVE}$ for
 2078 all documents would achieve 99% accuracy, but would be practically useless. We need
 2079 metrics that are capable of detecting the classifier's ability to discriminate between classes,
 2080 even when the distribution is skewed.

2081 One solution is to build a **balanced test set**, in which each possible label is equally rep-
 2082 resented. But in the EHR example, this would mean throwing away 98% of the original
 2083 dataset! Furthermore, the detection threshold itself might be a design consideration: in
 2084 health-related applications, we might prefer a very sensitive classifier, which returned a
 2085 positive prediction if there is even a small chance that $y^{(i)} = \text{POSITIVE}$. In other applica-
 2086 tions, a positive result might trigger a costly action, so we would prefer a classifier that
 2087 only makes positive predictions when absolutely certain. We need additional metrics to
 2088 capture these characteristics.

2089 4.4.1 Precision, recall, and F-MEASURE

2090 For any label (e.g., positive for presence of symptoms of a disease), there are two possible
 2091 errors:

- 2092 • **False positive**: the system incorrectly predicts the label.
- 2093 • **False negative**: the system incorrectly fails to predict the label.

2094 Similarly, for any label, there are two ways to be correct:

- 2095 • **True positive**: the system correctly predicts the label.
- 2096 • **True negative**: the system correctly predicts that the label does not apply to this
 2097 instance.

Classifiers that make a lot of false positives are too sensitive; classifiers that make a lot of false negatives are not sensitive enough. These two conditions are captured by the

metrics of **recall** and **precision**:

$$\text{RECALL}(\mathbf{y}, \hat{\mathbf{y}}, k) = \frac{\text{TP}}{\text{TP} + \text{FN}} \quad [4.5]$$

$$\text{PRECISION}(\mathbf{y}, \hat{\mathbf{y}}, k) = \frac{\text{TP}}{\text{TP} + \text{FP}}. \quad [4.6]$$

2098 Recall and precision are both conditional likelihoods of a correct prediction, which is why
 2099 their numerators are the same. Recall is conditioned on k being the correct label, $y^{(i)} = k$,
 2100 so the denominator sums over true positive and false negatives. Precision is conditioned
 2101 on k being the prediction, so the denominator sums over true positives and false positives.
 2102 Note that true negatives are not considered in either statistic. The classifier that labels
 2103 every document as “negative” would achieve zero recall; precision would be $\frac{0}{0}$.

2104 Recall and precision are complementary. A high-recall classifier is preferred when
 2105 false negatives are cheaper than false positives: for example, in a preliminary screening
 2106 for symptoms of a disease, the cost of a false positive might be an additional test, while a
 2107 false negative would result in the disease going untreated. Conversely, a high-precision
 2108 classifier is preferred when false positives are more expensive: for example, in spam de-
 2109tection, a false negative is a relatively minor inconvenience, while a false positive might
 2110 mean that an important message goes unread.

The ***F*-MEASURE** combines recall and precision into a single metric, using the harmonic mean:

$$\text{F-MEASURE}(\mathbf{y}, \hat{\mathbf{y}}, k) = \frac{2rp}{r + p}, \quad [4.7]$$

2111 where r is recall and p is precision.⁷

Evaluating multi-class classification Recall, precision, and ***F*-MEASURE** are defined with respect to a specific label k . When there are multiple labels of interest (e.g., in word sense disambiguation or emotion classification), it is necessary to combine the ***F*-MEASURE** across each class. **Macro *F*-MEASURE** is the average ***F*-MEASURE** across several classes,

$$\text{Macro-}F(\mathbf{y}, \hat{\mathbf{y}}) = \frac{1}{|\mathcal{K}|} \sum_{k \in \mathcal{K}} \text{F-MEASURE}(\mathbf{y}, \hat{\mathbf{y}}, k) \quad [4.8]$$

2112 In multi-class problems with unbalanced class distributions, the macro ***F*-MEASURE** is a
 2113 balanced measure of how well the classifier recognizes each class. In **micro *F*-MEASURE**,
 2114 we compute true positives, false positives, and false negatives for each class, and then add
 2115 them up to compute a single recall, precision, and ***F*-MEASURE**. This metric is balanced
 2116 across instances rather than classes, so it weights each class in proportion to its frequency
 2117 — unlike macro ***F*-MEASURE**, which weights each class equally.

⁷ F -MEASURE is sometimes called F_1 , and generalizes to $F_\beta = \frac{(1+\beta^2)rp}{\beta^2p+r}$. The β parameter can be tuned to emphasize recall or precision.

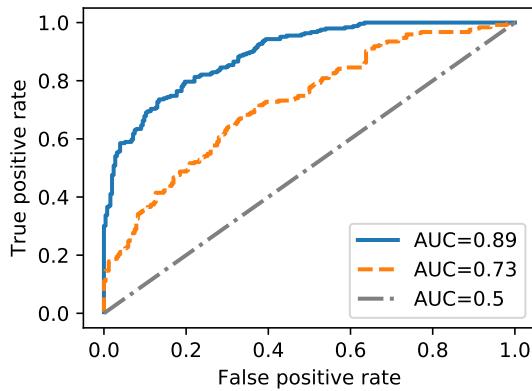


Figure 4.4: ROC curves for three classifiers of varying discriminative power, measured by AUC (area under the curve)

2118 4.4.2 Threshold-free metrics

2119 In binary classification problems, it is possible to trade off between recall and precision by
 2120 adding a constant “threshold” to the output of the scoring function. This makes it possible
 2121 to trace out a curve, where each point indicates the performance at a single threshold. In
 2122 the **receiver operating characteristic (ROC)** curve,⁸ the *x*-axis indicates the **false positive**
 2123 **rate**, $\frac{FP}{FP+TN}$, and the *y*-axis indicates the recall, or **true positive rate**. A perfect classifier
 2124 attains perfect recall without any false positives, tracing a “curve” from the origin (0,0) to
 2125 the upper left corner (0,1), and then to (1,1). In expectation, a non-discriminative classifier
 2126 traces a diagonal line from the origin (0,0) to the upper right corner (1,1). Real classifiers
 2127 tend to fall between these two extremes. Examples are shown in Figure 4.4.

2128 The ROC curve can be summarized in a single number by taking its integral, the **area**
 2129 **under the curve (AUC)**. The AUC can be interpreted as the probability that a randomly-
 2130 selected positive example will be assigned a higher score by the classifier than a randomly-
 2131 selected negative example. A perfect classifier has AUC = 1 (all positive examples score
 2132 higher than all negative examples); a non-discriminative classifier has AUC = 0.5 (given
 2133 a randomly selected positive and negative example, either could score higher with equal
 2134 probability); a perfectly wrong classifier would have AUC = 0 (all negative examples score
 2135 higher than all positive examples). One advantage of AUC in comparison to *F*-MEASURE
 2136 is that the baseline rate of 0.5 does not depend on the label distribution.

⁸The name “receiver operator characteristic” comes from the metric’s origin in signal processing applications (Peterson et al., 1954). Other threshold-free metrics include **precision-recall curves**, **precision-at-*k***, and **balanced *F*-MEASURE**; see Manning et al. (2008) for more details.

2137 **4.4.3 Classifier comparison and statistical significance**

2138 Natural language processing research and engineering often involves comparing different
 2139 classification techniques. In some cases, the comparison is between algorithms, such as
 2140 logistic regression versus averaged perceptron, or L_2 regularization versus L_1 . In other
 2141 cases, the comparison is between feature sets, such as the bag-of-words versus positional
 2142 bag-of-words (see § 4.2.2). **Ablation testing** involves systematically removing (ablating)
 2143 various aspects of the classifier, such as feature groups, and testing the **null hypothesis**
 2144 that the ablated classifier is as good as the full model.

2145 A full treatment of hypothesis testing is beyond the scope of this text, but this section
 2146 contains a brief summary of the techniques necessary to compare classifiers. The main
 2147 aim of hypothesis testing is to determine whether the difference between two statistics
 2148 — for example, the accuracies of two classifiers — is likely to arise by chance. We will
 2149 be concerned with chance fluctuations that arise due to the finite size of the test set.⁹ An
 2150 improvement of 10% on a test set with ten instances may reflect a random fluctuation that
 2151 makes the test set more favorable to classifier c_1 than c_2 ; on another test set with a different
 2152 ten instances, we might find that c_2 does better than c_1 . But if we observe the same 10%
 2153 improvement on a test set with 1000 instances, this is highly unlikely to be explained
 2154 by chance. Such a finding is said to be **statistically significant** at a level p , which is the
 2155 probability of observing an effect of equal or greater magnitude when the null hypothesis
 2156 is true. The notation $p < .05$ indicates that the likelihood of an equal or greater effect is
 2157 less than 5%, assuming the null hypothesis is true.¹⁰

2158 **4.4.3.1 The binomial test**

2159 The statistical significance of a difference in accuracy can be evaluated using classical tests,
 2160 such as the **binomial test**.¹¹ Suppose that classifiers c_1 and c_2 disagree on N instances in a
 2161 test set with binary labels, and that c_1 is correct on k of those instances. Under the null hy-
 2162 pothesis that the classifiers are equally accurate, we would expect k/N to be roughly equal
 2163 to 1/2, and as N increases, k/N should be increasingly close to this expected value. These
 2164 properties are captured by the **binomial distribution**, which is a probability over counts

⁹Other sources of variance include the initialization of non-convex classifiers such as neural networks, and the ordering of instances in online learning such as stochastic gradient descent and perceptron.

¹⁰Statistical hypothesis testing is useful only to the extent that the existing test set is representative of the instances that will be encountered in the future. If, for example, the test set is constructed from news documents, no hypothesis test can predict which classifier will perform best on documents from another domain, such as electronic health records.

¹¹A well-known alternative to the binomial test is **McNemar's test**, which computes a **test statistic** based on the number of examples that are correctly classified by one system and incorrectly classified by the other. The null hypothesis distribution for this test statistic is known to be drawn from a chi-squared distribution with a single degree of freedom, so a p -value can be computed from the cumulative density function of this distribution (Dietterich, 1998). Both tests give similar results in most circumstances, but the binomial test is easier to understand from first principles.

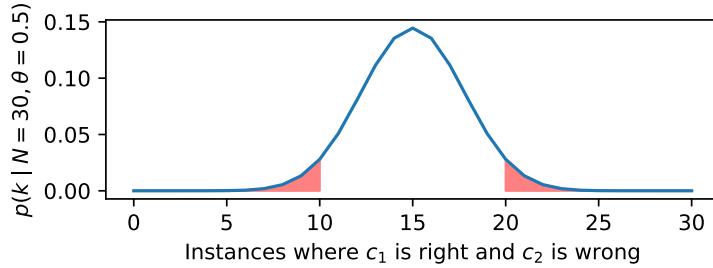


Figure 4.5: Probability mass function for the binomial distribution. The pink highlighted areas represent the cumulative probability for a significance test on an observation of $k = 10$ and $N = 30$.

of binary random variables. We write $k \sim \text{Binom}(\theta, N)$ to indicate that k is drawn from a binomial distribution, with parameter N indicating the number of random “draws”, and θ indicating the probability of “success” on each draw. Each draw is an example on which the two classifiers disagree, and a “success” is a case in which c_1 is right and c_2 is wrong. (The label space is assumed to be binary, so if the classifiers disagree, exactly one of them is correct. The test can be generalized to multi-class classification by focusing on the examples in which exactly one classifier is correct.)

The probability mass function (PMF) of the binomial distribution is,

$$p_{\text{Binom}}(k; N, \theta) = \binom{N}{k} \theta^k (1 - \theta)^{N-k}, \quad [4.9]$$

with θ^k representing the probability of the k successes, $(1 - \theta)^{N-k}$ representing the probability of the $N - k$ unsuccessful draws. The expression $\binom{N}{k} = \frac{N!}{k!(N-k)!}$ is a binomial coefficient, representing the number of possible orderings of events; this ensures that the distribution sums to one over all $k \in \{0, 1, 2, \dots, N\}$.

Under the null hypothesis, when the classifiers disagree, each classifier is equally likely to be right, so $\theta = \frac{1}{2}$. Now suppose that among N disagreements, c_1 is correct $k < \frac{N}{2}$ times. The probability of c_1 being correct k or fewer times is the **one-tailed p-value**, because it is computed from the area under the binomial probability mass function from 0 to k , as shown in the left tail of Figure 4.5. This **cumulative probability** is computed as a sum over all values $i \leq k$,

$$\Pr_{\text{Binom}} \left(\text{count}(\hat{y}_2^{(i)} = y^{(i)} \neq \hat{y}_1^{(i)}) \leq k; N, \theta = \frac{1}{2} \right) = \sum_{i=0}^k p_{\text{Binom}} \left(i; N, \theta = \frac{1}{2} \right). \quad [4.10]$$

The one-tailed p-value applies only to the asymmetric null hypothesis that c_1 is at least as accurate as c_2 . To test the **two-tailed** null hypothesis that c_1 and c_2 are equally accu-

Algorithm 7 Bootstrap sampling for classifier evaluation. The original test set is $\{\mathbf{x}^{(1:N)}, \mathbf{y}^{(1:N)}\}$, the metric is $\delta(\cdot)$, and the number of samples is M .

```

procedure BOOTSTRAP-SAMPLE( $\mathbf{x}^{(1:N)}, \mathbf{y}^{(1:N)}, \delta(\cdot), M$ )
    for  $t \in \{1, 2, \dots, M\}$  do
        for  $i \in \{1, 2, \dots, N\}$  do
             $j \sim \text{UniformInteger}(1, N)$ 
             $\tilde{\mathbf{x}}^{(i)} \leftarrow \mathbf{x}^{(j)}$ 
             $\tilde{\mathbf{y}}^{(i)} \leftarrow \mathbf{y}^{(j)}$ 
             $d^{(t)} \leftarrow \delta(\tilde{\mathbf{x}}^{(1:N)}, \tilde{\mathbf{y}}^{(1:N)})$ 
    return  $\{d^{(t)}\}_{t=1}^M$ 
```

2179 rate, we would take the sum of one-tailed p -values, where the second term is computed
 2180 from the right tail of Figure 4.5. The binomial distribution is symmetric, so this can be
 2181 computed by simply doubling the one-tailed p -value.

2182 Two-tailed tests are more stringent, but they are necessary in cases in which there is
 2183 no prior intuition about whether c_1 or c_2 is better. For example, in comparing logistic
 2184 regression versus averaged perceptron, a two-tailed test is appropriate. In an ablation
 2185 test, c_2 may contain a superset of the features available to c_1 . If the additional features are
 2186 thought to be likely to improve performance, then a one-tailed test would be appropriate,
 2187 if chosen in advance. However, such a test can only prove that c_2 is more accurate than
 2188 c_1 , and not the reverse.

2189 **4.4.3.2 *Randomized testing**

2190 The binomial test is appropriate for accuracy, but not for more complex metrics such as
 2191 F -MEASURE. To compute statistical significance for arbitrary metrics, we can apply ran-
 2192 domization. Specifically, draw a set of M **bootstrap samples** (Efron and Tibshirani, 1993),
 2193 by resampling instances from the original test set with replacement. Each bootstrap sam-
 2194 ple is itself a test set of size N . Some instances from the original test set will not appear
 2195 in any given bootstrap sample, while others will appear multiple times; but overall, the
 2196 sample will be drawn from the same distribution as the original test set. We can then com-
 2197 pute any desired evaluation on each bootstrap sample, which gives a distribution over the
 2198 value of the metric. Algorithm 7 shows how to perform this computation.

2199 To compare the F -MEASURE of two classifiers c_1 and c_2 , we set the function $\delta(\cdot)$ to
 2200 compute the difference in F -MEASURE on the bootstrap sample. If the difference is less
 2201 than or equal to zero in at least 5% of the samples, then we cannot reject the one-tailed
 2202 null hypothesis that c_2 is at least as good as c_1 (Berg-Kirkpatrick et al., 2012). We may
 2203 also be interested in the 95% **confidence interval** around a metric of interest, such as
 2204 the F -MEASURE of a single classifier. This can be computed by sorting the output of

2205 Algorithm 7, and then setting the top and bottom of the 95% confidence interval to the
 2206 values at the 2.5% and 97.5% percentiles of the sorted outputs. Alternatively, you can fit
 2207 a normal distribution to the set of differences across bootstrap samples, and compute a
 2208 Gaussian confidence interval from the mean and variance.

2209 As the number of bootstrap samples goes to infinity, $M \rightarrow \infty$, the bootstrap estimate
 2210 is increasingly accurate. A typical choice for M is 10^4 or 10^5 ; larger numbers of samples
 2211 are necessary for smaller p -values. One way to validate your choice of M is to run the test
 2212 multiple times, and ensure that the p -values are similar; if not, increase M by an order of
 2213 magnitude. This is a heuristic measure of the **variance** of the test, which can decrease
 2214 with the square root \sqrt{M} (Robert and Casella, 2013).

2215 4.4.4 *Multiple comparisons

2216 Sometimes it is necessary to perform multiple hypothesis tests, such as when compar-
 2217 ing the performance of several classifiers on multiple datasets. Suppose you have five
 2218 datasets, and you compare four versions of your classifier against a baseline system, for a
 2219 total of 20 comparisons. Even if none of your classifiers is better than the baseline, there
 2220 will be some chance variation in the results, and in expectation you will get one statis-
 2221 tically significant improvement at $p = 0.05 = \frac{1}{20}$. It is therefore necessary to adjust the
 2222 p -values when reporting the results of multiple comparisons.

2223 One approach is to require a threshold of $\frac{\alpha}{m}$ to report a p value of $p < \alpha$ when per-
 2224 forming m tests. This is known as the **Bonferroni correction**, and it limits the overall
 2225 probability of incorrectly rejecting the null hypothesis at α . Another approach is to bound
 2226 the **false discovery rate** (FDR), which is the fraction of null hypothesis rejections that are
 2227 incorrect. Benjamini and Hochberg (1995) propose a p -value correction that bounds the
 2228 fraction of false discoveries at α : sort the p -values of each individual test in ascending
 2229 order, and set the significance threshold equal to largest k such that $p_k \leq \frac{k}{m}\alpha$. If $k > 1$, the
 2230 FDR adjustment is more permissive than the Bonferroni correction.

2231 4.5 Building datasets

2232 Sometimes, if you want to build a classifier, you must first build a dataset of your own.
 2233 This includes selecting a set of documents or instances to annotate, and then performing
 2234 the annotations. The scope of the dataset may be determined by the application: if you
 2235 want to build a system to classify electronic health records, then you must work with a
 2236 corpus of records of the type that your classifier will encounter when deployed. In other
 2237 cases, the goal is to build a system that will work across a broad range of documents. In
 2238 this case, it is best to have a *balanced* corpus, with contributions from many styles and
 2239 genres. For example, the Brown corpus draws from texts ranging from government doc-
 2240 uments to romance novels (Francis, 1964), and the Google Web Treebank includes an-

2241 notations for five “domains” of web documents: question answers, emails, newsgroups,
2242 reviews, and blogs (Petrov and McDonald, 2012).

2243 **4.5.1 Metadata as labels**

2244 Annotation is difficult and time-consuming, and most people would rather avoid it. It
2245 is sometimes possible to exploit existing metadata to obtain labels for training a classi-
2246 fier. For example, reviews are often accompanied by a numerical rating, which can be
2247 converted into a classification label (see § 4.1). Similarly, the nationalities of social media
2248 users can be estimated from their profiles (Dredze et al., 2013) or even the time zones of
2249 their posts (Gouws et al., 2011). More ambitiously, we may try to classify the political af-
2250 filiations of social media profiles based on their social network connections to politicians
2251 and major political parties (Rao et al., 2010).

2252 The convenience of quickly constructing large labeled datasets without manual an-
2253 notation is appealing. However this approach relies on the assumption that unlabeled
2254 instances — for which metadata is unavailable — will be similar to labeled instances.
2255 Consider the example of labeling the political affiliation of social media users based on
2256 their network ties to politicians. If a classifier attains high accuracy on such a test set,
2257 is it safe to assume that it accurately predicts the political affiliation of all social media
2258 users? Probably not. Social media users who establish social network ties to politicians
2259 may be more likely to mention politics in the text of their messages, as compared to the
2260 average user, for whom no political metadata is available. If so, the accuracy on a test set
2261 constructed from social network metadata would give an overly optimistic picture of the
2262 method’s true performance on unlabeled data.

2263 **4.5.2 Labeling data**

2264 In many cases, there is no way to get ground truth labels other than manual annotation.
2265 An annotation protocol should satisfy several criteria: the annotations should be *expressive*
2266 enough to capture the phenomenon of interest; they should be *replicable*, meaning that
2267 another annotator or team of annotators would produce very similar annotations if given
2268 the same data; and they should be *scalable*, so that they can be produced relatively quickly.
2269 Hovy and Lavid (2010) propose a structured procedure for obtaining annotations that
2270 meet these criteria, which is summarized below.

- 2271 1. **Determine what the annotations are to include.** This is usually based on some
2272 theory of the underlying phenomenon: for example, if the goal is to produce an-
2273 notations about the emotional state of a document’s author, one should start with a
2274 theoretical account of the types or dimensions of emotion (e.g., Mohammad and Tur-
2275 ney, 2013). At this stage, the tradeoff between expressiveness and scalability should

2276 be considered: a full instantiation of the underlying theory might be too costly to
2277 annotate at scale, so reasonable approximations should be considered.

- 2278 2. Optionally, one may **design or select a software tool to support the annotation**
2279 **effort**. Existing general-purpose annotation tools include BRAT (Stenetorp et al.,
2280 2012) and MMAX2 (Müller and Strube, 2006).
- 2281 3. **Formalize the instructions for the annotation task.** To the extent that the instruc-
2282 tions are not explicit, the resulting annotations will depend on the intuitions of the
2283 annotators. These intuitions may not be shared by other annotators, or by the users
2284 of the annotated data. Therefore explicit instructions are critical to ensuring the an-
2285 notations are replicable and usable by other researchers.
- 2286 4. **Perform a pilot annotation** of a small subset of data, with multiple annotators for
2287 each instance. This will give a preliminary assessment of both the replicability and
2288 scalability of the current annotation instructions. Metrics for computing the rate of
2289 agreement are described below. Manual analysis of specific disagreements should
2290 help to clarify the instructions, and may lead to modifications of the annotation task
2291 itself. For example, if two labels are commonly conflated by annotators, it may be
2292 best to merge them.
- 2293 5. **Annotate the data.** After finalizing the annotation protocol and instructions, the
2294 main annotation effort can begin. Some, if not all, of the instances should receive
2295 multiple annotations, so that inter-annotator agreement can be computed. In some
2296 annotation projects, instances receive many annotations, which are then aggregated
2297 into a “consensus” label (e.g., Danescu-Niculescu-Mizil et al., 2013). However, if the
2298 annotations are time-consuming or require significant expertise, it may be preferable
2299 to maximize scalability by obtaining multiple annotations for only a small subset of
2300 examples.
- 2301 6. **Compute and report inter-annotator agreement, and release the data.** In some
2302 cases, the raw text data cannot be released, due to concerns related to copyright or
2303 privacy. In these cases, one solution is to publicly release **stand-off annotations**,
2304 which contain links to document identifiers. The documents themselves can be re-
2305 leased under the terms of a licensing agreement, which can impose conditions on
2306 how the data is used. It is important to think through the potential consequences of
2307 releasing data: people may make personal data publicly available without realizing
2308 that it could be redistributed in a dataset and publicized far beyond their expecta-
2309 tions (boyd and Crawford, 2012).

2310 **4.5.2.1 Measuring inter-annotator agreement**

2311 To measure the replicability of annotations, a standard practice is to compute the extent to
 2312 which annotators agree with each other. If the annotators frequently disagree, this casts
 2313 doubt on either their reliability or on the annotation system itself. For classification, one
 2314 can compute the frequency with which the annotators agree; for rating scales, one can
 2315 compute the average distance between ratings. These raw agreement statistics must then
 2316 be compared with the rate of **chance agreement** — the level of agreement that would be
 2317 obtained between two annotators who ignored the data.

2318 **Cohen's Kappa** is widely used for quantifying the agreement on discrete labeling
 2319 tasks (Cohen, 1960; Carletta, 1996),¹²

$$\kappa = \frac{\text{agreement} - E[\text{agreement}]}{1 - E[\text{agreement}]}.$$
 [4.11]

2320 The numerator is the difference between the observed agreement and the chance agree-
 2321 ment, and the denominator is the difference between perfect agreement and chance agree-
 2322 ment. Thus, $\kappa = 1$ when the annotators agree in every case, and $\kappa = 0$ when the annota-
 2323 tors agree only as often as would happen by chance. Various heuristic scales have been
 2324 proposed for determining when κ indicates “moderate”, “good”, or “substantial” agree-
 2325 ment; for reference, Lee and Narayanan (2005) report $\kappa \approx 0.45 - 0.47$ for annotations
 2326 of emotions in spoken dialogues, which they describe as “moderate agreement”; Stolcke
 2327 et al. (2000) report $\kappa = 0.8$ for annotations of **dialogue acts**, which are labels for the pur-
 2328 pose of each turn in a conversation.

2329 When there are two annotators, the expected chance agreement is computed as,

$$E[\text{agreement}] = \sum_k \hat{\Pr}(Y = k)^2,$$
 [4.12]

2330 where k is a sum over labels, and $\hat{\Pr}(Y = k)$ is the empirical probability of label k across
 2331 all annotations. The formula is derived from the expected number of agreements if the
 2332 annotations were randomly shuffled. Thus, in a binary labeling task, if one label is applied
 2333 to 90% of instances, chance agreement is $.9^2 + .1^2 = .82$.

2334 **4.5.2.2 Crowdsourcing**

2335 Crowdsourcing is often used to rapidly obtain annotations for classification problems.
 2336 For example, **Amazon Mechanical Turk** makes it possible to define “human intelligence
 2337 tasks (hits)”, such as labeling data. The researcher sets a price for each set of annotations
 2338 and a list of minimal qualifications for annotators, such as their native language and their

¹² For other types of annotations, Krippendorff's alpha is a popular choice (Hayes and Krippendorff, 2007; Artstein and Poesio, 2008).

2339 satisfaction rate on previous tasks. The use of relatively untrained “crowdworkers” con-
 2340 trasts with earlier annotation efforts, which relied on professional linguists (Marcus et al.,
 2341 1993). However, crowdsourcing has been found to produce reliable annotations for many
 2342 language-related tasks (Snow et al., 2008). Crowdsourcing is part of the broader field of
 2343 **human computation** (Law and Ahn, 2011).

2344 Additional resources

2345 Many of the preprocessing issues discussed in this chapter also arise in information re-
 2346 trieval. See (Manning et al., 2008, chapter 2) for discussion of tokenization and related
 2347 algorithms.

2348 Exercises

2349 1. As noted in § 4.3.3, words tend to appear in clumps, with subsequent occurrences
 2350 of a word being more probable. More concretely, if word j has probability $\phi_{y,j}$
 2351 of appearing in a document with label y , then the probability of two appearances
 2352 ($x_j^{(i)} = 2$) is greater than $\phi_{y,j}^2$.

2353 Suppose you are applying Naïve Bayes to a binary classification. Focus on a word j
 2354 which is more probable under label $y = 1$, so that,

$$\Pr(w = j \mid y = 1) > \Pr(w = j \mid y = 0). \quad [4.13]$$

2355 Now suppose that $x_j^{(i)} > 1$. All else equal, will the classifier overestimate or under-
 2356 estimate the posterior $\Pr(y = 1 \mid x)$?

2357 2. Prove that F-measure is never greater than the arithmetic mean of recall and pre-
 2358 cision, $\frac{r+p}{2}$. Your solution should also show that F-measure is equal to $\frac{r+p}{2}$ iff $r = p$.

2359 3. Given a binary classification problem in which the probability of the “positive” label
 2360 is equal to α , what is the expected F-MEASURE of a random classifier which ignores
 2361 the data, and selects $\hat{y} = +1$ with probability $\frac{1}{2}$? (Assume that $p(\hat{y}) \perp p(y)$.) What is
 2362 the expected F-MEASURE of a classifier that selects $\hat{y} = +1$ with probability α (also
 2363 independent of $y^{(i)}$)? Depending on α , which random classifier will score better?

2364 4. Suppose that binary classifiers c_1 and c_2 disagree on $N = 30$ cases, and that c_1 is
 2365 correct in $k = 10$ of those cases.

- 2366 • Write a program that uses primitive functions such as `exp` and `factorial` to com-
 2367 pute the **two-tailed** p -value — you may use an implementation of the “choose”
 2368 function if one is available. Verify your code against the output of a library for

2369 computing the binomial test or the binomial CDF, such as `scipy.stats.binom`
 2370 in Python.

- 2371 • Then use a randomized test to try to obtain the same p -value. In each sample,
 2372 draw from a binomial distribution with $N = 30$ and $\theta = \frac{1}{2}$. Count the fraction
 2373 of samples in which $k \leq 10$. This is the one-tailed p -value; double this to
 2374 compute the two-tailed p -value.
- 2375 • Try this with varying numbers of bootstrap samples: $M \in \{100, 1000, 5000, 10000\}$.
 2376 For $M = 100$ and $M = 1000$, run the test 10 times, and plot the resulting p -
 2377 values.
- 2378 • Finally, perform the same tests for $N = 70$ and $k = 25$.

2379 5. SemCor 3.0 is a labeled dataset for word sense disambiguation. You can download
 2380 it,¹³ or access it in `nltk.corpora.semcor`.

2381 Choose a word that appears at least ten times in SemCor (*find*), and annotate its
 2382 WordNet senses across ten randomly-selected examples, without looking at the ground
 2383 truth. Use online WordNet to understand the definition of each of the senses.¹⁴ Have
 2384 a partner do the same annotations, and compute the raw rate of agreement, expected
 2385 chance rate of agreement, and Cohen's kappa.

2386 6. Download the Pang and Lee movie review data, currently available from <http://www.cs.cornell.edu/people/pabo/movie-review-data/>. Hold out a
 2387 randomly-selected 400 reviews as a test set.

2388 Download a sentiment lexicon, such as the one currently available from Bing Liu,
 2389 <https://www.cs.uic.edu/~liub/FBS/sentiment-analysis.html>. Tokenize
 2390 the data, and classify each document as positive iff it has more positive sentiment
 2391 words than negative sentiment words. Compute the accuracy and *F*-MEASURE on
 2392 detecting positive reviews on the test set, using this lexicon-based classifier.

2393 Then train a discriminative classifier (averaged perceptron or logistic regression) on
 2394 the training set, and compute its accuracy and *F*-MEASURE on the test set.

2395 Determine whether the differences are statistically significant, using two-tailed hy-
 2396 pothesis tests: Binomial for the difference in accuracy, and bootstrap for the differ-
 2397 ence in macro-*F*-MEASURE.

2398 2399 The remaining problems will require you to build a classifier and test its properties. Pick
 2400 a multi-class text classification dataset, such as RCV1¹⁵). Divide your data into training

¹³e.g., https://github.com/google-research-datasets/word_sense_disambiguation_corpora or <http://globalwordnet.org/wordnet-annotated-corpora/>

¹⁴<http://wordnetweb.princeton.edu/perl/webwn>

¹⁵http://www.ai.mit.edu/projects/jmlr/papers/volume5/lewis04a/lyrl2004_rcv1v2_README.htm

2401 (60%), development (20%), and test sets (20%), if no such division already exists. [todo:
2402 this dataset is already tokenized, find something else]

2403 7. Compare various vocabulary sizes of $10^2, 10^3, 10^4, 10^5$, using the most frequent words
2404 in each case (you may use any reasonable tokenizer). Train logistic regression clas-
2405 sifiers for each vocabulary size, and apply them to the development set. Plot the
2406 accuracy and Macro-*F*-MEASURE with the increasing vocabulary size. For each vo-
2407 cabulary size, tune the regularizer to maximize accuracy on a subset of data that is
2408 held out from the training set.

2409 8. Compare the following tokenization algorithms:

- 2410 • Whitespace, using a regular expression
2411 • Penn Treebank
2412 • Split input into five-character units, regardless of whitespace or punctuation

2413 Compute the token/type ratio for each tokenizer on the training data, and explain
2414 what you find. Train your classifier on each tokenized dataset, tuning the regularizer
2415 on a subset of data that is held out from the training data. Tokenize the development
2416 set, and report accuracy and Macro-*F*-MEASURE.

2417 9. Apply the Porter and Lancaster stemmers to the training set, using any reasonable
2418 tokenizer, and compute the token/type ratios. Train your classifier on the stemmed
2419 data, and compute the accuracy and Macro-*F*-MEASURE on stemmed development
2420 data, again using a held-out portion of the training data to tune the regularizer.

2421 10. Identify the best combination of vocabulary filtering, tokenization, and stemming
2422 from the previous three problems. Apply this preprocessing to the test set, and
2423 compute the test set accuracy and Macro-*F*-MEASURE. Compare against a baseline
2424 system that applies no vocabulary filtering, whitespace tokenization, and no stem-
2425 ming.

2426 Use the binomial test to determine whether your best-performing system is signifi-
2427 cantly more accurate than the baseline.

2428 Use the bootstrap test with $M = 10^4$ to determine whether your best-performing
2429 system achieves significantly higher macro-*F*-MEASURE.

2430 Chapter 5

2431 Learning without supervision

2432 So far we've assumed the following setup:

- 2433 a **training set** where you get observations x and labels y ;
- 2434 a **test set** where you only get observations x .

2435 Without labeled data, is it possible to learn anything? This scenario is known as **unsu-**
2436 **pervised learning**, and we will see that indeed it is possible to learn about the underlying
2437 structure of unlabeled observations. This chapter will also explore some related scenarios:
2438 **semi-supervised learning**, in which only some instances are labeled, and **domain adap-**
2439 **tation**, in which the training data differs from the data on which the trained system will
2440 be deployed.

2441 5.1 Unsupervised learning

2442 To motivate unsupervised learning, consider the problem of word sense disambiguation
2443 (§ 4.2). Our goal is to classify each instance of a word, such as *bank* into a sense,

- 2444 bank#1: a financial institution
- 2445 bank#2: the land bordering a river

2446 It is difficult to obtain sufficient training data for word sense disambiguation, because
2447 even a large corpus will contain only a few instances of all but the most common words.
2448 Is it possible to learn anything about these different senses without labeled data?

2449 Word sense disambiguation is usually performed using feature vectors constructed
2450 from the local context of the word to be disambiguated. For example, for the word




Figure 5.1: Counts of words from two different context groups

2451 *bank*, the immediate context might typically include words from one of the following two
 2452 groups:

- 2453 1. *financial, deposits, credit, lending, capital, markets, regulated, reserve, liquid, assets*
 2454 2. *land, water, geography, stream, river, flow, deposits, discharge, channel, ecology*

2455 Now consider a scatterplot, in which each point is a document containing the word *bank*.
 2456 The location of the document on the x -axis is the count of words in group 1, and the
 2457 location on the y -axis is the count for group 2. In such a plot, shown in Figure 5.1, two
 2458 “blobs” might emerge, and these blobs correspond to the different senses of *bank*.

2459 Here’s a related scenario, from a different problem. Suppose you download thousands
 2460 of news articles, and make a scatterplot, where each point corresponds to a document:
 2461 the x -axis is the frequency of the group of words (*hurricane, winds, storm*); the y -axis is the
 2462 frequency of the group (*election, voters, vote*). This time, three blobs might emerge: one
 2463 for documents that are largely about a hurricane, another for documents largely about a
 2464 election, and a third for documents about neither topic.

2465 These clumps represent the underlying structure of the data. But the two-dimensional
 2466 scatter plots are based on groupings of context words, and in real scenarios these word
 2467 lists are unknown. Unsupervised learning applies the same basic idea, but in a high-
 2468 dimensional space with one dimension for every context word. This space can’t be di-
 2469 rectly visualized, but the idea is the same: try to identify the underlying structure of the
 2470 observed data, such that there are a few clusters of points, each of which is internally
 2471 coherent. **Clustering** algorithms are capable of finding such structure automatically.

2472 5.1.1 **K-means** clustering

2473 Clustering algorithms assign each data point to a discrete cluster, $z_i \in 1, 2, \dots, K$. One of
 2474 the best known clustering algorithms is ***K-means***, an iterative algorithm that maintains

Algorithm 8 K -means clustering algorithm

```

1: procedure  $K$ -MEANS( $\mathbf{x}_{1:N}, K$ )
2:   for  $i \in 1 \dots N$  do                                 $\triangleright$  initialize cluster memberships
3:      $z^{(i)} \leftarrow \text{RandomInt}(1, K)$ 
4:   repeat
5:     for  $k \in 1 \dots K$  do                           $\triangleright$  recompute cluster centers
6:        $\boldsymbol{\nu}_k \leftarrow \frac{1}{\delta(z^{(i)}=k)} \sum_{i=1}^N \delta(z^{(i)} = k) \mathbf{x}^{(i)}$ 
7:     for  $i \in 1 \dots N$  do                       $\triangleright$  reassigned instances to nearest clusters
8:        $z^{(i)} \leftarrow \operatorname{argmin}_k \|\mathbf{x}^{(i)} - \boldsymbol{\nu}_k\|^2$ 
9:   until converged
10:  return  $\{z^{(i)}\}$                                  $\triangleright$  return cluster assignments

```

2475 a cluster assignment for each instance, and a central (“mean”) location for each cluster.
 2476 K -means iterates between updates to the assignments and the centers:

2477 1. each instance is placed in the cluster with the closest center;

2478 2. each center is recomputed as the average over points in the cluster.

2479 This is formalized in Algorithm 8. The term $\|\mathbf{x}^{(i)} - \boldsymbol{\nu}\|^2$ refers to the squared Euclidean
 2480 norm, $\sum_{j=1}^V (x_j^{(i)} - \nu_j)^2$.

2481 **Soft K -means** is a particularly relevant variant. Instead of directly assigning each
 2482 point to a specific cluster, soft K -means assigns each point a **distribution** over clusters
 2483 $\mathbf{q}^{(i)}$, so that $\sum_{k=1}^K q^{(i)}(k) = 1$, and $\forall_k, q^{(i)}(k) \geq 0$. The soft weight $q^{(i)}(k)$ is computed from
 2484 the distance of $\mathbf{x}^{(i)}$ to the cluster center $\boldsymbol{\nu}_k$. In turn, the center of each cluster is computed
 2485 from a **weighted average** of the points in the cluster,

$$\boldsymbol{\nu}_k = \frac{1}{\sum_{i=1}^N q^{(i)}(k)} \sum_{i=1}^N q^{(i)}(k) \mathbf{x}^{(i)}. \quad [5.1]$$

2486 We will now explore a probabilistic version of soft K -means clustering, based on **expectation**
 2487 **maximization** (EM). Because EM clustering can be derived as an approximation to
 2488 maximum-likelihood estimation, it can be extended in a number of useful ways.

2489 **5.1.2 Expectation Maximization (EM)**

Expectation maximization combines the idea of soft K -means with Naïve Bayes classification. To review, Naïve Bayes defines a probability distribution over the data,

$$\log p(\mathbf{x}, \mathbf{y}; \boldsymbol{\phi}, \boldsymbol{\mu}) = \sum_{i=1}^N \log \left(p(\mathbf{x}^{(i)} | y^{(i)}; \boldsymbol{\phi}) \times p(y^{(i)}; \boldsymbol{\mu}) \right) \quad [5.2]$$

Now suppose that you never observe the labels. To indicate this, we'll refer to the label of each instance as $z^{(i)}$, rather than $y^{(i)}$, which is usually reserved for observed variables. By marginalizing over the **latent** variables \mathbf{z} , we compute the marginal probability of the observed instances \mathbf{x} :

$$\log p(\mathbf{x}; \boldsymbol{\phi}, \boldsymbol{\mu}) = \sum_{i=1}^N \log p(\mathbf{x}^{(i)}; \boldsymbol{\phi}, \boldsymbol{\mu}) \quad [5.3]$$

$$= \sum_{i=1}^N \log \sum_{z=1}^K p(\mathbf{x}^{(i)}, z; \boldsymbol{\phi}, \boldsymbol{\mu}) \quad [5.4]$$

$$= \sum_{i=1}^N \log \sum_{z=1}^K p(\mathbf{x}^{(i)} | z; \boldsymbol{\phi}) \times p(z; \boldsymbol{\mu}). \quad [5.5]$$

2490 To estimate the parameters $\boldsymbol{\phi}$ and $\boldsymbol{\mu}$, we can maximize the marginal likelihood in Equa-
 2491 tion 5.5. Why is this the right thing to maximize? Without labels, discriminative learning
 2492 is impossible — there's nothing to discriminate. So maximum likelihood is all we have.

2493 When the labels are observed, we can estimate the parameters of the Naïve Bayes
 2494 probability model separately for each label. But marginalizing over the labels couples
 2495 these parameters, making direct optimization of $\log p(\mathbf{x})$ intractable. We will approximate
 2496 the log-likelihood by introducing an *auxiliary variable* $\mathbf{q}^{(i)}$, which is a distribution over the
 2497 label set $\mathcal{Z} = \{1, 2, \dots, K\}$. The optimization procedure will alternate between updates to
 2498 \mathbf{q} and updates to the parameters $(\boldsymbol{\phi}, \boldsymbol{\mu})$. Thus, $\mathbf{q}^{(i)}$ plays here as in soft K -means.

To derive the updates for this optimization, multiply the right side of Equation 5.5 by

the ratio $\frac{q^{(i)}(z)}{q^{(i)}(z)} = 1$,

$$\log p(\mathbf{x}; \phi, \mu) = \sum_{i=1}^M \log \sum_{z=1}^K p(\mathbf{x}^{(i)} | z; \phi) \times p(z; \mu) \times \frac{q^{(i)}(z)}{q^{(i)}(z)} \quad [5.6]$$

$$= \sum_{i=1}^M \log \sum_{z=1}^K q^{(i)}(z) \times p(\mathbf{x}^{(i)} | z; \phi) \times p(z; \mu) \times \frac{1}{q^{(i)}(z)} \quad [5.7]$$

$$= \sum_{i=1}^M \log E_{\mathbf{q}^{(i)}} \left[\frac{p(\mathbf{x}^{(i)} | z; \phi) p(z; \mu)}{q^{(i)}(z)} \right], \quad [5.8]$$

where $E_{\mathbf{q}^{(i)}} [f(z)] = \sum_{z=1}^K q^{(i)}(z) \times f(z)$ refers to the expectation of the function f under the distribution $z \sim \mathbf{q}^{(i)}$.

Jensen's inequality says that because \log is a concave function, we can push it inside the expectation, and obtain a lower bound.

$$\log p(\mathbf{x}; \phi, \mu) \geq \sum_{i=1}^N E_{\mathbf{q}^{(i)}} \left[\log \frac{p(\mathbf{x}^{(i)} | z; \phi) p(z; \mu)}{q^{(i)}(z)} \right] \quad [5.9]$$

$$J \triangleq \sum_{i=1}^N E_{\mathbf{q}^{(i)}} \left[\log p(\mathbf{x}^{(i)} | z; \phi) + \log p(z; \mu) - \log q^{(i)}(z) \right] \quad [5.10]$$

$$= \sum_{i=1}^N E_{\mathbf{q}^{(i)}} \left[\log p(\mathbf{x}^{(i)}, z; \phi, \mu) \right] + H(\mathbf{q}^{(i)}) \quad [5.11]$$

We will focus on Equation 5.10, which is the lower bound on the marginal log-likelihood of the observed data, $\log p(\mathbf{x})$. Equation 5.11 shows the connection to the information theoretic concept of **entropy**, $H(\mathbf{q}^{(i)}) = -\sum_{z=1}^K q^{(i)}(z) \log q^{(i)}(z)$, which measures the average amount of information produced by a draw from the distribution $q^{(i)}$. The lower bound J is a function of two groups of arguments:

- the distributions $\mathbf{q}^{(i)}$ for each instance;
- the parameters μ and ϕ .

The expectation-maximization (EM) algorithm maximizes the bound with respect to each of these arguments in turn, while holding the other fixed.

5.1.2.1 The E-step

The step in which we update $\mathbf{q}^{(i)}$ is known as the **E-step**, because it updates the distribution under which the expectation is computed. To derive this update, first write out the

expectation in the lower bound as a sum,

$$J = \sum_{i=1}^N \sum_{z=1}^K q^{(i)}(z) \left[\log p(\mathbf{x}^{(i)} | z; \boldsymbol{\phi}) + \log p(z; \boldsymbol{\mu}) - \log q^{(i)}(z) \right]. \quad [5.12]$$

When optimizing this bound, we must also respect a set of “sum-to-one” constraints, $\sum_{z=1}^K q^{(i)}(z) = 1$ for all i . Just as in Naïve Bayes, this constraint can be incorporated into a Lagrangian:

$$J_q = \sum_{i=1}^N \sum_{z=1}^K q^{(i)}(z) \left(\log p(\mathbf{x}^{(i)} | z; \boldsymbol{\phi}) + \log p(z; \boldsymbol{\mu}) - \log q^{(i)}(z) \right) + \lambda^{(i)} \left(1 - \sum_{z=1}^K q^{(i)}(z) \right), \quad [5.13]$$

where $\lambda^{(i)}$ is the Lagrange multiplier for instance i .

The Lagrangian is maximized by taking the derivative and solving for $q^{(i)}$:

$$\frac{\partial J_q}{\partial q^{(i)}(z)} = \log p(\mathbf{x}^{(i)} | z; \boldsymbol{\phi}) + \log p(z; \boldsymbol{\mu}) - \log q^{(i)}(z) - 1 - \lambda^{(i)} \quad [5.14]$$

$$\log q^{(i)}(z) = \log p(\mathbf{x}^{(i)} | z; \boldsymbol{\phi}) + \log p(z; \boldsymbol{\mu}) - 1 - \lambda^{(i)} \quad [5.15]$$

$$q^{(i)}(z) \propto p(\mathbf{x}^{(i)} | z; \boldsymbol{\phi}) \times p(z; \boldsymbol{\mu}). \quad [5.16]$$

Applying the sum-to-one constraint gives an exact solution,

$$q^{(i)}(z) = \frac{p(\mathbf{x}^{(i)} | z; \boldsymbol{\phi}) \times p(z; \boldsymbol{\mu})}{\sum_{z'=1}^K p(\mathbf{x}^{(i)} | z'; \boldsymbol{\phi}) \times p(z'; \boldsymbol{\mu})} \quad [5.17]$$

$$= p(z | \mathbf{x}^{(i)}; \boldsymbol{\phi}, \boldsymbol{\mu}). \quad [5.18]$$

After normalizing, each $q^{(i)}$ — which is the soft distribution over clusters for data $\mathbf{x}^{(i)}$ — is set to the posterior probability $p(z | \mathbf{x}^{(i)}; \boldsymbol{\phi}, \boldsymbol{\mu})$ under the current parameters. Although the Lagrange multipliers $\lambda^{(i)}$ were introduced as additional parameters, they drop out during normalization.

5.1.2.2 The M-step

Next, we hold fixed the soft assignments $q^{(i)}$, and maximize with respect to the parameters, $\boldsymbol{\phi}$ and $\boldsymbol{\mu}$. Let’s focus on the parameter $\boldsymbol{\phi}$, which parametrizes the likelihood $p(\mathbf{x} | z; \boldsymbol{\phi})$, and leave $\boldsymbol{\mu}$ for an exercise. The parameter $\boldsymbol{\phi}$ is a distribution over words for each cluster, so it is optimized under the constraint that $\sum_{j=1}^V \phi_{z,j} = 1$. To incorporate this

constraint, we introduce a set of Lagrange multipliers $\{\lambda_z\}_{z=1}^K$, and from the Lagrangian,

$$J_\phi = \sum_{i=1}^N \sum_{z=1}^K q^{(i)}(z) \left(\log p(\mathbf{x}^{(i)} | z; \boldsymbol{\phi}) + \log p(z; \mu) - \log q^{(i)}(z) \right) + \sum_{z=1}^K \lambda_z \left(1 - \sum_{j=1}^V \phi_{z,j} \right). \quad [5.19]$$

2517 The term $\log p(\mathbf{x}^{(i)} | z; \boldsymbol{\phi})$ is the conditional log-likelihood for the multinomial, which
2518 expands to,

$$\log p(\mathbf{x}^{(i)} | z, \boldsymbol{\phi}) = C + \sum_{j=1}^V x_j \log \phi_{z,j}, \quad [5.20]$$

2519 where C is a constant with respect to $\boldsymbol{\phi}$ — see Equation 2.12 in § 2.1 for more discussion
2520 of this probability function.

Setting the derivative of J_ϕ equal to zero,

$$\frac{\partial J_\phi}{\partial \phi_{z,j}} = \sum_{i=1}^N q^{(i)}(z) \times \frac{x_j^{(i)}}{\phi_{z,j}} - \lambda_z \quad [5.21]$$

$$\phi_{z,j} \propto \sum_{i=1}^N q^{(i)}(z) \times x_j^{(i)}. \quad [5.22]$$

Because ϕ_z is constrained to be a probability distribution, the exact solution is computed as,

$$\phi_{z,j} = \frac{\sum_{i=1}^N q^{(i)}(z) \times x_j^{(i)}}{\sum_{j'=1}^V \sum_{i=1}^N q^{(i)}(z) \times x_{j'}^{(i)}} = \frac{E_q [\text{count}(z, j)]}{\sum_{j'=1}^V E_q [\text{count}(z, j')]} \quad [5.23]$$

2521 where the counter $j \in \{1, 2, \dots, V\}$ indexes over base features, such as words.

2522 This update sets ϕ_z equal to the relative frequency estimate of the *expected counts* under
2523 the distribution q . As in supervised Naïve Bayes, we can smooth these counts by adding
2524 a constant α . The update for μ is similar: $\mu_z \propto \sum_{i=1}^N q^{(i)}(z) = E_q [\text{count}(z)]$, which is the
2525 expected frequency of cluster z . These probabilities can also be smoothed. In sum, the
2526 M-step is just like Naïve Bayes, but with expected counts rather than observed counts.

2527 The multinomial likelihood $p(\mathbf{x} | z)$ can be replaced with other probability distribu-
2528 tions: for example, for continuous observations, a Gaussian distribution can be used. In
2529 some cases, there is no closed-form update to the parameters of the likelihood. One ap-
2530 proach is to run gradient-based optimization at each M-step; another is to simply take a
2531 single step along the gradient step and then return to the E-step (Berg-Kirkpatrick et al.,
2532 2010).

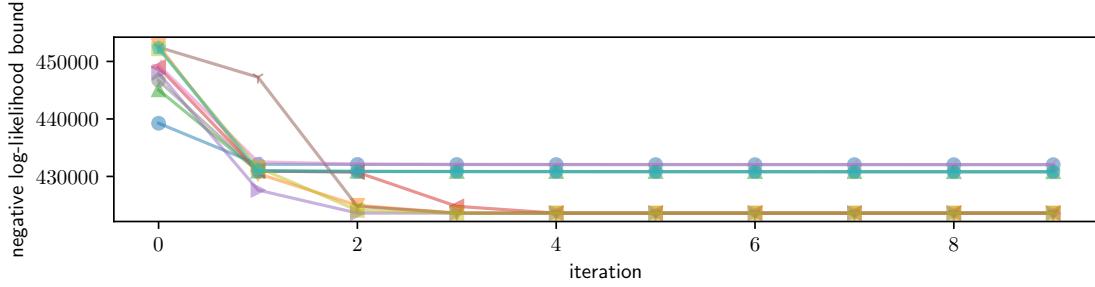


Figure 5.2: Sensitivity of expectation maximization to initialization. Each line shows the progress of optimization from a different random initialization.

2533 5.1.3 EM as an optimization algorithm

2534 Algorithms that alternate between updating subsets of the parameters are called **coordi-**
 2535 **nate ascent** algorithms. The objective J (the lower bound on the marginal likelihood of
 2536 the data) is separately convex in q and (μ, ϕ) , but it is not jointly convex in all terms; this
 2537 condition is known as **biconvexity**. Each step of the expectation-maximization algorithm
 2538 is guaranteed not to decrease the lower bound J , which means that EM will converge
 2539 towards a solution at which no nearby points yield further improvements. This solution
 2540 is a **local optimum** — it is as good or better than any of its immediate neighbors, but is
 2541 *not* guaranteed to be optimal among all possible configurations of (q, μ, ϕ) .

2542 The fact that there is no guarantee of global optimality means that initialization is
 2543 important: where you start can determine where you finish. To illustrate this point,
 2544 Figure 5.2 shows the objective function for EM with ten different random initializations:
 2545 while the objective function improves monotonically in each run, it converges to several
 2546 different values.¹ For the convex objectives that we encountered in chapter 2, it was not
 2547 necessary to worry about initialization, because gradient-based optimization guaranteed
 2548 to reach the global minimum. But in expectation-maximization — and in the deep neural
 2549 networks from chapter 3 — initialization matters.

2550 In **hard EM**, each $q^{(i)}$ distribution assigns probability of 1 to a single label $\hat{z}^{(i)}$, and zero
 2551 probability to all others (Neal and Hinton, 1998). This is similar in spirit to K -means clus-
 2552 tering, and can outperform standard EM in some cases (Spitkovsky et al., 2010). Another
 2553 variant of expectation maximization incorporates stochastic gradient descent (SGD): after
 2554 performing a local E-step at each instance $x^{(i)}$, we immediately make a gradient update
 2555 to the parameters (μ, ϕ) . This algorithm has been called **incremental expectation maxi-**
 2556 **mization** (Neal and Hinton, 1998) and **online expectation maximization** (Sato and Ishii,
 2557 2000; Cappé and Moulines, 2009), and is especially useful when there is no closed-form

¹The figure shows the upper bound on the negative log-likelihood, because optimization is typically framed as minimization rather than maximization.

2558 optimum for the likelihood $p(\mathbf{x} \mid z)$, and in online settings where new data is constantly
 2559 streamed in (see Liang and Klein, 2009, for a comparison for online EM variants).

2560 **5.1.4 How many clusters?**

2561 So far, we have assumed that the number of clusters K is given. In some cases, this as-
 2562 sumption is valid. For example, a lexical semantic resource like WordNet might define the
 2563 number of senses for a word. In other cases, the number of clusters could be a parameter
 2564 for the user to tune: some readers want a coarse-grained clustering of news stories into
 2565 three or four clusters, while others want a fine-grained clustering into twenty or more.
 2566 But many times there is little extrinsic guidance for how to choose K .

2567 One solution is to choose the number of clusters to maximize a metric of clustering
 2568 quality. The other parameters μ and ϕ are chosen to maximize the log-likelihood bound
 2569 J , so this might seem a potential candidate for tuning K . However, J will never decrease
 2570 with K : if it is possible to obtain a bound of J_K with K clusters, then it is always possible
 2571 to do at least as well with $K + 1$ clusters, by simply ignoring the additional cluster and
 2572 setting its probability to zero in q and μ . It is therefore necessary to introduce a penalty
 2573 for model complexity, so that fewer clusters are preferred. For example, the Akaike Infor-
 2574 mation Crition (AIC; Akaike, 1974) is the linear combination of the number of parameters
 2575 and the log-likelihood,

$$\text{AIC} = 2M - 2J, \quad [5.24]$$

2576 where M is the number of parameters. In an expectation-maximization clustering algo-
 2577 rithm, $M = K \times V + K$. Since the number of parameters increases with the number of
 2578 clusters K , the AIC may prefer more parsimonious models, even if they do not fit the data
 2579 quite as well.

2580 Another choice is to maximize the **predictive likelihood** on heldout data. This data
 2581 is not used to estimate the model parameters ϕ and μ , and so it is not the case that the
 2582 likelihood on this data is guaranteed to increase with K . Figure 5.3 shows the negative
 2583 log-likelihood on training and heldout data, as well as the AIC.

2584 ***Bayesian nonparametrics** An alternative approach is to treat the number of clusters as
 2585 another latent variable. This requires statistical inference over a set of models with a vari-
 2586 able number of clusters. This is not possible within the framework of expectation max-
 2587 imization, but there are several alternative inference procedures which can be applied,
 2588 including **Markov Chain Monte Carlo (MCMC)**, which is briefly discussed in § 5.5 (for
 2589 more details, see Chapter 25 of Murphy, 2012). Bayesian nonparametrics have been ap-
 2590 plied to the problem of unsupervised word sense induction, learning not only the word
 2591 senses but also the number of senses per word (Reisinger and Mooney, 2010).




Figure 5.3: The negative log-likelihood and AIC for several runs of expectation maximization, on synthetic data. Although the data was generated from a model with $K = 10$, the optimal number of clusters is $\hat{K} = 15$, according to AIC and the heldout log-likelihood. The training set log-likelihood continues to improve as K increases.

2592 5.2 Applications of expectation-maximization

2593 EM is not really an “algorithm” like, say, quicksort. Rather, it is a framework for learning
2594 with missing data. The recipe for using EM on a problem of interest is:

- 2595 • Introduce latent variables z , such that it is easy to write the probability $P(x, z)$. It
2596 should also be easy to estimate the associated parameters, given knowledge of z .
- 2597 • Derive the E-step updates for $q(z)$, which is typically factored as $q(z) = \prod_{i=1}^N q_{z^{(i)}}(z^{(i)})$,
2598 where i is an index over instances.
- 2599 • The M-step updates typically correspond to the soft version of a probabilistic super-
2600 vised learning algorithm, like Naïve Bayes.

2601 This section discusses a few of the many applications of this general framework.

2602 5.2.1 Word sense induction

2603 The chapter began by considering the problem of word sense disambiguation when the
2604 senses are not known in advance. Expectation-maximization can be applied to this prob-
2605 lem by treating each cluster as a word sense. Each instance represents the use of an
2606 ambiguous word, and $x^{(i)}$ is a vector of counts for the other words that appear nearby:
2607 Schütze (1998) uses all words within a 50-word window. The probability $p(x^{(i)} | z)$ can be
2608 set to the multinomial distribution, as in Naïve Bayes. The EM algorithm can be applied
2609 directly to this data, yielding clusters that (hopefully) correspond to the word senses.

Better performance can be obtained by first applying truncated **singular value decom-
position (SVD)** to the matrix of context-counts $C_{ij} = \text{count}(i, j)$, where $\text{count}(i, j)$ is the

count of word j in the context of instance i . Truncated singular value decomposition approximates the matrix \mathbf{C} as a product of three matrices, $\mathbf{U}, \mathbf{S}, \mathbf{V}$, under the constraint that \mathbf{U} and \mathbf{V} are orthonormal, and \mathbf{S} is diagonal:

$$\begin{aligned} & \min_{\mathbf{U}, \mathbf{S}, \mathbf{V}} \|\mathbf{C} - \mathbf{USV}^\top\|_F \\ & \text{s.t. } \mathbf{U} \in \mathbb{R}^{V \times K}, \mathbf{UU}^\top = \mathbb{I} \\ & \quad \mathbf{S} = \text{Diag}(s_1, s_2, \dots, s_K) \\ & \quad \mathbf{V}^\top \in \mathbb{R}^{N_p \times K}, \mathbf{VV}^\top = \mathbb{I}, \end{aligned} \quad [5.25]$$

where $\|\cdot\|_F$ is the Frobenius norm, $\|X\|_F = \sqrt{\sum_{i,j} X_{i,j}^2}$. The matrix \mathbf{U} contains the left singular vectors of \mathbf{C} , and the rows of this matrix can be used as low-dimensional representations of the count vectors \mathbf{c}_i . EM clustering can be made more robust by setting the instance descriptions $\mathbf{x}^{(i)}$ equal to these rows, rather than using raw counts (Schütze, 1998). However, because the instances are now dense vectors of continuous numbers, the probability $p(\mathbf{x}^{(i)} | z)$ must be defined as a multivariate Gaussian distribution.

In truncated singular value decomposition, the hyperparameter K is the truncation limit: when K is equal to the rank of \mathbf{C} , the norm of the difference between the original matrix \mathbf{C} and its reconstruction \mathbf{USV}^\top will be zero. Lower values of K increase the reconstruction error, but yield vector representations that are smaller and easier to learn from. Singular value decomposition is discussed in more detail in chapter 14.

5.2.2 Semi-supervised learning

Expectation-maximization can also be applied to the problem of **semi-supervised learning**: learning from both labeled and unlabeled data in a single model. Semi-supervised learning makes use of ground truth annotations, ensuring that each label y corresponds to the desired concept. By adding unlabeled data, it is possible cover a greater fraction of the features than would be possible using labeled data alone. Other methods for semi-supervised learning are discussed in § 5.3, but for now, let's approach the problem within the framework of expectation-maximization (Nigam et al., 2000).

Suppose we have labeled data $\{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^{N_\ell}$, and unlabeled data $\{\mathbf{x}^{(i)}\}_{i=N_\ell+1}^{N_\ell+N_u}$, where N_ℓ is the number of labeled instances and N_u is the number of unlabeled instances. We can learn from the combined data by maximizing a lower bound on the joint log-likelihood,

$$\mathcal{L} = \sum_{i=1}^{N_\ell} \log p(\mathbf{x}^{(i)}, y^{(i)}; \boldsymbol{\mu}, \boldsymbol{\phi}) + \sum_{j=N_\ell+1}^{N_\ell+N_u} \log p(\mathbf{x}^{(j)}; \boldsymbol{\mu}, \boldsymbol{\phi}) \quad [5.26]$$

$$= \sum_{i=1}^{N_\ell} \left(\log p(\mathbf{x}^{(i)} | y^{(i)}; \boldsymbol{\phi}) + \log p(y^{(i)}; \boldsymbol{\mu}) \right) + \sum_{j=N_\ell+1}^{N_\ell+N_u} \log \sum_{y=1}^K p(\mathbf{x}^{(j)}, y; \boldsymbol{\mu}, \boldsymbol{\phi}). \quad [5.27]$$

Algorithm 9 Generative process for the Naïve Bayes classifier with hidden components

for Document $i \in \{1, 2, \dots, N\}$ **do**:

Draw the label $y^{(i)} \sim \text{Categorical}(\mu)$;

Draw the component $z^{(i)} \sim \text{Categorical}(\beta_{y^{(i)}})$;

Draw the word counts $x^{(i)} | y^{(i)}, z^{(i)} \sim \text{Multinomial}(\phi_{z^{(i)}})$.

2629 The left sum is identical to the objective in Naïve Bayes; the right sum is the marginal log-
 2630 likelihood for expectation-maximization clustering, from Equation 5.5. We can construct a
 2631 lower bound on this log-likelihood by introducing distributions $q^{(j)}$ for all $j \in \{N_\ell + 1, \dots, N_\ell + N_u\}$.
 2632 The E-step updates these distributions; the M-step updates the parameters ϕ and μ , us-
 2633 ing the expected counts from the unlabeled data and the observed counts from the labeled
 2634 data.

2635 A critical issue in semi-supervised learning is how to balance the impact of the labeled
 2636 and unlabeled data on the classifier weights, especially when the unlabeled data is much
 2637 larger than the labeled dataset. The risk is that the unlabeled data will dominate, caus-
 2638 ing the parameters to drift towards a “natural clustering” of the instances — which may
 2639 not correspond to a good classifier for the labeled data. One solution is to heuristically
 2640 reweight the two components of Equation 5.26, tuning the weight of the two components
 2641 on a heldout development set (Nigam et al., 2000).

2642 **5.2.3 Multi-component modeling**

2643 As a final application, let’s return to fully supervised classification. A classic dataset for
 2644 text classification is 20 newsgroups, which contains posts to a set of online forums, called
 2645 newsgroups. One of the newsgroups is `comp.sys.mac.hardware`, which discusses Ap-
 2646 ple computing hardware. Suppose that within this newsgroup there are two kinds of
 2647 posts: reviews of new hardware, and question-answer posts about hardware problems.
 2648 The language in these *components* of the `mac.hardware` class might have little in com-
 2649 mon; if so, it would be better to model these components separately, rather than treating
 2650 their union as a single class. However, the component responsible for each instance is not
 2651 directly observed.

2652 Recall that Naïve Bayes is based on a generative process, which provides a stochastic
 2653 explanation for the observed data. In Naïve Bayes, each label is drawn from a categorical
 2654 distribution with parameter μ , and each vector of word counts is drawn from a multi-
 2655 nomial distribution with parameter ϕ_y . For multi-component modeling, we envision a
 2656 slightly different generative process, incorporating both the observed label $y^{(i)}$ and the
 2657 latent component $z^{(i)}$. This generative process is shown in Algorithm 9. A new parameter
 2658 $\beta_{y^{(i)}}$ defines the distribution of components, conditioned on the label $y^{(i)}$. The component,
 2659 and not the class label, then parametrizes the distribution over words.

-
- (5.1) ☺ Villeneuve a bel et bien **réussi** son pari de changer de perspectives tout en assurant une cohérence à la franchise.²
- (5.2) ☺ Il est également trop **long** et bancal dans sa narration, tiède dans ses intentions, et tirailé entre deux personnages et directions qui ne parviennent pas à coexister en harmonie.³
- (5.3) Denis Villeneuve a **réussi** une suite **parfaitemment** maîtrisée⁴
- (5.4) **Long, bavard**, hyper design, à peine agité (le comble de l'action : une bagarre dans la flotte), métaphysique et, surtout, ennuyeux jusqu'à la catalepsie.⁵
- (5.5) Une suite d'une écrasante puissance, mêlant **parfaitemment** le contemplatif au narratif.⁶
- (5.6) Le film impitoyablement **bavard** finit quand même par se taire quand se lève l'espèce de bouquet final où semble se déchaîner, comme en libre parcours de poulets décapiés, l'armée des graphistes numériques griffant nerveusement la palette graphique entre agonie et orgasme.⁷

Table 5.1: Labeled and unlabeled reviews of the films *Blade Runner 2049* and *Transformers: The Last Knight*.

The labeled data includes $(\mathbf{x}^{(i)}, y^{(i)})$, but not $z^{(i)}$, so this is another case of missing data. Again, we sum over the missing data, applying Jensen's inequality to as to obtain a lower bound on the log-likelihood,

$$\log p(\mathbf{x}^{(i)}, y^{(i)}) = \log \sum_{z=1}^{K_z} p(\mathbf{x}^{(i)}, y^{(i)}, z; \boldsymbol{\mu}, \boldsymbol{\phi}, \boldsymbol{\beta}) \quad [5.28]$$

$$\geq \log p(y^{(i)}; \boldsymbol{\mu}) + E_{q_{Z|Y}^{(i)}} [\log p(\mathbf{x}^{(i)} | z; \boldsymbol{\phi}) + \log p(z | y^{(i)}; \boldsymbol{\beta}) - \log q^{(i)}(z)]. \quad [5.29]$$

We are now ready to apply expectation maximization. As usual, the E-step updates the distribution over the missing data, $q_{Z|Y}^{(i)}$. The M-step updates the parameters,

$$\beta_{y,z} = \frac{E_q [\text{count}(y, z)]}{\sum_{z'=1}^{K_z} E_q [\text{count}(y, z')]} \quad [5.30]$$

$$\phi_{z,j} = \frac{E_q [\text{count}(z, j)]}{\sum_{j'=1}^V E_q [\text{count}(z, j')]} \quad [5.31]$$

2660 5.3 Semi-supervised learning

2661 In semi-supervised learning, the learner makes use of both labeled and unlabeled data.
 2662 To see how this could help, suppose you want to do sentiment analysis in French. In Ta-

ble 5.1, there are two labeled examples, one positive and one negative. From this data, a learner could conclude that *réussi* is positive and *long* is negative. This isn't much! However, we can propagate this information to the unlabeled data, and potentially learn more.

- If we are confident that *réussi* is positive, then we might guess that (5.3) is also positive.
- That suggests that *parfaitement* is also positive.
- We can then propagate this information to (5.5), and learn from this words in this example.
- Similarly, we can propagate from the labeled data to (5.4), which we guess to be negative because it shares the word *long*. This suggests that *bavard* is also negative, which we propagate to (5.6).

Instances (5.3) and (5.4) were “similar” to the labeled examples for positivity and negativity, respectively. By using these instances to expand the models for each class, it became possible to correctly label instances (5.5) and (5.6), which didn't share any important features with the original labeled data. This requires a key assumption: that similar instances will have similar labels.

In § 5.2.2, we discussed how expectation maximization can be applied to semi-supervised learning. Using the labeled data, the initial parameters ϕ would assign a high weight for *réussi* in the positive class, and a high weight for *long* in the negative class. These weights helped to shape the distributions q for instances (5.3) and (5.4) in the E-step. In the next iteration of the M-step, the parameters ϕ are updated with counts from these instances, making it possible to correctly label the instances (5.5) and (5.6).

However, expectation-maximization has an important disadvantage: it requires using a generative classification model, which restricts the features that can be used for classification. In this section, we explore non-probabilistic approaches, which impose fewer restrictions on the classification model.

5.3.1 Multi-view learning

EM semi-supervised learning can be viewed as **self-training**: the labeled data guides the initial estimates of the classification parameters; these parameters are used to compute a label distribution over the unlabeled instances, $q^{(i)}$; the label distributions are used to update the parameters. The risk is that self-training drifts away from the original labeled data. This problem can be ameliorated by **multi-view learning**. Here we take the assumption that the features can be decomposed into multiple “views”, each of which is conditionally independent, given the label. For example, consider the problem of classifying a name as a person or location: one view is the name itself; another is the context in which it appears. This situation is illustrated in Table 5.2.

	$\mathbf{x}^{(1)}$	$\mathbf{x}^{(2)}$	y
1.	Peachtree Street	located on	LOC
2.	Dr. Walker	said	PER
3.	Zanzibar	located in	? → LOC
4.	Zanzibar	flew to	? → LOC
5.	Dr. Robert	recommended	? → PER
6.	Oprah	recommended	? → PER

Table 5.2: Example of multiview learning for named entity classification

2699 **Co-training** is an iterative multi-view learning algorithm, in which there are separate
 2700 classifiers for each view (Blum and Mitchell, 1998). At each iteration of the algorithm, each
 2701 classifier predicts labels for a subset of the unlabeled instances, using only the features
 2702 available in its view. These predictions are then used as ground truth to train the classifiers
 2703 associated with the other views. In the example shown in Table 5.2, the classifier on $\mathbf{x}^{(1)}$
 2704 might correctly label instance #5 as a person, because of the feature *Dr*; this instance would
 2705 then serve as training data for the classifier on $\mathbf{x}^{(2)}$, which would then be able to correctly
 2706 label instance #6, thanks to the feature *recommended*. If the views are truly independent,
 2707 this procedure is robust to drift. Furthermore, it imposes no restrictions on the classifiers
 2708 that can be used for each view.

2709 Word-sense disambiguation is particularly suited to multi-view learning, thanks to the
 2710 heuristic of “one sense per discourse”: if a polysemous word is used more than once in
 2711 a given text or conversation, all usages refer to the same sense (Gale et al., 1992). This
 2712 motivates a multi-view learning approach, in which one view corresponds to the local
 2713 context (the surrounding words), and another view corresponds to the global context at
 2714 the document level (Yarowsky, 1995). The local context view is first trained on a small
 2715 seed dataset. We then identify its most confident predictions on unlabeled instances. The
 2716 global context view is then used to extend these confident predictions to other instances
 2717 within the same documents. These new instances are added to the training data to the
 2718 local context classifier, which is retrained and then applied to the remaining unlabeled
 2719 data.

2720 5.3.2 Graph-based algorithms

2721 Another family of approaches to semi-supervised learning begins by constructing a graph,
 2722 in which pairs of instances are linked with symmetric weights $\omega_{i,j}$, e.g.,

$$\omega_{i,j} = \exp(-\alpha \times \|\mathbf{x}^{(i)} - \mathbf{x}^{(j)}\|^2). \quad [5.32]$$

2723 The goal is to use this weighted graph to propagate labels from a small set of labeled
 2724 instances to larger set of unlabeled instances.

2725 In **label propagation**, this is done through a series of matrix operations (Zhu et al.,
 2726 Let \mathbf{Q} be a matrix of size $N \times K$, in which each row $\mathbf{q}^{(i)}$ describes the labeling
 2727 of instance i . When ground truth labels are available, then $\mathbf{q}^{(i)}$ is an indicator vector,
 2728 with $q_{y^{(i)}}^{(i)} = 1$ and $q_{y' \neq y^{(i)}}^{(i)} = 0$. Let us refer to the submatrix of rows containing labeled
 2729 instances as \mathbf{Q}_L , and the remaining rows as \mathbf{Q}_U . The rows of \mathbf{Q}_U are initialized to assign
 2730 equal probabilities to all labels, $q_{i,k} = \frac{1}{K}$.

2731 Now, let $T_{i,j}$ represent the “transition” probability of moving from node j to node i ,

$$T_{i,j} \triangleq \Pr(j \rightarrow i) = \frac{\omega_{i,j}}{\sum_{k=1}^N \omega_{k,j}}. \quad [5.33]$$

We compute values of $T_{i,j}$ for all instances j and all *unlabeled* instances i , forming a matrix of size $N_U \times N$. If the dataset is large, this matrix may be expensive to store and manipulate; a solution is to sparsify it, by keeping only the κ largest values in each row, and setting all other values to zero. We can then “propagate” the label distributions to the unlabeled instances,

$$\tilde{\mathbf{Q}}_U \leftarrow \mathbf{T}\mathbf{Q} \quad [5.34]$$

$$\mathbf{s} \leftarrow \tilde{\mathbf{Q}}_U \mathbf{1} \quad [5.35]$$

$$\mathbf{Q}_U \leftarrow \text{Diag}(\mathbf{s})^{-1} \tilde{\mathbf{Q}}_U. \quad [5.36]$$

2732 The expression $\tilde{\mathbf{Q}}_U \mathbf{1}$ indicates multiplication of $\tilde{\mathbf{Q}}_U$ by a column vector of ones, which is
 2733 equivalent to computing the sum of each row of $\tilde{\mathbf{Q}}_U$. The matrix $\text{Diag}(\mathbf{s})$ is a diagonal
 2734 matrix with the elements of \mathbf{s} on the diagonals. The product $\text{Diag}(\mathbf{s})^{-1} \tilde{\mathbf{Q}}_U$ has the effect
 2735 of normalizing the rows of $\tilde{\mathbf{Q}}_U$, so that each row of \mathbf{Q}_U is a probability distribution over
 2736 labels.

2737 5.4 Domain adaptation

2738 In many practical scenarios, the labeled data differs in some key respect from the data
 2739 to which the trained model is to be applied. A classic example is in consumer reviews:
 2740 we may have labeled reviews of movies (the **source domain**), but we want to predict the
 2741 reviews of appliances (the **target domain**). A similar issues arise with genre differences:
 2742 most linguistically-annotated data is news text, but application domains range from social
 2743 media to electronic health records. In general, there may be several source and target
 2744 domains, each with their own properties; however, for simplicity, this discussion will
 2745 focus mainly on the case of a single source and target domain.

2746 The simplest approach is “direct transfer”: train a classifier on the source domain,
 2747 and apply it directly to the target domain. The accuracy of this approach depends on the
 2748 extent to which features are shared across domains. In review text, words like *outstanding*

and *disappointing* will apply across both movies and appliances; but others, like *terrifying*, may have meanings that are domain-specific. **Domain adaptation** algorithms attempt to do better than direct transfer, by learning from data in both domains. There are two main families of domain adaptation algorithms, depending on whether any labeled data is available in the target domain.

5.4.1 Supervised domain adaptation

In supervised domain adaptation, there is a small amount of labeled data in the target domain, and a large amount of data in the source domain. The simplest approach would be to ignore domain differences, and simply merge the training data from the source and target domains. There are several other baseline approaches to dealing with this scenario (Daumé III, 2007):

Interpolation. Train a classifier for each domain, and combine their predictions. For example,

$$\hat{y} = \operatorname{argmax}_y \lambda_s \Psi_s(\mathbf{x}, y) + (1 - \lambda_s) \Psi_t(\mathbf{x}, y), \quad [5.37]$$

where Ψ_s and Ψ_t are the scoring functions from the source and target domain classifiers respectively, and λ_s is the interpolation weight.

Prediction. Train a classifier on the source domain data, use its prediction as an additional feature in a classifier trained on the target domain data.

Priors. Train a classifier on the source domain data, and use its weights as a prior distribution on the weights of the classifier for the target domain data. This is equivalent to regularizing the target domain weights towards the weights of the source domain classifier (Chelba and Acero, 2006),

$$\ell(\boldsymbol{\theta}_t) = \sum_{i=1}^N \ell^{(i)}(\mathbf{x}^{(i)}, y^{(i)}; \boldsymbol{\theta}_t) + \lambda \|\boldsymbol{\theta}_t - \boldsymbol{\theta}_s\|_2^2, \quad [5.38]$$

where $\ell^{(i)}$ is the prediction loss on instance i , and λ is the regularization weight.

An effective and “frustratingly simple” alternative is EasyAdapt (Daumé III, 2007), which creates copies of each feature: one for each domain and one for the cross-domain setting. For example, a negative review of the film *Wonder Woman* begins, *As boring and flavorless as a three-day-old grilled cheese sandwich....*⁸ The resulting bag-of-words feature

⁸<http://www.colesmithey.com/capsules/2017/06/wonder-woman.HTML>, accessed October 9, 2017.

vector would be,

$$\begin{aligned} \mathbf{f}(\mathbf{x}, y, d) = & \{(boring, -, \text{MOVIE}) : 1, (boring, -, *) : 1, \\ & (flavorless, -, \text{MOVIE}) : 1, (flavorless, -, *) : 1, \\ & (three-day-old, -, \text{MOVIE}) : 1, (three-day-old, -, *) : 1, \\ & \dots\}, \end{aligned}$$

with $(boring, -, \text{MOVIE})$ indicating the word *boring* appearing in a negative labeled document in the MOVIE domain, and $(boring, -, *)$ indicating the same word in a negative labeled document in *any* domain. It is up to the learner to allocate weight between the domain-specific and cross-domain features: for words that facilitate prediction in both domains, the learner will use the cross-domain features; for words that are relevant only to a single domain, the domain-specific features will be used. Any discriminative classifier can be used with these augmented features.⁹

5.4.2 Unsupervised domain adaptation

In unsupervised domain adaptation, there is no labeled data in the target domain. Unsupervised domain adaptation algorithms cope with this problem by trying to make the data from the source and target domains as similar as possible. This is typically done by learning a **projection function**, which puts the source and target data in a shared space, in which a learner can generalize across domains. This projection is learned from data in both domains, and is applied to the base features — for example, the bag-of-words in text classification. The projected features can then be used both for training and for prediction.

5.4.2.1 Linear projection

In linear projection, the cross-domain representation is constructed by a matrix-vector product,

$$\mathbf{g}(\mathbf{x}^{(i)}) = \mathbf{U}\mathbf{x}^{(i)}. \quad [5.39]$$

The projected vectors $\mathbf{g}(\mathbf{x}^{(i)})$ can then be used as base features during both training (from the source domain) and prediction (on the target domain).

The projection matrix \mathbf{U} can be learned in a number of different ways, but many approaches focus on compressing and reconstructing the base features (Ando and Zhang, 2005). For example, we can define a set of **pivot features**, which are typically chosen because they appear in both domains: in the case of review documents, pivot features might include evaluative adjectives like *outstanding* and *disappointing* (Blitzer et al., 2007). For each pivot feature j , we define an auxiliary problem of predicting whether the feature is

⁹EasyAdapt can be explained as a hierarchical Bayesian model, in which the weights for each domain are drawn from a shared prior (Finkel and Manning, 2009).

present in each example, using the remaining base features. Let ϕ_j denote the weights of this classifier, and us horizontally concatenate the weights for each of the N_p pivot features into a matrix $\Phi = [\phi_1, \phi_2, \dots, \phi_{N_p}]$.

We then perform truncated singular value decomposition on Φ , as described in § 5.2.1, obtaining $\Phi \approx \mathbf{U}\mathbf{S}\mathbf{V}^\top$. The rows of the matrix \mathbf{U} summarize information about each base feature: indeed, the truncated singular value decomposition identifies a low-dimension basis for the weight matrix Φ , which in turn links base features to pivot features. Suppose that a base feature *reliable* occurs only in the target domain of appliance reviews. Nonetheless, it will have a positive weight towards some pivot features (e.g., *outstanding*, *recommended*), and a negative weight towards others (e.g., *worthless*, *unpleasant*). A base feature such as *watchable* might have the same associations with the pivot features, and therefore, $\mathbf{u}_{\text{reliable}} \approx \mathbf{u}_{\text{watchable}}$. The matrix \mathbf{U} can thus project the base features into a space in which this information is shared.

5.4.2.2 Non-linear projection

Non-linear transformations of the base features can be accomplished by implementing the transformation function as a deep neural network, which is trained from an auxiliary objective.

Denoising objectives One possibility is to train a projection function to reconstruct a corrupted version of the original input. The original input can be corrupted in various ways: by the addition of random noise (Glorot et al., 2011; Chen et al., 2012), or by the deletion of features (Chen et al., 2012; Yang and Eisenstein, 2015). Denoising objectives share many properties of the linear projection method described above: they enable the projection function to be trained on large amounts of unlabeled data from the target domain, and allow information to be shared across the feature space, thereby reducing sensitivity to rare and domain-specific features.

Adversarial objectives The ultimate goal is for the transformed representations $\mathbf{g}(\mathbf{x}^{(i)})$ to be domain-general. This can be made an explicit optimization criterion by computing the similarity of transformed instances both within and between domains (Tzeng et al., 2015), or by formulating an auxiliary classification task, in which the domain itself is treated as a label (Ganin et al., 2016). This setting is **adversarial**, because we want to learn a representation that makes this classifier perform poorly. At the same time, we want $\mathbf{g}(\mathbf{x}^{(i)})$ to enable accurate predictions of the labels $y^{(i)}$.

To formalize this idea, let $d^{(i)}$ represent the domain of instance i , and let $\ell_d(\mathbf{g}(\mathbf{x}^{(i)}), d^{(i)}; \theta_d)$ represent the loss of a classifier (typically a deep neural network) trained to predict $d^{(i)}$ from the transformed representation $\mathbf{g}(\mathbf{x}^{(i)})$, using parameters θ_d . Analogously, let $\ell_y(\mathbf{g}(\mathbf{x}^{(i)}), y^{(i)}; \theta_y)$ represent the loss of a classifier trained to predict the label $y^{(i)}$ from $\mathbf{g}(\mathbf{x}^{(i)})$, using param-

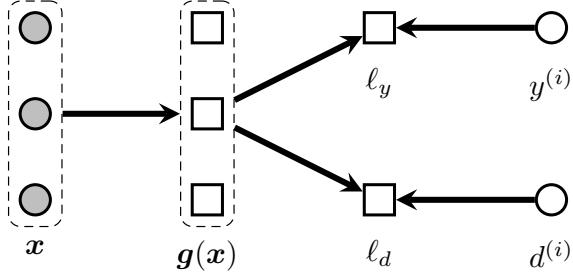


Figure 5.4: A schematic view of adversarial domain adaptation. The loss ℓ_y is computed only for instances from the source domain, where labels $y^{(i)}$ are available.

eters θ_y . The transformation g can then be trained from two criteria: it should yield accurate predictions of the labels $y^{(i)}$, while making *inaccurate* predictions of the domains $d^{(i)}$. This can be formulated as a joint optimization problem,

$$\min_{f, \theta_g, \theta_y, \theta_d} \sum_{i=1}^{N_\ell+N_u} \ell_d(g(\mathbf{x}^{(i)}; \theta_g), d^{(i)}; \theta_d) - \sum_{i=1}^{N_\ell} \ell_y(g(\mathbf{x}^{(i)}), y^{(i)}; \theta_y), \quad [5.40]$$

where N_ℓ is the number of labeled instances and N_u is the number of unlabeled instances, with the labeled instances appearing first in the dataset. This setup is shown in Figure 5.4. The loss can be optimized by stochastic gradient descent, jointly training the parameters of the non-linear transformation θ_g , and the parameters of the prediction models θ_d and θ_y .

5.5 *Other approaches to learning with latent variables

Expectation maximization provides a general approach to learning with latent variables, but it has limitations. One is the sensitivity to initialization; in practical applications, considerable attention may need to be devoted to finding a good initialization. A second issue is that EM tends to be easiest to apply in cases where the latent variables have a clear decomposition (in the cases we have considered, they decompose across the instances). For these reasons, it is worth briefly considering some alternatives to EM.

5.5.1 Sampling

In EM clustering, there is a distribution $q^{(i)}$ for the missing data related to each instance. The M-step consists of updating the parameters of this distribution. An alternative is to draw samples of the latent variables. If the sampling distribution is designed correctly, this procedure will eventually converge to drawing samples from the true posterior over the missing data, $p(z^{(1:N_z)} | \mathbf{x}^{(1:N_x)})$. For example, in the case of clustering, the missing

2854 data $\mathbf{z}^{(1:N_z)}$ is the set of cluster memberships, $\mathbf{y}^{(1:N)}$, so we draw samples from the pos-
 2855 terior distribution over clusterings of the data. If a single clustering is required, we can
 2856 select the one with the highest conditional likelihood, $\hat{\mathbf{z}} = \operatorname{argmax}_{\mathbf{z}} p(\mathbf{z}^{(1:N_z)} | \mathbf{x}^{(1:N_x)})$.

This general family of algorithms is called **Markov Chain Monte Carlo (MCMC)**: “Monte Carlo” because it is based on a series of random draws; “Markov Chain” because the sampling procedure must be designed such that each sample depends only on the previous sample, and not on the entire sampling history. **Gibbs sampling** is an MCMC algorithm in which each latent variable is sampled from its posterior distribution,

$$\mathbf{z}^{(n)} | \mathbf{x}, \mathbf{z}^{(-n)} \sim p(\mathbf{z}^{(n)} | \mathbf{x}, \mathbf{z}^{(-n)}), \quad [5.41]$$

where $\mathbf{z}^{(-n)}$ indicates $\{\mathbf{z} \setminus \mathbf{z}^{(n)}\}$, the set of all latent variables except for $\mathbf{z}^{(n)}$. Repeatedly drawing samples over all latent variables constructs a Markov chain, and which is guaranteed to converge to a sequence of samples from, $p(\mathbf{z}^{(1:N_z)} | \mathbf{x}^{(1:N_x)})$. In probabilistic clustering, the sampling distribution has the following form,

$$p(\mathbf{z}^{(i)} | \mathbf{x}, \mathbf{z}^{(-i)}) = \frac{p(\mathbf{x}^{(i)} | z^{(i)}; \boldsymbol{\phi}) \times p(z^{(i)}; \boldsymbol{\mu})}{\sum_{z=1}^K p(\mathbf{x}^{(i)} | z; \boldsymbol{\phi}) \times p(z; \boldsymbol{\mu})} \quad [5.42]$$

$$\propto \text{Multinomial}(\mathbf{x}^{(i)}; \boldsymbol{\phi}_{z^{(i)}}) \times \boldsymbol{\mu}_{z^{(i)}}. \quad [5.43]$$

2857 In this case, the sampling distribution does not depend on the other instances $\mathbf{x}^{(-i)}, \mathbf{z}^{(-i)}$:
 2858 given the parameters $\boldsymbol{\phi}$ and $\boldsymbol{\mu}$, the posterior distribution over each $z^{(i)}$ can be computed
 2859 from $\mathbf{x}^{(i)}$ alone.

2860 In sampling algorithms, there are several choices for how to deal with the parameters.
 2861 One possibility is to sample them too. To do this, we must add them to the generative
 2862 story, by introducing a prior distribution. For the multinomial and categorical parameters
 2863 in the EM clustering model, the **Dirichlet distribution** is a typical choice, since it defines
 2864 a probability on exactly the set of vectors that can be parameters: vectors that sum to one
 2865 and include only non-negative numbers.¹⁰

2866 To incorporate this prior, the generative model must augmented to indicate that each
 2867 $\boldsymbol{\phi}_z \sim \text{Dirichlet}(\boldsymbol{\alpha}_\phi)$, and $\boldsymbol{\mu} \sim \text{Dirichlet}(\boldsymbol{\alpha}_\mu)$. The hyperparameters $\boldsymbol{\alpha}$ are typically set to

¹⁰If $\sum_i^K \theta_i = 1$ and $\theta_i \geq 0$ for all i , then $\boldsymbol{\theta}$ is said to be on the $K - 1$ **simplex**. A Dirichlet distribution with parameter $\boldsymbol{\alpha} \in \mathbb{R}_+^K$ has support over the $K - 1$ simplex,

$$p_{\text{Dirichlet}}(\boldsymbol{\theta} | \boldsymbol{\alpha}) = \frac{1}{B(\boldsymbol{\alpha})} \prod_{i=1}^K \theta_i^{\alpha_i - 1} \quad [5.44]$$

$$B(\boldsymbol{\alpha}) = \frac{\prod_{i=1}^K \Gamma(\alpha_i)}{\Gamma(\sum_{i=1}^K \alpha_i)}, \quad [5.45]$$

with $\Gamma(\cdot)$ indicating the gamma function, a generalization of the factorial function to non-negative reals.

2868 a constant vector $\alpha = [\alpha, \alpha, \dots, \alpha]$. When α is large, the Dirichlet distribution tends to
 2869 generate vectors that are nearly uniform; when α is small, it tends to generate vectors that
 2870 assign most of their probability mass to a few entries. Given prior distributions over ϕ
 2871 and μ , we can now include them in Gibbs sampling, drawing values for these parameters
 2872 from posterior distributions that are conditioned on the other variables in the model.

2873 Unfortunately, sampling ϕ and μ usually leads to slow convergence, meaning that a
 2874 large number of samples is required before the Markov chain breaks free from the initial
 2875 conditions. The reason is that the sampling distributions for these parameters are tightly
 2876 constrained by the cluster memberships $y^{(i)}$, which in turn are tightly constrained by the
 2877 parameters. There are two solutions that are frequently employed:

- 2878 • **Empirical Bayesian** methods maintain ϕ and μ as parameters rather than latent
 2879 variables. They still employ sampling in the E-step of the EM algorithm, but they
 2880 update the parameters using expected counts that are computed from the samples
 2881 rather than from parametric distributions. This EM-MCMC hybrid is also known
 2882 as Monte Carlo Expectation Maximization (MCEM; Wei and Tanner, 1990), and is
 2883 well-suited for cases in which it is difficult to compute $q^{(i)}$ directly.
- 2884 • In **collapsed Gibbs sampling**, we analytically integrate ϕ and μ out of the model.
 2885 The cluster memberships $y^{(i)}$ are the only remaining latent variable; we sample them
 2886 from the compound distribution,

$$p(y^{(i)} | \mathbf{x}^{(1:N)}, \mathbf{y}^{(-i)}; \alpha_\phi, \alpha_\mu) = \int_{\phi, \mu} p(\phi, \mu | \mathbf{y}^{(-i)}, \mathbf{x}^{(1:N)}; \alpha_\phi, \alpha_\mu) p(y^{(i)} | \mathbf{x}^{(1:N)}, \mathbf{y}^{(-i)}, \phi, \mu) d\phi d\mu. \quad [5.46]$$

2887 For multinomial and Dirichlet distributions, the sampling distribution can be com-
 2888 puted in closed form.

2889 MCMC algorithms are guaranteed to converge to the true posterior distribution over
 2890 the latent variables, but there is no way to know how long this will take. In practice, the
 2891 rate of convergence depends on initialization, just as expectation-maximization depends
 2892 on initialization to avoid local optima. Thus, while Gibbs Sampling and other MCMC
 2893 algorithms provide a powerful and flexible array of techniques for statistical inference in
 2894 latent variable models, they are not a panacea for the problems experienced by EM.

2895 5.5.2 Spectral learning

Another approach to learning with latent variables is based on the **method of moments**, which makes it possible to avoid the problem of non-convex log-likelihood. Write $\bar{\mathbf{x}}^{(i)}$ for the normalized vector of word counts in document i , so that $\bar{\mathbf{x}}^{(i)} = \mathbf{x}^{(i)} / \sum_{j=1}^V x_j^{(i)}$. Then

we can form a matrix of word-word co-occurrence probabilities,

$$\mathbf{C} = \sum_{i=1}^N \bar{\mathbf{x}}^{(i)} (\bar{\mathbf{x}}^{(i)})^\top. \quad [5.47]$$

The expected value of this matrix under $p(\mathbf{x} | \phi, \mu)$, as

$$E[\mathbf{C}] = \sum_{i=1}^N \sum_{k=1}^K \Pr(Z^{(i)} = k; \boldsymbol{\mu}) \phi_k \phi_k^\top \quad [5.48]$$

$$= \sum_k^K N \mu_k \phi_k \phi_k^\top \quad [5.49]$$

$$= \Phi \text{Diag}(N\mu) \Phi^\top, \quad [5.50]$$

where Φ is formed by horizontally concatenating $\phi_1 \dots \phi_K$, and $\text{Diag}(N\mu)$ indicates a diagonal matrix with values $N\mu_k$ at position (k, k) . Setting \mathbf{C} equal to its expectation gives,

$$\mathbf{C} = \Phi \text{Diag}(N\mu) \Phi^\top, \quad [5.51]$$

which is similar to the eigendecomposition $\mathbf{C} = \mathbf{Q}\Lambda\mathbf{Q}^\top$. This suggests that simply by finding the eigenvectors and eigenvalues of \mathbf{C} , we could obtain the parameters ϕ and μ , and this is what motivates the name **spectral learning**.

While moment-matching and eigendecomposition are similar in form, they impose different constraints on the solutions: eigendecomposition requires orthonormality, so that $\mathbf{Q}\mathbf{Q}^\top = \mathbb{I}$; in estimating the parameters of a text clustering model, we require that μ and the columns of Φ are probability vectors. Spectral learning algorithms must therefore include a procedure for converting the solution into vectors that are non-negative and sum to one. One approach is to replace eigendecomposition (or the related singular value decomposition) with non-negative matrix factorization (Xu et al., 2003), which guarantees that the solutions are non-negative (Arora et al., 2013).

After obtaining the parameters ϕ and μ , the distribution over clusters can be computed from Bayes' rule:

$$p(z^{(i)} | \mathbf{x}^{(i)}; \phi, \mu) \propto p(\mathbf{x}^{(i)} | z^{(i)}; \phi) \times p(z^{(i)}; \mu). \quad [5.52]$$

Spectral learning yields provably good solutions without regard to initialization, and can be quite fast in practice. However, it is more difficult to apply to a broad family of generative models than more generic techniques like EM and Gibbs Sampling. For more on applying spectral learning across a range of latent variable models, see Anandkumar et al. (2014).

2914 **Additional resources**

2915 There are a number of other learning paradigms that deviate from supervised learning.

- 2916 • **Active learning:** the learner selects unlabeled instances and requests annotations (Set-
- 2917 tles, 2012).
- 2918 • **Multiple instance learning:** labels are applied to bags of instances, with a positive
- 2919 label applied if at least one instance in the bag meets the criterion (Dietterich et al.,
- 2920 1997; Maron and Lozano-Pérez, 1998).
- 2921 • **Constraint-driven learning:** supervision is provided in the form of explicit con-
- 2922 straints on the learner (Chang et al., 2007; Ganchev et al., 2010).
- 2923 • **Distant supervision:** noisy labels are generated from an external resource (Mintz
- 2924 et al., 2009, also see § 17.2.3).
- 2925 • **Multitask learning:** the learner induces a representation that can be used to solve
- 2926 multiple classification tasks (Collobert et al., 2011).
- 2927 • **Transfer learning:** the learner must solve a classification task that differs from the
- 2928 labeled data (Pan and Yang, 2010).

2929 Expectation maximization was introduced by Dempster et al. (1977), and is discussed

2930 in more detail by Murphy (2012). Like most machine learning treatments, Murphy focus

2931 on continuous observations and Gaussian likelihoods, rather than the discrete observa-

2932 tions typically encountered in natural language processing. Murphy (2012) also includes

2933 an excellent chapter on MCMC; for a textbook-length treatment, see Robert and Casella

2934 (2013). For still more on Bayesian latent variable models, see Barber (2012), and for ap-

2935 plications of Bayesian models to natural language processing, see Cohen (2016). Surveys

2936 are available for semi-supervised learning (Zhu and Goldberg, 2009) and domain adapta-

2937 tion (Søgaard, 2013), although both pre-date the current wave of interest in deep learning.

2938 **Exercises**

- 2939 1. Derive the expectation maximization update for the parameter μ in the EM cluster-
- 2940 ing model.
- 2941 2. The expectation maximization lower bound \mathcal{J} is defined in Equation 5.10. Prove
- 2942 that the inverse $-\mathcal{J}$ is convex in q . You can use the following facts about convexity:

 - 2943 • $f(\mathbf{x})$ is convex in \mathbf{x} iff $\alpha f(\mathbf{x}_1) + (1 - \alpha)f(\mathbf{x}_2) \geq f(\alpha\mathbf{x}_1 + (1 - \alpha)\mathbf{x}_2)$ for all
 - 2944 $\alpha \in [0, 1]$.
 - 2945 • If $f(\mathbf{x})$ and $g(\mathbf{x})$ are both convex in \mathbf{x} , then $f(\mathbf{x}) + g(\mathbf{x})$ is also convex in \mathbf{x} .

- 2946 • $\log(x + y) \leq \log x + \log y.$

2947 3. Derive the E-step and M-step updates for the following generative model. You may
2948 assume that the labels $y^{(i)}$ are observed, but $z_m^{(i)}$ is not.

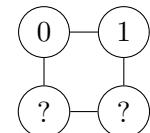
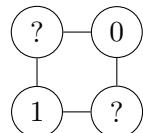
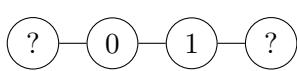
- 2949 • For each instance i ,

- 2950 – Draw label $y^{(i)} \sim \text{Categorical}(\boldsymbol{\mu})$
- 2951 – For each token $m \in \{1, 2, \dots, M^{(i)}\}$
 - 2952 * Draw $z_m^{(i)} \sim \text{Categorical}(\pi)$
 - 2953 * If $z_m^{(i)} = 0$, draw the current token from a label-specific distribution,
2954 $w_m^{(i)} \sim \phi_{y^{(i)}}$
 - 2955 * If $z_m^{(i)} = 1$, draw the current token from a document-specific distribu-
2956 tion, $w_m^{(i)} \sim \nu^{(i)}$

2957 4. Use expectation-maximization clustering to train a word-sense induction system,
2958 applied to the word *say*.

- 2959 • Import `nltk`, run `nltk.download()` and select `semcor`. Import `semcor`
2960 from `nltk.corpus`.
- 2961 • The command `semcor.tagged_sentences(tag='sense')` returns an iterator
2962 over sense-tagged sentences in the corpus. Each sentence can be viewed as
2963 an iterator over `tree` objects. For `tree` objects that are sense-annotated words,
2964 you can access the annotation as `tree.label()`, and the word itself with
2965 `tree.leaves()`. So `semcor.tagged_sentences(tag='sense')[0][2].label()`
2966 would return the sense annotation of the third word in the first sentence.
- 2967 • Extract all sentences containing the senses `say.v.01` and `say.v.02`.
- 2968 • Build bag-of-words vectors $x^{(i)}$, containing the counts of other words in those
2969 sentences, including all words that occur in at least two sentences.
- 2970 • Implement and run expectation-maximization clustering on the merged data.
- 2971 • Compute the frequency with which each cluster includes instances of `say.v.01`
2972 and `say.v.02`.

2973 5. Using the iterative updates in Equations 5.34-5.36, compute the outcome of the label
2974 propagation algorithm for the following examples.



2975 The value inside the node indicates the label, $y^{(i)} \in \{0, 1\}$, with $y^{(i)} = ?$ for unlabeled
 2976 nodes. The presence of an edge between two nodes indicates $w_{i,j} = 1$, and the
 2977 absence of an edge indicates $w_{i,j} = 0$. For the third example, you need only compute
 2978 the first three iterations, and then you can guess at the solution in the limit.

2979 In the remaining exercises, you will try out some approaches for semisupervised learning
 2980 and domain adaptation. You will need datasets in multiple domains. You can obtain
 2981 product reviews in multiple domains here: https://www.cs.jhu.edu/~mdredze/datasets/sentiment/processed_acl.tar.gz. Choose a source and target domain,
 2982 e.g. dvds and books, and divide the data for the target domain into training and test sets
 2983 of equal size.

- 2985 6. First, quantify the cost of cross-domain transfer.
- 2986 • Train a logistic regression classifier on the source domain training set, and eval-
 uate it on the target domain test set.
 - 2988 • Train a logistic regression classifier on the target domain training set, and eval-
 uate it on the target domain test set. This is the “direct transfer” baseline.

2990 Compute the difference in accuracy, which is a measure of the transfer loss across
 2991 domains.

- 2992 7. Next, apply the **label propagation** algorithm from § 5.3.2.

2993 As a baseline, using only 5% of the target domain training set, train a classifier, and
 2994 compute its accuracy on the target domain test set.

2995 Next, apply label propagation:

- 2996 • Compute the label matrix \mathbf{Q}_L for the labeled data (5% of the target domain
 training set), with each row equal to an indicator vector for the label (positive
 or negative).
- 2999 • Iterate through the target domain instances, including both test and training
 data. At each instance i , compute all w_{ij} , using Equation 5.32, with $\alpha = 0.01$.
 Use these values to fill in column i of the transition matrix \mathbf{T} , setting all but the
 ten largest values to zero for each column i . Be sure to normalize the column
 so that the remaining values sum to one. You may need to use a sparse matrix
 for this to fit into memory.
- 3005 • Apply the iterative updates from Equations 5.34-5.36 to compute the outcome
 of the label propagation algorithm for the unlabeled examples.

3007 Select the test set instances from \mathbf{Q}_U , and compute the accuracy of this method.
 3008 Compare with the supervised classifier trained only on the 5% sample of the target
 3009 domain training set.

- 3010 8. Using only 5% of the target domain training data (and all of the source domain train-
3011 ing data), implement one of the supervised domain adaptation baselines in § 5.4.1.
3012 See if this improves on the “direct transfer” baseline from the previous problem
- 3013 9. Implement EasyAdapt (§ 5.4.1), again using 5% of the target domain training data
3014 and all of the source domain data.
- 3015 10. Now try unsupervised domain adaptation, using the “linear projection” method
3016 described in § 5.4.2. Specifically:
- 3017 • Identify 500 pivot features as the words with the highest frequency in the (com-
3018 plete) training data for the source and target domains. Specifically, let x_i^d be the
3019 count of the word i in domain d : choose the 500 words with the largest values
3020 of $\min(x_i^{\text{source}}, x_i^{\text{target}})$.
 - 3021 • Train a classifier to predict each pivot feature from the remaining words in the
3022 document.
 - 3023 • Arrange the features of these classifiers into a matrix Φ , and perform truncated
3024 singular value decomposition, with $k = 20$
 - 3025 • Train a classifier from the source domain data, using the combined features
3026 $\mathbf{x}^{(i)} \oplus \mathbf{U}^\top \mathbf{x}^{(i)}$ — these include the original bag-of-words features, plus the pro-
3027 jected features.
 - 3028 • Apply this classifier to the target domain test set, and compute the accuracy.

3029

Part II

3030

Sequences and trees

3031 Chapter 6

3032 Language models

3033 In probabilistic classification, the problem is to compute the probability of a label, conditioned on the text. Let's now consider the inverse problem: computing the probability of text itself. Specifically, we will consider models that assign probability to a sequence of word tokens, $p(w_1, w_2, \dots, w_M)$, with $w_m \in \mathcal{V}$. The set \mathcal{V} is a discrete vocabulary,

$$\mathcal{V} = \{aardvark, abacus, \dots, zither\}. \quad [6.1]$$

3037 Why would you want to compute the probability of a word sequence? In many applications, the goal is to produce word sequences as output:

- 3039 • In **machine translation** (chapter 18), we convert from text in a source language to text in a target language.
- 3040 • In **speech recognition**, we convert from audio signal to text.
- 3041 • In **summarization** (§ 16.3.4.1; § 19.2), we convert from long texts into short texts.
- 3042 • In **dialogue systems** (§ 19.3), we convert from the user's input (and perhaps an external knowledge base) into a text response.

3045 In many of the systems for performing these tasks, there is a subcomponent that computes the probability of the output text. The purpose of this component is to generate texts that are more **fluent**. For example, suppose we want to translate a sentence from Spanish to English.

3049 (6.1) El cafe negro me gusta mucho.

3050 Here is a literal word-for-word translation (a **gloss**):

3051 (6.2) The coffee black me pleases much.

3052 A good language model of English will tell us that the probability of this translation is
 3053 low, in comparison with more grammatical alternatives,

$$p(\text{The coffee black me pleases much}) < p(\text{I love dark coffee}). \quad [6.2]$$

3054 How can we use this fact? Warren Weaver, one of the early leaders in machine trans-
 3055 lation, viewed it as a problem of breaking a secret code (Weaver, 1955):

3056 When I look at an article in Russian, I say: 'This is really written in English,
 3057 but it has been coded in some strange symbols. I will now proceed to decode.'

3058 This observation motivates a generative model (like Naïve Bayes):

3059 • The English sentence $w^{(e)}$ is generated from a **language model**, $p_e(w^{(e)})$.

3060 • The Spanish sentence $w^{(s)}$ is then generated from a **translation model**, $p_{s|e}(w^{(s)} | w^{(e)})$.

Given these two distributions, we can then perform translation by Bayes rule:

$$p_{e|s}(w^{(e)} | w^{(s)}) \propto p_{e,s}(w^{(e)}, w^{(s)}) \quad [6.3]$$

$$= p_{s|e}(w^{(s)} | w^{(e)}) \times p_e(w^{(e)}). \quad [6.4]$$

3061 This is sometimes called the **noisy channel model**, because it envisions English text
 3062 turning into Spanish by passing through a noisy channel, $p_{s|e}$. What is the advantage of
 3063 modeling translation this way, as opposed to modeling $p_{e|s}$ directly? The crucial point is
 3064 that the two distributions $p_{s|e}$ (the translation model) and p_e (the language model) can be
 3065 estimated from separate data. The translation model requires examples of correct trans-
 3066 lations, but the language model requires only text in English. Such monolingual data is
 3067 much more widely available. Furthermore, once estimated, the language model p_e can be
 3068 reused in any application that involves generating English text, from summarization to
 3069 speech recognition.

3070 6.1 *N*-gram language models

A simple approach to computing the probability of a sequence of tokens is to use a **relative frequency estimate**. For example, consider the quote, attributed to Picasso, "computers are useless, they can only give you answers." We can estimate the probability of this sentence,

$$\begin{aligned} p(\text{Computers are useless, they can only give you answers}) \\ = \frac{\text{count}(\text{Computers are useless, they can only give you answers})}{\text{count}(\text{all sentences ever spoken})} \end{aligned} \quad [6.5]$$

3071 This estimator is **unbiased**: in the theoretical limit of infinite data, the estimate will
 3072 be correct. But in practice, we are asking for accurate counts over an infinite number of
 3073 events, since sequences of words can be arbitrarily long. Even with an aggressive upper
 3074 bound of, say, $M = 20$ tokens in the sequence, the number of possible sequences is V^{20} . A
 3075 small vocabulary for English would have $V = 10^4$, so there are 10^{80} possible sequences.
 3076 Clearly, this estimator is very data-hungry, and suffers from high variance: even gram-
 3077 matical sentences will have probability zero if have not occurred in the training data.¹ We
 3078 therefore need to introduce bias to have a chance of making reliable estimates from finite
 3079 training data. The language models that follow in this chapter introduce bias in various
 3080 ways.

We begin with n -gram language models, which compute the probability of a sequence as the product of probabilities of subsequences. The probability of a sequence $p(\mathbf{w}) = p(w_1, w_2, \dots, w_M)$ can be refactored using the chain rule (see § A.2):

$$p(\mathbf{w}) = p(w_1, w_2, \dots, w_M) \quad [6.6]$$

$$= p(w_1) \times p(w_2 | w_1) \times p(w_3 | w_2, w_1) \times \dots \times p(w_M | w_{M-1}, \dots, w_1) \quad [6.7]$$

Each element in the product is the probability of a word given all its predecessors. We can think of this as a *word prediction* task: given the context *Computers are*, we want to compute a probability over the next token. The relative frequency estimate of the probability of the word *useless* in this context is,

$$\begin{aligned} p(\text{useless} | \text{computers are}) &= \frac{\text{count}(\text{computers are useless})}{\sum_{x \in \mathcal{V}} \text{count}(\text{computers are } x)} \\ &= \frac{\text{count}(\text{computers are useless})}{\text{count}(\text{computers are})}. \end{aligned}$$

3081 We haven't made any approximations yet, and we could have just as well applied the
 3082 chain rule in reverse order,

$$p(\mathbf{w}) = p(w_M) \times p(w_{M-1} | w_M) \times \dots \times p(w_1 | w_2, \dots, w_M), \quad [6.8]$$

3083 or in any other order. But this means that we also haven't really made any progress:
 3084 to compute the conditional probability $p(w_M | w_{M-1}, w_{M-2}, \dots, w_1)$, we would need to
 3085 model V^{M-1} contexts. Such a distribution cannot be estimated from any realistic sample
 3086 of text.

¹Chomsky has famously argued that this is evidence against the very concept of probabilistic language models: no such model could distinguish the grammatical sentence *colorless green ideas sleep furiously* from the ungrammatical permutation *furiously sleep ideas green colorless*. Indeed, even the bigrams in these two examples are unlikely to occur — at least, not in texts written before Chomsky proposed this example.

To solve this problem, n -gram models make a crucial simplifying approximation: condition on only the past $n - 1$ words.

$$p(w_m | w_{m-1} \dots w_1) \approx p(w_m | w_{m-1}, \dots, w_{m-n+1}) \quad [6.9]$$

This means that the probability of a sentence w can be approximated as

$$p(w_1, \dots, w_M) \approx \prod_{m=1}^M p(w_m | w_{m-1}, \dots, w_{m-n+1}) \quad [6.10]$$

To compute the probability of an entire sentence, it is convenient to pad the beginning and end with special symbols \square and \blacksquare . Then the bigram ($n = 2$) approximation to the probability of *I like black coffee* is:

$$p(I \text{ like black coffee}) = p(I | \square) \times p(\text{like} | I) \times p(\text{black} | \text{like}) \times p(\text{coffee} | \text{black}) \times p(\blacksquare | \text{coffee}). \quad [6.11]$$

3087 This model requires estimating and storing the probability of only V^n events, which is
 3088 exponential in the order of the n -gram, and not V^M , which is exponential in the length of
 3089 the sentence. The n -gram probabilities can be computed by relative frequency estimation,

$$p(w_m | w_{m-1}, w_{m-2}) = \frac{\text{count}(w_{m-2}, w_{m-1}, w_m)}{\sum_{w'} \text{count}(w_{m-2}, w_{m-1}, w')} \quad [6.12]$$

3090 The hyperparameter n controls the size of the context used in each conditional proba-
 3091 bility. If this is misspecified, the language model will perform poorly. Let's consider the
 3092 potential problems concretely.

3093 **When n is too small.** Consider the following sentences:

3094 (6.3) **Gorillas** always like to groom **their** friends.

3095 (6.4) The **computer** that's on the 3rd floor of our office building **crashed**.

3096 In each example, the bolded words depend on each other: the likelihood of *their*
 3097 depends on knowing that *gorillas* is plural, and the likelihood of *crashed* depends on
 3098 knowing that the subject is a *computer*. If the n -grams are not big enough to capture
 3099 this context, then the resulting language model would offer probabilities that are too
 3100 low for these sentences, and too high for sentences that fail basic linguistic tests like
 3101 number agreement.

3102 **When n is too big.** In this case, it is hard to get good estimates of the n -gram parameters from
 3103 our dataset, because of data sparsity. To handle the *gorilla* example, it is necessary to
 3104 model 6-grams, which means accounting for V^6 events. Under a very small vocabulary of $V = 10^4$, this means estimating the probability of 10^{24} distinct events.

3106 These two problems point to another **bias-variance tradeoff** (see § 2.1.4). A small n -
 3107 gram size introduces high bias, and a large n -gram size introduces high variance. But
 3108 in reality we often have both problems at the same time! Language is full of long-range
 3109 dependencies that we cannot capture because n is too small; at the same time, language
 3110 datasets are full of rare phenomena, whose probabilities we fail to estimate accurately
 3111 because n is too large. One solution is to try to keep n large, while still making low-
 3112 variance estimates of the underlying parameters. To do this, we will introduce a different
 3113 sort of bias: **smoothing**.

3114 6.2 Smoothing and discounting

3115 Limited data is a persistent problem in estimating language models. In § 6.1, we presented
 3116 n -grams as a partial solution. sparse data can be a problem even for low-order n -grams;
 3117 at the same time, many linguistic phenomena, like subject-verb agreement, cannot be in-
 3118 corporated into language models without high-order n -grams. It is therefore necessary to
 3119 add additional inductive biases to n -gram language models. This section covers some of
 3120 the most intuitive and common approaches, but there are many more (Chen and Good-
 3121 man, 1999).

3122 6.2.1 Smoothing

3123 A major concern in language modeling is to avoid the situation $p(w) = 0$, which could
 3124 arise as a result of a single unseen n-gram. A similar problem arose in Naïve Bayes, and
 3125 the solution was **smoothing**: adding imaginary “pseudo” counts. The same idea can be
 3126 applied to n -gram language models, as shown here in the bigram case,

$$P_{\text{smooth}}(w_m \mid w_{m-1}) = \frac{\text{count}(w_{m-1}, w_m) + \alpha}{\sum_{w' \in \mathcal{V}} \text{count}(w_{m-1}, w') + V\alpha}. \quad [6.13]$$

3127 This basic framework is called **Lidstone smoothing**, but special cases have other names:

- 3128 • **Laplace smoothing** corresponds to the case $\alpha = 1$.
- 3129 • **Jeffreys-Perks law** corresponds to the case $\alpha = 0.5$. Manning and Schütze (1999)
 3130 offer more insight on the justifications for this setting.

3131 To maintain normalization, anything that we add to the numerator (α) must also ap-
 3132 pear in the denominator ($V\alpha$). This idea is reflected in the concept of **effective counts**:

$$c_i^* = (c_i + \alpha) \frac{M}{M + V\alpha}, \quad [6.14]$$

	counts	unsmoothed probability	Lidstone smoothing, $\alpha = 0.1$		Discounting, $d = 0.1$	
			effective counts	smoothed probability	effective counts	smoothed probability
<i>impropriety</i>	8	0.4	7.826	0.391	7.9	0.395
<i>offense</i>	5	0.25	4.928	0.246	4.9	0.245
<i>damage</i>	4	0.2	3.961	0.198	3.9	0.195
<i>deficiencies</i>	2	0.1	2.029	0.101	1.9	0.095
<i>outbreak</i>	1	0.05	1.063	0.053	0.9	0.045
<i>infirmity</i>	0	0	0.097	0.005	0.25	0.013
<i>cephalopods</i>	0	0	0.097	0.005	0.25	0.013

Table 6.1: Example of Lidstone smoothing and absolute discounting in a bigram language model, for the context $(\text{alleged}, \cdot)$, for a toy corpus with a total of twenty counts over the seven words shown. Note that discounting decreases the probability for all but the unseen words, while Lidstone smoothing increases the effective counts and probabilities for *deficiencies* and *outbreak*.

where c_i is the count of event i , c_i^* is the effective count, and $M = \sum_{i=1}^V c_i$ is the total number of tokens in the dataset (w_1, w_2, \dots, w_M) . This term ensures that $\sum_{i=1}^V c_i^* = \sum_{i=1}^V c_i = M$. The **discount** for each n-gram is then computed as,

$$d_i = \frac{c_i^*}{c_i} = \frac{(c_i + \alpha)}{c_i} \frac{M}{(M + V\alpha)}.$$

3133 6.2.2 Discounting and backoff

3134 Discounting “borrows” probability mass from observed n -grams and redistributes it. In
 3135 Lidstone smoothing, the borrowing is done by increasing the denominator of the relative
 3136 frequency estimates. The borrowed probability mass is then redistributed by increasing
 3137 the numerator for all n -grams. Another approach would be to borrow the same amount
 3138 of probability mass from all observed n -grams, and redistribute it among only the unob-
 3139 served n -grams. This is called **absolute discounting**. For example, suppose we set an
 3140 absolute discount $d = 0.1$ in a bigram model, and then redistribute this probability mass
 3141 equally over the unseen words. The resulting probabilities are shown in Table 6.1.

Discounting reserves some probability mass from the observed data, and we need not redistribute this probability mass equally. Instead, we can **backoff** to a lower-order language model: if you have trigrams, use trigrams; if you don’t have trigrams, use bigrams; if you don’t even have bigrams, use unigrams. This is called **Katz backoff**. In the simple

case of backing off from bigrams to unigrams, the bigram probabilities are computed as,

$$c^*(i, j) = c(i, j) - d \quad [6.15]$$

$$p_{\text{Katz}}(i | j) = \begin{cases} \frac{c^*(i, j)}{c(j)} & \text{if } c(i, j) > 0 \\ \alpha(j) \times \frac{p_{\text{unigram}}(i)}{\sum_{i': c(i', j)=0} p_{\text{unigram}}(i')} & \text{if } c(i, j) = 0. \end{cases} \quad [6.16]$$

3142 The term $\alpha(j)$ indicates the amount of probability mass that has been discounted for
 3143 context j . This probability mass is then divided across all the unseen events, $\{i' : c(i', j) =$
 3144 $0\}$, proportional to the unigram probability of each word i' . The discount parameter d can
 3145 be optimized to maximize performance (typically held-out log-likelihood) on a develop-
 3146 ment set.

3147 6.2.3 *Interpolation

3148 Backoff is one way to combine different order n -gram models. An alternative approach
 3149 is **interpolation**: setting the probability of a word in context to a weighted sum of its
 3150 probabilities across progressively shorter contexts.

Instead of choosing a single n for the size of the n -gram, we can take the weighted average across several n -gram probabilities. For example, for an interpolated trigram model,

$$\begin{aligned} p_{\text{Interpolation}}(w_m | w_{m-1}, w_{m-2}) &= \lambda_3 p_3^*(w_m | w_{m-1}, w_{m-2}) \\ &\quad + \lambda_2 p_2^*(w_m | w_{m-1}) \\ &\quad + \lambda_1 p_1^*(w_m). \end{aligned}$$

3151 In this equation, p_n^* is the unsmoothed empirical probability given by an n -gram lan-
 3152 guage model, and λ_n is the weight assigned to this model. To ensure that the interpolated
 3153 $p(w)$ is still a valid probability distribution, the values of λ must obey the constraint,
 3154 $\sum_{n=1}^{n_{\max}} \lambda_n = 1$. But how to find the specific values?

3155 An elegant solution is **expectation maximization**. Recall from chapter 5 that we can
 3156 think about EM as learning with *missing data*: we just need to choose missing data such
 3157 that learning would be easy if it weren't missing. What's missing in this case? Think of
 3158 each word w_m as drawn from an n -gram of unknown size, $z_m \in \{1 \dots n_{\max}\}$. This z_m is
 3159 the missing data that we are looking for. Therefore, the application of EM to this problem
 3160 involves the following **generative process**:

3161 **for** Each token $w_m, m = 1, 2, \dots, M$ **do**:
 3162 draw the n -gram size $z_m \sim \text{Categorical}(\lambda)$;
 3163 draw $w_m \sim p_{z_m}^*(w_m | w_{m-1}, \dots, w_{m-z_m})$.

If the missing data $\{Z_m\}$ were known, then λ could be estimated as the relative frequency,

$$\lambda_z = \frac{\text{count}(Z_m = z)}{M} \quad [6.17]$$

$$\propto \sum_{m=1}^M \delta(Z_m = z). \quad [6.18]$$

But since we do not know the values of the latent variables Z_m , we impute a distribution q_m in the E-step, which represents the degree of belief that word token w_m was generated from a n -gram of order z_m ,

$$q_m(z) \triangleq \Pr(Z_m = z \mid \mathbf{w}_{1:m}; \lambda) \quad [6.19]$$

$$= \frac{p(w_m \mid \mathbf{w}_{1:m-1}, Z_m = z) \times p(z)}{\sum_{z'} p(w_m \mid \mathbf{w}_{1:m-1}, Z_m = z') \times p(z')} \quad [6.20]$$

$$\propto p_z^*(w_m \mid \mathbf{w}_{1:m-1}) \times \lambda_z. \quad [6.21]$$

In the M-step, λ is computed by summing the expected counts under q ,

$$\lambda_z \propto \sum_{m=1}^M q_m(z). \quad [6.22]$$

3165 A solution is obtained by iterating between updates to q and λ . The complete algorithm
 3166 is shown in Algorithm 10.

Algorithm 10 Expectation-maximization for interpolated language modeling

```

1: procedure ESTIMATE INTERPOLATED  $n$ -GRAM ( $\mathbf{w}_{1:M}, \{p_n^*\}_{n \in 1:n_{\max}}$ )
2:   for  $z \in \{1, 2, \dots, n_{\max}\}$  do ▷ Initialization
3:      $\lambda_z \leftarrow \frac{1}{n_{\max}}$ 
4:   repeat
5:     for  $m \in \{1, 2, \dots, M\}$  do ▷ E-step
6:       for  $z \in \{1, 2, \dots, n_{\max}\}$  do
7:          $q_m(z) \leftarrow p_z^*(w_m \mid \mathbf{w}_{1:m-1}) \times \lambda_z$ 
8:        $q_m \leftarrow \text{Normalize}(q_m)$ 
9:     for  $z \in \{1, 2, \dots, n_{\max}\}$  do ▷ M-step
10:       $\lambda_z \leftarrow \frac{1}{M} \sum_{m=1}^M q_m(z)$ 
11:    until tired
12:    return  $\lambda$ 
  
```

3167 **6.2.4 *Kneser-Ney smoothing**

3168 Kneser-Ney smoothing is based on absolute discounting, but it redistributes the result-
 3169 ing probability mass in a different way from Katz backoff. Empirical evidence points
 3170 to Kneser-Ney smoothing as the state-of-art for n -gram language modeling (Goodman,
 3171 2001). To motivate Kneser-Ney smoothing, consider the example: *I recently visited ..*
 3172 Which of the following is more likely?

- 3173 • *Francisco*
 3174 • *Duluth*

3175 Now suppose that both bigrams *visited Duluth* and *visited Francisco* are unobserved in
 3176 the training data, and furthermore, the unigram probability $p_1^*(\text{Francisco})$ is greater than
 3177 $p^*(\text{Duluth})$. Nonetheless we would still guess that $p(\text{visited Duluth}) > p(\text{visited Francisco})$,
 3178 because *Duluth* is a more “versatile” word: it can occur in many contexts, while *Francisco*
 3179 usually occurs in a single context, following the word *San*. This notion of versatility is the
 3180 key to Kneser-Ney smoothing.

Writing u for a context of undefined length, and $\text{count}(w, u)$ as the count of word w in
 context u , we define the Kneser-Ney bigram probability as

$$p_{KN}(w | u) = \begin{cases} \frac{\text{count}(w, u) - d}{\text{count}(u)}, & \text{count}(w, u) > 0 \\ \alpha(u) \times p_{\text{continuation}}(w), & \text{otherwise} \end{cases} \quad [6.23]$$

$$p_{\text{continuation}}(w) = \frac{|u : \text{count}(w, u) > 0|}{\sum_{w' \in \mathcal{V}} |u' : \text{count}(w', u') > 0|}. \quad [6.24]$$

First, note that we reserve probability mass using absolute discounting d , which is taken from all unobserved n -grams. The total amount of discounting in context u is $d \times |w : \text{count}(w, u) > 0|$, and we divide this probability mass equally among the unseen n -grams,

$$\alpha(u) = |w : \text{count}(w, u) > 0| \times \frac{d}{\text{count}(u)}. \quad [6.25]$$

3181 This is the amount of probability mass left to account for versatility, which we define via
 3182 the *continuation probability* $p_{\text{continuation}}(w)$ as proportional to the number of observed con-
 3183 texts in which w appears. The numerator of the continuation probability is the number of
 3184 contexts u in which w appears; the denominator normalizes the probability by summing
 3185 the same quantity over all words w' .

3186 The idea of modeling versatility by counting contexts may seem heuristic, but there is
 3187 an elegant theoretical justification from Bayesian nonparametrics (Teh, 2006). Kneser-Ney
 3188 smoothing on n -grams was the dominant language modeling technique before the arrival
 3189 of neural language models.

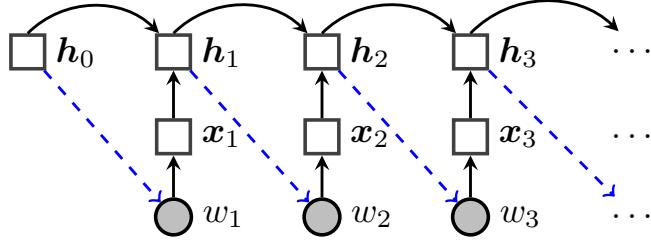


Figure 6.1: The recurrent neural network language model, viewed as an “unrolled” computation graph. Solid lines indicate direct computation, dotted blue lines indicate probabilistic dependencies, circles indicate random variables, and squares indicate computation nodes.

3190 6.3 Recurrent neural network language models

3191 N -gram language models have been largely supplanted by **neural networks**. These mod-
 3192 els do not make the n -gram assumption of restricted context; indeed, they can incorpo-
 3193 rate arbitrarily distant contextual information, while remaining computationally and statisti-
 3194 cally tractable.

3195 The first insight behind neural language models is to treat word prediction as a *dis-
 3196 criminative* learning task.² The goal is to compute the probability $p(w | u)$, where $w \in \mathcal{V}$ is
 3197 a word, and u is the context, which depends on the previous words. Rather than directly
 3198 estimating the word probabilities from (smoothed) relative frequencies, we can treat
 3199 language modeling as a machine learning problem, and estimate parameters that maxi-
 3200 mize the log conditional probability of a corpus.

3201 The second insight is to reparametrize the probability distribution $p(w | u)$ as a func-
 3202 tion of two dense K -dimensional numerical vectors, $\beta_w \in \mathbb{R}^K$, and $v_u \in \mathbb{R}^K$,

$$p(w | u) = \frac{\exp(\beta_w \cdot v_u)}{\sum_{w' \in \mathcal{V}} \exp(\beta_{w'} \cdot v_u)}, \quad [6.26]$$

3203 where $\beta_w \cdot v_u$ represents a dot product. As usual, the denominator ensures that the prob-
 3204 ability distribution is properly normalized. This vector of probabilities is equivalent to
 3205 applying the **softmax** transformation (see § 3.1) to the vector of dot-products,

$$p(\cdot | u) = \text{SoftMax}([\beta_1 \cdot v_u, \beta_2 \cdot v_u, \dots, \beta_V \cdot v_u]). \quad [6.27]$$

The word vectors β_w are parameters of the model, and are estimated directly. The context vectors v_u can be computed in various ways, depending on the model. A simple

²This idea predates neural language models (e.g., Rosenfeld, 1996; Roark et al., 2007).

but effective neural language model can be built from a **recurrent neural network** (RNN; Mikolov et al., 2010). The basic idea is to recurrently update the context vectors while moving through the sequence. Let \mathbf{h}_m represent the contextual information at position m in the sequence. RNN language models are defined,

$$\mathbf{x}_m \triangleq \phi_{w_m} \quad [6.28]$$

$$\mathbf{h}_m = \text{RNN}(\mathbf{x}_m, \mathbf{h}_{m-1}) \quad [6.29]$$

$$p(w_{m+1} | w_1, w_2, \dots, w_m) = \frac{\exp(\beta_{w_{m+1}} \cdot \mathbf{h}_m)}{\sum_{w' \in \mathcal{V}} \exp(\beta_{w'} \cdot \mathbf{h}_m)}, \quad [6.30]$$

where ϕ is a matrix of **input word embeddings**, and \mathbf{x}_m denotes the embedding for word w_m . The conversion of w_m to \mathbf{x}_m is sometimes known as a **lookup layer**, because we simply lookup the embeddings for each word in a table; see § 3.2.4.

The **Elman unit** defines a simple recurrent operation (Elman, 1990),

$$\text{RNN}(\mathbf{x}_m, \mathbf{h}_{m-1}) \triangleq g(\Theta \mathbf{h}_{m-1} + \mathbf{x}_m), \quad [6.31]$$

where $\Theta \in \mathbb{R}^{K \times K}$ is the recurrence matrix and g is a non-linear transformation function, often defined as the elementwise hyperbolic tangent \tanh (see § 3.1).³ The \tanh acts as a **squashing function**, ensuring that each element of \mathbf{h}_m is constrained to the range $[-1, 1]$.

Although each w_m depends on only the context vector \mathbf{h}_{m-1} , this vector is in turn influenced by *all* previous tokens, w_1, w_2, \dots, w_{m-1} , through the recurrence operation: w_1 affects \mathbf{h}_1 , which affects \mathbf{h}_2 , and so on, until the information is propagated all the way to \mathbf{h}_{m-1} , and then on to w_m (see Figure 6.1). This is an important distinction from n -gram language models, where any information outside the n -word window is ignored. In principle, the RNN language model can handle long-range dependencies, such as number agreement over long spans of text — although it would be difficult to know where exactly in the vector \mathbf{h}_m this information is represented. The main limitation is that information is attenuated by repeated application of the squashing function g . **Long short-term memories** (LSTMs), described below, are a variant of RNNs that address this issue, using memory cells to propagate information through the sequence without applying nonlinearities (Hochreiter and Schmidhuber, 1997).

The denominator in Equation 6.30 is a computational bottleneck, because it involves a sum over the entire vocabulary. One solution is to use a **hierarchical softmax** function, which computes the sum more efficiently by organizing the vocabulary into a tree (Mikolov et al., 2011). Another strategy is to optimize an alternative metric, such as **noise-contrastive estimation** (Gutmann and Hyvärinen, 2012), which learns by distinguishing observed instances from artificial instances generated from a noise distribution (Mnih and Teh, 2012). Both of these strategies are described in § 14.5.3.

³In the original Elman network, the sigmoid function was used in place of \tanh . For an illuminating mathematical discussion of the advantages and disadvantages of various nonlinearities in recurrent neural networks, see the lecture notes from Cho (2015).

3232 **6.3.1 Backpropagation through time**

3233 The recurrent neural network language model has the following parameters:

- 3234 • $\phi_i \in \mathbb{R}^K$, the “input” word vectors (these are sometimes called **word embeddings**,
3235 since each word is embedded in a K -dimensional space);
- 3236 • $\beta_i \in \mathbb{R}^K$, the “output” word vectors;
- 3237 • $\Theta \in \mathbb{R}^{K \times K}$, the recurrence operator;
- 3238 • \mathbf{h}_0 , the initial state.

3239 Each of these parameters can be estimated by formulating an objective function over the
3240 training corpus, $L(\mathbf{w})$, and then applying **backpropagation** to obtain gradients on the
3241 parameters from a minibatch of training examples (see § 3.3.1). Gradient-based updates
3242 can be computed from an online learning algorithm such as stochastic gradient descent
3243 (see § 2.5.2).

3244 The application of backpropagation to recurrent neural networks is known as **back-**
3245 **propagation through time**, because the gradients on units at time m depend in turn on the
3246 gradients of units at earlier times $n < m$. Let ℓ_{m+1} represent the negative log-likelihood
3247 of word $m + 1$,

$$\ell_{m+1} = -\log p(w_{m+1} | w_1, w_2, \dots, w_m). \quad [6.32]$$

We require the gradient of this loss with respect to each parameter, such as $\theta_{k,k'}$, an individual element in the recurrence matrix Θ . Since the loss depends on the parameters only through \mathbf{h}_m , we can apply the chain rule of differentiation,

$$\frac{\partial \ell_{m+1}}{\partial \theta_{k,k'}} = \frac{\partial \ell_{m+1}}{\partial \mathbf{h}_m} \frac{\partial \mathbf{h}_m}{\partial \theta_{k,k'}}. \quad [6.33]$$

The vector \mathbf{h}_m depends on Θ in several ways. First, \mathbf{h}_m is computed by multiplying Θ by the previous state \mathbf{h}_{m-1} . But the previous state \mathbf{h}_{m-1} also depends on Θ :

$$\mathbf{h}_m = g(\mathbf{x}_m, \mathbf{h}_{m-1}) \quad [6.34]$$

$$\frac{\partial h_{m,k}}{\partial \theta_{k,k'}} = g'(\mathbf{x}_{m,k} + \boldsymbol{\theta}_k \cdot \mathbf{h}_{m-1})(h_{m-1,k'} + \boldsymbol{\theta}_k \cdot \frac{\partial \mathbf{h}_{m-1}}{\partial \theta_{k,k'}}), \quad [6.35]$$

3248 where g' is the local derivative of the nonlinear function g . The key point in this equation
3249 is that the derivative $\frac{\partial \mathbf{h}_m}{\partial \theta_{k,k'}}$ depends on $\frac{\partial \mathbf{h}_{m-1}}{\partial \theta_{k,k'}}$, which will depend in turn on $\frac{\partial \mathbf{h}_{m-2}}{\partial \theta_{k,k'}}$, and
3250 so on, until reaching the initial state \mathbf{h}_0 .

3251 Each derivative $\frac{\partial \mathbf{h}_m}{\partial \theta_{k,k'}}$ will be reused many times: it appears in backpropagation from
3252 the loss ℓ_m , but also in all subsequent losses $\ell_{n>m}$. Neural network toolkits such as
3253 Torch (Collobert et al., 2011) and DyNet (Neubig et al., 2017) compute the necessary

derivatives automatically, and cache them for future use. An important distinction from the feedforward neural networks considered in chapter 3 is that the size of the computation graph is not fixed, but varies with the length of the input. This poses difficulties for toolkits that are designed around static computation graphs, such as TensorFlow (Abadi et al., 2016).⁴

6.3.2 Hyperparameters

The RNN language model has several hyperparameters that must be tuned to ensure good performance. The model capacity is controlled by the size of the word and context vectors K , which play a role that is somewhat analogous to the size of the n -gram context. For datasets that are large with respect to the vocabulary (i.e., there is a large token-to-type ratio), we can afford to estimate a model with a large K , which enables more subtle distinctions between words and contexts. When the dataset is relatively small, then K must be smaller too, or else the model may “memorize” the training data, and fail to generalize. Unfortunately, this general advice has not yet been formalized into any concrete formula for choosing K , and trial-and-error is still necessary. Overfitting can also be prevented by **dropout**, which involves randomly setting some elements of the computation to zero (Srivastava et al., 2014), forcing the learner not to rely too much on any particular dimension of the word or context vectors. The dropout rate must also be tuned on development data.

6.3.3 Gated recurrent neural networks

In principle, recurrent neural networks can propagate information across infinitely long sequences. But in practice, repeated applications of the nonlinear recurrence function causes this information to be quickly attenuated. The same problem affects learning: back-propagation can lead to **vanishing gradients** that decay to zero, or **exploding gradients** that increase towards infinity (Bengio et al., 1994). The exploding gradient problem can be addressed by clipping gradients at some maximum value (Pascanu et al., 2013). The other issues must be addressed by altering the model itself.

The **long short-term memory (LSTM)** (Hochreiter and Schmidhuber, 1997) is a popular variant of RNNs that is more robust to these problems. This model augments the hidden state \mathbf{h}_m with a **memory cell** c_m . The value of the memory cell at each time m is a gated sum of two quantities: its previous value c_{m-1} , and an “update” \tilde{c}_m , which is computed from the current input x_m and the previous hidden state \mathbf{h}_{m-1} . The next state \mathbf{h}_m is then computed from the memory cell. Because the memory cell is not passed through a non-linear squashing function during the update, it is possible for information to propagate through the network over long distances.

⁴See <https://www.tensorflow.org/tutorials/recurrent> (retrieved Feb 8, 2018).

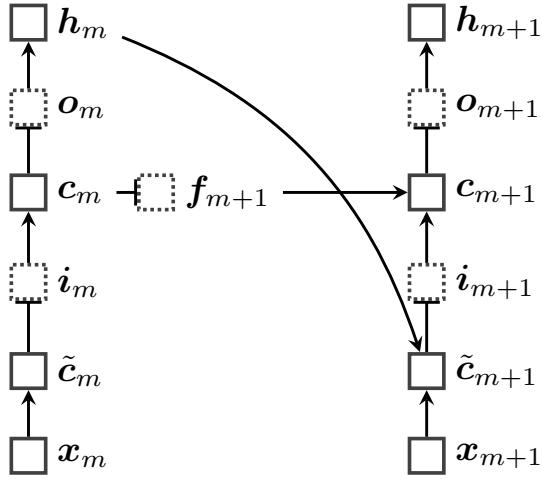


Figure 6.2: The long short-term memory (LSTM) architecture. Gates are shown in boxes with dotted edges. In an LSTM language model, each h_m would be used to predict the next word w_{m+1} .

The gates are functions of the input and previous hidden state. They are computed from elementwise sigmoid activations, $\sigma(x) = (1 + \exp(-x))^{-1}$, ensuring that their values will be in the range $[0, 1]$. They can therefore be viewed as soft, differentiable logic gates. The LSTM architecture is shown in Figure 6.2, and the complete update equations are:

$$f_{m+1} = \sigma(\Theta^{(h \rightarrow f)} h_m + \Theta^{(x \rightarrow f)} x_{m+1} + b_f) \quad \text{forget gate} \quad [6.36]$$

$$i_{m+1} = \sigma(\Theta^{(h \rightarrow i)} h_m + \Theta^{(x \rightarrow i)} x_{m+1} + b_i) \quad \text{input gate} \quad [6.37]$$

$$\tilde{c}_{m+1} = \tanh(\Theta^{(h \rightarrow c)} h_m + \Theta^{(x \rightarrow c)} x_{m+1}) \quad \text{update candidate} \quad [6.38]$$

$$c_{m+1} = f_{m+1} \odot c_m + i_{m+1} \odot \tilde{c}_{m+1} \quad \text{memory cell update} \quad [6.39]$$

$$o_{m+1} = \sigma(\Theta^{(h \rightarrow o)} h_m + \Theta^{(x \rightarrow o)} x_{m+1} + b_o) \quad \text{output gate} \quad [6.40]$$

$$h_{m+1} = o_{m+1} \odot \tanh(c_{m+1}) \quad \text{output.} \quad [6.41]$$

3288 The operator \odot is an elementwise (Hadamard) product. Each gate is controlled by a vec-
 3289 tor of weights, which parametrize the previous hidden state (e.g., $\Theta^{(h \rightarrow f)}$) and the current
 3290 input (e.g., $\Theta^{(x \rightarrow f)}$), plus a vector offset (e.g., b_f). The overall operation can be infor-
 3291 mally summarized as $(h_m, c_m) = \text{LSTM}(x_m, (h_{m-1}, c_{m-1}))$, with (h_m, c_m) representing
 3292 the LSTM state after reading token m .

3293 The LSTM outperforms standard recurrent neural networks across a wide range of
 3294 problems. It was first used for language modeling by Sundermeyer et al. (2012), but can
 3295 be applied more generally: the vector h_m can be treated as a complete representation of

3296 the input sequence up to position m , and can be used for any labeling task on a sequence
 3297 of tokens, as we will see in the next chapter.

3298 There are several LSTM variants, of which the Gated Recurrent Unit (Cho et al., 2014)
 3299 is one of the more well known. Many software packages implement a variety of RNN
 3300 architectures, so choosing between them is simple from a user’s perspective. Jozefowicz
 3301 et al. (2015) provide an empirical comparison of various modeling choices circa 2015.

3302 6.4 Evaluating language models

3303 Language modeling is not usually an application in itself: language models are typically
 3304 components of larger systems, and they would ideally be evaluated **extrinsically**. This
 3305 means evaluating whether the language model improves performance on the application
 3306 task, such as machine translation or speech recognition. But this is often hard to do, and
 3307 depends on details of the overall system which may be irrelevant to language modeling.
 3308 In contrast, **intrinsic evaluation** is task-neutral. Better performance on intrinsic metrics
 3309 may be expected to improve extrinsic metrics across a variety of tasks, but there is always
 3310 the risk of over-optimizing the intrinsic metric. This section discusses some intrinsic met-
 3311 rics, but keep in mind the importance of performing extrinsic evaluations to ensure that
 3312 intrinsic performance gains carry over to the applications that we care about.

3313 6.4.1 Held-out likelihood

The goal of probabilistic language models is to accurately measure the probability of sequences of word tokens. Therefore, an intrinsic evaluation metric is the likelihood that the language model assigns to **held-out data**, which is not used during training. Specifically, we compute,

$$\ell(\mathbf{w}) = \sum_{m=1}^M \log p(w_m | w_{m-1}, \dots, w_1), \quad [6.42]$$

3314 treating the entire held-out corpus as a single stream of tokens.

3315 Typically, unknown words are mapped to the $\langle \text{UNK} \rangle$ token. This means that we have
 3316 to estimate some probability for $\langle \text{UNK} \rangle$ on the training data. One way to do this is to fix
 3317 the vocabulary \mathcal{V} to the $V - 1$ words with the highest counts in the training data, and then
 3318 convert all other tokens to $\langle \text{UNK} \rangle$. Other strategies for dealing with out-of-vocabulary
 3319 terms are discussed in § 6.5.

3320 **6.4.2 Perplexity**

Held-out likelihood is usually presented as **perplexity**, which is a deterministic transformation of the log-likelihood into an information-theoretic quantity,

$$\text{Perplex}(\mathbf{w}) = 2^{-\frac{\ell(\mathbf{w})}{M}}, \quad [6.43]$$

3321 where M is the total number of tokens in the held-out corpus.

3322 Lower perplexities correspond to higher likelihoods, so lower scores are better on this
3323 metric — it is better to be less perplexed. Here are some special cases:

- 3324 • In the limit of a perfect language model, probability 1 is assigned to the held-out
3325 corpus, with $\text{Perplex}(\mathbf{w}) = 2^{-\frac{1}{M} \log_2 1} = 2^0 = 1$.
- 3326 • In the opposite limit, probability zero is assigned to the held-out corpus, which cor-
3327 responds to an infinite perplexity, $\text{Perplex}(\mathbf{w}) = 2^{-\frac{1}{M} \log_2 0} = 2^\infty = \infty$.
- 3328 • Assume a uniform, unigram model in which $p(w_i) = \frac{1}{V}$ for all words in the vocab-
3329 ular. Then,

$$\begin{aligned} \log_2(\mathbf{w}) &= \sum_{m=1}^M \log_2 \frac{1}{V} = - \sum_{m=1}^M \log_2 V = -M \log_2 V \\ \text{Perplex}(\mathbf{w}) &= 2^{\frac{1}{M} M \log_2 V} \\ &= 2^{\log_2 V} \\ &= V. \end{aligned}$$

3328 This is the “worst reasonable case” scenario, since you could build such a language
3329 model without even looking at the data.

3330 In practice, language models tend to give perplexities in the range between 1 and V .
3331 A small benchmark dataset is the **Penn Treebank**, which contains roughly a million to-
3332 kens; its vocabulary is limited to 10,000 words, with all other tokens mapped a special
3333 $\langle \text{UNK} \rangle$ symbol. On this dataset, a well-smoothed 5-gram model achieves a perplexity of
3334 141 (Mikolov and Zweig, Mikolov and Zweig), and an LSTM language model achieves
3335 perplexity of roughly 80 (Zaremba, Sutskever, and Vinyals, Zaremba et al.). Various en-
3336 hancements to the LSTM architecture can bring the perplexity below 60 (Merity et al.,
3337 2018). A larger-scale language modeling dataset is the 1B Word Benchmark (Chelba et al.,
3338 2013), which contains text from Wikipedia. On this dataset, a perplexities of around 25
3339 can be obtained by averaging together multiple LSTM language models (Jozefowicz et al.,
3340 2016).

3341 **6.5 Out-of-vocabulary words**

3342 So far, we have assumed a **closed-vocabulary** setting — the vocabulary \mathcal{V} is assumed to be
 3343 a finite set. In realistic application scenarios, this assumption may not hold. Consider, for
 3344 example, the problem of translating newspaper articles. The following sentence appeared
 3345 in a Reuters article on January 6, 2017:⁵

3346 The report said U.S. intelligence agencies believe Russian military intelligence,
 3347 the **GRU**, used intermediaries such as **WikiLeaks**, **DCLeaks.com** and the **Guc-**
 3348 **cifer** 2.0 "persona" to release emails...

3349 Suppose that you trained a language model on the Gigaword corpus,⁶ which was released
 3350 in 2003. The bolded terms either did not exist at this date, or were not widely known; they
 3351 are unlikely to be in the vocabulary. The same problem can occur for a variety of other
 3352 terms: new technologies, previously unknown individuals, new words (e.g., *hashtag*), and
 3353 numbers.

3354 One solution is to simply mark all such terms with a special token, $\langle \text{UNK} \rangle$. While
 3355 training the language model, we decide in advance on the vocabulary (often the K most
 3356 common terms), and mark all other terms in the training data as $\langle \text{UNK} \rangle$. If we do not want
 3357 to determine the vocabulary size in advance, an alternative approach is to simply mark
 3358 the first occurrence of each word type as $\langle \text{UNK} \rangle$.

3359 But it often better to make distinctions about the likelihood of various unknown words.
 3360 This is particularly important in languages that have rich morphological systems, with
 3361 many inflections for each word. For example, Portuguese is only moderately complex
 3362 from a morphological perspective, yet each verb has dozens of inflected forms (see Fig-
 3363 ure 4.3b). In such languages, there will be many word types that we do not encounter in a
 3364 corpus, which are nonetheless predictable from the morphological rules of the language.
 3365 To use a somewhat contrived English example, if *transfenestrate* is in the vocabulary, our
 3366 language model should assign a non-zero probability to the past tense *transfenestrated*,
 3367 even if it does not appear in the training data.

3368 One way to accomplish this is to supplement word-level language models with **character-**
 3369 **level language models**. Such models can use n -grams or RNNs, but with a fixed vocab-
 3370 uary equal to the set of ASCII or Unicode characters. For example Ling et al. (2015)
 3371 propose an LSTM model over characters, and Kim (2014) employ a **convolutional neural**
 3372 **network** (LeCun and Bengio, 1995). A more linguistically motivated approach is to seg-
 3373 ment words into meaningful subword units, known as **morphemes** (see chapter 9). For

⁵Bayoumy, Y. and Strobel, W. (2017, January 6). U.S. intel report: Putin directed cy-
 ber campaign to help Trump. *Reuters*. Retrieved from <http://www.reuters.com/article/us-usa-russia-cyber-idUSKBN14Q1T8> on January 7, 2017.

⁶<https://catalog.ldc.upenn.edu/LDC2003T05>

3374 example, Botha and Blunsom (2014) induce vector representations for morphemes, which
 3375 they build into a log-bilinear language model; Bhatia et al. (2016) incorporate morpheme
 3376 vectors into an LSTM.

3377 Additional resources

3378 A variety of neural network architectures have been applied to language modeling. No-
 3379 table earlier non-recurrent architectures include the neural probabilistic language model (Ben-
 3380 gio et al., 2003) and the log-bilinear language model (Mnih and Hinton, 2007). Much more
 3381 detail on these models can be found in the text by Goodfellow et al. (2016).

3382 Exercises

3383 1. Prove that n -gram language models give valid probabilities if the n -gram probabili-
 3384 ties are valid. Specifically, assume that,

$$\sum_{w_m}^{\mathcal{V}} p(w_m | w_{m-1}, w_{m-2}, \dots, w_{m-n+1}) = 1 \quad [6.44]$$

3385 for all contexts $(w_{m-1}, w_{m-2}, \dots, w_{m-n+1})$. Prove that $\sum_w p_n(w) = 1$ for all $w \in \mathcal{V}^*$,
 3386 where p_n is the probability under an n -gram language model. Your proof should
 3387 proceed by induction. You should handle the start-of-string case $p(w_1 | \underbrace{\square, \dots, \square}_{n-1})$,

3388 but you need not handle the end-of-string token.

3389 2. • Can you show that RNN language models are valid using a similar proof tech-
 3390 nique to the one in the previous problem?
 3391 • Let $p_r(w)$ indicate the probability of w under RNN r . An ensemble of RNN
 3392 language models computes the probability,

$$p(w) = \frac{1}{R} \sum_{r=1}^R p_r(w). \quad [6.45]$$

3393 Does an ensemble of RNN language models compute a valid probability?

3394 3. Consider a unigram language model over a vocabulary of size V . Suppose that a
 3395 word appears m times in a corpus with M tokens in total. With Lidstone smoothing
 3396 of α , for what values of m is the smoothed probability greater than the unsmoothed
 3397 probability?

3398 **Chapter 7**

3399 **Sequence labeling**

3400 The goal of sequence labeling is to assign tags to words, or more generally, to assign dis-
3401 crete labels to discrete elements in a sequence. There are many applications of sequence
3402 labeling in natural language processing, and chapter 8 presents an overview. A classic ap-
3403 plication is **part-of-speech tagging**, which involves tagging each word by its grammatical
3404 category. Coarse-grained grammatical categories include **NOUNs**, which describe things,
3405 properties, or ideas, and **VERBs**, which describe actions and events. Consider a simple
3406 input:

3407 (7.1) They can fish.

3408 A dictionary of coarse-grained part-of-speech tags might include **NOUN** as the only valid
3409 tag for *they*, but both **NOUN** and **VERB** as potential tags for *can* and *fish*. An accurate se-
3410 quence labeling algorithm should select the verb tag for both *can* and *fish* in (7.1), but it
3411 should select the noun tags for the same two words in the phrase *can of fish*.

3412 **7.1 Sequence labeling as classification**

One way to solve a tagging problem is to turn it into a classification problem. Let $f((\mathbf{w}, m), y)$ indicate the feature function for tag y at position m in the sequence $\mathbf{w} = (w_1, w_2, \dots, w_M)$. A simple tagging model would have a single base feature, the word itself:

$$f((\mathbf{w} = \text{they can fish}, m = 1), \text{N}) = (\text{they}, \text{N}) \quad [7.1]$$

$$f((\mathbf{w} = \text{they can fish}, m = 2), \text{V}) = (\text{can}, \text{V}) \quad [7.2]$$

$$f((\mathbf{w} = \text{they can fish}, m = 3), \text{V}) = (\text{fish}, \text{V}). \quad [7.3]$$

3413 Here the feature function takes three arguments as input: the sentence to be tagged (e.g.,
3414 *they can fish*), the proposed tag (e.g., N or V), and the index of the token to which this tag

3415 is applied. This simple feature function then returns a single feature: a tuple including
 3416 the word to be tagged and the tag that has been proposed. If the vocabulary size is V
 3417 and the number of tags is K , then there are $V \times K$ features. Each of these features must
 3418 be assigned a weight. These weights can be learned from a labeled dataset using a clas-
 3419 sification algorithm such as perceptron, but this isn't necessary in this case: it would be
 3420 equivalent to define the classification weights directly, with $\theta_{w,y} = 1$ for the tag y most
 3421 frequently associated with word w , and $\theta_{w,y} = 0$ for all other tags.

However, it is easy to see that this simple classification approach cannot correctly tag both *they can fish* and *can of fish*, because *can* and *fish* are grammatically ambiguous. To handle both of these cases, the tagger must rely on context, such as the surrounding words. We can build context into the feature set by incorporating the surrounding words as additional features:

$$\begin{aligned} f((\mathbf{w} = \text{they can fish}, 1), \mathbf{N}) = & \{(w_m = \text{they}, y_m = \mathbf{N}), \\ & (w_{m-1} = \square, y_m = \mathbf{N}), \\ & (w_{m+1} = \text{can}, y_m = \mathbf{N})\} \end{aligned} \quad [7.4]$$

$$\begin{aligned} f((\mathbf{w} = \text{they can fish}, 2), \mathbf{V}) = & \{(w_m = \text{can}, y_m = \mathbf{V}), \\ & (w_{m-1} = \text{they}, y_m = \mathbf{V}), \\ & (w_{m+1} = \text{fish}, y_m = \mathbf{V})\} \end{aligned} \quad [7.5]$$

$$\begin{aligned} f((\mathbf{w} = \text{they can fish}, 3), \mathbf{V}) = & \{(w_m = \text{fish}, y_m = \mathbf{V}), \\ & (w_{m-1} = \text{can}, y_m = \mathbf{V}), \\ & (w_{m+1} = \blacksquare, y_m = \mathbf{V})\}. \end{aligned} \quad [7.6]$$

3422 These features contain enough information that a tagger should be able to choose the
 3423 right tag for the word *fish*: words that come after *can* are likely to be verbs, so the feature
 3424 $(w_{m-1} = \text{can}, y_m = \mathbf{V})$ should have a large positive weight.

3425 However, even with this enhanced feature set, it may be difficult to tag some se-
 3426 quences correctly. One reason is that there are often relationships between the tags them-
 3427 selves. For example, in English it is relatively rare for a verb to follow another verb —
 3428 particularly if we differentiate MODAL verbs like *can* and *should* from more typical verbs,
 3429 like *give*, *transcend*, and *befuddle*. We would like to incorporate preferences against tag se-
 3430 quences like VERB-VERB, and in favor of tag sequences like NOUN-VERB. The need for
 3431 such preferences is best illustrated by a **garden path sentence**:

3432 (7.2) The old man the boat.

3433 Grammatically, the word *the* is a DETERMINER. When you read the sentence, what
 3434 part of speech did you first assign to *old*? Typically, this word is an ADJECTIVE — abbrevi-
 3435 ated as J — which is a class of words that modify nouns. Similarly, *man* is usually a noun.
 3436 The resulting sequence of tags is D J N D N. But this is a mistaken “garden path” inter-
 3437 pretation, which ends up leading nowhere. It is unlikely that a determiner would directly

follow a noun,¹ and it is particularly unlikely that the entire sentence would lack a verb. The only possible verb in (7.2) is the word *man*, which can refer to the act of maintaining and piloting something — often boats. But if *man* is tagged as a verb, then *old* is seated between a determiner and a verb, and must be a noun. And indeed, adjectives often have a second interpretation as nouns when used in this way (e.g., *the young*, *the restless*). This reasoning, in which the labeling decisions are intertwined, cannot be applied in a setting where each tag is produced by an independent classification decision.

7.2 Sequence labeling as structure prediction

As an alternative, think of the entire sequence of tags as a label itself. For a given sequence of words $\mathbf{w} = (w_1, w_2, \dots, w_M)$, there is a set of possible taggings $\mathcal{Y}(\mathbf{w}) = \mathcal{Y}^M$, where $\mathcal{Y} = \{\text{N, V, D, ...}\}$ refers to the set of individual tags, and \mathcal{Y}^M refers to the set of tag sequences of length M . We can then treat the sequence labeling problem as a classification problem in the label space $\mathcal{Y}(\mathbf{w})$,

$$\hat{\mathbf{y}} = \underset{\mathbf{y} \in \mathcal{Y}(\mathbf{w})}{\operatorname{argmax}} \Psi(\mathbf{w}, \mathbf{y}), \quad [7.7]$$

where $\mathbf{y} = (y_1, y_2, \dots, y_M)$ is a sequence of M tags, and Ψ is a scoring function on pairs of sequences, $V^M \times \mathcal{Y}^M \mapsto \mathbb{R}$. Such a function can include features that capture the relationships between tagging decisions, such as the preference that determiners not follow nouns, or that all sentences have verbs.

Given that the label space is exponentially large in the length of the sequence M , can it ever be practical to perform tagging in this way? The problem of making a series of interconnected labeling decisions is known as **inference**. Because natural language is full of interrelated grammatical structures, inference is a crucial aspect of natural language processing. In English, it is not unusual to have sentences of length $M = 20$; part-of-speech tag sets vary in size from 10 to several hundred. Taking the low end of this range, we have $|\mathcal{Y}(\mathbf{w}_{1:M})| \approx 10^{20}$, one hundred billion billion possible tag sequences. Enumerating and scoring each of these sequences would require an amount of work that is exponential in the sequence length, so inference is intractable.

However, the situation changes when we restrict the scoring function. Suppose we choose a function that decomposes into a sum of local parts,

$$\Psi(\mathbf{w}, \mathbf{y}) = \sum_{m=1}^{M+1} \psi(\mathbf{w}, y_m, y_{m-1}, m), \quad [7.8]$$

where each $\psi(\cdot)$ scores a local part of the tag sequence. Note that the sum goes up to $M+1$, so that we can include a score for a special end-of-sequence tag, $\psi(\mathbf{w}_{1:M}, \diamond, y_M, M+1)$. We also define a special tag to begin the sequence, $y_0 \triangleq \diamond$.

¹The main exception occurs with ditransitive verbs, such as *They gave the winner a trophy*.

3467 In a linear model, local scoring function can be defined as a dot product of weights
 3468 and features,

$$\psi(\mathbf{w}_{1:M}, y_m, y_{m-1}, m) = \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, y_m, y_{m-1}, m). \quad [7.9]$$

3469 The feature vector \mathbf{f} can consider the entire input \mathbf{w} , and can look at pairs of adjacent
 3470 tags. This is a step up from per-token classification: the weights can assign low scores
 3471 to infelicitous tag pairs, such as noun-determiner, and high scores for frequent tag pairs,
 3472 such as determiner-noun and noun-verb.

In the example *they can fish*, a minimal feature function would include features for word-tag pairs (sometimes called **emission features**) and tag-tag pairs (sometimes called **transition features**):

$$\begin{aligned} \mathbf{f}(\mathbf{w} = \text{they can fish}, \mathbf{y} = \text{N V V}) &= \sum_{m=1}^{M+1} \mathbf{f}(\mathbf{w}, y_m, y_{m-1}, m) \\ &= \mathbf{f}(\mathbf{w}, \text{N}, \diamond, 1) \\ &\quad + \mathbf{f}(\mathbf{w}, \text{V}, \text{N}, 2) \\ &\quad + \mathbf{f}(\mathbf{w}, \text{V}, \text{V}, 3) \\ &\quad + \mathbf{f}(\mathbf{w}, \blacklozenge, \text{V}, 4) \end{aligned} \quad [7.10]$$

$$\begin{aligned} &= (w_m = \text{they}, y_m = \text{N}) + (y_m = \text{N}, y_{m-1} = \diamond) \\ &\quad + (w_m = \text{can}, y_m = \text{V}) + (y_m = \text{V}, y_{m-1} = \text{N}) \\ &\quad + (w_m = \text{fish}, y_m = \text{V}) + (y_m = \text{V}, y_{m-1} = \text{V}) \\ &\quad + (y_m = \blacklozenge, y_{m-1} = \text{V}). \end{aligned} \quad [7.11]$$

$$[7.12]$$

3473 There are seven active features for this example: one for each word-tag pair, and one
 3474 for each tag-tag pair, including a final tag $y_{M+1} = \blacklozenge$. These features capture the two main
 3475 sources of information for part-of-speech tagging in English: which tags are appropriate
 3476 for each word, and which tags tend to follow each other in sequence. Given appropriate
 3477 weights for these features, taggers can achieve high accuracy, even for difficult cases like
 3478 *the old man the boat*. We will now discuss how this restricted scoring function enables
 3479 efficient inference, through the **Viterbi algorithm** (Viterbi, 1967).

3480 **7.3 The Viterbi algorithm**

By decomposing the scoring function into a sum of local parts, it is possible to rewrite the tagging problem as follows:

$$\hat{\mathbf{y}} = \operatorname{argmax}_{\mathbf{y} \in \mathcal{Y}(\mathbf{w})} \Psi(\mathbf{w}, \mathbf{y}) \quad [7.13]$$

$$= \operatorname{argmax}_{\mathbf{y}_{1:M}} \sum_{m=1}^{M+1} \psi(\mathbf{w}, y_m, y_{m-1}, m) \quad [7.14]$$

$$= \operatorname{argmax}_{\mathbf{y}_{1:M}} \sum_{m=1}^{M+1} s_m(y_m, y_{m-1}), \quad [7.15]$$

3481 where the final line simplifies the notation with the shorthand,

$$s_m(y_m, y_{m-1}) \triangleq \psi(\mathbf{w}_{1:M}, y_m, y_{m-1}, m). \quad [7.16]$$

This inference problem can be solved efficiently using **dynamic programming**, a algorithmic technique for reusing work in recurrent computations. As is often the case in dynamic programming, we begin by solving an auxiliary problem: rather than finding the best tag sequence, we simply compute the *score* of the best tag sequence,

$$\max_{\mathbf{y}_{1:M}} \Psi(\mathbf{w}, \mathbf{y}_{1:M}) = \max_{\mathbf{y}_{1:M}} \sum_{m=1}^{M+1} s_m(y_m, y_{m-1}). \quad [7.17]$$

This score involves a maximization over all tag sequences of length M , written $\max_{\mathbf{y}_{1:M}}$. This maximization can be broken into two pieces,

$$\max_{\mathbf{y}_{1:M}} \Psi(\mathbf{w}, \mathbf{y}_{1:M}) = \max_{y_M} \max_{\mathbf{y}_{1:M-1}} \sum_{m=1}^{M+1} s_m(y_m, y_{m-1}), \quad [7.18]$$

which simply says that we maximize over the final tag y_M , and we maximize over all “prefixes”, $\mathbf{y}_{1:M-1}$. But within the sum of scores, only the final term $s_{M+1}(\blacklozenge, y_M)$ depends on y_M . We can pull this term out of the second maximization,

$$\max_{\mathbf{y}_{1:M}} \Psi(\mathbf{w}, \mathbf{y}_{1:M}) = \max_{y_M} s_{M+1}(\blacklozenge, y_M) + \max_{\mathbf{y}_{1:M-1}} \sum_{m=1}^M s_m(y_m, y_{m-1}). \quad [7.19]$$

This same reasoning can be applied recursively to the second term of Equation 7.19, pulling out $s_M(y_M, y_{M-1})$, and so on. We can formalize this idea by defining an auxiliary

Algorithm 11 The Viterbi algorithm. Each $s_m(k, k')$ is a local score for tag $y_m = k$ and $y_{m-1} = k'$.

```

for  $k \in \{0, \dots, K\}$  do
     $v_1(k) = s_1(k, \diamond)$ 
for  $m \in \{2, \dots, M\}$  do
    for  $k \in \{0, \dots, K\}$  do
         $v_m(k) = \max_{k'} s_m(k, k') + v_{m-1}(k')$ 
         $b_m(k) = \operatorname{argmax}_{k'} s_m(k, k') + v_{m-1}(k')$ 
     $y_M = \operatorname{argmax}_k s_{M+1}(\blacklozenge, k) + v_M(k)$ 
    for  $m \in \{M-1, \dots, 1\}$  do
         $y_m = b_m(y_{m+1})$ 
return  $\mathbf{y}_{1:M}$ 
```

Viterbi variable,

$$v_m(y_m) \triangleq \max_{\mathbf{y}_{1:m-1}} \sum_{n=1}^m s_n(y_n, y_{n-1}) \quad [7.20]$$

$$= \max_{y_{m-1}} s_m(y_m, y_{m-1}) + \max_{\mathbf{y}_{1:m-2}} \sum_{n=1}^{m-1} s_n(y_n, y_{n-1}) \quad [7.21]$$

$$= \max_{y_{m-1}} s_m(y_m, y_{m-1}) + v_{m-1}(y_{m-1}). \quad [7.22]$$

3482 The variable $v_m(k)$ represents the score of the best sequence of length m ending in tag k .

Each set of Viterbi variables is computed from the local score $s_m(y_m, y_{m-1})$, and from the previous set of Viterbi variables. The initial condition of the recurrence is simply the first score,

$$v_1(y_1) \triangleq s_1(y_1, \diamond). \quad [7.23]$$

The maximum overall score for the sequence is then the final Viterbi variable,

$$\max_{\mathbf{y}_{1:M}} \Psi(\mathbf{w}_{1:M}, \mathbf{y}_{1:M}) = v_{M+1}(\blacklozenge). \quad [7.24]$$

3483 Thus, the score of the best labeling for the sequence can be computed in a single forward
 3484 sweep: first compute all variables $v_1(\cdot)$ from Equation 7.23, and then compute all variables
 3485 $v_2(\cdot)$ from the recurrence Equation 7.22, and continue until reaching the final variable
 3486 $v_{M+1}(\blacklozenge)$.

3487 Graphically, it is customary to arrange these variables in a structure known as a **trellis**,
 3488 shown in Figure 7.1. Each column indexes a token m in the sequence, and each row

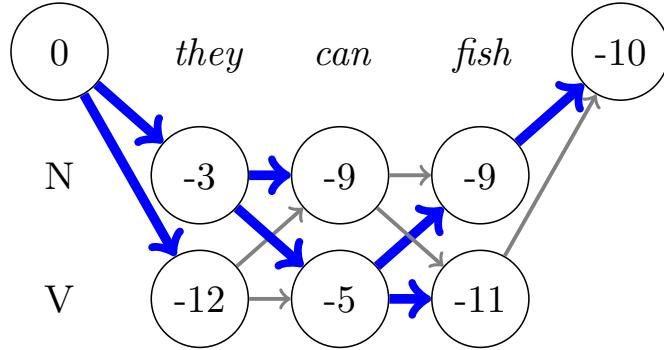


Figure 7.1: The trellis representation of the Viterbi variables, for the example *they can fish*, using the weights shown in Table 7.1.

3489 indexes a tag in \mathcal{Y} ; every $v_{m-1}(k)$ is connected to every $v_m(k')$, that $v_m(k')$ is computed
 3490 from $v_{m-1}(k)$. Special nodes are set aside for the start and end states.

3491 Our real goal is to find the best scoring sequence, not simply to compute its score.
 3492 But solving the auxiliary problem gets us almost all the way there. Recall that each $v_m(k)$
 3493 represents the score of the best tag sequence ending in that tag k in position m . To compute
 3494 this, we maximize over possible values of y_{m-1} . If we keep track of the “argmax” tag that
 3495 maximizes this choice at each step, then we can walk backwards from the final tag, and
 3496 recover the optimal tag sequence. This is indicated in Figure 7.1 by the solid blue lines,
 3497 which we trace back from the final position. These “back-pointers” are written $b_m(k)$,
 3498 indicating the optimal tag y_{m-1} on the path to $Y_m = k$.

3499 The complete Viterbi algorithm is shown in Algorithm 11. When computing the initial
 3500 Viterbi variables $v_1(\cdot)$, we use a special tag, \diamond , to indicate the start of the sequence. When
 3501 computing the final tag Y_M , we use another special tag, \blacklozenge , to indicate the end of the
 3502 sequence. Linguistically, these special tags enable the use of transition features for the tags
 3503 that begin and end the sequence: for example, conjunctions are unlikely to end sentences
 3504 in English, so we would like a low score for $s_{M+1}(\blacklozenge, CC)$; nouns are relatively likely to
 3505 appear at the beginning of sentences, so we would like a high score for $s_1(N, \diamond)$, assuming
 3506 the noun tag is compatible with the first word token w_1 .

3507 **Complexity** If there are K tags and M positions in the sequence, then there are $M \times K$
 3508 Viterbi variables to compute. Computing each variable requires finding a maximum over
 3509 K possible predecessor tags. The total time complexity of populating the trellis is therefore
 3510 $\mathcal{O}(MK^2)$, with an additional factor for the number of active features at each position.
 3511 After completing the trellis, we simply trace the backwards pointers to the beginning of
 3512 the sequence, which takes $\mathcal{O}(M)$ operations.

	<i>they</i>	<i>can</i>	<i>fish</i>	
N	-2	-3	-3	
V	-10	-1	-3	

(a) Weights for emission features.

	N	V	♦
◊	-1	-2	$-\infty$
N	-3	-1	-1
V	-1	-3	-1

(b) Weights for transition features. The “from” tags are on the columns, and the “to” tags are on the rows.

Table 7.1: Feature weights for the example trellis shown in Figure 7.1. Emission weights from \diamond and ♦ are implicitly set to $-\infty$.3513

7.3.1 Example

3514 Consider the minimal tagset $\{N, V\}$, corresponding to nouns and verbs. Even in this
 3515 tagset, there is considerable ambiguity: for example, the words *can* and *fish* can each take
 3516 both tags. Of the $2 \times 2 \times 2 = 8$ possible taggings for the sentence *they can fish*, four are
 3517 possible given these possible tags, and two are grammatical.²

3518 The values in the trellis in Figure 7.1 are computed from the feature weights defined in
 3519 Table 7.1. We begin with $v_1(N)$, which has only one possible predecessor, the start tag \diamond .
 3520 This score is therefore equal to $s_1(N, \diamond) = -2 - 1 = -3$, which is the sum of the scores for
 3521 the emission and transition features respectively; the backpointer is $b_1(N) = \diamond$. The score
 3522 for $v_1(V)$ is computed in the same way: $s_1(V, \diamond) = -10 - 2 = -12$, and again $b_1(V) = \diamond$.
 3523 The backpointers are represented in the figure by thick lines.

Things get more interesting at $m = 2$. The score $v_2(N)$ is computed by maximizing over the two possible predecessors,

$$v_2(N) = \max(v_1(N) + s_2(N, N), v_1(V) + s_2(N, V)) \quad [7.25]$$

$$= \max(-3 - 3 - 3, -12 - 3 - 1) = -9 \quad [7.26]$$

$$b_2(N) = N. \quad [7.27]$$

This continues until reaching $v_4(\diamond)$, which is computed as,

$$v_4(\diamond) = \max(v_3(N) + s_4(\diamond, N), v_3(V) + s_4(\diamond, V)) \quad [7.28]$$

$$= \max(-9 + 0 - 1, -11 + 0 - 1) \quad [7.29]$$

$$= -10, \quad [7.30]$$

3524 so $b_4(\diamond) = N$. As there is no emission w_4 , the emission features have scores of zero.

²The tagging *they/N can/V fish/N* corresponds to the scenario of putting fish into cans, or perhaps of firing them.

3525 To compute the optimal tag sequence, we walk backwards from here, next checking
 3526 $b_3(N) = V$, and then $b_2(V) = N$, and finally $b_1(N) = \diamond$. This yields $y = (N, V, N)$, which
 3527 corresponds to the linguistic interpretation of the fishes being put into cans.

3528 **7.3.2 Higher-order features**

3529 The Viterbi algorithm was made possible by a restriction of the scoring function to local
 3530 parts that consider only pairs of adjacent tags. We can think of this as a bigram language
 3531 model over tags. A natural question is how to generalize Viterbi to tag trigrams, which
 3532 would involve the following decomposition:

$$\Psi(\mathbf{w}, \mathbf{y}) = \sum_{m=1}^{M+2} f(\mathbf{w}, y_m, y_{m-1}, y_{m-2}, m), \quad [7.31]$$

3533 where $y_{-1} = \diamond$ and $y_{M+2} = \blacklozenge$.

3534 One solution is to create a new tagset $\mathcal{Y}^{(2)}$ from the Cartesian product of the original
 3535 tagset with itself, $\mathcal{Y}^{(2)} = \mathcal{Y} \times \mathcal{Y}$. The tags in this product space are ordered pairs, rep-
 3536 resenting adjacent tags at the token level: for example, the tag (N, V) would represent a
 3537 noun followed by a verb. Transitions between such tags must be consistent: we can have a
 3538 transition from (N, V) to (V, N) (corresponding to the tag sequence $N V N$), but not from
 3539 (N, V) to (N, N) , which would not correspond to any coherent tag sequence. This con-
 3540 straint can be enforced in feature weights, with $\theta_{((a,b),(c,d))} = -\infty$ if $b \neq c$. The remaining
 3541 feature weights can encode preferences for and against various tag trigrams.

3542 In the Cartesian product tag space, there are K^2 tags, suggesting that the time com-
 3543 plexity will increase to $\mathcal{O}(MK^4)$. However, it is unnecessary to max over predecessor tag
 3544 bigrams that are incompatible with the current tag bigram. By exploiting this constraint,
 3545 it is possible to limit the time complexity to $\mathcal{O}(MK^3)$. The space complexity grows to
 3546 $\mathcal{O}(MK^2)$, since the trellis must store all possible predecessors of each tag. In general, the
 3547 time and space complexity of higher-order Viterbi grows exponentially with the order of
 3548 the tag n -grams that are considered in the feature decomposition.

3549 **7.4 Hidden Markov Models**

3550 Let us now consider how to learn the scores $s_m(y, y')$ that parametrize the Viterbi sequence
 3551 labeling algorithm, beginning with a probabilistic approach. Recall from § 2.1 that the
 3552 probabilistic Naïve Bayes classifier selects the label y to maximize $p(y | \mathbf{x}) \propto p(y, \mathbf{x})$. In
 3553 probabilistic sequence labeling, our goal is similar: select the tag sequence that maximizes
 3554 $p(y | \mathbf{w}) \propto p(\mathbf{y}, \mathbf{w})$. The locality restriction in Equation 7.8 can be viewed as a conditional
 3555 independence assumption on the random variables \mathbf{y} .

Algorithm 12 Generative process for the hidden Markov model

```

 $y_0 \leftarrow \diamond,$     $m \leftarrow 1$ 
repeat
     $y_m \sim \text{Categorical}(\lambda_{y_{m-1}})$             $\triangleright$  sample the current tag
     $w_m \sim \text{Categorical}(\phi_{y_m})$             $\triangleright$  sample the current word
until  $y_m = \blacklozenge$             $\triangleright$  terminate when the stop symbol is generated

```

3556 Naïve Bayes was introduced as a generative model — a probabilistic story that ex-
 3557 plains the observed data as well as the hidden label. A similar story can be constructed
 3558 for probabilistic sequence labeling: first, the tags are drawn from a prior distribution; next,
 3559 the tokens are drawn from a conditional likelihood. However, for inference to be tractable,
 3560 additional independence assumptions are required. First, the probability of each token
 3561 depends only on its tag, and not on any other element in the sequence:

$$p(w | y) = \prod_{m=1}^M p(w_m | y_m). \quad [7.32]$$

3562 Second, each tag y_m depends only on its predecessor,

$$p(y) = \prod_{m=1}^M p(y_m | y_{m-1}), \quad [7.33]$$

3563 where $y_0 = \diamond$ in all cases. Due to this **Markov assumption**, probabilistic sequence labeling
 3564 models are known as **hidden Markov models** (HMMs).

3565 The generative process for the hidden Markov model is shown in Algorithm 12. Given
 3566 the parameters λ and ϕ , we can compute $p(w, y)$ for any token sequence w and tag se-
 3567 quence y . The HMM is often represented as a **graphical model** (Wainwright and Jordan,
 3568 2008), as shown in Figure 7.2. This representation makes the independence assumptions
 3569 explicit: if a variable v_1 is probabilistically conditioned on another variable v_2 , then there
 3570 is an arrow $v_2 \rightarrow v_1$ in the diagram. If there are no arrows between v_1 and v_2 , they
 3571 are **conditionally independent**, given each variable's **Markov blanket**. In the hidden
 3572 Markov model, the Markov blanket for each tag y_m includes the “parent” y_{m-1} , and the
 3573 “children” y_{m+1} and w_m .³

3574 It is important to reflect on the implications of the HMM independence assumptions.
 3575 A non-adjacent pair of tags y_m and y_n are conditionally independent; if $m < n$ and we
 3576 are given y_{n-1} , then y_m offers no additional information about y_n . However, if we are
 3577 not given any information about the tags in a sequence, then all tags are probabilistically
 3578 coupled.

³In general graphical models, a variable's Markov blanket includes its parents, children, and its children's other parents (Murphy, 2012).

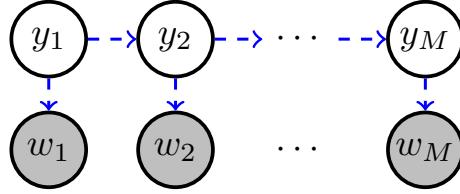


Figure 7.2: Graphical representation of the hidden Markov model. Arrows indicate probabilistic dependencies.

3579 7.4.1 Estimation

3580 The hidden Markov model has two groups of parameters:

3581 **Emission probabilities.** The probability $p_e(w_m | y_m; \phi)$ is the emission probability, since
3582 the words are treated as probabilistically “emitted”, conditioned on the tags.

3583 **Transition probabilities.** The probability $p_t(y_m | y_{m-1}; \lambda)$ is the transition probability,
3584 since it assigns probability to each possible tag-to-tag transition.

Both of these groups of parameters are typically computed from smoothed relative frequency estimation on a labeled corpus (see § 6.2 for a review of smoothing). The unsmoothed probabilities are,

$$\begin{aligned}\phi_{k,i} &\triangleq \Pr(W_m = i | Y_m = k) = \frac{\text{count}(W_m = i, Y_m = k)}{\text{count}(Y_m = k)} \\ \lambda_{k,k'} &\triangleq \Pr(Y_m = k' | Y_{m-1} = k) = \frac{\text{count}(Y_m = k', Y_{m-1} = k)}{\text{count}(Y_{m-1} = k)}.\end{aligned}$$

3585 Smoothing is more important for the emission probability than the transition probability,
3586 because the vocabulary is much larger than the number of tags.

3587 7.4.2 Inference

3588 The goal of inference in the hidden Markov model is to find the highest probability tag
3589 sequence,

$$\hat{y} = \underset{y}{\operatorname{argmax}} p(y | w). \quad [7.34]$$

3590 As in Naïve Bayes, it is equivalent to find the tag sequence with the highest *log*-probability,
3591 since the logarithm is a monotonically increasing function. It is furthermore equivalent
3592 to maximize the joint probability $p(y, w) = p(y | w) \times p(w) \propto p(y | w)$, which is pro-
3593 portional to the conditional probability. Putting these observations together, the inference

3594 problem can be reformulated as,

$$\hat{\mathbf{y}} = \operatorname{argmax}_{\mathbf{y}} \log p(\mathbf{y}, \mathbf{w}). \quad [7.35]$$

We can now apply the HMM independence assumptions:

$$\log p(\mathbf{y}, \mathbf{w}) = \log p(\mathbf{y}) + \log p(\mathbf{w} \mid \mathbf{y}) \quad [7.36]$$

$$= \sum_{m=1}^{M+1} \log p_Y(y_m \mid y_{m-1}) + \log p_{W|Y}(w_m \mid y_m) \quad [7.37]$$

$$= \sum_{m=1}^{M+1} \log \lambda_{y_m, y_{m-1}} + \log \phi_{y_m, w_m} \quad [7.38]$$

$$= \sum_{m=1}^{M+1} s_m(y_m, y_{m-1}), \quad [7.39]$$

where,

$$s_m(y_m, y_{m-1}) \triangleq \log \lambda_{y_m, y_{m-1}} + \log \phi_{y_m, w_m}, \quad [7.40]$$

3595 and,

$$\phi_{\diamond, w} = \begin{cases} 1, & w = \blacksquare \\ 0, & \text{otherwise,} \end{cases} \quad [7.41]$$

3596 which ensures that the stop tag \diamond can only be applied to the final token \blacksquare .

This derivation shows that HMM inference can be viewed as an application of the Viterbi decoding algorithm, given an appropriately defined scoring function. The local score $s_m(y_m, y_{m-1})$ can be interpreted probabilistically,

$$s_m(y_m, y_{m-1}) = \log p_y(y_m \mid y_{m-1}) + \log p_{w|y}(w_m \mid y_m) \quad [7.42]$$

$$= \log p(y_m, w_m \mid y_{m-1}). \quad [7.43]$$

Now recall the definition of the Viterbi variables,

$$v_m(y_m) = \max_{y_{m-1}} s_m(y_m, y_{m-1}) + v_{m-1}(y_{m-1}) \quad [7.44]$$

$$= \max_{y_{m-1}} \log p(y_m, w_m \mid y_{m-1}) + v_{m-1}(y_{m-1}). \quad [7.45]$$

By setting $v_{m-1}(y_{m-1}) = \max_{\mathbf{y}_{1:m-2}} \log p(\mathbf{y}_{1:m-1}, \mathbf{w}_{1:m-1})$, we obtain the recurrence,

$$v_m(y_m) = \max_{y_{m-1}} \log p(y_m, w_m \mid y_{m-1}) + \max_{\mathbf{y}_{1:m-2}} \log p(\mathbf{y}_{1:m-1}, \mathbf{w}_{1:m-1}) \quad [7.46]$$

$$= \max_{\mathbf{y}_{1:m-1}} \log p(y_m, w_m \mid y_{m-1}) + \log p(\mathbf{y}_{1:m-1}, \mathbf{w}_{1:m-1}) \quad [7.47]$$

$$= \max_{\mathbf{y}_{1:m-1}} \log p(\mathbf{y}_{1:m}, \mathbf{w}_{1:m}). \quad [7.48]$$

In words, the Viterbi variable $v_m(y_m)$ is the log probability of the best tag sequence ending in y_m , joint with the word sequence $w_{1:m}$. The log probability of the best complete tag sequence is therefore,

$$\max_{\mathbf{y}_{1:M}} \log p(\mathbf{y}_{1:M+1}, \mathbf{w}_{1:M+1}) = v_{M+1}(\spadesuit) \quad [7.49]$$

***Viterbi as an example of the max-product algorithm** The Viterbi algorithm can also be implemented using probabilities, rather than log-probabilities. In this case, each $v_m(y_m)$ is equal to,

$$v_m(y_m) = \max_{\mathbf{y}_{1:m-1}} p(\mathbf{y}_{1:m-1}, y_m, \mathbf{w}_{1:m}) \quad [7.50]$$

$$= \max_{y_{m-1}} p(y_m, w_m | y_{m-1}) \times \max_{\mathbf{y}_{1:m-2}} p(\mathbf{y}_{1:m-2}, y_{m-1}, \mathbf{w}_{1:m-1}) \quad [7.51]$$

$$= \max_{y_{m-1}} p(y_m, w_m | y_{m-1}) \times v_{m-1}(y_{m-1}) \quad [7.52]$$

$$= p_{w|y}(w_m | y_m) \times \max_{y_{m-1}} p_y(y_m | y_{m-1}) \times v_{m-1}(y_{m-1}). \quad [7.53]$$

3597 Each Viterbi variable is computed by *maximizing* over a set of *products*. Thus, the Viterbi
 3598 algorithm is a special case of the **max-product algorithm** for inference in graphical mod-
 3599 els (Wainwright and Jordan, 2008). However, the product of probabilities tends towards
 3600 zero over long sequences, so the log-probability version of Viterbi is recommended in
 3601 practical implementations.

3602 7.5 Discriminative sequence labeling with features

3603 Today, hidden Markov models are rarely used for supervised sequence labeling. This is
 3604 because HMMs are limited to only two phenomena:

- 3605 • word-tag compatibility, via the emission probability $p_{W|Y}(w_m | y_m)$;
- 3606 • local context, via the transition probability $p_Y(y_m | y_{m-1})$.

3607 The Viterbi algorithm permits the inclusion of richer information in the local scoring func-
 3608 tion $\psi(\mathbf{w}_{1:M}, y_m, y_{m-1}, m)$, which can be defined as a weighted sum of arbitrary local *fea-*
 3609 *tures*,

$$\psi(\mathbf{w}, y_m, y_{m-1}, m) = \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, y_m, y_{m-1}, m), \quad [7.54]$$

3610 where \mathbf{f} is a locally-defined feature function, and $\boldsymbol{\theta}$ is a vector of weights.

The local decomposition of the scoring function Ψ is reflected in a corresponding decomposition of the feature function:

$$\Psi(\mathbf{w}, \mathbf{y}) = \sum_{m=1}^{M+1} \psi(\mathbf{w}, y_m, y_{m-1}, m) \quad [7.55]$$

$$= \sum_{m=1}^{M+1} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, y_m, y_{m-1}, m) \quad [7.56]$$

$$= \boldsymbol{\theta} \cdot \sum_{m=1}^{M+1} \mathbf{f}(\mathbf{w}, y_m, y_{m-1}, m) \quad [7.57]$$

$$= \boldsymbol{\theta} \cdot \mathbf{f}^{(\text{global})}(\mathbf{w}, \mathbf{y}_{1:M}), \quad [7.58]$$

3611 where $\mathbf{f}^{(\text{global})}(\mathbf{w}, \mathbf{y})$ is a global feature vector, which is a sum of local feature vectors,

$$\mathbf{f}^{(\text{global})}(\mathbf{w}, \mathbf{y}) = \sum_{m=1}^{M+1} \mathbf{f}(\mathbf{w}_{1:M}, y_m, y_{m-1}, m), \quad [7.59]$$

3612 with $y_{M+1} = \diamond$ and $y_0 = \diamond$ by construction.

3613 Let's now consider what additional information these features might encode.

3614 **Word affix features.** Consider the problem of part-of-speech tagging on the first four
3615 lines of the poem *Jabberwocky* (Carroll, 1917):

3616 (7.3) 'Twas brillig, and the slithy toves
3617 Did gyre and gimble in the wabe:
3618 All mimsy were the borogoves,
3619 And the mome raths outgrabe.

3620 Many of these words were made up by the author of the poem, so a corpus would offer
3621 no information about their probabilities of being associated with any particular part of
3622 speech. Yet it is not so hard to see what their grammatical roles might be in this passage.
3623 Context helps: for example, the word *slithy* follows the determiner *the*, so it is probably a
3624 noun or adjective. Which do you think is more likely? The suffix *-thy* is found in a number
3625 of adjectives, like *frothy*, *healthy*, *pithy*, *worthy*. It is also found in a handful of nouns — e.g.,
3626 *apathy*, *sympathy* — but nearly all of these have the longer coda *-pathy*, unlike *slithy*. So the
3627 suffix gives some evidence that *slithy* is an adjective, and indeed it is: later in the text we
3628 find that it is a combination of the adjectives *lithe* and *slimy*.⁴

⁴Morphology is the study of how words are formed from smaller linguistic units. Computational approaches to morphological analysis are touched on in chapter 9; Bender (2013) provides a good overview of the underlying linguistic principles.

3629 **Fine-grained context.** The hidden Markov model captures contextual information in the
 3630 form of part-of-speech tag bigrams. But sometimes, the necessary contextual information
 3631 is more specific. Consider the noun phrases *this fish* and *these fish*. Many part-of-speech
 3632 tagsets distinguish between singular and plural nouns, but do not distinguish between
 3633 singular and plural determiners.⁵ A hidden Markov model would be unable to correctly
 3634 label *fish* as singular or plural in both of these cases, because it only has access to two
 3635 features: the preceding tag (determiner in both cases) and the word (*fish* in both cases).
 3636 The classification-based tagger discussed in § 7.1 had the ability to use preceding and suc-
 3637 ceeding words as features, and it can also be incorporated into a Viterbi-based sequence
 3638 labeler as a local feature.

Example Consider the tagging D J N (determiner, adjective, noun) for the sequence *the slithy toves*, so that

$$\mathbf{w} = \text{the slithy toves}$$

$$\mathbf{y} = \text{D J N}.$$

Let's create the feature vector for this example, assuming that we have word-tag features (indicated by W), tag-tag features (indicated by T), and suffix features (indicated by M). You can assume that you have access to a method for extracting the suffix *-thy* from *slithy*, *-es* from *toves*, and \emptyset from *the*, indicating that this word has no suffix.⁶ The resulting feature vector is,

$$\begin{aligned} f(\text{the slithy toves, D J N}) &= f(\text{the slithy toves, D}, \diamond, 1) \\ &\quad + f(\text{the slithy toves, J}, D, 2) \\ &\quad + f(\text{the slithy toves, N}, J, 3) \\ &\quad + f(\text{the slithy toves}, \blacklozenge, N, 4) \\ &= \{(T : \diamond, D), (W : \text{the}, D), (M : \emptyset, D), \\ &\quad (T : D, J), (W : \text{slithy}, J), (M : -thy, J), \\ &\quad (T : J, N), (W : \text{toves}, N), (M : -es, N) \\ &\quad (T : N, \blacklozenge)\}. \end{aligned}$$

3639 These examples show that local features can incorporate information that lies beyond
 3640 the scope of a hidden Markov model. Because the features are local, it is possible to apply
 3641 the Viterbi algorithm to identify the optimal sequence of tags. The remaining question

⁵For example, the Penn Treebank tagset follows these conventions.

⁶Such a system is called a **morphological segmenter**. The task of morphological segmentation is briefly described in § 9.1.4.4; a well known segmenter is Morfessor (Creutz and Lagus, 2007). In real applications, a typical approach is to include features for all orthographic suffixes up to some maximum number of characters: for *slithy*, we would have suffix features for *-y*, *-hy*, and *-thy*.

3642 is how to estimate the weights on these features. § 2.2 presented three main types of
 3643 discriminative classifiers: perceptron, support vector machine, and logistic regression.
 3644 Each of these classifiers has a structured equivalent, enabling it to be trained from labeled
 3645 sequences rather than individual tokens.

3646 **7.5.1 Structured perceptron**

The perceptron classifier is trained by increasing the weights for features that are associated with the correct label, and decreasing the weights for features that are associated with incorrectly predicted labels:

$$\hat{y} = \operatorname{argmax}_{y \in \mathcal{Y}} \theta \cdot f(\mathbf{x}, y) \quad [7.60]$$

$$\theta^{(t+1)} \leftarrow \theta^{(t)} + f(\mathbf{x}, y) - f(\mathbf{x}, \hat{y}). \quad [7.61]$$

We can apply exactly the same update in the case of structure prediction,

$$\hat{\mathbf{y}} = \operatorname{argmax}_{\mathbf{y} \in \mathcal{Y}(\mathbf{w})} \theta \cdot f(\mathbf{w}, \mathbf{y}) \quad [7.62]$$

$$\theta^{(t+1)} \leftarrow \theta^{(t)} + f(\mathbf{w}, \mathbf{y}) - f(\mathbf{w}, \hat{\mathbf{y}}). \quad [7.63]$$

3647 This learning algorithm is called **structured perceptron**, because it learns to predict the
 3648 structured output \mathbf{y} . The only difference is that instead of computing \hat{y} by enumerating
 3649 the entire set \mathcal{Y} , the Viterbi algorithm is used to efficiently search the set of possible tag-
 3650 gings, \mathcal{Y}^M . Structured perceptron can be applied to other structured outputs as long as
 3651 efficient inference is possible. As in perceptron classification, weight averaging is crucial
 3652 to get good performance (see § 2.2.2).

Example For the example *they can fish*, suppose that the reference tag sequence is $\mathbf{y}^{(i)} =$
 N V V, but the tagger incorrectly returns the tag sequence $\hat{\mathbf{y}} = \text{N V N}$. Assuming a model
 with features for emissions (w_m, y_m) and transitions (y_{m-1}, y_m) , the corresponding structured
 perceptron update is:

$$\theta_{(fish,V)} \leftarrow \theta_{(fish,V)} + 1, \quad \theta_{(fish,N)} \leftarrow \theta_{(fish,N)} - 1 \quad [7.64]$$

$$\theta_{(V,V)} \leftarrow \theta_{(V,V)} + 1, \quad \theta_{(V,N)} \leftarrow \theta_{(V,N)} - 1 \quad [7.65]$$

$$\theta_{(V,\blacklozenge)} \leftarrow \theta_{(V,\blacklozenge)} + 1, \quad \theta_{(N,\blacklozenge)} \leftarrow \theta_{(N,\blacklozenge)} - 1. \quad [7.66]$$

3653 **7.5.2 Structured support vector machines**

3654 Large-margin classifiers such as the support vector machine improve on the perceptron by
 3655 pushing the classification boundary away from the training instances. The same idea can

3656 be applied to sequence labeling. A support vector machine in which the output is a struc-
 3657 tured object, such as a sequence, is called a **structured support vector machine** (Tsochan-
 3658 taridis et al., 2004).⁷

3659 In classification, we formalized the large-margin constraint as,

$$\forall \mathbf{y} \neq \mathbf{y}^{(i)}, \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}, \mathbf{y}^{(i)}) - \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}, \mathbf{y}) \geq 1, \quad [7.67]$$

3660 requiring a margin of at least 1 between the scores for all labels \mathbf{y} that are not equal to the
 3661 correct label $\mathbf{y}^{(i)}$. The weights $\boldsymbol{\theta}$ are then learned by constrained optimization (see § 2.3.2).

3662 This idea can be applied to sequence labeling by formulating an equivalent set of con-
 3663 straints for all possible labelings $\mathcal{Y}(\mathbf{w})$ for an input \mathbf{w} . However, there are two problems.
 3664 First, in sequence labeling, some predictions are more wrong than others: we may miss
 3665 only one tag out of fifty, or we may get all fifty wrong. We would like our learning algo-
 3666 rithm to be sensitive to this difference. Second, the number of constraints is equal to the
 3667 number of possible labelings, which is exponentially large in the length of the sequence.

3668 The first problem can be addressed by adjusting the constraint to require larger mar-
 3669 gins for more serious errors. Let $c(\mathbf{y}^{(i)}, \hat{\mathbf{y}}) \geq 0$ represent the *cost* of predicting label $\hat{\mathbf{y}}$ when
 3670 the true label is $\mathbf{y}^{(i)}$. We can then generalize the margin constraint,

$$\forall \mathbf{y}, \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}^{(i)}, \mathbf{y}^{(i)}) - \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}^{(i)}, \mathbf{y}) \geq c(\mathbf{y}^{(i)}, \mathbf{y}). \quad [7.68]$$

3671 This cost-augmented margin constraint specializes to the constraint in Equation 7.67 if we
 3672 choose the delta function $c(\mathbf{y}^{(i)}, \mathbf{y}) = \delta((\mathbf{y}^{(i)} \neq \mathbf{y}))$. A more expressive cost function is
 3673 the **Hamming cost**,

$$c(\mathbf{y}^{(i)}, \mathbf{y}) = \sum_{m=1}^M \delta(y_m^{(i)} \neq y_m), \quad [7.69]$$

3674 which computes the number of errors in \mathbf{y} . By incorporating the cost function as the
 3675 margin constraint, we require that the true labeling be separated from the alternatives by
 3676 a margin that is proportional to the number of incorrect tags in each alternative labeling.

The second problem is that the number of constraints is exponential in the length
 of the sequence. This can be addressed by focusing on the prediction $\hat{\mathbf{y}}$ that *maximally*
 violates the margin constraint. This prediction can be identified by solving the following
cost-augmented decoding problem:

$$\hat{\mathbf{y}} = \operatorname{argmax}_{\mathbf{y} \neq \mathbf{y}^{(i)}} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}^{(i)}, \mathbf{y}) - \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}^{(i)}, \mathbf{y}^{(i)}) + c(\mathbf{y}^{(i)}, \mathbf{y}) \quad [7.70]$$

$$= \operatorname{argmax}_{\mathbf{y} \neq \mathbf{y}^{(i)}} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}^{(i)}, \mathbf{y}) + c(\mathbf{y}^{(i)}, \mathbf{y}), \quad [7.71]$$

⁷This model is also known as a **max-margin Markov network** (Taskar et al., 2003), emphasizing that the scoring function is constructed from a sum of components, which are Markov independent.

3677 where in the second line we drop the term $\theta \cdot f(\mathbf{w}^{(i)}, \mathbf{y}^{(i)})$, which is constant in \mathbf{y} .

We can now reformulate the margin constraint for sequence labeling,

$$\theta \cdot f(\mathbf{w}^{(i)}, \mathbf{y}^{(i)}) - \max_{\mathbf{y} \in \mathcal{Y}(\mathbf{w})} (\theta \cdot f(\mathbf{w}^{(i)}, \mathbf{y}) + c(\mathbf{y}^{(i)}, \mathbf{y})) \geq 0. \quad [7.72]$$

3678 If the score for $\theta \cdot f(\mathbf{w}^{(i)}, \mathbf{y}^{(i)})$ is greater than the cost-augmented score for all alternatives,
 3679 then the constraint will be met. The name “cost-augmented decoding” is due to the fact
 3680 that the objective includes the standard decoding problem, $\max_{\hat{\mathbf{y}} \in \mathcal{Y}(\mathbf{w})} \theta \cdot f(\mathbf{w}, \hat{\mathbf{y}})$, plus
 3681 an additional term for the cost. Essentially, we want to train against predictions that are
 3682 strong and wrong: they should score highly according to the model, yet incur a large loss
 3683 with respect to the ground truth. Training adjusts the weights to reduce the score of these
 3684 predictions.

3685 For cost-augmented decoding to be tractable, the cost function must decompose into
 3686 local parts, just as the feature function $f(\cdot)$ does. The Hamming cost, defined above,
 3687 obeys this property. To perform cost-augmented decoding using the Hamming cost, we
 3688 need only to add features $f_m(y_m) = \delta(y_m \neq y_m^{(i)})$, and assign a constant weight of 1 to
 3689 these features. Decoding can then be performed using the Viterbi algorithm.⁸

As with large-margin classifiers, it is possible to formulate the learning problem in an unconstrained form, by combining a regularization term on the weights and a Lagrangian for the constraints:

$$\min_{\theta} \frac{1}{2} \|\theta\|_2^2 - C \left(\sum_i \theta \cdot f(\mathbf{w}^{(i)}, \mathbf{y}^{(i)}) - \max_{\mathbf{y} \in \mathcal{Y}(\mathbf{w}^{(i)})} [\theta \cdot f(\mathbf{w}^{(i)}, \mathbf{y}) + c(\mathbf{y}^{(i)}, \mathbf{y})] \right), \quad [7.73]$$

3690 In this formulation, C is a parameter that controls the tradeoff between the regularization
 3691 term and the margin constraints. A number of optimization algorithms have been
 3692 proposed for structured support vector machines, some of which are discussed in § 2.3.2.
 3693 An empirical comparison by Kummerfeld et al. (2015) shows that stochastic subgradient
 3694 descent — which is essentially a cost-augmented version of the structured perceptron —
 3695 is highly competitive.

3696 7.5.3 Conditional random fields

3697 The **conditional random field** (CRF; Lafferty et al., 2001) is a conditional probabilistic
 3698 model for sequence labeling; just as structured perceptron is built on the perceptron clas-
 3699 sifier, conditional random fields are built on the logistic regression classifier.⁹ The basic

⁸Are there cost functions that do not decompose into local parts? Suppose we want to assign a constant loss c to any prediction $\hat{\mathbf{y}}$ in which k or more predicted tags are incorrect, and zero loss otherwise. This loss function is combinatorial over the predictions, and thus we cannot decompose it into parts.

⁹The name “Conditional Random Field” is derived from **Markov random fields**, a general class of models in which the probability of a configuration of variables is proportional to a product of scores across pairs (or

3700 probability model is,

$$p(\mathbf{y} \mid \mathbf{w}) = \frac{\exp(\Psi(\mathbf{w}, \mathbf{y}))}{\sum_{\mathbf{y}' \in \mathcal{Y}(\mathbf{w})} \exp(\Psi(\mathbf{w}, \mathbf{y}'))}. \quad [7.74]$$

3701 This is almost identical to logistic regression, but because the label space is now tag
 3702 sequences, we require efficient algorithms for both **decoding** (searching for the best tag
 3703 sequence given a sequence of words \mathbf{w} and a model θ) and for **normalizing** (summing
 3704 over all tag sequences). These algorithms will be based on the usual locality assumption
 3705 on the scoring function, $\Psi(\mathbf{w}, \mathbf{y}) = \sum_{m=1}^{M+1} \psi(\mathbf{w}, y_m, y_{m-1}, m)$.

3706 **7.5.3.1 Decoding in CRFs**

Decoding — finding the tag sequence $\hat{\mathbf{y}}$ that maximizes $p(\mathbf{y} \mid \mathbf{w})$ — is a direct application of the Viterbi algorithm. The key observation is that the decoding problem does not depend on the denominator of $p(\mathbf{y} \mid \mathbf{w})$,

$$\begin{aligned} \hat{\mathbf{y}} &= \operatorname{argmax}_{\mathbf{y}} \log p(\mathbf{y} \mid \mathbf{w}) \\ &= \operatorname{argmax}_{\mathbf{y}} \Psi(\mathbf{y}, \mathbf{w}) - \log \sum_{\mathbf{y}' \in \mathcal{Y}(\mathbf{w})} \exp \Psi(\mathbf{y}', \mathbf{w}) \\ &= \operatorname{argmax}_{\mathbf{y}} \Psi(\mathbf{y}, \mathbf{w}) = \operatorname{argmax}_{\mathbf{y}} \sum_{m=1}^{M+1} s_m(y_m, y_{m-1}). \end{aligned}$$

3707 This is identical to the decoding problem for structured perceptron, so the same Viterbi
 3708 recurrence as defined in Equation 7.22 can be used.

3709 **7.5.3.2 Learning in CRFs**

As with logistic regression, the weights θ are learned by minimizing the regularized negative log-probability,

$$\ell = \frac{\lambda}{2} \|\theta\|^2 - \sum_{i=1}^N \log p(\mathbf{y}^{(i)} \mid \mathbf{w}^{(i)}; \theta) \quad [7.75]$$

$$= \frac{\lambda}{2} \|\theta\|^2 - \sum_{i=1}^N \theta \cdot f(\mathbf{w}^{(i)}, \mathbf{y}^{(i)}) + \log \sum_{\mathbf{y}' \in \mathcal{Y}(\mathbf{w}^{(i)})} \exp (\theta \cdot f(\mathbf{w}^{(i)}, \mathbf{y}')), \quad [7.76]$$

more generally, cliques) of variables in a **factor graph**. In sequence labeling, the pairs of variables include all adjacent tags (y_m, y_{m-1}). The probability is *conditioned* on the words \mathbf{w} , which are always observed, motivating the term “conditional” in the name.

3710 where λ controls the amount of regularization. The final term in Equation 7.76 is a sum
 3711 over all possible labelings. This term is the log of the denominator in Equation 7.74, some-
 3712 times known as the **partition function**.¹⁰ There are $|\mathcal{Y}|^M$ possible labelings of an input of
 3713 size M , so we must again exploit the decomposition of the scoring function to compute
 3714 this sum efficiently.

The sum $\sum_{\mathbf{y} \in \mathcal{Y}^{w(i)}} \exp \Psi(\mathbf{y}, \mathbf{w})$ can be computed efficiently using the **forward recurrence**, which is closely related to the Viterbi recurrence. We first define a set of **forward variables**, $\alpha_m(y_m)$, which is equal to the sum of the scores of all paths leading to tag y_m at position m :

$$\alpha_m(y_m) \triangleq \sum_{\mathbf{y}_{1:m-1}} \exp \sum_{n=1}^m s_n(y_n, y_{n-1}) \quad [7.77]$$

$$= \sum_{\mathbf{y}_{1:m-1}} \prod_{n=1}^m \exp s_n(y_n, y_{n-1}). \quad [7.78]$$

Note the similarity to the definition of the Viterbi variable, $v_m(y_m) = \max_{\mathbf{y}_{1:m-1}} \sum_{n=1}^m s_n(y_n, y_{n-1})$. In the hidden Markov model, the Viterbi recurrence had an alternative interpretation as the max-product algorithm (see Equation 7.53); analogously, the forward recurrence is known as the **sum-product algorithm**, because of the form of [7.78]. The forward variable can also be computed through a recurrence:

$$\alpha_m(y_m) = \sum_{\mathbf{y}_{1:m-1}} \prod_{n=1}^m \exp s_n(y_n, y_{n-1}) \quad [7.79]$$

$$= \sum_{y_{m-1}} (\exp s_m(y_m, y_{m-1})) \sum_{\mathbf{y}_{1:m-2}} \prod_{n=1}^{m-1} \exp s_n(y_n, y_{n-1}) \quad [7.80]$$

$$= \sum_{y_{m-1}} (\exp s_m(y_m, y_{m-1})) \times \alpha_{m-1}(y_{m-1}). \quad [7.81]$$

Using the forward recurrence, it is possible to compute the denominator of the conditional probability,

$$\sum_{\mathbf{y} \in \mathcal{Y}(\mathbf{w})} \Psi(\mathbf{w}, \mathbf{y}) = \sum_{\mathbf{y}_{1:M}} s_{M+1}(\blacklozenge, y_M) \prod_{m=1}^M s_m(y_m, y_{m-1}) \quad [7.82]$$

$$= \alpha_{M+1}(\blacklozenge). \quad [7.83]$$

¹⁰The terminology of “potentials” and “partition functions” comes from statistical mechanics (Bishop, 2006).

The conditional log-likelihood can be rewritten,

$$\ell = \frac{\lambda}{2} \|\boldsymbol{\theta}\|^2 - \sum_{i=1}^N \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}^{(i)}, \mathbf{y}^{(i)}) + \log \alpha_{M+1}(\blacklozenge). \quad [7.84]$$

- 3715 Probabilistic programming environments, such as `Torch` (Collobert et al., 2011) and `dynet` (Neu-
 3716 big et al., 2017), can compute the gradient of this objective using automatic differentiation.
 3717 The programmer need only implement the forward algorithm as a computation graph.

As in logistic regression, the gradient of the likelihood with respect to the parameters is a difference between observed and expected feature counts:

$$\frac{d\ell}{d\theta_j} = \lambda \theta_j + \sum_{i=1}^N E[f_j(\mathbf{w}^{(i)}, \mathbf{y})] - f_j(\mathbf{w}^{(i)}, \mathbf{y}^{(i)}), \quad [7.85]$$

- 3718 where $f_j(\mathbf{w}^{(i)}, \mathbf{y}^{(i)})$ refers to the count of feature j for token sequence $\mathbf{w}^{(i)}$ and tag se-
 3719 quence $\mathbf{y}^{(i)}$. The expected feature counts are computed “under the hood” when automatic
 3720 differentiation is applied to Equation 7.84 (Eisner, 2016).

- 3721 Before the widespread use of automatic differentiation, it was common to compute
 3722 the feature expectations from marginal tag probabilities $p(y_m | \mathbf{w})$. These marginal prob-
 3723 abilities are sometimes useful on their own, and can be computed using the **forward-**
 3724 **backward algorithm**. This algorithm combines the forward recurrence with an equivalent
 3725 **backward recurrence**, which traverses the input from w_M back to w_1 .

3726 7.5.3.3 *Forward-backward algorithm

Marginal probabilities over tag bigrams can be written as,¹¹

$$\Pr(Y_{m-1} = k', Y_m = k | \mathbf{w}) = \frac{\sum_{\mathbf{y}: Y_m=k, Y_{m-1}=k'} \prod_{n=1}^M \exp s_n(y_n, y_{n-1})}{\sum_{\mathbf{y}'} \prod_{n=1}^M \exp s_n(y'_n, y'_{n-1})}. \quad [7.86]$$

The numerator sums over all tag sequences that include the transition $(Y_{m-1} = k') \rightarrow (Y_m = k)$. Because we are only interested in sequences that include the tag bigram, this sum can be decomposed into three parts: the *prefixes* $\mathbf{y}_{1:m-1}$, terminating in $Y_{m-1} = k'$; the

¹¹Recall the notational convention of upper-case letters for random variables, e.g. Y_m , and lower case letters for specific values, e.g., y_m , so that $Y_m = k$ is interpreted as the event of random variable Y_m taking the value k .

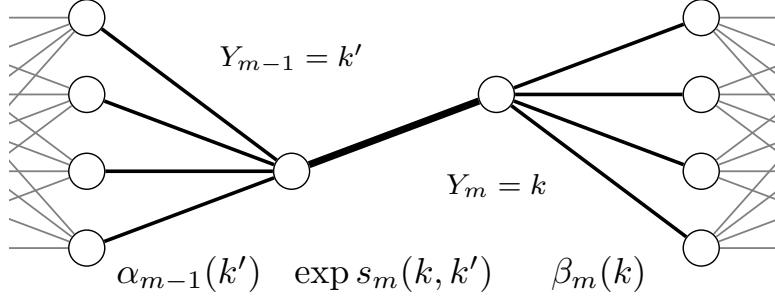


Figure 7.3: A schematic illustration of the computation of the marginal probability $\Pr(Y_{m-1} = k', Y_m = k)$, using the forward score $\alpha_{m-1}(k')$ and the backward score $\beta_m(k)$.

transition $(Y_{m-1} = k') \rightarrow (Y_m = k)$; and the *suffixes* $\mathbf{y}_{m:M}$, beginning with the tag $Y_m = k$:

$$\sum_{\mathbf{y}: Y_m=k, Y_{m-1}=k'} \prod_{n=1}^M \exp s_n(y_n, y_{n-1}) = \sum_{\mathbf{y}_{1:m-1}: Y_{m-1}=k'} \prod_{n=1}^{m-1} \exp s_n(y_n, y_{n-1}) \times \exp s_m(k, k') \times \sum_{\mathbf{y}_{m:M}: Y_m=k} \prod_{n=m+1}^{M+1} \exp s_n(y_n, y_{n-1}). \quad [7.87]$$

The result is product of three terms: a score that sums over all the ways to get to the position $(Y_{m-1} = k')$, a score for the transition from k' to k , and a score that sums over all the ways of finishing the sequence from $(Y_m = k)$. The first term of Equation 7.87 is equal to the **forward variable**, $\alpha_{m-1}(k')$. The third term — the sum over ways to finish the sequence — can also be defined recursively, this time moving over the trellis from right to left, which is known as the **backward recurrence**:

$$\beta_m(k) \triangleq \sum_{\mathbf{y}_{m:M}: Y_m=k} \prod_{n=m}^{M+1} \exp s_n(y_n, y_{n-1}) \quad [7.88]$$

$$= \sum_{k' \in \mathcal{Y}} \exp s_{m+1}(k', k) \sum_{\mathbf{y}_{m+1:M}: Y_m=k'} \prod_{n=m+1}^{M+1} \exp s_n(y_n, y_{n-1}) \quad [7.89]$$

$$= \sum_{k' \in \mathcal{Y}} \exp s_{m+1}(k', k) \times \beta_{m+1}(k'). \quad [7.90]$$

³⁷²⁷ To understand this computation, compare with the forward recurrence in Equation 7.81.

In practice, numerical stability demands that we work in the log domain,

$$\log \alpha_m(k) = \log \sum_{k' \in \mathcal{Y}} \exp (\log s_m(k, k') + \log \alpha_{m-1}(k')) \quad [7.91]$$

$$\log \beta_{m-1}(k) = \log \sum_{k' \in \mathcal{Y}} \exp (\log s_m(k', k) + \log \beta_m(k')). \quad [7.92]$$

The application of the forward and backward probabilities is shown in Figure 7.3. Both the forward and backward recurrences operate on the trellis, which implies a space complexity $\mathcal{O}(MK)$. Because both recurrences require computing a sum over K terms at each node in the trellis, their time complexity is $\mathcal{O}(MK^2)$.

7.6 Neural sequence labeling

In neural network approaches to sequence labeling, we construct a vector representation for each tagging decision, based on the word and its context. Neural networks can perform tagging as a per-token classification decision, or they can be combined with the Viterbi algorithm to tag the entire sequence globally.

7.6.1 Recurrent neural networks

Recurrent neural networks (RNNs) were introduced in chapter 6 as a language modeling technique, in which the context at token m is summarized by a recurrently-updated vector,

$$\mathbf{h}_m = g(\mathbf{x}_m, \mathbf{h}_{m-1}), \quad m = 1, 2, \dots, M,$$

where \mathbf{x}_m is the vector **embedding** of the token w_m and the function g defines the recurrence. The starting condition \mathbf{h}_0 is an additional parameter of the model. The long short-term memory (LSTM) is a more complex recurrence, in which a memory cell is through a series of gates, avoiding repeated application of the non-linearity. Despite these bells and whistles, both models share the basic architecture of recurrent updates across a sequence, and both will be referred to as RNNs here.

A straightforward application of RNNs to sequence labeling is to score each tag y_m as a linear function of \mathbf{h}_m :

$$\psi_m(y) = \beta_y \cdot \mathbf{h}_m \quad [7.93]$$

$$\hat{y}_m = \underset{y}{\operatorname{argmax}} \psi_m(y). \quad [7.94]$$

The score $\psi_m(y)$ can also be converted into a probability distribution using the usual softmax operation,

$$p(y | \mathbf{w}_{1:m}) = \frac{\exp \psi_m(y)}{\sum_{y' \in \mathcal{Y}} \exp \psi_m(y')}. \quad [7.95]$$

3746 Using this transformation, it is possible to train the tagger from the negative log-likelihood
 3747 of the tags, as in a conditional random field. Alternatively, a hinge loss or margin loss
 3748 objective can be constructed from the raw scores $\psi_m(y)$.

The hidden state \mathbf{h}_m accounts for information in the input leading up to position m , but it ignores the subsequent tokens, which may also be relevant to the tag y_m . This can be addressed by adding a second RNN, in which the input is reversed, running the recurrence from w_M to w_1 . This is known as a **bidirectional recurrent neural network** (Graves and Schmidhuber, 2005), and is specified as:

$$\overleftarrow{\mathbf{h}}_m = g(\mathbf{x}_m, \overleftarrow{\mathbf{h}}_{m+1}), \quad m = 1, 2, \dots, M. \quad [7.96]$$

3749 The hidden states of the left-to-right RNN are denoted $\overrightarrow{\mathbf{h}}_m$. The left-to-right and right-to-
 3750 left vectors are concatenated, $\mathbf{h}_m = [\overleftarrow{\mathbf{h}}_m; \overrightarrow{\mathbf{h}}_m]$. The scoring function in Equation 7.93 is
 3751 applied to this concatenated vector.

3752 Bidirectional RNN tagging has several attractive properties. Ideally, the representa-
 3753 tion \mathbf{h}_m summarizes the useful information from the surrounding context, so that it is not
 3754 necessary to design explicit features to capture this information. If the vector \mathbf{h}_m is an ad-
 3755 equate summary of this context, then it may not even be necessary to perform the tagging
 3756 jointly: in general, the gains offered by joint tagging of the entire sequence are diminished
 3757 as the individual tagging model becomes more powerful. Using backpropagation, the
 3758 word vectors \mathbf{x} can be trained “end-to-end”, so that they capture word properties that are
 3759 useful for the tagging task. Alternatively, if limited labeled data is available, we can use
 3760 word embeddings that are “pre-trained” from unlabeled data, using a language modeling
 3761 objective (as in § 6.3) or a related word embedding technique (see chapter 14). It is even
 3762 possible to combine both fine-tuned and pre-trained embeddings in a single model.

3763 **Neural structure prediction** The bidirectional recurrent neural network incorporates in-
 3764 formation from throughout the input, but each tagging decision is made independently.
 3765 In some sequence labeling applications, there are very strong dependencies between tags:
 3766 it may even be impossible for one tag to follow another. In such scenarios, the tagging
 3767 decision must be made jointly across the entire sequence.

3768 Neural sequence labeling can be combined with the Viterbi algorithm by defining the
 3769 local scores as:

$$s_m(y_m, y_{m-1}) = \beta_{y_m} \cdot \mathbf{h}_m + \eta_{y_{m-1}, y_m}, \quad [7.97]$$

3770 where \mathbf{h}_m is the RNN hidden state, β_{y_m} is a vector associated with tag y_m , and η_{y_{m-1}, y_m}
 3771 is a scalar parameter for the tag transition (y_{m-1}, y_m) . These local scores can then be
 3772 incorporated into the Viterbi algorithm for inference, and into the forward algorithm for
 3773 training. This model is shown in Figure 7.4. It can be trained from the conditional log-
 3774 likelihood objective defined in Equation 7.76, backpropagating to the tagging parameters

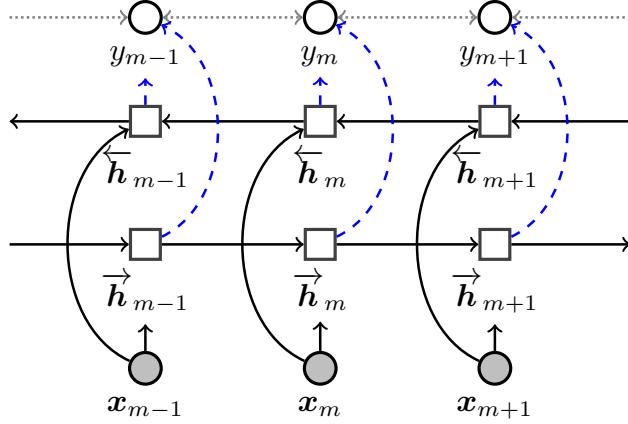


Figure 7.4: Bidirectional LSTM for sequence labeling. The solid lines indicate computation, the dashed lines indicate probabilistic dependency, and the dotted lines indicate the optional additional probabilistic dependencies between labels in the biLSTM-CRF.

3775 β and η , as well as the parameters of the RNN. This model is called the **LSTM-CRF**, due
 3776 to its combination of aspects of the long short-term memory and conditional random field
 3777 models (Huang et al., 2015).

3778 The LSTM-CRF is especially effective on the task of **named entity recognition** (Lample
 3779 et al., 2016), a sequence labeling task that is described in detail in § 8.3. This task has strong
 3780 dependencies between adjacent tags, so structure prediction is especially important.

3781 7.6.2 Character-level models

3782 As in language modeling, rare and unseen words are a challenge: if we encounter a word
 3783 that was not in the training data, then there is no obvious choice for the word embed-
 3784 ding x_m . One solution is to use a generic **unseen word** embedding for all such words.
 3785 However, in many cases, properties of unseen words can be guessed from their spellings.
 3786 For example, *whimsical* does not appear in the Universal Dependencies (UD) English Tree-
 3787 bank, yet the suffix *-al* makes it likely to be adjective; by the same logic, *unflinchingly* is
 3788 likely to be an adverb, and *barnacle* is likely to be a noun.

3789 In feature-based models, these morphological properties were handled by suffix fea-
 3790 tures; in a neural network, they can be incorporated by constructing the embeddings of
 3791 unseen words from their spellings or morphology. One way to do this is to incorporate
 3792 an additional layer of bidirectional RNNs, one for each word in the vocabulary (Ling
 3793 et al., 2015). For each such character-RNN, the inputs are the characters, and the output
 3794 is the concatenation of the final states of the left-facing and right-facing passes, $\phi_w =$

[$\vec{h}_{N_w}^{(w)}; \overleftarrow{h}_0^{(w)}$], where $\vec{h}_{N_w}^{(w)}$ is the final state of the right-facing pass for word w , and N_w is the number of characters in the word. The character RNN model is trained by back-propagation from the tagging objective. On the test data, the trained RNN is applied to out-of-vocabulary words (or all words), yielding inputs to the word-level tagging RNN. Other approaches to compositional word embeddings are described in § 14.7.1.

7.6.3 Convolutional Neural Networks for Sequence Labeling

One disadvantage of recurrent neural networks is that the architecture requires iterating through the sequence of inputs and predictions: each hidden vector h_m must be computed from the previous hidden vector h_{m-1} , before predicting the tag y_m . These iterative computations are difficult to parallelize, and fail to exploit the speedups offered by **graphics processing units (GPUs)** on operations such as matrix multiplication. **Convolutional neural networks** achieve better computational performance by predicting each label y_m from a set of matrix operations on the neighboring word embeddings, $x_{m-k:m+k}$ (Collobert et al., 2011). Because there is no hidden state to update, the predictions for each y_m can be computed in parallel. For more on convolutional neural networks, see § 3.4. Character-based word embeddings can also be computed using convolutional neural networks (Santos and Zadrozny, 2014).

7.7 *Unsupervised sequence labeling

In unsupervised sequence labeling, the goal is to induce a hidden Markov model from a corpus of *unannotated* text ($w^{(1)}, w^{(2)}, \dots, w^{(N)}$), where each $w^{(i)}$ is a sequence of length $M^{(i)}$. This is an example of the general problem of **structure induction**, which is the unsupervised version of structure prediction. The tags that result from unsupervised sequence labeling might be useful for some downstream task, or they might help us to better understand the language’s inherent structure. For part-of-speech tagging, it is common to use a tag dictionary that lists the allowed tags for each word, simplifying the problem (Christodoulopoulos et al., 2010).

Unsupervised learning in hidden Markov models can be performed using the **Baum-Welch algorithm**, which combines the forward-backward algorithm (§ 7.5.3.3) with expectation-maximization (EM; § 5.1.2). In the M-step, the HMM parameters from expected counts:

$$\Pr(W = i \mid Y = k) = \phi_{k,i} = \frac{E[\text{count}(W = i, Y = k)]}{E[\text{count}(Y = k)]}$$

$$\Pr(Y_m = k \mid Y_{m-1} = k') = \lambda_{k',k} = \frac{E[\text{count}(Y_m = k, Y_{m-1} = k')]}{E[\text{count}(Y_{m-1} = k')]} \quad 3820$$

3821 The expected counts are computed in the E-step, using the forward and backward
 3822 recurrences. The local scores follow the usual definition for hidden Markov models,

$$s_m(k, k') = \log p_E(w_m | Y_m = k; \phi) + \log p_T(Y_m = k | Y_{m-1} = k'; \lambda). \quad [7.98]$$

The expected transition counts for a single instance are,

$$E[\text{count}(Y_m = k, Y_{m-1} = k') | \mathbf{w}] = \sum_{m=1}^M \Pr(Y_{m-1} = k', Y_m = k | \mathbf{w}) \quad [7.99]$$

$$= \frac{\sum_{\mathbf{y}: Y_m=k, Y_{m-1}=k'} \prod_{n=1}^M \exp s_n(y_n, y_{n-1})}{\sum_{\mathbf{y}'} \prod_{n=1}^M \exp s_n(y'_n, y'_{n-1})}. \quad [7.100]$$

As described in § 7.5.3.3, these marginal probabilities can be computed from the forward-backward recurrence,

$$\Pr(Y_{m-1} = k', Y_m = k | \mathbf{w}) = \frac{\alpha_{m-1}(k') \times s_m(k, k') \times \beta_m(k)}{\alpha_{M+1}(\blacklozenge)}. \quad [7.101]$$

In a hidden Markov model, each element of the forward-backward computation has a special interpretation:

$$\alpha_{m-1}(k') = p(Y_{m-1} = k', \mathbf{w}_{1:m-1}) \quad [7.102]$$

$$s_m(k, k') = p(Y_m = k, w_m | Y_{m-1} = k') \quad [7.103]$$

$$\beta_m(k) = p(\mathbf{w}_{m+1:M} | Y_m = k). \quad [7.104]$$

Applying the conditional independence assumptions of the hidden Markov model (defined in Algorithm 12), the product is equal to the joint probability of the tag bigram and the entire input,

$$\begin{aligned} \alpha_{m-1}(k') \times s_m(k, k') \times \beta_m(k) &= p(Y_{m-1} = k', \mathbf{w}_{1:m-1}) \\ &\quad \times p(Y_m = k, w_m | Y_{m-1} = k') \\ &\quad \times p(\mathbf{w}_{m+1:M} | Y_m = k) \\ &= p(Y_{m-1} = k', Y_m = k, \mathbf{w}_{1:M}). \end{aligned} \quad [7.105]$$

Dividing by $\alpha_{M+1}(\blacklozenge) = p(\mathbf{w}_{1:M})$ gives the desired probability,

$$\frac{\alpha_{m-1}(k') \times s_m(k, k') \times \beta_m(k)}{\alpha_{M+1}(\blacklozenge)} = \frac{p(Y_{m-1} = k', Y_m = k, \mathbf{w}_{1:M})}{p(\mathbf{w}_{1:M})} \quad [7.106]$$

$$= \Pr(Y_{m-1} = k', Y_m = k | \mathbf{w}_{1:M}). \quad [7.107]$$

3823 The expected emission counts can be computed in a similar manner, using the product
 3824 $\alpha_m(k) \times \beta_m(k)$.

3825 **7.7.1 Linear dynamical systems**

3826 The forward-backward algorithm can be viewed as Bayesian state estimation in a discrete
 3827 state space. In a continuous state space, $\mathbf{y}_m \in \mathbb{R}^K$, the equivalent algorithm is the **Kalman**
 3828 **smoother**. It also computes marginals $p(\mathbf{y}_m | \mathbf{x}_{1:M})$, using a similar two-step algorithm
 3829 of forward and backward passes. Instead of computing a trellis of values at each step, the
 3830 Kalman smoother computes a probability density function $q_{\mathbf{y}_m}(\mathbf{y}_m; \boldsymbol{\mu}_m, \Sigma_m)$, character-
 3831 ized by a mean $\boldsymbol{\mu}_m$ and a covariance Σ_m around the latent state. Connections between the
 3832 Kalman Smoother and the forward-backward algorithm are elucidated by Minka (1999)
 3833 and Murphy (2012).

3834 **7.7.2 Alternative unsupervised learning methods**

As noted in § 5.5, expectation-maximization is just one of many techniques for structure induction. One alternative is to use **Markov Chain Monte Carlo (MCMC)** sampling algorithms, which are briefly described in § 5.5.1. For the specific case of sequence labeling, Gibbs sampling can be applied by iteratively sampling each tag y_m conditioned on all the others (Finkel et al., 2005):

$$p(y_m | \mathbf{y}_{-m}, \mathbf{w}_{1:M}) \propto p(w_m | y_m) p(y_m | \mathbf{y}_{-m}). \quad [7.108]$$

3835 Gibbs Sampling has been applied to unsupervised part-of-speech tagging by Goldwater
 3836 and Griffiths (2007). **Beam sampling** is a more sophisticated sampling algorithm, which
 3837 randomly draws entire sequences $\mathbf{y}_{1:M}$, rather than individual tags y_m ; this algorithm
 3838 was applied to unsupervised part-of-speech tagging by Van Gael et al. (2009). Spectral
 3839 learning (see § 5.5.2) can also be applied to sequence labeling. By factoring matrices of
 3840 co-occurrence counts of word bigrams and trigrams (Song et al., 2010; Hsu et al., 2012), it
 3841 is possible to obtain globally optimal estimates of the transition and emission parameters,
 3842 under mild assumptions.

3843 **7.7.3 Semiring Notation and the Generalized Viterbi Algorithm**

The Viterbi and Forward recurrences can each be performed over probabilities or log probabilities, yielding a total of four closely related recurrences. These four recurrence scan in fact be expressed as a single recurrence in a more general notation, known as **semiring algebra**. Let the symbol \oplus represent generalized addition, and the symbol \otimes represent generalized multiplication.¹² Given these operators, we can denote a general-

¹²In a semiring, the addition and multiplication operators must both obey associativity, and multiplication must distribute across addition; the addition operator must be commutative; there must be additive and multiplicative identities $\bar{0}$ and $\bar{1}$, such that $a \oplus \bar{0} = a$ and $a \otimes \bar{1} = a$; and there must be a multiplicative annihilator $\bar{0}$, such that $a \otimes \bar{0} = \bar{0}$.

ized Viterbi recurrence as,

$$v_m(k) = \bigoplus_{k' \in \mathcal{Y}} s_m(k, k') \otimes v_{m-1}(k'). \quad [7.109]$$

3844 Each recurrence that we have seen so far is a special case of this generalized Viterbi
 3845 recurrence:

- 3846 • In the max-product Viterbi recurrence over probabilities, the \oplus operation corre-
 3847 sponds to maximization, and the \otimes operation corresponds to multiplication.
- 3848 • In the forward recurrence over probabilities, the \oplus operation corresponds to addi-
 3849 tion, and the \otimes operation corresponds to multiplication.
- 3850 • In the max-product Viterbi recurrence over log-probabilities, the \oplus operation corre-
 3851 sponds to maximization, and the \otimes operation corresponds to addition.¹³
- 3852 • In the forward recurrence over log-probabilities, the \oplus operation corresponds to log-
 3853 addition, $a \oplus b = \log(e^a + e^b)$. The \otimes operation corresponds to addition.

3854 The mathematical abstraction offered by semiring notation can be applied to the soft-
 3855 ware implementations of these algorithms, yielding concise and modular implemen-
 3856 tations. The OPENFST library (Allauzen et al., 2007) is an example of a software package in
 3857 which the algorithms are parametrized by the choice of semiring.

3858 Exercises

- 3859 1. Consider the garden path sentence, *The old man the boat*. Given word-tag and tag-tag
 3860 features, what inequality in the weights must hold for the correct tag sequence to
 3861 outscore the garden path tag sequence for this example?
- 3862 2. Sketch out an algorithm for a variant of Viterbi that returns the top- n label se-
 3863 quences. What is the time and space complexity of this algorithm?
- 3864 3. Show how to compute the marginal probability $\Pr(y_{m-2} = k, y_m = k' \mid \mathbf{w}_{1:M})$, in
 3865 terms of the forwards and backward variables, and the potentials $s_n(y_n, y_{n-1})$.
- 3866 4. Suppose you receive a stream of text, where some of tokens have been replaced at
 3867 random with *NOISE*. For example:
 - 3868 • Source: *I try all things, I achieve what I can*
 - 3869 • Message received: *I try NOISE NOISE, I NOISE what I NOISE*

¹³This is sometimes called the **tropical semiring**, in honor of the Brazilian mathematician Imre Simon.

3870 Assume you have access to a pre-trained bigram language model, which gives prob-
3871 abilities $p(w_m \mid w_{m-1})$. These probabilities can be assumed to be non-zero for all
3872 bigrams.

- 3873 a) Show how to use the Viterbi algorithm to try to recover the source by maxi-
3874 mizing the bigram language model log-probability. Specifically, set the scores
3875 $s_m(y_m, y_{m-1})$ so that the Viterbi algorithm selects a sequence of words that
3876 maximizes the bigram language model log-probability, *while leaving the non-*
3877 *noise tokens intact*. Your solution should not modify the logic of the Viterbi
3878 algorithm, it should only set the scores $s_m(y_m, y_{m-1})$.
- 3879 b) An alternative solution is to iterate through the text from $m \in \{1, 2, \dots, M\}$,
3880 replacing each noise token with the word that maximizes $P(w_m \mid w_{m-1})$ ac-
3881 cording to the bigram language model. Given an upper bound on the expected
3882 fraction of tokens for which the two approaches will disagree.
- 3883 5. Consider an RNN tagging model with a tanh activation function on the hidden
3884 layer, and a hinge loss on the output. (The problem also works for the margin loss
3885 and negative log-likelihood.) Suppose you initialize all parameters to zero: this
3886 includes the word embeddings that make up \mathbf{x} , the transition matrix Θ , the out-
3887 put weights β , and the initial hidden state \mathbf{h}_0 . Prove that for any data and for any
3888 gradient-based learning algorithm, all parameters will be stuck at zero.
3889 Extra credit: would a sigmoid activation function avoid this problem?

3890 Chapter 8

3891 Applications of sequence labeling

3892 Sequence labeling has applications throughout natural language processing. This chap-
3893 ter focuses on part-of-speech tagging, morpho-syntactic attribute tagging, named entity
3894 recognition, and tokenization. It also touches briefly on two applications to interactive
3895 settings: dialogue act recognition and the detection of code-switching points between
3896 languages.

3897 8.1 Part-of-speech tagging

3898 The **syntax** of a language is the set of principles under which sequences of words are
3899 judged to be grammatically acceptable by fluent speakers. One of the most basic syntactic
3900 concepts is the **part-of-speech** (POS), which refers to the syntactic role of each word in a
3901 sentence. This concept was used informally in the previous chapter, and you may have
3902 some intuitions from your own study of English. For example, in the sentence *We like*
3903 *vegetarian sandwiches*, you may already know that *we* and *sandwiches* are nouns, *like* is a
3904 verb, and *vegetarian* is an adjective. These labels depend on the context in which the word
3905 appears: in *she eats like a vegetarian*, the word *like* is a preposition, and the word *vegetarian*
3906 is a noun.

3907 Parts-of-speech can help to disentangle or explain various linguistic problems. Recall
3908 Chomsky's proposed distinction in chapter 6:

- 3909 (8.1) Colorless green ideas sleep furiously.
- 3910 (8.2) *Ideas colorless furiously green sleep.

3911 One difference between these two examples is that the first contains part-of-speech transitions
3912 that are typical in English: adjective to adjective, adjective to noun, noun to verb, and verb
3913 to adverb. The second example contains transitions that are unusual: noun to adjective
3914 and adjective to verb. The ambiguity in a headline like,

3915 (8.3) Teacher Strikes Idle Children

3916 can also be explained in terms of parts of speech: in the interpretation that was likely
 3917 intended, *strikes* is a noun and *idle* is a verb; in the alternative explanation, *strikes* is a verb
 3918 and *idle* is an adjective.

3919 Part-of-speech tagging is often taken as a early step in a natural language processing
 3920 pipeline. Indeed, parts-of-speech provide features that can be useful for many of the
 3921 tasks that we will encounter later, such as parsing (chapter 10), coreference resolution
 3922 (chapter 15), and relation extraction (chapter 17).

3923 **8.1.1 Parts-of-Speech**

3924 The **Universal Dependencies** project (UD) is an effort to create syntactically-annotated
 3925 corpora across many languages, using a single annotation standard (Nivre et al., 2016). As
 3926 part of this effort, they have designed a part-of-speech **tagset**, which is meant to capture
 3927 word classes across as many languages as possible.¹ This section describes that inventory,
 3928 giving rough definitions for each of tags, along with supporting examples.

3929 Part-of-speech tags are **morphosyntactic**, rather than **semantic**, categories. This means
 3930 that they describe words in terms of how they pattern together and how they are inter-
 3931 nally constructed (e.g., what suffixes and prefixes they include). For example, you may
 3932 think of a noun as referring to objects or concepts, and verbs as referring to actions or
 3933 events. But events can also be nouns:

3934 (8.4) ... the **howling** of the **shrieking** storm.

3935 Here *howling* and *shrieking* are events, but grammatically they act as a noun and adjective
 3936 respectively.

3937 **8.1.1.1 The Universal Dependency part-of-speech tagset**

3938 The UD tagset is broken up into three groups: open class tags, closed class tags, and
 3939 “others.”

3940 **Open class tags** Nearly all languages contain nouns, verbs, adjectives, and adverbs.²
 3941 These are all **open word classes**, because new words can easily be added to them. The
 3942 UD tagset includes two other tags that are open classes: proper nouns and interjections.

3943 • **Nouns** (UD tag: NOUN) tend to describe entities and concepts, e.g.,

¹The UD tagset builds on earlier work from Petrov et al. (2012), in which a set of twelve universal tags was identified by creating mappings from tagsets for individual languages.

²One prominent exception is Korean, which some linguists argue does not have adjectives Kim (2002).

3944 (8.5) **Toes** are scarce among veteran **blubber men**.

3945 In English, nouns tend to follow determiners and adjectives, and can play the subject
3946 role in the sentence. They can be marked for the plural number by an -s suffix.

3947 • **Proper nouns** (PROPN) are tokens in names, which uniquely specify a given entity,

3948 (8.6) “**Moby Dick?**” shouted **Ahab**.

3949 • **Verbs** (VERB), according to the UD guidelines, “typically signal events and ac-
3950 tions.” But they are also defined grammatically: they “can constitute a minimal
3951 predicate in a clause, and govern the number and types of other constituents which
3952 may occur in a clause.”³

3953 (8.7) “**Moby Dick?**” shouted Ahab.

3954 (8.8) Shall we **keep chasing** this murderous fish?

3955 English verbs tend to come in between the subject and some number of direct ob-
3956 jects, depending on the verb. They can be marked for **tense** and **aspect** using suffixes
3957 such as *-ed* and *-ing*. (These suffixes are an example of **inflectional morphology**,
3958 which is discussed in more detail in § 9.1.4.)

3959 • **Adjectives** (ADJ) describe properties of entities,

3960 (8.9) Shall we keep chasing this **murderous** fish?

3961 (8.10) Toes are **scarce** among **veteran** blubber men.

3962 In the second example, *scarce* is a predicative adjective, linked to the subject by the
3963 **copula verb** *are*. This means that In contrast, *murderous* and *veteran* are attribute
3964 adjectives, modifying the noun phrase in which they are embedded.

3965 • **Adverbs** (ADV) describe properties of events, and may also modify adjectives or
3966 other adverbs:

3967 (8.11) It is not down on any map; true places **never** are.

3968 (8.12) ... **treacherously** hidden beneath the loveliest tints of azure

3969 (8.13) Not drowned **entirely**, though.

3970 • **Interjections** (INTJ) are used in exclamations, e.g.,

3971 (8.14) **Aye aye!** it was that accursed white whale that razed me.

³<http://universaldependencies.org/u/pos/VERB.html>

3972 **Closed class tags** Closed word classes rarely receive new members. They are sometimes
 3973 referred to as **function words** — as opposed to **content words** — as they have little lexical
 3974 meaning of their own, but rather, help to organize the components of the sentence.

- 3975 • **Adpositions** (ADP) describe the relationship between a complement (usually a noun
 3976 phrase) and another unit in the sentence, typically a noun or verb phrase.

- 3977 (8.15) Toes are scarce **among** veteran blubber men.
 3978 (8.16) It is not **down on** any map.
 3979 (8.17) Give not thyself **up** then.

3980 As the examples show, English generally uses prepositions, which are adpositions
 3981 that appear before their complement. (An exception is *ago*, as in, *we met three days*
 3982 *ago*). Postpositions are used in other languages, such as Japanese and Turkish.

- 3983 • **Auxiliary verbs** (AUX) are a closed class of verbs that add information such as
 3984 tense, aspect, person, and number.

- 3985 (8.18) **Shall** we keep chasing this murderous fish?
 3986 (8.19) What the white whale was to Ahab, **has been** hinted.
 3987 (8.20) Ahab **must** use tools.
 3988 (8.21) Meditation and water **are** wedded forever.
 3989 (8.22) Toes **are** scarce among veteran blubber men.

3990 The final example is a copula verb, which is also tagged as an auxiliary in the UD
 3991 corpus.

- 3992 • **Coordinating conjunctions** (CCONJ) express relationships between two words or
 3993 phrases, which play a parallel role:

- 3994 (8.23) Meditation **and** water are wedded forever.

- 3995 • **Subordinating conjunctions** (SCONJ) link two elements, making one syntactically
 3996 subordinate to the other:

- 3997 (8.24) There is wisdom **that** is woe.

- 3998 • **Pronouns** (PRON) are words that substitute for nouns or noun phrases.

- 3999 (8.25) Be **it what it will**, I'll go to **it** laughing.
 4000 (8.26) **I** try all things, **I** achieve **what I can**.

4001 The example includes the personal pronouns *I* and *it*, as well as the relative pronoun
4002 *what*. Other pronouns include *myself*, *somebody*, and *nothing*.

- 4003 • **Determiners** (DET) provide additional information about the nouns or noun phrases
4004 that they modify:

4005 (8.27) What **the** white whale was to Ahab, has been hinted.

4006 (8.28) It is not down on **any** map.

4007 (8.29) I try **all** things ...

4008 (8.30) Shall we keep chasing **this** murderous fish?

4009 Determiners include articles (*the*), possessive determiners (*their*), demonstratives
4010 (*this murderous fish*), and quantifiers (*any map*).

- 4011 • **Numerals** (NUM) are an infinite but closed class, which includes integers, fractions,
4012 and decimals, regardless of whether spelled out or written in numerical form.

4013 (8.31) How then can this **one** small heart beat.

4014 (8.32) I am going to put him down for the **three hundredth**.

- 4015 • **Particles** (PART) are a catch-all of function words that combine with other words or
4016 phrases, but do not meet the conditions of the other tags. In English, this includes
4017 the infinitival *to*, the possessive marker, and negation.

4018 (8.33) Better **to** sleep with a sober cannibal than a drunk Christian.

4019 (8.34) So man's insanity is heaven's sense

4020 (8.35) It is **not** down on any map

4021 As the second example shows, the possessive marker is not considered part of the
4022 same token as the word that it modifies, so that *man's* is split into two tokens. (Tok-
4023 enization is described in more detail in § 8.4.) A non-English example of a particle
4024 is the Japanese question marker *ka*, as in,⁴

4025 (8.36) *Sensei desu ka*

Teacher are ?

4026 Is she a teacher?

⁴In this notation, the first line is the transliterated Japanese text, the second line is a token-to-token **gloss**, and the third line is the translation.

4027 **Other** The remaining UD tags include punctuation (PUN) and symbols (SYM). Punc-
 4028 tuation is purely structural — e.g., commas, periods, colons — while symbols can carry
 4029 content of their own. Examples of symbols include dollar and percentage symbols, math-
 4030 ematical operators, emoticons, emojis, and internet addresses. A final catch-all tag is X,
 4031 which is used for words that cannot be assigned another part-of-speech category. The X
 4032 tag is also used in cases of **code switching** (between languages), described in § 8.5.

4033 **8.1.1.2 Other tagsets**

4034 Prior to the Universal Dependency treebank, part-of-speech tagging was performed us-
 4035 ing language-specific tagsets. The dominant tagset for English was designed as part of
 4036 the **Penn Treebank** (PTB), and it includes 45 tags — more than three times as many as
 4037 the UD tagset. This granularity is reflected in distinctions between singular and plural
 4038 nouns, verb tenses and aspects, possessive and non-possessive pronouns, comparative
 4039 and superlative adjectives and adverbs (e.g., *faster, fastest*), and so on. The Brown corpus
 4040 includes a tagset that is even more detailed, with 87 tags Francis (1964), including special
 4041 tags for individual auxiliary verbs such as *be, do, and have*.

4042 Different languages make different distinctions, and so the PTB and Brown tagsets are
 4043 not appropriate for a language such as Chinese, which does not mark the verb tense (Xia,
 4044 2000); nor for Spanish, which marks every combination of person and number in the
 4045 verb ending; nor for German, which marks the case of each noun phrase. Each of these
 4046 languages requires more detail than English in some areas of the tagset, and less in other
 4047 areas. The strategy of the Universal Dependencies corpus is to design a coarse-grained
 4048 tagset to be used across all languages, and then to additionally annotate language-specific
 4049 **morphosyntactic attributes**, such as number, tense, and case. The attribute tagging task
 4050 is described in more detail in § 8.2.

4051 Social media such as Twitter have been shown to require tagsets of their own (Gimpel
 4052 et al., 2011). Such corpora contain some tokens that are not equivalent to anything en-
 4053 countered in a typical written corpus: e.g., emoticons, URLs, and hashtags. Social media
 4054 also includes dialectal words like *gonna* ('going to', e.g. *We gonna be fine*) and *Ima* ('I'm
 4055 going to', e.g., *Ima tell you one more time*), which can be analyzed either as non-standard
 4056 orthography (making tokenization impossible), or as lexical items in their own right. In
 4057 either case, it is clear that existing tags like NOUN and VERB cannot handle cases like *Ima*,
 4058 which combine aspects of the noun and verb. Gimpel et al. (2011) therefore propose a new
 4059 set of tags to deal with these cases.

4060 **8.1.2 Accurate part-of-speech tagging**

4061 Part-of-speech tagging is the problem of selecting the correct tag for each word in a sen-
 4062 tence. Success is typically measured by accuracy on an annotated test set, which is simply
 4063 the fraction of tokens that were tagged correctly.

4064 8.1.2.1 Baselines

4065 A simple baseline for part-of-speech tagging is to choose the most common tag for each
4066 word. For example, in the Universal Dependencies treebank, the word *talk* appears 96
4067 times, and 85 of those times it is labeled as a VERB: therefore, this baseline will always
4068 predict VERB for this word. For words that do not appear in the training corpus, the base-
4069 line simply guesses the most common tag overall, which is NOUN. In the Penn Treebank,
4070 this simple baseline obtains accuracy above 92%. A more rigorous evaluation is the accu-
4071 racy on **out-of-vocabulary words**, which are not seen in the training data. Tagging these
4072 words correctly requires attention to the context and the word's internal structure.

4073 8.1.2.2 Contemporary approaches

4074 Conditional random fields and structured perceptron perform at or near the state-of-the-
4075 art for part-of-speech tagging in English. For example, (Collins, 2002) achieved 97.1%
4076 accuracy on the Penn Treebank, using a structured perceptron with the following base
4077 features (originally introduced by Ratnaparkhi (1996)):

- 4078 • current word, w_m
- 4079 • previous words, w_{m-1}, w_{m-2}
- 4080 • next words, w_{m+1}, w_{m+2}
- 4081 • previous tag, y_{m-1}
- 4082 • previous two tags, (y_{m-1}, y_{m-2})
- 4083 • for rare words:
 - 4084 – first k characters, up to $k = 4$
 - 4085 – last k characters, up to $k = 4$
 - 4086 – whether w_m contains a number, uppercase character, or hyphen.

4087 Similar results for the PTB data have been achieved using conditional random fields (CRFs;
4088 Toutanova et al., 2003).

4089 More recent work has demonstrated the power of neural sequence models, such as the
4090 **long short-term memory (LSTM)** (§ 7.6). Plank et al. (2016) apply a CRF and a bidirec-
4091 tional LSTM to twenty-two languages in the UD corpus, achieving an average accuracy
4092 of 94.3% for the CRF, and 96.5% with the bi-LSTM. Their neural model employs three
4093 types of embeddings: fine-tuned word embeddings, which are updated during training;
4094 pre-trained word embeddings, which are never updated, but which help to tag out-of-
4095 vocabulary words; and character-based embeddings. The character-based embeddings
4096 are computed by running an LSTM on the individual characters in each word, thereby
4097 capturing common orthographic patterns such as prefixes, suffixes, and capitalization.
4098 Extensive evaluations show that these additional embeddings are crucial to their model's
4099 success.

word	PTB tag	UD tag	UD attributes
<i>The</i>	DT	DET	DEFINITE=DEF PRONTYPE=ART
<i>German</i>	JJ	ADJ	DEGREE=POS
<i>Expressionist</i>	NN	NOUN	NUMBER=SING
<i>movement</i>	NN	NOUN	NUMBER=SING
<i>was</i>	VBD	AUX	MOOD=IND NUMBER=SING PERSON=3 TENSE=PAST VERBFORM=FIN
<i>destroyed</i>	VBN	VERB	TENSE=PAST VERBFORM=PART VOICE=PASS
<i>as</i>	IN	ADP	
<i>a</i>	DT	DET	DEFINITE=IND PRONTYPE=ART
<i>result</i>	NN	NOUN	NUMBER=SING
.	.	PUNCT	

Figure 8.1: UD and PTB part-of-speech tags, and UD morphosyntactic attributes. Example selected from the UD 1.4 English corpus.

4100 8.2 Morphosyntactic Attributes

4101 There is considerably more to say about a word than whether it is a noun or a verb: in En-
 4102 glish, verbs are distinguish by features such tense and aspect, nouns by number, adjectives
 4103 by degree, and so on. These features are language-specific: other languages distinguish
 4104 other features, such as **case** (the role of the noun with respect to the action of the sen-
 4105 tence, which is marked in languages such as Latin and German⁵) and **evidentiality** (the
 4106 source of information for the speaker’s statement, which is marked in languages such as
 4107 Turkish). In the UD corpora, these attributes are annotated as feature-value pairs for each
 4108 token.⁶

4109 An example is shown in Figure 8.1. The determiner *the* is marked with two attributes:
 4110 **PRONTYPE=ART**, which indicates that it is an **article** (as opposed to another type of deter-

⁵Case is marked in English for some personal pronouns, e.g., *She saw her, They saw them*.

⁶The annotation and tagging of morphosyntactic attributes can be traced back to earlier work on Turkish (Oflazer and Kuruöz, 1994) and Czech (Hajič and Hladká, 1998). MULTEXT-East was an early multilingual corpus to include morphosyntactic attributes (Dimitrova et al., 1998).

miner or pronominal modifier), and DEFINITE=DEF, which indicates that it is a **definite article** (referring to a specific, known entity). The verbs are each marked with several attributes. The auxiliary verb *was* is third-person, singular, past tense, finite (conjugated), and indicative (describing an event that has happened or is currently happenings); the main verb *destroyed* is in participle form (so there is no additional person and number information), past tense, and passive voice. Some, but not all, of these distinctions are reflected in the PTB tags VBD (past-tense verb) and VBN (past participle).

While there are thousands of papers on part-of-speech tagging, there is comparatively little work on automatically labeling morphosyntactic attributes. Faruqui et al. (2016) train a support vector machine classification model, using a minimal feature set that includes the word itself, its prefixes and suffixes, and type-level information listing all possible morphosyntactic attributes for each word and its neighbors. Mueller et al. (2013) use a conditional random field (CRF), in which the tag space consists of all observed combinations of morphosyntactic attributes (e.g., the tag would be DEF+ART for the word *the* in Figure 8.1). This massive tag space is managed by decomposing the feature space over individual attributes, and pruning paths through the trellis. More recent work has employed bidirectional LSTM sequence models. For example, Pinter et al. (2017) train a bidirectional LSTM sequence model. The input layer and hidden vectors in the LSTM are shared across attributes, but each attribute has its own output layer, culminating in a softmax over all attribute values, e.g. $y_t^{\text{NUMBER}} \in \{\text{SING}, \text{PLURAL}, \dots\}$. They find that character-level information is crucial, especially when the amount of labeled data is limited.

Evaluation is performed by first computing recall and precision for each attribute. These scores can then be averaged at either the type or token level to obtain micro- or macro-*F*-MEASURE. Pinter et al. (2017) evaluate on 23 languages in the UD treebank, reporting a median micro-*F*-MEASURE of 0.95. Performance is strongly correlated with the size of the labeled dataset for each language, with a few outliers: for example, Chinese is particularly difficult, because although the dataset is relatively large (10^5 tokens in the UD 1.4 corpus), only 6% of tokens have any attributes, offering few useful labeled instances.

8.3 Named Entity Recognition

A classical problem in information extraction is to recognize and extract mentions of **named entities** in text. In news documents, the core entity types are people, locations, and organizations; more recently, the task has been extended to include amounts of money, percentages, dates, and times. In item 8.37 (Figure 8.2), the named entities include: *The U.S. Army*, an organization; *Atlanta*, a location; and *May 14, 1864*, a date. Named entity recognition is also a key task in **biomedical natural language processing**, with entity types including proteins, DNA, RNA, and cell lines (e.g., Collier et al., 2000; Ohta et al., 2002). Figure 8.2 shows an example from the GENIA corpus of biomedical research ab-

- (8.37) *The U.S. Army captured Atlanta on May 14, 1864*
 B-ORG I-ORG I-ORG O B-LOC O B-DATE I-DATE I-DATE I-DATE
 (8.38) *Number of glucocorticoid receptors in lymphocytes and ...*
 O O B-PROTEIN I-PROTEIN O B-CELLTYPE O ...

Figure 8.2: BIO notation for named entity recognition. Example (8.38) is drawn from the GENIA corpus of biomedical documents (Ohta et al., 2002).

4149 stracts.

4150 A standard approach to tagging named entity spans is to use discriminative sequence
 4151 labeling methods such as conditional random fields. However, the named entity recogni-
 4152 tion (NER) task would seem to be fundamentally different from sequence labeling tasks
 4153 like part-of-speech tagging: rather than tagging each token, the goal is to recover *spans*
 4154 of tokens, such as *The United States Army*.

4155 This is accomplished by the **BIO notation**, shown in Figure 8.2. Each token at the
 4156 beginning of a name span is labeled with a B- prefix; each token within a name span is la-
 4157 beled with an I- prefix. These prefixes are followed by a tag for the entity type, e.g. B-LOC
 4158 for the beginning of a location, and I-PROTEIN for the inside of a protein name. Tokens
 4159 that are not parts of name spans are labeled as O. From this representation, the entity
 4160 name spans can be recovered unambiguously. This tagging scheme is also advantageous
 4161 for learning: tokens at the beginning of name spans may have different properties than
 4162 tokens within the name, and the learner can exploit this. This insight can be taken even
 4163 further, with special labels for the last tokens of a name span, and for unique tokens in
 4164 name spans, such as *Atlanta* in the example in Figure 8.2. This is called BILOU notation,
 4165 and it can yield improvements in supervised named entity recognition (Ratinov and Roth,
 4166 2009).

Feature-based sequence labeling Named entity recognition was one of the first applications of conditional random fields (McCallum and Li, 2003). The use of Viterbi decoding restricts the feature function $f(\mathbf{w}, \mathbf{y})$ to be a sum of local features, $\sum_m f(\mathbf{w}, y_m, y_{m-1}, m)$, so that each feature can consider only local adjacent tags. Typical features include tag transitions, word features for w_m and its neighbors, character-level features for prefixes and suffixes, and “word shape” features for capitalization and other orthographic properties. As an example, base features for the word *Army* in the example in (8.37) include:

(CURR-WORD:*Army*, PREV-WORD:*U.S.*, NEXT-WORD:*captured*, PREFIX-1:*A-*,
 PREFIX-2:*Ar-*, SUFFIX-1:*-y*, SUFFIX-2:*-my*, SHAPE:*Xxxx*)

4167 Another source of features is to use **gazetteers**: lists of known entity names. For example,
 4168 the U.S. Social Security Administration provides a list of tens of thousands of given names

- (1) 日文 章魚 怎麼 說?
 Japanese octopus how say
 How to say octopus in Japanese?
- (2) 日 文章 魚 怎麼 說?
 Japan essay fish how say

Figure 8.3: An example of tokenization ambiguity in Chinese (Sproat et al., 1996)

4169 — more than could be observed in any annotated corpus. Tokens or spans that match an
 4170 entry in a gazetteer can receive special features; this provides a way to incorporate hand-
 4171 crafted resources such as name lists in a learning-driven framework.

4172 **Neural sequence labeling for NER** Current research has emphasized neural sequence
 4173 labeling, using similar LSTM models to those employed in part-of-speech tagging (Ham-
 4174 merton, 2003; Huang et al., 2015; Lample et al., 2016). The bidirectional LSTM-CRF (Fig-
 4175 ure 7.4 in § 7.6) does particularly well on this task, due to its ability to model tag-to-tag
 4176 dependencies. However, Strubell et al. (2017) show that **convolutional neural networks**
 4177 can be equally accurate, with significant improvement in speed due to the efficiency of
 4178 implementing ConvNets on **graphics processing units (GPUs)**. The key innovation in
 4179 this work was the use of **dilated convolution**, which is described in more detail in § 3.4.

4180 8.4 Tokenization

4181 A basic problem for text analysis, first discussed in § 4.3.1, is to break the text into a se-
 4182 quence of discrete tokens. For alphabetic languages such as English, deterministic scripts
 4183 suffice to achieve accurate tokenization. However, in logographic writing systems such
 4184 as Chinese script, words are typically composed of a small number of characters, with-
 4185 out intervening whitespace. The tokenization must be determined by the reader, with
 4186 the potential for occasional ambiguity, as shown in Figure 8.3. One approach is to match
 4187 character sequences against a known dictionary (e.g., Sproat et al., 1996), using additional
 4188 statistical information about word frequency. However, no dictionary is completely com-
 4189 prehensive, and dictionary-based approaches can struggle with such out-of-vocabulary
 4190 words.

4191 Chinese tokenization has therefore been approached as a supervised sequence label-
 4192 ing problem. Xue et al. (2003) train a logistic regression classifier to make independent
 4193 segmentation decisions while moving a sliding window across the document. A set of
 4194 rules is then used to convert these individual classification decisions into an overall tok-
 4195 enization of the input. However, these individual decisions may be globally suboptimal,
 4196 motivating a structure prediction approach. Peng et al. (2004) train a conditional random

4197 field to predict labels of START or NONSTART on each character. More recent work has
 4198 employed neural network architectures. For example, Chen et al. (2015) use an LSTM-
 4199 CRF architecture, as described in § 7.6: they construct a trellis, in which each tag is scored
 4200 according to the hidden state of an LSTM, and tag-tag transitions are scored according
 4201 to learned transition weights. The best-scoring segmentation is then computed by the
 4202 Viterbi algorithm.

4203 8.5 Code switching

4204 Multilingual speakers and writers do not restrict themselves to a single language. **Code**
4205 **switching** is the phenomenon of switching between languages in speech and text (Auer,
4206 2013; Poplack, 1980). Written code switching has become more common in online social
4207 media, as in the following extract from Justin Trudeau's website:⁷

- 4208 (8.39) *Although everything written on this site est disponible en anglais
is available in English
and in French, my personal videos seront bilingues
will be bilingual*

4210 Accurately analyzing such texts requires first determining which languages are being
4211 used. Furthermore, quantitative analysis of code switching can provide insights on the
4212 languages themselves and their relative social positions.

Code switching can be viewed as a sequence labeling problem, where the goal is to label each token as a candidate switch point. In the example above, the words *est*, *and*, and *seront* would be labeled as switch points. Solorio and Liu (2008) detect English-Spanish switch points using a supervised classifier, with features that include the word, its part-of-speech in each language (according to a supervised part-of-speech tagger), and the probabilities of the word and part-of-speech in each language. Nguyen and Dogruöz (2013) apply a conditional random field to the problem of detecting code switching between Turkish and Dutch.

Code switching is a special case of the more general problem of word level language identification, which Barman et al. (2014) address in the context of trilingual code switching between Bengali, English, and Hindi. They further observe an even more challenging phenomenon: intra-word code switching, such as the use of English suffixes with Bengali roots. They therefore mark each token as either (1) belonging to one of the three languages; (2) a mix of multiple languages; (3) “universal” (e.g., symbols, numbers, emoticons); or (4) undefined.

⁷As quoted in <http://blogues.lapresse.ca/lagace/2008/09/08/justin-trudeau-really-parfait-bilingue/>, accessed August 21, 2017.

Speaker	Dialogue Act	Utterance
A	YES-NO-QUESTION	<i>So do you go college right now?</i>
A	ABANDONED	<i>Are yo-</i>
B	YES-ANSWER	<i>Yeah,</i>
B	STATEMENT	<i>It's my last year [laughter].</i>
A	DECLARATIVE-QUESTION	<i>You're a, so you're a senior now.</i>
B	YES-ANSWER	<i>Yeah,</i>
B	STATEMENT	<i>I'm working on my projects trying to graduate [laughter]</i>
A	APPRECIATION	<i>Oh, good for you.</i>
B	BACKCHANNEL	<i>Yeah.</i>

Figure 8.4: An example of dialogue act labeling (Stolcke et al., 2000)

4228 8.6 Dialogue acts

4229 The sequence labeling problems that we have discussed so far have been over sequences
 4230 of word tokens or characters (in the case of tokenization). However, sequence labeling
 4231 can also be performed over higher-level units, such as **utterances**. **Dialogue acts** are la-
 4232 bels over utterances in a dialogue, corresponding roughly to the speaker’s intention —
 4233 the utterance’s **illocutionary force** (Austin, 1962). For example, an utterance may state a
 4234 proposition (*it is not down on any map*), pose a question (*shall we keep chasing this murderous*
 4235 *fish?*), or provide a response (*aye aye!*). Stolcke et al. (2000) describe how a set of 42 dia-
 4236 logue acts were annotated for the 1,155 conversations in the Switchboard corpus (Godfrey
 4237 et al., 1992).⁸

4238 An example is shown in Figure 8.4. The annotation is performed over UTTERANCES,
 4239 with the possibility of multiple utterances per **conversational turn** (in cases such as inter-
 4240 ruptions, an utterance may split over multiple turns). Some utterances are clauses (e.g., *So*
 4241 *do you go to college right now?*), while others are single words (e.g., *yeah*). Stolcke et al. (2000)
 4242 report that hidden Markov models (HMMs) achieve 96% accuracy on supervised utter-
 4243 ance segmentation. The labels themselves reflect the conversational goals of the speaker:
 4244 the utterance *yeah* functions as an answer in response to the question *you’re a senior now*,
 4245 but in the final line of the excerpt, it is a **backchannel** (demonstrating comprehension).

4246 For task of dialogue act labeling, Stolcke et al. (2000) apply a hidden Markov model.
 4247 The probability $p(w_m | y_m)$ must generate the entire sequence of words in the utterance,
 4248 and it is modeled as a trigram language model (§ 6.1). Stolcke et al. (2000) also account
 4249 for acoustic features, which capture the **prosody** of each utterance — for example, tonal
 4250 and rhythmic properties of speech, which can be used to distinguish dialogue acts such

⁸Dialogue act modeling is not restricted to speech; it is relevant in any interactive conversation. For example, Jeong et al. (2009) annotate a more limited set of **speech acts** in a corpus of emails and online forums.

4251 as questions and answers. These features are handled with an additional emission distri-
4252 bution, $p(a_m | y_m)$, which is modeled with a probabilistic decision tree (Murphy, 2012).
4253 While acoustic features yield small improvements overall, they play an important role in
4254 distinguish questions from statements, and agreements from backchannels.

4255 Recurrent neural architectures for dialogue act labeling have been proposed by Kalch-
4256 brenner and Blunsom (2013) and Ji et al. (2016), with strong empirical results. Both models
4257 are recurrent at the utterance level, so that each complete utterance updates a hidden state.
4258 The recurrent-convolutional network of Kalchbrenner and Blunsom (2013) uses convolu-
4259 tion to obtain a representation of each individual utterance, while Ji et al. (2016) use a
4260 second level of recurrence, over individual words. This enables their method to also func-
4261 tion as a language model, giving probabilities over sequences of words in a document.

4262 Exercises

- 4263 1. [todo: exercises tk]

4264

Chapter 9

4265

Formal language theory

4266 We have now seen methods for learning to label individual words, vectors of word counts,
4267 and sequences of words; we will soon proceed to more complex structural transfor-
4268 mations. Most of these techniques could apply to counts or sequences from any discrete vo-
4269 cabulary; there is nothing fundamentally linguistic about, say, a hidden Markov model.
4270 This raises a basic question that this text has not yet considered: what is a language?

4271 This chapter will take the perspective of **formal language theory**, in which a language
4272 is defined as a set of **strings**, each of which is a sequence of elements from a finite alphabet.
4273 For interesting languages, there are an infinite number of strings that are in the language,
4274 and an infinite number of strings that are not. For example:

- 4275 • the set of all even-length sequences from the alphabet $\{a, b\}$, e.g., $\{\emptyset, aa, ab, ba, bb, aaaa, aaab, \dots\}$;
4276 • the set of all sequences from the alphabet $\{a, b\}$ that contain *aaa* as a substring, e.g.,
4277 $\{aaa, aaaa, baaa, aaab, \dots\}$;
4278 • the set of all sequences of English words (drawn from a finite dictionary) that con-
4279 tain at least one verb (a finite subset of the dictionary);
4280 • the `python` programming language.

4281 Formal language theory defines classes of languages and their computational prop-
4282 erties. Of particular interest is the computational complexity of solving the **membership**
4283 **problem** — determining whether a string is in a language. The chapter will focus on
4284 three classes of formal languages: regular, context-free, and “mildly” context-sensitive
4285 languages.

4286 A key insight of 20th century linguistics is that formal language theory can be usefully
4287 applied to natural languages such as English, by designing formal languages that cap-
4288 ture as many properties of the natural language as possible. For many such formalisms, a
4289 useful linguistic analysis comes as a byproduct of solving the membership problem. The

4290 membership problem can be generalized to the problems of *scoring* strings for their ac-
 4291 ceptability (as in language modeling), and of **transducing** one string into another (as in
 4292 translation).

4293 9.1 Regular languages

4294 Sooner or later, most computer scientists will write a **regular expression**. If you have,
 4295 then you have defined a **regular language**, which is any language that can be defined by
 4296 a regular expression. Formally, a regular expression can include the following elements:

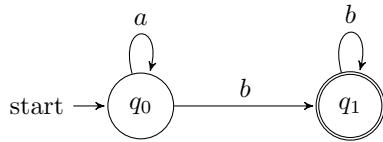
- 4297 • A **literal character** drawn from some finite alphabet Σ .
- 4298 • The **empty string** ϵ .
- 4299 • The concatenation of two regular expressions RS , where R and S are both regular
 4300 expressions. The resulting expression accepts any string that can be decomposed
 4301 $x = yz$, where y is accepted by R and z is accepted by S .
- 4302 • The alternation $R \mid S$, where R and S are both regular expressions. The resulting
 4303 expression accepts a string x if it is accepted by R or it is accepted by S .
- 4304 • The **Kleene star** R^* , which accepts any string x that can be decomposed into a se-
 4305 quence of strings which are all accepted by R .
- 4306 • Parenthesization ((R)), which is used to limit the scope of the concatenation, alterna-
 4307 tion, and Kleene star operators.

4308 Here are some example regular expressions:

- 4309 • The set of all even length strings on the alphabet $\{a, b\}$: $((aa)|(ab)|(ba)|(bb))^*$
- 4310 • The set of all sequences of the alphabet $\{a, b\}$ that contain aaa as a substring: $(a|b)^*aaa(a|b)^*$
- 4311 • The set of all sequences of English words that contain at least one verb: W^*VW^* ,
 4312 where W is an alternation between all words in the dictionary, and V is an alterna-
 4313 tion between all verbs ($V \subseteq W$).

4314 This list does not include a regular expression for the Python programming language,
 4315 because this language is not regular — there is no regular expression that can capture its
 4316 syntax. We will discuss why towards the end of this section.

4317 Regular languages are **closed** under union, intersection, and concatenation. This means,
 4318 for example, that if two languages L_1 and L_2 are regular, then so are the languages $L_1 \cup L_2$,
 4319 $L_1 \cap L_2$, and the language of strings that can be decomposed as $s = tu$, with $s \in L_1$ and
 4320 $t \in L_2$. Regular languages are also closed under negation: if L is regular, then so is the
 4321 language $\bar{L} = \{s \notin L\}$.

Figure 9.1: State diagram for the finite state acceptor M_1 .4322 **9.1.1 Finite state acceptors**

4323 A regular expression defines a regular language, but does not give an algorithm for de-
 4324 termining whether a string is in the language that it defines. **Finite state automata** are
 4325 theoretical models of computation on regular languages, which involve transitions be-
 4326 tween a finite number of states. The most basic type of finite state automaton is the **finite**
 4327 **state acceptor (FSA)**, which describes the computation involved in testing if a string is
 4328 a member of a language. Formally, a finite state acceptor is a tuple $M = (Q, \Sigma, q_0, F, \delta)$,
 4329 consisting of:

- 4330 • a finite alphabet Σ of input symbols;
- 4331 • a finite set of states $Q = \{q_0, q_1, \dots, q_n\}$;
- 4332 • a start state $q_0 \in Q$;
- 4333 • a set of final states $F \subseteq Q$;
- 4334 • a transition function $\delta : Q \times (\Sigma \cup \{\epsilon\}) \rightarrow 2^Q$. The transition function maps from a
 4335 state and an input symbol (or empty string ϵ) to a *set* of possible resulting states.

4336 A **path** in M is a sequence of transitions, $\pi = t_1, t_2, \dots, t_N$, where each t_i traverses an
 4337 arc in the transition function δ . The finite state acceptor M accepts a string ω if there is
 4338 a **accepting path**, in which the initial transition t_1 begins at the start state q_0 , the final
 4339 transition t_N terminates in a final state in Q , and the entire input ω is consumed.

4340 **9.1.1.1 Example**

Consider the following FSA, M_1 .

$$\Sigma = \{a, b\} \quad [9.1]$$

$$Q = \{q_0, q_1\} \quad [9.2]$$

$$F = \{q_1\} \quad [9.3]$$

$$\delta = \{(q_0, a) \rightarrow q_0, (q_0, b) \rightarrow q_1, (q_1, b) \rightarrow q_1\}. \quad [9.4]$$

4341 This FSA defines a language over an alphabet of two symbols, a and b . The transition
 4342 function δ is written as a set of arcs: $(q_0, a) \rightarrow q_0$ says that if the machine is in state

4343 q_0 and reads symbol a , it stays in q_0 . Figure 9.1 provides a graphical representation of
 4344 M_1 . Because each pair of initial state and symbol has at most one resulting state, M_1 is
 4345 **deterministic**: each string ω induces at most one accepting path. Note that there are no
 4346 transitions for the symbol a in state q_1 ; if a is encountered in q_1 , then the acceptor is stuck,
 4347 and the input string is rejected.

4348 What strings does M_1 accept? The start state is q_0 , and we have to get to q_1 , since this
 4349 is the only final state. Any number of a symbols can be consumed in q_0 , but a b symbol is
 4350 required to transition to q_1 . Once there, any number of b symbols can be consumed, but
 4351 an a symbol cannot. So the regular expression corresponding to the language defined by
 4352 M_1 is a^*bb^* .

4353 9.1.1.2 Computational properties of finite state acceptors

4354 The key computational question for finite state acceptors is: how fast can we determine
 4355 whether a string is accepted? For deterministic FSAs, this computation can be performed
 4356 by Dijkstra's algorithm, with time complexity $\mathcal{O}(V \log V + E)$, where V is the number of
 4357 vertices in the FSA, and E is the number of edges (Cormen et al., 2009). Non-deterministic
 4358 FSAs (NFSAs) can include multiple transitions from a given symbol and state. Any NSFA
 4359 can be converted into a deterministic FSA, but the resulting automaton may have a num-
 4360 ber of states that is exponential in the number of size of the original NFSFA (Mohri et al.,
 4361 2002).

4362 9.1.2 Morphology as a regular language

4363 Many words have internal structure, such as prefixes and suffixes that shape their mean-
 4364 ing. The study of word-internal structure is the domain of **morphology**, of which there
 4365 are two main types:

- 4366 • **Derivational morphology** describes the use of affixes to convert a word from one
 4367 grammatical category to another (e.g., from the noun *grace* to the adjective *graceful*),
 4368 or to change the meaning of the word (e.g., from *grace* to *disgrace*).
- 4369 • **Inflectional morphology** describes the addition of details such as gender, number,
 4370 person, and tense (e.g., the *-ed* suffix for past tense in English).

4371 Morphology is a rich topic in linguistics, deserving of a course in its own right.¹ The
 4372 focus here will be on the use of finite state automata for morphological analysis. The

¹A good starting point would be a chapter from a linguistics textbook (e.g., Akmajian et al., 2010; Bender, 2013). A key simplification in this chapter is the focus on affixes at the sole method of derivation and inflection. English makes use of affixes, but also incorporates **apophony**, such as the inflection of *foot* to *feet*. Semitic languages like Arabic and Hebrew feature a template-based system of morphology, in which roots are triples of consonants (e.g., *ktb*), and words are created by adding vowels: *kataba* (Arabic: he wrote), *kutub* (books), *maktab* (desk). For more detail on morphology, see texts from Haspelmath and Sims (2013) and Lieber (2015).

4373 current section deals with derivational morphology; inflectional morphology is discussed
4374 in § 9.1.4.3.

4375 Suppose that we want to write a program that accepts only those words that are con-
4376 structed in accordance with the rules of English derivational morphology:

- 4377 (9.1) grace, graceful, gracefully, *gracelyful
4378 (9.2) disgrace, *ungrace, disgraceful, disgracefully
4379 (9.3) allure, *allureful, alluring, alluringly
4380 (9.4) fairness, unfair, *disfair, fairly

4381 (Recall that the asterisk indicates that a linguistic example is judged unacceptable by flu-
4382 ent speakers of a language.) These examples cover only a tiny corner of English deriva-
4383 tional morphology, but a number of things stand out. The suffix *-ful* converts the nouns
4384 *grace* and *disgrace* into adjectives, and the suffix *-ly* converts adjectives into adverbs. These
4385 suffixes must be applied in the correct order, as shown by the unacceptability of **grace-*
4386 *lyful*. The *-ful* suffix works for only some words, as shown by the use of *alluring* as the
4387 adjectival form of *allure*. Other changes are made with prefixes, such as the derivation
4388 of *disgrace* from *grace*, which roughly corresponds to a negation; however, *fair* is negated
4389 with the *un-* prefix instead. Finally, while the first three examples suggest that the direc-
4390 tion of derivation is noun → adjective → adverb, the example of *fair* suggests that the
4391 adjective can also be the base form, with the *-ness* suffix performing the conversion to a
4392 noun.

4393 Can we build a computer program that accepts only well-formed English words, and
4394 rejects all others? This might at first seem trivial to solve with a brute-force attack: simply
4395 make a dictionary of all valid English words. But such an approach fails to account for
4396 morphological **productivity** — the applicability of existing morphological rules to new
4397 words and names, such as *Trump* to *Trumpy* and *Trumpkin*, and *Clinton* to *Clintonian* and
4398 *Clintonite*. We need an approach that represents morphological rules explicitly, and for
4399 this we will try a finite state acceptor.

4400 The dictionary approach can be implemented as a finite state acceptor, with the vo-
4401 cabulary Σ equal to the vocabulary of English, and a transition from the start state to the
4402 accepting state for each word. But this would of course fail to generalize beyond the origi-
4403 nal vocabulary, and would not capture anything about the **morphotactic** rules that govern
4404 derivations from new words. The first step towards a more general approach is shown in
4405 Figure 9.2, which is the state diagram for a finite state acceptor in which the vocabulary
4406 consists of **morphemes**, which include **stems** (e.g., *grace*, *allure*) and **affixes** (e.g., *dis-*, *-ing*,
4407 *-ly*). This finite state acceptor consists of a set of paths leading away from the start state,
4408 with derivational affixes added along the path. Except for q_{neg} , the states on these paths
4409 are all final, so the FSA will accept *disgrace*, *disgraceful*, and *disgracefully*, but not *dis-*.

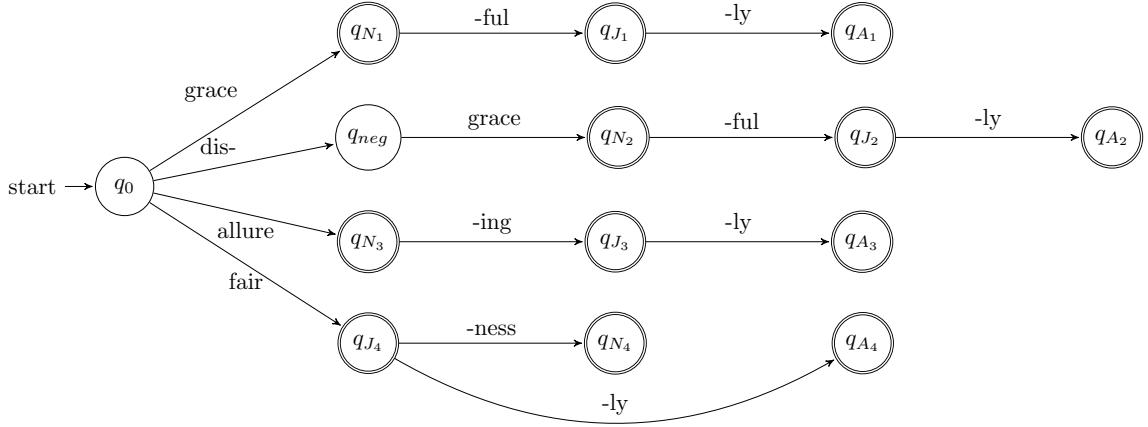


Figure 9.2: A finite state acceptor for a fragment of English derivational morphology. Each path represents possible derivations from a single root form.

4410 This FSA can be **minimized** to the form shown in Figure 9.3, which makes the generality of the finite state approach more apparent. For example, the transition from q_0 to
 4411 q_{J_2} can be made to accept not only *fair* but any single-morpheme (**monomorphemic**) adjective that takes *-ness* and *-ly* as suffixes. In this way, the finite state acceptor can easily
 4412 be extended: as new word stems are added to the vocabulary, their derived forms will be
 4413 accepted automatically. Of course, this FSA would still need to be extended considerably
 4414 to cover even this small fragment of English morphology. As shown by cases like *music*
 4415 → *musical*, *athlete* → *athletic*, English includes several classes of nouns, each with its own
 4416 rules for derivation.

4419 The FSAs shown in Figure 9.2 and 9.3 accept *allureing*, not *alluring*. This reflects a distinction between morphology — the question of which morphemes to use, and in what order — and **orthography** — the question of how the morphemes are rendered in written language. Just as orthography requires dropping the *e* preceding the *-ing* suffix, **phonology** imposes a related set of constraints on how words are rendered in speech. As we will see soon, these issues are handled through **finite state transducers**, which are finite state automata that take inputs and produce outputs.

4426 9.1.3 Weighted finite state acceptors

4427 According to the FSA treatment of morphology, every word is either in or out of the language, with no wiggle room. Perhaps you agree that *musicky* and *fishful* are not valid
 4428 English words; but if forced to choose, you probably find *a fishful stew* or *a musicky trib-
 4429 ute* preferable to *behaving disgracelyful*. Rather than asking whether a word is acceptable,
 4430 we might like to ask how acceptable it is. Aronoff (1976, page 36) puts it another way:




Figure 9.3: Minimization of the finite state acceptor shown in Figure 9.2.

4432 “Though many things are possible in morphology, some are more possible than others.”
 4433 But finite state acceptors give no way to express preferences among technically valid
 4434 choices.

4435 **Weighted finite state acceptors (WFSAs)** are generalizations of FSAs, in which each
 4436 accepting path is assigned a score, computed from the transitions, the initial state, and the
 4437 final state. Formally, a weighted finite state acceptor $M = (Q, \Sigma, \lambda, \rho, \delta)$ consists of:

- 4438 • a finite set of states $Q = \{q_0, q_1, \dots, q_n\}$;
- 4439 • a finite alphabet Σ of input symbols;
- 4440 • an initial weight function, $\lambda : Q \mapsto \mathbb{R}$;
- 4441 • a final weight function $\rho : Q \mapsto \mathbb{R}$;
- 4442 • a transition function $\delta : Q \times \Sigma \times Q \mapsto \mathbb{R}$.

4443 WFSAs depart from the FSA formalism in three ways: every state can be an initial
 4444 state, with score $\lambda(q)$; every state can be an accepting state, with score $\rho(q)$; transitions are
 4445 possible between any pair of states on any input, with a score $\delta(q_i, \omega, q_j)$. Nonetheless,
 4446 FSAs can be viewed as a special case: for any FSA M we can build an equivalent WFA
 4447 by setting $\lambda(q) = \infty$ for all $q \neq q_0$, $\rho(q) = \infty$ for all $q \notin F$, and $\delta(q_i, \omega, q_j) = \infty$ for all
 4448 transitions $\{(q_1, \omega) \rightarrow q_2\}$ that are not permitted by the transition function of M .

4449 The total score for any path $\pi = t_1, t_2, \dots, t_N$ is equal to the sum of these scores,

$$d(\pi) = \lambda(\text{from-state}(t_1)) + \sum_n^N \delta(t_n) + \rho(\text{to-state}(t_N)). \quad [9.5]$$

4450 A **shortest-path algorithm** is used to find the minimum-cost path through a WFA for
 4451 string ω , with time complexity $\mathcal{O}(E + V \log V)$, where E is the number of edges and V is
 4452 the number of vertices (Cormen et al., 2009).²

²Shortest-path algorithms find the path with the minimum cost. In many cases, the path weights are log

4453 **9.1.3.1 N-gram language models as WFSAs**

4454 In **n-gram language models** (see § 6.1), the probability of a sequence of tokens w_1, w_2, \dots, w_M
 4455 is modeled as,

$$p(w_1, \dots, w_M) \approx \prod_{m=1}^M p_n(w_m | w_{m-1}, \dots, w_{m-n+1}). \quad [9.6]$$

The log probability under an n -gram language model can be modeled in a WFSA. First consider a unigram language model. We need only a single state q_0 , with transition scores $\delta(q_0, \omega, q_0) = \log p_1(\omega)$. The initial and final scores can be set to zero. Then the path score for w_1, w_2, \dots, w_M is equal to,

$$0 + \sum_m^M \delta(q_0, w_m, q_0) + 0 = \sum_m^M \log p_1(w_m). \quad [9.7]$$

For an n -gram language model with $n > 1$, we need probabilities that condition on the past history. For example, in a bigram language model, the transition weights must represent $\log p_2(w_m | w_{m-1})$. The transition scoring function must somehow “remember” the previous word or words. This can be done by adding more states: to model the bigram probability $p_2(w_m | w_{m-1})$, we need a state for every possible w_{m-1} — a total of V states. The construction indexes each state q_i by a context event $w_{m-1} = i$. The weights are then assigned as follows:

$$\begin{aligned} \delta(q_i, \omega, q_j) &= \begin{cases} \log \Pr(w_m = j | w_{m-1} = i), & \omega = j \\ -\infty, & \omega \neq j \end{cases} \\ \lambda(q_i) &= \log \Pr(w_1 = i | w_0 = \square) \\ \rho(q_i) &= \log \Pr(w_{M+1} = \blacksquare | w_M = i). \end{aligned}$$

4456 The transition function is designed to ensure that the context is recorded accurately:
 4457 we can move to state j on input ω only if $\omega = j$; otherwise, transitioning to state j is
 4458 forbidden by the weight of $-\infty$. The initial weight function $\lambda(q_i)$ is the log probability of
 4459 receiving i as the first token, and the final weight function $\rho(q_i)$ is the log probability of
 4460 receiving an “end-of-string” token after observing $w_M = i$.

4461 **9.1.3.2 *Semiring weighted finite state acceptors**

4462 The n -gram language model WFSA is deterministic: each input has exactly one accepting
 4463 path, for which the WFSA computes a score. In non-deterministic WFSAs, a given input

probabilities, so we want the path with the maximum score, which can be accomplished by making each local score into a *negative* log-probability. The remainder of this section will refer to **best-path algorithms**, which are assumed to “do the right thing.”

4464 may have multiple accepting paths. In some applications, the score for the input is ag-
 4465 gregated across all such paths. Such aggregate scores can be computed by generalizing
 4466 WFSAs with **semiring notation**, first introduced in § 7.7.3.

4467 Let $d(\pi)$ represent the total score for path $\pi = t_1, t_2, \dots, t_N$, which is computed as,

$$d(\pi) = \lambda(\text{from-state}(t_1)) \otimes \delta(t_1) \otimes \delta(t_2) \otimes \dots \otimes \delta(t_N) \otimes \rho(\text{to-state}(t_N)). \quad [9.8]$$

4468 This is a generalization of Equation 9.5 to semiring notation, using the semiring multipli-
 4469 cation operator \otimes in place of addition.

4470 Now let $s(\omega)$ represent the total score for all paths $\Pi(\omega)$ that consume input ω ,

$$s(\omega) = \bigoplus_{\pi \in \Pi(\omega)} d(\pi). \quad [9.9]$$

4471 Here, semiring addition (\oplus) is used to combine the scores of multiple paths.

4472 The generalization to semirings covers a number of useful special cases. In the log-
 4473 probability semiring, multiplication is defined as $\log p(x) \otimes \log p(y) = \log p(x) + \log p(y)$,
 4474 and addition is defined as $\log p(x) \oplus \log p(y) = \log(p(x) + p(y))$. Thus, $s(\omega)$ represents
 4475 the log-probability of accepting input ω , marginalizing over all paths $\pi \in \Pi(\omega)$. In the
 4476 **boolean semiring**, the \otimes operator is logical conjunction, and the \oplus operator is logical
 4477 disjunction. This reduces to the special case of unweighted finite state acceptors, where
 4478 the score $s(\omega)$ is a boolean indicating whether there exists any accepting path for ω . In
 4479 the **tropical semiring**, the \oplus operator is a maximum, so the resulting score is the score of
 4480 the best-scoring path through the WFSAs. The OpenFST toolkit uses semirings and poly-
 4481 morphism to implement general algorithms for weighted finite state automata (Allauzen
 4482 et al., 2007).

4483 9.1.3.3 *Interpolated n -gram language models

4484 Recall from § 6.2.3 that an **interpolated n -gram language model** combines the probabili-
 4485 ties from multiple n -gram models. For example, an interpolated bigram language model
 4486 computes probability,

$$\hat{p}(w_m | w_{m-1}) = \lambda_1 p_1(w_m) + \lambda_2 p_2(w_m | w_{m-1}), \quad [9.10]$$

4487 with \hat{p} indicating the interpolated probability, p_2 indicating the bigram probability, and
 4488 p_1 indicating the unigram probability. We set $\lambda_2 = (1 - \lambda_1)$ so that the probabilities sum
 4489 to one.

4490 Interpolated bigram language models can be implemented using a non-deterministic
 4491 WFSAs (Knight and May, 2009). The basic idea is shown in Figure 9.4. In an interpolated
 4492 bigram language model, there is one state for each element in the vocabulary — in this

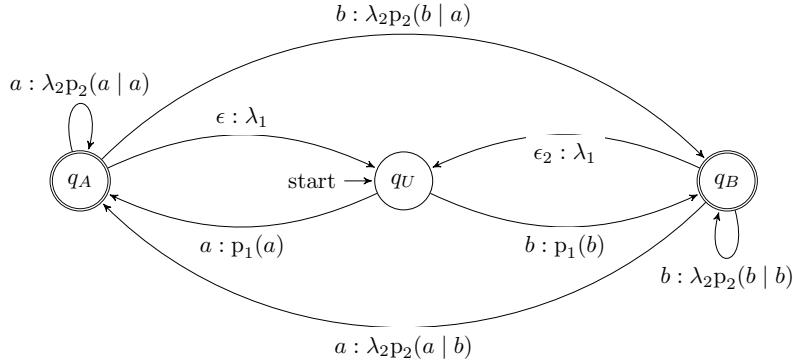


Figure 9.4: WFSA implementing an interpolated bigram/unigram language model, on the alphabet $\Sigma = \{a, b\}$. For simplicity, the WFSA is constrained to force the first token to be generated from the unigram model, and does not model the emission of the end-of-sequence token.

4493 case, the states q_A and q_B — which capture the contextual conditioning in the bigram
 4494 probabilities. To model unigram probabilities, there is an additional state q_U , which “for-
 4495 gets” the context. Transitions out of q_U involve unigram probabilities, $p_1(a)$ and $p_2(b)$;
 4496 transitions into q_U emit the empty symbol ϵ , and have probability λ_1 , reflecting the inter-
 4497 polation weight for the unigram model. The interpolation weight for the bigram model is
 4498 included in the weight of the transition $q_A \rightarrow q_B$.

4499 The epsilon transitions into q_U make this WFSA non-deterministic. Consider the score
 4500 for the sequence (a, b, b) . The initial state is q_U , so the symbol a is generated with score
 4501 $p_1(a)$ ³ Next, we can generate b from the unigram model by taking the transition $q_A \rightarrow q_B$,
 4502 with score $\lambda_2 p_2(b | a)$. Alternatively, we can take a transition back to q_U with score λ_1 ,
 4503 and then emit b from the unigram model with score $p_1(b)$. To generate the final b token,
 4504 we face the same choice: emit it directly from the self-transition to q_B , or transition to q_U
 4505 first.

The total score for the sequence (a, b, b) is the semiring sum over all accepting paths,

$$\begin{aligned}
 s(a, b, b) &= (p_1(a) \otimes \lambda_2 p_2(b | a) \otimes \lambda_2 p_2(b | b)) \\
 &\oplus (p_1(a) \otimes \lambda_1 \otimes p_1(b) \otimes \lambda_2 p_2(b | b)) \\
 &\oplus (p_1(a) \otimes \lambda_2 p_2(b | a) \otimes p_1(b) \otimes p_1(b)) \\
 &\oplus (p_1(a) \otimes \lambda_1 \otimes p_1(b) \otimes p_1(b) \otimes p_1(b)). \tag{[9.11]}
 \end{aligned}$$

4506 Each line in Equation 9.11 represents the probability of a specific path through the WFSA.
 4507 In the probability semiring, \otimes is multiplication, so that each path is the product of each

³We could model the sequence-initial bigram probability $p_2(a | \square)$, but for simplicity the WFSA does not admit this possibility, which would require another state.

4508 transition weight, which are themselves probabilities. The \oplus operator is addition, so that
 4509 the total score is the sum of the scores (probabilities) for each path. This corresponds to
 4510 the probability under the interpolated bigram language model.

4511 **9.1.4 Finite state transducers**

4512 Finite state acceptors can determine whether a string is in a regular language, and weighted
 4513 finite state acceptors can compute a score for every string over a given alphabet. **Finite**
 4514 **state transducers** (FSTs) extend the formalism further, by adding an output symbol to each
 4515 transition. Formally, a finite state transducer is a tuple $T = (Q, \Sigma, \Omega, \lambda, \rho, \delta)$, with Ω repre-
 4516 senting an output vocabulary and the transition function $\delta : Q \times (\Sigma \cup \epsilon) \times (\Omega \cup \epsilon) \times Q \rightarrow \mathbb{R}$
 4517 mapping from states, input symbols, and output symbols to states. The remaining ele-
 4518 ments (Q, Σ, λ, ρ) are identical to their definition in weighted finite state acceptors (§ 9.1.3).
 4519 Thus, each path through the FST T transduces the input string into an output.

4520 **9.1.4.1 String edit distance**

The **edit distance** between two strings s and t is a measure of how many operations are required to transform one string into another. There are several ways to compute edit distance, but one of the most popular is the **Levenshtein edit distance**, which counts the minimum number of insertions, deletions, and substitutions. This can be computed by a one-state weighted finite state transducer, in which the input and output alphabets are identical. For simplicity, consider the alphabet $\Sigma = \Omega = \{a, b\}$. The edit distance can be computed by a one-state transducer with the following transitions,

$$\delta(q, a, a, q) = \delta(q, b, b, q) = 0 \quad [9.12]$$

$$\delta(q, a, b, q) = \delta(q, b, a, q) = 1 \quad [9.13]$$

$$\delta(q, a, \epsilon, q) = \delta(q, b, \epsilon, q) = 1 \quad [9.14]$$

$$\delta(q, \epsilon, a, q) = \delta(q, \epsilon, b, q) = 1. \quad [9.15]$$

4521 The state diagram is shown in Figure 9.5.

4522 For a given string pair, there are multiple paths through the transducer: the best-
 4523 scoring path from *dessert* to *desert* involves a single deletion, for a total score of 1; the
 4524 worst-scoring path involves seven deletions and six additions, for a score of 13.

4525 **9.1.4.2 The Porter stemmer**

The Porter (1980) stemming algorithm is a “lexicon-free” algorithm for stripping suffixes from English words, using a sequence of character-level rules. Each rule can be described

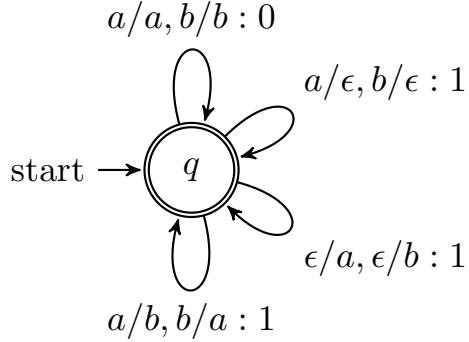


Figure 9.5: State diagram for the Levenshtein edit distance finite state transducer. The label $x/y : c$ indicates a cost of c for a transition with input x and output y .

by an unweighted finite state transducer. The first rule is:

-sses → -ss e.g., *dresses* → *dress* [9.16]

-ies → -i e.g., *parties* → *parti* [9.17]

-ss → -ss e.g., *dress* → *dress* [9.18]

-s → ε e.g., *cats* → *cat* [9.19]

4526 The final two lines appear to conflict; they are meant to be interpreted as an instruction
 4527 to remove a terminal *-s* unless it is part of an *-ss* ending. A state diagram to handle just
 4528 these final two lines is shown in Figure 9.6. Make sure you understand how this finite
 4529 state transducer handles *cats*, *steps*, *bass*, and *basses*.

4530 9.1.4.3 Inflectional morphology

4531 In **inflectional morphology**, word **lemmas** are modified to add grammatical information
 4532 such as tense, number, and case. For example, many English nouns are pluralized by the
 4533 suffix *-s*, and many verbs are converted to past tense by the suffix *-ed*. English's inflectional
 4534 morphology is considerably simpler than many of the world's languages. For example,
 4535 Romance languages (derived from Latin) feature complex systems of verb suffixes which
 4536 must agree with the person and number of the verb, as shown in Table 9.1.

4537 The task of **morphological analysis** is to read a form like *canto*, and output an analysis
 4538 like CANTAR+VERB+PRESIND+1P+SING, where +PRESIND describes the tense as present
 4539 indicative, +1P indicates the first-person, and +SING indicates the singular number. The
 4540 task of **morphological generation** is the reverse, going from CANTAR+VERB+PRESIND+1P+SING
 4541 to *canto*. Finite state transducers are an attractive solution, because they can solve both
 4542 problems with a single model (Beesley and Karttunen, 2003). As an example, Figure 9.7
 4543 shows a fragment of a finite state transducer for Spanish inflectional morphology. The

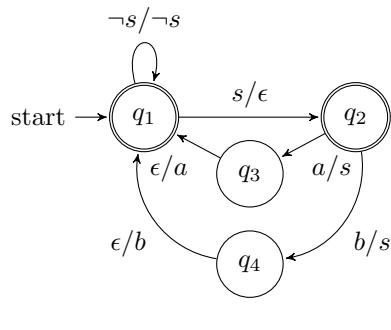


Figure 9.6: State diagram for final two lines of step 1a of the Porter stemming diagram. States q_3 and q_4 “remember” the observations a and b respectively; the ellipsis \dots represents additional states for each symbol in the input alphabet. The notation $\neg s / \neg s$ is not part of the FST formalism; it is a shorthand to indicate a set of self-transition arcs for every input/output symbol except s .

infinitive	cantar (to sing)	comer (to eat)	vivir (to live)
yo (1st singular)	canto	como	vivo
tu (2nd singular)	cantas	comes	vives
él, ella, usted (3rd singular)	canta	come	vive
nosotros (1st plural)	cantamos	comemos	vivimos
vosotros (2nd plural, informal)	cantáis	coméis	vívís
ellos, ellas (3rd plural); ustedes (2nd plural)	cantan	comen	viven

Table 9.1: Spanish verb inflections for the present indicative tense. Each row represents a person and number, and each column is a regular example from a class of verbs, as indicated by the ending of the infinitive form.

4544 input vocabulary Σ corresponds to the set of letters used in Spanish spelling, and the out-
 4545 put vocabulary Ω corresponds to these same letters, plus the vocabulary of morphological
 4546 features (e.g., +SING, +VERB). In Figure 9.7, there are two paths that take *canto* as input,
 4547 corresponding to the verb and noun meanings; the choice between these paths could be
 4548 guided by a part-of-speech tagger. By **inversion**, the inputs and outputs for each trans-
 4549 sition are switched, resulting in a finite state generator, capable of producing the correct
 4550 **surface form** for any morphological analysis.

4551 Finite state morphological analyzers and other unweighted transducers can be de-
 4552 signed by hand. The designer’s goal is to avoid **overgeneration** — accepting strings or
 4553 making transductions that are not valid in the language — as well as **undergeneration** —

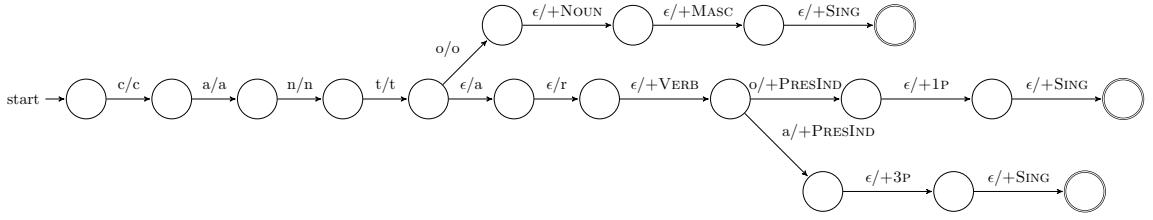


Figure 9.7: Fragment of a finite state transducer for Spanish morphology. There are two accepting paths for the input *canto*: *canto*+NOUN+MASC+SING (masculine singular noun, meaning a song), and *cantar*+VERB+PRESIND+1P+SING (I sing). There is also an accepting path for *canta*, with output *cantar*+VERB+PRESIND+3P+SING (he/she sings).

4554 failing to accept strings or transductions that are valid. For example, a pluralization trans-
 4555 ducer that does not accept *foot/feet* would undergenerate. Suppose we “fix” the transducer
 4556 to accept this example, but as a side effect, it now accepts *boot/beet*; the transducer would
 4557 then be said to overgenerate. A transducer that accepts *foot/foots* but not *foot/feet* would
 4558 both overgenerate and undergenerate.

4559 9.1.4.4 Finite state composition

4560 Designing finite state transducers to capture the full range of morphological phenomena
 4561 in any real language is a huge task. Modularization is a classic computer science approach
 4562 for this situation: decompose a large and unwieldy problem into a set of subproblems,
 4563 each of which will hopefully have a concise solution. Finite state automata can be mod-
 4564 ularized through **composition**: feeding the output of one transducer T_1 as the input to
 4565 another transducer T_2 , written $T_2 \circ T_1$. Formally, if there exists some y such that $(x, y) \in T_1$
 4566 (meaning that T_1 produces output y on input x), and $(y, z) \in T_2$, then $(x, z) \in (T_2 \circ T_1)$.
 4567 Because finite state transducers are closed under composition, there is guaranteed to be
 4568 a single finite state transducer that $T_3 = T_2 \circ T_1$, which can be constructed as a machine
 4569 with one state for each pair of states in T_1 and T_2 (Mohri et al., 2002).

4570 **Example: Morphology and orthography** In English morphology, the suffix *-ed* is added
 4571 to signal the past tense for many verbs: *cook*→*cooked*, *want*→*wanted*, etc. However, English
 4572 **orthography** dictates that this process cannot produce a spelling with consecutive e’s, so
 4573 that *bake*→*baked*, not *bakeed*. A modular solution is to build separate transducers for mor-
 4574 phology and orthography. The morphological transducer T_M transduces from *bake*+PAST
 4575 to *bake+ed*, with the + symbol indicating a segment boundary. The input alphabet of T_M
 4576 includes the lexicon of words and the set of morphological features; the output alphabet
 4577 includes the characters *a-z* and the + boundary marker. Next, an orthographic transducer
 4578 T_O is responsible for the transductions *cook+ed*→*cooked*, and *bake+ed*→*baked*. The input
 4579 alphabet of T_O must be the same as the output alphabet for T_M , and the output alphabet

4580 is simply the characters *a-z*. The composed transducer ($T_O \circ T_M$) then transduces from
 4581 *bake*+PAST to the spelling *baked*. The design of T_O is left as an exercise.

Example: Hidden Markov models Hidden Markov models (chapter 7) can be viewed as weighted finite state transducers, and they can be constructed by transduction. Recall that a hidden Markov model defines a joint probability over words and tags, $p(w, y)$, which can be computed as a path through a **trellis** structure. This trellis is itself a weighted finite state acceptor, with edges between all adjacent nodes $q_{m-1,i} \rightarrow q_{m,j}$ on input $Y_m = j$. The edge weights are log-probabilities,

$$\delta(q_{m-1,i}, Y_m = j, q_{m,j}) = \log p(w_m, Y_m = j \mid Y_{m-1} = i) \quad [9.20]$$

$$= \log p(w_m \mid Y_m = j) + \log \Pr(Y_m = j \mid Y_{m-1} = i). \quad [9.21]$$

4582 Because there is only one possible transition for each tag Y_m , this WFSA is deterministic.
 4583 The score for any tag sequence $\{y_m\}_{m=1}^M$ is the sum of these log-probabilities, correspond-
 4584 ing to the total log probability $\log p(w, y)$. Furthermore, the trellis can be constructed by
 4585 the composition of simpler FSTs.

- 4586 • First, construct a “transition” transducer to represent a bigram probability model
 4587 over tag sequences, T_T . This transducer is almost identical to the n -gram language
 4588 model acceptor in § 9.1.3.1: there is one state for each tag, and the edge weights
 4589 equal to the transition log-probabilities, $\delta(q_i, j, j, q_j) = \log \Pr(Y_m = j \mid Y_{m-1} = i)$.
 4590 Note that T_T is a transducer, with identical input and output at each arc; this makes
 4591 it possible to compose T_T with other transducers.
- 4592 • Next, construct an “emission” transducer to represent the probability of words given
 4593 tags, T_E . This transducer has only a single state, with arcs for each word/tag pair,
 4594 $\delta(q_0, i, j, q_0) = \log \Pr(W_m = j \mid Y_m = i)$. The input vocabulary is the set of all tags,
 4595 and the output vocabulary is the set of all words.
- 4596 • The composition $T_E \circ T_T$ is a finite state transducer with one state per tag, as shown
 4597 in Figure 9.8. Each state has $V \times K$ outgoing edges, representing transitions to each
 4598 of the K other states, with outputs for each of the V words in the vocabulary. The
 4599 weights for these edges are equal to,

$$\delta(q_i, Y_m = j, w_m, q_j) = \log p(w_m, Y_m = j \mid Y_{m-1} = i). \quad [9.22]$$

- 4600 • The trellis is a structure with $M \times K$ nodes, for each of the M words to be tagged and
 4601 each of the K tags in the tagset. It can be built by composition of $(T_E \circ T_T)$ against an
 4602 unweighted **chain FSA** $M_A(w)$ that is specially constructed to accept only a given
 4603 input w_1, w_2, \dots, w_M , shown in Figure 9.9. The trellis for input w is built from the
 4604 composition $M_A(w) \circ (T_E \circ T_T)$. Composing with the unweighted $M_A(w)$ does not
 4605 affect the edge weights from $(T_E \circ T_T)$, but it selects the subset of paths that generate
 4606 the word sequence w .




Figure 9.8: Finite state transducer for hidden Markov models, with a small tagset of nouns and verbs. For each pair of tags (including self-loops), there is an edge for every word in the vocabulary. For simplicity, input and output are only shown for the edges from the start state. Weights are also omitted from the diagram; for each edge from q_i to q_j , the weight is equal to $\log p(w_m, Y_m = j \mid Y_{m-1} = i)$, except for edges to the end state, which are equal to $\log \Pr(Y_m = \diamond \mid Y_{m-1} = i)$.




Figure 9.9: Chain finite state acceptor for the input *They can fish*.

4607 9.1.5 *Learning weighted finite state automata

4608 In generative models such as n -gram language models and hidden Markov models, the
 4609 edge weights correspond to log probabilities, which can be obtained from relative fre-
 4610 quency estimation. However, in other cases, we wish to learn the edge weights from in-
 4611 put/output pairs. This is difficult in non-deterministic finite state automata, because we
 4612 do not observe the specific arcs that are traversed in accepting the input, or in transducing
 4613 from input to output. The path through the automaton is a **latent variable**.

4614 Chapter 5 presented one method for learning with latent variables: expectation max-
 4615 imization (EM). This involves computing a distribution $q(\cdot)$ over the latent variable, and
 4616 iterating between updates to this distribution and updates to the parameters — in this
 4617 case, the arc weights. The **forward-backward algorithm** (§ 7.5.3.3) describes a dynamic
 4618 program for computing a distribution over arcs in the trellis structure of a hidden Markov

model, but this is a special case of the more general problem for finite state automata. Eisner (2002) describes an **expectation semiring**, which enables the expected number of transitions across each arc to be computed through a semiring shortest-path algorithm. Alternative approaches for generative models include Markov Chain Monte Carlo (Chiang et al., 2010) and spectral learning (Balle et al., 2011).

Further afield, we can take a perceptron-style approach, with each arc corresponding to a feature. The classic perceptron update would update the weights by subtracting the difference between the feature vector corresponding to the predicted path and the feature vector corresponding to the correct path. Since the path is not observed, we resort to a **hidden variable perceptron**. The model is described formally in § 12.4, but the basic idea is to compute an update from the difference between the features from the predicted path and the features for the best-scoring path that generates the correct output.

9.2 Context-free languages

Beyond the class of regular languages lie the context-free languages. An example of a language that is context-free but not finite state is the set of arithmetic expressions with balanced parentheses. Intuitively, to accept only strings in this language, an FSA would have to “count” the number of left parentheses, and make sure that they are balanced against the number of right parentheses. An arithmetic expression can be arbitrarily long, yet by definition an FSA has a finite number of states. Thus, for any FSA, there will be a string that with too many parentheses to count. More formally, the **pumping lemma** is a proof technique for showing that languages are not regular. It is typically demonstrated for the simpler case $a^n b^n$, the language of strings containing a sequence of a 's, and then an equal-length sequence of b 's.⁴

There are at least two arguments for the relevance of non-regular formal languages to linguistics. First, there are natural language phenomena that are argued to be isomorphic to $a^n b^n$. For English, the classic example is **center embedding**, shown in Figure 9.10. The initial expression *the dog* specifies a single dog. Embedding this expression into *the cat ... chased* specifies a particular cat — the one chased by the dog. This cat can then be embedded again to specify a goat, in the less felicitous but arguably grammatical expression, *the goat the cat the dog chased kissed*, which refers to the goat who was kissed by the cat which was chased by the dog. Chomsky (1957) argues that to be grammatical, a center-embedded construction must be balanced: if it contains n noun phrases (e.g., *the cat*), they must be followed by exactly $n - 1$ verbs. An FSA that could recognize such expressions would also be capable of recognizing the language $a^n b^n$. Because we can prove that no FSA exists for $a^n b^n$, no FSA can exist for center embedded constructions either. En-

⁴Details of the proof can be found in an introductory computer science theory textbook (e.g., Sipser, 2012).

			the dog	
	the cat	the dog	chased	
the goat	the cat	the dog	chased	kissed
			...	

Figure 9.10: Three levels of center embedding

4654 glish includes center embedding, and so the argument goes, English grammar as a whole
 4655 cannot be regular.⁵

4656 A more practical argument for moving beyond regular languages is modularity. Many
 4657 linguistic phenomena — especially in syntax — involve constraints that apply at long
 4658 distance. Consider the problem of determiner-noun number agreement in English: we
 4659 can say *the coffee* and *these coffees*, but not **these coffee*. By itself, this is easy enough to model
 4660 in an FSA. However, fairly complex modifying expressions can be inserted between the
 4661 determiner and the noun:

- 4662 (9.5) the burnt coffee
- 4663 (9.6) the badly-ground coffee
- 4664 (9.7) the burnt and badly-ground Italian coffee
- 4665 (9.8) these burnt and badly-ground Italian coffees
- 4666 (9.9) *these burnt and badly-ground Italian coffee

4667 Again, an FSA can be designed to accept modifying expressions such as *burnt and badly-*
 4668 *ground Italian*. Let's call this FSA F_M . To reject the final example, a finite state acceptor
 4669 must somehow "remember" that the determiner was plural when it reaches the noun *cof-*
 4670 *fee* at the end of the expression. The only way to do this is to make two identical copies
 4671 of F_M : one for singular determiners, and one for plurals. While this is possible in the
 4672 finite state framework, it is inconvenient — especially in languages where more than one
 4673 attribute of the noun is marked by the determiner. **Context-free languages** facilitate mod-
 4674 ularity across such long-range dependencies.

4675 9.2.1 Context-free grammars

4676 Context-free languages are specified by **context-free grammars (CFGs)**, which are tuples
 4677 (N, Σ, R, S) consisting of:

⁵The claim that arbitrarily deep center-embedded expressions are grammatical has drawn skepticism. Corpus evidence shows that embeddings of depth greater than two are exceedingly rare (Karlsson, 2007), and that embeddings of depth greater than three are completely unattested. If center-embedding is capped at some finite depth, then it is regular.

$$\begin{aligned}
 S &\rightarrow S \text{ OP } S \mid \text{NUM} \\
 \text{OP} &\rightarrow + \mid - \mid \times \mid \div \\
 \text{NUM} &\rightarrow \text{NUM DIGIT} \mid \text{DIGIT} \\
 \text{DIGIT} &\rightarrow 0 \mid 1 \mid 2 \mid \dots \mid 9
 \end{aligned}$$

Figure 9.11: A context-free grammar for arithmetic expressions

- 4678 • a finite set of **non-terminals** N ;
- 4679 • a finite alphabet Σ of **terminal symbols**;
- 4680 • a set of **production rules** R , each of the form $A \rightarrow \beta$, where $A \in N$ and $\beta \in (\Sigma \cup N)^*$;
- 4681 • a designated start symbol S .

4682 In the production rule $A \rightarrow \beta$, the left-hand side (LHS) A must be a non-terminal;
 4683 the right-hand side (RHS) can be a sequence of terminals or non-terminals, $\{n, \sigma\}^*, n \in$
 4684 $N, \sigma \in \Sigma$. A non-terminal can appear on the left-hand side of many production rules.
 4685 A non-terminal can appear on both the left-hand side and the right-hand side; this is a
 4686 **recursive production**, and is analogous to self-loops in finite state automata. The name
 4687 “context-free” is based on the property that the production rule depends only on the LHS,
 4688 and not on its ancestors or neighbors; this is analogous to Markov property of finite state
 4689 automata, in which the behavior at each step depends only on the current state, on not on
 4690 the path by which that state was reached.

4691 A **derivation** τ is a sequence of steps from the start symbol S to a surface string $w \in \Sigma^*$,
 4692 which is the **yield** of the derivation. A string w is in a context-free language if there is
 4693 some derivation from S yielding w . **Parsing** is the problem of finding a derivation for a
 4694 string in a grammar. Algorithms for parsing are described in chapter 10.

4695 Like regular expressions, context-free grammars define the language but not the com-
 4696 putation necessary to recognize it. The context-free analogues to finite state acceptors are
 4697 **pushdown automata**, a theoretical model of computation in which input symbols can be
 4698 pushed onto a stack with potentially infinite depth. For more details, see Sipser (2012).

4699 9.2.1.1 Example

4700 Figure 9.11 shows a context-free grammar for arithmetic expressions such as $1 + 2 \div 3 - 4$.
 4701 In this grammar, the terminal symbols include the digits $\{1, 2, \dots, 9\}$ and the op-
 4702 erators $\{+, -, \times, \div\}$. The rules include the $|$ symbol, a notational convenience that makes
 4703 it possible to specify multiple right-hand sides on a single line: the statement $A \rightarrow x | y$

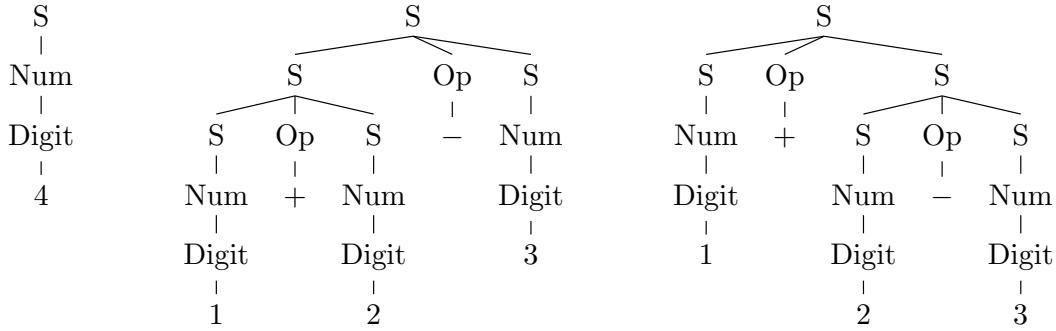


Figure 9.12: Some example derivations from the arithmetic grammar in Figure 9.11

4704 defines *two* productions, $A \rightarrow x$ and $A \rightarrow y$. This grammar is recursive: the non-termals S
4705 and NUM can produce themselves.

4706 Derivations are typically shown as trees, with production rules applied from the top
4707 to the bottom. The tree on the left in Figure 9.12 describes the derivation of a single digit,
4708 through the sequence of productions $S \rightarrow \text{NUM} \rightarrow \text{DIGIT} \rightarrow 4$ (these are all **unary produc-**
4709 **tions**, because the right-hand side contains a single element). The other two trees in
4710 Figure 9.12 show alternative derivations of the string $1 + 2 - 3$. The existence of multiple
4711 derivations for a string indicates that the grammar is **ambiguous**.

Context-free derivations can also be written out according to the pre-order tree traversal.⁶ For the two derivations of $1 + 2 - 3$ in Figure 9.12, the notation is:

$$(S (S (S (\text{Num} (Digit 1))) (\text{Op} +) (S (\text{Num} (Digit 2))))) (\text{Op} -) (S (\text{Num} (Digit 3)))) \quad [9.23]$$

$$(S (S (\text{Num} (Digit 1))) (\text{Op} +) (S (\text{Num} (Digit 2)) (\text{Op} -) (S (\text{Num} (Digit 3)))))). \quad [9.24]$$

4712 9.2.1.2 Grammar equivalence and Chomsky Normal Form

A single context-free language can be expressed by more than one context-free grammar. For example, the following two grammars both define the language $a^n b^n$ for $n > 0$.

$$\begin{aligned} S &\rightarrow aSb \mid ab \\ S &\rightarrow aSb \mid aabb \mid ab \end{aligned}$$

4713 Two grammars are **weakly equivalent** if they generate the same strings. Two grammars
4714 are **strongly equivalent** if they generate the same strings via the same derivations. The
4715 grammars above are only weakly equivalent.

⁶This is a depth-first left-to-right search that prints each node the first time it is encountered (Cormen et al., 2009, chapter 12).

In **Chomsky Normal Form (CNF)**, the right-hand side of every production includes either two non-terminals, or a single terminal symbol:

$$A \rightarrow BC$$

$$A \rightarrow a$$

- 4716 All CFGs can be converted into a CNF grammar that is weakly equivalent. To convert a
 4717 grammar into CNF, we first address productions that have more than two non-terminals
 4718 on the RHS by creating new “dummy” non-terminals. For example, if we have the pro-
 4719 duction,

$$W \rightarrow X Y Z, \quad [9.25]$$

it is replaced with two productions,

$$W \rightarrow X W \setminus X \quad [9.26]$$

$$W \setminus X \rightarrow Y Z. \quad [9.27]$$

- 4720 In these productions, $W \setminus X$ is a new dummy non-terminal. This transformation **binarizes**
 4721 the grammar, which is critical for efficient bottom-up parsing, as we will see in chapter 10.
 4722 Productions whose right-hand side contains a mix of terminal and non-terminal symbols
 4723 can be replaced in a similar fashion.

- 4724 Unary non-terminal productions $A \rightarrow B$ are replaced as follows: identify all produc-
 4725 tions $B \rightarrow \alpha$, and add $A \rightarrow \alpha$ to the grammar. For example, in the grammar described in
 4726 Figure 9.11, we would replace $\text{NUM} \rightarrow \text{DIGIT}$ with $\text{NUM} \rightarrow 1 \mid 2 \mid \dots \mid 9$. However, we
 4727 keep the production $\text{NUM} \rightarrow \text{NUM DIGIT}$, which is a valid binary production.

4728 9.2.2 Natural language syntax as a context-free language

- 4729 Context-free grammars are widely used to represent **syntax**, which is the set of rules that
 4730 determine whether an utterance is judged to be grammatical. If this representation were
 4731 perfectly faithful, then a natural language such as English could be transformed into a
 4732 formal language, consisting of exactly the (infinite) set of strings that would be judged to
 4733 be grammatical by a fluent English speaker. We could then build parsing software that
 4734 would automatically determine if a given utterance were grammatical.⁷

- 4735 Contemporary theories generally do *not* consider natural languages to be context-free
 4736 (see § 9.3), yet context-free grammars are widely used in natural language parsing. The
 4737 reason is that context-free representations strike a good balance: they cover a broad range
 4738 of syntactic phenomena, and they can be parsed efficiently. This section therefore de-
 4739 scribes how to handle a core fragment of English syntax in context-free form, following

⁷You are encouraged to move beyond this cursory treatment of syntax by consulting a textbook on linguistics (e.g., Akmajian et al., 2010; Bender, 2013).

4740 the conventions of the **Penn Treebank** (PTB; Marcus et al., 1993), a large-scale annotation
 4741 of English language syntax. The generalization to “mildly” context-sensitive languages is
 4742 discussed in § 9.3.

4743 The Penn Treebank annotation is a **phrase-structure grammar** of English. This means
 4744 that sentences are broken down into **constituents**, which are contiguous sequences of
 4745 words that function as coherent units for the purpose of linguistic analysis. Constituents
 4746 generally have a few key properties:

4747 **Movement.** Constituents can often be moved around sentences as units.

- 4748 (9.10) Abigail gave (her brother) (a fish).
 4749 (9.11) Abigail gave (a fish) to (her brother).

4750 In contrast, *gave her* and *brother a* cannot easily be moved while preserving gram-
 4751 maticality.

4752 **Substitution.** Constituents can be substituted by other phrases of the same type.

- 4753 (9.12) Max thanked (his older sister).
 4754 (9.13) Max thanked (her).

4755 In contrast, substitution is not possible for other contiguous units like *Max thanked*
 4756 and *thanked his*.

4757 **Coordination.** Coordinators like *and* and *or* can conjoin constituents.

- 4758 (9.14) (Abigail) and (her younger brother) bought a fish.
 4759 (9.15) Abigail (bought a fish) and (gave it to Max).
 4760 (9.16) Abigail (bought) and (greedily ate) a fish.

4761 Units like *brother bought* and *bought a* cannot easily be coordinated.

4762 These examples argue for units such as *her brother* and *bought a fish* to be treated as con-
 4763 stituents. Other sequences of words in these examples, such as *Abigail gave* and *brother*
a fish, cannot be moved, substituted, and coordinated in these ways. In phrase-structure
 4764 grammar, constituents are nested, so that *the senator from New Jersey* contains the con-
 4765 stituent *from New Jersey*, which in turn contains *New Jersey*. The sentence itself is the max-
 4766 imal constituent; each word is a minimal constituent, derived from a unary production
 4767 from a part-of-speech tag. Between part-of-speech tags and sentences are **phrases**. In
 4768 phrase-structure grammar, phrases have a type that is usually determined by their **head**
 4769 **word**: for example, a **noun phrase** corresponds to a noun and the group of words that

4771 modify it, such as *her younger brother*; a **verb phrase** includes the verb and its modifiers,
4772 such as *bought a fish* and *greedily ate it*.

4773 In context-free grammars, each phrase type is a non-terminal, and each constituent is
4774 the substring that the non-terminal yields. Grammar design involves choosing the right
4775 set of non-terminals. Fine-grained non-terminals make it possible to represent more fine-
4776 grained linguistic phenomena. For example, by distinguishing singular and plural noun
4777 phrases, it is possible to have a grammar of English that generates only sentences that
4778 obey subject-verb agreement. However, enforcing subject-verb agreement is considerably
4779 more complicated in languages like Spanish, where the verb must agree in both person
4780 and number with subject. In general, grammar designers must trade off between **over-**
4781 **generation** — a grammar that permits ungrammatical sentences — and **undergeneration**
4782 — a grammar that fails to generate grammatical sentences. Furthermore, if the grammar is
4783 to support manual annotation of syntactic structure, it must be simple enough to annotate
4784 efficiently.

4785 9.2.3 A phrase-structure grammar for English

4786 To better understand how phrase-structure grammar works, let's consider the specific
4787 case of the Penn Treebank grammar of English. The main phrase categories in the Penn
4788 Treebank (PTB) are based on the main part-of-speech classes: noun phrase (NP), verb
4789 phrase (VP), prepositional phrase (PP), adjectival phrase (ADJP), and adverbial phrase
4790 (ADVP). The top-level category is S, which conveniently stands in for both “sentence”
4791 and the “start” symbol. **Complement clauses** (e.g., *I take the good old fashioned ground that*
4792 *the whale is a fish*) are represented by the non-terminal SBAR. The terminal symbols in
4793 the grammar are individual words, which are generated from unary productions from
4794 part-of-speech tags (the PTB tagset is described in § 8.1).

4795 This section explores the productions from the major phrase-level categories, explaining
4796 how to generate individual tag sequences. The production rules are approached in a
4797 “theory-driven” manner: first the syntactic properties of each phrase type are described,
4798 and then some of the necessary production rules are listed. But it is important to keep
4799 in mind that the Penn Treebank was produced in a “data-driven” manner. After the set
4800 of non-terminals was specified, annotators were free to analyze each sentence in what-
4801 ever way seemed most linguistically accurate, subject to some high-level guidelines. The
4802 grammar of the Penn Treebank is simply the set of productions that were required to ana-
4803 lyze the several million words of the corpus. By design, the grammar overgenerates — it
4804 does not exclude ungrammatical sentences.

4805 **9.2.3.1 Sentences**

The most common production rule for sentences is,

$$S \rightarrow NP VP \quad [9.28]$$

which accounts for simple sentences like *Abigail ate the kimchi* — as we will see, the direct object *the kimchi* is part of the verb phrase. But there are more complex forms of sentences as well:

$$S \rightarrow ADVP NP VP \quad \text{Unfortunately } Abigail \text{ ate the kimchi.} \quad [9.29]$$

$$S \rightarrow S CC S \quad \text{Abigail ate the kimchi and Max had a burger.} \quad [9.30]$$

$$S \rightarrow VP \quad \text{Eat the kimchi.} \quad [9.31]$$

- 4806 where ADVP is an adverbial phrase (e.g., *unfortunately*, *very unfortunately*) and CC is a
 4807 coordinating conjunction (e.g., *and*, *but*).⁸

4808 **9.2.3.2 Noun phrases**

Noun phrases refer to entities, real or imaginary, physical or abstract: *Asha*, *the steamed dumpling*, *parts and labor*, *nobody*, *the whiteness of the whale*, and *the rise of revolutionary syndicalism in the early twentieth century*. Noun phrase productions include “bare” nouns, which may optionally follow determiners, as well as pronouns:

$$NP \rightarrow NN | NNS | NNP | PRP \quad [9.32]$$

$$NP \rightarrow DET NN | DET NNS | DET NNP \quad [9.33]$$

- 4809 The tags NN, NNS, and NNP refer to singular, plural, and proper nouns; PRP refers to
 4810 personal pronouns, and DET refers to determiners. The grammar also contains terminal
 4811 productions from each of these tags, e.g., $PRP \rightarrow I | you | we | \dots$.

Noun phrases may be modified by adjectival phrases (ADJP; e.g., *the small Russian dog*) and numbers (CD; e.g., *the five pastries*), each of which may optionally follow a determiner:

$$NP \rightarrow ADJP NN | ADJP NNS | DET ADJP NN | DET ADJP NNS \quad [9.34]$$

$$NP \rightarrow CD NNS | DET CD NNS | \dots \quad [9.35]$$

Some noun phrases include multiple nouns, such as *the liberation movement* and *an antelope horn*, necessitating additional productions:

$$NP \rightarrow NN NN | NN NNS | DET NN NN | \dots \quad [9.36]$$

⁸Notice that the grammar does not include the recursive production $S \rightarrow ADVP S$. It may be helpful to think about why this production would cause the grammar to overgenerate.

4812 These multiple noun constructions can be combined with adjectival phrases and cardinal
 4813 numbers, leading to a large number of additional productions.

Recursive noun phrase productions include coordination, prepositional phrase attachment, subordinate clauses, and verb phrase adjuncts:

$NP \rightarrow NP\ Cc\ NP$	<i>e.g., the red and the black</i>	[9.37]
$NP \rightarrow NP\ PP$	<i>e.g., the President of the Georgia Institute of Technology</i>	[9.38]
$NP \rightarrow NP\ SBAR$	<i>e.g., a whale which he had wounded</i>	[9.39]
$NP \rightarrow NP\ VP$	<i>e.g., a whale taken near Shetland</i>	[9.40]

4814 These recursive productions are a major source of ambiguity, because the VP and PP non-
 4815 terminals can also generate NP children. Thus, the *the President of the Georgia Institute of*
 4816 *Technology* can be derived in two ways, as can *a whale taken near Shetland in October*.

4817 But aside from these few recursive productions, the noun phrase fragment of the Penn
 4818 Treebank grammar is relatively flat, containing a large of number of productions that go
 4819 from NP directly to a sequence of parts-of-speech. If noun phrases had more internal
 4820 structure, the grammar would need fewer rules, which, as we will see, would make pars-
 4821 ing faster and machine learning easier. Vadas and Curran (2011) propose to add additional
 4822 structure in the form of a new non-terminal called a **nominal modifier** (NML), e.g.,

4823 (9.17) (NP (NN crude) (NN oil) (NNS prices)) (PTB analysis)
 4824 (NP (NML (NN crude) (NN oil)) (NNS prices)) (NML-style analysis)

4825 Another proposal is to treat the determiner as the head of a **determiner phrase** (DP;
 4826 Abney, 1987). There are linguistic arguments for and against determiner phrases (e.g.,
 4827 Van Eynde, 2006). From the perspective of context-free grammar, DPs enable more struc-
 4828 tured analyses of some constituents, e.g.,

4829 (9.18) (NP (DT the) (JJ white) (NN whale)) (PTB analysis)
 4830 (DP (DT the) (NP (JJ white) (NN whale))) (DP-style analysis).

4831 9.2.3.3 Verb phrases

Verb phrases describe actions, events, and states of being. The PTB tagset distinguishes several classes of verb inflections: base form (VB; *she likes to snack*), present-tense third-person singular (VBD; *she snacks*), present tense but not third-person singular (VBP; *they snack*), past tense (VBD; *they snacked*), present participle (VBG; *they are snacking*), and past participle (VBN; *they had snacked*).⁹ Each of these forms can constitute a verb phrase on its

⁹It bears emphasis the principles governing this tagset design are entirely English-specific: VBP is a meaningful category only because English morphology distinguishes third-person singular from all person-number combinations.

own:

$$\text{VP} \rightarrow \text{VB} \mid \text{VBZ} \mid \text{VBD} \mid \text{VBN} \mid \text{VBG} \mid \text{VBP} \quad [9.41]$$

More complex verb phrases can be formed by a number of recursive productions, including the use of coordination, modal verbs (MD; *she should snack*), and the infinitival *to* (TO):

$\text{VP} \rightarrow \text{MD VP}$	<i>She will snack</i>	[9.42]
$\text{VP} \rightarrow \text{VBD VP}$	<i>She had snacked</i>	[9.43]
$\text{VP} \rightarrow \text{VBZ VP}$	<i>She has been snacking</i>	[9.44]
$\text{VP} \rightarrow \text{VBN VP}$	<i>She has been snacking</i>	[9.45]
$\text{VP} \rightarrow \text{TO VP}$	<i>She wants to snack</i>	[9.46]
$\text{VP} \rightarrow \text{VP CC VP}$	<i>She buys and eats many snacks</i>	[9.47]

- 4832 Each of these productions uses recursion, with the VP non-terminal appearing in both the
 4833 LHS and RHS. This enables the creation of complex verb phrases, such as *She will have*
 4834 *wanted to have been snacking*.

Transitive verbs take noun phrases as direct objects, and ditransitive verbs take two direct objects:

$\text{VP} \rightarrow \text{VBZ NP}$	<i>She teaches algebra</i>	[9.48]
$\text{VP} \rightarrow \text{VBG NP}$	<i>She has been teaching algebra</i>	[9.49]
$\text{VP} \rightarrow \text{VBD NP NP}$	<i>She taught her brother algebra</i>	[9.50]

These productions are *not* recursive, so a unique production is required for each verb part-of-speech. They also do not distinguish transitive from intransitive verbs, so the resulting grammar overgenerates examples like **She sleeps sushi* and **She learns Boyang algebra*. Sentences can also be direct objects:

$\text{VP} \rightarrow \text{VBZ S}$	<i>Asha wants to eat the kimchi</i>	[9.51]
$\text{VP} \rightarrow \text{VBZ SBAR}$	<i>Asha knows that Boyang eats the kimchi</i>	[9.52]

- 4835 The first production overgenerates, licensing sentences like **Asha sees Boyang eats the kim-*
 4836 *chi*. This problem could be addressed by designing a more specific set of sentence non-
 4837 terminals, indicating whether the main verb can be conjugated.

Verbs can also be modified by prepositional phrases and adverbial phrases:

$\text{VP} \rightarrow \text{VBZ PP}$	<i>She studies at night</i>	[9.53]
$\text{VP} \rightarrow \text{VBZ ADVP}$	<i>She studies intensively</i>	[9.54]
$\text{VP} \rightarrow \text{ADVP VBG}$	<i>She is not studying</i>	[9.55]

4838 Again, because these productions are not recursive, the grammar must include productions
 4839 for every verb part-of-speech.

A special set of verbs, known as **copula**, can take **predicative adjectives** as direct objects:

$VP \rightarrow VBZ\ ADJP$	<i>She is hungry</i>	[9.56]
$VP \rightarrow VBP\ ADJP$	<i>Success seems increasingly unlikely</i>	[9.57]

4840 The PTB does not have a special non-terminal for copular verbs, so this production generates
 4841 non-grammatical examples such as **She eats tall*.

Particles (PRT as a phrase; RP as a part-of-speech) work to create phrasal verbs:

$VP \rightarrow VB\ PRT$	<i>She told them to fuck off</i>	[9.58]
$VP \rightarrow VBD\ PRT\ NP$	<i>They gave up their ill-gotten gains</i>	[9.59]

4842 As the second production shows, particle productions are required for all configurations
 4843 of verb parts-of-speech and direct objects.

4844 9.2.3.4 Other constituents

The remaining constituents require far fewer productions. **Prepositional phrases** almost always consist of a preposition and a noun phrase,

$PP \rightarrow IN\ NP$	<i>the whiteness of the whale</i>	[9.60]
$PP \rightarrow TO\ NP$	<i>What the white whale was to Ahab, has been hinted.</i>	[9.61]

Similarly, complement clauses consist of a complementizer (usually a preposition, possibly null) and a sentence,

$SBAR \rightarrow IN\ S$	<i>She said that it was spicy</i>	[9.62]
$SBAR \rightarrow S$	<i>She said it was spicy</i>	[9.63]

Adverbial phrases are usually bare adverbs ($ADVP \rightarrow RB$), with a few exceptions:

$ADVP \rightarrow RB\ RBR$	<i>They went considerably further</i>	[9.64]
$ADVP \rightarrow ADVP\ PP$	<i>They went considerably further than before</i>	[9.65]

4845 The tag RBR is a comparative adverb.

Adjectival phrases extend beyond bare adjectives ($\text{ADJP} \rightarrow \text{JJ}$) in a number of ways:

$\text{ADJP} \rightarrow \text{RB JJ}$	<i>very hungry</i>	[9.66]
$\text{ADJP} \rightarrow \text{RBR JJ}$	<i>more hungry</i>	[9.67]
$\text{ADJP} \rightarrow \text{JJS JJ}$	<i>best possible</i>	[9.68]
$\text{ADJP} \rightarrow \text{RB JJR}$	<i>even bigger</i>	[9.69]
$\text{ADJP} \rightarrow \text{JJ CC JJ}$	<i>high and mighty</i>	[9.70]
$\text{ADJP} \rightarrow \text{JJ JJ}$	<i>West German</i>	[9.71]
$\text{ADJP} \rightarrow \text{RB VBN}$	<i>previously reported</i>	[9.72]

4846 The tags JJR and JJS refer to comparative and superlative adjectives respectively.

All of these phrase types can be coordinated:

$\text{PP} \rightarrow \text{PP CC PP}$	<i>on time and under budget</i>	[9.73]
$\text{ADVP} \rightarrow \text{ADVP CC ADVP}$	<i>now and two years ago</i>	[9.74]
$\text{ADJP} \rightarrow \text{ADJP CC ADJP}$	<i>quaint and rather deceptive</i>	[9.75]
$\text{SBAR} \rightarrow \text{SBAR CC SBAR}$	<i>whether they want control</i>	[9.76]
	<i>or whether they want exports</i>	

4847 9.2.4 Grammatical ambiguity

4848 Context-free parsing is useful not only because it determines whether a sentence is grammatical, but mainly because the constituents and their relations can be applied to tasks
 4849 such as information extraction (chapter 17) and sentence compression (Jing, 2000; Clarke
 4850 and Lapata, 2008). However, the **ambiguity** of wide-coverage natural language grammars
 4851 poses a serious problem for such potential applications. As an example, Figure 9.13 shows
 4852 two possible analyses for the simple sentence *We eat sushi with chopsticks*, depending on
 4853 whether the *chopsticks* modify *eat* or *sushi*. Realistic grammars can license thousands or
 4854 even millions of parses for individual sentences. **Weighted context-free grammars** solve
 4855 this problem by attaching weights to each production, and selecting the derivation with
 4856 the highest score. This is the focus of chapter 10.
 4857

4858 9.3 *Mildly context-sensitive languages

4859 Beyond context-free languages lie **context-sensitive languages**, in which the expansion
 4860 of a non-terminal depends on its neighbors. In the general class of context-sensitive
 4861 languages, computation becomes much more challenging: the membership problem for
 4862 context-sensitive languages is PSPACE-complete. Since PSPACE contains the complexity
 4863 class NP (problems that can be solved in polynomial time on a non-deterministic Turing

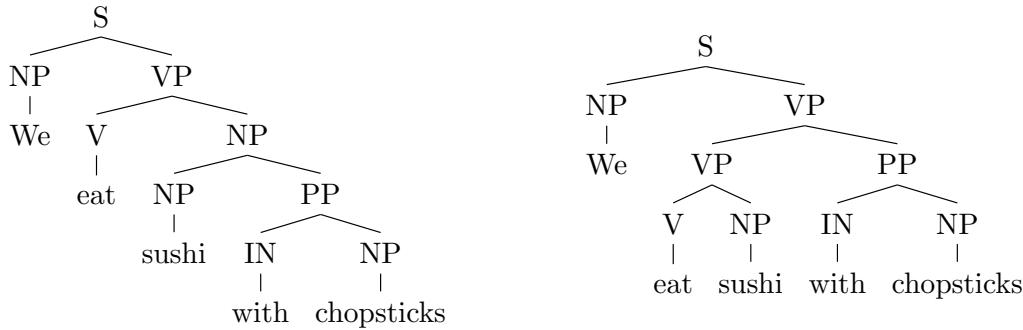


Figure 9.13: Two derivations of the same sentence

4864 machine), PSPACE-complete problems cannot be solved efficiently if $P \neq NP$. Thus, de-
 4865 signing an efficient parsing algorithm for the full class of context-sensitive languages is
 4866 probably hopeless.¹⁰

4867 However, Joshi (1985) identifies a set of properties that define **mildly context-sensitive**
 4868 **languages**, which are a strict subset of context-sensitive languages. Like context-free lan-
 4869 guages, mildly context-sensitive languages are efficiently parseable. However, the mildly
 4870 context-sensitive languages include non-context-free languages, such as the “copy lan-
 4871 guage” $\{ww \mid w \in \Sigma^*\}$ and the language $a^m b^n c^m d^n$. Both are characterized by **cross-**
 4872 **serial dependencies**, linking symbols at long distance across the string.¹¹ For example, in
 4873 the language $a^n b^m c^n d^m$, each a symbol is linked to exactly one c symbol, regardless of the
 4874 number of intervening b symbols.

4875 9.3.1 Context-sensitive phenomena in natural language

4876 Such phenomena are occasionally relevant to natural language. A classic example is found
 4877 in Swiss-German (Shieber, 1985), in which sentences such as *we let the children help Hans*
 4878 *paint the house* are realized by listing all nouns before all verbs, i.e., *we the children Hans the*
 4879 *house let help paint*. Furthermore, each noun’s determiner is dictated by the noun’s **case**
 4880 **marking** (the role it plays with respect to the verb). Using an argument that is analogous
 4881 to the earlier discussion of center-embedding (§ 9.2), Shieber argues that these case mark-
 4882 ing constraints are a cross-serial dependency, homomorphic to $a^m b^n c^m d^n$, and therefore
 4883 not context-free.

¹⁰If $P \neq NP$, then it contains problems that cannot be solved in polynomial time on a non-deterministic Turing machine; equivalently, solutions to these problems cannot even be checked in polynomial time (Arora and Barak, 2009).

¹¹A further condition of the set of mildly-context-sensitive languages is *constant growth*: if the strings in the language are arranged by length, the gap in length between any pair of adjacent strings is bounded by some language specific constant. This condition excludes languages such as $\{a^{2^n} \mid n \geq 0\}$.

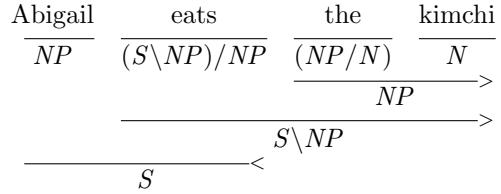


Figure 9.14: A syntactic analysis in CCG involving forward and backward function application

As with the move from regular to context-free languages, mildly context-sensitive languages can be motivated by expedience. While infinite sequences of cross-serial dependencies cannot be handled by context-free grammars, even finite sequences of cross-serial dependencies are more convenient to handle using a mildly context-sensitive formalism like **tree-adjoining grammar** (TAG) and **combinatory categorial grammar** (CCG). Furthermore, TAG-inspired parsers have been shown to be particularly effective in parsing the Penn Treebank (Collins, 1997; Carreras et al., 2008), and CCG plays a leading role in current research on semantic parsing (Zettlemoyer and Collins, 2005). Furthermore, these two formalisms are weakly equivalent: any language that can be specified in TAG can also be specified in CCG, and vice versa (Joshi et al., 1991). The remainder of the chapter gives a brief overview of CCG, but you are encouraged to consult Joshi and Schabes (1997) and Steedman and Baldridge (2011) for more detail on TAG and CCG respectively.

9.3.2 Combinatory categorial grammar

In combinatory categorial grammar, structural analyses are built up through a small set of generic combinatorial operations, which apply to immediately adjacent sub-structures. These operations act on the categories of the sub-structures, producing a new structure with a new category. The basic categories include S (sentence), NP (noun phrase), VP (verb phrase) and N (noun). The goal is to label the entire span of text as a sentence, S .

Complex categories, or types, are constructed from the basic categories, parentheses, and forward and backward slashes: for example, S/NP is a complex type, indicating a sentence that is lacking a noun phrase to its right; $S\backslash NP$ is a sentence lacking a noun phrase to its left. Complex types act as functions, and the most basic combinatory operations are function application to either the right or left neighbor. For example, the type of a verb phrase, such as *eats*, would be $S\backslash NP$. Applying this function to a subject noun phrase to its left results in an analysis of *Abigail eats* as category S , indicating a successful parse.

Transitive verbs must first be applied to the direct object, which in English appears to the right of the verb, before the subject, which appears on the left. They therefore have the more complex type $(S\backslash NP)/NP$. Similarly, the application of a determiner to the noun at

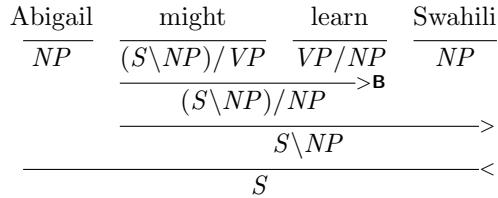


Figure 9.15: A syntactic analysis in CCG involving function composition (example modified from Steedman and Baldridge, 2011)

4913 its right results in a noun phrase, so determiners have the type NP/N. Figure 9.14 pro-
 4914 vides an example involving a transitive verb and a determiner. A key point from this
 4915 example is that it can be trivially transformed into phrase-structure tree, by treating each
 4916 function application as a constituent phrase. Indeed, when CCG's only combinatory op-
 4917 erators are forward and backward function application, it is equivalent to context-free
 4918 grammar. However, the location of the "effort" has changed. Rather than designing good
 4919 productions, the grammar designer must focus on the **lexicon** — choosing the right cate-
 4920 gories for each word. This makes it possible to parse a wide range of sentences using only
 4921 a few generic combinatory operators.

4922 Things become more interesting with the introduction of two additional operators:
 4923 **composition** and **type-raising**. Function composition enables the combination of com-
 4924 plex types: $X/Y \circ Y/Z \Rightarrow_B X/Z$ (forward composition) and $Y\backslash Z \circ X\backslash Y \Rightarrow_B X\backslash Z$ (back-
 4925 ward composition).¹² Composition makes it possible to "look inside" complex types, and
 4926 combine two adjacent units if the "input" for one is the "output" for the other. Figure 9.15
 4927 shows how function composition can be used to handle modal verbs. While this sen-
 4928 tence can be parsed using only function application, the composition-based analysis is
 4929 preferable because the unit *might learn* functions just like a transitive verb, as in the exam-
 4930 ple *Abigail studies Swahili*. This in turn makes it possible to analyze conjunctions such as
 4931 *Abigail studies and might learn Swahili*, attaching the direct object *Swahili* to the entire con-
 4932 joined verb phrase *studies and might learn*. The Penn Treebank grammar fragment from
 4933 § 9.2.3 would be unable to handle this case correctly: the direct object *Swahili* could attach
 4934 only to the second verb *learn*.

4935 Type raising converts an element of type X to a more complex type: $X \Rightarrow_T T/(T\backslash X)$
 4936 (forward type-raising to type T), and $X \Rightarrow_T T\backslash(T/X)$ (backward type-raising to type
 4937 T). Type-raising makes it possible to reverse the relationship between a function and its
 4938 argument — by transforming the argument into a function over functions over arguments!
 4939 An example may help. Figure 9.15 shows how to analyze an object relative clause, *a story*
 4940 *that Abigail tells*. The problem is that *tells* is a transitive verb, expecting a direct object to
 4941 its right. As a result, *Abigail tells* is not a valid constituent. The issue is resolved by raising

¹²The subscript **B** follows notation from Curry and Feys (1958).

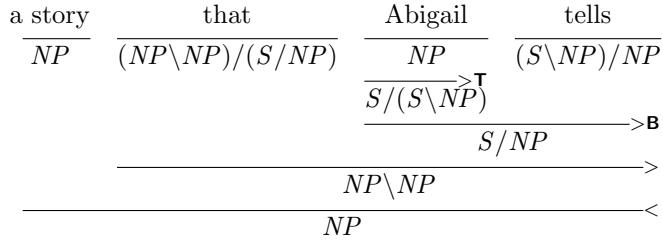


Figure 9.16: A syntactic analysis in CCG involving an object relative clause (based on slides from Alex Clark)

4942 *Abigail* from NP to the complex type $(S / NP) \setminus NP$. This function can then be combined
 4943 with the transitive verb *tells* by forward composition, resulting in the type (S / NP) , which
 4944 is a sentence lacking a direct object to its right.¹³ From here, we need only design the
 4945 lexical entry for the complementizer *that* to expect a right neighbor of type (S / NP) , and
 4946 the remainder of the derivation can proceed by function application.

4947 Composition and type-raising give CCG considerable power and flexibility, but at a price.
 4948 The simple sentence *Abigail tells Max* can be parsed in two different ways: by function
 4949 application (first forming the verb phrase *tells Max*), and by type-raising and compo-
 4950 sition (first forming the non-constituent *Abigail tells*). This **derivational ambiguity** does
 4951 not affect the resulting linguistic analysis, so it is sometimes known as **spurious ambi-**
 4952 **guity**. Hockenmaier and Steedman (2007) present a translation algorithm for converting
 4953 the Penn Treebank into CCG derivations, using composition and type-raising only when
 4954 necessary.

4955 Exercises

- 4956 1. Sketch out the state diagram for finite-state acceptors for the following languages
 4957 on the alphabet $\{a, b\}$.
- 4958 a) Even-length strings. (Be sure to include 0 as an even number.)
- 4959 b) Strings that contain *aaa* as a substring.
- 4960 c) Strings containing an even number of *a* and an odd number of *b* symbols.
- 4961 d) Strings in which the substring *bbb* must be terminal if it appears — the string
 4962 need not contain *bbb*, but if it does, nothing can come after it.
- 4963 2. Levenshtein edit distance is the number of insertions, substitutions, or deletions
 4964 required to convert one string to another.

¹³The missing direct object would be analyzed as a **trace** in CFG-like approaches to syntax, including the Penn Treebank.

- 4965 a) Define a finite-state acceptor that accepts all strings with edit distance 1 from
 4966 the target string, *target*.
 4967 b) Now think about how to generalize your design to accept all strings with edit
 4968 distance from the target string equal to d . If the target string has length ℓ , what
 4969 is the minimal number of states required?
- 4970 3. Construct an FSA in the style of Figure 9.3, which handles the following examples:

- 4971 • *nation*/N, *national*/ADJ, *nationalize*/V, *nationalizer*/N
 4972 • *America*/N, *American*/ADJ, *Americanize*/V, *Americanizer*/N

4973 Be sure that your FSA does not accept any further derivations, such as **nationalizeral*
 4974 and **Americanizern*.

- 4975 4. Show how to construct a trigram language model in a weighted finite-state acceptor.
 4976 Make sure that you handle the edge cases at the beginning and end of the sequence
 4977 accurately.
- 4978 5. Extend the FST in Figure 9.6 to handle the other two parts of rule 1a of the Porter
 4979 stemmer: *-sses* → *ss*, and *-ies* → *-i*.

- 4980 6. § 9.1.4.4 describes T_O , a transducer that captures English orthography by transduc-
 4981 ing *cook + ed* → *cooked* and *bake + ed* → *baked*. Design an unweighted finite-state
 4982 transducer that captures this property of English orthography.

4983 Next, augment the transducer to appropriately model the suffix *-s* when applied to
 4984 words ending in *s*, e.g. *kiss+s* → *kisses*.

- 4985 7. Add parenthesization to the grammar in Figure 9.11 so that it is no longer ambigu-
 4986 ous.
- 4987 8. Construct three examples — a noun phrase, a verb phrase, and a sentence — which
 4988 can be derived from the Penn Treebank grammar fragment in § 9.2.3, yet are not
 4989 grammatical. Avoid reusing examples from the text. Optionally, propose corrections
 4990 to the grammar to avoid generating these cases.
- 4991 9. Produce parses for the following sentences, using the Penn Treebank grammar frag-
 4992 ment from § 9.2.3.

- 4993 (9.19) This aggression will not stand.
 4994 (9.20) I can get you a toe.
 4995 (9.21) Sometimes you eat the bar and sometimes the bar eats you.

4996 Then produce parses for three short sentences from a news article from this week.

4997 10. * One advantage of CCG is its flexibility in handling coordination:

4998 (9.22) *Abigail and Max speak Swahili*

4999 (9.23) *Abigail speaks and Max understands Swahili*

Define the lexical entry for *and* as

$$\text{and} := (X/X) \setminus X, \quad [9.77]$$

5000 where X can refer to any type. Using this lexical entry, show how to parse the two
5001 examples above. In the second example, *Swahili* should be combined with the coor-
5002 dination *Abigail speaks and Max understands*, and not just with the verb *understands*.

5003 **Chapter 10**

5004 **Context-free parsing**

5005 Parsing is the task of determining whether a string can be derived from a given context-
5006 free grammar, and if so, how. The parse structure can answer basic questions of who-did-
5007 what-to-whom, and is useful for various downstream tasks, such as semantic analysis
5008 (chapter 12 and 13) and information extraction (chapter 17).

For a given input and grammar, how many parse trees are there? Consider a minimal context-free grammar with only one non-terminal, X , and the following productions:

$$\begin{aligned} X &\rightarrow X \ X \\ X &\rightarrow aardvark \mid abacus \mid \dots \mid zyther \end{aligned}$$

The second line indicates unary productions to every nonterminal in Σ . In this grammar, the number of possible derivations for a string w is equal to the number of binary bracketings, e.g.,

$$(((w_1 w_2) w_3) w_4) w_5), \quad (((w_1 (w_2 w_3)) w_4) w_5), \quad ((w_1 (w_2 (w_3 w_4))) w_5), \quad \dots$$

5009 The number of such bracketings is a **Catalan number**, which grows super-exponentially
5010 in the length of the sentence, $C_n = \frac{(2n)!}{(n+1)n!}$. As with sequence labeling, it is only possible to
5011 exhaustively search the space of parses by resorting to locality assumptions, which make it
5012 possible to search efficiently by reusing shared substructures with dynamic programming.
5013 This chapter focuses on a bottom-up dynamic programming algorithm, which enables
5014 exhaustive search of the space of possible parses, but imposes strict limitations on the
5015 form of scoring function. These limitations can be relaxed by abandoning exhaustive
5016 search. Non-exact search methods will be briefly discussed at the end of this chapter, and
5017 one of them — **transition-based parsing** — will be the focus of chapter 11.

S	\rightarrow	NP VP
NP	\rightarrow	NP PP <i>we</i> <i>sushi</i> <i>chopsticks</i>
PP	\rightarrow	IN NP
IN	\rightarrow	<i>with</i>
VP	\rightarrow	V NP VP PP
V	\rightarrow	<i>eat</i>

Table 10.1: A toy example context-free grammar

5018 **10.1 Deterministic bottom-up parsing**

5019 The **CKY algorithm**¹ is a bottom-up approach to parsing in a context-free grammar. It
 5020 efficiently tests whether a string is in a language, without enumerating all possible parses.
 5021 The algorithm first forms small constituents, and then tries to merge them into larger
 5022 constituents.

5023 To understand the algorithm, consider the input, *We eat sushi with chopsticks*. According-
 5024 ing to the toy grammar in Table 10.1, each terminal symbol can be generated by exactly
 5025 one unary production, resulting in the sequence NP V NP IN NP. Next, we try to apply
 5026 binary productions to merge adjacent symbols into larger constituents: for example, V
 5027 NP can be merged into a verb phrase (VP), and IN NP can be merged into a prepositional
 5028 phrase (PP). Bottom-up parsing tries to find some series of mergers that ultimately results
 5029 in the start symbol S covering the entire input.

5030 The CKY algorithm systematizes this approach, incrementally constructing a table t in
 5031 which each cell $t[i, j]$ contains the set of nonterminals that can derive the span $w_{i+1:j}$. The
 5032 algorithm fills in the upper right triangle of the table; it begins with the diagonal, which
 5033 corresponds to substrings of length 1, and then computes derivations for progressively
 5034 larger substrings, until reaching the upper right corner $t[0, M]$, which corresponds to the
 5035 entire input, $w_{1:M}$. If the start symbol S is in $t[0, M]$, then the string w is in the language
 5036 defined by the grammar. This process is detailed in Algorithm 13, and the resulting data
 5037 structure is shown in Figure 10.1. Informally, here's how it works:

- 5038 • Begin by filling in the diagonal: the cells $t[m - 1, m]$ for all $m \in \{1, 2, \dots, M\}$. These
 5039 cells are filled with terminal productions that yield the individual tokens; for the
 5040 word $w_2 = \text{sushi}$, we fill in $t[1, 2] = \{\text{NP}\}$, and so on.
- 5041 • Then fill in the next diagonal, in which each cell corresponds to a subsequence of
 5042 length two: $t[0, 2], t[1, 3], \dots, t[M - 2, M]$. These cells are filled in by looking for
 5043 binary productions capable of producing at least one entry in each of the cells corre-

¹The name is for Cocke-Kasami-Younger, the inventors of the algorithm. It is a special case **chart parsing**, because its stores reusable computations in a chart-like data structure.

Algorithm 13 The CKY algorithm for parsing a sequence $w \in \Sigma^*$ in a context-free grammar $G = (N, \Sigma, R, S)$, with non-terminals N , production rules R , and start symbol S . The grammar is assumed to be in Chomsky normal form (§ 9.2.1.2). The function $\text{PICKFROM}(b[i, j, X])$ selects an element of the set $b[i, j, X]$ arbitrarily. All values of t and b are initialized to \emptyset .

```

1: procedure CKY( $w, G = (N, \Sigma, R, S)$ )
2:   for  $m \in \{1 \dots M\}$  do
3:      $t[m - 1, m] \leftarrow \{X : (X \rightarrow w_m) \in R\}$ 
4:   for  $\ell \in \{2, 3, \dots, M\}$  do                                 $\triangleright$  Iterate over constituent lengths
5:     for  $m \in \{0, 1, \dots, M - \ell\}$  do           $\triangleright$  Iterate over left endpoints
6:       for  $k \in \{m + 1, m + 2, \dots, m + \ell - 1\}$  do       $\triangleright$  Iterate over split points
7:         for  $(X \rightarrow Y Z) \in R$  do           $\triangleright$  Iterate over rules
8:           if  $Y \in t[m, k] \wedge Z \in t[k, m + \ell]$  then
9:              $t[m, m + \ell] \leftarrow t[m, m + \ell] \cup X$            $\triangleright$  Add non-terminal to table
10:             $b[m, m + \ell, X] \leftarrow b[m, m + \ell, X] \cup (Y, Z, k)$        $\triangleright$  Add back-pointers
11:   if  $S \in t[0, M]$  then
12:     return TRACEBACK( $S, 0, M, b$ )
13:   else
14:     return  $\emptyset$ 
15: procedure TRACEBACK( $X, i, j, b$ )
16:   if  $j = i + 1$  then
17:     return  $X$ 
18:   else
19:      $(Y, Z, k) \leftarrow \text{PICKFROM}(b[i, j, X])$ 
20:     return  $X \rightarrow (\text{TRACEBACK}(Y, i, k, b), \text{TRACEBACK}(Z, k, j, b))$ 

```

5044 sponding to left and right children. For example, the cell $t[1, 3]$ includes VP because
 5045 the grammar includes the production $\text{VP} \rightarrow \text{V NP}$, and the chart contains $\text{V} \in t[1, 2]$
 5046 and $\text{NP} \in t[2, 3]$.

- 5047 • At the next diagonal, the entries correspond to spans of length three. At this level,
 5048 there is an additional decision at each cell: where to split the left and right children.
 5049 The cell $t[i, j]$ corresponds to the subsequence $w_{i+1:j}$, and we must choose some
 5050 *split point* $i < k < j$, so that $w_{i+1:k}$ is the left child and $w_{k+1:j}$ is the right child. We
 5051 consider all possible k , looking for productions that generate elements in $t[i, k]$ and
 5052 $t[k, j]$; the left-hand side of all such productions can be added to $t[i, j]$. When it is
 5053 time to compute $t[i, j]$, the cells $t[i, k]$ and $t[k, j]$ are guaranteed to be complete, since
 5054 these cells correspond to shorter sub-strings of the input.

- 5055 • The process continues until we reach $t[0, M]$.

5056 Figure 10.1 shows the chart that arises from parsing the sentence *We eat sushi with chop-*
 5057 *sticks* using the grammar defined above.

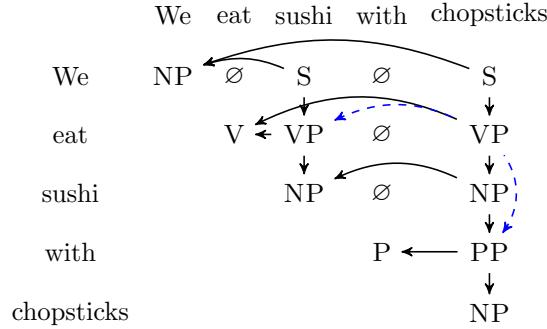


Figure 10.1: An example completed CKY chart. The solid and dashed lines show the back pointers resulting from the two different derivations of VP in position $t[1, 5]$.

5058 10.1.1 Recovering the parse tree

5059 As with the Viterbi algorithm, it is possible to identify a successful parse by storing and
 5060 traversing an additional table of back-pointers. If we add an entry X to cell $t[i, j]$ by using
 5061 the production $X \rightarrow YZ$ and the split point k , then we store the back-pointer $b[i, j, X] =$
 5062 (Y, Z, k) . Once the table is complete, we can recover a parse by tracing this pointers,
 5063 starting at $b[0, M, S]$, and stopping when they ground out at terminal productions.

5064 For ambiguous sentences, there will be multiple paths to reach $S \in t[0, M]$. For exam-
 5065 ple, in Figure 10.1, the goal state $S \in t[0, M]$ is reached through the state $VP \in t[1, 5]$, and
 5066 there are two different ways to generate this constituent: one with *(eat sushi)* and *(with
 5067 chopsticks)* as children, and another with *(eat)* and *(sushi with chopsticks)* as children. The
 5068 presence of multiple paths indicates that the input can be generated by the grammar in
 5069 more than one way. In Algorithm 13, one of these derivations is selected arbitrarily. As
 5070 discussed in § 10.3, **weighted context-free grammars** can select a single parse that maxi-
 5071 mizes a scoring function.

5072 10.1.2 Non-binary productions

5073 The CKY algorithm assumes that all productions with non-terminals on the right-hand
 5074 side (RHS) are binary. But in real grammars, such as the one considered in chapter 9,
 5075 there will be productions with more than two elements on the right-hand side, and other
 5076 productions with only a single element.

- 5077 • Productions with more than two elements on the right-hand side can be **binarized**
 5078 by creating additional non-terminals, as described in § 9.2.1.2. For example, given
 5079 the production $VP \rightarrow V NP NP$ (for ditransitive verbs), we can convert to $VP \rightarrow$
 5080 $VP_{ditrans}/NP NP$, and then add the production $VP_{ditrans}/NP \rightarrow V NP$.

- 5081 • What about unary productions like $VP \rightarrow V$? In practice, this is handled by making
5082 a second pass on each diagonal, in which each cell $t[i, j]$ is augmented with all
5083 possible unary productions capable of generating each item already in the cell —
5084 formally, $t[i, j]$ is extended to its **unary closure**. Suppose the example grammar in
5085 Table 10.1 were extended to include the production $VP \rightarrow V$, enabling sentences
5086 with intransitive verb phrases, like *we eat*. Then the cell $t[1, 2]$ — corresponding to
5087 the word *eat* — would first include the set $\{V\}$, and would be augmented to the set
5088 $\{V, VP\}$ during this second pass.

5089 **10.1.3 Complexity**

5090 For an input of length M and a grammar with R productions and N non-terminals, the
5091 space complexity of the CKY algorithm is $\mathcal{O}(M^2N)$: the number of cells in the chart is
5092 $\mathcal{O}(M^2)$, and each cell must hold $\mathcal{O}(N)$ elements. The time complexity is $\mathcal{O}(M^3R)$: each
5093 cell is computed by searching over $\mathcal{O}(M)$ split points, with R possible productions for
5094 each split point. Both the time and space complexity are considerably worse than the
5095 Viterbi algorithm, which is linear in the length of the input.

5096 **10.2 Ambiguity**

5097 Syntactic ambiguity is endemic to natural language. Here are a few broad categories:

- 5098 • **Attachment ambiguity**: e.g., *We eat sushi with chopsticks, I shot an elephant in my pajamas*. In these examples, the prepositions (*with, in*) can attach to either the verb
5099 or the direct object.
5100
- 5101 • **Modifier scope**: e.g., *southern food store, plastic cup holder*. In these examples, the first
5102 word could be modifying the subsequent adjective, or the final noun.
5103
- 5104 • **Particle versus preposition**: e.g., *The puppy tore up the staircase*. Phrasal verbs like
5105 *tore up* often include particles which could also act as prepositions. This has struc-
5106 tural implications: if *up* is a preposition, then *up the staircase* is a prepositional
5107 phrase; if *up* is a particle, then *the staircase* is the direct object to the verb.
5108
- 5109 • **Complement structure**: e.g., *The students complained to the professor that they didn't
5110 understand*. This is another form of attachment ambiguity, where the complement
5111 *that they didn't understand* could attach to the main verb (*complained*), or to the indi-
5112 rect object (*the professor*).
5113
- 5114 • **Coordination scope**: e.g., *"I see," said the blind man, as he picked up the hammer and
5115 saw*. In this example, the lexical ambiguity for *saw* enables it to be coordinated either
5116 with the noun *hammer* or the verb *picked up*.

These forms of ambiguity can combine, so that seemingly simple headlines like *Fed raises interest rates* have dozens of possible analyses even in a minimal grammar. In a broad coverage grammar, typical sentences can have millions of parses. While careful grammar design can chip away at this ambiguity, a better strategy is to combine broad coverage parsers with data driven strategies for identifying the correct analysis.

10.2.1 Parser evaluation

Before continuing to parsing algorithms that are able to handle ambiguity, we stop to consider how to measure parsing performance. Suppose we have a set of *reference parses* — the ground truth — and a set of *system parses* that we would like to score. A simple solution would be per-sentence accuracy: the parser is scored by the proportion of sentences on which the system and reference parses exactly match.² But as any good student knows, it is better to get *partial credit*, which we can assign to analyses that correctly match parts of the reference parse. The PARSEval metrics (Grishman et al., 1992) score each system parse via:

Precision: the fraction of constituents in the system parse that match a constituent in the reference parse.

Recall: the fraction of constituents in the reference parse that match a constituent in the system parse.

In **labeled precision** and **recall**, the system must also match the phrase type for each constituent; in **unlabeled precision** and **recall**, it is only required to match the constituent structure. As in chapter 4, the precision and recall can be combined into an *F*-MEASURE, $F = \frac{2 \times P \times R}{P + R}$.

In Figure 10.2, suppose that the left tree is the system parse and the right tree is the reference parse. We have the following spans:

- $S \rightarrow w_{1:5}$ is *true positive*, because it appears in both trees.
- $VP \rightarrow w_{2:5}$ is *true positive* as well.
- $NP \rightarrow w_{3:5}$ is *false positive*, because it appears only in the system output.
- $PP \rightarrow w_{4:5}$ is *true positive*, because it appears in both trees.
- $VP \rightarrow w_{2:3}$ is *false negative*, because it appears only in the reference.

²Most parsing papers do not report results on this metric, but Finkel et al. (2008) find that a strong parser finds the exact correct parse on 35% of sentences of length ≤ 40 , and on 62% of parses of length ≤ 15 in the Penn Treebank.

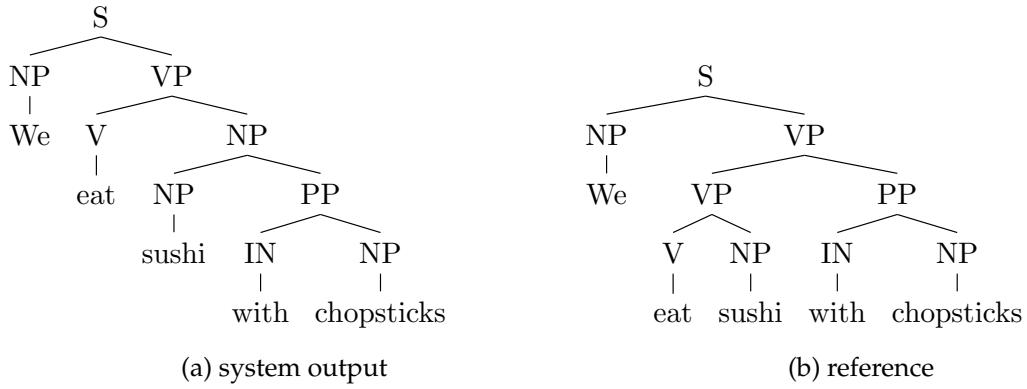


Figure 10.2: Two possible analyses from the grammar in Table 10.1

5143 The labeled and unlabeled precision of this parse is $\frac{3}{4} = 0.75$, and the recall is $\frac{3}{4} = 0.75$, for
 5144 an F-measure of 0.75. For an example in which precision and recall are not equal, suppose
 5145 the reference parse instead included the production $VP \rightarrow V NP PP$. In this parse, the
 5146 reference does not contain the constituent $w_{2:3}$, so the recall would be 1.³

5147 10.2.2 Local solutions

5148 Some ambiguity can be resolved locally. Consider the following examples,

5149 (10.1) We met the President on Monday.

5150 (10.2) We met the President of Mexico.

Each case ends with a preposition, which can be attached to the verb *met* or the noun phrase *the president*. This ambiguity can be resolved by using a labeled corpus to compare the likelihood of observing the preposition alongside each candidate attachment point,

$$p(on \mid met) \geq p(on \mid President) \quad [10.1]$$

$$p(of \mid met) \geq p(of \mid President). \quad [10.2]$$

5151 A comparison of these probabilities would successfully resolve this case (Hindle and
5152 Rooth, 1993). Other cases, such as the example ... *eat sushi with chopsticks*, require consider-
5153 ing the object of the preposition — consider the alternative ... *eat sushi with soy sauce*. With
5154 sufficient labeled data, the problem of prepositional phrase attachment can be treated as
5155 a classification task (Ratnaparkhi et al., 1994).

³While the grammar must be binarized before applying the CKY algorithm, evaluation is performed on the original parses. It is therefore necessary to “unbinarize” the output of a CKY-based parser, converting it back to the original grammar.

5156 However, there are inherent limitations to local solutions. While toy examples may
 5157 have just a few ambiguities to resolve, realistic sentences have thousands or millions of
 5158 possible parses. Furthermore, attachment decisions are interdependent, as shown in the
 5159 garden path example:

5160 (10.3) Cats scratch people with claws with knives.

5161 We may want to attach *with claws* to *scratch*, as would be correct in the shorter sentence
 5162 in *cats scratch people with claws*. But this leaves nowhere to attach *with knives*. The cor-
 5163 rect interpretation can be identified only by considering the attachment decisions jointly.
 5164 The huge number of potential parses may seem to make exhaustive search impossible.
 5165 But as with sequence labeling, locality assumptions make it possible to search this space
 5166 efficiently.

5167 10.3 Weighted Context-Free Grammars

5168 Let us define a derivation τ as a set of **anchored productions**,

$$\tau = \{X \rightarrow \alpha, (i, j, k)\}, \quad [10.3]$$

5169 with X corresponding to the left-hand side non-terminal and α corresponding to the right-
 5170 hand side. For grammars in Chomsky normal form, α is either a pair of non-terminals or
 5171 a terminal symbol. The indices i, j, k anchor the production in the input, with X deriving
 5172 the span $w_{i+1:j}$. For binary productions, $w_{i+1:k}$ indicates the span of the left child, and
 5173 $w_{k+1:j}$ indicates the span of the right child; for unary productions, k is ignored. For an
 5174 input w , the optimal parse is then,

$$\hat{\tau} = \underset{\tau \in \mathcal{T}(w)}{\operatorname{argmax}} \Psi(\tau), \quad [10.4]$$

5175 where $\mathcal{T}(w)$ is the set of derivations that yield the input w .

5176 The scoring function Ψ decomposes across anchored productions,

$$\Psi(\tau) = \sum_{(X \rightarrow \alpha, (i, j, k)) \in \tau} \psi(X \rightarrow \alpha, (i, j, k)). \quad [10.5]$$

5177 This is a locality assumption, akin to the assumption in Viterbi sequence labeling. In this
 5178 case, the assumption states that the overall score is a sum over scores of productions,
 5179 which are computed independently. In a **weighted context-free grammar** (WCFG), the
 5180 score of each anchored production $X \rightarrow (\alpha, i, j, k)$ is simply $\psi(X \rightarrow \alpha)$, ignoring the
 5181 anchors (i, j, k) . In other parsing models, the anchors can be used to access features of the
 5182 input, while still permitting efficient bottom-up parsing.

		$\psi(\cdot)$	$\exp \psi(\cdot)$
S	$\rightarrow \text{NP VP}$	0	1
NP	$\rightarrow \text{NP PP}$	-1	$\frac{1}{2}$
	$\rightarrow \text{we}$	-2	$\frac{1}{4}$
	$\rightarrow \text{sushi}$	-3	$\frac{1}{8}$
	$\rightarrow \text{chopsticks}$	-3	$\frac{1}{8}$
PP	$\rightarrow \text{IN NP}$	0	1
IN	$\rightarrow \text{with}$	0	1
VP	$\rightarrow \text{V NP}$	-1	$\frac{1}{2}$
	$\rightarrow \text{VP PP}$	-2	$\frac{1}{4}$
	$\rightarrow \text{MD V}$	-2	$\frac{1}{4}$
V	$\rightarrow \text{eat}$	0	1

Table 10.2: An example weighted context-free grammar (WCFG). The weights are chosen so that $\exp \psi(\cdot)$ sums to one over right-hand sides for each non-terminal; this is required by probabilistic context-free grammars, but not by WCFGs in general.

Example Consider the weighted grammar shown in Table 10.2, and the analysis in Figure 10.2b.

$$\begin{aligned} \Psi(\tau) = & \psi(S \rightarrow \text{NP VP}) + \psi(VP \rightarrow \text{VP PP}) + \psi(VP \rightarrow \text{V NP}) + \psi(PP \rightarrow \text{IN NP}) \\ & + \psi(\text{NP} \rightarrow \text{We}) + \psi(\text{V} \rightarrow \text{eat}) + \psi(\text{NP} \rightarrow \text{sushi}) + \psi(\text{IN} \rightarrow \text{with}) + \psi(\text{NP} \rightarrow \text{chopsticks}) \end{aligned} \quad [10.6]$$

$$= 0 - 2 - 1 + 0 - 2 + 0 - 3 + 0 - 3 = -11. \quad [10.7]$$

5183 In the alternative parse in Figure 10.2a, the production $VP \rightarrow VP PP$ (with score -2) is
 5184 replaced with the production $NP \rightarrow NP PP$ (with score -1); all other productions are the
 5185 same. As a result, the score for this parse is -10.

5186 This example hints at a big problem with WCFG parsing on non-terminals such as
 5187 NP, VP, and PP: a WCFG will *always* prefer either VP or NP attachment, without regard
 5188 to what is being attached! This problem is addressed in § 10.5.

5189 10.3.1 Parsing with weighted context-free grammars

5190 The optimization problem in Equation 10.4 can be solved by modifying the CKY algo-
 5191 rithm. In the deterministic CKY algorithm, each cell $t[i, j]$ stored a set of non-terminals
 5192 capable of deriving the span $w_{i+1:j}$. We now augment the table so that the cell $t[i, j, X]$
 5193 is the *score of the best derivation of $w_{i+1:j}$ from non-terminal X* . This score is computed
 5194 recursively: for the anchored binary production $(X \rightarrow Y Z, (i, j, k))$, we compute:

Algorithm 14 CKY algorithm for parsing a string $w \in \Sigma^*$ in a weighted context-free grammar (N, Σ, R, S) , where N is the set of non-terminals and R is the set of weighted productions. The grammar is assumed to be in Chomsky normal form (§ 9.2.1.2). The function TRACEBACK is defined in Algorithm 13.

```

procedure WCKY( $w, G = (N, \Sigma, R, S)$ )
  for all  $i, j, X$  do ▷ Initialization
     $t[i, j, X] \leftarrow 0$ 
     $b[i, j, X] \leftarrow \emptyset$ 
  for  $m \in \{1, 2, \dots, M\}$  do
    for all  $X \in N$  do
       $t[m, m + 1, X] \leftarrow \psi(X \rightarrow w_m, (m, m + 1, m))$ 
  for  $\ell \in \{2, 3, \dots, M\}$  do
    for  $m \in \{0, 1, \dots, M - \ell\}$  do
      for  $k \in \{m + 1, m + 2, \dots, m + \ell - 1\}$  do
         $t[m, m + \ell, X] \leftarrow \max_{k, Y, Z} \psi(X \rightarrow Y Z, (m, m + \ell, k)) + t[m, k, Y] + t[k, m + \ell, Z]$ 
         $b[m, m + \ell, X] \leftarrow \operatorname{argmax}_{k, Y, Z} \psi(X \rightarrow Y Z, (m + \ell, k)) + t[m, k, Y] + t[k, m + \ell, Z]$ 
  return TRACEBACK( $S, 0, M, b$ )

```

- 5195 • the score of the anchored production, $\psi(X \rightarrow Y Z, (i, j, k))$;
- 5196 • the score of the best derivation of the left child, $t[i, k, Y]$;
- 5197 • the score of the best derivation of the right child, $t[k, j, Z]$.

5198 These scores are combined by addition. As in the unscored CKY algorithm, the table
 5199 is constructed by considering spans of increasing length, so the scores for spans $t[i, k, Y]$
 5200 and $t[k, j, Z]$ are guaranteed to be available at the time we compute the score $t[i, j, X]$. The
 5201 value $t[0, M, S]$ is the score of the best derivation of w from the grammar. Algorithm 14
 5202 formalizes this procedure.

5203 As in unweighted CKY, the parse is recovered from the table of back pointers b , where
 5204 each $b[i, j, X]$ stores the argmax split point k and production $X \rightarrow Y Z$ in the derivation of
 5205 $w_{i+1:j}$ from X . The best parse can be obtained by tracing these pointers backwards from
 5206 $b[0, M, S]$, all the way to the terminal symbols. This is analogous to the computation of the
 5207 best sequence of labels in the Viterbi algorithm by tracing pointers backwards from the
 5208 end of the trellis. Note that we need only store back-pointers for the *best* path to $t[i, j, X]$;
 5209 this follows from the locality assumption that the global score for a parse is a combination
 5210 of the local scores of each production in the parse.

Example Let's revisit the parsing table in Figure 10.1. In a weighted CFG, each cell would include a score for each non-terminal; non-terminals that cannot be generated are

Algorithm 15 Generative model for derivations from probabilistic context-free grammars in Chomsky Normal Form (CNF).

```

procedure DRAWSUBTREE(X)
    sample  $(X \rightarrow \alpha) \sim p(\alpha | X)$ 
    if  $\alpha = (Y Z)$  then
        return DRAWSUBTREE(Y)  $\cup$  DRAWSUBTREE(Z)
    else
        return  $(X \rightarrow \alpha)$             $\triangleright$  In CNF, all unary productions yield terminal symbols

```

assumed to have a score of $-\infty$. The first diagonal contains the scores of unary productions: $t[0, 1, \text{NP}] = -2$, $t[1, 2, \text{V}] = 0$, and so on. At the next diagonal, we compute the scores for spans of length 2: $t[1, 3, \text{VP}] = -1 + 0 - 3 = -4$, $t[3, 5, \text{PP}] = 0 + 0 - 3 = -3$, and so on. Things get interesting when we reach the cell $t[1, 5, \text{VP}]$, which contains the score for the derivation of the span $w_{2:5}$ from the non-terminal VP. This score is computed as a max over two alternatives,

$$t[1, 5, \text{VP}] = \max(\psi(\text{VP} \rightarrow \text{VP PP}, (1, 3, 5)) + t[1, 3, \text{VP}] + t[3, 5, \text{PP}], \\ \psi(\text{VP} \rightarrow \text{V NP}, (1, 2, 5)) + t[1, 2, \text{V}] + t[2, 5, \text{NP}]) \quad [10.8]$$

$$= \max(-2 - 4 - 3, -1 + 0 - 7) = -8. \quad [10.9]$$

5211 Since the second case is the argmax, we set the back-pointer $b[1, 5, \text{VP}] = (\text{V}, \text{NP}, 2)$, enabling the optimal derivation to be recovered.

5213 **10.3.2 Probabilistic context-free grammars**

5214 **Probabilistic context-free grammars (PCFGs)** are a special case of weighted context-
5215 free grammars that arises when the weights correspond to probabilities. Specifically, the
5216 weight $\psi(X \rightarrow \alpha, (i, j, k)) = \log p(\alpha | X)$, where the probability of the right-hand side
5217 α is conditioned on the non-terminal X . These probabilities must be normalized over all
5218 possible right-hand sides, so that $\sum_\alpha p(\alpha | X) = 1$, for all X . For a given parse τ , the prod-
5219 uct of the probabilities of the productions is equal to $p(\tau)$, under the **generative model**
5220 $\tau \sim \text{DRAWSUBTREE}(S)$, where the function DRAWSUBTREE is defined in Algorithm 15.

5221 The conditional probability of a parse given a string is,

$$p(\tau | w) = \frac{p(\tau)}{\sum_{\tau' \in \mathcal{T}(w)} p(\tau')} = \frac{\exp \Psi(\tau)}{\sum_{\tau' \in \mathcal{T}(w)} \exp \Psi(\tau')}, \quad [10.10]$$

5222 where $\Psi(\tau) = \sum_{X \rightarrow \alpha, (i, j, k) \in \tau} \psi(X \rightarrow \alpha)$; the anchor is ignored. Because the probability
5223 is monotonic in the score $\Psi(\tau)$, the maximum likelihood parse can be identified by the
5224 CKY algorithm without modification. If a normalized probability $p(\tau | w)$ is required,
5225 the denominator of Equation 10.10 can be computed by the **inside recurrence**, described
5226 below.

Example The WCFG in Table 10.2 is designed so that the weights are log-probabilities, satisfying the constraint $\sum_{\alpha} \exp \psi(X \rightarrow \alpha) = 1$. As noted earlier, there are two parses in \mathcal{T} (*we eat sushi with chopsticks*), with scores $\Psi(\tau_1) = \log p(\tau_1) = -10$ and $\Psi(\tau_2) = \log p(\tau_2) = -11$. Therefore, the conditional probability $p(\tau_1 | \mathbf{w})$ is equal to,

$$p(\tau_1 | \mathbf{w}) = \frac{p(\tau_1)}{p(\tau_1) + p(\tau_2)} = \frac{\exp \Psi(\tau_1)}{\exp \Psi(\tau_1) + \exp \Psi(\tau_2)} = \frac{2^{-10}}{2^{-10} + 2^{-11}} = \frac{2}{3}. \quad [10.11]$$

5227 **The inside recurrence** The denominator of Equation 10.10 can be viewed as a language
5228 model, summing over all valid derivations of the string \mathbf{w} ,

$$p(\mathbf{w}) = \sum_{\tau': \text{yield}(\tau') = \mathbf{w}} p(\tau'). \quad [10.12]$$

Just as the CKY algorithm makes it possible to maximize over all such analyses, with a few modifications it can also compute their sum. Each cell $t[i, j, X]$ must store the log probability of deriving $\mathbf{w}_{i+1:j}$ from non-terminal X . To compute this, we replace the maximization over split points k and productions $X \rightarrow Y Z$ with a “log-sum-exp” operation, which exponentiates the log probabilities of the production and the children, sums them in probability space, and then converts back to the log domain:

$$t[i, j, X] = \log \sum_{k, Y, Z} \exp (\psi(X \rightarrow Y Z) + t[i, k, Y] + t[k, j, Z]) \quad [10.13]$$

$$= \log \sum_{k, Y, Z} \exp (\log p(Y Z | X) + \log p(Y \rightarrow \mathbf{w}_{i+1:k}) + \log p(Z \rightarrow \mathbf{w}_{k+1:j})) \quad [10.14]$$

$$= \log \sum_{k, Y, Z} p(Y Z | X) \times p(Y \rightarrow \mathbf{w}_{i+1:k}) \times p(Z \rightarrow \mathbf{w}_{k+1:j}) \quad [10.15]$$

$$= \log \sum_{k, Y, Z} p(Y Z, \mathbf{w}_{i+1:k}, \mathbf{w}_{k+1:j} | X) \quad [10.16]$$

$$= \log p(X \rightarrow \mathbf{w}_{i+1:j}). \quad [10.17]$$

5229 This is called the **inside recurrence**, because it computes the probability of each subtree
5230 as a combination of the probabilities of the smaller subtrees that are inside of it. The
5231 name implies a corresponding **outside recurrence**, which computes the probability of
5232 a non-terminal X spanning $\mathbf{w}_{i+1:j}$, joint with the outside context $(\mathbf{w}_{1:i}, \mathbf{w}_{j+1:M})$. This
5233 recurrence is described in § 10.4.3. The inside and outside recurrences are analogous to the
5234 forward and backward recurrences in probabilistic sequence labeling (see § 7.5.3.3). They
5235 can be used to compute the marginal probabilities of individual anchored productions,
5236 $p(X \rightarrow \alpha, (i, j, k) | \mathbf{w})$, summing over all possible derivations of \mathbf{w} .

5237 **10.3.3 *Semiring weighted context-free grammars**

The weighted and unweighted CKY algorithms can be unified with the inside recurrence using the same semiring notation described in § 7.7.3. The generalized recurrence is:

$$t[i, j, X] = \bigoplus_{k, Y, Z} \psi(X \rightarrow Y Z, (i, j, k)) \otimes t[i, k, Y] \otimes t[k, j, Z]. \quad [10.18]$$

5238 This recurrence subsumes all of the algorithms that we have encountered in this chapter.

5239 **Unweighted CKY.** When $\psi(X \rightarrow \alpha, (i, j, k))$ is a *Boolean truth value* $\{\top, \perp\}$, \otimes is logical
5240 conjunction, and \bigoplus is logical disjunction, then we derive the CKY recurrence for
5241 unweighted context-free grammars, discussed in § 10.1 and Algorithm 13.

5242 **Weighted CKY.** When $\psi(X \rightarrow \alpha, (i, j, k))$ is a scalar score, \otimes is addition, and \bigoplus is maxi-
5243 mization, then we derive the CKY recurrence for weighted context-free grammars,
5244 discussed in § 10.3 and Algorithm 14. When $\psi(X \rightarrow \alpha, (i, j, k)) = \log p(\alpha \mid X)$,
5245 this same setting derives the CKY recurrence for finding the maximum likelihood
5246 derivation in a probabilistic context-free grammar.

5247 **Inside recurrence.** When $\psi(X \rightarrow \alpha, (i, j, k))$ is a log probability, \otimes is addition, and $\bigoplus =$
5248 $\log \sum \exp$, then we derive the inside recurrence for probabilistic context-free gram-
5249 mmars, discussed in § 10.3.2. It is also possible to set $\psi(X \rightarrow \alpha, (i, j, k))$ directly equal
5250 to the probability $p(\alpha \mid X)$. In this case, \otimes is multiplication, and \bigoplus is addition.
5251 While this may seem more intuitive than working with log probabilities, there is the
5252 risk of underflow on long inputs.

5253 Regardless of how the scores are combined, the key point is the locality assumption:
5254 the score for a derivation is the combination of the independent scores for each anchored
5255 production, and these scores do not depend on any other part of the derivation. For exam-
5256 ple, if two non-terminals are siblings, the scores of productions from these non-terminals
5257 are computed independently. This locality assumption is analogous to the first-order
5258 Markov assumption in sequence labeling, where the score for transitions between tags
5259 depends only on the previous tag and current tag, and not on the history. As with se-
5260 quence labeling, this assumption makes it possible to find the optimal parse efficiently; its
5261 linguistic limitations are discussed in § 10.5.

5262 **10.4 Learning weighted context-free grammars**

5263 Like sequence labeling, context-free parsing is a form of structure prediction. As a result,
5264 WCFGs can be learned using the same set of algorithms: generative probabilistic models,
5265 structured perceptron, maximum conditional likelihood, and maximum margin learning.

5266 In all cases, learning requires a **treebank**, which is a dataset of sentences labeled with
 5267 context-free parses. Parsing research was catalyzed by the **Penn Treebank** (Marcus et al.,
 5268 1993), the first large-scale dataset of this type (see § 9.2.2). Phrase structure treebanks exist
 5269 for roughly two dozen other languages, with coverage mainly restricted to European and
 5270 East Asian languages, plus Arabic and Urdu.

5271 **10.4.1 Probabilistic context-free grammars**

Probabilistic context-free grammars are similar to hidden Markov models, in that they are generative models of text. In this case, the parameters of interest correspond to probabilities of productions, conditional on the left-hand side. As with hidden Markov models, these parameters can be estimated by relative frequency:

$$\psi(X \rightarrow \alpha) = \log p(X \rightarrow \alpha) \quad [10.19]$$

$$\hat{p}(X \rightarrow \alpha) = \frac{\text{count}(X \rightarrow \alpha)}{\text{count}(X)}. \quad [10.20]$$

5272 For example, the probability of the production $NP \rightarrow DET\ NN$ is the corpus count of
 5273 this production, divided by the count of the non-terminal NP . This estimator applies
 5274 to terminal productions as well: the probability of $NN \rightarrow whale$ is the count of how often
 5275 *whale* appears in the corpus as generated from an NN tag, divided by the total count of the
 5276 NN tag. Even with the largest treebanks — currently on the order of one million tokens
 5277 — it is difficult to accurately compute probabilities of even moderately rare events, such
 5278 as $NN \rightarrow whale$. Therefore, smoothing is critical for making PCFGs effective.

5279 **10.4.2 Feature-based parsing**

5280 The scores for each production can be computed as an inner product of weights and fea-
 5281 tures,

$$\psi(X \rightarrow \alpha) = \boldsymbol{\theta} \cdot \mathbf{f}(X, \alpha, (i, j, k), \mathbf{w}), \quad [10.21]$$

5282 where the feature vector $\mathbf{f}(X, \alpha)$ is a function of the left-hand side X , the right-hand side
 5283 α , the anchor indices (i, j, k) , and the input \mathbf{w} .

5284 The basic feature $\mathbf{f}(X, \alpha, (i, j, k)) = \{(X, \alpha)\}$ encodes only the identity of the pro-
 5285 duction itself, which is a discriminatively-trained model with the same expressiveness as
 5286 a PCFG. Features on anchored productions can include the words that border the span
 5287 w_i, w_{j+1} , the word at the split point w_{k+1} , the presence of a verb or noun in the left child
 5288 span $w_{i+1:k}$, and so on (Durrett and Klein, 2015). Scores on anchored productions can be
 5289 incorporated into CKY parsing without any modification to the algorithm, because it is
 5290 still possible to compute each element of the table $t[i, j, X]$ recursively from its immediate
 5291 children.

5292 Other features can be obtained by grouping elements on either the left-hand or right-
 5293 hand side: for example it can be particularly beneficial to compute additional features
 5294 by clustering terminal symbols, with features corresponding to groups of words with
 5295 similar syntactic properties. The clustering can be obtained from unlabeled datasets that
 5296 are much larger than any treebank, improving coverage. Such methods are described in
 5297 chapter 14.

Feature-based parsing models can be estimated using the usual array of discriminative learning techniques. For example, a structure perceptron update can be computed as (Carreras et al., 2008),

$$\mathbf{f}(\tau, \mathbf{w}^{(i)}) = \sum_{(X \rightarrow \alpha, (i, j, k)) \in \tau} \mathbf{f}(X, \alpha, (i, j, k), \mathbf{w}^{(i)}) \quad [10.22]$$

$$\hat{\tau} = \operatorname{argmax}_{\tau \in \mathcal{T}(\mathbf{w})} \mathbf{f}(\tau, \mathbf{w}^{(i)}) \quad [10.23]$$

$$\boldsymbol{\theta} \leftarrow \mathbf{f}(\tau^{(i)}, \mathbf{w}^{(i)}) - \mathbf{f}(\hat{\tau}, \mathbf{w}^{(i)}). \quad [10.24]$$

5298 A margin-based objective can be optimized by selecting $\hat{\tau}$ through cost-augmented decoding (§ 2.3.2), enforcing a margin of $\Delta(\hat{\tau}, \tau)$ between the hypothesis and the reference parse,
 5299 where Δ is a non-negative cost function, such as the Hamming loss (Stern et al., 2017). It
 5300 is also possible to train feature-based parsing models by conditional log-likelihood, as
 5301 described in the next section.

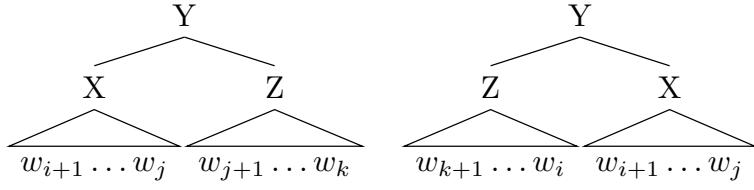
5303 10.4.3 *Conditional random field parsing

5304 The score of a derivation $\Psi(\tau)$ can be converted into a probability by normalizing over all
 5305 possible derivations,

$$p(\tau | \mathbf{w}) = \frac{\exp \Psi(\tau)}{\sum_{\tau' \in \mathcal{T}(\mathbf{w})} \exp \Psi(\tau')}. \quad [10.25]$$

5306 Using this probability, a WCFG can be trained by maximizing the conditional log-likelihood
 5307 of a labeled corpus.

5308 Just as in logistic regression and the conditional random field over sequences, the
 5309 gradient of the conditional log-likelihood is the difference between the observed and ex-
 5310 pected counts of each feature. The expectation $E_{\tau|\mathbf{w}}[\mathbf{f}(\tau, \mathbf{w}^{(i)}); \boldsymbol{\theta}]$ requires summing over
 5311 all possible parses, and computing the marginal probabilities of anchored productions,
 5312 $p(X \rightarrow \alpha, (i, j, k) | \mathbf{w})$. In CRF sequence labeling, marginal probabilities over tag bigrams
 5313 are computed by the two-pass **forward-backward algorithm** (§ 7.5.3.3). The analogue for
 5314 context-free grammars is the **inside-outside algorithm**, in which marginal probabilities
 5315 are computed from terms generated by an upward and downward pass over the parsing
 5316 chart:

Figure 10.3: The two cases faced by the outside recurrence in the computation of $\beta(i, j, X)$

- The upward pass is performed by the **inside recurrence**, which is described in § 10.3.2. Each inside variable $\alpha(i, j, X)$ is the score of deriving $w_{i+1:j}$ from the non-terminal X . In a PCFG, this corresponds to the log-probability $\log p(w_{i+1:j} \mid X)$. This is computed by the recurrence,

$$\alpha(i, j, X) \triangleq \log \sum_{(X \rightarrow Y \ Z)} \sum_{k=i+1}^j \exp (\psi(X \rightarrow Y \ Z, (i, j, k)) + \alpha(i, k, Y) + \alpha(k, j, Z)). \quad [10.26]$$

5317 The initial condition of this recurrence is $\alpha(m - 1, m, X) = \psi(X \rightarrow w_m)$. The de-
5318 nominator $\sum_{\tau \in \mathcal{T}(w)} \exp \Psi(\tau)$ is equal to $\exp \alpha(0, M, S)$.

- The downward pass is performed by the **outside recurrence**, which recursively populates the same table structure, starting at the root of the tree. Each outside variable $\beta(i, j, X)$ is the score of having a phrase of type X covering the span $(i + 1 : j)$, joint with the exterior context $w_{1:i}$ and $w_{j+1:M}$. In a PCFG, this corresponds to the log probability $\log p((X, i + 1, j), w_{1:i}, w_{j+1:M})$. Each outside variable is computed by the recurrence,

$$\exp \beta(i, j, X) \triangleq \sum_{(Y \rightarrow X \ Z)} \sum_{k=j+1}^M \exp [\psi(Y \rightarrow X \ Z, (i, k, j)) + \alpha(j, k, Z) + \beta(i, k, Y)] \quad [10.27]$$

$$+ \sum_{(Y \rightarrow Z \ X)} \sum_{k=0}^{i-1} \exp [\psi(Y \rightarrow Z \ X, (k, i, j)) + \alpha(k, i, Z) + \beta(k, j, Y)]. \quad [10.28]$$

5319 The first line of Equation 10.28 is the score under the condition that X is a left child
5320 of its parent, which spans $w_{i+1:k}$, with $k > j$; the second line is the score under the
5321 condition that X is a right child of its parent Y , which spans $w_{k+1:j}$, with $k < i$.
5322 The two cases are shown in Figure 10.3. In each case, we sum over all possible
5323 productions with X on the right-hand side. The parent Y is bounded on one side

5324 by either i or j , depending on whether X is a left or right child of Y ; we must sum
 5325 over all possible values for the other boundary. The initial conditions for the outside
 5326 recurrence are $\beta(0, M, S) = 0$ and $\beta(0, M, X \neq S) = -\infty$.

The marginal probability of a non-terminal X over span $w_{i+1:j}$ is written $p(X \rightsquigarrow w_{i+1:j} | w)$, and can be computed from the inside and outside scores,

$$p(X \rightsquigarrow w_{i+1:j} | w) = \frac{p(X \rightsquigarrow w_{i+1:j}, w)}{p(w)} \quad [10.29]$$

$$= \frac{p(w_{i+1:j} | X) \times p(X, w_{1:i}, w_{j+1:M})}{p(w)} \quad [10.30]$$

$$= \frac{\exp(\alpha(i, j, X) + \beta(i, j, X))}{\exp \alpha(0, M, S)}. \quad [10.31]$$

5327 Marginal probabilities of individual productions can be computed similarly (see exercise
 5328 2). These marginal probabilities can be used for training a conditional random field parser,
 5329 and also for the task of unsupervised **grammar induction**, in which a PCFG is estimated
 5330 from a dataset of unlabeled text (Lari and Young, 1990; Pereira and Schabes, 1992).

5331 10.4.4 Neural context-free grammars

5332 Recent work has applied neural representations to parsing, representing each span with
 5333 a dense numerical vector (Socher et al., 2013; Durrett and Klein, 2015; Cross and Huang,
 5334 2016).⁴ For example, the anchor (i, j, k) and sentence w can be associated with a fixed-
 5335 length column vector,

$$\mathbf{v}_{(i,j,k)} = [\mathbf{u}_{w_{i-1}}; \mathbf{u}_{w_i}; \mathbf{u}_{w_{j-1}}; \mathbf{u}_{w_j}; \mathbf{u}_{w_{k-1}}; \mathbf{u}_{w_k}], \quad [10.32]$$

where \mathbf{u}_{w_i} is a word embedding associated with the word w_i . The vector $\mathbf{v}_{i,j,k}$ can then be passed through a feedforward neural network, and used to compute the score of the anchored production. For example, this score can be computed as a bilinear product (Durrett and Klein, 2015),

$$\tilde{\mathbf{v}}_{(i,j,k)} = \text{FeedForward}(\mathbf{v}_{(i,j,k)}) \quad [10.33]$$

$$\psi(X \rightarrow \alpha, (i, j, k)) = \tilde{\mathbf{v}}_{(i,j,k)}^\top \Theta \mathbf{f}(X \rightarrow \alpha), \quad [10.34]$$

5336 where $\mathbf{f}(X \rightarrow \alpha)$ is a vector of discrete features of the production, and Θ is a parameter
 5337 matrix. The matrix Θ and the parameters of the feedforward network can be learned by
 5338 backpropagating from an objective such as the margin loss or the negative conditional
 5339 log-likelihood.

⁴Earlier work on neural constituent parsing used transition-based parsing algorithms (§ 10.6.2) rather than CKY-style chart parsing (Henderson, 2004; Titov and Henderson, 2007).

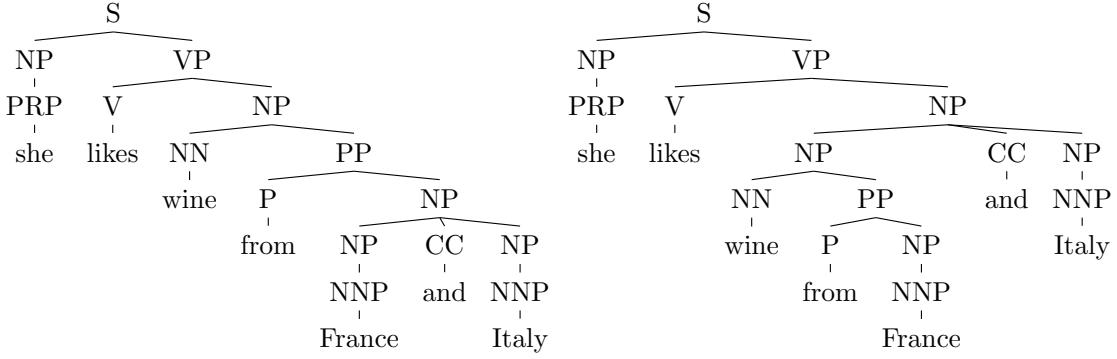


Figure 10.4: The left parse is preferable because of the conjunction of phrases headed by *France* and *Italy*, but these parses cannot be distinguished by a WCFG.

5340 10.5 Grammar refinement

5341 The locality assumptions underlying CFG parsing depend on the granularity of the non-
 5342 terminals. For the Penn Treebank non-terminals, there are several reasons to believe that
 5343 these assumptions are too strong to enable accurate parsing (Johnson, 1998):

- 5344 • The context-free assumption is too strict: for example, the probability of the produc-
 5345 tion $NP \rightarrow NP\ PP$ is much higher (in the PTB) if the parent of the noun phrase is a
 5346 verb phrase (indicating that the NP is a direct object) than if the parent is a sentence
 5347 (indicating that the NP is the subject of the sentence).
- 5348 • The Penn Treebank non-terminals are too coarse: there are many kinds of noun
 5349 phrases and verb phrases, and accurate parsing sometimes requires knowing the
 5350 difference. As we have already seen, when faced with prepositional phrase at-
 5351 tachment ambiguity, a weighted CFG will either always choose NP attachment (if
 5352 $\psi(NP \rightarrow NP\ PP) > \psi(VP \rightarrow VP\ PP)$), or it will always choose VP attachment. To
 5353 get more nuanced behavior, more fine-grained non-terminals are needed.
- 5354 • More generally, accurate parsing requires some amount of **semantics** — understand-
 5355 ing the meaning of the text to be parsed. Consider the example *cats scratch people with*
 5356 *claws*: knowledge of about *cats*, *claws*, and scratching is necessary to correctly resolve
 5357 the attachment ambiguity.

5358 An extreme example is shown in Figure 10.4. The analysis on the left is preferred
 5359 because of the conjunction of similar entities *France* and *Italy*. But given the non-terminals
 5360 shown in the analyses, there is no way to differentiate these two parses, since they include
 5361 exactly the same productions. What is needed seems to be more precise non-terminals.
 5362 One possibility would be to rethink the linguistics behind the Penn Treebank, and ask

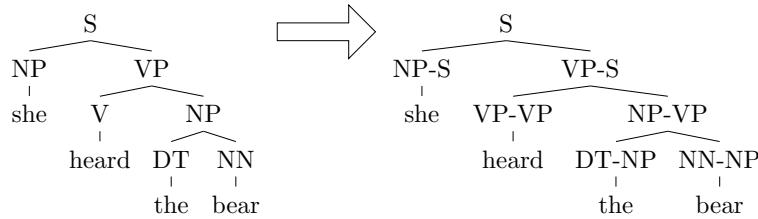


Figure 10.5: Parent annotation in a CFG derivation

5363 the annotators to try again. But the original annotation effort took five years, and there
 5364 is a little appetite for another annotation effort of this scope. Researchers have therefore
 5365 turned to automated techniques.

5366 10.5.1 Parent annotations and other tree transformations

The key assumption underlying context-free parsing is that productions depend only on the identity of the non-terminal on the left-hand side, and not on its ancestors or neighbors. The validity of this assumption is an empirical question, and it depends on the non-terminals themselves: ideally, every noun phrase (and verb phrase, etc) would be distributionally identical, so the assumption would hold. But in the Penn Treebank, the observed probability of productions often depends on the parent of the left-hand side. For example, noun phrases are more likely to be modified by prepositional phrases when they are in the object position (e.g., *they amused the students from Georgia*) than in the subject position (e.g., *the students from Georgia amused them*). This means that the $\text{NP} \rightarrow \text{NP PP}$ production is more likely if the entire constituent is the child of a VP than if it is the child of S. The observed statistics are (Johnson, 1998):

$$\Pr(\text{NP} \rightarrow \text{NP PP}) = 11\% \quad [10.35]$$

$$\Pr(\text{NP under S} \rightarrow \text{NP PP}) = 9\% \quad [10.36]$$

$$\Pr(\text{NP under VP} \rightarrow \text{NP PP}) = 23\%. \quad [10.37]$$

5367 This phenomenon can be captured by **parent annotation** (Johnson, 1998), in which each
 5368 non-terminal is augmented with the identity of its parent, as shown in Figure 10.5). This is
 5369 sometimes called **vertical Markovization**, since a Markov dependency is introduced be-
 5370 tween each node and its parent (Klein and Manning, 2003). It is analogous to moving from
 5371 a bigram to a trigram context in a hidden Markov model. In principle, parent annotation
 5372 squares the size of the set of non-terminals, which could make parsing considerably less
 5373 efficient. But in practice, the increase in the number of non-terminals that actually appear
 5374 in the data is relatively modest (Johnson, 1998).

5375 Parent annotation weakens the WCFG locality assumptions. This improves accuracy
 5376 by enabling the parser to make more fine-grained distinctions, which better capture real
 5377 linguistic phenomena. However, each production is more rare, and so careful smoothing
 5378 or regularization is required to control the variance over production scores.

5379 10.5.2 Lexicalized context-free grammars

5380 The examples in § 10.2.2 demonstrate the importance of individual words in resolving
 5381 parsing ambiguity: the preposition *on* is more likely to attach to *met*, while the preposition
 5382 *of* is more likely to attachment to *President*. But of all word pairs, which are relevant to
 5383 attachment decisions? Consider the following variants on the original examples:

- 5384 (10.4) We met the President of Mexico.
- 5385 (10.5) We met the first female President of Mexico.
- 5386 (10.6) They had supposedly met the President on Monday.

5387 The underlined words are the **head words** of their respective phrases: *met* heads the verb
 5388 phrase, and *President* heads the direct object noun phrase. These heads provide useful
 5389 semantic information. But they break the context-free assumption, which states that the
 5390 score for a production depends only on the parent and its immediate children, and not
 5391 the substructure under each child.

The incorporation of head words into context-free parsing is known as **lexicalization**,
 and is implemented in rules of the form,

$$\text{NP}(\text{President}) \rightarrow \text{NP}(\text{President}) \text{ PP}(of) \quad [10.38]$$

$$\text{NP}(\text{President}) \rightarrow \text{NP}(\text{President}) \text{ PP}(on). \quad [10.39]$$

5392 Lexicalization was a major step towards accurate PCFG parsing. It requires solving three
 5393 problems: identifying the heads of all constituents in a treebank; parsing efficiently while
 5394 keeping track of the heads; and estimating the scores for lexicalized productions.

5395 10.5.2.1 Identifying head words

5396 The head of a constituent is the word that is the most useful for determining how that
 5397 constituent is integrated into the rest of the sentence.⁵ The head word of a constituent is
 5398 determined recursively: for any non-terminal production, the head of the left-hand side
 5399 must be the head of one of the children. The head is typically selected according to a set of
 5400 deterministic rules, sometimes called **head percolation rules**. In many cases, these rules
 5401 are straightforward: the head of a noun phrase in a $\text{NP} \rightarrow \text{DET NN}$ production is the head

⁵This is a pragmatic definition, befitting our goal of using head words to improve parsing; for a more formal definition, see (Bender, 2013, chapter 7).

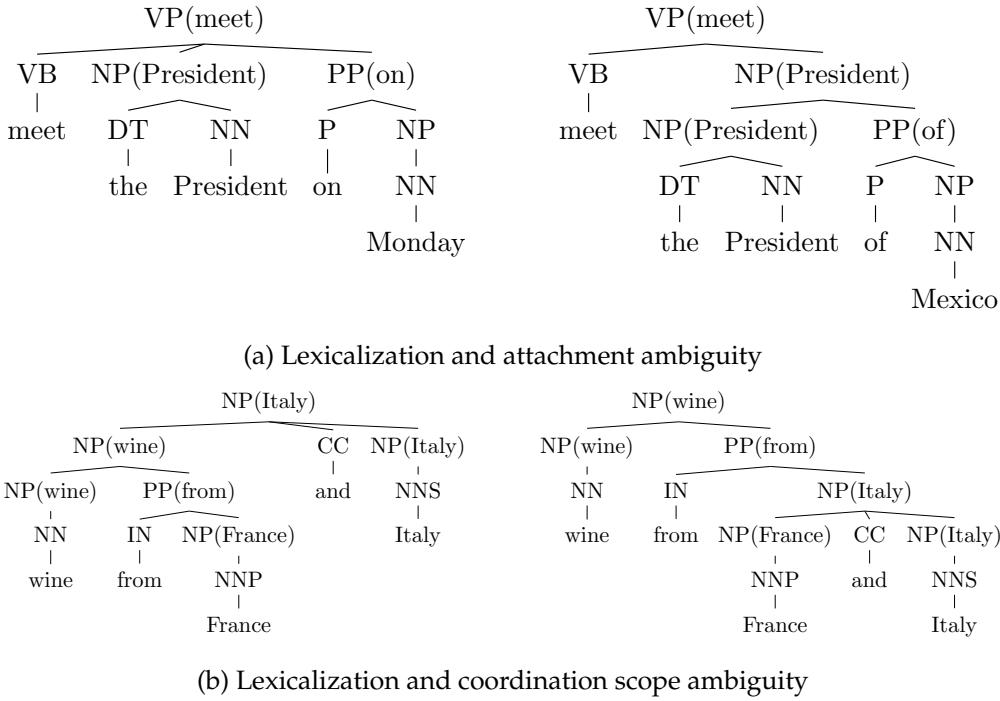


Figure 10.6: Examples of lexicalization

5402 of the noun; the head of a sentence in a $S \rightarrow NP\ VP$ production is the head of the verb
 5403 phrase.

5404 Table 10.3 shows a fragment of the head percolation rules used in many English pars-
 5405 ing systems. The meaning of the first rule is that to find the head of an S constituent, first
 5406 look for the rightmost VP child; if you don't find one, then look for the rightmost $SBAR$
 5407 child, and so on down the list. Verb phrases are headed by left verbs (the head of *can plan*
 5408 *on walking* is *planned*, since the modal verb *can* is tagged *MD*); noun phrases are headed by
 5409 the rightmost noun-like non-terminal (so the head of *the red cat* is *cat*),⁶ and prepositional
 5410 phrases are headed by the preposition (the head of *at Georgia Tech* is *at*). Some of these
 5411 rules are somewhat arbitrary — there's no particular reason why the head of *cats and dogs*
 5412 should be *dogs* — but the point here is just to get some lexical information that can support
 5413 parsing, not to make deep claims about syntax. Figure 10.6 shows the application of these
 5414 rules to two of the running examples.

⁶The noun phrase non-terminal is sometimes treated as a special case. Collins (1997) uses a heuristic that looks for the rightmost child which is a noun-like part-of-speech (e.g., *NN*, *NNP*), a possessive marker, or a superlative adjective (e.g., *the greatest*). If no such child is found, the heuristic then looks for the *leftmost* NP . If there is no child with tag NP , the heuristic then applies another priority list, this time from right to left.

Non-terminal	Direction	Priority
S	right	VP SBAR ADJP UCP NP
VP	left	VBD VBN MD VBZ TO VB VP VBG VBP ADJP NP
NP	right	N* EX \$ CD QP PRP ...
PP	left	IN TO FW

Table 10.3: A fragment of head percolation rules for English, from <http://www.cs.columbia.edu/~mcollins/papers/heads>

5415 10.5.2.2 Parsing lexicalized context-free grammars

5416 A naïve application of lexicalization would simply increase the set of non-terminals by
 5417 taking the cross-product with the set of terminal symbols, so that the non-terminals now
 5418 include symbols like $NP(President)$ and $VP(meet)$. Under this approach, the CKY parsing
 5419 algorithm could be applied directly to the lexicalized production rules. However, the
 5420 complexity would be cubic in the size of the vocabulary of terminal symbols, which would
 5421 clearly be intractable.

Another approach is to augment the CKY table with an additional index, keeping track of the head of each constituent. The cell $t[i, j, h, X]$ stores the score of the best derivation in which non-terminal X spans $w_{i+1:j}$ with head word h , where $i < h \leq j$. To compute such a table recursively, we must consider the possibility that each phrase gets its head from either its left or right child. The scores of the best derivations in which the head comes from the left and right child are denoted t_ℓ and t_r respectively, leading to the following recurrence:

$$t_\ell[i, j, h, X] = \max_{(X \rightarrow Y Z)} \max_{k > h} \max_{k < h' \leq j} t[i, k, h, Y] + t[k, j, h', Z] + \psi(X(h) \rightarrow Y(h)Z(h')) \quad [10.40]$$

$$t_r[i, j, h, X] = \max_{(X \rightarrow Y Z)} \max_{k < h} \max_{i < h' \leq k} t[i, k, h', Y] + t[k, j, h, Z] + (\psi(X(h) \rightarrow Y(h')Z(h))) \quad [10.41]$$

$$t[i, j, h, X] = \max(t_\ell[i, j, h, X], t_r[i, j, h, X]). \quad [10.42]$$

5422 To compute t_ℓ , we maximize over all split points $k > h$, since the head word must be in
 5423 the left child. We then maximize again over possible head words h' for the right child. An
 5424 analogous computation is performed for t_r . The size of the table is now $\mathcal{O}(M^3N)$, where
 5425 M is the length of the input and N is the number of non-terminals. Furthermore, each
 5426 cell is computed by performing $\mathcal{O}(M^2)$ operations, since we maximize over both the split
 5427 point k and the head h' . The time complexity of the algorithm is therefore $\mathcal{O}(RM^5N)$,
 5428 where R is the number of rules in the grammar. Fortunately, more efficient solutions are
 5429 possible. In general, the complexity of parsing can be reduced to $\mathcal{O}(M^4)$ in the length of

5430 the input; for a broad class of lexicalized CFGs, the complexity can be made cubic in the
 5431 length of the input, just as in unlexicalized CFGs (Eisner, 2000).

5432 **10.5.2.3 Estimating lexicalized context-free grammars**

5433 The final problem for lexicalized parsing is how to estimate weights for lexicalized pro-
 5434 ductions $X(i) \rightarrow Y(j) Z(k)$. These productions are said to be **bilexical**, because they
 5435 involve scores over pairs of words: in the example *meet the President of Mexico*, we hope
 5436 to choose the correct attachment point by modeling the bilexical affinities of (*meet, of*) and
 5437 (*President, of*). The number of such word pairs is quadratic in the size of the vocabulary,
 5438 making it difficult to estimate the weights of lexicalized production rules directly from
 5439 data. This is especially true for probabilistic context-free grammars, in which the weights
 5440 are obtained from smoothed relative frequency. In a treebank with a million tokens, a
 5441 vanishingly small fraction of the possible lexicalized productions will be observed more
 5442 than once.⁷ The Charniak (1997) and Collins (1997) parsers therefore focus on approxi-
 5443 mating the probabilities of lexicalized productions, using various smoothing techniques
 5444 and independence assumptions.

In discriminatively-trained weighted context-free grammars, the scores for each production can be computed from a set of features, which can be made progressively more fine-grained (Finkel et al., 2008). For example, the score of the lexicalized production $\text{NP}(\text{President}) \rightarrow \text{NP}(\text{President}) \text{ PP}(of)$ can be computed from the following features:

$$\begin{aligned} f(\text{NP}(\text{President}) \rightarrow \text{NP}(\text{President}) \text{ PP}(of)) = & \{\text{NP}(*) \rightarrow \text{NP}(*) \text{ PP}(*), \\ & \text{NP}(\text{President}) \rightarrow \text{NP}(\text{President}) \text{ PP}(*), \\ & \text{NP}(*) \rightarrow \text{NP}(*) \text{ PP}(of), \\ & \text{NP}(\text{President}) \rightarrow \text{NP}(\text{President}) \text{ PP}(of)\} \end{aligned}$$

5445 The first feature scores the unlexicalized production $\text{NP} \rightarrow \text{NP PP}$; the next two features
 5446 lexicalize only one element of the production, thereby scoring the appropriateness of NP
 5447 attachment for the individual words *President* and *of*; the final feature scores the specific
 5448 bilexical affinity of *President* and *of*. For bilexical pairs that are encountered frequently in
 5449 the treebank, this bilexical feature can play an important role in parsing; for pairs that are
 5450 absent or rare, regularization will drive its weight to zero, forcing the parser to rely on the
 5451 more coarse-grained features.

5452 In chapter 14, we will encounter techniques for clustering words based on their **distribu-**
 5453 **tional** properties — the contexts in which they appear. Such a clustering would group
 5454 rare and common words, such as *whale*, *shark*, *Leviathan*. Word clusters can be used

⁷The real situation is even more difficult, because non-binary context-free grammars can involve **trilexical** or higher-order dependencies, between the head of the constituent and multiple of its children (Carreras et al., 2008).

5455 as features in discriminative lexicalized parsing, striking a middle ground between full
 5456 lexicalization and non-terminals (Finkel et al., 2008). In this way, labeled examples con-
 5457 taining relatively common words like *whale* can help to improve parsing for rare words
 5458 like *beluga*, as long as those two words are clustered together.

5459 **10.5.3 *Refinement grammars**

5460 Lexicalization improves on context-free parsing by adding detailed information in the
 5461 form of lexical heads. However, estimating the scores of lexicalized productions is dif-
 5462 ficult. Klein and Manning (2003) argue that the right level of linguistic detail is some-
 5463 where between treebank categories and individual words. Some parts-of-speech and non-
 5464 terminals are truly substitutable: for example, *cat*/N and *dog*/N. But others are not: for
 5465 example, the preposition *of* exclusively attaches to nouns, while the preposition *as* is more
 5466 likely to modify verb phrases. Klein and Manning (2003) obtained a 2% improvement in
 5467 *F*-MEASURE on a parent-annotated PCFG parser by making a single change: splitting the
 5468 preposition category into six subtypes. They propose a series of linguistically-motivated
 5469 refinements to the Penn Treebank annotations, which in total yielded a 40% error reduc-
 5470 tion.

5471 Non-terminal refinement process can be automated by treating the refined categories
 5472 as latent variables. For example, we might split the noun phrase non-terminal into NP1, NP2, NP3, ...,
 5473 without defining in advance what each refined non-terminal corresponds to. This can
 5474 be treated as **partially supervised learning**, similar to the multi-component document
 5475 classification model described in § 5.2.3. A latent variable PCFG can be estimated by
 5476 expectation-maximization (Matsuzaki et al., 2005):

- 5477 • In the E-step, estimate a marginal distribution q over the refinement type of each
 5478 non-terminal in each derivation. These marginals are constrained by the original
 5479 annotation: an NP can be reannotated as NP4, but not as VP3. Marginal probabili-
 5480 ties over refined productions can be computed from the **inside-outside algorithm**,
 5481 as described in § 10.4.3, where the E-step enforces the constraints imposed by the
 5482 original annotations.
- 5483 • In the M-step, recompute the parameters of the grammar, by summing over the
 5484 probabilities of anchored productions that were computed in the E-step:

$$E[\text{count}(X \rightarrow Y Z)] = \sum_{i=0}^M \sum_{j=i}^M \sum_{k=i}^j p(X \rightarrow Y Z, (i, j, k) | \mathbf{w}). \quad [10.43]$$

5485 As usual, this process can be iterated to convergence. To determine the number of re-
 5486 finement types for each tag, Petrov et al. (2006) apply a split-merge heuristic; Liang et al.
 5487 (2007) and Finkel et al. (2007) apply **Bayesian nonparametrics** (Cohen, 2016).

Proper nouns			
NNP-14	<i>Oct.</i>	<i>Nov.</i>	<i>Sept.</i>
NNP-12	<i>John</i>	<i>Robert</i>	<i>James</i>
NNP-2	<i>J.</i>	<i>E.</i>	<i>L.</i>
NNP-1	<i>Bush</i>	<i>Noriega</i>	<i>Peters</i>
NNP-15	<i>New</i>	<i>San</i>	<i>Wall</i>
NNP-3	<i>York</i>	<i>Francisco</i>	<i>Street</i>
Personal Pronouns			
PRP-0	<i>It</i>	<i>He</i>	<i>I</i>
PRP-1	<i>it</i>	<i>he</i>	<i>they</i>
PRP-2	<i>it</i>	<i>them</i>	<i>him</i>

Table 10.4: Examples of automatically refined non-terminals and some of the words that they generate (Petrov et al., 2006).

5488 Some examples of refined non-terminals are shown in Table 10.4. The proper nouns
 5489 differentiate months, first names, middle initials, last names, first names of places, and
 5490 second names of places; each of these will tend to appear in different parts of grammatical
 5491 productions. The personal pronouns differentiate grammatical role, with PRP-0 appear-
 5492 ing in subject position at the beginning of the sentence (note the capitalization), PRP-1
 5493 appearing in subject position but not at the beginning of the sentence, and PRP-2 appear-
 5494 ing in object position.

5495 10.6 Beyond context-free parsing

5496 In the context-free setting, the score for a parse is a combination of the scores of individual
 5497 productions. As we have seen, these models can be improved by using finer-grained non-
 5498 terminals, via parent-annotation, lexicalization, and automated refinement. However, the
 5499 inherent limitations to the expressiveness of context-free parsing motivate the consider-
 5500 ation of other search strategies. These strategies abandon the optimality guaranteed by
 5501 bottom-up parsing, in exchange for the freedom to consider arbitrary properties of the
 5502 proposed parses.

5503 10.6.1 Reranking

5504 A simple way to relax the restrictions of context-free parsing is to perform a two-stage pro-
 5505 cess, in which a context-free parser generates a k -best list of candidates, and a **reranker**
 5506 then selects the best parse from this list (Charniak and Johnson, 2005; Collins and Koo,
 5507 2005). The reranker can be trained from an objective that is similar to multi-class classi-
 5508 fication: the goal is to learn weights that assign a high score to the reference parse, or to

5509 the parse on the k -best list that has the lowest error. In either case, the reranker need only
 5510 evaluate the K best parses, and so no context-free assumptions are necessary. This opens
 5511 the door to more expressive scoring functions:

- 5512 • It is possible to incorporate arbitrary non-local features, such as the structural par-
 5513 allelism and right-branching orientation of the parse (Charniak and Johnson, 2005).
 5514 • Reranking enables the use of **recursive neural networks**, in which each constituent
 5515 span $w_{i+1:j}$ receives a vector $\mathbf{u}_{i,j}$ which is computed from the vector representa-
 5516 tions of its children, using a composition function that is linked to the production
 5517 rule (Socher et al., 2013), e.g.,

$$\mathbf{u}_{i,j} = f \left(\Theta_{X \rightarrow Y} Z \begin{bmatrix} \mathbf{u}_{i,k} \\ \mathbf{u}_{k,j} \end{bmatrix} \right) \quad [10.44]$$

5518 The overall score of the parse can then be computed from the final vector, $\Psi(\tau) =$
 5519 $\theta \mathbf{u}_{0,M}$.

5520 Reranking can yield substantial improvements in accuracy. The main limitation is that it
 5521 can only find the best parse among the K -best offered by the generator, so it is inherently
 5522 limited by the ability of the bottom-up parser to find high-quality candidates.

5523 10.6.2 Transition-based parsing

5524 Structure prediction can be viewed as a form of search. An alternative to bottom-up pars-
 5525 ing is to read the input from left-to-right, gradually building up a parse structure through
 5526 a series of **transitions**. Transition-based parsing is described in more detail in the next
 5527 chapter, in the context of dependency parsing. However, it can also be applied to CFG
 5528 parsing, as briefly described here.

5529 For any context-free grammar, there is an equivalent **pushdown automaton**, a model
 5530 of computation that accepts exactly those strings that can be derived from the grammar.
 5531 This computational model consumes the input from left to right, while pushing and pop-
 5532 ping elements on a stack. This architecture provides a natural transition-based parsing
 5533 framework for context-free grammars, known as **shift-reduce parsing**.

5534 Shift-reduce parsing is a type of transition-based parsing, in which the parser can take
 5535 the following actions:

- 5536 • *shift* the next terminal symbol onto the stack;
 5537 • *unary-reduce* the top item on the stack, using a unary production rule in the gram-
 5538 mar;
 5539 • *binary-reduce* the top two items onto the stack, using a binary production rule in the
 5540 grammar.

5541 The set of available actions is constrained by the situation: the parser can only shift if
 5542 there are remaining terminal symbols in the input, and it can only reduce if an applicable
 5543 production rule exists in the grammar. If the parser arrives at a state where the input
 5544 has been completely consumed, and the stack contains only the element S, then the input
 5545 is accepted. If the parser arrives at a non-accepting state where there are no possible
 5546 actions, the input is rejected. A parse error occurs if there is some action sequence that
 5547 would accept an input, but the parser does not find it.

5548 **Example** Consider the input *we eat sushi* and the grammar in Table 10.1. The input can
 5549 be parsed through the following sequence of actions:

- 5550 1. **Shift** the first token *we* onto the stack.
- 5551 2. **Reduce** the top item on the stack to NP, using the production $NP \rightarrow we$.
- 5552 3. **Shift** the next token *eat* onto the stack, and **reduce** it to V with the production $V \rightarrow eat$.
- 5554 4. **Shift** the final token *sushi* onto the stack, and **reduce** it to NP. The input has been
 5555 completely consumed, and the stack contains [NP, V, NP].
- 5556 5. **Reduce** the top two items using the production $VP \rightarrow V NP$. The stack now con-
 5557 tains [VP, NP].
- 5558 6. **Reduce** the top two items using the production $S \rightarrow NP VP$. The stack now contains
 5559 [S]. Since the input is empty, this is an accepting state.

5560 One thing to notice from this example is that the number of shift actions is equal to the length of the input. The number of reduce actions is equal to the number of non-terminals in the analysis, which grows linearly in the length of the input. Thus, the overall time complexity of shift-reduce parsing is linear in the length of the input (assuming the complexity of each individual classification decision is constant in the length of the input).
 5565 This is far better than the cubic time complexity required by CKY parsing.

5566 **Transition-based parsing as inference** In general, it is not possible to guarantee that a transition-based parser will find the optimal parse, $\text{argmax}_\tau \Psi(\tau; \mathbf{w})$, even under the usual CFG independence assumptions. We could assign a score to each anchored parsing action in each context, with $\psi(a, c)$ indicating the score of performing action a in context c .
 5567 One might imagine that transition-based parsing could efficiently find the derivation that maximizes the sum of such scores. But this too would require backtracking and searching over an exponentially large number of possible action sequences: if a bad decision is made at the beginning of the derivation, then it may be impossible to recover the optimal action sequence without backtracking to that early mistake. This is known as a **search error**. Transition-based parsers can incorporate arbitrary features, without the restrictive

5576 independence assumptions required by chart parsing; search errors are the price that must
 5577 be paid for this flexibility.

5578 **Learning transition-based parsing** Transition-based parsing can be combined with ma-
 5579 chine learning by training a classifier to select the correct action in each situation. This
 5580 classifier is free to choose any feature of the input, the state of the parser, and the parse
 5581 history. However, there is no optimality guarantee: the parser may choose a suboptimal
 5582 parse, due to a mistake at the beginning of the analysis. Nonetheless, some of the strongest
 5583 CFG parsers are based on the shift-reduce architecture, rather than CKY. A recent genera-
 5584 tion of models links shift-reduce parsing with recurrent neural networks, updating a
 5585 hidden state vector while consuming the input (e.g., Cross and Huang, 2016; Dyer et al.,
 5586 2016). Learning algorithms for transition-based parsing are discussed in more detail in
 5587 § 11.3.

5588 Exercises

1. Consider the following PCFG:

$$p(X \rightarrow X X) = \frac{1}{2} \quad [10.45]$$

$$p(X \rightarrow Y) = \frac{1}{2} \quad [10.46]$$

$$p(Y \rightarrow \sigma) = \frac{1}{|\Sigma|}, \forall \sigma \in \Sigma \quad [10.47]$$

5589 a) Compute the probability $p(\hat{\tau})$ of the maximum probability parse for a string
 5590 $w \in \Sigma^M$.

5591 b) Compute the marginal probability $p(w) = \sum_{\tau: \text{yield}(\tau)=w} p(\tau)$.

5592 c) Compute the conditional probability $p(\hat{\tau} | w)$.

2. Use the inside and outside scores to compute the marginal probability $p(X_{i:j} \rightarrow Y_{i:k-1} Z_{k:j} | w)$, indicating that Y spans $w_{i:k-1}$, Z spans $w_{k:j}$, and X is the parent of Y and Z , spanning $w_{i:j}$.
3. Suppose that the potentials $\Psi(X \rightarrow \alpha)$ are log-probabilities, so that $\sum_{\alpha} \exp \Psi(X \rightarrow \alpha) = 1$ for all X . Verify that the semiring inside recurrence from Equation 10.26 generates the log-probability $\log p(w) = \log \sum_{\tau: \text{yield}(\tau)=w} p(\tau)$.
4. more exercises tk

5600 Chapter 11

5601 Dependency parsing

5602 The previous chapter discussed algorithms for analyzing sentences in terms of nested con-
5603 stituents, such as noun phrases and verb phrases. However, many of the key sources of
5604 ambiguity in phrase-structure analysis relate to questions of **attachment**: where to attach a
5605 prepositional phrase or complement clause, how to scope a coordinating conjunction, and
5606 so on. These attachment decisions can be represented with a more lightweight structure:
5607 a directed graph over the words in the sentence, known as a **dependency parse**. Syntac-
5608 tic annotation has shifted its focus to such dependency structures: at the time of this
5609 writing, the **Universal Dependencies** project offers more than 100 dependency treebanks
5610 for more than 60 languages.¹ This chapter will describe the linguistic ideas underlying
5611 dependency grammar, and then discuss exact and transition-based parsing algorithms.
5612 The chapter will also discuss recent research on **learning to search** in transition-based
5613 structure prediction.

5614 11.1 Dependency grammar

5615 While **dependency grammar** has a rich history of its own (Tesnière, 1966; Kübler et al.,
5616 2009), it can be motivated by extension from the lexicalized context-free grammars that
5617 we encountered in previous chapter (§ 10.5.2). Recall that lexicalization augments each
5618 non-terminal with a **head word**. The head of a constituent is identified recursively, using
5619 a set of **head rules**, as shown in Table 10.3. An example of a lexicalized context-free parse
5620 is shown in Figure 11.1a. In this sentence, the head of the S constituent is the main verb,
5621 *scratch*; this non-terminal then produces the noun phrase *the cats*, whose head word is
5622 *cats*, and from which we finally derive the word *the*. Thus, the word *scratch* occupies the
5623 central position for the sentence, with the word *cats* playing a supporting role. In turn, *cats*

¹universaldependencies.org

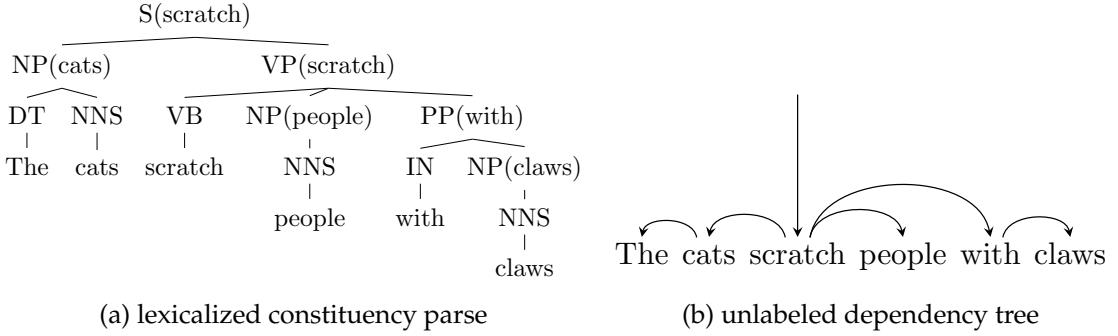


Figure 11.1: Dependency grammar is closely linked to lexicalized context free grammars: each lexical head has a dependency path to every other word in the constituent. (This example is based on the lexicalization rules from § 10.5.2, which make the preposition the head of a prepositional phrase. In the more contemporary Universal Dependencies annotations, the head of *with claws* would be *claws*, so there would be an edge *scratch* → *claws*.)

5624 occupies the central position for the noun phrase, with the word *the* playing a supporting
5625 role.

5626 The relationships between words in a sentence can be formalized in a directed graph,
5627 based on the lexicalized phrase-structure parse: create an edge (i, j) iff word i is the head
5628 of a phrase whose child is a phrase headed by word j . Thus, in our example, we would
5629 have *scratch* → *cats* and *cats* → *the*. We would not have the edge *scratch* → *the*, because
5630 although $S(\text{scratch})$ dominates $\text{DET}(\text{the})$ in the phrase-structure parse tree, it is not its im-
5631 mediate parent. These edges describe **syntactic dependencies**, a bilexical relationship
5632 between a **head** and a **dependent**, which is at the heart of dependency grammar.

5633 Continuing to build out this **dependency graph**, we will eventually reach every word
5634 in the sentence, as shown in Figure 11.1b. In this graph — and in all graphs constructed
5635 in this way — every word has exactly one incoming edge, except for the root word, which
5636 is indicated by a special incoming arrow from above. Furthermore, the graph is *weakly*
5637 *connected*: if the directed edges were replaced with undirected edges, there would be a
5638 path between all pairs of nodes. From these properties, it can be shown that there are no
5639 cycles in the graph (or else at least one node would have to have more than one incoming
5640 edge), and therefore, the graph is a tree. Because the graph includes all vertices, it is a
5641 **spanning tree**.

5642 11.1.1 Heads and dependents

5643 A dependency edge implies an asymmetric syntactic relationship between the head and
5644 dependent words, sometimes called **modifiers**. For a pair like *the cats* or *cats scratch*, how

5645 do we decide which is the head? Here are some possible criteria:

- 5646 • The head sets the syntactic category of the construction: for example, nouns are the
5647 heads of noun phrases, and verbs are the heads of verb phrases.
- 5648 • The modifier may be optional while the head is mandatory: for example, in the
5649 sentence *cats scratch people with claws*, the subtrees *cats scratch* and *cats scratch people*
5650 are grammatical sentences, but *with claws* is not.
- 5651 • The head determines the morphological form of the modifier: for example, in lan-
5652 guages that require gender agreement, the gender of the noun determines the gen-
5653 der of the adjectives and determiners.
- 5654 • Edges should first connect content words, and then connect function words.

5655 As always, these guidelines sometimes conflict. The Universal Dependencies (UD)
5656 project has attempted to identify a set of principles that can be applied to dozens of dif-
5657 ferent languages (Nivre et al., 2016).² These guidelines are based on the universal part-
5658 of-speech tags from chapter 8. They differ somewhat from the head rules described in
5659 § 10.5.2: for example, on the principle that dependencies should relate content words, the
5660 prepositional phrase *with claws* would be headed by *claws*, resulting in an edge *scratch* →
5661 *claws*, and another edge *claws* → *with*.

5662 One objection to dependency grammar is that not all syntactic relations are asymmet-
5663 ric. Coordination is one of the most obvious examples (Popel et al., 2013): in the sentence,
5664 *Abigail and Max like kimchi* (Figure 11.2), which word is the head of the coordinated noun
5665 phrase *Abigail and Max*? Choosing either *Abigail* or *Max* seems arbitrary; fairness argues
5666 for making *and* the head, but this seems like the least important word in the noun phrase,
5667 and selecting it would violate the principle of linking content words first. The Universal
5668 Dependencies annotation system arbitrarily chooses the left-most item as the head — in
5669 this case, *Abigail* — and includes edges from this head to both *Max* and the coordinating
5670 conjunction *and*. These edges are distinguished by the labels CONJ (for the thing begin
5671 conjoined) and CC (for the coordinating conjunction). The labeling system is discussed
5672 next.

5673 11.1.2 Labeled dependencies

5674 Edges may be **labeled** to indicate the nature of the syntactic relation that holds between
5675 the two elements. For example, in Figure 11.2, the label NSUBJ on the edge from *like* to
5676 *Abigail* indicates that the subtree headed by *Abigail* is the noun subject of the verb *like*;
5677 similarly, the label OBJ on the edge from *like* to *kimchi* indicates that the subtree headed by

²The latest and most specific guidelines are available at universaldependencies.org/guidelines.html

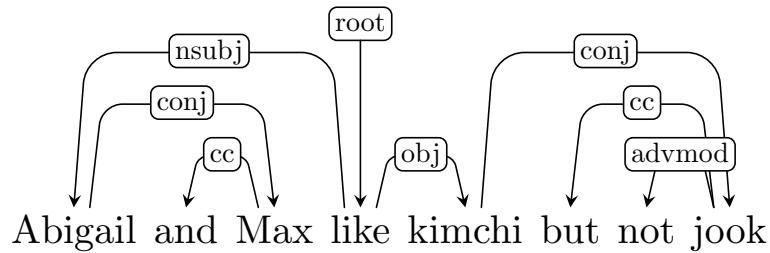


Figure 11.2: In the Universal Dependencies annotation system, the left-most item of a coordination is the head.

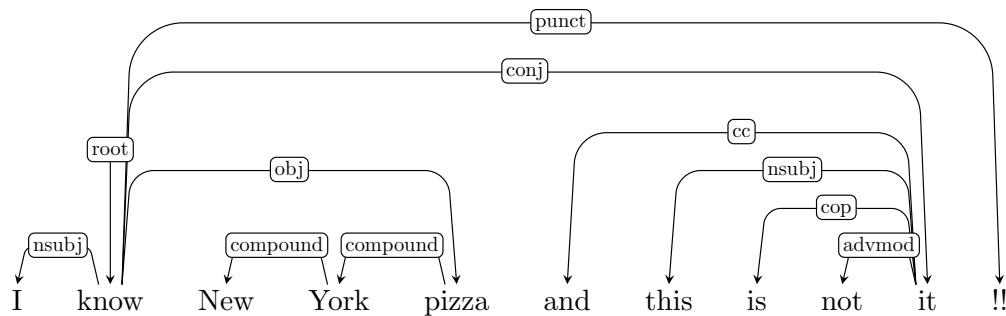


Figure 11.3: A labeled dependency parse from the English UD Treebank (reviews-361348-0006)

5678 *kimchi* is the object.³ The negation *not* is treated as an adverbial modifier (ADVMOD) on
5679 the noun *jook*.

5680 A slightly more complex example is shown in Figure 11.3. The multiword expression
5681 *New York pizza* is treated as a “flat” unit of text, with the elements linked by the COM-
5682 POUND relation. The sentence includes two clauses that are conjoined in the same way
5683 that noun phrases are conjoined in Figure 11.2. The second clause contains a **copula** verb
5684 (see § 8.1.1). For such clauses, we treat the “object” of the verb as the root — in this case,
5685 *it* — and label the verb as a dependent, with the COP relation. This example also shows
5686 how punctuation are treated, with label PUNCT.

5687 11.1.3 Dependency subtrees and constituents

5688 Dependency trees hide information that would be present in a CFG parse. Often what
5689 is hidden is in fact irrelevant: for example, Figure 11.4 shows three different ways of

³Earlier work distinguished direct and indirect objects (De Marneffe and Manning, 2008), but this has been dropped in version 2.0 of the Universal Dependencies annotation system.

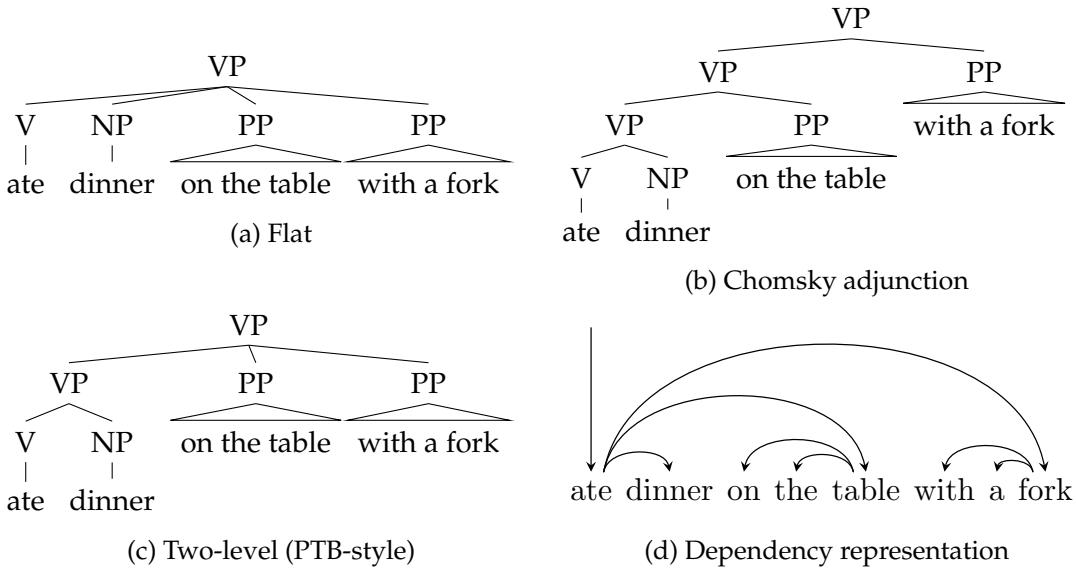


Figure 11.4: The three different CFG analyses of this verb phrase all correspond to a single dependency structure.

representing prepositional phrase adjuncts to the verb *ate*. Because there is apparently no meaningful difference between these analyses, the Penn Treebank decides by convention to use the two-level representation (see Johnson, 1998, for a discussion). As shown in Figure 11.4d, these three cases all look the same in a dependency parse.

But dependency grammar imposes its own set of annotation decisions, such as the identification of the head of a coordination (§ 11.1.1); without lexicalization, context-free grammar does not require either element in a coordination to be privileged in this way. Dependency parses can be disappointingly flat: for example, in the sentence *Yesterday, Abigail was reluctantly giving Max kimchi*, the root *giving* is the head of every dependency! The constituent parse arguably offers a more useful structural analysis for such cases.

Projectivity Thus far, we have defined dependency trees as spanning trees over a graph in which each word is a vertex. As we have seen, one way to construct such trees is by connecting the heads in a lexicalized constituent parse. However, there are spanning trees that cannot be constructed in this way. Syntactic constituents are *contiguous spans*. In a spanning tree constructed from a lexicalized constituent parse, the head h of any constituent that spans the nodes from i to j must have a path to every node in this span. This property is known as **projectivity**, and projective dependency parses are a restricted class of spanning trees. Informally, projectivity means that “crossing edges” are prohibited. The formal definition follows:

	% non-projective edges	% non-projective sentences
Czech	1.86%	22.42%
English	0.39%	7.63%
German	2.33%	28.19%

Table 11.1: Frequency of non-projective dependencies in three languages (Kuhlmann and Nivre, 2010)

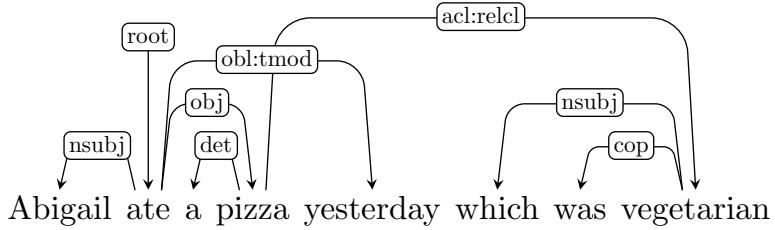


Figure 11.5: An example of a non-projective dependency parse. The “crossing edge” arises from the relative clause *which was vegetarian* and the oblique temporal modifier *yesterday*.

5709 **Definition 2** (Projectivity). *An edge from i to j is projective iff all k between i and j are descendants of i . A dependency parse is projective iff all its edges are projective.*

5711 Figure 11.5 gives an example of a non-projective dependency graph in English. This
 5712 dependency graph does not correspond to any constituent parse. As shown in Table 11.1,
 5713 non-projectivity is more common in languages such as Czech and German. Even though
 5714 relatively few dependencies are non-projective in these languages, many sentences have
 5715 at least one such dependency. As we will soon see, projectivity has important algorithmic
 5716 consequences.

5717 11.2 Graph-based dependency parsing

5718 Let $\mathbf{y} = \{i \xrightarrow{r} j\}$ represent a dependency graph, in which each edge is a relation r from
 5719 head word $i \in \{1, 2, \dots, M, \text{ROOT}\}$ to modifier $j \in \{1, 2, \dots, M\}$. The special node ROOT
 5720 indicates the root of the graph, and M is the length of the input $|\mathbf{w}|$. Given a scoring
 5721 function $\Psi(\mathbf{y}, \mathbf{w}; \theta)$, the optimal parse is,

$$\hat{\mathbf{y}} = \underset{\mathbf{y} \in \mathcal{Y}(\mathbf{w})}{\operatorname{argmax}} \Psi(\mathbf{y}, \mathbf{w}; \theta), \quad [11.1]$$

5722 where $\mathcal{Y}(\mathbf{w})$ is the set of valid dependency parses on the input \mathbf{w} . As usual, the number
 5723 of possible labels $|\mathcal{Y}(\mathbf{w})|$ is exponential in the length of the input (Wu and Chao, 2004).

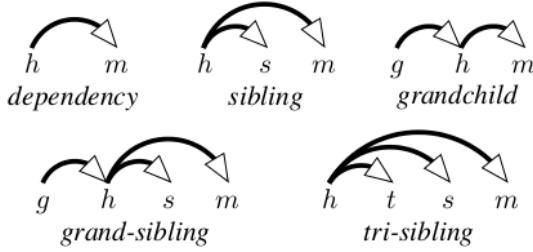


Figure 11.6: Feature templates for higher-order dependency parsing (Koo and Collins, 2010) [todo: permission]

5724 Algorithms that search over this space of possible graphs are known as **graph-based de-**
5725 **pendency parsers.**

In sequence labeling and constituent parsing, it was possible to search efficiently over an exponential space by choosing a feature function that decomposes into a sum of local feature vectors. A similar approach is possible for dependency parsing, by requiring the scoring function to decompose across dependency arcs $i \rightarrow j$:

$$\Psi(\mathbf{y}, \mathbf{w}; \boldsymbol{\theta}) = \sum_{i \xrightarrow{r} j \in \mathbf{y}} \psi(i \xrightarrow{r} j, \mathbf{w}; \boldsymbol{\theta}). \quad [11.2]$$

5726 Dependency parsers that operate under this assumption are known as **arc-factored**, since
5727 the overall score is a product of scores over all arcs.

Higher-order dependency parsing The arc-factored decomposition can be relaxed to allow higher-order dependencies. In **second-order dependency parsing**, the scoring function may include grandparents and siblings, as shown by the templates in Figure 11.6. The scoring function is,

$$\begin{aligned} \Psi(\mathbf{y}, \mathbf{w}; \boldsymbol{\theta}) = & \sum_{i \xrightarrow{r} j \in \mathbf{y}} \sum_{k \xrightarrow{r'} i \in \mathbf{y}} \psi_{\text{grandparent}}(i \xrightarrow{r} j, k, r', \mathbf{w}; \boldsymbol{\theta}) \\ & \sum_{\substack{i \xrightarrow{r'} s \in \mathbf{y} \\ s \neq j}} \psi_{\text{sibling}}(i \xrightarrow{r} j, s, r', \mathbf{w}; \boldsymbol{\theta}). \end{aligned} \quad [11.3]$$

5728 The top line scores computes a scoring function that includes the grandparent *k*; the
5729 bottom line computes a scoring function for each sibling *s*. For projective dependency
5730 graphs, there are efficient algorithms for second-order and third-order dependency pars-
5731 ing (Eisner, 1996; McDonald and Pereira, 2006; Koo and Collins, 2010); for non-projective
5732 dependency graphs, second-order dependency parsing is NP-hard (McDonald and Pereira,
5733 2006). The specific algorithms are discussed in the next section.

5734 **11.2.1 Graph-based parsing algorithms**

5735 The distinction between projective and non-projective dependency trees (§ 11.1.3) plays
 5736 a key role in the choice of algorithms. Because projective dependency trees are closely
 5737 related to (and can be derived from) lexicalized constituent trees, lexicalized parsing al-
 5738 gorithms can be applied directly. For the more general problem of parsing to arbitrary
 5739 spanning trees, a different class of algorithms is required. In both cases, arc-factored de-
 5740 pendency parsing relies on precomputing the scores $\psi(i \xrightarrow{r} j, w; \theta)$ for each potential
 5741 edge. There are $\mathcal{O}(M^2 R)$ such scores, where M is the length of the input and R is the
 5742 number of dependency relation types, and this is a lower bound on the time and space
 5743 complexity of any exact algorithm for arc-factored dependency parsing.

5744 **11.2.1.1 Projective dependency parsing**

5745 Any lexicalized constituency tree can be converted into a projective dependency tree by
 5746 creating arcs between the heads of constituents and their parents, so any algorithm for
 5747 lexicalized constituent parsing can be converted into an algorithm for projective depen-
 5748 dency parsing, by converting arc scores into scores for lexicalized productions. As noted
 5749 in § 10.5.2, there are cubic time algorithms for lexicalized constituent parsing, which are
 5750 extensions of the CKY algorithm. Therefore, arc-factored projective dependency parsing
 5751 can be performed in cubic time in the length of the input.

5752 Second-order projective dependency parsing can also be performed in cubic time, with
 5753 minimal modifications to the lexicalized parsing algorithm (Eisner, 1996). It is possible to
 5754 go even further, to **third-order dependency parsing**, in which the scoring function may
 5755 consider great-grandparents, grand-siblings, and “tri-siblings”, as shown in Figure 11.6.
 5756 Third-order dependency parsing can be performed in $\mathcal{O}(M^4)$ time, which can be made
 5757 practical through the use of pruning to eliminate unlikely edges (Koo and Collins, 2010).

5758 **11.2.1.2 Non-projective dependency parsing**

5759 In non-projective dependency parsing, the goal is to identify the highest-scoring span-
 5760 ning tree over the words in the sentence. The arc-factored assumption ensures that the
 5761 score for each spanning tree will be computed as a sum over scores for the edges, which
 5762 are precomputed. Based on these scores, we build a weighted connected graph. Arc-
 5763 factored non-projective dependency parsing is then equivalent to finding the spanning
 5764 tree that achieves the maximum total score, $\Psi(y, w) = \sum_{i \xrightarrow{r} j \in y} \psi(i \xrightarrow{r} j, w)$. The **Chu-**
 5765 **Liu-Edmonds algorithm** (Chu and Liu, 1965; Edmonds, 1967) computes this **maximum**
 5766 **spanning tree** efficiently. It does this by first identifying the best incoming edge $i \xrightarrow{r} j$ for
 5767 each vertex j . If the resulting graph does not contain cycles, it is the maximum spanning
 5768 tree. If there is a cycle, it is collapsed into a super-vertex, whose incoming and outgoing
 5769 edges are based on the edges to the vertices in the cycle. The algorithm is then applied

5770 recursively to the resulting graph, and process repeats until a graph without cycles is
 5771 obtained.

5772 The time complexity of identifying the best incoming edge for each vertex is $\mathcal{O}(M^2R)$,
 5773 where M is the length of the input and R is the number of relations; in the worst case, the
 5774 number of cycles is $\mathcal{O}(M)$. Therefore, the complexity of the Chu-Liu-Edmonds algorithm
 5775 is $\mathcal{O}(M^3R)$. This complexity can be reduced to $\mathcal{O}(M^2N)$ by storing the edge scores in a
 5776 Fibonacci heap (Gabow et al., 1986). For more detail on graph-based parsing algorithms,
 5777 see Eisner (1997) and Kübler et al. (2009).

5778 **Higher-order non-projective dependency parsing** Given the tractability of higher-order
 5779 projective dependency parsing, you may be surprised to learn that non-projective second-
 5780 order dependency parsing is NP-Hard. This can be proved by reduction from the vertex
 5781 cover problem (Neuhaus and Bröker, 1997). A heuristic solution is to do projective pars-
 5782 ing first, and then post-process the projective dependency parse to add non-projective
 5783 edges (Nivre and Nilsson, 2005). More recent work has applied techniques for approxi-
 5784 mate inference in graphical models, including belief propagation (Smith and Eisner, 2008),
 5785 integer linear programming (Martins et al., 2009), variational inference (Martins et al.,
 5786 2010), and Markov Chain Monte Carlo (Zhang et al., 2014).

5787 11.2.2 Computing scores for dependency arcs

The arc-factored scoring function $\psi(i \xrightarrow{r} j, \mathbf{w}; \boldsymbol{\theta})$ can be defined in several ways:

$$\text{Linear} \quad \psi(i \xrightarrow{r} j, \mathbf{w}; \boldsymbol{\theta}) = \boldsymbol{\theta} \cdot \mathbf{f}(i \xrightarrow{r} j, \mathbf{w}) \quad [11.4]$$

$$\text{Neural} \quad \psi(i \xrightarrow{r} j, \mathbf{w}; \boldsymbol{\theta}) = \text{Feedforward}([\mathbf{u}_{w_i}; \mathbf{u}_{w_j}]; \boldsymbol{\theta}) \quad [11.5]$$

$$\text{Generative} \quad \psi(i \xrightarrow{r} j, \mathbf{w}; \boldsymbol{\theta}) = \log p(w_j, r | w_i). \quad [11.6]$$

5788 11.2.2.1 Linear feature-based arc scores

5789 Linear models for dependency parsing incorporate many of the same features used in
 5790 sequence labeling and discriminative constituent parsing. These include:

- 5791 • the length and direction of the arc;
- 5792 • the words w_i and w_j linked by the dependency relation;
- 5793 • the prefixes, suffixes, and parts-of-speech of these words;
- 5794 • the neighbors of the dependency arc, $w_{i-1}, w_{i+1}, w_{j-1}, w_{j+1}$;
- 5795 • the prefixes, suffixes, and part-of-speech of these neighbor words.

5796 Each of these features can be conjoined with the dependency edge label r . Note that
 5797 features in an arc-factored parser can refer to words other than w_i and w_j . The restriction
 5798 is that the features consider only a single arc.

Bilexical features (e.g., *sushi* → *chopsticks*) are powerful but rare, so it is useful to augment them with coarse-grained alternatives, by “backing off” to the part-of-speech or affix. For example, the following features are created by backing off to part-of-speech tags in an unlabeled dependency parser:

$$\begin{aligned} f(3 \rightarrow 5, \text{we eat sushi with chopsticks}) = & \langle \text{sushi} \rightarrow \text{chopsticks}, \\ & \text{sushi} \rightarrow \text{NNS}, \\ & \text{NN} \rightarrow \text{chopsticks}, \\ & \text{NNS} \rightarrow \text{NN} \rangle. \end{aligned}$$

5799 Regularized discriminative learning algorithms can then trade off between features at
 5800 varying levels of detail. McDonald et al. (2005) take this approach as far as *tetralexical*
 5801 features (e.g., $(w_i, w_{i+1}, w_{j-1}, w_j)$). Such features help to avoid choosing arcs that are un-
 5802 likely due to the intervening words: for example, there is unlikely to be an edge between
 5803 two nouns if the intervening span contains a verb. A large list of first and second-order
 5804 features is provided by Bohnet (2010), who uses a hashing function to store these features
 5805 efficiently.

5806 11.2.2.2 Neural arc scores

Given vector representations \mathbf{x}_i for each word w_i in the input, a set of arc scores can be computed from a feedforward neural network:

$$\psi(i \xrightarrow{r} j, \mathbf{w}; \boldsymbol{\theta}) = \text{FeedForward}([\mathbf{x}_i; \mathbf{x}_j]; \boldsymbol{\theta}_r), \quad [11.7]$$

where unique weights $\boldsymbol{\theta}_r$ are available for each arc type (Pei et al., 2015; Kiperwasser and Goldberg, 2016). Kiperwasser and Goldberg (2016) use a feedforward network with a single hidden layer,

$$\mathbf{z} = g(\boldsymbol{\Theta}_r[\mathbf{x}_i; \mathbf{x}_j] + b_r^{(z)}) \quad [11.8]$$

$$\psi(i \xrightarrow{r} j) = \boldsymbol{\beta}_r \mathbf{z} + b_r^{(y)}, \quad [11.9]$$

5807 where $\boldsymbol{\Theta}_r$ is a matrix, $\boldsymbol{\beta}_r$ is a vector, each b_r is a scalar, and the function g is an elementwise
 5808 tanh activation function.

5809 The vector \mathbf{x}_i can be set equal to the word embedding, which may be pre-trained or
 5810 learned by backpropagation (Pei et al., 2015). Alternatively, contextual information can
 5811 be incorporated by applying a bidirectional recurrent neural network across the input, as
 5812 described in § 7.6. The RNN hidden states at each word can be used as inputs to the arc
 5813 scoring function (Kiperwasser and Goldberg, 2016).

5814 **11.2.2.3 Probabilistic arc scores**

If each arc score is equal to the log probability $\log p(w_j, r \mid w_i)$, then the sum of scores gives the log probability of the sentence and arc labels, by the chain rule. For example, consider the unlabeled parse of *we eat sushi with rice*,

$$\mathbf{y} = \{(ROOT, 2), (2, 1), (2, 3), (3, 5), (5, 4)\} \quad [11.10]$$

$$\log p(\mathbf{w} \mid \mathbf{y}) = \sum_{(i \rightarrow j) \in \mathbf{y}} \log p(w_j \mid w_i) \quad [11.11]$$

$$\begin{aligned} &= \log p(eat \mid ROOT) + \log p(we \mid eat) + \log p(sushi \mid eat) \\ &\quad + \log p(rice \mid sushi) + \log p(with \mid rice). \end{aligned} \quad [11.12]$$

5815 Probabilistic generative models are used in combination with expectation-maximization
5816 (chapter 5) for unsupervised dependency parsing (Klein and Manning, 2004).

5817 **11.2.3 Learning**

Having formulated graph-based dependency parsing as a structure prediction problem, we can apply similar learning algorithms to those used in sequence labeling. Given a loss function $\ell(\boldsymbol{\theta}; \mathbf{w}^{(i)}, \mathbf{y}^{(i)})$, we can compute gradient-based updates to the parameters. For a model with feature-based arc scores and a perceptron loss, we obtain the usual structured perceptron update,

$$\hat{\mathbf{y}} = \operatorname{argmax}_{\mathbf{y}' \in \mathcal{Y}(\mathbf{w})} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{w}, \mathbf{y}') \quad [11.13]$$

$$\boldsymbol{\theta} = \boldsymbol{\theta} + \mathbf{f}(\mathbf{w}, \mathbf{y}) - \mathbf{f}(\mathbf{w}, \hat{\mathbf{y}}) \quad [11.14]$$

5818 In this case, the argmax requires a maximization over all dependency trees for the sen-
5819 tence, which can be computed using the algorithms described in § 11.2.1. We can apply
5820 all the usual tricks from § 2.2: weight averaging, a large margin objective, and regular-
5821 ization. McDonald et al. (2005) were the first to treat dependency parsing as a structure
5822 prediction problem, using MIRA, an online margin-based learning algorithm. Neural arc
5823 scores can be learned in the same way, backpropagating from a margin loss to updates on
5824 the feedforward network that computes the score for each edge.

A conditional random field for arc-factored dependency parsing is built on the probability model,

$$p(\mathbf{y} \mid \mathbf{w}) = \frac{\exp \sum_{i \xrightarrow{r} j \in \mathbf{y}} \psi(i \xrightarrow{r} j, \mathbf{w}; \boldsymbol{\theta})}{\sum_{\mathbf{y}' \in \mathcal{Y}(\mathbf{w})} \exp \sum_{i \xrightarrow{r} j \in \mathbf{y}'} \psi(i \xrightarrow{r} j, \mathbf{w}; \boldsymbol{\theta})} \quad [11.15]$$

5825 Such a model is trained to minimize the negative log conditional-likelihood. Just as in
5826 CRF sequence models (§ 7.5.3) and the logistic regression classifier (§ 2.4), the gradients

5827 involve marginal probabilities $p(i \xrightarrow{r} j \mid \mathbf{w}; \theta)$, which in this case are probabilities over
 5828 individual dependencies. In arc-factored models, these probabilities can be computed
 5829 in polynomial time. For projective dependency trees, the marginal probabilities can be
 5830 computed in cubic time, using a variant of the inside-outside algorithm (Lari and Young,
 5831 1990). For non-projective dependency parsing, marginals can also be computed in cubic
 5832 time, using the **matrix-tree theorem** (Koo et al., 2007; McDonald et al., 2007; Smith and
 5833 Smith, 2007). Details of these methods are described by Kübler et al. (2009).

5834 11.3 Transition-based dependency parsing

5835 Graph-based dependency parsing offers exact inference, meaning that it is possible to re-
 5836 cover the best-scoring parse for any given model. But this comes at a price: the scoring
 5837 function is required to decompose into local parts — in the case of non-projective parsing,
 5838 these parts are restricted to individual arcs. These limitations are felt more keenly in de-
 5839 pendency parsing than in sequence labeling, because second-order dependency features
 5840 are critical to correctly identify some types of attachments. For example, prepositional
 5841 phrase attachment depends on the attachment point, the object of the preposition, and
 5842 the preposition itself; arc-factored scores cannot account for all three of these features si-
 5843 multaneously. Graph-based dependency parsing may also be criticized on the basis of
 5844 intuitions about human language processing: people read and listen to sentences *sequen-*
 5845 *tially*, incrementally building mental models of the sentence structure and meaning before
 5846 getting to the end (Jurafsky, 1996). This seems hard to reconcile with graph-based algo-
 5847 rithms, which perform bottom-up operations on the entire sentence, requiring the parser
 5848 to keep every word in memory. Finally, from a practical perspective, graph-based depen-
 5849 dency parsing is relatively slow, running in cubic time in the length of the input.

5850 Transition-based algorithms address all three of these objections. They work by mov-
 5851 ing through the sentence sequentially, while performing actions that incrementally up-
 5852 date a stored representation of what has been read thus far. As with the shift-reduce
 5853 parser from § 10.6.2, this representation consists of a stack, onto which parsing substruc-
 5854 tures can be pushed and popped. In shift-reduce, these substructures were constituents;
 5855 in the transition systems that follow, they will be projective dependency trees over partial
 5856 spans of the input.⁴ Parsing is complete when the input is consumed and there is only
 5857 a single structure on the stack. The sequence of actions that led to the parse is known as
 5858 the **derivation**. One problem with transition-based systems is that there may be multiple
 5859 derivations for a single parse structure — a phenomenon known as **spurious ambiguity**.

⁴Transition systems also exist for non-projective dependency parsing (e.g., Nivre, 2008).

5860 **11.3.1 Transition systems for dependency parsing**

5861 A **transition system** consists of a representation for describing configurations of the parser,
 5862 and a set of transition actions, which manipulate the configuration. There are two main
 5863 transition systems for dependency parsing: **arc-standard**, which is closely related to shift-
 5864 reduce, and **arc-eager**, which adds an additional action that can simplify derivations (Ab-
 5865 ney and Johnson, 1991). In both cases, transitions are between **configurations** that are
 5866 represented as triples, $C = (\sigma, \beta, A)$, where σ is the stack, β is the input buffer, and A is
 5867 the list of arcs that have been created (Nivre, 2008). In the initial configuration,

$$C_{\text{initial}} = ([\text{ROOT}], \mathbf{w}, \emptyset), \quad [11.16]$$

5868 indicating that the stack contains only the special node ROOT, the entire input is on the
 5869 buffer, and the set of arcs is empty. An accepting configuration is,

$$C_{\text{accept}} = ([\text{ROOT}], \emptyset, A), \quad [11.17]$$

5870 where the stack contains only ROOT, the buffer is empty, and the arcs A define a spanning
 5871 tree over the input. The arc-standard and arc-eager systems define a set of transitions
 5872 between configurations, which are capable of transforming an initial configuration into
 5873 an accepting configuration. In both of these systems, the number of actions required to
 5874 parse an input grows linearly in the length of the input, making transition-based parsing
 5875 considerably more efficient than graph-based methods.

5876 **11.3.1.1 Arc-standard**

5877 The **arc-standard** transition system is closely related to shift-reduce, and to the LR algo-
 5878 rithm that is used to parse programming languages (Aho et al., 2006). It includes the
 5879 following classes of actions:

- 5880 • SHIFT: move the first item from the input buffer on to the top of the stack,

$$(\sigma, i|\beta, A) \Rightarrow (\sigma|i, \beta, A), \quad [11.18]$$

5881 where we write $i|\beta$ to indicate that i is the leftmost item in the input buffer, and $\sigma|i$
 5882 to indicate the result of pushing i on to stack σ .

- 5883 • ARC-LEFT: create a new left-facing arc of type r between the item on the top of the
 5884 stack and the first item in the input buffer. The head of this arc is j , which remains
 5885 at the front of the input buffer. The arc $j \xrightarrow{r} i$ is added to A . Formally,

$$(\sigma|i, j|\beta, A) \Rightarrow (\sigma, j|\beta, A \oplus j \xrightarrow{r} i), \quad [11.19]$$

5886 where r is the label of the dependency arc, and \oplus concatenates the new arc $j \xrightarrow{r} i$ to
 5887 the list A .

σ	β	action	arc added to \mathcal{A}
1. [ROOT]	<i>they like bagels with lox</i>	SHIFT	
2. [ROOT, <i>they</i>]	<i>like bagels with lox</i>	ARC-LEFT	(<i>they</i> \leftarrow <i>like</i>)
3. [ROOT]	<i>like bagels with lox</i>	SHIFT	
4. [ROOT, <i>like</i>]	<i>bagels with lox</i>	SHIFT	
5. [ROOT, <i>like</i> , <i>bagels</i>]	<i>with lox</i>	SHIFT	
6. [ROOT, <i>like</i> , <i>bagels</i> , <i>with</i>]	<i>lox</i>	ARC-LEFT	(<i>with</i> \leftarrow <i>lox</i>)
7. [ROOT, <i>like</i> , <i>bagels</i>]	<i>lox</i>	ARC-RIGHT	(<i>bagels</i> \rightarrow <i>lox</i>)
8. [ROOT, <i>like</i>]	<i>bagels</i>	ARC-RIGHT	(<i>like</i> \rightarrow <i>bagels</i>)
9. [ROOT]	<i>like</i>	ARC-RIGHT	(ROOT \rightarrow <i>like</i>)
10. [ROOT]	\emptyset	DONE	

Table 11.2: Arc-standard derivation of the unlabeled dependency parse for the input *they like bagels with lox*.

- ARC-RIGHT: creates a new right-facing arc of type r between the item on the top of the stack and the first item in the input buffer. The head of this arc is i , which is “popped” from the stack and pushed to the front of the input buffer. The arc $i \xrightarrow{r} j$ is added to A . Formally,

$$(\sigma | i, j | \beta, A) \Rightarrow (\sigma, i | \beta, A \oplus i \xrightarrow{r} j), \quad [11.20]$$

where again r is the label of the dependency arc.

Each action has preconditions. The SHIFT action can be performed only when the buffer has at least one element. The ARC-LEFT action cannot be performed when the root node ROOT is on top of the stack, since this node must be the root of the entire tree. The ARC-LEFT and ARC-RIGHT remove the modifier words from the stack (in the case of ARC-LEFT) and from the buffer (in the case of ARC-RIGHT), so it is impossible for any word to have more than one parent. Furthermore, the end state can only be reached when every word is removed from the buffer and stack, so the set of arcs is guaranteed to constitute a spanning tree. An example arc-standard derivation is shown in Table 11.2.

11.3.1.2 Arc-eager dependency parsing

In the arc-standard transition system, a word is completely removed from the parse once it has been made the modifier in a dependency arc. At this time, any dependents of this word must have already been identified. Right-branching structures are common in English (and many other languages), with words often modified by units such as prepositional phrases to their right. In the arc-standard system, this means that we must first shift all the units of the input onto the stack, and then work backwards, creating a series of

5908 arcs, as occurs in Table 11.2. Note that the decision to shift *bagels* onto the stack guarantees
 5909 that the prepositional phrase *with lox* will attach to the noun phrase, and that this decision
 5910 must be made before the prepositional phrase is itself parsed. This has been argued to be
 5911 cognitively implausible (Abney and Johnson, 1991); from a computational perspective, it
 5912 means that a parser may need to look several steps ahead to make the correct decision.

5913 **Arc-eager dependency parsing** changes the ARC-RIGHT action so that right depen-
 5914 dents can be attached before all of their dependents have been found. Rather than re-
 5915 moving the modifier from both the buffer and stack, the ARC-RIGHT action pushes the
 5916 modifier on to the stack, on top of the head. Because the stack can now contain elements
 5917 that already have parents in the partial dependency graph, two additional changes are
 5918 necessary:

- 5919 • A precondition is required to ensure that the ARC-LEFT action cannot be applied
 5920 when the top element on the stack already has a parent in A .
- 5921 • A new REDUCE action is introduced, which can remove elements from the stack if
 5922 they already have a parent in A :

$$(\sigma|i, \beta, A) \Rightarrow (\sigma, \beta, A). \quad [11.21]$$

5923 As a result of these changes, it is now possible to create the arc *like* \rightarrow *bagels* before parsing
 5924 the prepositional phrase *with lox*. Furthermore, this action does not imply a decision about
 5925 whether the prepositional phrase will attach to the noun or verb. Noun attachment is
 5926 chosen in the parse in Table 11.3, but verb attachment could be achieved by applying the
 5927 REDUCE action at step 5 or 7.

5928 11.3.1.3 Projectivity

5929 The arc-standard and arc-eager transition systems are guaranteed to produce projective
 5930 dependency trees, because all arcs are between the word at the top of the stack and the
 5931 left-most edge of the buffer (Nivre, 2008). Non-projective transition systems can be con-
 5932 structed by adding actions that create arcs with words that are second or third in the
 5933 stack (Attardi, 2006), or by adopting an alternative configuration structure, which main-
 5934 tains a list of all words that do not yet have heads (Covington, 2001). In **pseudo-projective**
 5935 **dependency parsing**, a projective dependency parse is generated first, and then a set of
 5936 graph transformation techniques are applied, producing non-projective edges (Nivre and
 5937 Nilsson, 2005).

5938 11.3.1.4 Beam search

5939 In “greedy” transition-based parsing, the parser tries to make the best decision at each
 5940 configuration. This can lead to search errors, when an early decision locks the parser into

σ	β	action	arc added to \mathcal{A}
1. [ROOT]	<i>they like bagels with lox</i>	SHIFT	
2. [ROOT, <i>they</i>]	<i>like bagels with lox</i>	ARC-LEFT	(<i>they</i> \leftarrow <i>like</i>)
3. [ROOT]	<i>like bagels with lox</i>	ARC-RIGHT	(ROOT \rightarrow <i>like</i>)
4. [ROOT, <i>like</i>]	<i>bagels with lox</i>	ARC-RIGHT	(<i>like</i> \rightarrow <i>bagels</i>)
5. [ROOT, <i>like</i> , <i>bagels</i>]	<i>with lox</i>	SHIFT	
6. [ROOT, <i>like</i> , <i>bagels</i> , <i>with</i>]	<i>lox</i>	ARC-LEFT	(<i>with</i> \leftarrow <i>lox</i>)
7. [ROOT, <i>like</i> , <i>bagels</i>]	<i>lox</i>	ARC-RIGHT	(<i>bagels</i> \rightarrow <i>lox</i>)
8. [ROOT, <i>like</i> , <i>bagels</i> , <i>lox</i>]	\emptyset	REDUCE	
9. [ROOT, <i>like</i> , <i>bagels</i>]	\emptyset	REDUCE	
10. [ROOT, <i>like</i>]	\emptyset	REDUCE	
11. [ROOT]	\emptyset	DONE	

Table 11.3: Arc-eager derivation of the unlabeled dependency parse for the input *they like bagels with lox*.

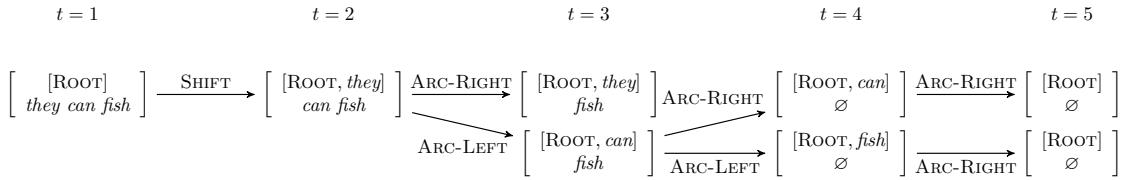


Figure 11.7: Beam search for unlabeled dependency parsing, with beam size $K = 2$. The arc lists for each configuration are not shown, but can be computed from the transitions.

5941 a poor derivation. For example, in Table 11.2, if ARC-RIGHT were chosen at step 4, then
 5942 the parser would later be forced to attach the prepositional phrase *with lox* to the verb
 5943 *likes*. Note that the *likes* \rightarrow *bagels* arc is indeed part of the correct dependency parse, but
 5944 the arc-standard transition system requires it to be created later in the derivation.

Beam search addresses this issue by maintaining a set of hypothetical derivations, called a beam. At step t of the derivation, there is a set of k hypotheses, each of which is a tuple of a score and a sequence of actions,

$$h_t^{(k)} = (s_t^{(k)}, A_t^{(k)}) \quad [11.22]$$

5945 Each hypothesis is then “expanded” by considering the set of all valid actions from the
 5946 current configuration $c_t^{(k)}$, written $\mathcal{A}(c_t^{(k)})$. This yields a large set of new hypotheses. For
 5947 each action $a \mathcal{A}(c_t^{(k)})$, we score the new hypothesis $A_t^{(k)} \oplus a$. The top k hypotheses by
 5948 this scoring metric are kept, and parsing proceeds to the next step (Zhang and Clark,

5949 2008). Note that beam search requires a scoring function for action *sequences*, rather than
 5950 individual actions. This issue will be revisited in the next section.

5951 An example of beam search is shown in Figure 11.7, with a beam size of $K = 2$. For the
 5952 first transition, the only valid action is SHIFT, so there is only one possible configuration
 5953 at $t = 2$. From this configuration, there are three possible actions. The top two are ARC-
 5954 RIGHT and ARC-LEFT, and so the resulting hypotheses from these actions are on the beam
 5955 at $t = 3$. From these configurations, there are three possible actions each, but the best
 5956 two are expansions of the bottom hypothesis at $t = 3$. Parsing continues until $t = 5$, at
 5957 which point both hypotheses reach an accepting state. The best-scoring hypothesis is then
 5958 selected as the parse.

5959 11.3.2 Scoring functions for transition-based parsers

Transition-based parsing requires selecting a series of actions. In greedy transition-based
 parsing, this can be done by training a classifier,

$$\hat{a} = \underset{a \in \mathcal{A}(c)}{\operatorname{argmax}} \Psi(a, c, \mathbf{w}; \boldsymbol{\theta}), \quad [11.23]$$

5960 where $\mathcal{A}(c)$ is the set of admissible actions in the current configuration c , \mathbf{w} is the input,
 5961 and Ψ is a scoring function with parameters $\boldsymbol{\theta}$ (Yamada and Matsumoto, 2003).

5962 A feature-based score can be computed, $\Psi(a, c, \mathbf{w}) = \boldsymbol{\theta} \cdot \mathbf{f}(a, c, \mathbf{w})$, using features that
 5963 may consider any aspect of the current configuration and input sequence. Typical features
 5964 for transition-based dependency parsing include: the word and part-of-speech of the top
 5965 element on the stack; the word and part-of-speech of the first, second, and third elements
 5966 on the input buffer; pairs and triples of words and parts-of-speech from the top of the
 5967 stack and the front of the buffer; the distance (in tokens) between the element on the top
 5968 of the stack and the element in the front of the input buffer; the number of modifiers of
 5969 each of these elements; and higher-order dependency features as described above in the
 5970 section on graph-based dependency parsing (see, e.g., Zhang and Nivre, 2011).

5971 Parse actions can also be scored by neural networks. For example, Chen and Manning
 5972 (2014) build a feedforward network in which the input layer consists of the concatenation
 5973 of embeddings of several words and tags:

- 5974 • the top three words on the stack, and the first three words on the buffer;
- 5975 • the first and second leftmost and rightmost children (dependents) of the top two
 5976 words on the stack;
- 5977 • the leftmost and right most grandchildren of the top two words on the stack;
- 5978 • embeddings of the part-of-speech tags of these words.

Let us call this base layer $\mathbf{x}(c, \mathbf{w})$, defined as,

$$c = (\sigma, \beta, A)$$

$$\mathbf{x}(c, \mathbf{w}) = [\mathbf{v}_{w_{\sigma_1}}, \mathbf{v}_{t_{\sigma_1}} \mathbf{v}_{w_{\sigma_2}}, \mathbf{v}_{t_{\sigma_2}}, \mathbf{v}_{w_{\sigma_3}}, \mathbf{v}_{t_{\sigma_3}}, \mathbf{v}_{w_{\beta_1}}, \mathbf{v}_{t_{\beta_1}}, \mathbf{v}_{w_{\beta_2}}, \mathbf{v}_{t_{\beta_2}}, \dots],$$

where $\mathbf{v}_{w_{\sigma_1}}$ is the embedding of the first word on the stack, $\mathbf{v}_{t_{\beta_2}}$ is the embedding of the part-of-speech tag of the second word on the buffer, and so on. Given this base encoding of the parser state, the score for the set of possible actions is computed through a feedforward network,

$$\mathbf{z} = g(\Theta^{(x \rightarrow z)} \mathbf{x}(c, \mathbf{w})) \quad [11.24]$$

$$\psi(a, c, \mathbf{w}; \boldsymbol{\theta}) = \Theta_a^{(z \rightarrow y)} \mathbf{z}, \quad [11.25]$$

5979 where the vector \mathbf{z} plays the same role as the features $\mathbf{f}(a, c, \mathbf{w})$, but is a learned representation.
 5980 Chen and Manning (2014) use a cubic elementwise activation function, $g(x) = x^3$,
 5981 so that the hidden layer models products across all triples of input features. The learning
 5982 algorithm updates the embeddings as well as the parameters of the feedforward network.

5983 11.3.3 Learning to parse

5984 Transition-based dependency parsing suffers from a mismatch between the supervision,
 5985 which comes in the form of dependency trees, and the classifier's prediction space, which
 5986 is a set of parsing actions. One solution is to create new training data by converting parse
 5987 trees into action sequences; another is to derive supervision directly from the parser's
 5988 performance.

5989 11.3.3.1 Oracle-based training

5990 A transition system can be viewed as a function from action sequences (also called **derivations**)
 5991 to parse trees. The inverse of this function is a mapping from parse trees to derivations,
 5992 which is called an **oracle**. For the arc-standard and arc-eager parsing system, an
 5993 oracle can be computed in linear time in the length of the derivation (Kübler et al., 2009,
 5994 page 32). Both the arc-standard and arc-eager transition systems suffer from **spurious**
 5995 **ambiguity**: there exist dependency parses for which multiple derivations are possible,
 5996 such as $1 \leftarrow 2 \rightarrow 3$. The oracle must choose between these different derivations. For exam-
 5997 ple, the algorithm described by Kübler et al. (2009) would first create the left arc ($1 \leftarrow 2$),
 5998 and then create the right arc, $(1 \leftarrow 2) \rightarrow 3$; another oracle might begin by shifting twice,
 5999 resulting in the derivation $1 \leftarrow (2 \rightarrow 3)$.

Given such an oracle, a dependency treebank can be converted into a set of oracle action sequences $\{A^{(i)}\}_{i=1}^N$. The parser can be trained by stepping through the oracle action sequences, and optimizing on an classification-based objective that rewards selecting the

oracle action. For transition-based dependency parsing, maximum conditional likelihood is a typical choice (Chen and Manning, 2014; Dyer et al., 2015):

$$p(a | c, \mathbf{w}) = \frac{\exp \Psi(a, c, \mathbf{w}; \boldsymbol{\theta})}{\sum_{a' \in \mathcal{A}(c)} \exp \Psi(a', c, \mathbf{w}; \boldsymbol{\theta})} \quad [11.26]$$

$$\hat{\boldsymbol{\theta}} = \underset{\boldsymbol{\theta}}{\operatorname{argmax}} \sum_{i=1}^N \sum_{t=1}^{|A^{(i)}|} \log p(a_t^{(i)} | c_t^{(i)}, \mathbf{w}), \quad [11.27]$$

where $|A^{(i)}|$ is the length of the action sequence $A^{(i)}$.

Recall that beam search requires a scoring function for action sequences. Such a score can be obtained by adding the log-likelihoods (or hinge losses) across all actions in the sequence (Chen and Manning, 2014).

11.3.3.2 Global objectives

The objective in Equation 11.27 is **locally-normalized**: it is the product of normalized probabilities over individual actions. A similar characterization could be made of non-probabilistic algorithms in which hinge-loss objectives are summed over individual actions. In either case, training on individual actions can be sub-optimal with respect to global performance, due to the **label bias problem** (Lafferty et al., 2001; Andor et al., 2016).

As a stylized example, suppose that a given configuration appears 100 times in the training data, with action a_1 as the oracle action in 51 cases, and a_2 as the oracle action in the other 49 cases. However, in cases where a_2 is correct, choosing a_1 results in a cascade of subsequent errors, while in cases where a_1 is correct, choosing a_2 results in only a single error. A classifier that is trained on a local objective function will learn to always choose a_1 , but choosing a_2 would minimize the overall number of errors.

This observation motivates a global objective, such as the globally-normalized conditional likelihood,

$$p(A^{(i)} | \mathbf{w}; \boldsymbol{\theta}) = \frac{\exp \sum_{t=1}^{|A^{(i)}|} \Psi(a_t^{(i)}, c_t^{(i)}, \mathbf{w})}{\sum_{A' \in \mathbb{A}(\mathbf{w})} \exp \sum_{t=1}^{|A'|} \Psi(a'_t, c'_t, \mathbf{w})}, \quad [11.28]$$

where the denominator sums over the set of all possible action sequences, $\mathbb{A}(\mathbf{w})$.⁵ In the conditional random field model for sequence labeling (§ 7.5.3), it was possible to compute

⁵Andor et al. (2016) prove that the set of globally-normalized conditional distributions is a strict superset of the set of locally-normalized conditional distributions, and that globally-normalized conditional models are therefore strictly more expressive.

this sum explicitly, using dynamic programming. In transition-based parsing, this is not possible. However, the sum can be approximated using beam search,

$$\sum_{A' \in \mathbb{A}(\mathbf{w})} \exp \sum_{t=1}^{|A'|} \Psi(a'_t, c'_t, \mathbf{w}) \approx \sum_{k=1}^K \exp \sum_{t=1}^{|A^{(k)}|} \Psi(a_t^{(k)}, c_t^{(k)}, \mathbf{w}), \quad [11.29]$$

where $A^{(k)}$ is an action sequence on a beam of size K . This gives rise to the following loss function,

$$L(\boldsymbol{\theta}) = - \sum_{t=1}^{|A^{(i)}|} \Psi(a_t^{(i)}, c_t^{(i)}, \mathbf{w}) + \log \sum_{k=1}^K \exp \sum_{t=1}^{|A^{(k)}|} \Psi(a_t^{(k)}, c_t^{(k)}, \mathbf{w}). \quad [11.30]$$

6019 The derivatives of this loss involve expectations with respect to a probability distribution
6020 over action sequences on the beam.

6021 11.3.3.3 *Early update and the incremental perceptron

6022 When learning in the context of beam search, the goal is to learn a decision function so that
6023 the gold dependency parse is always reachable from at least one of the partial derivations
6024 on the beam. (The combination of a transition system (such as beam search) and a scoring
6025 function for actions is known as a **policy**.) To achieve this, we can make an **early update**
6026 as soon as the oracle action sequence “falls off” the beam, even before a complete analysis
6027 is available (Collins and Roark, 2004; Daumé III and Marcu, 2005). The loss can be based
6028 on the best-scoring hypothesis on the beam, or the sum of all hypotheses (Huang et al.,
6029 2012).

6030 For example, consider the beam search in Figure 11.7. In the correct parse, *fish* is the
6031 head of dependency arcs to both of the other two words. In the arc-standard system,
6032 this can be achieved only by using SHIFT for the first two actions. At $t = 3$, the oracle
6033 action sequence has fallen off the beam. The parser should therefore stop, and update the
6034 parameters by the gradient $\frac{\partial}{\partial \boldsymbol{\theta}} L(A_{1:3}^{(i)}, \{A_{1:3}^{(k)}\}; \boldsymbol{\theta})$, where $A_{1:3}^{(i)}$ is the first three actions of the
6035 oracle sequence, and $\{A_{1:3}^{(k)}\}$ is the beam.

6036 This integration of incremental search and learning was first developed in the **incremental**
6037 **perceptron** (Collins and Roark, 2004). This method updates the parameters with
6038 respect to a hinge loss, which compares the top-scoring hypothesis and the gold action
6039 sequence, up to the current point t . Several improvements to this basic protocol are pos-
6040 sible:

- 6041 • As noted earlier, the gold dependency parse can be derived by multiple action se-
6042 quences. Rather than checking for the presence of a single oracle action sequence on
6043 the beam, we can check if the gold dependency parse is *reachable* from the current
6044 beam, using a **dynamic oracle** (Goldberg and Nivre, 2012).

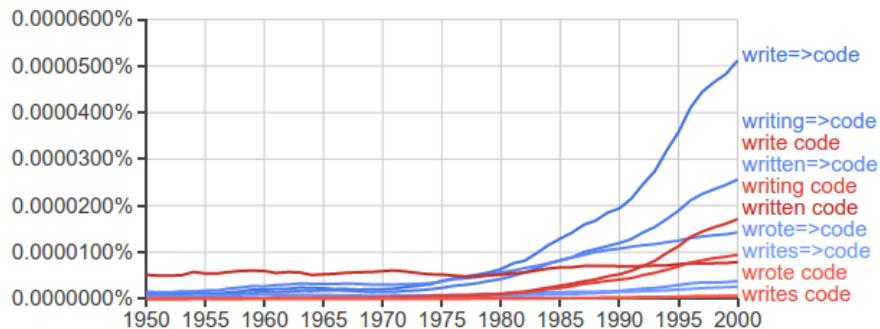


Figure 11.8: Google n-grams results for the bigram *write code* and the dependency arc *write => code* (and their morphological variants)

- 6045 • By maximizing the score of the gold action sequence, we are training a decision
- 6046 function to find the correct action given the gold context. But in reality, the parser
- 6047 will make errors, and the parser is not trained to find the best action given a context
- 6048 that may not itself be optimal. This issue is addressed by various generalizations of
- 6049 incremental perceptron, known as **learning to search** (Daumé III et al., 2009). Some
- 6050 of these methods are discussed in chapter 15.

6051 11.4 Applications

6052 Dependency parsing is used in many real-world applications: any time you want to know
 6053 about pairs of words which might not be adjacent, you can use dependency arcs instead
 6054 of regular expression search patterns. For example, you may want to match strings like
 6055 *delicious pastries*, *delicious French pastries*, and *the pastries are delicious*.

6056 It is possible to search the Google *n*-gramscorpus by dependency edges, finding the
 6057 trend in how often a dependency edge appears over time. For example, we might be inter-
 6058 ested in knowing when people started talking about *writing code*, but we also want *write*
 6059 *some code*, *write good code*, *write all the code*, etc. The result of a search on the dependency
 6060 edge *write → code* is shown in Figure 11.8. This capability has been applied to research
 6061 in digital humanities, such as the analysis of gender in Shakespeare Muralidharan and
 6062 Hearst (2013).

A classic application of dependency parsing is **relation extraction**, which is described

in chapter 17. The goal of relation extraction is to identify entity pairs, such as

(MELVILLE, MOBY-DICK)
 (TOLSTOY, WAR AND PEACE)
 (MARQUÉZ, 100 YEARS OF SOLITUDE)
 (SHAKESPEARE, A MIDSUMMER NIGHT'S DREAM),

6063 which stand in some relation to each other (in this case, the relation is authorship). Such
 6064 entity pairs are often referenced via consistent chains of dependency relations. Therefore,
 6065 dependency paths are often a useful feature in supervised systems which learn to detect
 6066 new instances of a relation, based on labeled examples of other instances of the same
 6067 relation type (Culotta and Sorensen, 2004; Fundel et al., 2007; Mintz et al., 2009).

6068 Cui et al. (2005) show how dependency parsing can improve automated question an-
 6069 swering. Suppose you receive the following query:

6070 (11.1) What percentage of the nation's cheese does Wisconsin produce?

6071 The corpus contains this sentence:

6072 (11.2) In Wisconsin, where farmers produce 28% of the nation's cheese, ...

6073 The location of *Wisconsin* in the surface form of this string makes it a poor match for the
 6074 query. However, in the dependency graph, there is an edge from *produce* to *Wisconsin* in
 6075 both the question and the potential answer, raising the likelihood that this span of text is
 6076 relevant to the question.

6077 A final example comes from sentiment analysis. As discussed in chapter 4, the polarity
 6078 of a sentence can be reversed by negation, e.g.

6079 (11.3) *There is no reason at all to believe the polluters will suddenly become reasonable.*

6080 By tracking the sentiment polarity through the dependency parse, we can better iden-
 6081 tify the overall polarity of the sentence, determining when key sentiment words are re-
 6082 versed (Wilson et al., 2005; Nakagawa et al., 2010).

6083 Additional resources

6084 More details on dependency grammar and parsing algorithms can be found in the manuscript
 6085 by Kübler et al. (2009). For a comprehensive but whimsical overview of graph-based de-
 6086 pendency parsing algorithms, see Eisner (1997). Jurafsky and Martin (2018) describe an
 6087 **agenda-based** version of beam search, in which the beam contains hypotheses of varying
 6088 lengths. New hypotheses are added to the beam only if their score is better than the worst

item currently on the beam. Another search algorithm for transition-based parsing is **easy-first**, which abandons the left-to-right traversal order, and adds the highest-scoring edges first, regardless of where they appear (Goldberg and Elhadad, 2010). Goldberg et al. (2013) note that although transition-based methods can be implemented in linear time in the length of the input, naïve implementations of beam search will require quadratic time, due to the cost of copying each hypothesis when it is expanded on the beam. This issue can be addressed by using a more efficient data structure for the stack.

Exercises

1. The dependency structure $1 \leftarrow 2 \rightarrow 3$, with 2 as the root, can be obtained from more than one set of actions in arc-standard parsing. List both sets of actions that can obtain this parse.
2. Suppose you have a set of unlabeled arc scores $\psi(i \rightarrow j)$, where the score depends only on the identity of the two words. The scores include $\psi(\text{ROOT} \rightarrow j)$.
 - Assuming each word occurs only once in the sentence ($(i \neq j) \Leftarrow (w_i \neq w_j)$), how would you construct a weighted lexicalized context-free grammar so that the score of *any* projective dependency tree is equal to the score of some equivalent derivation in the lexicalized context-free grammar?
 - Verify that your method works for a simple example like *they eat fish*.
 - How would you adapt your method to handle the case an individual word may appear multiple times in the sentence?
3. Provide the UD-style dependency parse for the sentence *Xi-Lan eats shoots and leaves*, assuming *leaves* is a verb. Provide arc-standard and arc-eager derivations for this dependency parse.

6112

Part III

6113

Meaning

6114 Chapter 12

6115 Logical semantics

6116 The previous few chapters have focused on building systems that reconstruct the **syntax**
6117 of natural language — its structural organization — through tagging and parsing. But
6118 some of the most exciting and promising potential applications of language technology
6119 involve going beyond syntax to **semantics** — the underlying meaning of the text:

- 6120 • Answering questions, such as *where is the nearest coffeeshop?* or *what is the middle name*
6121 *of the mother of the 44th President of the United States?*.
- 6122 • Building a robot that can follow natural language instructions to execute tasks.
- 6123 • Translating a sentence from one language into another, while preserving the under-
6124 lying meaning.
- 6125 • Fact-checking an article by searching the web for contradictory evidence.
- 6126 • Logic-checking an argument by identifying contradictions, ambiguity, and unsup-
6127 ported assertions.

6128 Semantic analysis involves converting natural language into a **meaning representa-**
6129 **tion**. To be useful, a meaning representation must meet several criteria:

- 6130 • **c1**: it should be unambiguous: unlike natural language, there should be exactly one
6131 meaning per statement;
- 6132 • **c2**: it should provide a way to link language to external knowledge, observations,
6133 and actions;
- 6134 • **c3**: it should support computational **inference**, so that meanings can be combined
6135 to derive additional knowledge;
- 6136 • **c4**: it should be expressive enough to cover the full range of things that people talk
6137 about in natural language.

6138 Much more than this can be said about the question of how best to represent knowledge
 6139 for computation (e.g., Sowa, 2000), but this chapter will focus on these four criteria.

6140 12.1 Meaning and denotation

6141 The first criterion for a meaning representation is that statements in the representation
 6142 should be unambiguous — they should have only one possible interpretation. Natural
 6143 language does not have this property: as we saw in chapter 10, sentences like *cats scratch*
 6144 *people with claws* have multiple interpretations.

6145 But what does it mean for a statement to be unambiguous? Programming languages
 6146 provide a useful example: the output of a program is completely specified by the rules of
 6147 the language and the properties of the environment in which the program is run. For ex-
 6148 ample, the python code $5 + 3$ will have the output 8, as will the codes $(4 * 4) - (3 * 3) + 1$
 6149 and $((8))$. This output is known as the **denotation** of the program, and can be written
 6150 as,

$$\llbracket 5+3 \rrbracket = \llbracket (4 * 4) - (3 * 3) + 1 \rrbracket = \llbracket ((8)) \rrbracket = 8. \quad [12.1]$$

6151 The denotations of these arithmetic expressions are determined by the meaning of the
 6152 **constants** (e.g., 5, 3) and the **relations** (e.g., $+$, $*$, $(,)$). Now let's consider another snippet
 6153 of python code, `double(4)`. The denotation of this code could be, $\llbracket \text{double}(4) \rrbracket = 8$, or
 6154 it could be $\llbracket \text{double}(4) \rrbracket = 44$ — it depends on the meaning of `double`. This meaning
 6155 is defined in a **world model** \mathcal{M} as an infinite set of pairs. We write the denotation with
 6156 respect to model \mathcal{M} as $\llbracket \cdot \rrbracket_{\mathcal{M}}$, e.g., $\llbracket \text{double} \rrbracket_{\mathcal{M}} = \{(0, 0), (1, 2), (2, 4), \dots\}$. The world
 6157 model would also define the (infinite) list of constants, e.g., $\{0, 1, 2, \dots\}$. As long as the
 6158 denotation of string ϕ in model \mathcal{M} can be computed unambiguously, the language can be
 6159 said to be unambiguous.

6160 This approach to meaning is known as **model-theoretic semantics**, and it addresses
 6161 not only criterion *c1* (no ambiguity), but also *c2* (connecting language to external knowl-
 6162 edge, observations, and actions). For example, we can connect a representation of the
 6163 meaning of a statement like *the capital of Georgia* with a world model that includes knowl-
 6164 edge base of geographical facts, obtaining the denotation `Atlanta`. We might populate
 6165 a world model by applying an image analysis algorithm to Figure 12.1, and then use this
 6166 world model to evaluate **propositions** like *a man is riding a moose*. Another desirable prop-
 6167 erty of model-theoretic semantics is that when the facts change, the denotations change
 6168 too: the meaning representation of *President of the USA* would have a different denotation
 6169 in the model \mathcal{M}_{2014} as it would in \mathcal{M}_{2022} .



Figure 12.1: A (doctored) image, which could be the basis for a world model

6170 12.2 Logical representations of meaning

6171 Criterion *c3* requires that the meaning representation support inference — for example,
 6172 automatically deducing new facts from known premises. While many representations
 6173 have been proposed that meet these criteria, the most mature is the language of first-order
 6174 logic.¹

6175 12.2.1 Propositional logic

6176 The bare bones of logical meaning representation are Boolean operations on propositions:

6177 **Propositional symbols.** Greek symbols like ϕ and ψ will be used to represent **proposi-**
 6178 **tions**, which are statements that are either true or false. For example, ϕ may corre-
 6179 **spond to the proposition, *bagels are delicious*.**

6180 **Boolean operators.** We can build up more complex propositional formulas from Boolean
 6181 operators. These include:

- 6182 • Negation $\neg\phi$, which is true if ϕ is false.

¹Alternatives include the “variable-free” representation used in semantic parsing of geographical queries (Zelle and Mooney, 1996) and robotic control (Ge and Mooney, 2005), and dependency-based compositional semantics (Liang et al., 2013).

- 6183 • Conjunction, $\phi \wedge \psi$, which is true if both ϕ and ψ are true.
- 6184 • Disjunction, $\phi \vee \psi$, which is true if at least one of ϕ and ψ is true
- 6185 • Implication, $\phi \Rightarrow \psi$, which is true unless ϕ is true and ψ is false. Implication
6186 has identical truth conditions to $\neg\phi \vee \psi$.
- 6187 • Equivalence, $\phi \Leftrightarrow \psi$, which is true if ϕ and ψ are both true or both false. Equiv-
6188 alence has identical truth conditions to $(\phi \Rightarrow \psi) \wedge (\psi \Rightarrow \phi)$.

6189 It is not strictly necessary to have all five Boolean operators: readers familiar with
6190 Boolean logic will know that it is possible to construct all other operators from either the
6191 NAND (not-and) or NOR (not-or) operators. Nonetheless, it is clearest to use all five
6192 operators. From the truth conditions for these operators, it is possible to define a number
6193 of “laws” for these Boolean operators, such as,

- 6194 • *Commutativity*: $\phi \wedge \psi = \psi \wedge \phi$, $\phi \vee \psi = \psi \vee \phi$
- 6195 • *Associativity*: $\phi \wedge (\psi \wedge \chi) = (\phi \wedge \psi) \wedge \chi$, $\phi \vee (\psi \vee \chi) = (\phi \vee \psi) \vee \chi$
- 6196 • *Complementation*: $\phi \wedge \neg\phi = \perp$, $\phi \vee \neg\phi = \top$, where \top indicates a true proposition
6197 and \perp indicates a false proposition.

These laws can be combined to derive further equivalences, which can support logical inferences. For example, suppose $\phi = \text{The music is loud}$ and $\psi = \text{Max can't sleep}$. Then if we are given,

$$\begin{aligned} \phi \Rightarrow \psi & \quad \text{If the music is loud, Max can't sleep.} \\ \phi & \quad \text{The music is loud.} \end{aligned}$$

6198 we can derive ψ (*Max can't sleep*) by application of **modus ponens**, which is one of a
6199 set of **inference rules** that can be derived from more basic laws and used to manipulate
6200 propositional formulas. **Automated theorem provers** are capable of applying inference
6201 rules to a set of premises to derive desired propositions (Loveland, 2016).

6202 12.2.2 First-order logic

6203 Propositional logic is so named because it treats propositions as its base units. However,
6204 the criterion *c4* states that our meaning representation should be sufficiently expressive.
6205 Now consider the sentence pair,

- 6206 (12.1) If anyone is making noise, then Max can't sleep.
6207 Abigail is making noise.

6208 People are capable of making inferences from this sentence pair, but such inferences re-
6209 quire formal tools that are beyond propositional logic. To understand the relationship

6210 between the statement *anyone is making noise* and the statement *Abigail is making noise*, our
 6211 meaning representation requires the additional machinery of **first-order logic** (FOL).

6212 In FOL, logical propositions can be constructed from relationships between entities.
 6213 Specifically, FOL extends propositional logic with the following classes of terms:

6214 **Constants.** These are elements that name individual entities in the model, such as MAX
 6215 and ABIGAIL. The denotation of each constant in a model \mathcal{M} is an element in the
 6216 model, e.g., $[\![\text{MAX}]\!] = m$ and $[\![\text{ABIGAIL}]\!] = a$.

6217 **Relations.** Relations can be thought of as sets of entities, or sets of tuples. For example,
 6218 the relation CAN-SLEEP is defined as the set of entities who can sleep, and has the
 6219 denotation $[\![\text{CAN-SLEEP}]\!] = \{a, m, \dots\}$. To test the truth value of the proposition
 6220 CAN-SLEEP(MAX), we ask whether $[\![\text{MAX}]\!] \in [\![\text{CAN-SLEEP}]\!]$. Logical relations that are
 6221 defined over sets of entities are sometimes called **properties**.

6222 Relations may also be ordered tuples of entities. For example BROTHER(MAX,ABIGAIL)
 6223 expresses the proposition that MAX is the brother of ABIGAIL. The denotation of
 6224 such relations is a set of tuples, $[\![\text{BROTHER}]\!] = \{(m, a), (x, y), \dots\}$. To test the
 6225 truth value of the proposition BROTHER(MAX,ABIGAIL), we ask whether the tuple
 6226 $([\![\text{MAX}]\!], [\![\text{ABIGAIL}]\!])$ is in the denotation $[\![\text{BROTHER}]\!]$.

Using constants and relations, it is possible to express statements like *Max can't sleep* and *Max is Abigail's brother*:

$$\neg \text{CAN-SLEEP}(\text{MAX}) \\ \text{BROTHER}(\text{MAX}, \text{ABIGAIL}).$$

These statements can also be combined using Boolean operators, such as,

$$(\text{BROTHER}(\text{MAX}, \text{ABIGAIL}) \vee \text{BROTHER}(\text{MAX}, \text{STEVE})) \Rightarrow \neg \text{CAN-SLEEP}(\text{MAX}).$$

6227 This fragment of first-order logic permits only statements about specific entities. To
 6228 support inferences about statements like *If anyone is making noise, then Max can't sleep*,
 6229 two more elements must be added to the meaning representation:

6230 **Variables.** Variables are mechanisms for referring to entities that are not locally specified.
 6231 We can then write CAN-SLEEP(x) or BROTHER(x , ABIGAIL). In these cases, x is a **free**
 6232 **variable**, meaning that we have not committed to any particular assignment.

6233 **Quantifiers.** Variables are bound by quantifiers. There are two quantifiers in first-order
 6234 logic.²

- 6235 • The **existential quantifier** \exists , which indicates that there must be at least one en-
 6236 tity to which the variable can bind. For example, the statement $\exists x \text{MAKES-NOISE}(x)$
 6237 indicates that there is at least one entity for which MAKES-NOISE is true.
 6238 • The **universal quantifier** \forall , which indicates that the variable must be able to
 6239 bind to any entity in the model. For example, the statement,

$$\text{MAKES-NOISE(ABIGAIL)} \Rightarrow (\forall x \neg \text{CAN-SLEEP}(x)) \quad [12.3]$$

6240 asserts that if Abigail makes noise, no one can sleep.

6241 The expressions $\exists x$ and $\forall x$ make x into a **bound variable**. A formula that contains
 6242 no free variables is a **sentence**.

6243 **Functions.** Functions map from entities to entities, e.g., $\llbracket \text{CAPITAL-OF(GEORGIA)} \rrbracket = \llbracket \text{ATLANTA} \rrbracket$.
 6244 With functions, it is convenient to add an equality operator, supporting statements
 6245 like,

$$\forall x \exists y \text{MOTHER-OF}(x) = \text{DAUGHTER-OF}(y). \quad [12.4]$$

6246 Note that MOTHER-OF is a functional analogue of the relation MOTHER, so that
 6247 $\text{MOTHER-OF}(x) = y$ if $\text{MOTHER}(x, y)$. Any logical formula that uses functions can be
 6248 rewritten using only relations and quantification. For example,

$$\text{MAKES-NOISE}(\text{MOTHER-OF(ABIGAIL)}) \quad [12.5]$$

6249 can be rewritten as $\exists x \text{MAKES-NOISE}(x) \wedge \text{MOTHER}(x, \text{ABIGAIL})$.

An important property of quantifiers is that the order can matter. Unfortunately, natural language is rarely clear about this! The issue is demonstrated by examples like *everyone speaks a language*, which has the following interpretations:

$$\forall x \exists y \text{ SPEAKS}(x, y) \quad [12.6]$$

$$\exists y \forall x \text{ SPEAKS}(x, y). \quad [12.7]$$

6250 In the first case, y may refer to several different languages, while in the second case, there
 6251 is a single y that is spoken by everyone.

²In first-order logic, it is possible to quantify only over entities. In **second-order logic**, it is possible to quantify over properties, supporting statements like *Butch has every property that a good boxer has* (example from Blackburn and Bos, 2005),

$$\forall P \forall x ((\text{GOOD-BOXER}(x) \Rightarrow P(x)) \Rightarrow P(\text{BUTCH})). \quad [12.2]$$

6252 **12.2.2.1 Truth-conditional semantics**

6253 One way to look at the meaning of an FOL sentence ϕ is as a set of **truth conditions**,
 6254 or models under which ϕ is satisfied. But how to determine whether a sentence is true
 6255 or false in a given model? We will approach this inductively, starting with a predicate
 6256 applied to a tuple of constants. The truth of such a sentence depends on whether the
 6257 tuple of denotations of the constants is in the denotation of the predicate. For example,
 6258 CAPITAL(GEORGIA,ATLANTA) is true in model \mathcal{M} iff,

$$(\llbracket \text{GEORGIA} \rrbracket_{\mathcal{M}}, \llbracket \text{ATLANTA} \rrbracket_{\mathcal{M}}) \in \llbracket \text{CAPITAL} \rrbracket_{\mathcal{M}}. \quad [12.8]$$

6259 The Boolean operators \wedge, \vee, \dots provide ways to construct more complicated sentences,
 6260 and the truth of such statements can be assessed based on the truth tables associated with
 6261 these operators. The statement $\exists x\phi$ is true if there is some assignment of the variable x
 6262 to an entity in the model such that ϕ is true; the statement $\forall x\phi$ is true if ϕ is true under
 6263 all possible assignments of x . More formally, we would say that ϕ is **satisfied** under \mathcal{M} ,
 6264 written as $\mathcal{M} \models \phi$.

6265 Truth conditional semantics allows us to define several other properties of sentences
 6266 and pairs of sentences. Suppose that in every \mathcal{M} under which ϕ is satisfied, another
 6267 formula ψ is also satisfied; then ϕ **entails** ψ , which is also written as $\phi \models \psi$. For example,

$$\text{CAPITAL(GEORGIA,ATLANTA)} \models \exists x \text{CAPITAL(GEORGIA, } x\text{)}. \quad [12.9]$$

6268 A statement that is satisfied under any model, such as $\phi \vee \neg\phi$, is **valid**, written $\models (\phi \vee$
 6269 $\neg\phi)$. A statement that is not satisfied under any model, such as $\phi \wedge \neg\phi$, is **unsatisfiable**,
 6270 or **inconsistent**. A **model checker** is a program that determines whether a sentence ϕ
 6271 is satisfied in \mathcal{M} . A **model builder** is a program that constructs a model in which ϕ
 6272 is satisfied. The problems of checking for consistency and validity in first-order logic
 6273 are **undecidable**, meaning that there is no algorithm that can automatically determine
 6274 whether an FOL formula is valid or inconsistent.

6275 **12.2.2.2 Inference in first-order logic**

6276 Our original goal was to support inferences that combine general statements *If anyone is*
making noise, then Max can't sleep with specific statements like *Abigail is making noise*. We
 6277 can now represent such statements in first-order logic, but how are we to perform the
 6278 inference that *Max can't sleep*? One approach is to use “generalized” versions of proposi-
 6279 tional inference rules like modus ponens, which can be applied to FOL formulas. By
 6280 repeatedly applying such inference rules to a knowledge base of facts, it is possible to
 6281 produce proofs of desired propositions. To find the right sequence of inferences to derive
 6282 a desired theorem, classical artificial intelligence search algorithms like backward chain-
 6283 ing can be applied. Such algorithms are implemented in interpreters for the `prolog` logic
 6284 programming language (Pereira and Shieber, 2002).

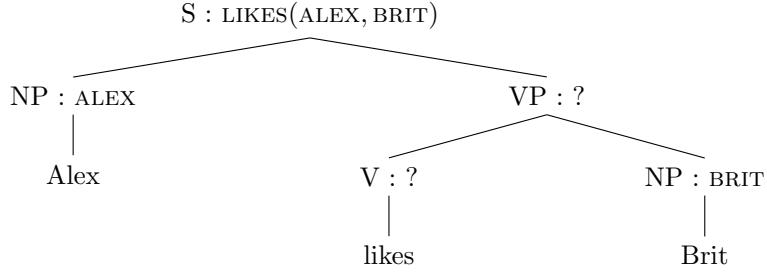


Figure 12.2: The principle of compositionality requires that we identify meanings for the constituents *likes* and *likes Brit* that will make it possible to compute the meaning for the entire sentence.

6286 12.3 Semantic parsing and the lambda calculus

6287 The previous section laid out a lot of formal machinery; the remainder of this chapter
 6288 links these formalisms back to natural language. Given an English sentence like *Alex likes*
 6289 *Brit*, how can we obtain the desired first-order logical representation, $\text{LIKES}(\text{ALEX}, \text{BRIT})$?
 6290 This is the task of **semantic parsing**. Just as a syntactic parser is a function from a natu-
 6291 ral language sentence to a syntactic structure such as a phrase structure tree, a semantic
 6292 parser is a function from natural language to logical formulas.

6293 As in syntactic analysis, semantic parsing is difficult because the space of inputs and
 6294 outputs is very large, and their interaction is complex. Our best hope is that, like syntactic
 6295 parsing, semantic parsing can somehow be decomposed into simpler sub-problems. This
 6296 idea, usually attributed to the German philosopher Gottlob Frege, is called the **principle**
 6297 **of compositionality**: the meaning of a complex expression is a function of the meanings of
 6298 that expression's constituent parts. We will define these "constituent parts" as syntactic
 6299 constituents: noun phrases and verb phrases. These constituents are combined using
 6300 function application: if the syntactic parse contains the production $x \rightarrow y z$, then the
 6301 semantics of x , written $x.\text{sem}$, will be computed as a function of the semantics of the
 6302 constituents, $y.\text{sem}$ and $z.\text{sem}$.³ ⁴

³§ 9.3.2 briefly discusses Combinatory Categorial Grammar (CCG) as an alternative to a phrase-structure analysis of syntax. CCG is argued to be particularly well-suited to semantic parsing (Hockenmaier and Steedman, 2007), and is used in much of the contemporary work on machine learning for semantic parsing, summarized in § 12.4.

⁴The approach of algorithmically building up meaning representations from a series of operations on the syntactic structure of a sentence is generally attributed to the philosopher Richard Montague, who published a series of influential papers on the topic in the early 1970s (e.g., Montague, 1973).

6303 **12.3.1 The lambda calculus**

6304 Let's see how this works for a simple sentence like *Alex likes Brit*, whose syntactic structure
 6305 is shown in Figure 12.2. Our goal is the formula, LIKES(ALEX,BRIT), and it is clear that the
 6306 meaning of the constituents *Alex* and *Brit* should be ALEX and BRIT. That leaves two more
 6307 constituents: the verb *likes*, and the verb phrase *likes Brit*. The meanings of these units
 6308 must be defined in a way that makes it possible to recover the desired meaning for the
 6309 entire sentence by function application. If the meanings of *Alex* and *Brit* are constants,
 6310 then the meanings of *likes* and *likes Brit* must be functional expressions, which can be
 6311 applied to their siblings to produce the desired analyses.

6312 Modeling these partial analyses requires extending the first-order logic meaning rep-
 6313 resentation. We do this by adding **lambda expressions**, which are descriptions of anonym-
 6314 ous functions,⁵ e.g.,

$$\lambda x.\text{LIKES}(x, \text{BRIT}). \quad [12.10]$$

6315 This functional expression is the meaning of the verb phrase *likes Brit*; it takes a single
 6316 argument, and returns the result of substituting that argument for x in the expression
 6317 $\text{LIKES}(x, \text{BRIT})$. We write this substitution as,

$$(\lambda x.\text{LIKES}(x, \text{BRIT}))@\text{ALEX} = \text{LIKES}(\text{ALEX}, \text{BRIT}), \quad [12.11]$$

6318 with the symbol "@" indicating function application. Function application in the lambda
 6319 calculus is sometimes called **β -reduction** or **β -conversion**. The expression $\phi@\psi$ indicates
 6320 a function application to be performed by β -reduction, and $\phi(\psi)$ indicates a function or
 6321 predicate in the final logical form.

6322 Equation 12.11 shows how to obtain the desired semantics for the sentence *Alex likes*
 6323 *Brit*: by applying the lambda expression $\lambda x.\text{LIKES}(x, \text{BRIT})$ to the logical constant ALEX.
 6324 This rule of composition can be specified in a **syntactic-semantic grammar**, in which
 6325 syntactic productions are paired with semantic operations. For the syntactic production
 6326 $S \rightarrow NP VP$, we have the semantic rule $VP.sem @ NP.sem$.

The meaning of the transitive verb phrase *likes Brit* can also be obtained by function application on its syntactic constituents. For the syntactic production $VP \rightarrow V NP$, we apply the semantic rule,

$$VP.sem = (V.sem) @ NP.sem \quad [12.12]$$

$$= (\lambda y. \lambda x. \text{LIKES}(x, y)) @ (\text{BRIT}) \quad [12.13]$$

$$= \lambda x. \text{LIKES}(x, \text{BRIT}). \quad [12.14]$$

⁵Formally, all first-order logic formulas are lambda expressions; in addition, if ϕ is a lambda expression, then $\lambda x.\phi$ is also a lambda expression. Readers who are familiar with functional programming will recognize lambda expressions from their use in programming languages such as Lisp and Python.

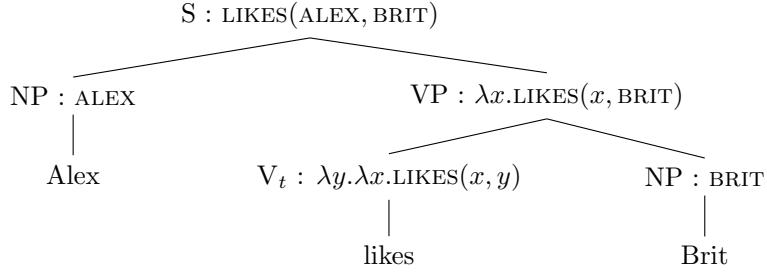


Figure 12.3: Derivation of the semantic representation for *Alex likes Brit* in the grammar G_1 .

S	\rightarrow	NP VP	VP.sem@NP.sem
VP	\rightarrow	V _t NP	V _t .sem@NP.sem
VP	\rightarrow	V _i	V _i .sem
V _t	\rightarrow	likes	$\lambda y. \lambda x. \text{LIKES}(x, y)$
V _i	\rightarrow	sleeps	$\lambda x. \text{SLEEPS}(x)$
NP	\rightarrow	Alex	ALEX
NP	\rightarrow	Brit	BRIT

Table 12.1: G_1 , a minimal syntactic-semantic context-free grammar

6327 Thus, the meaning of the transitive verb *likes* is a lambda expression whose output is
 6328 *another* lambda expression: it takes y as an argument to fill in one of the slots in the LIKES
 6329 relation, and returns a lambda expression that is ready to take an argument to fill in the
 6330 other slot.⁶

6331 Table 12.1 shows a minimal syntactic-semantic grammar fragment, G_1 . The complete
 6332 **derivation** of *Alex likes Brit* in G_1 is shown in Figure 12.3. In addition to the transitive
 6333 verb *likes*, the grammar also includes the intransitive verb *sleeps*; it should be clear how
 6334 to derive the meaning of sentences like *Alex sleeps*. For verbs that can be either transitive
 6335 or intransitive, such as *eats*, we would have two terminal productions, one for each sense
 6336 (terminal productions are also called the **lexical entries**). Indeed, most of the grammar is
 6337 in the **lexicon** (the terminal productions), since these productions select the basic units of
 6338 the semantic interpretation.

⁶This can be written in a few different ways. The notation $\lambda y. x. \text{LIKES}(x, y)$ is a somewhat informal way to indicate a lambda expression that takes two arguments; this would be acceptable in functional programming. Logicians (e.g., Carpenter, 1997) often prefer the more formal notation $\lambda y. \lambda x. \text{LIKES}(x)(y)$, indicating that each lambda expression takes exactly one argument.

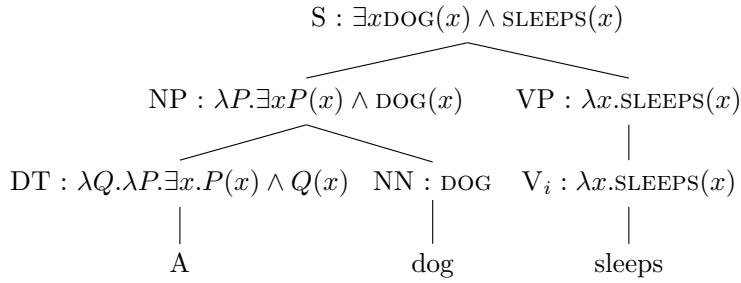


Figure 12.4: Derivation of the semantic representation for *A dog sleeps*, in grammar G_2

6339 12.3.2 Quantification

6340 Things get more complicated when we move from sentences about named entities to sen-
 6341 tences that involve more general noun phrases. Let's consider the example, *A dog sleeps*,
 6342 which has the meaning $\exists x\text{DOG}(x) \wedge \text{SLEEPS}(x)$. Clearly, the DOG relation will be intro-
 6343 duced by the word *dog*, and the SLEEP relation will be introduced by the word *sleeps*.⁷
 6344 The existential quantifier \exists must be introduced by the lexical entry for the determiner *a*.⁷
 6345 However, this seems problematic for the compositional approach taken in the grammar
 6346 G_1 : if the semantics of the noun phrase *a dog* is an existentially quantified expression, how
 6347 can it be the argument to the semantics of the verb *sleeps*, which expects an entity? And
 6348 where does the logical conjunction come from?

6349 There are a few different approaches to handling these issues.⁸ We will begin by re-
 6350 versing the semantic relationship between subject NPs and VPs, so that the production
 6351 $S \rightarrow \text{NP VP}$ has the semantics $\text{NP.sem}@\text{VP.sem}$: the meaning of the sentence is now the
 6352 semantics of the noun phrase applied to the verb phrase. The implications of this change
 6353 are best illustrated by exploring the derivation of the example, shown in Figure 12.4. Let's
 6354 start with the indefinite article *a*, to which we assign the rather intimidating semantics,

$$\lambda P. \lambda Q. \exists x P(x) \wedge Q(x). \quad [12.15]$$

This is a lambda expression that takes two **relations** as arguments, P and Q . The relation P is scoped to the outer lambda expression, so it will be provided by the immediately

⁷Conversely, the sentence *Every dog sleeps* would involve a universal quantifier, $\forall x\text{DOG}(x) \Rightarrow \text{SLEEPS}(x)$. The definite article *the* requires more consideration, since *the dog* must refer to some dog which is uniquely identifiable, perhaps from contextual information external to the sentence. Carpenter (1997, pp. 96-100) summarizes recent approaches to handling definite descriptions.

⁸Carpenter (1997) offers an alternative treatment based on combinatory categorial grammar.

adjacent noun, which in this case is DOG. Thus, the noun phrase *a dog* has the semantics,

$$\text{NP.sem} = \text{DET.sem} @ \text{NN.sem} \quad [12.16]$$

$$= (\lambda P. \lambda Q. \exists x P(x) \wedge Q(x)) @ (\text{DOG}) \quad [12.17]$$

$$= \lambda Q. \exists x \text{DOG}(x) \wedge Q(x). \quad [12.18]$$

6355 This is a lambda expression that is expecting another relation, Q , which will be provided
 6356 by the verb phrase, SLEEPS. This gives the desired analysis, $\exists x \text{DOG}(x) \wedge \text{SLEEPS}(x)$.⁹

6357 If noun phrases like *a dog* are interpreted as lambda expressions, then proper nouns
 6358 like *Alex* must be treated in the same way. This is achieved by **type-raising** from con-
 6359 stants to lambda expressions, $x \Rightarrow \lambda P. P(x)$. After type-raising, the semantics of *Alex* is
 6360 $\lambda P. P(\text{ALEX})$ — a lambda expression that expects a relation to tell us something about
 6361 *ALEX*.¹⁰ Again, make sure you see how the analysis in Figure 12.4 can be applied to the
 6362 sentence *Alex sleeps*.

6363 Direct objects are handled by applying the same type-raising operation to transitive
 6364 verbs: the meaning of verbs such as *likes* is raised to,

$$\lambda P. \lambda x. P(\lambda y. \text{LIKES}(x, y)) \quad [12.19]$$

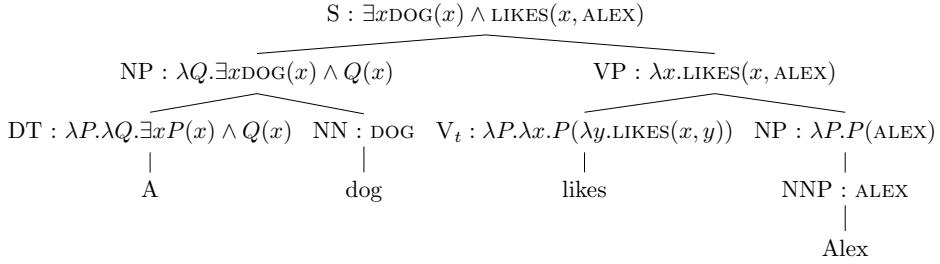
As a result, we can keep the verb phrase production $\text{VP.sem} = \text{V.sem} @ \text{NP.sem}$, knowing
 that the direct object will provide the function P in Equation 12.19. To see how this works,
 let's analyze the verb phrase *likes a dog*. After uniquely relabeling each lambda variable,
 we have,

$$\begin{aligned} \text{VP.sem} &= \text{V.sem} @ \text{NP.sem} \\ &= (\lambda P. \lambda x. P(\lambda y. \text{LIKES}(x, y))) @ (\lambda Q. \exists z \text{DOG}(z) \wedge Q(z)) \\ &= \lambda x. (\lambda Q. \exists z \text{DOG}(z) \wedge Q(z)) @ (\lambda y. \text{LIKES}(x, y)) \\ &= \lambda x. \exists z \text{DOG}(z) \wedge (\lambda y. \text{LIKES}(x, y)) @ z \\ &= \lambda x. \exists z \text{DOG}(z) \wedge \text{LIKES}(x, z). \end{aligned}$$

6365 These changes are summarized in the revised grammar G_2 , shown in Table 12.2. Fig-
 6366 ure 12.5 shows a derivation that involves a transitive verb, an indefinite noun phrase, and
 6367 a proper noun.

⁹When applying β -reduction to arguments that are themselves lambda expressions, be sure to use unique variable names to avoid confusion. For example, it is important to distinguish the x in the semantics for *a* from the x in the semantics for *likes*. Variable names are abstractions, and can always be changed — this is known as **α -conversion**. For example, $\lambda x. P(x)$ can be converted to $\lambda y. P(y)$, etc.

¹⁰Compositional semantic analysis is often supported by **type systems**, which make it possible to check whether a given function application is valid. The base types are entities e and truth values t . A property, such as DOG, is a function from entities to truth values, so its type is written $\langle e, t \rangle$. A transitive verb has type

Figure 12.5: Derivation of the semantic representation for *A dog likes Alex*.

S	\rightarrow NP VP	NP.sem@VP.sem
VP	\rightarrow V _t NP	V _t .sem@NP.sem
VP	\rightarrow V _i	V _i .sem
NP	\rightarrow DET NN	DET.sem@NN.sem
NP	\rightarrow NNP	$\lambda P. P(\text{NNP.sem})$
DET	$\rightarrow a$	$\lambda P. \lambda Q. \exists x P(x) \wedge Q(x)$
DET	\rightarrow every	$\lambda P. \lambda Q. \forall x (P(x) \Rightarrow Q(x))$
V _t	\rightarrow likes	$\lambda P. \lambda x. P(\lambda y. \text{LIKES}(x, y))$
V _i	\rightarrow sleeps	$\lambda x. \text{SLEEPS}(x)$
NN	\rightarrow dog	DOG
NNP	\rightarrow Alex	ALEX
NNP	\rightarrow Brit	BRIT

Table 12.2: G_2 , a syntactic-semantic context-free grammar fragment, which supports quantified noun phrases

6368 12.4 Learning semantic parsers

6369 As with syntactic parsing, any syntactic-semantic grammar with sufficient coverage risks
 6370 producing many possible analyses for any given sentence. Machine learning is the dom-
 6371 inant approach to selecting a single analysis. We will focus on algorithms that learn to
 6372 score logical forms by attaching weights to features of their derivations (Zettlemoyer
 6373 and Collins, 2005). Alternative approaches include transition-based parsing (Zelle and
 6374 Mooney, 1996; Misra and Artzi, 2016) and methods inspired by machine translation (Wong
 6375 and Mooney, 2006). Methods also differ in the form of supervision used for learning,

$\langle e, \langle e, t \rangle \rangle$: after receiving the first entity (the direct object), it returns a function from entities to truth values, which will be applied to the subject of the sentence. The type-raising operation $x \Rightarrow \lambda P. P(x)$ corresponds to a change in type from e to $\langle \langle e, t \rangle, t \rangle$: it expects a function from entities to truth values, and returns a truth value.

which can range from complete derivations to much more limited training signals. We will begin with the case of complete supervision, and then consider how learning is still possible even when seemingly key information is missing.

Datasets Early work on semantic parsing focused on natural language expressions of geographical database queries, such as *What states border Texas*. The GeoQuery dataset of Zelle and Mooney (1996) was originally coded in prolog, but has subsequently been expanded and converted into the SQL database query language by Popescu et al. (2003) and into first-order logic with lambda calculus by Zettlemoyer and Collins (2005), providing logical forms like $\lambda x.\text{STATE}(x) \wedge \text{BORDERS}(x, \text{TEXAS})$. Another early dataset consists of instructions for RoboCup robot soccer teams (Kate et al., 2005). More recent work has focused on broader domains, such as the Freebase database (Bollacker et al., 2008), for which queries have been annotated by Krishnamurthy and Mitchell (2012) and Cai and Yates (2013). Other recent datasets include child-directed speech (Kwiatkowski et al., 2012) and elementary school science exams (Krishnamurthy, 2016).

12.4.1 Learning from derivations

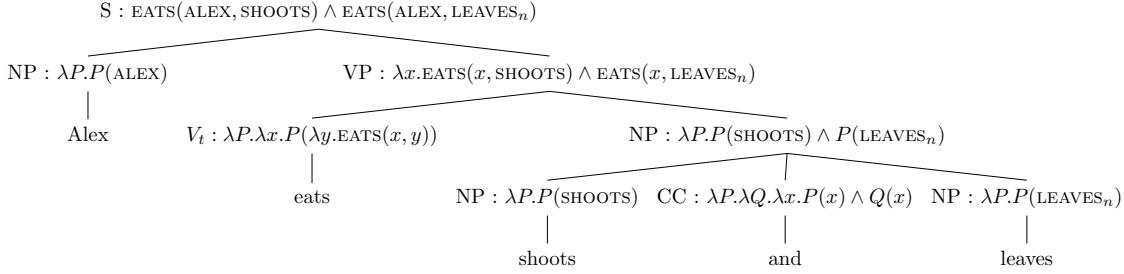
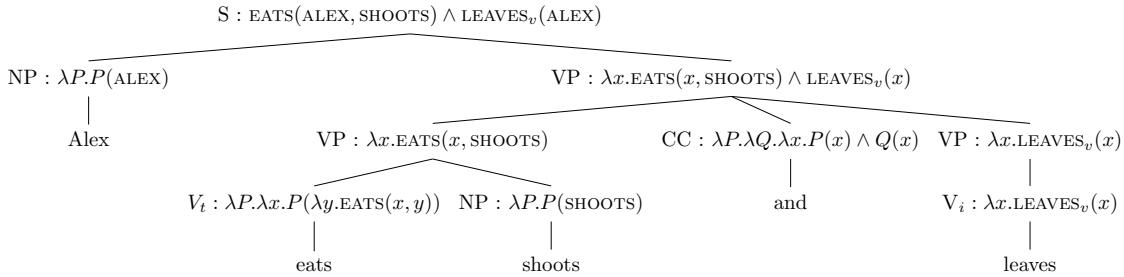
Let $w^{(i)}$ indicate a sequence of text, and let $y^{(i)}$ indicate the desired logical form. For example:

$$\begin{aligned} w^{(i)} &= \text{Alex eats shoots and leaves} \\ y^{(i)} &= \text{EATS(ALEX,SHOOTS)} \wedge \text{EATS(ALEX,LEAVES)} \end{aligned}$$

In the standard supervised learning paradigm that was introduced in § 2.2, we first define a feature function, $f(w, y)$, and then learn weights on these features, so that $y^{(i)} = \operatorname{argmax}_y \theta \cdot f(w, y)$. The weight vector θ is learned by comparing the features of the true label $f(w^{(i)}, y^{(i)})$ against either the features of the predicted label $f(w^{(i)}, \hat{y})$ (perceptron, support vector machine) or the expected feature vector $E_{y|w}[f(w^{(i)}, y)]$ (logistic regression).

While this basic framework seems similar to discriminative syntactic parsing, there is a crucial difference. In (context-free) syntactic parsing, the annotation $y^{(i)}$ contains all of the syntactic productions; indeed, the task of identifying the correct set of productions is identical to the task of identifying the syntactic structure. In semantic parsing, this is not the case: the logical form $\text{EATS(ALEX,SHOOTS)} \wedge \text{EATS(ALEX,LEAVES)}$ does not reveal the syntactic-semantic productions that were used to obtain it. Indeed, there may be **spurious ambiguity**, so that a single logical form can be reached by multiple derivations. (We previously encountered spurious ambiguity in transition-based dependency parsing, § 11.3.2.)

These ideas can be formalized by introducing an additional variable z , representing the **derivation** of the logical form y from the text w . Assume that the feature function de-

Figure 12.6: Derivation for gold semantic analysis of *Alex eats shoots and leaves*Figure 12.7: Derivation for incorrect semantic analysis of *Alex eats shoots and leaves*

6408 composes across the productions in the derivation, $f(\mathbf{w}, \mathbf{z}, \mathbf{y}) = \sum_{t=1}^T f(\mathbf{w}, z_t, \mathbf{y})$, where
 6409 z_t indicates a single syntactic-semantic production. For example, we might have a feature
 6410 for the production $S \rightarrow NP VP : NP.sem@VP.sem$, as well as for terminal productions
 6411 like $NNP \rightarrow Alex : ALEX$. Under this decomposition, it is possible to compute scores
 6412 for each semantically-annotated subtree in the analysis of \mathbf{w} , so that bottom-up parsing
 6413 algorithms like CKY (§ 10.1) can be applied to find the best-scoring semantic analysis.

6414 Figure 12.6 shows a derivation of the correct semantic analysis of the sentence *Alex*
 6415 *eats shoots and leaves*, in a simplified grammar in which the plural noun phrases *shoots*
 6416 and *leaves* are interpreted as logical constants *SHOOTS* and *LEAVES_n*. Figure 12.7 shows a
 6417 derivation of an incorrect analysis. Assuming one feature per production, the perceptron
 6418 update is shown in Table 12.3. From this update, the parser would learn to prefer the
 6419 noun interpretation of *leaves* over the verb interpretation. It would also learn to prefer
 6420 noun phrase coordination over verb phrase coordination.

6421 While the update is explained in terms of the perceptron, it would be easy to replace
 6422 the perceptron with a conditional random field. In this case, the online updates would be
 6423 based on feature expectations, which can be computed using the inside-outside algorithm
 6424 (§ 10.6).

$NP_1 \rightarrow NP_2 \ CC \ NP_3$	$(CC.sem @ (NP_2.sem)) @ (NP_3.sem)$	+1
$VP_1 \rightarrow VP_2 \ CC \ VP_3$	$(CC.sem @ (VP_2.sem)) @ (VP_3.sem)$	-1
$NP \rightarrow leaves$	$LEAVES_n$	+1
$VP \rightarrow V_i$	$V_i.sem$	-1
$V_i \rightarrow leaves$	$\lambda x.LEAVES_v$	-1

Table 12.3: Perceptron update for analysis in Figure 12.6 (gold) and Figure 12.7 (predicted)

6425 12.4.2 Learning from logical forms

Complete derivations are expensive to annotate, and are rarely available.¹¹ One solution is to focus on learning from logical forms directly, while treating the derivations as **latent variables** (Zettlemoyer and Collins, 2005). In a conditional probabilistic model over logical forms y and derivations z , we have,

$$p(y, z | w) = \frac{\exp(\theta \cdot f(w, z, y))}{\sum_{y', z'} \exp(\theta \cdot f(w, z', y'))}, \quad [12.20]$$

6426 which is the standard log-linear model, applied to the logical form y and the derivation
6427 z .

Since the derivation z unambiguously determines the logical form y , it may seem silly to model the joint probability over y and z . However, since z is unknown, it can be marginalized out,

$$p(y | w) = \sum_z p(y, z | w). \quad [12.21]$$

The semantic parser can then select the logical form with the maximum log marginal probability,

$$\log \sum_z p(y, z | w) = \log \sum_z \frac{\exp(\theta \cdot f(w, z, y))}{\sum_{y', z'} \exp(\theta \cdot f(w, z', y'))} \quad [12.22]$$

$$\propto \log \sum_z \exp(\theta \cdot f(w, z', y')) \quad [12.23]$$

$$\geq \max_z \theta \cdot f(w, z, y). \quad [12.24]$$

6428 It is impossible to push the log term inside the sum over z , so our usual linear scoring
6429 function does not apply. We can recover this scoring function only in approximation, by
6430 taking the max (rather than the sum) over derivations z , which provides a lower bound.

¹¹An exception is the work of Ge and Mooney (2005), who annotate the meaning of each syntactic constituents for several hundred sentences.

Learning can be performed by maximizing the log marginal likelihood,

$$\ell(\boldsymbol{\theta}) = \sum_{i=1}^N \log p(\mathbf{y}^{(i)} \mid \mathbf{w}^{(i)}; \boldsymbol{\theta}) \quad [12.25]$$

$$= \sum_{i=1}^N \log \sum_z p(\mathbf{y}^{(i)}, \mathbf{z}^{(i)} \mid \mathbf{w}^{(i)}; \boldsymbol{\theta}). \quad [12.26]$$

6431 This log-likelihood is not **convex** in $\boldsymbol{\theta}$, unlike the log-likelihood of a fully-observed conditional random field. This means that learning can give different results depending on the
6432 initialization.
6433

The derivative of Equation 12.26 is,

$$\frac{\partial \ell_i}{\partial \boldsymbol{\theta}} = \sum_z p(z \mid \mathbf{y}, \mathbf{w}; \boldsymbol{\theta}) \mathbf{f}(\mathbf{w}, z, \mathbf{y}) - \sum_{z'} p(z', \mathbf{y}' \mid \mathbf{w}; \boldsymbol{\theta}) \mathbf{f}(\mathbf{w}, z', \mathbf{y}') \quad [12.27]$$

$$= E_{z|\mathbf{y}, \mathbf{w}} \mathbf{f}(\mathbf{w}, z, \mathbf{y}) - E_{y,z|\mathbf{w}} \mathbf{f}(\mathbf{w}, z, \mathbf{y}) \quad [12.28]$$

6434 Both expectations can be computed via bottom-up algorithms like inside-outside. Al-
6435 ternatively, we can again maximize rather than marginalize over derivations for an ap-
6436 proximate solution. In either case, the first term of the gradient requires us to identify
6437 derivations z that are compatible with the logical form \mathbf{y} . This can be done in a bottom-
6438 up dynamic programming algorithm, by having each cell in the table $t[i, j, X]$ include the
6439 set of all possible logical forms for $X \rightsquigarrow \mathbf{w}_{i+1:j}$. The resulting table may therefore be much
6440 larger than in syntactic parsing. This can be controlled by using pruning to eliminate in-
6441 termediate analyses that are incompatible with the final logical form \mathbf{y} (Zettlemoyer and
6442 Collins, 2005), or by using beam search and restricting the size of each cell to some fixed
6443 constant (Liang et al., 2013).

6444 If we replace each expectation in Equation 12.28 with argmax and then apply stochastic
6445 gradient descent to learn the weights, we obtain the **latent variable perceptron**, a simple
6446 and general algorithm for learning with missing data. The algorithm is shown in its most
6447 basic form in Algorithm 16, but the usual tricks such as averaging and margin loss can
6448 be applied (Yu and Joachims, 2009). Aside from semantic parsing, the latent variable
6449 perceptron has been used in tasks such as machine translation (Liang et al., 2006) and
6450 named entity recognition (Sun et al., 2009). In **latent conditional random fields**, we use
6451 the full expectations rather than maximizing over the hidden variable. This model has
6452 also been employed in a range of problems beyond semantic parsing, including parse
6453 reranking (Koo and Collins, 2005) and gesture recognition (Quattoni et al., 2007).

6454 12.4.3 Learning from denotations

Logical forms are easier to obtain than complete derivations, but the annotation of logical forms still requires considerable expertise. However, it is relatively easy to obtain deno-

Algorithm 16 Latent variable perceptron

```

1: procedure LATENTVARIABLEPERCEPTRON( $\mathbf{w}^{(1:N)}, \mathbf{y}^{(1:N)}$ )
2:    $\theta \leftarrow 0$ 
3:   repeat
4:     Select an instance  $i$ 
5:      $\mathbf{z}^{(i)} \leftarrow \text{argmax}_{\mathbf{z}} \theta \cdot \mathbf{f}(\mathbf{w}^{(i)}, \mathbf{z}, \mathbf{y}^{(i)})$ 
6:      $\hat{\mathbf{y}}, \hat{\mathbf{z}} \leftarrow \text{argmax}_{\mathbf{y}', \mathbf{z}'} \theta \cdot \mathbf{f}(\mathbf{w}^{(i)}, \mathbf{z}', \mathbf{y}')$ 
7:      $\theta \leftarrow \theta + \mathbf{f}(\mathbf{w}^{(i)}, \mathbf{z}^{(i)}, \mathbf{y}^{(i)}) - \mathbf{f}(\mathbf{w}^{(i)}, \hat{\mathbf{z}}, \hat{\mathbf{y}})$ 
8:   until tired
9:   return  $\theta$ 

```

tations for many natural language sentences. For example, in the geography domain, the denotation of a question would be its answer (Clarke et al., 2010; Liang et al., 2013):

Text :*What states border Georgia?*
Logical form : $\lambda x.\text{STATE}(x) \wedge \text{BORDER}(x, \text{GEORGIA})$
Denotation :{Alabama, Florida, North Carolina,
South Carolina, Tennessee}

6455 Similarly, in a robotic control setting, the denotation of a command would be an action or
6456 sequence of actions (Artzi and Zettlemoyer, 2013). In both cases, the idea is to reward the
6457 semantic parser for choosing an analysis whose denotation is correct: the right answer to
6458 the question, or the right action.

Learning from logical forms was made possible by summing or maxing over derivations. This idea can be carried one step further, summing or maxing over all logical forms with the correct denotation. Let $v_i(\mathbf{y}) \in \{0, 1\}$ be a **validation function**, which assigns a binary score indicating whether the denotation $[\mathbf{y}]$ for the text $\mathbf{w}^{(i)}$ is correct. We can then learn by maximizing a conditional-likelihood objective,

$$\ell^{(i)}(\boldsymbol{\theta}) = \log \sum_{\mathbf{y}} v_i(\mathbf{y}) \times p(\mathbf{y} \mid \mathbf{w}; \boldsymbol{\theta}) \quad [12.29]$$

$$= \log \sum_{\mathbf{y}} v_i(\mathbf{y}) \times \sum_{\mathbf{z}} p(\mathbf{y}, \mathbf{z} \mid \mathbf{w}; \boldsymbol{\theta}), \quad [12.30]$$

6459 which sums over all derivations \mathbf{z} of all valid logical forms, $\{\mathbf{y} : v_i(\mathbf{y}) = 1\}$. This cor-
6460 responds to the log-probability that the semantic parser produces a logical form with a
6461 valid denotation.

Differentiating with respect to θ , we obtain,

$$\frac{\partial \ell^{(i)}}{\partial \theta} = \sum_{\mathbf{y}, \mathbf{z}: v_i(\mathbf{y})=1} p(\mathbf{y}, \mathbf{z} | \mathbf{w}) \mathbf{f}(\mathbf{w}, \mathbf{z}, \mathbf{y}) - \sum_{\mathbf{y}', \mathbf{z}'} p(\mathbf{y}', \mathbf{z}' | \mathbf{w}) \mathbf{f}(\mathbf{w}, \mathbf{z}', \mathbf{y}'), \quad [12.31]$$

which is the usual difference in feature expectations. The positive term computes the expected feature expectations conditioned on the denotation being valid, while the second term computes the expected feature expectations according to the current model, without regard to the ground truth. Large-margin learning formulations are also possible for this problem. For example, Artzi and Zettlemoyer (2013) generate a set of valid and invalid derivations, and then impose a constraint that all valid derivations should score higher than all invalid derivations. This constraint drives a perceptron-like learning rule.

Additional resources

A key issue not considered here is how to handle **semantic underspecification**: cases in which there are multiple semantic interpretations for a single syntactic structure. Quantifier scope ambiguity is a classic example. Blackburn and Bos (2005) enumerate a number of approaches to this issue, and also provide links between natural language semantics and computational inference techniques. Much of the contemporary research on semantic parsing uses the framework of combinatory categorial grammar (CCG). Carpenter (1997) provides a comprehensive treatment of how CCG can support compositional semantic analysis. Another recent area of research is the semantics of multi-sentence texts. This can be handled with models of **dynamic semantics**, such as dynamic predicate logic (Groenendijk and Stokhof, 1991).

Alternative readings on formal semantics include an “informal” reading from Levy and Manning (2009), and a more involved introduction from Briscoe (2011). To learn more about ongoing research on data-driven semantic parsing, readers may consult the survey article by Liang and Potts (2015), tutorial slides and videos by Artzi and Zettlemoyer (2013),¹² and the source code by Yoav Artzi¹³ and Percy Liang.¹⁴

Exercises

- Derive the **modus ponens** inference rule, which states that if we know $\phi \Rightarrow \psi$ and ϕ , then ψ must be true. The derivation can be performed using the definition of the \Rightarrow operator and some of the laws provided in § 12.2.1, plus one additional identity: $\perp \vee \phi = \phi$.

¹²Videos are currently available at <http://yoavartzi.com/tutorial/>

¹³<http://yoavartzi.com/spf>

¹⁴<https://github.com/percyliang/sempre>

- 6490 2. Convert the following examples into first-order logic, using the relations CAN-SLEEP,
 6491 MAKES-NOISE, and BROTHER.
- 6492 • If Abigail makes noise, no one can sleep.
 6493 • If Abigail makes noise, someone cannot sleep.
 6494 • None of Abigail's brothers can sleep.
 6495 • If one of Abigail's brothers makes noise, Abigail cannot sleep.
- 6496 3. Extend the grammar fragment G_1 to include the ditransitive verb *teaches* and the
 6497 proper noun *Swahili*. Show how to derive the interpretation for the sentence *Alex*
 6498 *teaches Brit Swahili*, which should be $\text{TEACHES}(\text{ALEX}, \text{BRIT}, \text{SWAHILI})$. The grammar
 6499 need not be in Chomsky Normal Form. For the ditransitive verb, use NP_1 and NP_2
 6500 to indicate the two direct objects.
- 6501 4. Derive the semantic interpretation for the sentence *Alex likes every dog*, using gram-
 6502 mar fragment G_2 .
- 6503 5. Extend the grammar fragment G_2 to handle adjectives, so that the meaning of *an
 6504 angry dog* is $\lambda P. \exists x \text{DOG}(x) \wedge \text{ANGRY}(x) \wedge P(x)$. Specifically, you should supply the
 6505 lexical entry for the adjective *angry*, and you should specify the syntactic-semantic
 6506 productions $\text{NP} \rightarrow \text{DET } \text{NOM}$, $\text{NOM} \rightarrow \text{JJ } \text{NOM}$, and $\text{NOM} \rightarrow \text{NN}$.
- 6507 6. Extend your answer to the previous question to cover copula constructions with
 6508 predicative adjectives, such as *Alex is angry*. The interpretation should be $\text{ANGRY}(\text{ALEX})$.
 6509 You should add a verb phrase production $\text{VP} \rightarrow \text{V}_{\text{cop}} \text{ JJ}$, and a terminal production
 6510 $\text{V}_{\text{cop}} \rightarrow \text{is}$. Show why your grammar extensions result in the correct interpretation.
- 6511 7. In Figure 12.6 and Figure 12.7, we treat the plurals *shoots* and *leaves* as entities. Revise
 6512 G_2 so that the interpretation of *Alex eats leaves* is $\forall x. (\text{LEAF}(x) \Rightarrow \text{EATS}(\text{ALEX}, x))$, and
 6513 show the resulting perceptron update.
- 6514 8. Statements like *every student eats a pizza* have two possible interpretations, depend-
 6515 ing on quantifier scope:

$$\forall x \exists y \text{PIZZA}(y) \wedge (\text{STUDENT}(x) \Rightarrow \text{EATS}(x, y)) \quad [12.32]$$

$$\exists y \forall x \text{PIZZA}(y) \wedge (\text{STUDENT}(x) \Rightarrow \text{EATS}(x, y)) \quad [12.33]$$

6514 Explain why these interpretations really are different, and modify the grammar G_2
 6515 so that it can produce both interpretations.

6516 9. Derive Equation 12.27.

6517 10. In the GeoQuery domain, give a natural language query that has multiple plausible
 6518 semantic interpretations with the same denotation. List both interpretaions and the
 6519 denotation.

6520 **Hint:** There are many ways to do this, but one approach involves using toponyms
6521 (place names) that could plausibly map to several different entities in the model.

6522

Chapter 13

6523

Predicate-argument semantics

6524 This chapter considers more “lightweight” semantic representations, which discard some
6525 aspects of first-order logic, but focus on predicate-argument structures. Let’s begin by
6526 thinking about the semantics of events, with a simple example:

6527 (13.1) Asha gives Boyang a book.

6528 A first-order logical representation of this sentence is,

$$\exists x. \text{BOOK}(x) \wedge \text{GIVE}(\text{ASHA}, \text{BOYANG}, x) \quad [13.1]$$

6529 In this representation, we define variable x for the book, and we link the strings *Asha* and
6530 *Boyang* to entities ASHA and BOYANG. Because the action of giving involves a giver, a
6531 recipient, and a gift, the predicate GIVE must take three arguments.

6532 Now suppose we have additional information about the event:

6533 (13.2) Yesterday, Asha reluctantly gave Boyang a book.

6534 One possible solution is to extend the predicate GIVE to take additional arguments,

$$\exists x. \text{BOOK}(x) \wedge \text{GIVE}(\text{ASHA}, \text{BOYANG}, x, \text{YESTERDAY}, \text{RELUCTANTLY}) \quad [13.2]$$

But this is clearly unsatisfactory: *yesterday* and *reluctantly* are optional arguments, and we would need a different version of the GIVE predicate for every possible combination of arguments. **Event semantics** solves this problem by **reifying** the event as an existentially quantified variable e ,

$$\begin{aligned} \exists e, x. & \text{GIVE-EVENT}(e) \wedge \text{GIVER}(e, \text{ASHA}) \wedge \text{GIFT}(e, x) \wedge \text{BOOK}(e, x) \wedge \text{RECIPIENT}(e, \text{BOYANG}) \\ & \wedge \text{TIME}(e, \text{YESTERDAY}) \wedge \text{MANNER}(e, \text{RELUCTANTLY}) \end{aligned}$$

6535 In this way, each argument of the event — the giver, the recipient, the gift — can be rep-
 6536 resented with a relation of its own, linking the argument to the event e . The expression
 6537 GIVER(e , ASHA) says that ASHA plays the **role** of GIVER in the event. This reformulation
 6538 handles the problem of optional information such as the time or manner of the event,
 6539 which are called **adjuncts**. Unlike arguments, adjuncts are not a mandatory part of the
 6540 relation, but under this representation, they can be expressed with additional logical rela-
 6541 tions that are conjoined to the semantic interpretation of the sentence.¹

6542 The event semantic representation can be applied to nested clauses, e.g.,

6543 (13.3) Chris sees Asha pay Boyang.

This is done by using the event variable as an argument:

$$\begin{aligned} \exists e_1 \exists e_2 \text{SEE-EVENT}(e_1) \wedge \text{SEER}(e_1, \text{CHRIS}) \wedge \text{SIGHT}(e_1, e_2) \\ \wedge \text{PAY-EVENT}(e_2) \wedge \text{PAYER}(e_2, \text{ASHA}) \wedge \text{PAYEE}(e_2, \text{BOYANG}) \end{aligned} \quad [13.3]$$

6544 As with first-order logic, the goal of event semantics is to provide a representation that
 6545 generalizes over many surface forms. Consider the following paraphrases of (13.1):

- 6546 (13.4) Asha gives a book to Boyang.
- 6547 (13.5) A book is given to Boyang by Asha.
- 6548 (13.6) A book is given by Asha to Boyang.
- 6549 (13.7) The gift of a book from Asha to Boyang ...

6550 All have the same event semantic meaning as Equation 13.1, but the ways in which the
 6551 meaning can be expressed are diverse. The final example does not even include a verb:
 6552 events are often introduced by verbs, but as shown by (13.7), the noun *gift* can introduce
 6553 the same predicate, with the same accompanying arguments.

6554 **Semantic role labeling** (SRL) is a relaxed form of semantic parsing, in which each
 6555 semantic role is filled by a set of tokens from the text itself. This is sometimes called
 6556 “shallow semantics” because, unlike model-theoretic semantic parsing, role fillers need
 6557 not be symbolic expressions with denotations in some world model. A semantic role
 6558 labeling system is required to identify all predicates, and then specify the spans of text
 6559 that fill each role. To give a sense of the task, here is a more complicated example:

- 6560 (13.8) Boyang wants Asha to give him a linguistics book.

¹This representation is often called **Neo-Davidsonian event semantics**. The use of existentially-quantified event variables was proposed by Davidson (1967) to handle the issue of optional adjuncts. In Neo-Davidsonian semantics, this treatment of adjuncts is extended to mandatory arguments as well (e.g., Parsons, 1990).

6561 In this example, there are two predicates, expressed by the verbs *want* and *give*. Thus, a
6562 semantic role labeler might return the following output:

- 6563 • (PREDICATE : *wants*, WANTED : *Boyang*, DESIRE : *Asha to give him a linguistics book*)
6564 • (PREDICATE : *give*, GIVER : *Asha*, RECIPIENT : *him*, GIFT : *a linguistics book*)

6565 *Boyang* and *him* may refer to the same person, but the semantic role labeling is not re-
6566 quired to resolve this reference. Other predicate-argument representations, such as **Ab-**
6567 **stract Meaning Representation (AMR)**, do require reference resolution. We will return to
6568 AMR in § 13.3, but first, let us further consider the definition of semantic roles.

6569 13.1 Semantic roles

6570 In event semantics, it is necessary to specify a number of additional logical relations to
6571 link arguments to events: GIVER, RECIPIENT, SEER, SIGHT, etc. Indeed, every predicate re-
6572 quires a set of logical relations to express its own arguments. In contrast, adjuncts such as
6573 TIME and MANNER are shared across many types of events. A natural question is whether
6574 it is possible to treat mandatory arguments more like adjuncts, by identifying a set of
6575 generic argument types that are shared across many event predicates. This can be further
6576 motivated by examples involving related verbs:

- 6577 (13.9) Asha gave Boyang a book.
6578 (13.10) Asha loaned Boyang a book.
6579 (13.11) Asha taught Boyang a lesson.
6580 (13.12) Asha gave Boyang a lesson.

6581 The respective roles of Asha, Boyang, and the book are nearly identical across the first
6582 two examples. The third example is slightly different, but the fourth example shows that
6583 the roles of GIVER and TEACHER can be viewed as related.

6584 One way to think about the relationship between roles such as GIVER and TEACHER is
6585 by enumerating the set of properties that an entity typically possesses when it fulfills these
6586 roles: givers and teachers are usually **animate** (they are alive and sentient) and **volitional**
6587 (they choose to enter into the action).² In contrast, the thing that gets loaned or taught is
6588 usually not animate or volitional; furthermore, it is unchanged by the event.

6589 Building on these ideas, **thematic roles** generalize across predicates by leveraging the
6590 shared semantic properties of typical role fillers (Fillmore, 1968). For example, in exam-
6591 ples (13.9-13.12), Asha plays a similar role in all four sentences, which we will call the

²There are always exceptions. For example, in the sentence *The C programming language has taught me a lot about perseverance*, the “teacher” is the *The C programming language*, which is presumably not animate or volitional.

	<i>Asha</i>	<i>gave</i>	<i>Boyang</i>	<i>a book</i>
VerbNet	AGENT		RECIPIENT	THEME
PropBank	ARG0: giver		ARG2: entity given to	ARG1: thing given
FrameNet	DONOR		RECIPIENT	THEME
	<i>Asha</i>	<i>taught</i>	<i>Boyang</i>	<i>algebra</i>
VerbNet	AGENT		RECIPIENT	TOPIC
PropBank	ARG0: teacher		ARG2: student	ARG1: subject
FrameNet	TEACHER		STUDENT	SUBJECT

Figure 13.1: Example semantic annotations according to VerbNet, PropBank, and FrameNet

6592 **agent.** This reflects several shared semantic properties: she is the one who is actively and
 6593 intentionally performing the action, while Boyang is a more passive participant; the book
 6594 and the lesson would play a different role, as non-animate participants in the event.

6595 Example annotations from three well known systems are shown in Figure 13.1. We
 6596 will now discuss these systems in more detail.

6597 13.1.1 VerbNet

6598 **VerbNet** (Kipper-Schuler, 2005) is a lexicon of verbs, and it includes thirty “core” thematic
 6599 roles played by arguments to these verbs. Here are some example roles, accompanied by
 6600 their definitions from the VerbNet Guidelines.³

- 6601 • AGENT: “ACTOR in an event who initiates and carries out the event intentionally or
 6602 consciously, and who exists independently of the event.”
- 6603 • PATIENT: “UNDERGOER in an event that experiences a change of state, location or
 6604 condition, that is causally involved or directly affected by other participants, and
 6605 exists independently of the event.”
- 6606 • RECIPIENT: “DESTINATION that is animate”
- 6607 • THEME: “UNDERGOER that is central to an event or state that does not have control
 6608 over the way the event occurs, is not structurally changed by the event, and/or is
 6609 characterized as being in a certain position or condition throughout the state.”
- 6610 • TOPIC: “THEME characterized by information content transferred to another partic-
 6611 ipant.”

³http://verbs.colorado.edu/verb-index/VerbNet_Guidelines.pdf

6612 VerbNet roles are organized in a hierarchy, so that a TOPIC is a type of THEME, which in
 6613 turn is a type of UNDERGOER, which is a type of PARTICIPANT, the top-level category.

6614 In addition, VerbNet organizes verb senses into a class hierarchy, in which verb senses
 6615 that have similar meanings are grouped together. Recall from § 4.2 that multiple meanings
 6616 of the same word are called **senses**, and that WordNet identifies senses for many English
 6617 words. VerbNet builds on WordNet, so that verb classes are identified by the WordNet
 6618 senses of the verbs that they contain. For example, the verb class give-13.1 includes
 6619 the first WordNet sense of *loan* and the second WordNet sense of *lend*.

6620 Each VerbNet class or subclass takes a set of thematic roles. For example, give-13.1
 6621 takes arguments with the thematic roles of AGENT, THEME, and RECIPIENT;⁴ the pred-
 6622 icate TEACH takes arguments with the thematic roles AGENT, TOPIC, RECIPIENT, and
 6623 SOURCE.⁵ So according to VerbNet, *Asha* and *Boyang* play the roles of AGENT and RECIP-
 6624 IENT in the sentences,

6625 (13.13) Asha gave Boyang a book.

6626 (13.14) Asha taught Boyang algebra.

6627 The *book* and *algebra* are both THEMES, but *algebra* is a subcategory of THEME — a TOPIC
 6628 — because it consists of information content that is given to the receiver.

6629 13.1.2 Proto-roles and PropBank

6630 Detailed thematic role inventories of the sort used in VerbNet are not universally accepted.
 6631 For example, Dowty (1991, pp. 547) notes that “Linguists have often found it hard to agree
 6632 on, and to motivate, the location of the boundary between role types.” He argues that a
 6633 solid distinction can be identified between just two **proto-roles**:

6634 **Proto-Agent.** Characterized by volitional involvement in the event or state; sentience
 6635 and/or perception; causing an event or change of state in another participant; move-
 6636 ment; exists independently of the event.

6637 **Proto-Patient.** Undergoes change of state; causally affected by another participant; sta-
 6638 tionary relative to the movement of another participant; does not exist indepen-
 6639 dently of the event.⁶

⁴<https://verbs.colorado.edu/verb-index/vn/give-13.1.php>

⁵https://verbs.colorado.edu/verb-index/vn/transfer_mesg-37.1.1.php

⁶Reisinger et al. (2015) ask crowd workers to annotate these properties directly, finding that annotators tend to agree on the properties of each argument. They also find that in English, arguments having more proto-agent properties tend to appear in subject position, while arguments with more proto-patient properties appear in object position.

6640 In the examples in Figure 13.1, Asha has most of the proto-agent properties: in giving
 6641 the book to Boyang, she is acting volitionally (as opposed to *Boyang got a book from Asha*, in
 6642 which it is not clear whether Asha gave up the book willingly); she is sentient; she causes
 6643 a change of state in Boyang; she exists independently of the event. Boyang has some
 6644 proto-agent properties: he is sentient and exists independently of the event. But he also
 6645 some proto-patient properties: he is the one who is causally affected and who undergoes
 6646 change of state. The book that Asha gives Boyang has even fewer of the proto-agent
 6647 properties: it is not volitional or sentient, and it has no causal role. But it also lacks many
 6648 of the proto-patient properties: it does not undergo change of state, exists independently
 6649 of the event, and is not stationary.

6650 The **Proposition Bank**, or PropBank (Palmer et al., 2005), builds on this basic agent-
 6651 patient distinction, as a middle ground between generic thematic roles and roles that are
 6652 specific to each predicate. Each verb is linked to a list of numbered arguments, with ARG0
 6653 as the proto-agent and ARG1 as the proto-patient. Additional numbered arguments are
 6654 verb-specific. For example, for the predicate TEACH,⁷ the arguments are:

- 6655 • ARG0: the teacher
- 6656 • ARG1: the subject
- 6657 • ARG2: the student(s)

6658 Verbs may have any number of arguments: for example, WANT and GET have five, while
 6659 EAT has only ARG0 and ARG1. In addition to the semantic arguments found in the frame
 6660 files, roughly a dozen general-purpose **adjuncts** may be used in combination with any
 6661 verb. These are shown in Table 13.1.

6662 PropBank-style semantic role labeling is annotated over the entire Penn Treebank. This
 6663 annotation includes the sense of each verbal predicate, as well as the argument spans.

6664 13.1.3 FrameNet

6665 Semantic **frames** are descriptions of situations or events. Frames may be *evoked* by one
 6666 of their **lexical units** (often a verb, but not always), and they include some number of
 6667 **frame elements**, which are like roles (Fillmore, 1976). For example, the act of teaching
 6668 is a frame, and can be evoked by the verb *taught*; the associated frame elements include
 6669 the teacher, the student(s), and the subject being taught. Frame semantics has played a
 6670 significant role in the history of artificial intelligence, in the work of Minsky (1974) and
 6671 Schank and Abelson (1977). In natural language processing, the theory of frame semantics
 6672 has been implemented in **FrameNet** (Fillmore and Baker, 2009), which consists of a lexicon

⁷<http://verbs.colorado.edu/propbank/framesets-english-aliases/teach.html>

TMP	time	<i>Boyang ate a bagel</i> [AM-TMP <i>yesterday</i>].
LOC	location	<i>Asha studies in</i> [AM-LOC <i>Stuttgart</i>]
MOD	modal verb	<i>Asha</i> [AM-MOD <i>will</i>] <i>study in Stuttgart</i>
ADV	general purpose	[AM-ADV <i>Luckily</i>], <i>Asha knew algebra</i> .
MNR	manner	<i>Asha ate</i> [AM-MNR <i>aggressively</i>].
DIS	discourse connective	[AM-DIS <i>However</i>], <i>Asha prefers algebra</i> .
PRP	purpose	<i>Barry studied</i> [AM-PRP <i>to pass the bar</i>].
DIR	direction	<i>Workers dumped burlap sacks</i> [AM-DIR <i>into a bin</i>].
NEG	negation	<i>Asha does</i> [AM-NEG <i>not</i>] <i>speak Albanian</i> .
EXT	extent	<i>Prices increased</i> [AM-EXT <i>4%</i>].
CAU	cause	<i>Boyang returned the book</i> [AM-CAU <i>because it was overdue</i>].

Table 13.1: PropBank adjuncts (Palmer et al., 2005), sorted by frequency in the corpus

6673 of roughly 1000 frames, and a corpus of more than 200,000 “exemplar sentences,” in which
 6674 the frames and their elements are annotated.⁸

6675 Rather than seeking to link semantic roles such as TEACHER and GIVER into the-
 6676 matic roles such as AGENT, FrameNet aggressively groups verbs into frames, and links
 6677 semantically-related roles across frames. For example, the following two sentences would
 6678 be annotated identically in FrameNet:

6679 (13.15) Asha taught Boyang algebra.

6680 (13.16) Boyang learned algebra from Asha.

6681 This is because *teach* and *learn* are both lexical units in the EDUCATION-TEACHING frame.
 6682 Furthermore, roles can be shared even when the frames are distinct, as in the following
 6683 two examples:

6684 (13.17) Asha gave Boyang a book.

6685 (13.18) Boyang got a book from Asha.

6686 The GIVING and GETTING frames both have RECIPIENT and THEME elements, so Boyang
 6687 and the book would play the same role. Asha’s role is different: she is the DONOR in the
 6688 GIVING frame, and the SOURCE in the GETTING frame. FrameNet makes extensive use of
 6689 multiple inheritance to share information across frames and frame elements: for example,
 6690 the COMMERCE-SELL and LENDING frames inherit from GIVING frame.

⁸Current details and data can be found at <https://framenet.icsi.berkeley.edu/>

6691 **13.2 Semantic role labeling**

6692 The task of semantic role labeling is to identify the parts of the sentence comprising the
 6693 semantic roles. In English, this task is typically performed on the PropBank corpus, with
 6694 the goal of producing outputs in the following form:

6695 (13.19) [ARG0 Asha] [GIVE.01 gave] [ARG2 Boyang's mom] [ARG1 a book] [AM-TMP yesterday].

6696 Note that a single sentence may have multiple verbs, and therefore a given word may be
 6697 part of multiple role-fillers:

6698 (13.20) [ARG0 Asha] [WANT.01 wanted]
 Asha wanted

6699 [ARG1 Boyang to give her the book].
 [ARG0 Boyang] [GIVE.01 to give] [ARG2 her] [ARG1 the book].

6700 **13.2.1 Semantic role labeling as classification**

6701 PropBank is annotated on the Penn Treebank, and annotators used phrasal constituents
 6702 (§ 9.2.2) to fill the roles. PropBank semantic role labeling can be viewed as the task of as-
 6703 signing to each phrase a label from the set $\mathcal{R} = \{\emptyset, \text{PRED}, \text{ARG0}, \text{ARG1}, \text{ARG2}, \dots, \text{AM-LOC}, \text{AM-TMP}, \dots\}$
 6704 with respect to each predicate. If we treat semantic role labeling as a classification prob-
 6705 lem, we obtain the following functional form:

$$\hat{y}_{(i,j)} = \underset{y}{\operatorname{argmax}} \psi(\mathbf{w}, y, i, j, \rho, \tau), \quad [13.4]$$

6706 where,

- 6707 • (i, j) indicates the span of a phrasal constituent $(w_{i+1}, w_{i+2}, \dots, w_j)$;⁹
- 6708 • \mathbf{w} represents the sentence as a sequence of tokens;
- 6709 • ρ is the index of the predicate verb in \mathbf{w} ;
- 6710 • τ is the structure of the phrasal constituent parse of \mathbf{w} .

6711 Early work on semantic role labeling focused on discriminative feature-based models,
 6712 where $\psi(\mathbf{w}, y, i, j, \rho, \tau) = \theta \cdot f(\mathbf{w}, y, i, j, \rho, \tau)$. Table 13.2 shows the features used in a sem-
 6713 inal paper on FrameNet semantic role labeling (Gildea and Jurafsky, 2002). By 2005 there

⁹PropBank roles can also be filled by **split constituents**, which are discontinuous spans of text. This situation most frequently in reported speech, e.g. [ARG1 *By addressing these problems*], *Mr. Maxwell said*, [ARG1 *the new funds have become extremely attractive.*] (example adapted from Palmer et al., 2005). This issue is typically addressed by defining “continuation arguments”, e.g. C-ARG1, which refers to the continuation of ARG1 after the split.

Predicate lemma and POS tag	The lemma of the predicate verb and its part-of-speech tag
Voice	Whether the predicate is in active or passive voice, as determined by a set of syntactic patterns for identifying passive voice constructions
Phrase type	The constituent phrase type for the proposed argument in the parse tree, e.g. NP, PP
Headword and POS tag	The head word of the proposed argument and its POS tag, identified using the Collins (1997) rules
Position	Whether the proposed argument comes before or after the predicate in the sentence
Syntactic path	The set of steps on the parse tree from the proposed argument to the predicate (described in detail in the text)
Subcategorization	The syntactic production from the first branching node above the predicate. For example, in Figure 13.2, the subcategorization feature around <i>taught</i> would be VP → VBD NP PP.

Table 13.2: Features used in semantic role labeling by Gildea and Jurafsky (2002).

6714 were several systems for PropBank semantic role labeling, and their approaches and fea-
 6715 ture sets are summarized by Carreras and Márquez (2005). Typical features include: the
 6716 phrase type, head word, part-of-speech, boundaries, and neighbors of the proposed argu-
 6717 ment $w_{i+1:j}$; the word, lemma, part-of-speech, and voice of the verb w_ρ (active or passive),
 6718 as well as features relating to its frameset; the distance and path between the verb and
 6719 the proposed argument. In this way, semantic role labeling systems are high-level “con-
 6720 sumers” in the NLP stack, using features produced from lower-level components such as
 6721 part-of-speech taggers and parsers. More comprehensive feature sets are enumerated by
 6722 Das et al. (2014) and Täckström et al. (2015).

6723 A particularly powerful class of features relate to the **syntactic path** between the ar-
 6724 gument and the predicate. These features capture the sequence of moves required to get
 6725 from the argument to the verb by traversing the phrasal constituent parse of the sentence.
 6726 The idea of these features is to capture syntactic regularities in how various arguments
 6727 are realized. Syntactic path features are best illustrated by example, using the parse tree
 6728 in Figure 13.2:

- 6729 • The path from *Asha* to the verb *taught* is NNP↑NP↑S↓VP↓VBD. The first part of
 6730 the path, NNP↑NP↑S, means that we must travel up the parse tree from the NNP
 6731 tag (proper noun) to the S (sentence) constituent. The second part of the path,
 6732 S↓VP↓VBD, means that we reach the verb by producing a VP (verb phrase) from

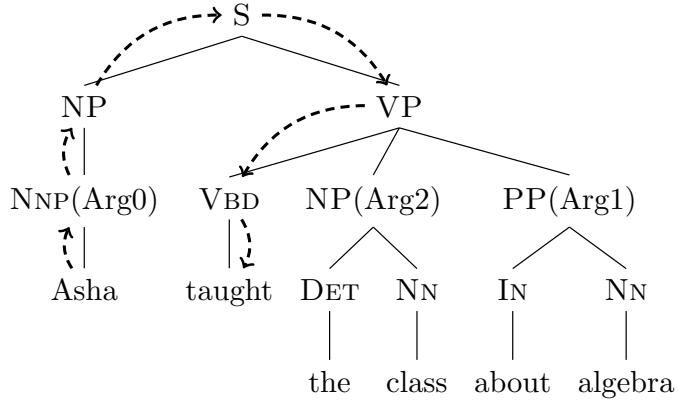


Figure 13.2: Semantic role labeling on the phrase-structure parse tree for a sentence. The dashed line indicates the syntactic path from *Asha* to the predicate verb *taught*.

6733 the S constituent, and then by producing a VBD (past tense verb). This feature is
 6734 consistent with *Asha* being in subject position, since the path includes the sentence
 6735 root S.

- 6736 • The path from *the class* to *taught* is NP↑VP↓VBD. This is consistent with *the class*
 6737 being in object position, since the path passes through the VP node that dominates
 6738 the verb *taught*.

6739 Because there are many possible path features, it can also be helpful to look at smaller
 6740 parts: for example, the upward and downward parts can be treated as separate features;
 6741 another feature might consider whether S appears anywhere in the path.

6742 Rather than using the constituent parse, it is also possible to build features from the
 6743 **dependency path** between the head word of each argument and the verb (Pradhan et al.,
 6744 2005). Using the Universal Dependency part-of-speech tagset and dependency relations (Nivre
 6745 et al., 2016), the dependency path from *Asha* to *taught* is PROPN $\xleftarrow[\text{NSUBJ}]{} \text{VERB}$, because *taught*
 6746 is the head of a relation of type $\xleftarrow[\text{NSUBJ}]{} \text{VERB}$ with *Asha*. Similarly, the dependency path from *class*
 6747 to *taught* is NOUN $\xleftarrow[\text{DOBJ}]{} \text{VERB}$, because *class* heads the noun phrase that is a direct object of
 6748 *taught*. A more interesting example is *Asha wanted to teach the class*, where the path from
 6749 *Asha* to *teach* is PROPN $\xleftarrow[\text{NSUBJ}]{} \text{VERB} \rightarrow[\text{XCOMP}] \text{VERB}$. The right-facing arrow in second relation
 6750 indicates that *wanted* is the head of its XCOMP relation with *teach*.

6751 **13.2.2 Semantic role labeling as constrained optimization**

6752 A potential problem with treating SRL as a classification problem is that there are a num-
 6753 ber of sentence-level **constraints**, which a classifier might violate.

- 6754 • For a given verb, there can be only one argument of each type (ARG0, ARG1, etc.)
 6755 • Arguments cannot overlap. This problem arises when we are labeling the phrases
 6756 in a constituent parse tree, as shown in Figure 13.2: if we label the PP *about algebra*
 6757 as an argument or adjunct, then its children *about* and *algebra* must be labeled as \emptyset .
 6758 The same constraint also applies to the syntactic ancestors of this phrase.

6759 These constraints introduce dependencies across labeling decisions. In structure pre-
 6760 diction problems such as sequence labeling and parsing, such dependencies are usually
 6761 handled by defining a scoring over the entire structure, \mathbf{y} . Efficient inference requires
 6762 that the global score decomposes into local parts: for example, in sequence labeling, the
 6763 scoring function decomposes into scores of pairs of adjacent tags, permitting the applica-
 6764 tion of the Viterbi algorithm for inference. But the constraints that arise in semantic role
 6765 labeling are less amenable to local decomposition.¹⁰ We therefore consider **constrained**
 6766 **optimization** as an alternative solution.

Let the set $\mathcal{C}(\tau)$ refer to all labelings that obey the constraints introduced by the parse τ . The semantic role labeling problem can be reformulated as a constrained optimization over $\mathbf{y} \in \mathcal{C}(\tau)$,

$$\begin{aligned} \max_{\mathbf{y}} \quad & \sum_{(i,j) \in \tau} \psi(\mathbf{w}, y_{i,j}, i, j, \rho, \tau) \\ \text{s.t. } \quad & \mathbf{y} \in \mathcal{C}(\tau). \end{aligned} \quad [13.5]$$

6767 In this formulation, the objective (shown on the first line) is a separable function of each
 6768 individual labeling decision, but the constraints (shown on the second line) apply to the
 6769 overall labeling. The sum $\sum_{(i,j) \in \tau}$ indicates that we are summing over all constituent
 6770 spans in the parse τ . The expression s.t. in the second line means that we maximize the
 6771 objective *subject to* the constraint $\mathbf{y} \in \mathcal{C}(\tau)$.

6772 A number of practical algorithms exist for restricted forms of constrained optimiza-
 6773 tion. One such restricted form is **integer linear programming**, in which the objective and
 6774 constraints are linear functions of integer variables. To formulate SRL as an integer linear
 6775 program, we begin by rewriting the labels as a set of binary variables $\mathbf{z} = \{z_{i,j,r}\}$ (Pun-
 6776 yakanok et al., 2008),

$$z_{i,j,r} = \begin{cases} 1, & y_{i,j} = r \\ 0, & \text{otherwise,} \end{cases} \quad [13.6]$$

¹⁰Dynamic programming solutions have been proposed by Tromble and Eisner (2006) and Täckström et al. (2015), but they involve creating a trellis structure whose size is exponential in the number of labels.

6777 where $r \in \mathcal{R}$ is a label in the set $\{\text{ARG0}, \text{ARG1}, \dots, \text{AM-LOC}, \dots, \emptyset\}$. Thus, the variables
 6778 \mathbf{z} are a binarized version of the semantic role labeling \mathbf{y} .

The objective can then be formulated as a linear function of \mathbf{z} .

$$\sum_{(i,j) \in \tau} \psi(\mathbf{w}, y_{i,j}, i, j, \rho, \tau) = \sum_{i,j,r} \psi(\mathbf{w}, r, i, j, \rho, \tau) \times z_{i,j,r}, \quad [13.7]$$

6779 which is the sum of the scores of all relations, as indicated by $z_{i,j,r}$.

Constraints Integer linear programming permits linear inequality constraints, of the general form $\mathbf{A}\mathbf{z} \leq \mathbf{b}$, where the parameters \mathbf{A} and \mathbf{b} define the constraints. To make this more concrete, let's start with the constraint that each non-null role type can occur only once in a sentence. This constraint can be written,

$$\forall r \neq \emptyset, \quad \sum_{(i,j) \in \tau} z_{i,j,r} \leq 1. \quad [13.8]$$

6780 Recall that $z_{i,j,r} = 1$ iff the span (i, j) has label r ; this constraint says that for each possible
 6781 label $r \neq \emptyset$, there can be at most one (i, j) such that $z_{i,j,r} = 1$. Rewriting this constraint
 6782 can be written in the form $\mathbf{A}\mathbf{z} \leq \mathbf{b}$, as you will find if you complete the exercises at the
 6783 end of the chapter.

Now consider the constraint that labels cannot overlap. Let's define the convenience function $o((i, j), (i', j')) = 1$ iff (i, j) overlaps (i', j') , and zero otherwise. Thus, o will indicate if a constituent (i', j') is either an ancestor or descendant of (i, j) . The constraint is that if two constituents overlap, only one can have a non-null label:

$$\forall (i, j) \in \tau, \quad \sum_{(i', j') \in \tau} \sum_{r \neq \emptyset} o((i, j), (i', j')) \times z_{i',j',r} \leq 1, \quad [13.9]$$

6784 where $o((i, j), (i, j)) = 1$.

In summary, the semantic role labeling problem can thus be rewritten as the following integer linear program,

$$\max_{\mathbf{z} \in \{0,1\}^{|\tau|}} \quad \sum_{(i,j) \in \tau} \sum_{r \in \mathcal{R}} z_{i,j,r} \psi_{i,j,r} \quad [13.10]$$

$$s.t. \quad \forall r \neq \emptyset, \quad \sum_{(i,j) \in \tau} z_{i,j,r} \leq 1. \quad [13.11]$$

$$\forall (i, j) \in \tau, \quad \sum_{(i', j') \in \tau} \sum_{r \neq \emptyset} o((i, j), (i', j')) \times z_{i',j',r} \leq 1. \quad [13.12]$$

6785 **Learning with constraints** Learning can be performed in the context of constrained op-
 6786 timization using the usual perceptron or large-margin classification updates. Because
 6787 constrained inference is generally more time-consuming, a key question is whether it is
 6788 necessary to apply the constraints during learning. Chang et al. (2008) find that better per-
 6789 formance can be obtained by learning *without* constraints, and then applying constraints
 6790 only when using the trained model to predict semantic roles for unseen data.

6791 **How important are the constraints?** Das et al. (2014) find that an unconstrained, classification-
 6792 based method performs nearly as well as constrained optimization for FrameNet parsing;
 6793 while it commits many violations of the “no-overlap” constraint, the overall F_1 score is
 6794 less than one point worse than the score at the constrained optimum. Similar results
 6795 were obtained for PropBank semantic role labeling by Punyakanok et al. (2008). He et al.
 6796 (2017) find that constrained inference makes a bigger impact if the constraints are based
 6797 on manually-labeled “gold” syntactic parses. This implies that errors from the syntac-
 6798 tic parser may limit the effectiveness of the constraints. Punyakanok et al. (2008) hedge
 6799 against parser error by including constituents from several different parsers; any con-
 6800 stituent can be selected from any parse, and additional constraints ensure that overlap-
 6801 ping constituents are not selected.

6802 **Implementation** Integer linear programming solvers such as `glpk`,¹¹ `cplex`,¹² and `Gurobi`¹³
 6803 allow inequality constraints to be expressed directly in the problem definition, rather than
 6804 in the matrix form $\mathbf{A}z \leq b$. The time complexity of integer linear programming is theoreti-
 6805 cally exponential in the number of variables $|z|$, but in practice these off-the-shelf solvers
 6806 obtain good solutions efficiently. Das et al. (2014) report that the `cplex` solver requires 43
 6807 seconds to perform inference on the FrameNet test set, which contains 4,458 predicates.

6808 Recent work has shown that many constrained optimization problems in natural lan-
 6809 guage processing can be solved in a highly parallelized fashion, using optimization tech-
 6810 niques such as **dual decomposition**, which are capable of exploiting the underlying prob-
 6811 lem structure (Rush et al., 2010). Das et al. (2014) apply this technique to FrameNet se-
 6812 mantic role labeling, obtaining an order-of-magnitude speedup over `cplex`.

6813 13.2.3 Neural semantic role labeling

6814 Neural network approaches to SRL have tended to treat it as a sequence labeling task,
 6815 using a labeling scheme such as the **BIO notation**, which we previously saw in named
 6816 entity recognition (§ 8.3). In this notation, the first token in a span of type ARG1 is labeled

¹¹<https://www.gnu.org/software/glpk/>

¹²<https://www-01.ibm.com/software/commerce/optimization/cplex-optimizer/>

¹³<http://www.gurobi.com/>

6817 B-ARG1; all remaining tokens in the span are *inside*, and are therefore labeled I-ARG1.
 6818 Tokens outside any argument are labeled O. For example:

- 6819 (13.21) *Asha taught Boyang 's mom about algebra*
 B-ARG0 PRED B-ARG2 I-ARG2 I-ARG2 B-ARG1 I-ARG1

Recurrent neural networks are a natural approach to this tagging task. For example, Zhou and Xu (2015) apply a deep bidirectional multilayer LSTM (see § 7.6) to PropBank semantic role labeling. In this model, each bidirectional LSTM serves as input for another, higher-level bidirectional LSTM, allowing complex non-linear transformations of the original input embeddings, $\mathbf{X} = [\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_M]$. The hidden state of the final LSTM is $\mathbf{Z}^{(K)} = [\mathbf{z}_1^{(K)}, \mathbf{z}_2^{(K)}, \dots, \mathbf{z}_M^{(K)}]$. The “emission” score for each tag $Y_m = y$ is equal to the inner product $\theta_y \cdot \mathbf{z}_m^{(K)}$, and there is also a transition score for each pair of adjacent tags. The complete model can be written,

$$\mathbf{Z}^{(1)} = \text{BiLSTM}(\mathbf{X}) \quad [13.13]$$

$$\mathbf{Z}^{(i)} = \text{BiLSTM}(\mathbf{Z}^{(i-1)}) \quad [13.14]$$

$$\hat{\mathbf{y}} = \underset{\mathbf{y}}{\text{argmax}} \sum_{m=1}^M \Theta^{(y)} \mathbf{z}_m^{(K)} + \psi_{y_{m-1}, y_m}. \quad [13.15]$$

6820 Note that the final step maximizes over the entire labeling \mathbf{y} , and includes a score for
 6821 each tag transition ψ_{y_{m-1}, y_m} . This combination of LSTM and pairwise potentials on tags
 6822 is an example of an **LSTM-CRF**. The maximization over \mathbf{y} is performed by the Viterbi
 6823 algorithm.

6824 This model strongly outperformed alternative approaches at the time, including con-
 6825 strained decoding and convolutional neural networks.¹⁴ More recent work has combined
 6826 recurrent neural network models with constrained decoding, using the A^* search algo-
 6827 rithm to search over labelings that are feasible with respect to the constraints (He et al.,
 6828 2017). This yields small improvements over the method of Zhou and Xu (2015). He et al.
 6829 (2017) obtain larger improvements by creating an **ensemble** of SRL systems, each trained
 6830 on an 80% subsample of the corpus. The average prediction across this ensemble is more
 6831 robust than any individual model.

6832 13.3 Abstract Meaning Representation

6833 Semantic role labeling transforms the task of semantic parsing to a labeling task. Consider
 6834 the sentence,

¹⁴The successful application of convolutional neural networks to semantic role labeling by Collobert and Weston (2008) was an influential early result in the most recent wave of neural networks in natural language processing.

```
(w / want-01
  :ARG0 (b / boy)
  :ARG1 (g / go-02
    :ARG0 b))
```

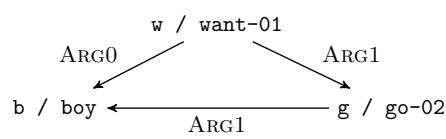


Figure 13.3: Two views of the AMR representation for the sentence *The boy wants to go.*

6835 (13.22) The boy wants to go.

6836 The PropBank semantic role labeling analysis is:

- 6837 • (PREDICATE : *wants*, ARG0 : *the boy*, ARG1 : *to go*)
- 6838 • (PREDICATE : *go*, ARG1 : *the boy*)

6839 The **Abstract Meaning Representation (AMR)** unifies this analysis into a graph structure, in which each node is a **variable**, and each edge indicates a **concept** (Banarescu et al., 2013). This can be written in two ways, as shown in Figure 13.3. On the left is the 6840 PENMAN notation (Matthiessen and Bateman, 1991), in which each set of parentheses 6841 introduces a variable. Each variable is an **instance** of a concept, which is indicated with 6842 the slash notation: for example, *w / want-01* indicates that the variable *w* is an instance 6843 of the concept *want-01*, which in turn refers to the PropBank frame for the first sense 6844 of the verb *want*. Relations are introduced with colons: for example, *:ARG0 (b / boy)* 6845 indicates a relation of type ARG0 with the newly-introduced variable *b*. Variables can be 6846 reused, so that when the variable *b* appears again as an argument to *g*, it is understood to 6847 refer to the same boy in both cases. This arrangement is indicated compactly in the graph 6848 structure on the right, with edges indicating concepts.

6849 One way in which AMR differs from PropBank-style semantic role labeling is that it 6850 reifies each entity as a variable: for example, *the boy* in (13.22) is reified in the variable 6851 *b*, which is reused as ARG0 in its relationship with *w / want-01*, and as ARG1 in its 6852 relationship with *g / go-02*. Reifying entities as variables also makes it possible to 6853 represent the substructure of noun phrases more explicitly. For example, *Asha borrowed* 6854 *the algebra book* would be represented as:

```
6857 (b / borrow-01
  :ARG0 (p / person
    :name (n / name
      :op1 "Asha"))
  :ARG1 (b2 / book
    :topic (a / algebra)))
```

6863 This indicates that the variable *p* is a person, whose name is the variable *n*; that name
 6864 has one token, the string *Asha*. Similarly, the variable *b2* is a book, and the topic of *b2*
 6865 is a variable *a* whose type is algebra. The relations name and topic are examples of
 6866 **non-core roles**, which are similar to adjunct modifiers in PropBank. However, AMR’s
 6867 inventory is more extensive, including more than 70 non-core roles, such as negation,
 6868 time, manner, frequency, and location. Lists and sequences — such as the list of tokens in
 6869 a name — are described using the roles *op1*, *op2*, etc.

6870 Another feature of AMR is that a semantic predicate can be introduced by any syntac-
 6871 tic element, as in the following examples from Banarescu et al. (2013):

- 6872 (13.23) The boy destroyed the room.
- 6873 (13.24) the destruction of the room by the boy ...
- 6874 (13.25) the boy’s destruction of the room ...

6875 All these examples have the same semantics in AMR,

```
6876 (d / destroy-01
6877   :ARG0 (b / boy)
6878   :ARG1 (r / room))
```

6879 The noun *destruction* is linked to the verb *destroy*, which is captured by the PropBank
 6880 frame *destroy-01*. This can happen with adjectives as well: in the phrase *the attractive*
 6881 *spy*, the adjective *attractive* is linked to the PropBank frame *attract-01*:

```
6882 (s / spy
6883   :ARG0-of (a / attract-01))
```

6884 In this example, *ARG0-of* is an **inverse relation**, indicating that *s* is the *ARG0* of the
 6885 predicate *a*. Inverse relations make it possible for all AMR parses to have a single root
 6886 concept, which should be the **focus** of the utterance.

6887 While AMR goes farther than semantic role labeling, it does not link semantically-
 6888 related frames such as *buy/sell* (as FrameNet does), does not handle quantification (as
 6889 first-order predicate calculus does), and makes no attempt to handle noun number and
 6890 verb tense (as PropBank does). A recent survey by Abend and Rappoport (2017) situ-
 6891 ates AMR with respect to several other semantic representation schemes. Other linguistic
 6892 features of AMR are summarized in the original paper (Banarescu et al., 2013) and the
 6893 tutorial slides by Schneider et al. (2015).

6894 13.3.1 AMR Parsing

6895 Abstract Meaning Representation is not a labeling of the original text — unlike PropBank
6896 semantic role labeling, and most of the other tagging and parsing tasks that we have
6897 encountered thus far. The AMR for a given sentence may include multiple concepts for
6898 single words in the sentence: as we have seen, the sentence *Asha likes algebra* contains both
6899 person and name concepts for the word *Asha*. Conversely, words in the sentence may not
6900 appear in the AMR: in *Boyang made a tour of campus*, the **light verb** *make* would not appear
6901 in the AMR, which would instead be rooted on the predicate *tour*. As a result, AMR
6902 is difficult to parse, and even evaluating AMR parsing involves considerable algorithmic
6903 complexity (Cai and Yates, 2013).

6904 A further complexity is that AMR labeled datasets do not explicitly show the **align-**
6905 **ment** between the AMR annotation and the words in the sentence. For example, the link
6906 between the word *wants* and the concept *want-01* is not annotated. To acquire train-
6907 ing data for learning-based parsers, it is therefore necessary to first perform an alignment
6908 between the training sentences and their AMR parses. Flanigan et al. (2014) introduce a
6909 rule-based parser, which links text to concepts through a series of increasingly high-recall
6910 steps.

6911 **Graph-based parsing** One family of approaches to AMR parsing is similar to the graph-
6912 based methods that we encountered in syntactic dependency parsing (chapter 11). For
6913 these systems (Flanigan et al., 2014), parsing is a two-step process:

- 6914 1. **Concept identification** (Figure 13.4a). This involves constructing concept subgraphs
6915 for individual words or spans of adjacent words. For example, in the sentence,
6916 *Asha likes algebra*, we would hope to identify the minimal subtree including just the
6917 concept *like-01* for the word *like*, and the subtree (*p / person :name (n /*
6918 *name :op1 Asha)*) for the word *Asha*.
- 6919 2. **Relation identification** (Figure 13.4b). This involves building a directed graph over
6920 the concepts, where the edges are labeled by the relation type. AMR imposes a
6921 number of constraints on the graph: all concepts must be included, the graph must
6922 be **connected** (there must be a path between every pair of nodes in the undirected
6923 version of the graph), and every node must have at most one outgoing edge of each
6924 type.

6925 Both of these problems are solved by structure prediction. Concept identification re-
6926 quires simultaneously segmenting the text into spans, and labeling each span with a graph
6927 fragment containing one or more concepts. This is done by computing a set of features
6928 for each candidate span *s* and concept labeling *c*, and then returning the labeling with the
6929 highest overall score.

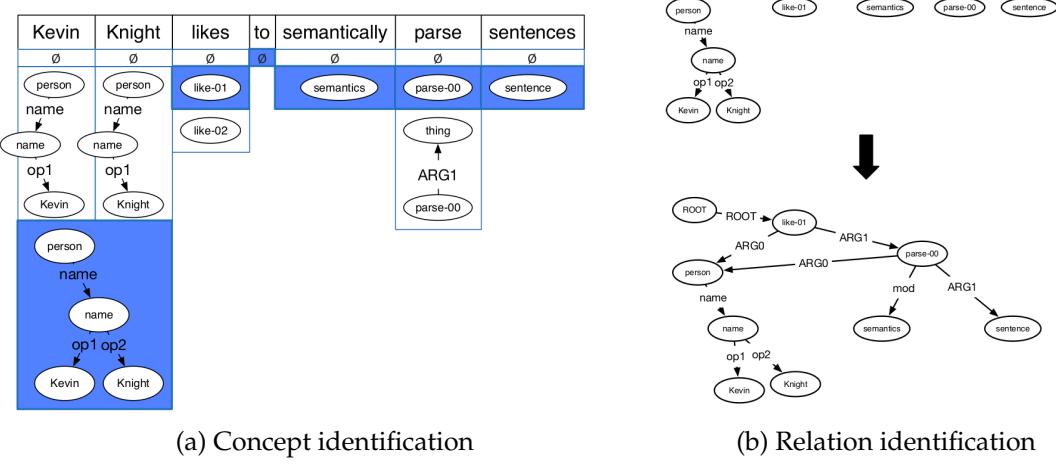


Figure 13.4: Subtasks for Abstract Meaning Representation parsing, from Schneider et al. (2015). [todo: permission]

6930 Relation identification can be formulated as search for the maximum spanning sub-
 6931 graph, under a set of constraints. Each labeled edge has a score, which is computed
 6932 from features of the concepts. We then search for the set of labeled edges that maximizes
 6933 the sum of these scores, under the constraint that the resulting graph is a well-formed
 6934 AMR (Flanigan et al., 2014). This constrained search can be performed by optimization
 6935 techniques such as integer linear programming, as described in § 13.2.2.

6936 **Transition-based parsing** In many cases, AMR parses are structurally similar to syn-
 6937 tactic dependency parses. Figure 13.5 shows one such example. This motivates an alter-
 6938 native approach to AMR parsing: modify the syntactic dependency parse until it looks
 6939 like a good AMR parse. Wang et al. (2015) propose a transition-based method, based on
 6940 incremental modifications to the syntactic dependency tree (transition-based dependency
 6941 parsing is discussed in § 11.3). At each step, the parser performs an action: for example,
 6942 adding an AMR relation label to the current dependency edge, swapping the direction of
 6943 a syntactic dependency edge, or cutting an edge and reattaching the orphaned subtree to
 6944 a new parent. The overall system is trained as a classifier, learning to choose the action as
 6945 would be given by an **oracle** that is capable of reproducing the ground-truth parse.

6946 13.4 Applications of Predicate-Argument Semantics

6947 **Question answering** Factoid questions have answers that are single words or phrases,
 6948 such as *who discovered prions?*, *where was Barack Obama born?*, and *in what year did the Knicks*
 6949 *last win the championship?* Semantic role labeling can be used to answer such questions,

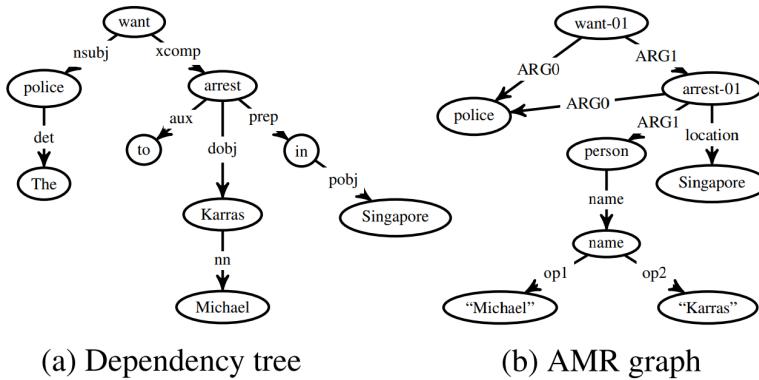


Figure 13.5: Syntactic dependency parse and AMR graph for the sentence *The police want to arrest Michael Karras in Singapore* (borrowed from Wang et al. (2015)) [todo: permission]

6950 by linking questions to sentences in a corpus of text. Shen and Lapata (2007) perform
 6951 FrameNet semantic role labeling on the query, and then construct a weighted **bipartite**
 6952 **graph**¹⁵ between FrameNet semantic roles and the words and phrases in the sentence.
 6953 This is done by first scoring all pairs of semantic roles and assignments, as shown in the
 6954 top half of Figure 13.6. They then find the bipartite edge cover, which is the minimum
 6955 weighted subset of edges such that each vertex has at least one edge, as shown in the
 6956 bottom half of Figure 13.6. After analyzing the question in this manner, Shen and Lapata
 6957 then find semantically-compatible sentences in the corpus, by performing graph matching
 6958 on the bipartite graphs for the question and candidate answer sentences. Finally, the
 6959 *expected answer phrase* in the question — typically the *wh*-word — is linked to a phrase in
 6960 the candidate answer source, and that phrase is returned as the answer.

6961 **Relation extraction** The task of **relation extraction** involves identifying pairs of entities
 6962 for which a given semantic relation holds (see § 17.2. For example, we might like to find
 6963 all pairs (i, j) such that i is the INVENTOR-OF j . PropBank semantic role labeling can
 6964 be applied to this task by identifying sentences whose verb signals the desired relation,
 6965 and then extracting ARG1 and ARG2 as arguments. (To fully solve this task, these argu-
 6966 ments must then be linked to entities, as described in chapter 17.) Christensen et al. (2010)
 6967 compare a semantic role labeling system against a simpler approach based on surface pat-
 6968 terns (Banko et al., 2007). They find that the SRL system is considerably more accurate,
 6969 but that it is several orders of magnitude slower. Conversely, Barnickel et al. (2009) apply
 6970 SENNA, a convolutional neural network SRL system (Collobert and Weston, 2008) to the
 6971 task of identifying biomedical relations (e.g., which genes inhibit or activate each other).

¹⁵A bipartite graph is one in which the vertices can be divided into two disjoint sets, and every edge connects a vertex in one set to a vertex in the other.

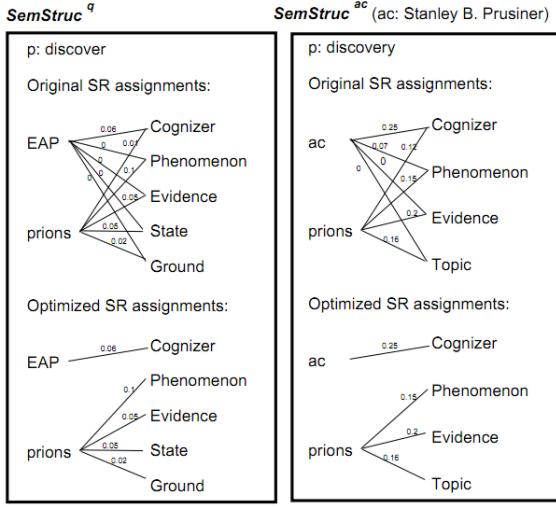


Figure 13.6: FrameNet semantic role labeling is used in factoid question answering, by aligning the semantic roles in the question (q) against those of sentences containing answer candidates (ac). “EAP” is the expected answer phrase, replacing the word *who* in the question. Figure reprinted from Shen and Lapata (2007) [todo: permission]




Figure 13.7: Fragment of AMR knowledge network for entity linking. Figure reprinted from Pan et al. (2015) [todo: permission]

6972 In comparison with a strong baseline that applies a set of rules to syntactic dependency
 6973 structures (Fundel et al., 2007), the SRL system is faster but less accurate. One possible
 6974 explanation for these divergent results is that Fundel et al. compare against a baseline
 6975 which is carefully tuned for performance in a relatively narrow domain, while the system
 6976 of Banko et al. is designed to analyze text across the entire web.

6977 **Entity linking** Another core task in information extraction is to link mentions of entities
6978 (e.g., *Republican candidates like Romney, Paul, and Johnson* ...) to entities in a knowledge
6979 base (e.g., LYNDON JOHNSON or GARY JOHNSON). This task, which is described in § 17.1,
6980 is often performed by examining nearby “collaborator” mentions — in this case, *Romney*
6981 and *Paul*. By jointly linking all such mentions, it is possible to arrive at a good overall
6982 solution. Pan et al. (2015) apply AMR to this problem. For each entity, they construct a
6983 knowledge network based on its semantic relations with other mentions within the same
6984 sentence. They then rerank a set of candidate entities, based on the overlap between
6985 the entity’s knowledge network and the semantic relations present in the sentence (Figure
6986 13.7).

6987 **Exercises**

- 6988 1. Write out an event semantic representation for the following sentences. You may
6989 make up your own predicates.
 - 6990 (13.26) *Abigail shares with Max.*
 - 6991 (13.27) *Abigail reluctantly shares a toy with Max.*
 - 6992 (13.28) *Abigail hates to share with Max.*
 - 6993 2. Find the PropBank framesets for *share* and *hate* at <http://verbs.colorado.edu/propbank/framesets-english-aliases/>, and rewrite your answers from the
6994 previous question, using the thematic roles ARG0, ARG1, and ARG2.
6995
 - 6996 3. Compute the syntactic path features for Abigail and Max in each of the example sentences (13.26) and (13.28) in Question 1, with respect to the verb *share*. If you’re not
6997 sure about the parse, you can try an online parser such as <http://nlp.stanford.edu:8080/parser/>.
6998
 - 7000 4. Compute the dependency path features for Abigail and Max in each of the example sentences (13.26) and (13.28) in Question 1, with respect to the verb *share*. Again, if
7001 you’re not sure about the parse, you can try an online parser such as <http://nlp.stanford.edu:8080/parser/>. As a hint, the dependency relation between *share*
7002 and *Max* is OBL according to the Universal Dependency treebank (version 2).
7003
 - 7004 5. PropBank semantic role labeling includes **reference arguments**, such as,
7005
- 7006 (13.29) [AM-LOC The bed] on [R-AM-LOC which] I slept broke.¹⁶

¹⁶Example from 2013 NAACL tutorial slides by Shumin Wu

7007 The label R-AM-LOC indicates that word *which* is a reference to *The bed*, which ex-
 7008 presses the location of the event. Reference arguments must have referents: the tag
 7009 R-AM-LOC can appear only when AM-LOC also appears in the sentence. Show how
 7010 to express this as a linear constraint, specifically for the tag R-AM-LOC. Be sure to
 7011 correctly handle the case in which neither AM-LOC nor R-AM-LOC appear in the
 7012 sentence.

- 7013 6. Explain how to express the constraints on semantic role labeling in Equation 13.8
 7014 and Equation 13.9 in the general form $Az \geq b$.
- 7015 7. Download the FrameNet sample data (<https://framenet.icsi.berkeley.edu/fndrupal/fulltextIndex>), and train a bag-of-words classifier to predict the
 7016 frame that is evoked by each verb in each example. Your classifier should build
 7017 a bag-of-words from the sentence in which the frame-evoking lexical unit appears.
 7019 [**todo: Somehow limit to one or a few lexical units.**] [**todo: use NLTK if possible**]
- 7020 8. Download the PropBank sample data, using NLTK (<http://www.nltk.org/howto/propbank.html>). Use a deep learning toolkit such as PyTorch or DyNet to train an
 7021 LSTM to predict tags. You will have to convert the downloaded instances to a BIO
 7022 sequence labeling representation first.
- 7024 9. Produce the AMR annotations for the following examples:

- 7025 (13.30) The girl likes the boy.
- 7026 (13.31) The girl was liked by the boy.
- 7027 (13.32) Abigail likes Maxwell Aristotle.
- 7028 (13.33) The spy likes the attractive boy.
- 7029 (13.34) The girl doesn't like the boy.
- 7030 (13.35) The girl likes her dog.

7031 For (13.32), recall that multi-token names are created using op1, op2, etc. You will
 7032 need to consult Banerjee et al. (2013) for (13.34), and Schneider et al. (2015) for
 7033 (13.35). You may assume that *her* refers to *the girl* in this example.

- 7034 10. Using an off-the-shelf PropBank SRL system,¹⁷ build a simplified question answer-
 7035 ing system in the style of Shen and Lapata (2007). Specifically, your system should
 7036 do the following:

¹⁷At the time of writing, the following systems are available: SENNA (<http://ronan.collobert.com/senna/>), Illinois Semantic Role Labeler (https://cogcomp.cs.illinois.edu/page/software_view/SRL), and mate-tools (<https://code.google.com/archive/p/mate-tools/>).

- 7037 • For each document in a collection, it should apply the semantic role labeler,
7038 and should store the output as a tuple.
- 7039 • For a question, your system should again apply the semantic role labeler. If
7040 any of the roles are filled by a *wh*-pronoun, you should mark that role as the
7041 expected answer phrase (EAP).
- 7042 • To answer the question, search for a stored tuple which matches the question as
7043 well as possible (same predicate, no incompatible semantic roles, and as many
7044 matching roles as possible). Align the EAP against its role filler in the stored
7045 tuple, and return this as the answer.

7046 To evaluate your system, download a set of three news articles on the same topic,
7047 and write down five factoid questions that should be answerable from the arti-
7048 cles. See if your system can answer these questions correctly. (If this problem is
7049 assigned to an entire class, you can build a large-scale test set and compare various
7050 approaches.)

7051 Chapter 14

7052 Distributional and distributed 7053 semantics

7054 A recurring theme in natural language processing is the complexity of the mapping from
7055 words to meaning. In chapter 4, we saw that a single word form, like *bank*, can have mul-
7056 tiple meanings; conversely, a single meaning may be created by multiple surface forms,
7057 a lexical semantic relationship known as **synonymy**. Despite this complex mapping be-
7058 tween words and meaning, natural language processing systems usually rely on words
7059 as the basic unit of analysis. This is especially true in semantics: the logical and frame
7060 semantic methods from the previous two chapters rely on hand-crafted lexicons that map
7061 from words to semantic predicates. But how can we analyze texts that contain words
7062 that we haven't seen before? This chapter describes methods that learn representations
7063 of word meaning by analyzing unlabeled data, vastly improving the generalizability of
7064 natural language processing systems. The theory that makes it possible to acquire mean-
7065 ingful representations from unlabeled data is the **distributional hypothesis**.

7066 14.1 The distributional hypothesis

7067 Here's a word you may not know: *tezgüino* (the example is from Lin, 1998). If you do not
7068 know the meaning of *tezgüino*, then you are in the same situation as a natural language
7069 processing system when it encounters a word that did not appear in its training data.
7070 Now suppose you see that *tezgüino* is used in the following contexts:

- 7071 (14.1) A bottle of _____ is on the table.
- 7072 (14.2) Everybody likes _____.
- 7073 (14.3) Don't have _____ before you drive.
- 7074 (14.4) We make _____ out of corn.

	(14.1)	(14.2)	(14.3)	(14.4)	...
<i>tezgüino</i>	1	1	1	1	
<i>loud</i>	0	0	0	0	
<i>motor oil</i>	1	0	0	1	
<i>tortillas</i>	0	1	0	1	
<i>choices</i>	0	1	0	0	
<i>wine</i>	1	1	1	0	

Table 14.1: Distributional statistics for *tezgüino* and five related terms

What other words fit into these contexts? How about: *loud*, *motor oil*, *tortillas*, *choices*, *wine*? Each row of Table 14.1 is a vector that summarizes the contextual properties for each word, with a value of one for contexts in which the word can appear, and a value of zero for contexts in which it cannot. Based on these vectors, we can conclude: *wine* is very similar to *tezgüino*; *motor oil* and *tortillas* are fairly similar to *tezgüino*; *loud* is completely different.

These vectors, which we will call **word representations**, describe the **distributional** properties of each word. Does vector similarity imply semantic similarity? This is the **distributional hypothesis**, stated by Firth (1957) as: “You shall know a word by the company it keeps.” The distributional hypothesis has stood the test of time: distributional statistics are a core part of language technology today, because they make it possible to leverage large amounts of unlabeled data to learn about rare words that do not appear in labeled training data.

Distributional statistics have a striking ability to capture lexical semantic relationships such as analogies. Figure 14.1 shows two examples, based on two-dimensional projections of distributional **word embeddings**, discussed later in this chapter. In each case, word-pair relationships correspond to regular linear patterns in this two dimensional space. No labeled data about the nature of these relationships was required to identify this underlying structure.

Distributional semantics are computed from context statistics. **Distributed** semantics are a related but distinct idea: that meaning can be represented by numerical vectors rather than symbolic structures. Distributed representations are often estimated from distributional statistics, as in latent semantic analysis and WORD2VEC, described later in this chapter. However, distributed representations can also be learned in a supervised fashion from labeled data, as in the neural classification models encountered in chapter 3.

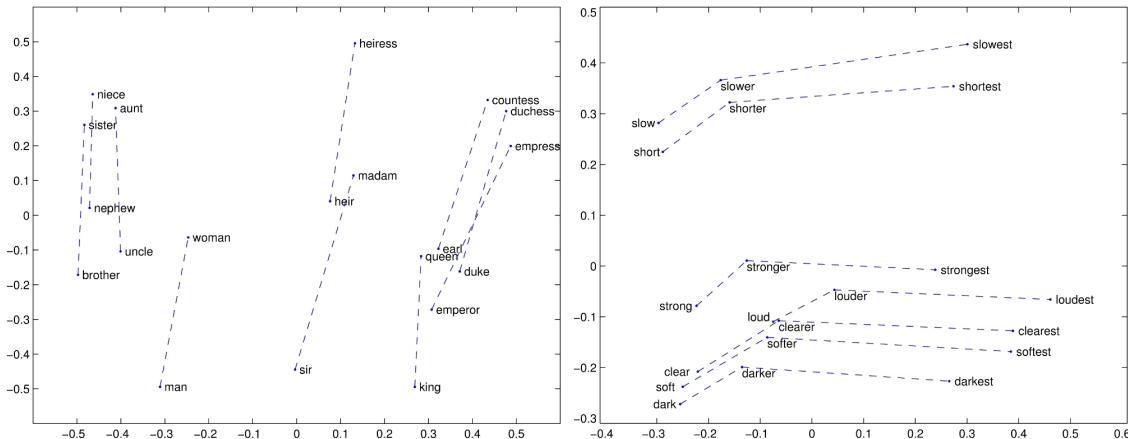


Figure 14.1: Lexical semantic relationships have regular linear structures in two dimensional projections of distributional statistics. From [http://nlp.stanford.edu/projects/glove/.\[todo: redo to make words bigger?\]](http://nlp.stanford.edu/projects/glove/.[todo: redo to make words bigger?])

14.2 Design decisions for word representations

There are many approaches for computing word representations, but most can be distinguished on three main dimensions: the nature of the representation, the source of contextual information, and the estimation procedure.

14.2.1 Representation

Today, the dominant word representations are k -dimensional vectors of real numbers, known as **word embeddings**. (The name is due to the fact that each discrete word is embedded in a continuous vector space.) This representation dates back at least to the late 1980s (Deerwester et al., 1990), and is used in popular techniques such as WORD2VEC (Mikolov et al., 2013).

Word embeddings are well suited for neural networks, where they can be plugged in as inputs. They can also be applied in linear classifiers and structure prediction models (Turian et al., 2010), although it can be difficult to learn linear models that employ real-valued features (Kummerfeld et al., 2015). A popular alternative is bit-string representations, such as **Brown clusters** (§ 14.4), in which each word is represented by a variable-length sequence of zeros and ones (Brown et al., 1992).

Another representational question is whether to estimate one embedding per surface form (e.g., *bank*), or to estimate distinct embeddings for each word sense or synset. Intuitively, if word representations are to capture the meaning of individual words, then words with multiple meanings should have multiple embeddings. This can be achieved

The moment one learns English, complications set in (Alfau, 1999)

Brown Clusters (Brown et al., 1992)	{one}
WORD2VEC (Mikolov et al., 2013) ($h = 2$)	{moment, one, English, complications}
Structured WORD2VEC (Ling et al., 2015) ($h = 2$)	$\{(moment, -2), (one, -1), (English, +1), (complications, +2)\}$
Dependency contexts (Levy and Goldberg, 2014)	$\{(one, \text{NSUBJ}), (\text{English}, \text{DOBJ}), (moment, \text{ACL}^{-1})\}$

Table 14.2: Contexts for the word *learns*, according to various word representations. For dependency context, *(one, NSUBJ)* means that there is a relation of type NSUBJ (nominal subject) **to** the word *one*, and *(moment, ACL⁻¹)* means that there is a relation of type ACL (adjectival clause) **from** the word *moment*.

7120 by integrating unsupervised clustering with word embedding estimation (Huang and
 7121 Yates, 2012; Li and Jurafsky, 2015). However, Arora et al. (2016) argue that it is unnec-
 7122 essary to model distinct word senses explicitly, because the embeddings for each surface
 7123 form are a linear combination of the embeddings of the underlying senses.

7124 14.2.2 Context

7125 The distributional hypothesis says that word meaning is related to the “contexts” in which
 7126 the word appears, but context can be defined in many ways. In the *tezgiiino* example, con-
 7127 texts are entire sentences, but in practice there are far too many sentences. At the oppo-
 7128 site extreme, the context could be defined as the immediately preceding word; this is the
 7129 context considered in Brown clusters. WORD2VEC takes an intermediate approach, using
 7130 local neighborhoods of words (e.g., $h = 5$) as contexts (Mikolov et al., 2013). Contexts
 7131 can also be much larger: for example, in **latent semantic analysis**, each word’s context
 7132 vector includes an entry per document, with a value of one if the word appears in the
 7133 document (Deerwester et al., 1990); in **explicit semantic analysis**, these documents are
 7134 Wikipedia pages (Gabrilovich and Markovitch, 2007).

7135 Words in context can be labeled by their position with respect to the target word w_m
 7136 (e.g., two words before, one word after), which makes the resulting word representations
 7137 more sensitive to syntactic differences (Ling et al., 2015). Another way to incorporate
 7138 syntax is to perform parsing as a preprocessing step, and then form context vectors from
 7139 the dependency edges (Levy and Goldberg, 2014) or predicate-argument relations (Lin,
 7140 1998). The resulting context vectors for several of these methods are shown in Table 14.2.

7141 The choice of context has a profound effect on the resulting representations, which

7142 can be viewed in terms of word similarity. Applying latent semantic analysis (§ 14.3) to
 7143 contexts of size $h = 2$ and $h = 30$ yields the following nearest-neighbors for the word
 7144 *dog*:¹

- 7145 • ($h = 2$): *cat, horse, fox, pet, rabbit, pig, animal, mongrel, sheep, pigeon*
 7146 • ($h = 30$): *kennel, puppy, pet, bitch, terrier, rottweiler, canine, cat, to bark, Alsatian*

7147 Which word list is better? Each word in the $h = 2$ list is an animal, reflecting the fact that
 7148 locally, the word *dog* tends to appear in the same contexts as other animal types (e.g., *pet*
 7149 *the dog, feed the dog*). In the $h = 30$ list, nearly everything is dog-related, including specific
 7150 breeds such as *rottweiler* and *Alsatian*. The list also includes words that are not animals
 7151 (*kennel*), and in one case (*to bark*), is not a noun at all. The 2-word context window is more
 7152 sensitive to syntax, while the 30-word window is more sensitive to topic.

7153 14.2.3 Estimation

7154 Word embeddings are estimated by optimizing some objective: the likelihood of a set of
 7155 unlabeled data (or a closely related quantity), or the reconstruction of a matrix of context
 7156 counts, similar to Table 14.1.

7157 **Maximum likelihood estimation** Likelihood-based optimization is derived from the
 7158 objective $\log p(\mathbf{w}; \mathbf{U})$, where $\mathbf{U} \in \mathbb{R}^{K \times V}$ is matrix of word embeddings, and $\mathbf{w} =$
 7159 $\{\mathbf{w}_m\}_{m=1}^M$ is a corpus, represented as a list of M tokens. Recurrent neural network lan-
 7160 guage models (§ 6.3) optimize this objective directly, backpropagating to the input word
 7161 embeddings through the recurrent structure. However, state-of-the-art word embeddings
 7162 employ huge corpora with hundreds of billions of tokens, and recurrent architectures are
 7163 difficult to scale to such data. As a result, likelihood-based word embeddings are usually
 7164 based on simplified likelihoods or heuristic approximations.

Matrix factorization The matrix $\mathbf{C} = \{\text{count}(i, j)\}$ stores the co-occurrence counts of
 word i and context j . Word representations can be obtained by approximately factoring
 this matrix, so that $\text{count}(i, j)$ is approximated by a function of a word embedding \mathbf{u}_i and
 a context embedding \mathbf{v}_j . These embeddings can be obtained by minimizing the norm of
 the reconstruction error,

$$\min_{\mathbf{u}, \mathbf{v}} \|\mathbf{C} - \tilde{\mathbf{C}}(\mathbf{u}, \mathbf{v})\|_F, \quad [14.1]$$

¹The example is from lecture slides by Marco Baroni, Alessandro Lenci, and Stefan Evert, who applied latent semantic analysis to the British National Corpus. You can find an online demo here: <http://clic.cimec.unitn.it/infomap-query/>

7165 where $\tilde{\mathbf{C}}(\mathbf{u}, \mathbf{v})$ is the approximate reconstruction resulting from the embeddings \mathbf{u} and
 7166 \mathbf{v} , and $\|\mathbf{X}\|_F$ indicates the Frobenius norm, $\sum_{i,j} x_{i,j}^2$. Rather than factoring the matrix of
 7167 word-context counts directly, it is often helpful to transform these counts using information-
 7168 theoretic metrics such as **pointwise mutual information** (PMI), described in the next sec-
 7169 tion.

7170 14.3 Latent semantic analysis

Latent semantic analysis (LSA) is one of the oldest approaches to distributed semantics (Deerwester et al., 1990). It induces continuous vector representations of words by factoring a matrix of word and context counts, using **truncated singular value decomposition** (SVD),

$$\min_{\mathbf{U} \in \mathbb{R}^{V \times K}, \mathbf{S} \in \mathbb{R}^{K \times K}, \mathbf{V} \in \mathbb{R}^{|\mathcal{C}| \times K}} \|\mathbf{C} - \mathbf{USV}^\top\|_F \quad [14.2]$$

$$\text{s.t. } \mathbf{U}^\top \mathbf{U} = \mathbb{I} \quad [14.3]$$

$$\mathbf{V}^\top \mathbf{V} = \mathbb{I} \quad [14.4]$$

$$\forall i \neq j, \mathbf{S}_{i,j} = 0, \quad [14.5]$$

7171 where V is the size of the vocabulary, $|\mathcal{C}|$ is the number of contexts, and K is size of the
 7172 resulting embeddings, which are set equal to the rows of the matrix \mathbf{U} . The matrix \mathbf{S} is
 7173 constrained to be diagonal (these diagonal elements are called the singular values), and
 7174 the columns of the product \mathbf{SV}^\top provide descriptions of the contexts. Each element $c_{i,j}$ is
 7175 then reconstructed as a **bilinear product**,

$$c_{i,j} \approx \sum_{k=1}^K u_{i,k} s_k v_{j,k}. \quad [14.6]$$

7176 The objective is to minimize the sum of squared approximation errors. The orthonormality
 7177 constraints $\mathbf{U}^\top \mathbf{U} = \mathbf{V}^\top \mathbf{V} = \mathbb{I}$ ensure that all pairs of dimensions in \mathbf{U} and \mathbf{V} are
 7178 uncorrelated, so that each dimension conveys unique information. Efficient implemen-
 7179 tations of truncated singular value decomposition are available in numerical computing
 7180 packages such as `scipy` and `matlab`.²

Latent semantic analysis is most effective when the count matrix is transformed before the application of SVD. One such transformation is **pointwise mutual information** (PMI; Church and Hanks, 1990), which captures the degree of association between word i and

²An important implementation detail is to represent \mathbf{C} as a **sparse matrix**, so that the storage cost is equal to the number of non-zero entries, rather than the size $V \times |\mathcal{C}|$.

context j ,

$$\text{PMI}(i, j) = \log \frac{p(i, j)}{p(i)p(j)} = \log \frac{p(i | j)p(j)}{p(i)p(j)} = \log \frac{p(i | j)}{p(i)} \quad [14.7]$$

$$= \log \text{count}(i, j) - \log \sum_{i'=1}^V \text{count}(i', j) \quad [14.8]$$

$$- \log \sum_{j' \in \mathcal{C}} \text{count}(i, j') + \log \sum_{i'=1}^V \sum_{j' \in \mathcal{C}} \text{count}(i', j'). \quad [14.9]$$

7181 The pointwise mutual information can be viewed as the logarithm of the ratio of the con-
 7182 ditional probability of word i in context j to the marginal probability of word i in all
 7183 contexts. When word i is statistically associated with context j , the ratio will be greater
 7184 than one, so $\text{PMI}(i, j) > 0$. The PMI transformation focuses latent semantic analysis on re-
 7185 constructing strong word-context associations, rather than on reconstructing large counts.

7186 The PMI is negative when a word and context occur together less often than if they
 7187 were independent, but such negative correlations are unreliable because counts of rare
 7188 events have high variance. Furthermore, the PMI is undefined when $\text{count}(i, j) = 0$. One
 7189 solution to these problems is to use the **Positive PMI** (PPMI),

$$\text{PPMI}(i, j) = \begin{cases} \text{PMI}(i, j), & p(i | j) > p(i) \\ 0, & \text{otherwise.} \end{cases} \quad [14.10]$$

7190 Bullinaria and Levy (2007) compare a range of matrix transformations for latent se-
 7191 mantic analysis, using a battery of tasks related to word meaning and word similarity
 7192 (for more on evaluation, see § 14.6). They find that PPMI-based latent semantic analysis
 7193 yields strong performance on a battery of tasks related to word meaning: for example,
 7194 PPMI-based LSA vectors can be used to solve multiple-choice word similarity questions
 7195 from the Test of English as a Foreign Language (TOEFL), obtaining 85% accuracy.

7196 14.4 Brown clusters

7197 Learning algorithms like perceptron and conditional random fields often perform better
 7198 with discrete feature vectors. A simple way to obtain discrete representations from distri-
 7199 butional statistics is by clustering (§ 5.1.1), so that words in the same cluster have similar
 7200 distributional statistics. This can help in downstream tasks, by sharing features between
 7201 all words in the same cluster. However, there is an obvious tradeoff: if the number of clus-
 7202 ters is too small, the words in each cluster will not have much in common; if the number
 7203 of clusters is too large, then the learner will not see enough examples from each cluster to
 7204 generalize.




Figure 14.2: Some subtrees produced by bottom-up Brown clustering (Miller et al., 2004) on news text [todo: permission]

7205 A solution to this problem is **hierarchical clustering**: using the distributional statistics
 7206 to induce a tree-structured representation. Fragments of **Brown cluster** trees are shown in
 7207 Figure 14.2 and Table 14.3. Each word’s representation consists of a binary string describ-
 7208 ing a path through the tree: 0 for taking the left branch, and 1 for taking the right branch.
 7209 In the subtree in the upper right of the figure, the representation of the word *conversation*
 7210 is 10; the representation of the word *assessment* is 0001. Bitstring prefixes capture simila-
 7211 rity at varying levels of specificity, and it is common to use the first eight, twelve, sixteen,
 7212 and twenty bits as features in tasks such as named entity recognition (Miller et al., 2004)
 7213 and dependency parsing (Koo et al., 2008).

Hierarchical trees can be induced from a likelihood-based objective, using a discrete latent variable $k_i \in \{1, 2, \dots, K\}$ to represent the cluster of word i :

$$\log p(\mathbf{w}; \mathbf{k}) \approx \sum_{m=1}^M \log p(w_m | w_{m-1}; \mathbf{k}) \quad [14.11]$$

$$\triangleq \sum_{m=1}^M \log p(w_m | k_{w_m}) + \log p(k_{w_m} | k_{w_{m-1}}). \quad [14.12]$$

7214 This is similar to a hidden Markov model, with the crucial difference that each word can

bitstring	ten most frequent words
01111010 0111	<i>excited thankful grateful stoked pumped anxious hyped psyched exited geeked</i>
01111010 100	<i>talking talkin complaining talkn bitching tlkn tlkin bragging rav- ing +k</i>
01111010 1010	<i>thinking thinkin dreaming worrying thinkn speakin reminiscing dreamin daydreaming fantasizing</i>
01111010 1011	<i>saying sayin suggesting stating sayn jokin talmbout implying insisting 5'2</i>
01111010 1100	<i>wonder dunno wondered duno donno dno doно wonda wounder dunnoe</i>
01111010 1101	<i>wondering wonders debating deciding pondering unsure won- derin debatin woundering wondern</i>
01111010 1110	<i>sure suree suuure suure sure- surre sures shuree</i>

Table 14.3: Fragment of a Brown clustering of Twitter data (Owoputi et al., 2013). Each row is a leaf in the tree, showing the ten most frequent words. This part of the tree emphasizes verbs of communicating and knowing, especially in the present participle. Each leaf node includes orthographic variants (*thinking*, *thinkin*, *thinkn*), semantically related terms (*excited*, *thankful*, *grateful*), and some outliers (5'2, +k). See http://www.cs.cmu.edu/~ark/TweetNLP/cluster_viewer.html for more.

7215 be emitted from only a single cluster: $\forall k \neq k_{w_m}, p(w_m | k) = 0$.

Using the objective in Equation 14.12, the Brown clustering tree can be constructed from the bottom up: begin with each word in its own cluster, and incrementally merge clusters until only a single cluster remains. At each step, we merge the pair of clusters such that the objective in Equation 14.12 is maximized. Although the objective seems to involve a sum over the entire corpus, the score for each merger can be computed from the cluster-to-cluster co-occurrence counts. These counts can be updated incrementally as the clustering proceeds. The optimal merge at each step can be shown to maximize the **average mutual information**,

$$I(\mathbf{k}) = \sum_{k_1=1}^K \sum_{k_2=1}^K p(k_1, k_2) \times \text{PMI}(k_1, k_2) \quad [14.13]$$

$$p(k_1, k_2) = \frac{\text{count}(k_1, k_2)}{\sum_{k_{1'}=1}^K \sum_{k_{2'}=1}^K \text{count}(k_{1'}, k_{2'})},$$

7216 where $p(k_1, k_2)$ is the joint probability of a bigram involving a word in cluster k_1 followed
7217 by a word in k_2 . This probability and the PMI are both computed from the co-occurrence

Algorithm 17 Exchange clustering algorithm. Assumes that words are sorted by frequency, and that MAXMI finds the cluster pair whose merger maximizes the mutual information, as defined in Equation 14.13.

```

procedure EXCHANGECLUSTERING({count(., .)}, K)
    for  $i \in 1 \dots K$  do ▷ Initialization
         $k_i \leftarrow i, \quad i = 1, 2, \dots, K$ 
        for  $j \in 1 \dots K$  do
             $c_{i,j} \leftarrow \text{count}(i, j)$ 
         $\tau \leftarrow \{(i)\}_{i=1}^K$ 
        for  $i \in \{K + 1, K + 2, \dots, V\}$  do ▷ Iteratively add each word to the clustering
             $\tau \leftarrow \tau \cup (i)$ 
            for  $k \in \tau$  do
                 $c_{k,i} \leftarrow \text{count}(k, i)$ 
                 $c_{i,k} \leftarrow \text{count}(i, k)$ 
             $\hat{i}, \hat{j} \leftarrow \text{MAXMI}(\mathbf{C})$ 
             $\tau, \mathbf{C} \leftarrow \text{MERGE}(\hat{i}, \hat{j}, \mathbf{C}, \tau)$ 
        repeat ▷ Merge the remaining clusters into a tree
             $\hat{i}, \hat{j} \leftarrow \text{MAXMI}(\mathbf{C}, \tau)$ 
             $\tau, \mathbf{C} \leftarrow \text{MERGE}(\hat{i}, \hat{j}, \mathbf{C}, \tau)$ 
        until  $|\tau| = 1$ 
        return  $\tau$ 
procedure MERGE( $i, j, \mathbf{C}, \tau$ )
     $\tau \leftarrow \tau \setminus i \setminus j \cup (i, j)$  ▷ Merge the clusters in the tree
    for  $k \in \tau$  do ▷ Aggregate the counts across the merged clusters
         $c_{k,(i,j)} \leftarrow c_{k,i} + c_{k,j}$ 
         $c_{(i,j),k} \leftarrow c_{i,k} + c_{j,k}$ 
    return  $\tau, \mathbf{C}$ 

```

7218 counts between clusters. After each merger, the co-occurrence vectors for the merged
 7219 clusters are simply added up, so that the next optimal merger can be found efficiently.

7220 This bottom-up procedure requires iterating over the entire vocabulary, and evaluating
 7221 K_t^2 possible mergers at each step, where K_t is the current number of clusters at step t
 7222 of the algorithm. Furthermore, computing the score for each merger involves a sum over
 7223 K_t^2 clusters. The maximum number of clusters is $K_0 = V$, which occurs when every word
 7224 is in its own cluster at the beginning of the algorithm. The time complexity is thus $\mathcal{O}(V^5)$.

7225 To avoid this complexity, practical implementations use a heuristic approximation
 7226 called **exchange clustering**. The K most common words are placed in clusters of their
 7227 own at the beginning of the process. We then consider the next most common word, and

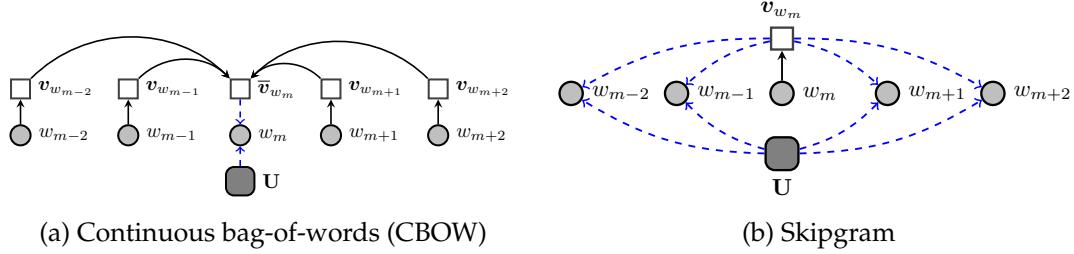


Figure 14.3: The CBOW and skipgram variants of WORD2VEC. The parameter \mathbf{U} is the matrix of word embeddings, and each \mathbf{v}_m is the context embedding for word w_m .

7228 merge it with one of the existing clusters. This continues until the entire vocabulary has
 7229 been incorporated, at which point the K clusters are merged down to a single cluster,
 7230 forming a tree. The algorithm never considers more than $K + 1$ clusters at any step, and
 7231 the complexity is $\mathcal{O}(VK + V \log V)$, with the second term representing the cost of sorting
 7232 the words at the beginning of the algorithm.

7233 14.5 Neural word embeddings

7234 Neural word embeddings combine aspects of the previous two methods: like latent se-
 7235 mantic analysis, they are a continuous vector representation; like Brown clusters, they are
 7236 trained from a likelihood-based objective. Let the vector \mathbf{u}_i represent the K -dimensional
 7237 **embedding** for word i , and let \mathbf{v}_j represent the K -dimensional embedding for context
 7238 j . The inner product $\mathbf{u}_i \cdot \mathbf{v}_j$ represents the compatibility between word i and context j .
 7239 By incorporating this inner product into an approximation to the log-likelihood of a cor-
 7240 pus, it is possible to estimate both parameters by backpropagation. WORD2VEC (Mikolov
 7241 et al., 2013) includes two such approximations: continuous bag-of-words (CBOW) and
 7242 skipgrams.

7243 14.5.1 Continuous bag-of-words (CBOW)

7244 In recurrent neural network language models, each word w_m is conditioned on a recurrently-
 7245 updated state vector, which is based on word representations going all the way back to the
 7246 beginning of the text. The **continuous bag-of-words (CBOW)** model is a simplification:
 7247 the local context is computed as an average of embeddings for words in the immediate
 7248 neighborhood $m - h, m - h + 1, \dots, m + h - 1, m + h$,

$$\bar{\mathbf{v}}_m = \frac{1}{2h} \sum_{n=1}^h \mathbf{v}_{w_{m+n}} + \mathbf{v}_{w_{m-n}}. \quad [14.14]$$

7249 Thus, CBOW is a bag-of-words model, because the order of the context words does not
 7250 matter; it is continuous, because rather than conditioning on the words themselves, we
 7251 condition on a continuous vector constructed from the word embeddings. The parameter
 7252 h determines the neighborhood size, which Mikolov et al. (2013) set to $h = 4$.

The CBOW model optimizes an approximation to the corpus log-likelihood,

$$\log p(\mathbf{w}) \approx \sum_{m=1}^M \log p(w_m | w_{m-h}, w_{m-h+1}, \dots, w_{m+h-1}, w_{m+h}) \quad [14.15]$$

$$= \sum_{m=1}^M \log \frac{\exp(\mathbf{u}_{w_m} \cdot \bar{\mathbf{v}}_m)}{\sum_{j=1}^V \exp(\mathbf{u}_j \cdot \bar{\mathbf{v}}_m)} \quad [14.16]$$

$$= \sum_{m=1}^M \mathbf{u}_{w_m} \cdot \bar{\mathbf{v}}_m - \log \sum_{j=1}^V \exp(\mathbf{u}_j \cdot \bar{\mathbf{v}}_m). \quad [14.17]$$

7253 14.5.2 Skipgrams

In the CBOW model, words are predicted from their context. In the **skipgram** model, the context is predicted from the word, yielding the objective:

$$\log p(\mathbf{w}) \approx \sum_{m=1}^M \sum_{n=1}^{h_m} \log p(w_{m-n} | w_m) + \log p(w_{m+n} | w_m) \quad [14.18]$$

$$= \sum_{m=1}^M \sum_{n=1}^{h_m} \log \frac{\exp(\mathbf{u}_{w_{m-n}} \cdot \mathbf{v}_{w_m})}{\sum_{j=1}^V \exp(\mathbf{u}_j \cdot \mathbf{v}_{w_m})} + \log \frac{\exp(\mathbf{u}_{w_{m+n}} \cdot \mathbf{v}_{w_m})}{\sum_{j=1}^V \exp(\mathbf{u}_j \cdot \mathbf{v}_{w_m})} \quad [14.19]$$

$$= \sum_{m=1}^M \sum_{n=1}^{h_m} \mathbf{u}_{w_{m-n}} \cdot \mathbf{v}_{w_m} + \mathbf{u}_{w_{m+n}} \cdot \mathbf{v}_{w_m} - 2 \log \sum_{j=1}^V \exp(\mathbf{u}_j \cdot \mathbf{v}_{w_m}). \quad [14.20]$$

7254 In the skipgram approximation, each word is generated multiple times; each time it is con-
 7255 ditioned only on a single word. This makes it possible to avoid averaging the word vec-
 7256 tors, as in the CBOW model. The local neighborhood size h_m is randomly sampled from
 7257 a uniform categorical distribution over the range $\{1, 2, \dots, h_{\max}\}$; Mikolov et al. (2013) set
 7258 $h_{\max} = 10$. Because the neighborhood grows outward with h , this approach has the effect
 7259 of weighting near neighbors more than distant ones. Skipgram performs better on most
 7260 evaluations than CBOW (see § 14.6 for details of how to evaluate word representations),
 7261 but CBOW is faster to train (Mikolov et al., 2013).

7262 14.5.3 Computational complexity

7263 The WORD2VEC models can be viewed as an efficient alternative to recurrent neural net-
 7264 work language models, which involve a recurrent state update whose time complexity

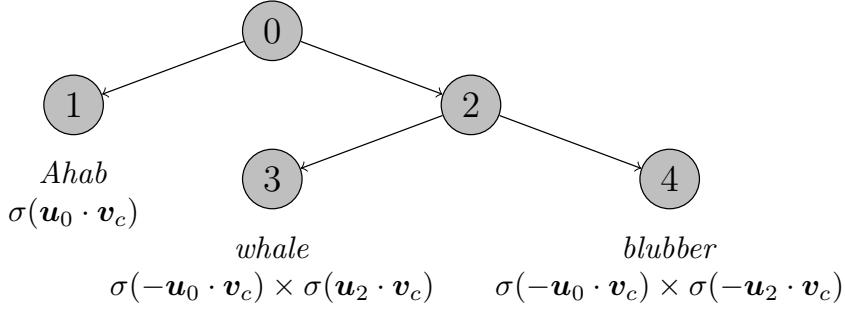


Figure 14.4: A fragment of a hierarchical softmax tree. The probability of each word is computed as a product of probabilities of local branching decisions in the tree.

is quadratic in the size of the recurrent state vector. CBOW and skipgram avoid this computation, and incur only a linear time complexity in the size of the word and context representations. However, all three models compute a normalized probability over word tokens; a naïve implementation of this probability requires summing over the entire vocabulary. The time complexity of this sum is $\mathcal{O}(V \times K)$, which dominates all other computational costs. There are two solutions: **hierarchical softmax**, a tree-based computation that reduces the cost to a logarithm of the size of the vocabulary; and **negative sampling**, an approximation that eliminates the dependence on vocabulary size. Both methods are also applicable to RNN language models.

14.5.3.1 Hierarchical softmax

In Brown clustering, the vocabulary is organized into a binary tree. Mnih and Hinton (2008) show that the normalized probability over words in the vocabulary can be reparametrized as a probability over paths through such a tree. This **hierarchical softmax** probability is computed as a product of binary decisions over whether to move left or right through the tree, with each binary decision represented as a sigmoid function of the inner product between the context embedding \mathbf{v}_c and an output embedding associated with the node \mathbf{u}_n ,

$$\Pr(\text{left at } n \mid c) = \sigma(\mathbf{u}_n \cdot \mathbf{v}_c) \quad [14.21]$$

$$\Pr(\text{right at } n \mid c) = 1 - \sigma(\mathbf{u}_n \cdot \mathbf{v}_c) = \sigma(-\mathbf{u}_n \cdot \mathbf{v}_c), \quad [14.22]$$

where σ refers to the sigmoid function, $\sigma(x) = \frac{1}{1+\exp(-x)}$. The range of the sigmoid is the interval $(0, 1)$, and $1 - \sigma(x) = \sigma(-x)$.

As shown in Figure 14.4, the probability of generating each word is redefined as the product of the probabilities across its path. The sum of all such path probabilities is guaranteed to be one, for any context vector $\mathbf{v}_c \in \mathbb{R}^K$. In a balanced binary tree, the depth is

7280 logarithmic in the number of leaf nodes, and thus the number of multiplications is equal
 7281 to $\mathcal{O}(\log V)$. The number of non-leaf nodes is equal to $\mathcal{O}(2V - 1)$, so the number of pa-
 7282 rameters to be estimated increases by only a small multiple. The tree can be constructed
 7283 using an incremental clustering procedure similar to hierarchical Brown clusters (Mnih
 7284 and Hinton, 2008), or by using the Huffman (1952) encoding algorithm for lossless com-
 7285 pression.

7286 **14.5.3.2 Negative sampling**

Likelihood-based methods are computationally intensive because each probability must be normalized over the vocabulary. These probabilities are based on scores for each word in each context, and it is possible to design an alternative objective that is based on these scores more directly: we seek word embeddings that maximize the score for the word that was really observed in each context, while minimizing the scores for a set of randomly selected **negative samples**:

$$\psi(i, j) = \log \sigma(\mathbf{u}_i \cdot \mathbf{v}_j) + \sum_{i' \in \mathcal{W}_{\text{neg}}} \log(1 - \sigma(\mathbf{u}_{i'} \cdot \mathbf{v}_j)), \quad [14.23]$$

7287 where $\psi(i, j)$ is the score for word i in context j , and \mathcal{W}_{neg} is the set of negative samples.
 7288 The objective is to maximize the sum over the corpus, $\sum_{m=1}^M \psi(w_m, c_m)$, where w_m is
 7289 token m and c_m is the associated context.

7290 The set of negative samples \mathcal{W}_{neg} is obtained by sampling from a unigram language
 7291 model. Mikolov et al. (2013) construct this unigram language model by exponentiating
 7292 the empirical word probabilities, setting $\hat{p}(i) \propto (\text{count}(i))^{\frac{3}{4}}$. This has the effect of redi-
 7293 tributing probability mass from common to rare words. The number of negative samples
 7294 increases the time complexity of training by a constant factor. Mikolov et al. (2013) report
 7295 that 5-20 negative samples works for small training sets, and that two to five samples
 7296 suffice for larger corpora.

7297 **14.5.4 Word embeddings as matrix factorization**

7298 The negative sampling objective in Equation 14.23 can be justified as an efficient approx-
 7299 imation to the log-likelihood, but it is also closely linked to the matrix factorization ob-
 7300 jective employed in latent semantic analysis. For a matrix of word-context pairs in which
 7301 all counts are non-zero, negative sampling is equivalent to factorization of the matrix M ,
 7302 where $M_{ij} = \text{PMI}(i, j) - \log k$: each cell in the matrix is equal to the pointwise mutual
 7303 information of the word and context, shifted by $\log k$, with k equal to the number of neg-
 7304 ative samples (Levy and Goldberg, 2014). For word-context pairs that are not observed in
 7305 the data, the pointwise mutual information is $-\infty$, but this can be addressed by consid-
 7306 ering only PMI values that are greater than $\log k$, resulting in a matrix of **shifted positive**

7307 **pointwise mutual information,**

$$M_{ij} = \max(0, \text{PMI}(i, j) - \log k). \quad [14.24]$$

7308 Word embeddings are obtained by factoring this matrix with truncated singular value
7309 decomposition.

GloVe (“global vectors”) are a closely related approach (Pennington et al., 2014), in which the matrix to be factored is constructed from log co-occurrence counts, $M_{ij} = \log \text{count}(i, j)$. The word embeddings are estimated by minimizing the sum of squares,

$$\begin{aligned} \min_{\mathbf{u}, \mathbf{v}, b, \tilde{b}} \quad & \sum_{j=1}^V \sum_{j \in C} f(M_{ij}) \left(\widehat{\log M_{ij}} - \log M_{ij} \right)^2 \\ \text{s.t.} \quad & \widehat{\log M_{ij}} = \mathbf{u}_i \cdot \mathbf{v}_j + b_i + \tilde{b}_j, \end{aligned} \quad [14.25]$$

7310 where b_i and \tilde{b}_j are offsets for word i and context j , which are estimated jointly with the
7311 embeddings \mathbf{u} and \mathbf{v} . The weighting function $f(M_{ij})$ is set to be zero at $M_{ij} = 0$, thus
7312 avoiding the problem of taking the logarithm of zero counts; it saturates at $M_{ij} = m_{\max}$,
7313 thus avoiding the problem of overcounting common word-context pairs. This heuristic
7314 turns out to be critical to the method’s performance.

7315 The time complexity of sparse matrix reconstruction is determined by the number of
7316 non-zero word-context counts. Pennington et al. (2014) show that this number grows
7317 sublinearly with the size of the dataset: roughly $\mathcal{O}(N^{0.8})$ for typical English corpora. In
7318 contrast, the time complexity of WORD2VEC is linear in the corpus size. Computing the co-
7319 occurrence counts also requires linear time in the size of the corpus, but this operation can
7320 easily be parallelized using MapReduce-style algorithms (Dean and Ghemawat, 2008).

7321 14.6 Evaluating word embeddings

7322 Distributed word representations can be evaluated in two main ways. **Intrinsic** evalua-
7323 tions test whether the representations cohere with our intuitions about word meaning.
7324 **Extrinsic** evaluations test whether they are useful for downstream tasks, such as sequence
7325 labeling.

7326 14.6.1 Intrinsic evaluations

7327 A basic question for word embeddings is whether the similarity of words i and j is re-
7328 flected in the similarity of the vectors \mathbf{u}_i and \mathbf{u}_j . **Cosine similarity** is typically used to
7329 compare two word embeddings,

$$\cos(\mathbf{u}_i, \mathbf{u}_j) = \frac{\mathbf{u}_i \cdot \mathbf{u}_j}{\|\mathbf{u}_i\|_2 \times \|\mathbf{u}_j\|_2}. \quad [14.26]$$

word 1	word 2	similarity
<i>love</i>	<i>sex</i>	6.77
<i>stock</i>	<i>jaguar</i>	0.92
<i>money</i>	<i>cash</i>	9.15
<i>development</i>	<i>issue</i>	3.97
<i>lad</i>	<i>brother</i>	4.46

Table 14.4: Subset of the WS-353 (Finkelstein et al., 2002) dataset of word similarity ratings (examples from Faruqui et al. (2016)).

7330 For any embedding method, we can evaluate whether the cosine similarity of word em-
 7331 beddings is correlated with human judgments of word similarity. The WS-353 dataset (Finkel-
 7332 stein et al., 2002) includes similarity scores for 353 word pairs (Table 14.4). To test the
 7333 accuracy of embeddings for rare and morphologically complex words, Luong et al. (2013)
 7334 introduce a dataset of “rare words.” Outside of English, word similarity resources are
 7335 limited, mainly consisting of translations of WS-353.

7336 Word analogies (e.g., *king:queen :: man:woman*) have also been used to evaluate word
 7337 embeddings (Mikolov et al., 2013). In this evaluation, the system is provided with the first
 7338 three parts of the analogy ($i_1 : j_1 :: i_2 : ?$), and the final element is predicted by finding the
 7339 word embedding most similar to $\mathbf{u}_{i_1} - \mathbf{u}_{j_1} + \mathbf{u}_{i_2}$. Another evaluation tests whether word
 7340 embeddings are related to broad lexical semantic categories called **supersenses** (Caramita
 7341 and Johnson, 2003): verbs of motion, nouns that describe animals, nouns that describe
 7342 body parts, and so on. These supersenses are annotated for English synsets in Word-
 7343 Net (Fellbaum, 2010). This evaluation is implemented in the `qvec` metric, which tests
 7344 whether the matrix of supersenses can be reconstructed from the matrix of word embed-
 7345 dings (Tsvetkov et al., 2015).

7346 Levy et al. (2015) compared several dense word representations for English — includ-
 7347 ing latent semantic analysis, WORD2VEC, and GloVe — using six word similarity metrics
 7348 and two analogy tasks. None of the embeddings outperformed the others on every task,
 7349 but skipgrams were the most broadly competitive. Hyperparameter tuning played a key
 7350 role: any method will perform badly if the wrong hyperparameters are used. Relevant
 7351 hyperparameters include the embedding size, as well as algorithm-specific details such
 7352 as the neighborhood size and the number of negative samples.

7353 14.6.2 Extrinsic evaluations

7354 Word representations contribute to downstream tasks like sequence labeling and docu-
 7355 ment classification by enabling generalization across words. The use of distributed repre-
 7356 sentations as features is a form of **semi-supervised learning**, in which performance on a

7357 supervised learning problem is augmented by learning distributed representations from
 7358 unlabeled data (Miller et al., 2004; Koo et al., 2008; Turian et al., 2010). These **pre-trained**
 7359 **word representations** can be used as features in a linear prediction model, or as the input
 7360 layer in a neural network, such as a Bi-LSTM tagging model (§ 7.6). Word representations
 7361 can be evaluated by the performance of the downstream systems that consume them:
 7362 for example, GloVe embeddings are convincingly better than Latent Semantic Analysis
 7363 as features in the downstream task of named entity recognition (Pennington et al., 2014).
 7364 Unfortunately, extrinsic and intrinsic evaluations do not always point in the same direc-
 7365 tion, and the best word representations for one downstream task may perform poorly on
 7366 another task (Schnabel et al., 2015).

7367 When word representations are updated from labeled data in the downstream task,
 7368 they are said to be **fine-tuned**. When labeled data is plentiful, pre-training may be un-
 7369 necessary; when labeled data is scarce, fine-tuning may lead to overfitting. Various com-
 7370 binations of pre-training and fine-tuning can be employed. Pre-trained embeddings can
 7371 be used as initialization before fine-tuning, and this can substantially improve perfor-
 7372 mance (Lample et al., 2016). Alternatively, both fine-tuned and pre-trained embeddings
 7373 can be used as inputs in a single model (Kim, 2014).

7374 In semi-supervised scenarios, pretrained word embeddings can be replaced by “con-
 7375 textualized” word representations (Peters et al., 2018). These contextualized represen-
 7376 tations are set to the hidden states of a deep bi-directional LSTM, which is trained as a
 7377 bi-directional language model, motivating the name **ELMo (embeddings from language**
 7378 **models**). Given a supervised learning problem, the language model generates contextu-
 7379 alized representations, which are then used as the base layer in a task-specific supervised
 7380 neural network. This approach yields significant gains over pretrained word embeddings
 7381 on several tasks, presumably because the contextualized embeddings use unlabeled data
 7382 to learn how to integrate linguistic context into the base layer of the supervised neural
 7383 network.

7384 14.7 Distributed representations beyond distributional statistics

7385 Distributional word representations can be estimated from huge unlabeled datasets, thereby
 7386 covering many words that do not appear in labeled data: for example, GloVe embeddings
 7387 are estimated from 800 billion tokens of web data,³ while the largest labeled datasets for
 7388 NLP tasks are on the order of millions of tokens. Nonetheless, even a dataset of hundreds
 7389 of billions of tokens will not cover every word that may be encountered in the future.
 7390 Furthermore, many words will appear only a few times, making their embeddings un-
 7391 reliable. Many languages exceed English in morphological complexity, and thus have
 7392 lower token-to-type ratios. When this problem is coupled with small training corpora, it

³<http://commoncrawl.org/>

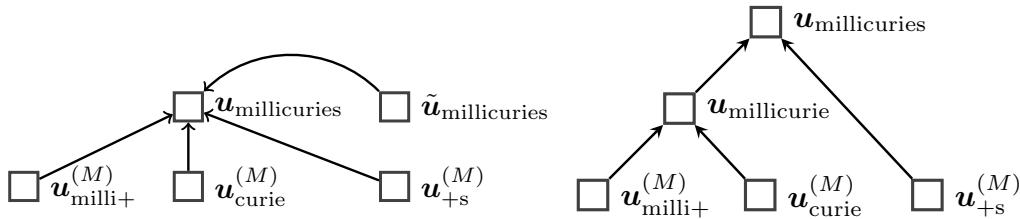


Figure 14.5: Two architectures for building word embeddings from subword units. On the left, morpheme embeddings $u^{(m)}$ are combined by addition with the non-compositional word embedding \tilde{u} (Botha and Blunsom, 2014). On the right, morpheme embeddings are combined in a recursive neural network (Luong et al., 2013).

7393 becomes especially important to leverage other sources of information beyond distributional statistics.

7395 14.7.1 Word-internal structure

7396 One solution is to incorporate word-internal structure into word embeddings. Purely
7397 distributional approaches consider words as atomic units, but in fact, many words have
7398 internal structure, so that their meaning can be **composed** from the representations of
7399 sub-word units. Consider the following terms, all of which are missing from Google's
7400 pre-trained WORD2VEC embeddings:⁴

7401 **millicuries** This word has **morphological** structure (see § 9.1.2 for more on morphology):
7402 the prefix *milli-* indicates an amount, and the suffix *-s* indicates a plural. (A *millicurie*
7403 is an unit of radioactivity.)

7404 **caesium** This word is a single morpheme, but the characters *-ium* are often associated
7405 with chemical elements. (*Caesium* is the British spelling of a chemical element,
7406 spelled *cesium* in American English.)

7407 **IAEA** This term is an acronym, as suggested by the use of capitalization. The prefix *I-* frequently
7408 refers to international organizations, and the suffix *-A* often refers to agencies or associations. (*IAEA* is the International Atomic Energy Agency.)

7410 **Zhezhan** This term is in title case, suggesting the name of a person or place, and the
7411 character bigram *zh* indicates that it is likely a transliteration. (*Zhezhan* is a mining
7412 facility in Kazakhstan.)

⁴<https://code.google.com/archive/p/word2vec/>, accessed September 20, 2017

7413 How can word-internal structure be incorporated into word representations? One
7414 approach is to construct word representations from embeddings of the characters or mor-
7415 phemes. For example, if word i has morphological segments \mathcal{M}_i , then its embedding can
7416 be constructed by addition (Botha and Blunsom, 2014),

$$\mathbf{u}_i = \tilde{\mathbf{u}}_i + \sum_{j \in \mathcal{M}_i} \mathbf{u}_j^{(M)}, \quad [14.27]$$

7417 where $\mathbf{u}_m^{(M)}$ is a morpheme embedding and $\tilde{\mathbf{u}}_i$ is a non-compositional embedding of the
7418 whole word, which is an additional free parameter of the model (Figure 14.5, left side).
7419 All embeddings are estimated from a **log-bilinear language model** (Mnih and Hinton,
7420 2007), which is similar to the CBOW model (§ 14.5), but includes only contextual informa-
7421 tion from preceding words. The morphological segments are obtained using an unsuper-
7422 vised segmenter (Creutz and Lagus, 2007). For words that do not appear in the training
7423 data, the embedding can be constructed directly from the morphemes, assuming that each
7424 morpheme appears in some other word in the training data. The free parameter $\tilde{\mathbf{u}}$ adds
7425 flexibility: words with similar morphemes are encouraged to have similar embeddings,
7426 but this parameter makes it possible for them to be different.

7427 Word-internal structure can be incorporated into word representations in various other
7428 ways. Here are some of the main parameters.

7429 **Subword units.** Examples like *IAEA* and *Zhezghan* are not based on morphological com-
7430 position, and a morphological segmenter is unlikely to identify meaningful sub-
7431 word units for these terms. Rather than using morphemes for subword embeddings,
7432 one can use characters (Santos and Zadrozny, 2014; Ling et al., 2015; Kim et al., 2016),
7433 character n -grams (Wieting et al., 2016; Bojanowski et al., 2017), and **byte-pair en-**
7434 **codings**, a compression technique which captures frequent substrings (Gage, 1994;
7435 Sennrich et al., 2016).

7436 **Composition.** Combining the subword embeddings by addition does not differentiate
7437 between orderings, nor does it identify any particular morpheme as the **root**. A
7438 range of more flexible compositional models have been considered, including re-
7439 currence (Ling et al., 2015), convolution (Santos and Zadrozny, 2014; Kim et al.,
7440 2016), and **recursive neural networks** (Luong et al., 2013), in which representa-
7441 tions of progressively larger units are constructed over a morphological parse, e.g.
7442 $((milli+curie)+s)$, $((in+flam)+able)$, $(in+(vis+ible))$. A recursive embedding model is
7443 shown in the right panel of Figure 14.5.

7444 **Estimation.** Estimating subword embeddings from a full dataset is computationally ex-
7445 pensive. An alternative approach is to train a subword model to match pre-trained
7446 word embeddings (Cotterell et al., 2016; Pinter et al., 2017). To train such a model, it
7447 is only necessary to iterate over the vocabulary, and the not the corpus.

7448 **14.7.2 Lexical semantic resources**

Resources such as WordNet provide another source of information about word meaning; if we know that *caesium* is a synonym of *cesium*, or that a *millicurie* is a type of *measurement unit*, then this should help to provide embeddings for the unknown words, and to smooth embeddings of rare words. One way to do this is to **retrofit** pre-trained word embeddings across a network of lexical semantic relationships (Faruqui et al., 2015) by minimizing the following objective,

$$\min_{\mathbf{U}} \sum_{j=1}^V \|\mathbf{u}_i - \hat{\mathbf{u}}_i\|_2 + \sum_{(i,j) \in \mathcal{L}} \beta_{ij} \|\mathbf{u}_i - \mathbf{u}_j\|_2, \quad [14.28]$$

7449 where $\hat{\mathbf{u}}_i$ is the pretrained embedding of word i , and $\mathcal{L} = \{(i,j)\}$ is a lexicon of word
 7450 relations. The hyperparameter β_{ij} controls the importance of adjacent words having
 7451 similar embeddings; Faruqui et al. (2015) set it to the inverse of the degree of word i ,
 7452 $\beta_{ij} = |\{j : (i,j) \in \mathcal{L}\}|^{-1}$. Retrofitting improves performance on a range of intrinsic evalua-
 7453 tions, and gives small improvements on an extrinsic document classification task.

7454 **14.8 Distributed representations of multiword units**

7455 Can distributed representations extend to phrases, sentences, paragraphs, and beyond?
 7456 Before exploring this possibility, recall the distinction between distributed and distri-
 7457 butional representations. Neural embeddings such as WORD2VEC are both distributed
 7458 (vector-based) and distributional (derived from counts of words in context). As we con-
 7459 sider larger units of text, the counts decrease: in the limit, a multi-paragraph span of text
 7460 would never appear twice, except by plagiarism. Thus, the meaning of a large span of
 7461 text cannot be determined from distributional statistics alone; it must be computed com-
 7462 positionally from smaller spans. But these considerations are orthogonal to the question
 7463 of whether distributed representations — dense numerical vectors — are sufficiently ex-
 7464 pressive to capture the meaning of phrases, sentences, and paragraphs.

7465 **14.8.1 Purely distributional methods**

7466 Some multiword phrases are non-compositional: the meaning of such phrases is not de-
 7467 rived from the meaning of the individual words using typical compositional semantics.
 7468 This includes proper nouns like *San Francisco* as well as idiomatic expressions like *kick*
 7469 *the bucket* (Baldwin and Kim, 2010). For these cases, purely distributional approaches
 7470 can work. A simple approach is to identify multiword units that appear together fre-
 7471 quently, and then treat these units as words, learning embeddings using a technique such
 7472 as WORD2VEC. The problem of identifying multiword units is sometimes called **colloca-**
 7473 **tion extraction**, and can be approached using metrics such as pointwise mutual informa-
 7474 tion: two-word units are extracted first, and then larger units are extracted. Mikolov et al.

7475 (2013) identify such units and then treat them as words when estimating skipgram em-
7476 beddings, showing that the resulting embeddings perform reasonably on a task of solving
7477 phrasal analogies, e.g. *New York : New York Times :: Baltimore : Baltimore Sun*.

7478 14.8.2 Distributional-compositional hybrids

7479 To move beyond short multiword phrases, composition is necessary. A simple but sur-
7480 prisingly powerful approach is to represent a sentence with the average of its word em-
7481 beddings (Mitchell and Lapata, 2010). This can be considered a hybrid of the distribu-
7482 tional and compositional approaches to semantics: the word embeddings are computed
7483 distributionally, and then the sentence representation is computed by composition.

7484 The WORD2VEC approach can be stretched considerably further, embedding entire
7485 sentences using a model similar to skipgrams, in the “skip-thought” model of Kiros et al.
7486 (2015). Each sentence is *encoded* into a vector using a recurrent neural network: the encod-
7487 ing of sentence t is set to the RNN hidden state at its final token, $h_{M_t}^{(t)}$. This vector is then
7488 a parameter in a *decoder* model that is used to generate the previous and subsequent sen-
7489 tences: the decoder is another recurrent neural network, which takes the encoding of the
7490 neighboring sentence as an additional parameter in its recurrent update. (This **encoder-**
7491 **decoder model** is discussed at length in chapter 18.) The encoder and decoder are trained
7492 simultaneously from a likelihood-based objective, and the trained encoder can be used to
7493 compute a distributed representation of any sentence. Skip-thought can also be viewed
7494 as a hybrid of distributional and compositional approaches: the vector representation of
7495 each sentence is computed compositionally from the representations of the individual
7496 words, but the training objective is distributional, based on sentence co-occurrence across
7497 a corpus.

7498 **Autoencoders** are a variant of encoder-decoder models in which the decoder is trained
7499 to produce the same text that was originally encoded, using only the distributed encod-
7500 ing vector (Li et al., 2015). The encoding acts as a bottleneck, so that generalization is
7501 necessary if the model is to successfully fit the training data. In **denoising autoencoders**,
7502 the input is a corrupted version of the original sentence, and the auto-encoder must re-
7503 construct the uncorrupted original (Vincent et al., 2010; Hill et al., 2016). By interpolating
7504 between distributed representations of two sentences, $\alpha \mathbf{u}_i + (1 - \alpha) \mathbf{u}_j$, it is possible to gen-
7505 erate sentences that combine aspects of the two inputs, as shown in Figure 14.6 (Bowman
7506 et al., 2016).

7507 Autoencoders can also be applied to longer texts, such as paragraphs and documents.
7508 This enables applications such as **question answering**, which can be performed by match-
7509 ing the encoding of the question with encodings of candidate answers (Miao et al., 2016).

this was the only way
 it was the only way
 it was her turn to blink
 it was hard to tell
 it was time to move on
 he had to do it again
 they all looked at each other
 they all turned to look back
 they both turned to face him
they both turned and walked away

Figure 14.6: By interpolating between the distributed representations of two sentences (in bold), it is possible to generate grammatical sentences that combine aspects of both (Bowman et al., 2016)

7510 14.8.3 Supervised compositional methods

7511 Given a supervision signal, such as a label describing the sentiment or meaning of a sen-
 7512 tence, a wide range of compositional methods can be applied to compute a distributed
 7513 representation that then predicts the label. The simplest is to average the embeddings
 7514 of each word in the sentence, and pass this average through a feedforward neural net-
 7515 work (Iyyer et al., 2015). Convolutional and recurrent neural networks go further, with
 7516 the ability to effectively capturing multiword phenomena such as negation (Kalchbrenner
 7517 et al., 2014; Kim, 2014; Li et al., 2015; Tang et al., 2015). Another approach is to incorpo-
 7518 rate the syntactic structure of the sentence into a **recursive neural networks**, in which the
 7519 representation for each syntactic constituent is computed from the representations of its
 7520 children (Socher et al., 2012). However, in many cases, recurrent neural networks perform
 7521 as well or better than recursive networks (Li et al., 2015).

7522 Whether convolutional, recurrent, or recursive, a key question is whether supervised
 7523 sentence representations are task-specific, or whether a single supervised sentence repre-
 7524 sentation model can yield useful performance on other tasks. Wieting et al. (2015) train a
 7525 variety of sentence embedding models for the task of labeling pairs of sentences as **para-**
 7526 **phrases**. They show that the resulting sentence embeddings give good performance for
 7527 sentiment analysis. The **Stanford Natural Language Inference corpus** classifies sentence
 7528 pairs as **entailments** (the truth of sentence i implies the truth of sentence j), **contradictions**
 7529 (the truth of sentence i implies the falsity of sentence j), and neutral (i neither entails nor
 7530 contradicts j). Sentence embeddings trained on this dataset transfer to a wide range of
 7531 classification tasks (Conneau et al., 2017).

7532 14.8.4 Hybrid distributed-symbolic representations

7533 The power of distributed representations is in their generality: the distributed represen-
7534 tation of a unit of text can serve as a summary of its meaning, and therefore as the input
7535 for downstream tasks such as classification, matching, and retrieval. For example, dis-
7536 tributed sentence representations can be used to recognize the paraphrase relationship
7537 between closely related sentences like the following:

- 7538 (14.5) Donald thanked Vlad profusely.
7539 (14.6) Donald conveyed to Vlad his profound appreciation.
7540 (14.7) Vlad was showered with gratitude by Donald.

7541 Symbolic representations are relatively brittle to this sort of variation, but are better
7542 suited to describe individual entities, the things that they do, and the things that are done
7543 to them. In examples (14.5)-(14.7), we not only know that somebody thanked someone
7544 else, but we can make a range of inferences about what has happened between the en-
7545 tities named *Donald* and *Vlad*. Because distributed representations do not treat entities
7546 symbolically, they lack the ability to reason about the roles played by entities across a sen-
7547 tence or larger discourse.⁵ A hybrid between distributed and symbolic representations
7548 might give the best of both worlds: robustness to the many different ways of describing
7549 the same event, plus the expressiveness to support inferences about entities and the roles
7550 that they play.

7551 A “top-down” hybrid approach is to begin with logical semantics (of the sort de-
7552 scribed in the previous two chapters), and but replace the predefined lexicon with a set
7553 of distributional word clusters (Poon and Domingos, 2009; Lewis and Steedman, 2013). A
7554 “bottom-up” approach is to add minimal symbolic structure to existing distributed repre-
7555 sentations, such as vector representations for each entity (Ji and Eisenstein, 2015; Wiseman
7556 et al., 2016). This has been shown to improve performance on two problems that we will
7557 encounter in the following chapters: classification of **discourse relations** between adjac-
7558 ent sentences (chapter 16; Ji and Eisenstein, 2015), and **coreference resolution** of entity
7559 mentions (chapter 15; Wiseman et al., 2016; Ji et al., 2017). Research on hybrid seman-
7560 tic representations is still in an early stage, and future representations may deviate more
7561 boldly from existing symbolic and distributional approaches.

7562 Additional resources

7563 Turney and Pantel (2010) survey a number of facets of vector word representations, fo-
7564 cusing on matrix factorization methods. Schnabel et al. (2015) highlight problems with

⁵At a 2014 workshop on semantic parsing, this critique of distributed representations was expressed by Ray Mooney — a leading researcher in computational semantics — in a now well-known quote, “you can’t cram the meaning of a whole sentence into a single vector!”

7565 similarity-based evaluations of word embeddings, and present a novel evaluation that
 7566 controls for word frequency. Baroni et al. (2014) address linguistic issues that arise in
 7567 attempts to combine distributed and compositional representations.

7568 In bilingual and multilingual distributed representations, embeddings are estimated
 7569 for translation pairs or tuples, such as (*dog, perro, chien*). These embeddings can improve
 7570 machine translation (Zou et al., 2013; Klementiev et al., 2012), transfer natural language
 7571 processing models across languages (Täckström et al., 2012), and make monolingual word
 7572 embeddings more accurate (Faruqui and Dyer, 2014). A typical approach is to learn a pro-
 7573 jection that maximizes the correlation of the distributed representations of each element
 7574 in a translation pair, which can be obtained from a bilingual dictionary. Distributed rep-
 7575 resentations can also be linked to perceptual information, such as image features. Bruni
 7576 et al. (2014) use textual descriptions of images to obtain visual contextual information for
 7577 various words, which supplements traditional distributional context. Image features can
 7578 also be inserted as contextual information in log bilinear language models (Kiros et al.,
 7579 2014), making it possible to automatically generate text descriptions of images.

7580 Exercises

- 7581 1. Prove that the sum of probabilities of paths through a hierarchical softmax tree is
 7582 equal to one.
- 7583 2. In skipgram word embeddings, the negative sampling objective can be written as,

$$\mathcal{L} = \sum_{i \in \mathcal{V}} \sum_{j \in \mathcal{C}} \text{count}(i, j) \psi(i, j), \quad [14.29]$$

7584 with $\psi(i, j)$ is defined in Equation 14.23.

7585 Suppose we draw the negative samples from the empirical unigram distribution
 $\hat{p}(i) = p_{\text{unigram}}(i)$. First, compute the expectation of \mathcal{L} with respect to this probability.

7586 Next, take the derivative of this expectation with respect to the score of a single word
 7587 context pair $\sigma(\mathbf{u}_i \cdot \mathbf{v}_j)$, and solve for the pointwise mutual information $\text{PMI}(i, j)$. You
 7588 should be able to show that at the optimum, the PMI is a simple function of $\sigma(\mathbf{u}_i \cdot \mathbf{v}_j)$
 7589 and the number of negative samples.

- 7590 3. * In Brown clustering, prove that the cluster merge that maximizes the average mu-
 7591 tual information (Equation 14.13) also maximizes the log-likelihood objective (Equa-
 7592 tion 14.12).
- 7593 4. A simple way to compute a distributed phrase representation is to add up the dis-
 7594 tributed representations of the words in the phrase. Consider a sentiment analysis
 7595 model in which the predicted sentiment is, $\psi(\mathbf{w}) = \boldsymbol{\theta} \cdot (\sum_{m=1}^M \mathbf{x}_m)$, where \mathbf{x}_m is

the vector representation of word m . Prove that in such a model, the following two inequalities cannot both hold:

$$\psi(\text{good}) > \psi(\text{not good}) \quad [14.30]$$

$$\psi(\text{bad}) < \psi(\text{not bad}). \quad [14.31]$$

7593 Then construct a similar example pair for the case in which phrase representations
 7594 are the *average* of the word representations.

5. Now let's consider a slight modification to the prediction model in the previous problem:

$$\psi(\mathbf{w}) = \boldsymbol{\theta} \cdot \text{ReLU}\left(\sum_{m=1}^M \mathbf{x}_m\right) \quad [14.32]$$

7595 Show that in this case, it *is* possible to achieve the inequalities above. Your solution
 7596 should provide the weights $\boldsymbol{\theta}$ and the embeddings \mathbf{x}_{good} , \mathbf{x}_{bad} , and \mathbf{x}_{not} .

7597 For the next two problems, download a set of pre-trained word embeddings, such as the
 7598 WORD2VEC or polyglot embeddings.

- 7599 6. Use cosine similarity to find the most similar words to: *dog*, *whale*, *before*, *however*,
 7600 *fabricate*.
- 7601 7. Use vector addition and subtraction to compute target vectors for the analogies be-
 7602 low. After computing each target vector, find the top three candidates by cosine
 7603 similarity.
- 7604 • *dog:puppy :: cat: ?*
 - 7605 • *speak:speaker :: sing: ?*
 - 7606 • *France:French :: England: ?*
 - 7607 • *France:wine :: England: ?*

7608 The remaining problems will require you to build a classifier and test its properties. Pick
 7609 a multi-class text classification dataset, such as RCV1⁶). Divide your data into training
 7610 (60%), development (20%), and test sets (20%), if no such division already exists.

⁶http://www.ai.mit.edu/projects/jmlr/papers/volume5/lewis04a/lyrl2004_rcv1v2_README.htm

- 7611 8. Train a convolutional neural network, with inputs set to pre-trained word embed-
7612 dings from the previous problem. Use a special, fine-tuned embedding for out-of-
7613 vocabulary words. Train until performance on the development set does not im-
7614 prove. You can also use the development set to tune the model architecture, such
7615 as the convolution width and depth. Report *F-MEASURE* and accuracy, as well as
7616 training time.
- 7617 9. Now modify your model from the previous problem to fine-tune the word embed-
7618 dings. Report *F-MEASURE*, accuracy, and training time.
- 7619 10. Try a simpler approach, in which word embeddings in the document are averaged,
7620 and then this average is passed through a feed-forward neural network. Again, use
7621 the development data to tune the model architecture. How close is the accuracy to
7622 the convolutional networks from the previous problems?

7623

Chapter 15

7624

Reference Resolution

7625 References are one of the most noticeable forms of linguistic ambiguity, afflicting not just
7626 automated natural language processing systems, but also fluent human readers. Warnings
7627 to avoid “ambiguous pronouns” are ubiquitous in manuals and tutorials on writing
7628 style. But referential ambiguity is not limited to pronouns, as shown in the text in Fig-
7629 ure 15.1. Each of the bracketed substrings refers to an entity that is introduced earlier
7630 in the passage. These references include the pronouns *he* and *his*, but also the shortened
7631 name *Cook*, and **nominals** such as *the firm* and *the firm’s biggest growth market*.

7632 **Reference resolution** subsumes several subtasks. This chapter will focus on **corefer-
7633 ence resolution**, which is the task of grouping spans of text that refer to a single underly-
7634 ing entity, or, in some cases, a single event: for example, the spans *Tim Cook*, *he*, and *Cook*
7635 are all **coreferent**. These individual spans are called **mentions**, because they mention an
7636 entity; the entity is sometimes called the **referent**. Each mention has a set of **antecedents**,
7637 which are preceding mentions that are coreferent; for the first mention of an entity, the an-
7638 tecedent set is empty. The task of **pronominal anaphora resolution** requires identifying
7639 only the antecedents of pronouns. In **entity linking**, references are resolved not to other
7640 spans of text, but to entities in a knowledge base. This task is discussed in chapter 17.

7641 Coreference resolution is a challenging problem for several reasons. Resolving differ-
7642 ent types of **referring expressions** requires different types of reasoning: the features and
7643 methods that are useful for resolving pronouns are different from those that are useful
7644 to resolve names and nominals. Coreference resolution involves not only linguistic rea-
7645 soning, but also world knowledge and pragmatics: you may not have known that China
7646 was Apple’s biggest growth market, but it is likely that you effortlessly resolved this ref-
7647 erence while reading the passage in Figure 15.1.¹ A further challenge is that coreference

¹This interpretation is based in part on the assumption that a **cooperative** author would not use the expression *the firm’s biggest growth market* to refer to an entity not yet mentioned in the article (Grice, 1975). **Pragmatics** is the discipline of linguistics concerned with the formalization of such assumptions (Huang,

- (15.1) *[[Apple Inc] Chief Executive Tim Cook] has jetted into [China] for talks with government officials as [he] seeks to clear up a pile of problems in [[the firm] 's biggest growth market] ... [Cook] is on [his] first trip to [the country] since taking over...*

Figure 15.1: Running example (Yee and Jones, 2012). Coreferring entity mentions are underlined and bracketed.

7648 resolution decisions are often entangled: each mention adds information about the entity,
 7649 which affects other coreference decisions. This means that coreference resolution must
 7650 be addressed as a structure prediction problem. But as we will see, there is no dynamic
 7651 program that allows the space of coreference decisions to be searched efficiently.

7652 15.1 Forms of referring expressions

7653 There are three main forms of referring expressions — pronouns, names, and nominals.

7654 15.1.1 Pronouns

7655 Pronouns are a closed class of words that are used for references. A natural way to think
 7656 about pronoun resolution is SMASH (Kehler, 2007):

- 7657 • Search for candidate antecedents;
- 7658 • Match against hard agreement constraints;
- 7659 • And Select using Heuristics, which are “soft” constraints such as recency, syntactic
 7660 prominence, and parallelism.

7661 15.1.1.1 Search

7662 In the search step, candidate antecedents are identified from the preceding text or speech.²
 7663 Any noun phrase can be a candidate antecedent, and pronoun resolution usually requires

2015).

²Pronouns whose referents come later are known as **cataphora**, as in this example from Márquez (1970):

- (15.1) Many years later, as [he] faced the firing squad, [Colonel Aureliano Buendía] was to remember that distant afternoon when his father took him to discover ice.

7664 parsing the text to identify all such noun phrases.³ Filtering heuristics can help to prune
 7665 the search space to noun phrases that are likely to be coreferent (Lee et al., 2013; Durrett
 7666 and Klein, 2013). In nested noun phrases, mentions are generally considered to be the
 7667 largest unit with a given head word: thus, *Apple Inc. Chief Executive Tim Cook* would be
 7668 included as a mention, but *Tim Cook* would not, since they share the same head word,
 7669 *Cook*.

7670 15.1.1.2 Matching constraints for pronouns

7671 References and their antecedents must agree on semantic features such as number, person,
 7672 gender, and animacy. Consider the pronoun *he* in this passage from the running example:

7673 (15.2) Tim Cook has jetted in for talks with officials as [he] seeks to clear up a pile of
 7674 problems...

7675 The pronoun and possible antecedents have the following features:

- 7676 • *he*: singular, masculine, animate, third person
- 7677 • *officials*: plural, animate, third person
- 7678 • *talks*: plural, inanimate, third person
- 7679 • *Tim Cook*: singular, masculine, animate, third person

7680 The SMASH method searches backwards from *he*, discarding *officials* and *talks* because they
 7681 do not satisfy the agreements constraints.

7682 Another source of constraints comes from syntax — specifically, from the phrase struc-
 7683 ture trees discussed in chapter 10. Consider a parse tree in which both *x* and *y* are phrasal
 7684 constituents. The constituent *x* **c-commands** the constituent *y* iff the first branching node
 7685 above *x* also dominates *y*. For example, in Figure 15.2a, *Abigail* c-commands *her*, because
 7686 the first branching node above *Abigail*, *S*, also dominates *her*. Now, if *x* c-commands *y*,
 7687 **government and binding theory** (Chomsky, 1982) states that *y* can refer to *x* only if it is
 7688 a **reflexive pronoun** (e.g., *herself*). Furthermore, if *y* is a reflexive pronoun, then its an-
 7689 tecedent must c-command it. Thus, in Figure 15.2a, *her* cannot refer to *Abigail*; conversely,
 7690 if we replace *her* with *herself*, then the reflexive pronoun *must* refer to *Abigail*, since this is
 7691 the only candidate antecedent that c-commands it.

7692 Now consider the example shown in Figure 15.2b. Here, *Abigail* does not c-command
 7693 *her*, but *Abigail's mom* does. Thus, *her* can refer to *Abigail* — and we cannot use reflexive

³In the OntoNotes coreference annotations, verbs can also be antecedents, if they are later referenced by nominals (Pradhan et al., 2011):

(15.1) Sales of passenger cars [grew] 22%. [The strong growth] followed year-to-year increases.




Figure 15.2: In (a), *Abigail* c-commands *her*; in (b), *Abigail* does not c-command *her*, but *Abigail's mom* does; in (c), the scope of *Abigail* is limited by the S non-terminal, so that *she* or *her* can bind to *Abigail*, but not both.

7694 *herself* in this context, unless we are talking about *Abigail*'s mom. However, *her* does not
 7695 have to refer to *Abigail*. Finally, Figure 15.2c shows how these constraints are limited.
 7696 In this case, the pronoun *she* can refer to *Abigail*, because the S non-terminal puts *Abigail*
 7697 outside the domain of *she*. Similarly, *her* can also refer to *Abigail*. But *she* and *her* cannot be
 7698 coreferent, because *she* c-commands *her*.

7699 15.1.1.3 Heuristics

7700 After applying constraints, heuristics are applied to select among the remaining candidates.
 7701 Recency is a particularly strong heuristic. All things equal, readers will prefer
 7702 the more recent referent for a given pronoun, particularly when comparing referents that
 7703 occur in different sentences. Jurafsky and Martin (2009) offer the following example:

- 7704 (15.3) The doctor found an old map in the captain's chest. Jim found an even older map
 7705 hidden on the shelf. [It] described an island.

7706 Readers are expected to prefer the older map as the referent for the pronoun *it*.

7707 However, subjects are often preferred over objects, and this can contradict the preference
 7708 for recency when two candidate referents are in the same sentence. For example,

- 7709 (15.4) Asha loaned Mei a book on Spanish. [She] is always trying to help people.

7710 Here, we may prefer to link *she* to *Asha* rather than *Mei*, because of *Asha*'s position in the
 7711 subject role of the preceding sentence. (Arguably, this preference would not be strong
 7712 enough to select *Asha* if the second sentence were *She is visiting Valencia next month*.)

7713 A third heuristic is parallelism:

- 7714 (15.5) Asha loaned Mei a book on Spanish. Olya loaned [her] a book on Portuguese.

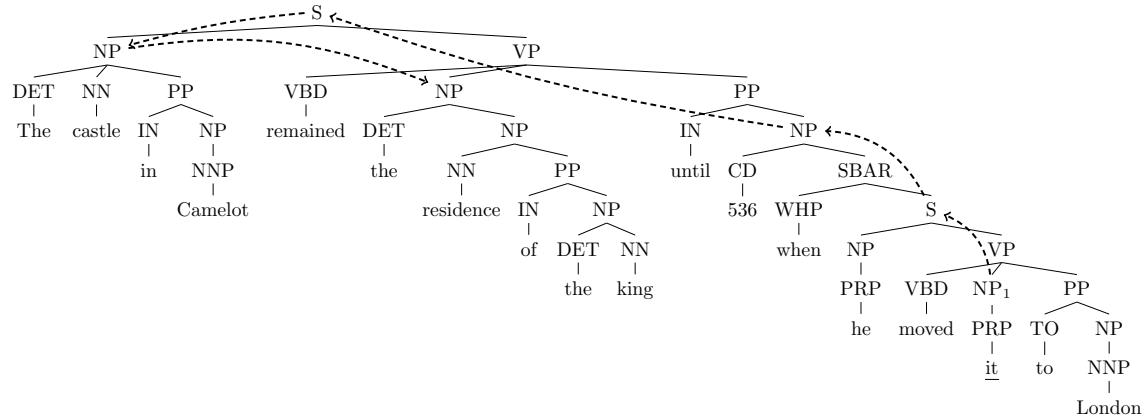


Figure 15.3: Left-to-right breadth-first tree traversal (Hobbs, 1978), indicating that the search for an antecedent for *it* (NP₁) would proceed in the following order: 536; *the castle in Camelot*; *the residence of the king*; *Camelot*; *the king*. Hobbs (1978) proposes semantic constraints to eliminate 536 and *the castle in Camelot* as candidates, since they are unlikely to be the direct object of the verb *move*.

7715 Here *Mei* is preferred as the referent for *her*, contradicting the preference for the subject
 7716 *Asha* in the preceding sentence.

7717 The recency and subject role heuristics can be unified by traversing the document in
 7718 a syntax-driven fashion (Hobbs, 1978): each preceding sentence is traversed breadth-first,
 7719 left-to-right (Figure 15.3). This heuristic successfully handles (15.4): *Asha* is preferred as
 7720 the referent for *she* because the subject NP is visited first. It also handles (15.3): the older
 7721 map is preferred as the referent for *it* because the more recent sentence is visited first. (An
 7722 alternative unification of recency and syntax is proposed by **centering theory** (Grosz et al.,
 7723 1995), which is discussed in detail in chapter 16.)

7724 In early work on reference resolution, the number of heuristics was small enough that
 7725 a set of numerical weights could be set by hand (Lappin and Leass, 1994). More recent
 7726 work uses machine learning to quantify the importance of each of these factors. However,
 7727 pronoun resolution cannot be completely solved by constraints and heuristics alone. This
 7728 is shown by the classic example pair (Winograd, 1972):

7729 (15.6) The [city council] denied [the protesters] a permit because [they] advocated / feared
 7730 violence.

7731 Without reasoning about the motivations of the city council and protesters, it is unlikely
 7732 that any system could correctly resolve both versions of this example.

7733 **15.1.1.4 Non-referential pronouns**

7734 While pronouns are generally used for reference, they need not refer to entities. The fol-
 7735 lowing examples show how pronouns can refer to propositions, events, and speech acts.

- 7736 (15.7) They told me that I was too ugly for show business, but I didn't believe [it].
 7737 (15.8) Asha saw Babak get angry, and I saw [it] too.
 7738 (15.9) Asha said she worked in security. I suppose [that]'s one way to put it.

7739 These forms of reference are generally not annotated in large-scale coreference resolution
 7740 datasets such as OntoNotes (Pradhan et al., 2011).

7741 Pronouns may also have **generic referents**:

- 7742 (15.10) A poor carpenter blames [her] tools.
 7743 (15.11) On the moon, [you] have to carry [your] own oxygen.
 7744 (15.12) Every farmer who owns a donkey beats [it]. (Geach, 1962)

7745 In the OntoNotes dataset, coreference is not annotated for generic referents, even in cases
 7746 like these examples, in which the same generic entity is mentioned multiple times.

7747 Some pronouns do not refer to anything at all:

- 7748 (15.13) *[It]'s raining.*
 [Il] pleut. (Fr)
 7749 (15.14) [It] 's money that she's really after.
 7750 (15.15) [It] is too bad that we have to work so hard.

7751 How can we automatically distinguish these usages of *it* from referential pronouns?
 7752 Consider the the difference between the following two examples (Bergsma et al., 2008):

- 7753 (15.16) You can make [it] in advance.
 7754 (15.17) You can make [it] in showbiz.

7755 In the second example, the pronoun *it* is non-referential. One way to see this is by substi-
 7756 tuting another pronoun, like *them*, into these examples:

- 7757 (15.18) You can make [them] in advance.
 7758 (15.19) ? You can make [them] in showbiz.

7759 The questionable grammaticality of the second example suggests that *it* is not referential.
 7760 Bergsma et al. (2008) operationalize this idea by comparing distributional statistics for the

7761 *n*-grams around the word *it*, testing how often other pronouns or nouns appear in the
7762 same context. In cases where nouns and other pronouns are infrequent, the *it* is unlikely
7763 to be referential.

7764 **15.1.2 Proper Nouns**

7765 If a proper noun is used as a referring expression, it often corefers with another proper
7766 noun, so that the coreference problem is simply to determine whether the two names
7767 match. Subsequent proper noun references often use a shortened form, as in the running
7768 example (Figure 15.1):

7769 (15.20) Apple Inc Chief Executive [Tim Cook] has jetted into China ... [Cook] is on his
7770 first business trip to the country ...

7771 A typical solution for proper noun coreference is to match the syntactic **head words**
7772 of the reference with the referent. In § 10.5.2, we saw that the head word of a phrase can
7773 be identified by applying head percolation rules to the phrasal parse tree; alternatively,
7774 the head can be identified as the root of the dependency subtree covering the name. For
7775 sequences of proper nouns, the head word will be the final token.

7776 There are a number of caveats to the practice of matching head words of proper nouns.

- 7777 • In the European tradition, family names tend to be more specific than given names,
7778 and family names usually come last. However, other traditions have other practices:
7779 for example, in Chinese names, the family name typically comes first; in Japanese,
7780 honorifics come after the name, as in *Nobu-San* (*Mr. Nobu*).
- 7781 • In organization names, the head word is often not the most informative, as in *Georgia*
7782 *Tech* and *Virginia Tech*. Similarly, *Lebanon* does not refer to the same entity as *Southern Lebanon*, necessitating special rules for the specific case of geographical modi-
7783 fiers (Lee et al., 2011).
- 7785 • Proper nouns can be nested, as in [*the CEO of [Microsoft]*], resulting in head word
7786 match without coreference.

7787 Despite these difficulties, proper nouns are the easiest category of references to re-
7788 solve (Stoyanov et al., 2009). In machine learning systems, one solution is to include a
7789 range of matching features, including exact match, head match, and string inclusion. In
7790 addition to matching features, competitive systems (e.g., Bengtson and Roth, 2008) in-
7791 clude large lists, or **gazetteers**, of acronyms (e.g., *the National Basketball Association/NBA*),
7792 demonyms (e.g., *the Israelis/Israel*), and other aliases (e.g., *the Georgia Institute of Technol-*
7793 *ogy/Georgia Tech*).

7794 **15.1.3 Nominals**

7795 In coreference resolution, noun phrases that are neither pronouns nor proper nouns are
 7796 referred to as **nominals**. In the running example (Figure 15.1), nominal references include:
 7797 *the firm (Apple Inc); the firm's biggest growth market (China); and the country (China)*.

7798 Nominals are especially difficult to resolve (Denis and Baldridge, 2007; Durrett and
 7799 Klein, 2013), and the examples above suggest why this may be the case: world knowledge
 7800 is required to identify *Apple Inc* as a *firm*, and *China* as a *growth market*. Other difficult
 7801 examples include the use of colloquial expressions, such as coreference between *Clinton*
 7802 *campaign officials* and *the Clinton camp* (Soon et al., 2001).

7803 **15.2 Algorithms for coreference resolution**

The ground truth training data for coreference resolution is a set of mention sets, where all mentions within each set refer to a single entity.⁴ In the running example from Figure 15.1, the ground truth coreference annotation is:

$$c_1 = \{Apple\ Inc_{1:2}, the\ firm_{27:28}\} \quad [15.1]$$

$$c_2 = \{Apple\ Inc\ Chief\ Executive\ Tim\ Cook_{1:6}, he_{17}, Cook_{33}, his_{36}\} \quad [15.2]$$

$$c_3 = \{China_{10}, the\ firm\ 's\ biggest\ growth\ market_{27:32}, the\ country_{40:41}\} \quad [15.3]$$

7804 Each row specifies the token spans that mention an entity. (“Singleton” entities, which are
 7805 mentioned only once (e.g., *talks, government officials*), are excluded from the annotations.)
 7806 Equivalently, if given a set of M mentions, $\{m_i\}_{i=1}^M$, each mention i can be assigned to a
 7807 cluster z_i , where $z_i = z_j$ if i and j are coreferent. The cluster assignments z are invariant
 7808 under permutation. The unique clustering associated with the assignment z is written
 7809 $c(z)$.

7810 **Mention identification** The task of identifying mention spans for coreference resolution
 7811 is often performed by applying a set of heuristics to the phrase structure parse of each
 7812 sentence. A typical approach is to start with all noun phrases and named entities, and
 7813 then apply filtering rules to remove nested noun phrases with the same head (e.g., [*Apple*
 7814 *CEO [Tim Cook]*]), numeric entities (e.g., [*100 miles*], [*97%*]), non-referential *it*, etc (Lee
 7815 et al., 2013; Durrett and Klein, 2013). In general, these deterministic approaches err in
 7816 favor of recall, since the mention clustering component can choose to ignore false positive
 7817 mentions, but cannot recover from false negatives. An alternative is to consider all spans

⁴In many annotations, the term **markable** is used to refer to spans of text that can *potentially* mention an entity. The set of markables includes non-referential pronouns, which does not mention any entity. Part of the job of the coreference system is to avoid incorrectly linking these non-referential markables to any mention chains.

7818 (up to some finite length) as candidate mentions, performing mention identification and
 7819 clustering jointly (Daumé III and Marcu, 2005; Lee et al., 2017).

7820 **Mention clustering** The overwhelming majority of research on coreference resolution
 7821 addresses the subtask of mention clustering, and this will be the focus of the remainder of
 7822 this chapter. There are two main sets of approaches. In *mention-based models*, the scoring
 7823 function for a coreference clustering decomposes over pairs of mentions. These pairwise
 7824 decisions are then aggregated, using a clustering heuristic. Mention-based coreference
 7825 clustering can be treated as a fairly direct application of supervised classification or rank-
 7826 ing. However, the mention-pair locality assumption can result in incoherent clusters, like
 7827 $\{Hillary Clinton \leftarrow Clinton \leftarrow Mr Clinton\}$, in which the pairwise links score well, but the
 7828 overall result is unsatisfactory. *Entity-based models* address this issue by scoring entities
 7829 holistically. This can make inference more difficult, since the number of possible entity
 7830 groupings is exponential in the number of mentions.

7831 15.2.1 Mention-pair models

7832 In the **mention-pair model**, a binary label $y_{i,j} \in \{0, 1\}$ is assigned to each pair of mentions
 7833 (i, j) , where $i < j$. If i and j corefer ($z_i = z_j$), then $y_{i,j} = 1$; otherwise, $y_{i,j} = 0$. The
 7834 mention *he* in Figure 15.1 is preceded by five other mentions: (1) *Apple Inc*; (2) *Apple Inc*
 7835 *Chief Executive Tim Cook*; (3) *China*; (4) *talks*; (5) *government officials*. The correct mention
 7836 pair labeling is $y_{2,6} = 1$ and $y_{i \neq 2,6} = 0$ for all other i . If a mention j introduces a new entity,
 7837 such as mention 3 in the example, then $y_{i,j} = 0$ for all i . The same is true for “mentions”
 7838 that do not refer to any entity, such as non-referential pronouns. If mention j refers to an
 7839 entity that has been mentioned more than once, then $y_{i,j} = 1$ for all $i < j$ that mention the
 7840 referent.

7841 By transforming coreference into a set of binary labeling problems, the mention-pair
 7842 model makes it possible to apply an off-the-shelf binary classifier (Soon et al., 2001). This
 7843 classifier is applied to each mention j independently, searching backwards from j until
 7844 finding an antecedent i which corefers with j with high confidence. After identifying a
 7845 single **antecedent**, the remaining mention pair labels can be computed by transitivity: if
 7846 $y_{i,j} = 1$ and $y_{j,k} = 1$, then $y_{i,k} = 1$.

7847 Since the ground truth annotations give entity chains c but not individual mention-
 7848 pair labels y , an additional heuristic must be employed to convert the labeled data into
 7849 training examples for classification. A typical approach is to generate at most one pos-
 7850 itive labeled instance $y_{a_j,j} = 1$ for mention j , where a_j is the index of the most recent
 7851 antecedent, $a_j = \max\{i : i < j \wedge z_i = z_j\}$. Negative labeled instances are generated for
 7852 all for all $i \in \{a_j + 1, \dots, j\}$. In the running example, the most recent antecedent of the
 7853 pronoun *he* is $a_6 = 2$, so the training data would be $y_{2,6} = 1$ and $y_{3,6} = y_{4,6} = y_{5,6} = 0$.

7854 The variable $y_{1,6}$ is not part of the training data, because the first mention appears before
 7855 the true antecedent $a_6 = 2$.

7856 **15.2.2 Mention-ranking models**

In **mention ranking** (Denis and Baldridge, 2007), the classifier learns to identify a single antecedent $a_i \in \{\epsilon, 1, 2, \dots, i-1\}$ for each referring expression i ,

$$\hat{a}_i = \operatorname{argmax}_{a \in \{\epsilon, 1, 2, \dots, i-1\}} \psi_M(a, i), \quad [15.4]$$

7857 where $\psi_M(a, i)$ is a score for the mention pair (a, i) . If $a = \epsilon$, then mention i does not refer
 7858 to any previously-introduced entity — it is not **anaphoric**. Mention-ranking is similar to
 7859 the mention-pair model, but all candidates are considered simultaneously, and at most
 7860 a single antecedent is selected. The mention-ranking model explicitly accounts for the
 7861 possibility that mention i is not anaphoric, through the score $\psi_M(\epsilon, i)$. The determination
 7862 of anaphoricity can be made by a special classifier in a preprocessing step, so that non- ϵ
 7863 antecedents are identified only for spans that are determined to be anaphoric (Denis and
 7864 Baldridge, 2008).

7865 As a learning problem, ranking can be trained using the same objectives as in dis-
 7866 criminative classification. For each mention i , we can define a gold antecedent a_i^* , and an
 7867 associated loss, such as the hinge loss, $\ell_i = (1 - \psi_M(a_i^*, i) + \psi_M(\hat{a}, i))_+$ or the negative
 7868 log-likelihood, $\ell_i = -\log p(a_i^* | i; \theta)$. (For more on learning to rank, see § 17.1.1.) But as
 7869 with the mention-pair model, there is a mismatch between the labeled data, which comes
 7870 in the form of mention sets, and the desired supervision, which would indicate the spe-
 7871 cific antecedent of each mention. The antecedent variables $\{a_i\}_{i=1}^M$ relate to the mention
 7872 sets in a many-to-one mapping: each set of antecedents induces a single clustering, but a
 7873 clustering can correspond to many different settings of antecedent variables.

A heuristic solution is to set $a_i^* = \max\{j : j < i \wedge z_j = z_i\}$, the most recent mention in
 the same cluster as i . But the most recent mention may not be the most informative: in the
 running example, the most recent antecedent of the mention *Cook* is the pronoun *he*, but
 a more useful antecedent is the earlier mention *Apple Inc Chief Executive Tim Cook*. Rather
 than selecting a specific antecedent to train on, the antecedent can be treated as a latent
 variable, in the manner of the **latent variable perceptron** from § 12.4.2 (Fernandes et al.,

2014):

$$\hat{\mathbf{a}} = \operatorname{argmax}_{\mathbf{a}} \sum_{i=1}^M \psi_M(a_i, i) \quad [15.5]$$

$$\mathbf{a}^* = \operatorname{argmax}_{\mathbf{a} \in \mathcal{A}(c)} \sum_{i=1}^M \psi_M(a_i, i) \quad [15.6]$$

$$\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} + \sum_{i=1}^M \frac{\partial L}{\partial \boldsymbol{\theta}} \psi_M(a_i^*, i) - \sum_{i=1}^M \frac{\partial L}{\partial \boldsymbol{\theta}} \psi_M(\hat{a}_i, i) \quad [15.7]$$

where $\mathcal{A}(c)$ is the set of antecedent structures that is compatible with the ground truth coreference clustering c . Another alternative is to sum over all the conditional probabilities of antecedent structures that are compatible with the ground truth clustering (Durrett and Klein, 2013; Lee et al., 2017). For the set of mention \mathbf{m} , we compute the following probabilities:

$$p(c | \mathbf{m}) = \sum_{\mathbf{a} \in \mathcal{A}(c)} p(\mathbf{a} | \mathbf{m}) = \sum_{\mathbf{a} \in \mathcal{A}(c)} \prod_{i=1}^M p(a_i | i, \mathbf{m}) \quad [15.8]$$

$$p(a_i | i, \mathbf{m}) = \frac{\exp(\psi_M(a_i, i))}{\sum_{a' \in \{\epsilon, 1, 2, \dots, i-1\}} \exp(\psi_M(a', i))}. \quad [15.9]$$

7874 This objective rewards models that assign high scores for all valid antecedent structures.
 7875 In the running example, this would correspond to summing the probabilities of the two
 7876 valid antecedents for *Cook, he* and *Apple Inc Chief Executive Tim Cook*. In one of the exer-
 7877 cises, you will compute the number of valid antecedent structures for a given clustering.

7878 15.2.3 Transitive closure in mention-based models

A problem for mention-based models is that individual mention-level decisions may be incoherent. Consider the following mentions:

$$m_1 = \text{Hillary Clinton} \quad [15.10]$$

$$m_2 = \text{Clinton} \quad [15.11]$$

$$m_3 = \text{Bill Clinton} \quad [15.12]$$

7879 A mention-pair system might predict $\hat{y}_{1,2} = 1, \hat{y}_{2,3} = 1, \hat{y}_{1,3} = 0$. Similarly, a mention-
 7880 ranking system might choose $\hat{a}_2 = 1$ and $\hat{a}_3 = 2$. Logically, if mentions 1 and 3 are both
 7881 coreferent with mention 2, then all three mentions must refer to the same entity. This
 7882 constraint is known as **transitive closure**.

Transitive closure can be applied *post hoc*, revising the independent mention-pair or mention-ranking decisions. However, there are many possible ways to enforce transitive closure: in the example above, we could set $\hat{y}_{1,3} = 1$, or $\hat{y}_{1,2} = 0$, or $\hat{y}_{2,3} = 0$. For documents with many mentions, there may be many violations of transitive closure, and many possible fixes. Transitive closure can be enforced by always adding edges, so that $\hat{y}_{1,3} = 1$ is preferred (e.g., Soon et al., 2001), but this can result in overclustering, with too many mentions grouped into too few entities.

Mention-pair coreference resolution can be viewed as a constrained optimization problem,

$$\begin{aligned} \max_{\mathbf{y} \in \{0,1\}^M} \quad & \sum_{j=1}^M \sum_{i=1}^j \psi_M(i, j) \times y_{i,j} \\ \text{s.t.} \quad & y_{i,j} + y_{j,k} - 1 \leq y_{i,k}, \quad \forall i < j < k, \end{aligned}$$

with the constraint enforcing transitive closure. This constrained optimization problem is equivalent to graph partitioning with positive and negative edge weights: construct a graph where the nodes are mentions, and the edges are the pairwise scores $\psi_M(i, j)$; the goal is to partition the graph so as to maximize the sum of the edge weights between all nodes within the same partition (McCallum and Wellner, 2004). This problem is NP-hard, motivating approximations such as correlation clustering (Bansal et al., 2004) and **integer linear programming** (Klenner, 2007; Finkel and Manning, 2008, also see § 13.2.2).

15.2.4 Entity-based models

A weakness of mention-based models is that they treat coreference resolution as a classification or ranking problem, when it is really a clustering problem: the goal is to group the mentions together into clusters that correspond to the underlying entities. Entity-based approaches attempt to identify these clusters directly. Such methods require a scoring function at the entity level, measuring whether each set of mentions is internally consistent. Coreference resolution can then be viewed as the following optimization,

$$\max_{\mathbf{z}} \quad \sum_{e=1} \psi_E(\{i : z_i = e\}), \tag{15.13}$$

where z_i indicates the entity referenced by mention i , and $\psi_E(\{i : z_i = e\})$ is a scoring function applied to all mentions i that are assigned to entity e .

Entity-based coreference resolution is conceptually similar to the unsupervised clustering problems encountered in chapter 5: the goal is to obtain clusters of mentions that are internally coherent. The number of possible clusterings is the **Bell number**, which is

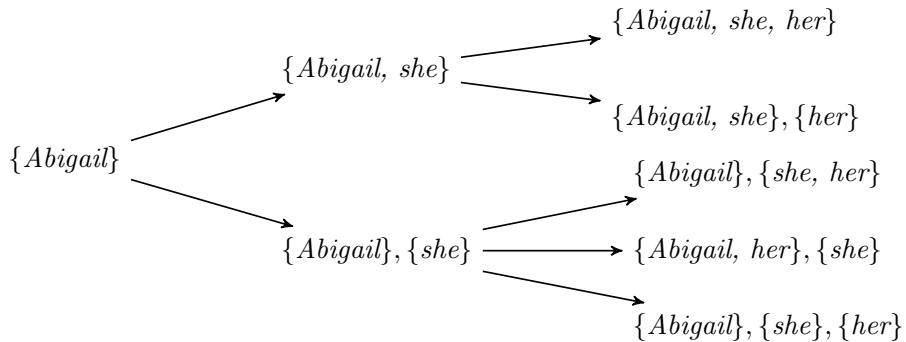


Figure 15.4: The Bell Tree for the sentence *Abigail hopes she speaks with her*. Which paths are excluded by the syntactic constraints mentioned in § 15.1.1?

defined by the following recurrence (Bell, 1934; Luo et al., 2004),

$$B_n = \sum_{k=0}^{n-1} B_k \binom{n-1}{k} = \frac{1}{e} \sum_{k=0}^{\infty} \frac{k^n}{k!}. \quad [15.14]$$

This recurrence is illustrated by the Bell tree, which is applied to a short coreference problem in Figure 15.4. The Bell number B_n grows exponentially with n , making exhaustive search of the space of clusterings impossible. For this reason, entity-based coreference resolution typically involves incremental search, in which clustering decisions are based on local evidence, in the hope of approximately optimizing the full objective in Equation 15.13. This approach is sometimes called **cluster ranking**, in contrast to mention ranking.

***Generative models of coreference** Entity-based cooreference can be approached through probabilistic **generative models**, in which the mentions in the document are conditioned on a set of latent entities (Haghghi and Klein, 2007, 2010). An advantage of these methods is that they can be learned from unlabeled data (Poon and Domingos, 2008, e.g.); a disadvantage is that probabilistic inference is required not just for learning, but also for prediction. Furthermore, generative models require independence assumptions that are difficult to apply in coreference resolution, where the diverse and heterogeneous features do not admit an easy decomposition into mutually independent subsets.

15.2.4.1 Incremental cluster ranking

The SMASH method (§ 15.1.1) can be extended to entity-based coreference resolution by building up coreference clusters while moving through the document (Cardie and Wagstaff, 1999). At each mention, the algorithm iterates backwards through possible antecedent

7919 clusters; but unlike SMASH, a cluster is selected only if *all* members of its cluster are compatible
 7920 with the current mention. As mentions are added to a cluster, so are their features
 7921 (e.g., gender, number, animacy). In this way, incoherent chains like *{Hillary Clinton, Clinton, Bill Clinton}*
 7922 can be avoided. However, an incorrect assignment early in the document — a **search error**
 7923 — might lead to a cascade of errors later on.

7924 More sophisticated search strategies can help to ameliorate the risk of search errors.
 7925 One approach is **beam search** (§ 11.3), in which a set of hypotheses is maintained through-
 7926 out search. Each hypothesis represents a path through the Bell tree (Figure 15.4). Hy-
 7927 potheses are “expanded” either by adding the next mention to an existing cluster, or by
 7928 starting a new cluster. Each expansion receives a score, based on Equation 15.13, and the
 7929 top K hypotheses are kept on the beam as the algorithm moves to the next step.

7930 Incremental cluster ranking can be made more accurate by performing multiple passes
 7931 over the document, applying rules (or “sieves”) with increasing recall and decreasing
 7932 precision at each pass (Lee et al., 2013). In the early passes, coreference links are pro-
 7933 posed only between mentions that are highly likely to corefer (e.g., exact string match
 7934 for full names and nominals). Information can then be shared among these mentions,
 7935 so that when more permissive matching rules are applied later, agreement is preserved
 7936 across the entire cluster. For example, in the case of *{Hillary Clinton, Clinton, she}*, the
 7937 name-matching sieve would link *Clinton* and *Hillary Clinton*, and the pronoun-matching
 7938 sieve would then link *she* to the combined cluster. A deterministic multi-pass system
 7939 won nearly every track of the 2011 CoNLL shared task on coreference resolution (Prad-
 7940 han et al., 2011). Given the dominance of machine learning in virtually all other areas
 7941 of natural language processing — and more than fifteen years of prior work on machine
 7942 learning for coreference — this was a surprising result, even if learning-based methods
 7943 have subsequently regained the upper hand (e.g., Lee et al., 2017, the state-of-the-art at
 7944 the time of this writing).

7945 15.2.4.2 Incremental perceptron

Incremental coreference resolution can be learned with the **incremental perceptron**, as described in § 11.3.2. At mention i , each hypothesis on the beam corresponds to a clustering of mentions $1 \dots i - 1$, or equivalently, a path through the Bell tree up to position $i - 1$. As soon as none of the hypotheses on the beam are compatible with the gold coreference clustering, a perceptron update is made (Daumé III and Marcu, 2005). For concreteness, consider a linear cluster ranking model,

$$\psi_E(\{i : z_i = e\}) = \sum_{i:z_i=e} \theta \cdot f(i, \{j : j < i \wedge z_j = e\}), \quad [15.15]$$

7946 where the score for each cluster is computed as the sum of scores of all mentions that are
 7947 linked into the cluster, and $f(i, \emptyset)$ is a set of features for the non-anaphoric mention that
 7948 initiates the cluster.

7949 Using Figure 15.4 as an example, suppose that the ground truth is,

$$\mathbf{c}^* = \{\text{Abigail}, \text{her}\}, \{\text{she}\}, \quad [15.16]$$

7950 but that with a beam of size one, the learner reaches the hypothesis,

$$\hat{\mathbf{c}} = \{\text{Abigail}, \text{she}\}. \quad [15.17]$$

This hypothesis is incompatible with \mathbf{c}^* , so an update is needed:

$$\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} + \mathbf{f}(\mathbf{c}^*) - \mathbf{f}(\hat{\mathbf{c}}) \quad [15.18]$$

$$= \boldsymbol{\theta} + (\mathbf{f}(\text{Abigail}, \emptyset) + \mathbf{f}(\text{she}, \emptyset)) - (\mathbf{f}(\text{Abigail}, \emptyset) + \mathbf{f}(\text{she}, \{\text{Abigail}\})) \quad [15.19]$$

$$= \boldsymbol{\theta} + \mathbf{f}(\text{she}, \emptyset) - \mathbf{f}(\text{she}, \{\text{Abigail}\}). \quad [15.20]$$

7951 This style of incremental update can also be applied to a margin loss between the gold
 7952 clustering and the top clustering on the beam. By backpropagating from this loss, it is also
 7953 possible to train a more complicated scoring function, such as a neural network in which
 7954 the score for each entity is a function of embeddings for the entity mentions (Wiseman
 7955 et al., 2015).

7956 15.2.4.3 Reinforcement learning

7957 **Reinforcement learning** is a topic worthy of a textbook of its own (Sutton and Barto,
 7958 1998),⁵ so this section will provide only a very brief overview, in the context of coreference
 7959 resolution. A stochastic **policy** assigns a probability to each possible **action**, conditional
 7960 on the context. The goal is to learn a policy that achieves a high expected reward, or
 7961 equivalently, a low expected cost.

7962 In incremental cluster ranking, a complete clustering on M mentions can be produced
 7963 by a sequence of M actions, in which the action z_i either merges mention i with an existing
 7964 cluster or begins a new cluster. We can therefore create a stochastic policy using the cluster
 7965 scores (Clark and Manning, 2016),

$$\Pr(z_i = e; \boldsymbol{\theta}) = \frac{\exp \psi_E(i \cup \{j : z_j = e\}; \boldsymbol{\theta})}{\sum_{e'} \exp \psi_E(i \cup \{j : z_j = e'\}; \boldsymbol{\theta})}, \quad [15.21]$$

7966 where $\psi_E(i \cup \{j : z_j = e\}; \boldsymbol{\theta})$ is the score under parameters $\boldsymbol{\theta}$ for assigning mention i to
 7967 cluster e . This score can be an arbitrary function of the mention i , the cluster e and its
 7968 (possibly empty) set of mentions; it can also include the history of actions taken thus far.

⁵A draft of the second edition can be found here: <http://incompleteideas.net/book/the-book-2nd.html>. Reinforcement learning has been used in spoken dialogue systems (Walker, 2000) and text-based game playing (Branavan et al., 2009), and was applied to coreference resolution by Clark and Manning (2015).

7969 If a policy assigns probability $p(c; \theta)$ to clustering c , then its expected loss is,

$$L(\theta) = \sum_{c \in \mathcal{C}(m)} p_\theta(c) \times \ell(c), \quad [15.22]$$

7970 where $\mathcal{C}(m)$ is the set of possible clusterings for mentions m . The loss $\ell(c)$ can be based on
 7971 any arbitrary scoring function, including the complex evaluation metrics used in corefer-
 7972 ence resolution (see § 15.4). This is an advantage of reinforcement learning, which can be
 7973 trained directly on the evaluation metric — unlike traditional supervised learning, which
 7974 requires a loss function that is differentiable and decomposable across individual deci-
 7975 sions.

Rather than summing over the exponentially many possible clusterings, we can ap-
 proximate the expectation by sampling trajectories of actions, $z = (z_1, z_2, \dots, z_M)$, from
 the current policy. Each action z_i corresponds to a step in the Bell tree: adding mention
 m_i to an existing cluster, or forming a new cluster. Each trajectory z corresponds to a
single clustering c , and so we can write the loss of an action sequence as $\ell(c(z))$. The
policy gradient algorithm computes the gradient of the expected loss as an expectation
over trajectories (Sutton et al., 2000),

$$\frac{\partial}{\partial \theta} L(\theta) = E_{z \sim \mathcal{Z}(m)} \ell(c(z)) \sum_{i=1}^M \frac{\partial}{\partial \theta} \log p(z_i | z_{1:i-1}, m) \quad [15.23]$$

$$\approx \frac{1}{K} \sum_{k=1}^K \ell(c(z^{(k)})) \sum_{i=1}^M \frac{\partial}{\partial \theta} \log p(z_i^{(k)} | z_{1:i-1}^{(k)}, m) \quad [15.24]$$

[15.25]

7976 where the action sequence $z^{(k)}$ is sampled from the current policy. Unlike the incremental
 7977 perceptron, an update is not made until the complete action sequence is available.

7978 15.2.4.4 Learning to search

7979 Policy gradient can suffer from high variance: while the average loss over K samples is
 7980 asymptotically equal to the expected reward of a given policy, this estimate may not be
 7981 accurate unless K is very large. This can make it difficult to allocate credit and blame to
 7982 individual actions. In **learning to search**, this problem is addressed through the addition
 7983 of an **oracle** policy, which is known to receive zero or small loss. The oracle policy can be
 7984 used in two ways:

- 7985 • The oracle can be used to generate partial hypotheses that are likely to score well,
 7986 by generating i actions from the initial state. These partial hypotheses are then used
 7987 as starting points for the learned policy. This is known as **roll-in**.

Algorithm 18 Learning to search for entity-based coreference resolution

```

1: procedure COMPUTE-GRADIENT(mentions  $m$ , loss function  $\ell$ , parameters  $\theta$ )
2:    $L(\theta) \leftarrow 0$ 
3:    $z \sim p(z | m; \theta)$                                  $\triangleright$  Sample a trajectory from the current policy
4:   for  $i \in \{1, 2, \dots, M\}$  do
5:     for action  $z \in \mathcal{Z}(z_{1:i-1}, m)$  do           $\triangleright$  All possible actions after history  $z_{1:i-1}$ 
6:        $h \leftarrow z_{1:i-1} \oplus z$                        $\triangleright$  Concatenate history  $z_{1:i-1}$  with action  $z$ 
7:       for  $j \in \{i+1, i+2, \dots, M\}$  do            $\triangleright$  Roll-out
8:          $h_j \leftarrow \operatorname{argmin}_h \ell(h_{1:j-1} \oplus h)$      $\triangleright$  Oracle selects action with minimum loss
9:        $L(\theta) \leftarrow L(\theta) + p(z | z_{1:i-1}, m; \theta) \times \ell(h)$        $\triangleright$  Update expected loss
10:      return  $\frac{\partial}{\partial \theta} L(\theta)$ 

```

- 7988 • The oracle can be used to compute the minimum possible loss from a given state, by
 7989 generating $M - i$ actions from the current state until completion. This is known as
 7990 **roll-out**.

7991 The oracle can be combined with the existing policy during both roll-in and roll-out, sam-
 7992 pling actions from each policy (Daumé III et al., 2009). One approach is to gradually
 7993 decrease the number of actions drawn from the oracle over the course of learning (Ross
 7994 et al., 2011).

7995 In the context of entity-based coreference resolution, Clark and Manning (2016) use
 7996 the learned policy for roll-in and the oracle policy for roll-out. Algorithm 18 shows how
 7997 the gradients on the policy weights are computed in this case. In this application, the
 7998 oracle is “noisy”, because it selects the action that minimizes only the *local* loss — the
 7999 accuracy of the coreference clustering up to mention i — rather than identifying the action
 8000 sequence that will lead to the best final coreference clustering on the entire document.
 8001 When learning from noisy oracles, it can be helpful to mix in actions from the current
 8002 policy with the oracle during roll-out (Chang et al., 2015).

8003 **15.3 Representations for coreference resolution**

8004 Historically, coreference resolution has employed an array of hand-engineered features
 8005 to capture the linguistic constraints and preferences described in § 15.1 (Soon et al., 2001).
 8006 Later work has documented the utility of lexical and bilexical features on mention pairs (Björkelund
 8007 and Nugues, 2011; Durrett and Klein, 2013). The most recent and successful methods re-
 8008 place many (but not all) of these features with distributed representations of mentions
 8009 and entities (Wiseman et al., 2015; Clark and Manning, 2016; Lee et al., 2017).

8010 **15.3.1 Features**

8011 Coreference features generally rely on a preprocessing pipeline to provide part-of-speech
 8012 tags and phrase structure parses. This pipeline makes it possible to design features that
 8013 capture many of the phenomena from § 15.1, and is also necessary for typical approaches
 8014 to mention identification. However, the pipeline may introduce errors that propagate
 8015 to the downstream coreference clustering system. Furthermore, the existence of such
 8016 a pipeline presupposes resources such as treebanks, which do not exist for many lan-
 8017 guages.⁶

8018 **15.3.1.1 Mention features**

8019 Features of individual mentions can help to predict anaphoricity. In systems where men-
 8020 tion detection is performed jointly with coreference resolution, these features can also
 8021 predict whether a span of text is likely to be a mention. For mention i , typical features
 8022 include:

8023 **Mention type.** Each span can be identified as a pronoun, name, or nominal, using the
 8024 part-of-speech of the head word of the mention: both the Penn Treebank and Uni-
 8025 versal Dependencies tagsets (§ 8.1.1) include tags for pronouns and proper nouns,
 8026 and all other heads can be marked as nominals (Haghghi and Klein, 2009).

8027 **Mention width.** The number of tokens in a mention is a rough predictor of its anaphor-
 8028 icity, with longer mentions being less likely to refer back to previously-defined enti-
 8029 ties.

8030 **Lexical features.** The first, last, and head words can help to predict anaphoricity; they are
 8031 also useful in conjunction with features such as mention type and part-of-speech,
 8032 providing a rough measure of agreement (Björkelund and Nugues, 2011). The num-
 8033 ber of lexical features can be very large, so it can be helpful to select only frequently-
 8034 occurring features (Durrett and Klein, 2013).

8035 **Morphosyntactic features.** These features include the part-of-speech, number, gender,
 8036 and dependency ancestors.

8037 The features for mention i and candidate antecedent a can be conjoined, producing
 8038 joint features that can help to assess the compatibility of the two mentions. For example,
 8039 Durrett and Klein (2013) conjoin each feature with the mention types of the anaphora
 8040 and the antecedent. Coreference resolution corpora such as ACE and OntoNotes contain

⁶The Universal Dependencies project has produced dependency treebanks for more than sixty languages. However, coreference features and mention detection are generally based on phrase structure trees, which exist for roughly two dozen languages. A list is available here: <https://en.wikipedia.org/wiki/Treebank>

8041 documents from various genres. By conjoining the genre with other features, it is possible
8042 to learn genre-specific feature weights.

8043 **15.3.1.2 Mention-pair features**

8044 For any pair of mentions i and j , typical features include:

8045 **Distance.** The number of intervening tokens, mentions, and sentences between i and j
8046 can all be used as distance features. These distances can be computed on the surface
8047 text, or on a transformed representation reflecting the breadth-first tree traversal
8048 (Figure 15.3). Rather than using the distances directly, they are typically binned,
8049 creating binary features.

8050 **String match.** A variety of string match features can be employed: exact match, suffix
8051 match, head match, and more complex matching rules that disregard irrelevant
8052 modifiers (Soon et al., 2001).

8053 **Compatibility.** Building on the model, features can measure the anaphor and antecedent
8054 agree with respect to morphosyntactic attributes such as gender, number, and ani-
8055 macy.

8056 **Nesting.** If one mention is nested inside another (e.g., *[The President of [France]]*), they
8057 generally cannot corefer.

8058 **Same speaker.** For documents with quotations, such as news articles, personal pronouns
8059 can be resolved only by determining the speaker for each mention (Lee et al., 2013).
8060 Coreference is also more likely between mentions from the same speaker.

8061 **Gazetteers.** These features indicate that the anaphor and candidate antecedent appear in
8062 a gazetteer of acronyms (e.g., *USA/United States*, *GATech/Georgia Tech*), demonymns
8063 (e.g., *Israel/Israeli*), or other aliases (e.g., *Knickerbockers/New York Knicks*).

8064 **Lexical semantics.** These features use a lexical resource such as WordNet to determine
8065 whether the head words of the mentions are related through synonymy, antonymy,
8066 and hypernymy (§ 4.2).

8067 **Dependency paths.** The dependency path between the anaphor and candidate antecedent
8068 can help to determine whether the pair can corefer, under the government and bind-
8069 ing constraints described in § 15.1.1.

8070 Comprehensive lists of mention-pair features are offered by Bengtson and Roth (2008) and
8071 Rahman and Ng (2011). Neural network approaches use far fewer mention-pair features:
8072 for example, Lee et al. (2017) include only speaker, genre, distance, and mention width
8073 features.

8074 **Semantics** In many cases, coreference seems to require knowledge and semantic in-
 8075 ferences, as in the running example, where we link *China* with a *country* and a *growth*
 8076 *market*. Some of this information can be gleaned from WordNet, which defines a graph
 8077 over **synsets** (see § 4.2). For example, one of the synsets of *China* is an instance of an
 8078 *Asian_nation#1*, which in turn is a hyponym of *country#2*, a synset that includes
 8079 *country*.⁷ Such paths can be used to measure the similarity between concepts (Pedersen
 8080 et al., 2004), and this similarity can be incorporated into coreference resolution as a fea-
 8081 ture (Ponzetto and Strube, 2006). Similar ideas can be applied to knowledge graphs in-
 8082 duced from Wikipedia (Ponzetto and Strube, 2007). But while such approaches improve
 8083 relatively simple classification-based systems, they have proven less useful when added
 8084 to the current generation of techniques.⁸ For example, Durrett and Klein (2013) employ
 8085 a range of semantics-based features — WordNet synonymy and hypernymy relations on
 8086 head words, named entity types (e.g., person, organization), and unsupervised clustering
 8087 over nominal heads — but find that these features give minimal improvement over a
 8088 baseline system using surface features.

8089 15.3.1.3 Entity features

8090 Many of the features for entity-mention coreference are generated by aggregating mention-
 8091 pair features over all mentions in the candidate entity (Culotta et al., 2007; Rahman and
 8092 Ng, 2011). Specifically, for each binary mention-pair feature $f(i, j)$, we compute the fol-
 8093 lowing entity-mention features for mention i and entity $e = \{j : j < i \wedge z_j = e\}$.

- 8094 • ALL-TRUE: Feature $f(i, j)$ holds for all mentions $j \in e$.
- 8095 • MOST-TRUE: Feature $f(i, j)$ holds for at least half and fewer than all mentions $j \in e$.
- 8096 • MOST-FALSE: Feature $f(i, j)$ holds for at least one and fewer than half of all men-
 8097 tions $j \in e$.
- 8098 • NONE: Feature $f(i, j)$ does not hold for any mention $j \in e$.

8099 For scalar mention-pair features (e.g., distance features), aggregation can be performed by
 8100 computing the minimum, maximum, and median values across all mentions in the cluster.
 8101 Additional entity-mention features include the number of mentions currently clustered in
 8102 the entity, and ALL-X and MOST-X features for each mention type.

8103 15.3.2 Distributed representations of mentions and entities

8104 Recent work has emphasized distributed representations of both mentions and entities.
 8105 One potential advantage is that pre-trained embeddings could help to capture the se-

⁷teletype font is used to indicate wordnet synsets, and *italics* is used to indicate strings.

⁸This point was made by Michael Strube at a 2015 workshop, noting that as the quality of the machine learning models in coreference has improved, the benefit of including semantics has become negligible.

8106 mantic compatibility underlying nominal coreference, helping with difficult cases like
 8107 (*Apple, the firm*) and (*China, the firm's biggest growth market*). Furthermore, a distributed
 8108 representation of entities can be trained to capture semantic features that are added by
 8109 each mention.

8110 **15.3.2.1 Mention embeddings**

8111 Entity mentions can be embedded into a vector space, providing the base layer for neural
 8112 networks that score coreference decisions (Wiseman et al., 2015).

8113 **Constructing the mention embedding** Various approaches for embedding multiword
 8114 units can be applied (see § 14.8). Figure 15.5 shows a recurrent neural network approach,
 8115 which begins by running a bidirectional LSTM over the entire text, obtaining hidden states
 8116 from the left-to-right and right-to-left passes, $\mathbf{h}_m = [\overleftarrow{\mathbf{h}}_m; \overrightarrow{\mathbf{h}}_m]$. Each candidate mention
 8117 span (s, t) is then represented by the vertical concatenation of four vectors:

$$\mathbf{u}^{(s,t)} = [\mathbf{u}_{\text{first}}^{(s,t)}; \mathbf{u}_{\text{last}}^{(s,t)}; \mathbf{u}_{\text{head}}^{(s,t)}; \phi^{(s,t)}], \quad [15.26]$$

8118 where $\mathbf{u}_{\text{first}}^{(s,t)} = \mathbf{h}_{s+1}$ is the embedding of the first word in the span, $\mathbf{u}_{\text{last}}^{(s,t)} = \mathbf{h}_t$ is the
 8119 embedding of the last word, $\mathbf{u}_{\text{head}}^{(s,t)}$ is the embedding of the “head” word, and $\phi^{(s,t)}$ is a
 8120 vector of surface features, such as the length of the span (Lee et al., 2017).

Attention over head words Rather than identifying the head word from the output of a parser, it can be computed from a neural **attention mechanism**:

$$\tilde{\alpha}_m = \theta_\alpha \cdot \mathbf{h}_m \quad [15.27]$$

$$\mathbf{a}^{(s,t)} = \text{SoftMax}([\tilde{\alpha}_{s+1}, \tilde{\alpha}_{s+2}, \dots, \tilde{\alpha}_t]) \quad [15.28]$$

$$\mathbf{u}_{\text{head}}^{(s,t)} = \sum_{m=s+1}^t a_m^{(s,t)} \mathbf{h}_m. \quad [15.29]$$

8121 Each token m gets a scalar score $\tilde{\alpha}_m = \theta_\alpha \cdot \mathbf{h}_m$, which is the dot product of the LSTM
 8122 hidden state \mathbf{h}_m and a vector of weights θ_α . The vector of scores for tokens in the span
 8123 $m \in \{s + 1, s + 2, \dots, t\}$ is then passed through a softmax layer, yielding a vector $\mathbf{a}^{(s,t)}$
 8124 that allocates one unit of attention across the span. This eliminates the need for syntactic
 8125 parsing to recover the head word; instead, the model learns to identify the most important
 8126 words in each span. Attention mechanisms were introduced in neural machine transla-
 8127 tion (Bahdanau et al., 2014), and are described in more detail in § 18.3.1.

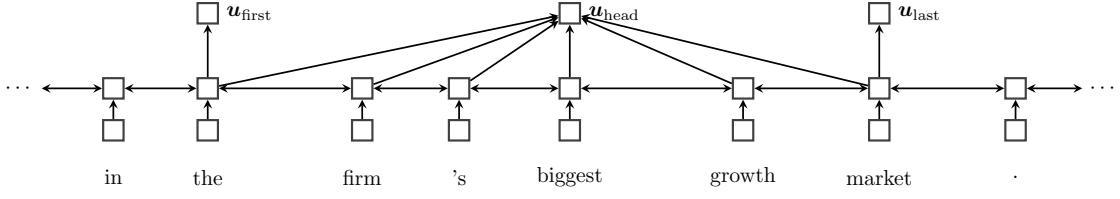


Figure 15.5: A bidirectional recurrent model of mention embeddings. The mention is represented by its first word, its last word, and an estimate of its head word, which is computed from a weighted average (Lee et al., 2017).

Using mention embeddings Given a set of mention embeddings, each mention i and candidate antecedent a is scored as,

$$\psi(a, i) = \psi_S(a) + \psi_S(i) + \psi_M(a, i) \quad [15.30]$$

$$\psi_S(a) = \text{FeedForward}_S(\mathbf{u}^{(a)}) \quad [15.31]$$

$$\psi_S(i) = \text{FeedForward}_S(\mathbf{u}^{(i)}) \quad [15.32]$$

$$\psi_M(a, i) = \text{FeedForward}_M([\mathbf{u}^{(a)}; \mathbf{u}^{(i)}; \mathbf{u}^{(a)} \odot \mathbf{u}^{(i)}; \mathbf{f}(a, i, \mathbf{w})]), \quad [15.33]$$

where $\mathbf{u}^{(a)}$ and $\mathbf{u}^{(i)}$ are the embeddings for spans a and i respectively, as defined in Equation 15.26.

- The scores $\psi_S(a)$ quantify whether span a is likely to be a coreferring mention, independent of what it corefers with. This allows the model to learn identify mentions directly, rather than identifying mentions with a preprocessing step.
- The score $\psi_M(a, i)$ computes the compatibility of spans a and i . Its base layer is a vector that includes the embeddings of spans a and i , their elementwise product $\mathbf{u}^{(a)} \odot \mathbf{u}^{(i)}$, and a vector of surface features $\mathbf{f}(a, i, \mathbf{w})$, including distance, speaker, and genre information.

Lee et al. (2017) provide an error analysis that shows how this method can correctly link a *blaze* and a *fire*, while incorrectly linking *pilots* and *fight attendants*. In each case, the coreference decision is based on similarities in the word embeddings.

Rather than embedding individual mentions, Clark and Manning (2016) embed mention pairs. At the base layer, their network takes embeddings of the words in and around each mention, as well as **one-hot** vectors representing a few surface features, such as the distance and string matching features. This base layer is then passed through a multilayer feedforward network with ReLU nonlinearities, resulting in a representation of the mention pair. The output of the mention pair encoder $\mathbf{u}_{i,j}$ is used in the scoring function of a mention-ranking model, $\psi_M(i, j) = \theta \cdot \mathbf{u}_{i,j}$. A similar approach is used to score cluster

8147 pairs, constructing a cluster-pair encoding by **pooling** over the mention-pair encodings
8148 for all pairs of mentions within the two clusters.

8149 **15.3.2.2 Entity embeddings**

8150 In entity-based coreference resolution, each entity should be represented by properties of
8151 its mentions. In a distributed setting, we maintain a set of vector entity embeddings, v_e .
8152 Each candidate mention receives an embedding u_i ; Wiseman et al. (2016) compute this
8153 embedding by a single-layer neural network, applied to a vector of surface features. The
8154 decision of whether to merge mention i with entity e can then be driven by a feedforward
8155 network, $\psi_E(i, e) = \text{Feedforward}([v_e; u_i])$. If i is added to entity e , then its representa-
8156 tion is updated recurrently, $v_e \leftarrow f(v_e, u_i)$, using a recurrent neural network such as a
8157 long short-term memory (LSTM; chapter 6). Alternatively, we can apply a **pooling** oper-
8158 ation, such as max-pooling or average-pooling (chapter 3), setting $v_e \leftarrow \text{Pool}(v_e, u_i)$. In
8159 either case, the update to the representation of entity e can be thought of as adding new
8160 information about the entity from mention i .

8161 **15.4 Evaluating coreference resolution**

8162 The state of coreference evaluation is aggravatingly complex. Early attempts at sim-
8163 ple evaluation metrics were found to under-penalize trivial baselines, such as placing
8164 each mention in its own cluster, or grouping all mentions into a single cluster. Follow-
8165 ing Denis and Baldridge (2009), the CoNLL 2011 shared task on coreference (Pradhan
8166 et al., 2011) formalized the practice of averaging across three different metrics: MUC (Vi-
8167 lain et al., 1995), B-CUBED (Bagga and Baldwin, 1998a), and CEAf (Luo, 2005). Refer-
8168 ence implementations of these metrics are available from Pradhan et al. (2014) at <https://github.com/conll/reference-coreference-scorers>.
8169

8170 **Additional resources**

8171 Ng (2010) surveys coreference resolution through 2010. Early work focused exclusively
8172 on pronoun resolution, with rule-based (Lappin and Leass, 1994) and probabilistic meth-
8173 ods (Ge et al., 1998). The full coreference resolution problem was popularized in a shared
8174 task associated with the sixth Message Understanding Conference, which included coref-
8175 erence annotations for training and test sets of thirty documents each (Grishman and
8176 Sundheim, 1996). An influential early paper was the decision tree approach of Soon et al.
8177 (2001), who introduced mention ranking. A comprehensive list of surface features for
8178 coreference resolution is offered by Bengtson and Roth (2008). Durrett and Klein (2013)
8179 improved on prior work by introducing a large lexicalized feature set; subsequent work
8180 has emphasized neural representations of entities and mentions (Wiseman et al., 2015).

8181 **Exercises**

- 8182 1. Select an article from today's news, and annotate coreference for the first twenty
8183 noun phrases that appear in the article (include nested noun phrases). That is,
8184 group the noun phrases into entities, where each entity corresponds to a set of noun
8185 phrases. Then specify the mention-pair training data that would result from the first
8186 five noun phrases.
 - 8187 2. Using your annotations from the preceding problem, compute the following statistics:
8188
 - 8189 • The number of times new entities are introduced by each of the three types of
8190 referring expressions: pronouns, proper nouns, and nominals. Include "single-
8191 ton" entities that are mentioned only once.
 - 8192 • For each type of referring expression, compute the fraction of mentions that are
8193 anaphoric.
 - 8194 3. Apply a simple heuristic to all pronouns in the article from the previous exercise.
8195 Specifically, link each pronoun to the closest preceding noun phrase that agrees in
8196 gender, number, animacy, and person. Compute the following evaluation:
8197
 - 8198 • True positive: a pronoun that is linked to a noun phrase with which it is coref-
erent, or is correctly labeled as the first mention of an entity.
 - 8199 • False positive: a pronoun that is linked to a noun phrase with which it is not
8200 coreferent. (This includes mistakenly linking singleton or non-referential pro-
8201 nouns.)
 - 8202 • False negative: a pronoun that is not linked to a noun phrase with which it is
8203 coreferent.
- 8204 Compute the *F-MEASURE* for your method, and for a trivial baseline in which ev-
8205 ery mention is its own entity. Are there any additional heuristics that would have
8206 improved the performance of this method?
- 8207 4. Durrett and Klein (2013) compute the probability of the gold coreference clustering
8208 by summing over all antecedent structures that are compatible with the clustering.
8209 Compute the number of antecedent structures for a single entity with K mentions.
 - 8210 5. Use the policy gradient algorithm to compute the gradient for the following sce-
8211 nario, based on the Bell tree in Figure 15.4:
 - 8212 • The gold clustering c^* is $\{Abigail, her\}, \{she\}$.

- Drawing a single sequence of actions ($K = 1$) from the current policy, you obtain the following incremental clusterings:

$$\begin{aligned}\mathbf{c}(a_1) &= \{\text{Abigail}\} \\ \mathbf{c}(\mathbf{a}_{1:2}) &= \{\text{Abigail}, \text{she}\} \\ \mathbf{c}(\mathbf{a}_{1:3}) &= \{\text{Abigail}, \text{she}\}, \{\text{her}\}.\end{aligned}$$

- 8213 • At each mention t , the action space \mathcal{A}_t is to merge the mention with each exist-
8214 ing cluster, or the empty cluster, with probability,

$$\Pr(\text{Merge}(m_t, \mathbf{c}(\mathbf{a}_{1:t-1}))) \propto \exp \psi_E(m_t \cup \mathbf{c}(\mathbf{a}_{1:t-1})), \quad [15.34]$$

8215 where the cluster score $\psi_E(m_t \cup c)$ is defined in Equation 15.15.

8216 Compute the gradient $\frac{\partial}{\partial \theta} L(\theta)$ in terms of the loss $\ell(c(a))$ and the features of each
8217 (potential) cluster. Explain the differences between the gradient-based update $\theta \leftarrow \theta - \frac{\partial}{\partial \theta} L(\theta)$
8218 and the incremental perceptron update from this sample example.

8219 **Chapter 16**

8220 **Discourse**

8221 Applications of natural language processing often concern multi-sentence documents:
8222 from paragraph-long restaurant reviews, to 500-word newspaper articles, to 500-page
8223 novels. Yet most of the methods that we have discussed thus far are concerned with
8224 individual sentences. This chapter discusses theories and methods for handling multi-
8225 sentence linguistic phenomena, known collectively as **discourse**. There are diverse char-
8226 acterizations of discourse structure, and no single structure is ideal for every computa-
8227 tional application. This chapter covers some of the most well studied discourse repre-
8228 sentations, while highlighting computational models for identifying and exploiting these
8229 structures.

8230 **16.1 Segments**

8231 A document or conversation can be viewed as a sequence of **segments**, each of which is
8232 **cohesive** in its content and/or function. In Wikipedia biographies, these segments often
8233 pertain to various aspects to the subject's life: early years, major events, impact on others,
8234 and so on. This segmentation is organized around **topics**. Alternatively, scientific research
8235 articles are often organized by **functional themes**: the introduction, a survey of previous
8236 research, experimental setup, and results.

8237 Written texts often mark segments with section headers and related formatting de-
8238 vices. However, such formatting may be too coarse-grained to support applications such
8239 as the retrieval of specific passages of text that are relevant to a query (Hearst, 1997).
8240 Unformatted speech transcripts, such as meetings and lectures, are also an application
8241 scenario for segmentation (Carletta, 2007; Glass et al., 2007; Janin et al., 2003).

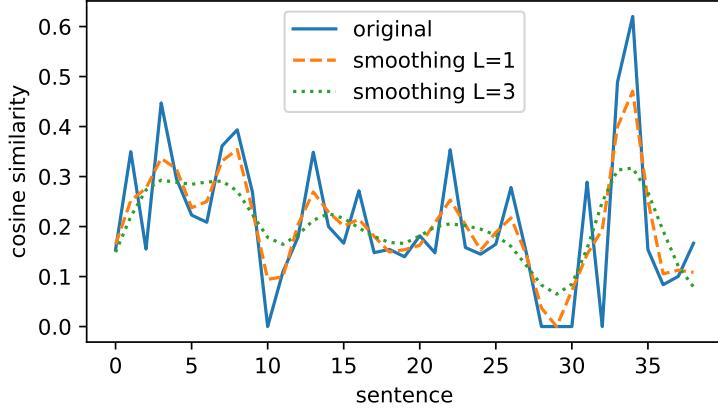


Figure 16.1: Smoothed cosine similarity among adjacent sentences in a news article. Local minima at $m = 10$ and $m = 29$ indicate likely segmentation points.

8242 16.1.1 Topic segmentation

A cohesive topic segment forms a unified whole, using various linguistic devices: repeated references to an entity or event; the use of conjunctions to link related ideas; and the repetition of meaning through lexical choices (Halliday and Hasan, 1976). Each of these cohesive devices can be measured, and then used as features for topic segmentation. A classical example is the use of lexical cohesion in the `TextTiling` method for topic segmentation (Hearst, 1997). The basic idea is to compute the textual similarity between each pair of adjacent blocks of text (sentences or fixed-length units), using a formula such as the smoothed **cosine similarity** of their bag-of-words vectors,

$$s_m = \frac{\mathbf{x}_m \cdot \mathbf{x}_{m+1}}{\|\mathbf{x}_m\|_2 \times \|\mathbf{x}_{m+1}\|_2} \quad [16.1]$$

$$\bar{s}_m = \sum_{\ell=0}^L k_\ell (s_{m+\ell} + s_{m-\ell}), \quad [16.2]$$

8243 with k_ℓ representing the value of a smoothing kernel of size L , e.g. $\mathbf{k} = [1, 0.5, 0.25]^\top$.
 8244 Segmentation points are then identified at local minima in the smoothed similarities \bar{s} ,
 8245 since these points indicate changes in the overall distribution of words in the text. An
 8246 example is shown in Figure 16.1.

8247 Text segmentation can also be formulated as a probabilistic model, in which each seg-
 8248 ment has a unique language model that defines the probability over the text in the seg-
 8249 ment (Utiyama and Isahara, 2001; Eisenstein and Barzilay, 2008; Du et al., 2013).¹ A good

¹There is a rich literature on how latent variable models (such as **latent Dirichlet allocation**) can track

8250 segmentation achieves high likelihood by grouping segments with similar word distribu-
8251 tions. This probabilistic approach can be extended to **hierarchical topic segmentation**, in
8252 which each topic segment is divided into subsegments (Eisenstein, 2009). All of these ap-
8253 proaches are unsupervised. While labeled data can be obtained from well-formatted texts
8254 such as textbooks, such annotations may not generalize to speech transcripts in alterna-
8255 tive domains. Supervised methods have been tried in cases where in-domain labeled data
8256 is available, substantially improving performance by learning weights on multiple types
8257 of features (Galley et al., 2003).

8258 16.1.2 Functional segmentation

8259 In some genres, there is a canonical set of communicative *functions*: for example, in sci-
8260 entific research articles, one such function is to communicate the general background for
8261 the article, another is to introduce a new contribution, or to describe the aim of the re-
8262 search (Teufel et al., 1999). A **functional segmentation** divides the document into con-
8263 tiguous segments, sometimes called **rhetorical zones**, in which each sentence has the same
8264 function. Teufel and Moens (2002) train a supervised classifier to identify the functional
8265 of each sentence in a set of scientific research articles, using features that describe the sen-
8266 tence's position in the text, its similarity to the rest of the article and title, tense and voice of
8267 the main verb, and the functional role of the previous sentence. Functional segmentation
8268 can also be performed without supervision. Noting that some types of Wikipedia arti-
8269 cles have very consistent functional segmentations (e.g., articles about cities or chemical
8270 elements), Chen et al. (2009) introduce an unsupervised model for functional segmenta-
8271 tion, which learns both the language model associated with each function and the typical
8272 patterning of functional segments across the article.

8273 16.2 Entities and reference

8274 Another dimension of discourse relates to which entities are mentioned throughout the
8275 text, and how. Consider the examples in Figure 16.2: Grosz et al. (1995) argue that the first
8276 discourse is more coherent. Do you agree? The examples differ in their choice of **refe-
8277 ring expressions** for the protagonist *John*, and in the syntactic constructions in sentences
8278 (b) and (d). The examples demonstrate the need for theoretical models to explain how
8279 referring expressions are chosen, and where they are placed within sentences. Such mod-
8280 els can then be used to help interpret the overall structure of the discourse, to measure
8281 discourse coherence, and to generate discourses in which referring expressions are used
8282 coherently.

topics across documents (Blei et al., 2003; Blei, 2012).

- | | |
|--|---|
| (16.1) a. John went to his favorite music store to buy a piano.
b. He had frequented the store for many years.
c. He was excited that he could finally buy a piano.
d. He arrived just as the store was closing for the day | (16.2) a. John went to his favorite music store to buy a piano.
b. It was a store John had frequented for many years.
c. He was excited that he could finally buy a piano.
d. It was closing just as John arrived. |
|--|---|

Figure 16.2: Two tellings of the same story (Grosz et al., 1995). The discourse on the left uses referring expressions coherently, while the one on the right does not.

8283 16.2.1 Centering theory

8284 The relationship between discourse and entity reference is most elaborated in **centering**
8285 **theory** (Grosz et al., 1995). According to the theory, every utterance in the discourse is
8286 characterized by a set of entities, known as *centers*.

- 8287 • The **forward-looking centers** in utterance m are all the entities that are mentioned
8288 in the utterance, $c_f(w_m) = \{e_1, e_2, \dots\}$. The forward-looking centers are partially
8289 ordered by their syntactic prominence, favoring subjects over other positions.
- 8290 • The **backward-looking center** $c_b(w_m)$ is the highest-ranked element in the set of
8291 forward-looking centers from the previous utterance $c_f(w_{m-1})$ that is also men-
8292 tioned in w_m .

8293 Given these two definitions, centering theory makes the following predictions about
8294 the form and position of referring expressions:

- 8295 1. If a pronoun appears in the utterance w_m , then the backward-looking center $c_b(w_m)$
8296 must also be realized as a pronoun. This rule argues against the use of *it* to refer
8297 to the piano store in Example (16.2d), since JOHN is the backward looking center of
8298 (16.2d), and he is mentioned by name and not by a pronoun.
- 8299 2. Sequences of utterances should retain the same backward-looking center if possible,
8300 and ideally, the backward-looking center should also be the top-ranked element in
8301 the list of forward-looking centers. This rule argues in favor of the preservation of
8302 JOHN as the backward-looking center throughout Example (16.1).

8303 Centering theory unifies aspects of syntax, discourse, and anaphora resolution. However,
8304 it can be difficult to clarify exactly how to rank the elements of each utterance, or even
8305 how to partition a text or dialog into utterances (Poesio et al., 2004).

	SKYLER	WALTER	DANGER	A GUY	THE DOOR
<i>You don't know who you're talking to,</i>	S	-	-	-	-
<i>so let me clue you in.</i>	O	O	-	-	-
<i>I am not in danger, Skyler.</i>	X	S	X	-	-
<i>I am the danger.</i>	-	S	O	-	-
<i>A guy opens his door and gets shot,</i>	-	-	-	S	O
<i>and you think that of me?</i>	S	X	-	-	-
<i>No. I am the one who knocks!</i>	-	S	-	-	-

Figure 16.3: The entity grid representation for a dialogue from the television show *Breaking Bad*.

8306 16.2.2 The entity grid

8307 One way to formalize the ideas of centering theory is to arrange the entities in a text or
 8308 conversation in an **entity grid**. This is a data structure with one row per sentence, and
 8309 one column per entity (Barzilay and Lapata, 2008). Each cell $c(m, i)$ can take the following
 8310 values:

$$c(m, i) = \begin{cases} S, & \text{entity } i \text{ is in subject position in sentence } m \\ O, & \text{entity } i \text{ is in object position in sentence } m \\ X, & \text{entity } i \text{ appears in sentence } m, \text{ in neither subject nor object position} \\ -, & \text{entity } i \text{ does not appear in sentence } m. \end{cases} \quad [16.3]$$

8311 To populate the entity grid, syntactic parsing is applied to identify subject and object
 8312 positions, and coreference resolution is applied to link multiple mentions of a single entity.
 8313 An example is shown in Figure 16.3.

8314 After the grid is constructed, the coherence of a document can be measured by the
 8315 transitions between adjacent cells in each column. For example, the transition $(S \rightarrow S)$
 8316 keeps an entity in subject position across adjacent sentences; the transition $(O \rightarrow S)$ pro-
 8317 motes an entity from object position to subject position; the transition $(S \rightarrow -)$ drops the
 8318 subject of one sentence from the next sentence. The probabilities of each transition can be
 8319 estimated from labeled data, and an entity grid can then be scored by the sum of the log-
 8320 probabilities across all columns and all transitions, $\sum_{i=1}^{N_e} \sum_{m=1}^M \log p(c(m, i) | c(m-1, i))$.
 8321 The resulting probability can be used as a proxy for the coherence of a text. This has been
 8322 shown to be useful for a range of tasks: determining which of a pair of articles is more
 8323 readable (Schwartz and Ostendorf, 2005), correctly ordering the sentences in a scrambled

8324 text (Lapata, 2003), and disentangling multiple conversational threads in an online multi-
 8325 party chat (Elsner and Charniak, 2010).

8326 **16.2.3 *Formal semantics beyond the sentence level**

8327 An alternative view of the role of entities in discourse focuses on formal semantics, and the
 8328 construction of meaning representations for multi-sentence units. Consider the following
 8329 two sentences (from Bird et al., 2009):

- 8330 (16.3) a. Angus owns a dog.
 8331 b. It bit Irene.

8332 We would like to recover the formal semantic representation,

$$\exists x. \text{DOG}(x) \wedge \text{OWN}(\text{ANGUS}, x) \wedge \text{BITE}(x, \text{IRENE}). \quad [16.4]$$

However, the semantic representations of each individual sentence are:

$$\exists x. \text{DOG}(x) \wedge \text{OWN}(\text{ANGUS}, x) \quad [16.5]$$

$$\text{BITE}(y, \text{IRENE}). \quad [16.6]$$

8333 Unifying these two representations into the form of Equation 16.4 requires linking the
 8334 unbound variable y from [16.6] with the quantified variable x in [16.5]. Discourse under-
 8335 standing therefore requires the reader to update a set of assignments, from variables
 8336 to entities. This update would (presumably) link the *dog* in the first sentence of [16.3]
 8337 with the unbound variable y in the second sentence, thereby licensing the conjunction in
 8338 [16.4].² This basic idea is at the root of **dynamic semantics** (Groenendijk and Stokhof,
 8339 1991). **Segmented discourse representation theory** links dynamic semantics with a set
 8340 of **discourse relations**, which explain how adjacent units of text are rhetorically or con-
 8341 ceptually related (Lascarides and Asher, 2007). The next section explores the theory of
 8342 discourse relations in more detail.

8343 **16.3 Relations**

8344 In dependency grammar, sentences are characterized by a graph (usually a tree) of syntac-
 8345 tic relations between words, such as NSUBJ and DET. A similar idea can be applied at the
 8346 document level, identifying relations between discourse units, such as clauses, sentences,
 8347 or paragraphs. The task of **discourse parsing** involves identifying discourse units and
 8348 the relations that hold between them. These relations can then be applied to tasks such as
 8349 document classification and summarization, as discussed in § 16.3.4.

²This linking task is similar to coreference resolution (see chapter 15), but here the connections are between semantic variables, rather than spans of text.

- TEMPORAL
 - Asynchronous
 - Synchronous: precedence, succession
- CONTINGENCY
 - Cause: result, reason
 - Pragmatic cause: justification
 - Condition: hypothetical, general, unreal present, unreal past, real present, real past
 - Pragmatic condition: relevance, implicit assertion
- COMPARISON
 - Contrast: juxtaposition, opposition
 - Pragmatic contrast
 - Concession: expectation, contra-expectation
 - Pragmatic concession
- EXPANSION
 - Conjunction
 - Instantiation
 - Restatement: specification, equivalence, generalization
 - Alternative: conjunctive, disjunctive, chosen alternative
 - Exception
 - List

Table 16.1: The hierarchy of discourse relation in the Penn Discourse Treebank annotations (Prasad et al., 2008). For example, PRECEDENCE is a subtype of SYNCHRONOUS, which is a type of TEMPORAL relation.

8350 16.3.1 Shallow discourse relations

8351 The existence of discourse relations is hinted by **discourse connectives**, such as *however*,
 8352 *moreover*, *meanwhile*, and *if ... then*. These connectives explicitly specify the relationship
 8353 between adjacent units of text: *however* signals a contrastive relationship, *moreover* signals
 8354 that the subsequent text elaborates or strengthens the point that was made immediately
 8355 beforehand, *meanwhile* indicates that two events are contemporaneous, and *if ... then* sets
 8356 up a conditional relationship. Discourse connectives can therefore be viewed as a starting
 8357 point for the analysis of discourse relations.

8358 In **lexicalized tree-adjoining grammar for discourse (D-LTAG)**, each connective an-
 8359 chors a relationship between two units of text (Webber, 2004). This model provides the
 8360 theoretical basis for the **Penn Discourse Treebank (PDTB)**, the largest corpus of discourse
 8361 relations in English (Prasad et al., 2008). It includes a hierarchical inventory of discourse
 8362 relations (shown in Table 16.1), which is created by abstracting the meanings implied by
 8363 the discourse connectives that appear in real texts (Knott, 1996). These relations are then
 8364 annotated on the same corpus of news text used in the Penn Treebank (see § 9.2.2), adding
 8365 the following information:

- Each connective is annotated for the discourse relation or relations that it expresses, if any — many discourse connectives have senses in which they do not signal a discourse relation (Pitler and Nenkova, 2009).
- For each discourse relation, the two arguments of the relation are specified as ARG1 and ARG2, where ARG2 is constrained to be adjacent to the connective. These arguments may be sentences, but they may also smaller or larger units of text.
- Adjacent sentences are annotated for **implicit discourse relations**, which are not marked by any connective. When a connective could be inserted between a pair of sentence, the annotator supplies it, and also labels its sense (e.g., example 16.5). In some cases, there is no relationship at all between a pair of adjacent sentences; in other cases, the only relation is that the adjacent sentences mention one or more shared entity. These phenomena are annotated as NOREL and ENTRREL (entity relation), respectively.

Examples of Penn Discourse Treebank annotations are shown in (16.4). In (16.4), the word *therefore* acts as an explicit discourse connective, linking the two adjacent units of text. The Treebank annotations also specify the “sense” of each relation, linking the connective to a relation in the sense inventory shown in Table 16.1: in (16.4), the relation is PRAGMATIC CAUSE:JUSTIFICATION because it relates to the author’s communicative intentions. The word *therefore* can also signal causes in the external world (e.g., *He was therefore forced to relinquish his plan*). In **discourse sense classification**, the goal is to determine which discourse relation, if any, is expressed by each connective. A related task is the classification of implicit discourse relations, as in (16.5). In this example, the relationship between the adjacent sentences could be expressed by the connective *because*, indicating a CAUSE:REASON relationship.

16.3.1.1 Classifying explicit discourse relations and their arguments

As suggested by the examples above, many connectives can be used to invoke multiple types of discourse relations. Similarly, some connectives have senses that are unrelated to discourse: for example, *and* functions as a discourse connective when it links propositions, but not when it links noun phrases (Lin et al., 2014). Nonetheless, the senses of explicitly-marked discourse relations in the Penn Treebank are relatively easy to classify, at least at the coarse-grained level. When classifying the four top-level PDTB relations, 90% accuracy can be obtained simply by selecting the most common relation for each connective (Pitler and Nenkova, 2009). At the more fine-grained levels of the discourse relation hierarchy, connectives are more ambiguous. This fact is reflected both in the accuracy of automatic sense classification (Versley, 2011) and in interannotator agreement, which falls to 80% for level-3 discourse relations (Prasad et al., 2008).

- (16.4) *...as this business of whaling has somehow come to be regarded among landsmen as a rather unpoetical and disreputable pursuit; therefore, I am all anxiety to convince ye, ye landsmen, of the injustice hereby done to us hunters of whales.*
- (16.5) But a few funds have taken other defensive steps. *Some have raised their cash positions to record levels. Implicit = BECAUSE High cash positions help buffer a fund when the market falls.*
- (16.6) Michelle lives in a hotel room, and although **she drives a canary-colored Porsche**, *she hasn't time to clean or repair it.*
- (16.7) *Most oil companies, when they set exploration and production budgets for this year, forecast revenue of \$15 for each barrel of crude produced.*

Figure 16.4: Example annotations of discourse relations. In the style of the Penn Discourse Treebank, the discourse connective is underlined, the first argument is shown in italics, and the second argument is shown in bold. Examples (16.5-16.7) are quoted from Prasad et al. (2008).

8402 A more challenging task for explicitly-marked discourse relations is to identify the
 8403 scope of the arguments. Discourse connectives need not be adjacent to ARG1, as shown
 8404 in item 16.6, where ARG1 follows ARG2; furthermore, the arguments need not be contigu-
 8405 ous, as shown in (16.7). For these reasons, recovering the arguments of each discourse
 8406 connective is a challenging subtask. Because intra-sentential arguments are often syn-
 8407 tactic constituents (see chapter 10), many approaches train a classifier to predict whether
 8408 each constituent is an appropriate argument for each explicit discourse connective (Well-
 8409 ner and Pustejovsky, 2007; Lin et al., 2014, e.g.,).

8410 16.3.1.2 Classifying implicit discourse relations

Implicit discourse relations are considerably more difficult to classify and to annotate.³ Most approaches are based on an encoding of each argument, which is then used as input to a non-linear classifier:

$$\mathbf{z}^{(i)} = \text{Encode}(\mathbf{w}^{(i)}) \quad [16.7]$$

$$\mathbf{z}^{(i+1)} = \text{Encode}(\mathbf{w}^{(i+1)}) \quad [16.8]$$

$$\hat{y}_i = \underset{y}{\operatorname{argmax}} \Psi(y, \mathbf{z}^{(i)}, \mathbf{z}^{(i+1)}). \quad [16.9]$$

³In the dataset for the 2015 shared task on shallow discourse parsing, the interannotator agreement was 91% for explicit discourse relations and 81% for non-explicit relations, across all levels of detail (Xue et al., 2015).

8411 This basic framework can be instantiated in several ways, including both feature-based
 8412 and neural encoders. Several recent approaches are compared in the 2015 and 2016 shared
 8413 tasks at the Conference on Natural Language Learning (Xue et al., 2015, 2016).

8414 **Feature-based approaches** Each argument can be encoded into a vector of surface fea-
 8415 tures. The encoding typically includes lexical features (all words, or all content words, or
 8416 a subset of words such as the first three and the main verb), Brown clusters of individ-
 8417 ual words (§ 14.4), and syntactic features such as terminal productions and dependency
 8418 arcs (Pitler et al., 2009; Lin et al., 2009; Rutherford and Xue, 2014). The classification func-
 8419 tion then has two parts. First, it creates a joint feature vector by combining the encodings
 8420 of each argument, typically by computing the cross-product of all features in each encod-
 8421 ing:

$$\mathbf{f}(y, \mathbf{z}^{(i)}, \mathbf{z}^{(i+1)}) = \{(a \times b \times y) : (\mathbf{z}_a^{(i)} \mathbf{z}_b^{(i+1)})\} \quad [16.10]$$

8422 The size of this feature set grows with the square of the size of the vocabulary, so it can be
 8423 helpful to select a subset of features that are especially useful on the training data (Park
 8424 and Cardie, 2012). After \mathbf{f} is computed, any classifier can be trained to compute the final
 8425 score, $\Psi(y, \mathbf{z}^{(i)}, \mathbf{z}^{(i+1)}) = \boldsymbol{\theta} \cdot \mathbf{f}(y, \mathbf{z}^{(i)}, \mathbf{z}^{(i+1)})$.

8426 **Neural network approaches** In neural network architectures, the encoder is learned
 8427 jointly with the classifier as an end-to-end model. Each argument can be encoded using
 8428 a variety of neural architectures (surveyed in § 14.8): recursive (§ 10.6.1; Ji and Eisenstein,
 8429 2015), recurrent (§ 6.3; Ji et al., 2016), and convolutional (§ 3.4; Qin et al., 2017). The clas-
 8430 sification function can then be implemented as a feedforward neural network on the two
 8431 encodings (chapter 3; for examples, see Rutherford et al., 2017; Qin et al., 2017), or as a
 8432 simple bilinear product, $\Psi(y, \mathbf{z}^{(i)}, \mathbf{z}^{(i+1)}) = (\mathbf{z}^{(i)})^\top \boldsymbol{\Theta}_y \mathbf{z}^{(i+1)}$ (Ji and Eisenstein, 2015). The
 8433 encoding model can be trained by backpropagation from the classification objective, such
 8434 as the margin loss. Rutherford et al. (2017) show that neural architectures outperform
 8435 feature-based approaches in most settings. While neural approaches require engineering
 8436 the network architecture (e.g., embedding size, number of hidden units in the classifier),
 8437 feature-based approaches also require significant engineering to incorporate linguistic re-
 8438 sources such as Brown clusters and parse trees, and to select a subset of relevant features.

8439 16.3.2 Hierarchical discourse relations

8440 In sentence parsing, adjacent phrases combine into larger constituents, ultimately pro-
 8441 ducing a single constituent for the entire sentence. The resulting tree structure enables
 8442 structured analysis of the sentence, with subtrees that represent syntactically coherent
 8443 chunks of meaning. **Rhetorical Structure Theory (RST)** extends this style of hierarchical
 8444 analysis to the discourse level (Mann and Thompson, 1988).

8445 The basic element of RST is the **discourse unit**, which refers to a contiguous span of
 8446 text. **Elementary discourse units** (EDUs) are the atomic elements in this framework, and
 8447 are typically (but not always) clauses.⁴ Each discourse relation combines two or more
 8448 adjacent discourse units into a larger, composite discourse unit; this process ultimately
 8449 unites the entire text into a tree-like structure.⁵

8450 **Nuclearity** In many discourse relations, one argument is primary. For example:

8451 (16.8) [LaShawn loves animals]_N
 8452 [She has nine dogs and one pig]_S

8453 In this example, the second sentence provides EVIDENCE for the point made in the first
 8454 sentence. The first sentence is thus the **nucleus** of the discourse relation, and the second
 8455 sentence is the **satellite**. The notion of **nuclearity** is analogous to the head-modifier struc-
 8456 ture of dependency parsing (see § 11.1.1). However, in RST, some relations have multiple
 8457 nuclei. For example, the arguments of the CONTRAST relation are equally important:

8458 (16.9) [The clash of ideologies survives this treatment]_N
 8459 [but the nuance and richness of Gorky's individual characters have vanished in the scuffle]_N⁶

8460 Relations that have multiple nuclei are called **coordinating**; relations with a single nu-
 8461 cleus are called **subordinating**. Subordinating relations are constrained to have only two
 8462 arguments, while coordinating relations (such as CONJUNCTION) may have more than
 8463 two.

8464 **RST Relations** Rhetorical structure theory features a large inventory of discourse rela-
 8465 tions, which are divided into two high-level groups: subject matter relations, and presen-
 8466 tational relations. Presentational relations are organized around the intended beliefs of
 8467 the reader. For example, in (16.8), the second discourse unit provides evidence intended
 8468 to increase the reader's belief in the proposition expressed by the first discourse unit, that
 8469 *LaShawn loves animals*. In contrast, subject-matter relations are meant to communicate ad-
 8470 dditional facts about the propositions contained in the discourse units that they relate:

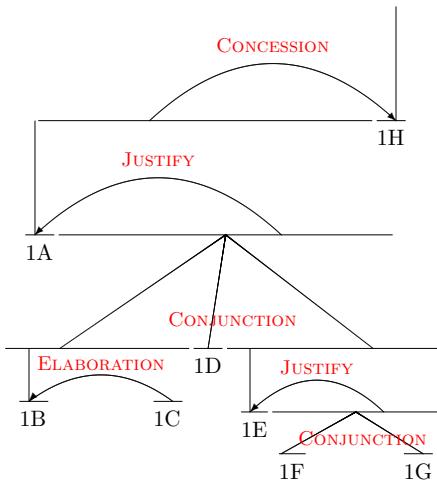
⁴Details of discourse segmentation can be found in the RST annotation manual (Carlson and Marcu, 2001).

⁵While RST analyses are typically trees, this should be taken as a strong theoretical commitment to the principle that all coherent discourses have a tree structure. Taboada and Mann (2006) write:

It is simply the case that trees are convenient, easy to represent, and easy to understand. There is, on the other hand, no theoretical reason to assume that trees are the only possible representation of discourse structure and of coherence relations.

The appropriateness of tree structures to discourse has been challenged, e.g., by Wolf and Gibson (2005), who propose a more general graph-structured representation.

⁶from the RST Treebank (Carlson et al., 2002)



[It could have been a great movie]^{1A} [It does have beautiful scenery,]^{1B} [some of the best since Lord of the Rings.]^{1C} [The acting is well done,]^{1D} [and I really liked the son of the leader of the Samurai.]^{1E} [He was a likable chap,]^{1F} [and I hated to see him die.]^{1G} [But, other than all that, this movie is nothing more than hidden rip-offs.]^{1H}

Figure 16.5: A rhetorical structure theory analysis of a short movie review, adapted from Voll and Taboada (2007). Positive and negative sentiment words are underlined, indicating RST’s potential utility in document-level sentiment analysis.

8471 (16.10) [the debt plan was rushed to completion]_N
 8472 [in order to be announced at the meeting]_S⁷

8473 In this example, the satellite describes a world state that is realized by the action described
 8474 in the nucleus. This relationship is about the world, and not about the author’s commu-
 8475 nicative intentions.

8476 **Example** Figure 16.5 depicts an RST analysis of a paragraph from a movie review. Asym-
 8477 metric (subordinating) relations are depicted with an arrow from the satellite to the nu-
 8478 cleus; symmetric (coordinating) relations are depicted with lines. The elementary dis-
 8479 course units 1F and 1G are combined into a larger discourse unit with the symmetric
 8480 CONJUNCTION relation. The resulting discourse unit is then the satellite in a JUSTIFY
 8481 relation with 1E.

⁷from the RST Treebank (Carlson et al., 2002)

8482 **16.3.2.1 Hierarchical discourse parsing**

8483 The goal of discourse parsing is to recover a hierarchical structural analysis from a doc-
 8484 ument text, such as the analysis in Figure 16.5. For now, let's assume a segmentation
 8485 of the document into elementary discourse units (EDUs); segmentation algorithms are
 8486 discussed below. After segmentation, discourse parsing can be viewed as a combination
 8487 of two components: the discourse relation classification techniques discussed in § 16.3.1.2,
 8488 and algorithms for phrase-structure parsing, such as chart parsing and shift-reduce, which
 8489 were discussed in chapter 10.

8490 Both chart parsing and shift-reduce require encoding composite discourse units, ei-
 8491 ther in a discrete feature vector or a dense neural representation.⁸ Some discourse parsers
 8492 rely on the **strong compositionality criterion** (Marcu, 1996), which states the assumption
 8493 that a composite discourse unit can be represented by its nucleus. This criterion is used in
 8494 feature-based discourse parsing to determine the feature vector for a composite discourse
 8495 unit (Hernault et al., 2010); it is used in neural approaches to setting the vector encod-
 8496 ing for a composite discourse unit equal to the encoding of its nucleus (Ji and Eisenstein,
 8497 2014). An alternative neural approach is to learn a composition function over the compo-
 8498 nents of a composite discourse unit (Li et al., 2014), using a recursive neural network (see
 8499 § 14.8.3).

8500 **Bottom-up discourse parsing** Assume a segmentation of the text into N elementary
 8501 discourse units with base representations $\{z^{(i)}\}_{i=1}^N$, and assume a composition function
 8502 COMPOSE $(z^{(i)}, z^{(j)}, \ell)$, which maps two encodings and a discourse relation ℓ into a new
 8503 encoding. The composition function can follow the strong compositionality criterion and
 8504 simply select the encoding of the nucleus, or it can do something more complex. We
 8505 also need a scoring function $\Psi(z^{(i,k)}, z^{(k,j)}, \ell)$, which computes a scalar score for the (bi-
 8506 narized) discourse relation ℓ with left child covering the span $i + 1 : k$, and the right
 8507 child covering the span $k + 1 : j$. Given these components, we can construct vector rep-
 8508 resentations for each span, and this is the basic idea underlying **compositional vector**
 8509 **grammars** (Socher et al., 2013).

8510 These same components can also be used in bottom-up parsing, in a manner that is
 8511 similar to the CKY algorithm for weighted context-free grammars (see § 10.1): compute
 8512 the score and best analysis for each possible span of increasing lengths, while storing
 8513 back-pointers that make it possible to recover the optimal parse of the entire input. How-
 8514 ever, there is an important distinction from CKY parsing: for each labeled span (i, j, ℓ) , we
 8515 must use the composition function to construct a representation $z^{(i,j,\ell)}$. This representa-
 8516 tion is then used to combine the discourse unit spanning $i + 1 : j$ in higher-level discourse
 8517 relations. The representation $z^{(i,j,\ell)}$ depends on the entire substructure of the unit span-

⁸To use these algorithms, is also necessary to binarize all discourse relations during parsing, and then to “unbinarize” them to reconstruct the desired structure (e.g., Hernault et al., 2010).

8518 ning $i + 1 : j$, and this violates the locality assumption that underlie CKY’s optimality
 8519 guarantee. Bottom-up parsing with recursively constructed span representations is gen-
 8520 erally not guaranteed to find the best-scoring discourse parse. This problem is explored
 8521 in an exercise at the end of the chapter.

8522 **Transition-based discourse parsing** One drawback of bottom-up parsing is its cubic
 8523 time complexity in the length of the input. For long documents, transition-based parsing
 8524 is an appealing alternative. The shift-reduce algorithm can be applied to discourse parsing
 8525 fairly directly (Sagae, 2009): the stack stores a set of discourse units and their repres-
 8526 entations, and each action is chosen by a function of these representations. This function
 8527 could be a linear product of weights and features, or it could be a neural network ap-
 8528 plied to encodings of the discourse units. The REDUCE action then performs composition
 8529 on the two discourse units at the top of the stack, yielding a larger composite discourse
 8530 unit, which goes on top of the stack. All of the techniques for integrating learning and
 8531 transition-based parsing, described in § 11.3, are applicable to discourse parsing.

8532 16.3.2.2 Segmenting discourse units

8533 In rhetorical structure theory, elementary discourse units do not cross the sentence bound-
 8534 ary, so discourse segmentation can be performed within sentences, assuming the sentence
 8535 segmentation is given. The segmentation of sentences into elementary discourse units is
 8536 typically performed using features of the syntactic analysis (Braud et al., 2017). One ap-
 8537 proach is to train a classifier to determine whether each syntactic constituent is an EDU,
 8538 using features such as the production, tree structure, and head words (Soricut and Marcu,
 8539 2003; Hernault et al., 2010). Another approach is to train a sequence labeling model, such
 8540 as a conditional random field (Sporleder and Lapata, 2005; Xuan Bach et al., 2012; Feng
 8541 et al., 2014). This is done using the BIO formalism for segmentation by sequence labeling,
 8542 described in § 8.3.

8543 16.3.3 Argumentation

8544 An alternative view of text-level relational structure focuses on **argumentation** (Stab and
 8545 Gurevych, 2014b). Each segment (typically a sentence or clause) may support or rebut
 8546 another segment, creating a graph structure over the text. In the following example (from
 8547 Peldszus and Stede, 2013), segment S_2 provides argumentative support for the proposi-
 8548 tion in the segment S_1 :

8549 (16.11) [We should tear the building down] $_{S1}$
 8550 [because it is full of asbestos] $_{S2}$.

8551 Assertions may also support or rebut proposed links between two other assertions, cre-
 8552 ating a **hypergraph**, which is a generalization of a graph to the case in which edges can

8553 join any number of vertices. This can be seen by introducing another sentence into the
8554 example:

8555 (16.12) [In principle it is possible to clean it up.]_{S3}
8556 [but according to the mayor that is too expensive.]_{S4}

8557 S3 acknowledges the validity of *S2*, but undercuts its support of *S1*. This can be repre-
8558 sented by introducing a hyperedge, $(S3, S2, S1)_{\text{undercut}}$, indicating that *S3* undercuts the
8559 proposed relationship between *S2* and *S1*. *S4* then undercuts the relevance of *S3*.

8560 **Argumentation mining** is the task of recovering such structures from raw texts. At
8561 present, annotations of argumentation structure are relatively small: Stab and Gurevych
8562 (2014a) have annotated a collection of 90 persuasive essays, and Peldszus and Stede (2015)
8563 have solicited and annotated a set of 112 paragraph-length “microtexts” in German.

8564 16.3.4 Applications of discourse relations

8565 The predominant application of discourse parsing is to select content within a document.
8566 In rhetorical structure theory, the nucleus is considered the more important element of
8567 the relation, and is more likely to be part of a summary of the document; it may also
8568 be more informative for document classification. The D-LTAG theory that underlies the
8569 Penn Discourse Treebank lacks this notion of nuclearity, but arguments may have varying
8570 importance, depending on the relation type. For example, the span of text constituting
8571 ARG1 of an expansion relation is more likely to appear in a summary, while the sentence
8572 constituting ARG2 of an implicit relation is less likely (Louis et al., 2010). Discourse rela-
8573 tions may also signal segmentation points in the document structure. Explicit discourse
8574 markers have been shown to correlate with changes in subjectivity, and identifying such
8575 change points can improve document-level sentiment classification, by helping the clas-
8576 sifier to focus on the subjective parts of the text (Trivedi and Eisenstein, 2013; Yang and
8577 Cardie, 2014).

8578 16.3.4.1 Extractive Summarization

8579 Text **summarization** is the problem of converting a longer text into a shorter one, while
8580 still conveying the key facts, events, ideas, and sentiments from the original. In **extractive**
8581 **summarization**, the summary is a subset of the original text; in **abstractive summariza-**
8582 **tion**, the summary is produced *de novo*, by paraphrasing the original, or by first encoding
8583 it into a semantic representation (see § 19.2). The main strategy for extractive summa-
8584 rization is to maximize **coverage**, choosing a subset of the document that best covers the
8585 concepts mentioned in the document as a whole; typically, coverage is approximated by
8586 bag-of-words overlap (Nenkova and McKeown, 2012). Coverage-based objectives can be
8587 supplemented by hierarchical discourse relations, using the principle of nuclearity: in any
8588 subordinating discourse relation, the nucleus is more critical to the overall meaning of the

8589 text, and is therefore more important to include in an extractive summary (Marcu, 1997a).⁹
 8590 This insight can be generalized from individual relations using the concept of **discourse**
 8591 **depth** (Hirao et al., 2013): for each elementary discourse unit e , the discourse depth d_e is
 8592 the number of relations in which a discourse unit containing e is the satellite.

8593 Both discourse depth and nuclearity can be incorporated into extractive summarization
 8594 using constrained optimization. Let \mathbf{x}_n be a bag-of-words vector representation of
 8595 elementary discourse unit n , let $y_n \in \{0, 1\}$ indicate whether n is included in the summary,
 8596 and let d_n be the depth of unit n . Furthermore, let each discourse unit have a “head” h ,
 8597 which is defined recursively:

- 8598 • if a discourse unit is produced by a subordinating relation, then its head is the head
 8599 of the (unique) nucleus;
- 8600 • if a discourse unit is produced by a coordinating relation, then its head is the head
 8601 of the left-most nucleus;
- 8602 • for each elementary discourse unit, its parent $\pi(n) \in \{\emptyset, 1, 2, \dots, N\}$ is the head of
 8603 the smallest discourse unit containing n whose head is not n ;
- 8604 • if n is the head of the discourse unit spanning the whole document, then $\pi(n) = \emptyset$.

With these definitions in place, discourse-driven extractive summarization can be formalized as (Hirao et al., 2013),

$$\begin{aligned} & \max_{\mathbf{y}=\{0,1\}^N} \sum_{n=1}^N y_n \frac{\Psi(\mathbf{x}_n, \{\mathbf{x}_{1:N}\})}{d_n} \\ & \text{s.t. } \sum_{n=1}^N y_n \left(\sum_{j=1}^V x_{n,j} \right) \leq L \\ & \quad y_{\pi(n)} \geq y_n, \quad \forall n \end{aligned} \tag{16.11}$$

8605 where $\Psi(\mathbf{x}_n, \{\mathbf{x}_{1:N}\})$ measures the coverage of elementary discourse unit n with respect
 8606 to the rest of the document, and $\sum_{j=1}^V x_{n,m}$ is the number of tokens in \mathbf{x}_n . The first con-
 8607 straint ensures that the number of tokens in the summary has an upper bound L . The
 8608 second constraint ensures that no elementary discourse unit is included unless its parent
 8609 is also included. In this way, the discourse structure is used twice: to downweight the
 8610 contributions of elementary discourse units that are not central to the discourse, and to
 8611 ensure that the resulting structure is a subtree of the original discourse parse. The opti-

⁹Conversely, the arguments of a multi-nuclear relation should either both be included in the summary, or both excluded (Durrett et al., 2016).

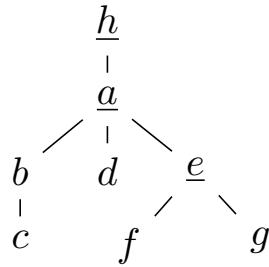


Figure 16.6: A **discourse depth tree** (Hirao et al., 2013) for the discourse parse from Figure 16.5, in which each elementary discourse unit is connected to its parent. The discourse units in one valid summary are underlined.

8612 mization problem in 16.11 can be solved with **integer linear programming**, described in
 8613 § 13.2.2.¹⁰

8614 Figure 16.6 shows a **discourse depth tree** for the RST analysis from Figure 16.5, in
 8615 which each elementary discourse is connected to (and below) its parent. The figure also
 8616 shows a valid summary, corresponding to:

8617 (16.13) It could have been a great movie, and I really liked the son of the leader of the
 8618 Samurai. But, other than all that, this movie is nothing more than hidden rip-offs.

8619 16.3.4.2 Document classification

8620 Hierarchical discourse structures lend themselves naturally to text classification: in a sub-
 8621 ordinating discourse relation, the nucleus should play a stronger role in the classification
 8622 decision than the satellite. Various implementations of this idea have been proposed.

- 8623 • Focusing on within-sentence discourse relations and lexicon-based classification (see
 8624 § 4.1.2), Voll and Taboada (2007) simply ignore the text in the satellites of each dis-
 8625 course relation.
- 8626 • At the document level, elements of each discourse relation argument can be reweighted,
 8627 favoring words in the nucleus, and disfavoring words in the satellite (Heerschop
 8628 et al., 2011; Bhatia et al., 2015). This approach can be applied recursively, computing
 8629 weights across the entire document. The weights can be relation-specific, so that the
 8630 features from the satellites of contrastive relations are discounted or even reversed.
- 8631 • Alternatively, the hierarchical discourse structure can define the structure of a **re-
 8632 cursive neural network** (see § 10.6.1). In this network, the representation of each

¹⁰Formally, 16.11 is a special case of the **knapsack problem**, in which the goal is to find a subset of items with maximum value, constrained by some maximum weight (Cormen et al., 2009).

8633 discourse unit is computed from its arguments and from a parameter corresponding
 8634 to the discourse relation (Ji and Smith, 2017).

8635 Shallow, non-hierarchical discourse relations have also been applied to document clas-
 8636 sification. One approach is to impose a set of constraints on the analyses of individual
 8637 discourse units, so that adjacent units have the same polarity when they are connected
 8638 by a discourse relation indicating agreement, and opposite polarity when connected by a
 8639 contrastive discourse relation, indicating disagreement (Somasundaran et al., 2009; Zirn
 8640 et al., 2011). Yang and Cardie (2014) apply explicitly-marked relations from the Penn
 8641 Discourse Treebank to the problem of sentence-level sentiment polarity classification (see
 8642 § 4.1). They impose the following soft constraints:

- 8643 • When a CONTRAST relation appears between two sentences, those sentences should
 8644 have opposite sentiment polarity.
- 8645 • When an EXPANSION or CONTINGENCY relation appears between two sentences,
 8646 they should have the same polarity.
- 8647 • When a CONTRAST relation appears *within* a sentence, it should have neutral polar-
 8648 ity, since it is likely to express both sentiments.

8649 These discourse-driven constraints are shown to improve performance on two datasets of
 8650 product reviews.

8651 16.3.4.3 Coherence

8652 Just as **grammaticality** is the property shared by well-structured sentences, **coherence** is
 8653 the property shared by well-structured discourses. One application of discourse process-
 8654 ing is to measure (and maximize) the coherence of computer-generated texts like transla-
 8655 tions and summaries (Kibble and Power, 2004). Coherence assessment is also used to eval-
 8656 uate human-generated texts, such as student essays (e.g., Miltsakaki and Kukich, 2004;
 8657 Burstein et al., 2013).

8658 Coherence subsumes a range of phenomena, many of which have been highlighted
 8659 earlier in this chapter: e.g., that adjacent sentences should be lexically cohesive (Foltz
 8660 et al., 1998; Ji et al., 2015; Li and Jurafsky, 2017), and that entity references should follow
 8661 the principles of centering theory (Barzilay and Lapata, 2008; Nguyen and Joty, 2017).
 8662 Discourse relations also bear on the coherence of a text in a variety of ways:

- 8663 • Hierarchical discourse relations tend to have a “canonical ordering” of the nucleus
 8664 and satellite (Mann and Thompson, 1988): for example, in the ELABORATION rela-
 8665 tion from rhetorical structure theory, the nucleus always comes first, while in the
 8666 JUSTIFICATION relation, the satellite tends to be first (Marcu, 1997b).

- Discourse relations should be signaled by connectives that are appropriate to the semantic or functional relationship between the arguments: for example, a coherent text would be more likely to use *however* to signal a COMPARISON relation than a *temporal* relation (Kibble and Power, 2004).
- Discourse relations tend to appear in predictable sequences: for example, COMPARISON relations tend to immediately precede CONTINGENCY relations (Pitler et al., 2008). This observation can be formalized by generalizing the entity grid model (§ 16.2.2), so that each cell (i, j) provides information about the role of the discourse argument containing a mention of entity j in sentence i (Lin et al., 2011). For example, if the first sentence is ARG1 of a comparison relation, then any entity mentions in the sentence would be labeled COMP.ARG1. This approach can also be applied to RST discourse relations (Feng et al., 2014).

Datasets One difficulty with evaluating metrics of discourse coherence is that human-generated texts usually meet some minimal threshold of coherence. For this reason, much of the research on measuring coherence has focused on synthetic data. A typical setting is to permute the sentences of a human-written text, and then determine whether the original sentence ordering scores higher according to the proposed coherence measure (Barzilay and Lapata, 2008). There are also small datasets of human evaluations of the coherence of machine summaries: for example, human judgments of the summaries from the participating systems in the 2003 Document Understanding Conference are available online.¹¹ Researchers from the Educational Testing Service (an organization which administers several national exams in the United States) have studied the relationship between discourse coherence and student essay quality (Burstein et al., 2003, 2010). A public dataset of essays from second-language learners, with quality annotations, has been made available by researchers at Cambridge University (Yannakoudakis et al., 2011). At the other extreme, Louis and Nenkova (2013) analyze the structure of professionally written scientific essays, finding that discourse relation transitions help to distinguish prize-winning essays from other articles in the same genre.

Additional resources

For a manuscript-length discussion of discourse processing, see Stede (2011). Article-length surveys are offered by Webber et al. (2012) and Webber and Joshi (2012).

¹¹<http://homepages.inf.ed.ac.uk/mlap/coherence/>

8698 **Exercises**

- 8699 1. • Implement the smoothed cosine similarity metric from Equation 16.2, using the
 8700 smoothing kernel $k = [.5, .3, .15, .05]$.
 8701 • Download the text of a news article with at least ten paragraphs.
 8702 • Compute and plot the smoothed similarity \bar{s} over the length of the article.
 8703 • Identify *local minima* in \bar{s} as follows: first find all sentences m such that $\bar{s}_m <$
 8704 $\bar{s}_{m \pm 1}$. Then search among these points to find the five sentences with the lowest
 8705 \bar{s}_m .
 8706 • How often do the five local minima correspond to paragraph boundaries?
 8707 – The fraction of local minima that are paragraph boundaries is the **precision-**
 8708 **at- k** , where in this case, $k = 5$.
 8709 – The fraction of paragraph boundaries which are local minima is the **recall-**
 8710 **at- k** .
 8711 – Compute precision-at- k and recall-at- k for $k = 3$ and $k = 10$.
- 8712 2. This exercise is to be done in pairs. Each participant selects an article from to-
 8713 day's news, and replaces all mentions of individual people with special tokens like
 8714 PERSON1, PERSON2, and so on. The other participant should then use the rules
 8715 of centering theory to guess each type of referring expression: full name (*Captain*
 8716 *Ahab*), partial name (e.g., *Ahab*), nominal (e.g., *the ship's captain*), or pronoun. Check
 8717 whether the predictions match the original article, and whether the original article
 8718 conforms to the rules of centering theory.
- 8719 3. In § 16.3.2.1, it is noted that bottom-up parsing with compositional representations
 8720 of each span is not guaranteed to be optimal. In this exercise, you will construct
 8721 a minimal example proving this point. Consider a discourse with four units, with
 8722 base representations $\{z^{(i)}\}_{i=1}^4$. Construct a scenario in which the parse selected by
 8723 bottom-up parsing is not optimal, and give the precise mathematical conditions that
 8724 must hold for this suboptimal parse to be selected. You may ignore the relation
 8725 labels ℓ for the purpose of this example.

8726

Part IV

8727

Applications

8728 Chapter 17

8729 Information extraction

8730 Computers offer powerful capabilities for searching and reasoning about structured records
8731 and relational data. Some even argue that the most important limitation of artificial intel-
8732 ligence is not inference or learning, but simply having too little knowledge (Lenat et al.,
8733 1990). Natural language processing provides an appealing solution: automatically con-
8734 struct a structured **knowledge base** by reading natural language text.

8735 For example, many Wikipedia pages have an “infobox” that provides structured in-
8736 formation about an entity or event. An example is shown in Figure 17.1a: each row rep-
8737 resents one or more properties of the entity IN THE AEROPLANE OVER THE SEA, a record
8738 album. The set of properties is determined by a predefined **schema**, which applies to all
8739 record albums in Wikipedia. As shown in Figure 17.1b, the values for many of these fields
8740 are indicated directly in the first few sentences of text on the same Wikipedia page.

8741 The task of automatically constructing (or “populating”) an infobox from text is an
8742 example of **information extraction**. Much of information extraction can be described in
8743 terms of **entities**, **relations**, and **events**.

8744 • **Entities** are uniquely specified objects in the world, such as people (JEFF MANGUM),
8745 places (ATHENS, GEORGIA), organizations (MERGE RECORDS), and times (FEBRUARY
8746 10, 1998). Chapter 8 described the task of **named entity recognition**, which labels
8747 tokens as parts of entity spans. Now we will see how to go further, **linking** each
8748 entity **mention** to an element in a **knowledge base**.

- 8749 • **Relations** include a **predicate** and two **arguments**: for example, CAPITAL(GEORGIA, ATLANTA).
- **Events** involve multiple typed arguments. For example, the production and release

Studio album by Neutral Milk Hotel	
Released	February 10, 1998
Recorded	July–September 1997
Studio	Pet Sounds Studio, Denver, Colorado
Genre	Indie rock • psychedelic folk • lo-fi
Length	39:55
Label	Merge • Domino
Producer	Robert Schneider

(a) A Wikipedia infobox

- (17.1) In the Aeroplane Over the Sea is the second and final studio album by the American indie rock band Neutral Milk Hotel.
- (17.2) It was released in the United States on February 10, 1998 on Merge Records and May 1998 on Blue Rose Records in the United Kingdom.
- (17.3) Jeff Mangum moved from Athens, Georgia to Denver, Colorado to prepare the bulk of the album's material with producer Robert Schneider, this time at Schneider's newly created Pet Sounds Studio at the home of Jim McIntyre.

- (b) The first few sentences of text. Strings that match fields or field names in the infobox are underlined; strings that mention other entities are wavy underlined.

Figure 17.1: From the Wikipedia page for the album “In the Aeroplane Over the Sea”, retrieved October 26, 2017.

of the album described in Figure 17.1 is described by the event,

```
<TITLE : IN THE AEROPLANE OVER THE SEA,
ARTIST : NEUTRAL MILK HOTEL,
RELEASE-DATE : 1998-FEB-10,...>
```

8750 The set of arguments for an event type is defined by a **schema**. Events often refer to
 8751 time-delimited occurrences: weddings, protests, purchases, terrorist attacks.

8752 Information extraction is similar to semantic role labeling (chapter 13): we may think
 8753 of predicates as corresponding to events, and the arguments as defining slots in the event
 8754 representation. However, the goals of information extraction are different. Rather than
 8755 accurately parsing every sentence, information extraction systems often focus on recog-
 8756 nizing a few key relation or event types, or on the task of identifying all properties of a
 8757 given entity. Information extraction is often evaluated by the correctness of the resulting
 8758 knowledge base, and not by how many sentences were accurately parsed. The goal is
 8759 sometimes described as **macro-reading**, as opposed to **micro-reading**, in which each sen-
 8760 tence must be analyzed correctly. Macro-reading systems are not penalized for ignoring
 8761 difficult sentences, as long as they can recover the same information from other, easier-
 8762 to-read sources. However, macro-reading systems must resolve apparent inconsistencies

8763 (was the album released on MERGE RECORDS or BLUE ROSE RECORDS?), requiring reasoning across the entire dataset.

8765 In addition to the basic tasks of recognizing entities, relations, and events, information
8766 extraction systems must handle negation, and must be able to distinguish statements of
8767 fact from hopes, fears, hunches, and hypotheticals. Finally, information extraction is often paired with the problem of **question answering**, which requires accurately parsing a
8769 query, and then selecting or generating a textual answer. Question answering systems can
8770 be built on knowledge bases that are extracted from large text corpora, or may attempt to
8771 identify answers directly from the source texts.

8772 17.1 Entities

8773 The starting point for information extraction is to identify mentions of entities in text.
8774 Consider the following example:

8775 (17.4) *The United States Army captured a hill overlooking Atlanta on May 14, 1864.*

8776 For this sentence, there are two goals:

- 8777 1. *Identify* the spans *United States Army*, *Atlanta*, and *May 14, 1864* as entity mentions.
8778 (The hill is not uniquely identified, so it is not a *named* entity.) We may also want to
8779 recognize the **named entity types**: organization, location, and date. This is **named**
8780 **entity recognition**, and is described in chapter 8.
- 8781 2. *Link* these spans to entities in a knowledge base: U.S. ARMY, ATLANTA, and 1864-
8782 MAY-14. This task is known as **entity linking**.

8783 The strings to be linked to entities are **mentions** — similar to the use of this term in
8784 coreference resolution. In some formulations of the entity linking task, only named entities
8785 are candidates for linking. This is sometimes called **named entity linking** (Ling et al.,
8786 2015). In other formulations, such as **Wikification** (Milne and Witten, 2008), any string
8787 can be a mention. The set of target entities often corresponds to Wikipedia pages, and
8788 Wikipedia is the basis for more comprehensive knowledge bases such as YAGO (Suchanek
8789 et al., 2007), DBPedia (Auer et al., 2007), and Freebase (Bollacker et al., 2008). Entity link-
8790 ing may also be performed in more “closed” settings, where a much smaller list of targets
8791 is provided in advance. The system must also determine if a mention does not refer to
8792 any entity in the knowledge base, sometimes called a **NIL entity** (McNamee and Dang,
8793 2009).

8794 Returning to (17.4), the three entity mentions may seem unambiguous. But the Wikipedia
8795 disambiguation page for the string *Atlanta* says otherwise:¹ there are more than twenty

¹[https://en.wikipedia.org/wiki/Atlanta_\(disambiguation\)](https://en.wikipedia.org/wiki/Atlanta_(disambiguation)), retrieved November 1, 2017.

8796 different towns and cities, five United States Navy vessels, a magazine, a television show,
 8797 a band, and a singer — each prominent enough to have its own Wikipedia page. We now
 8798 consider how to choose among these dozens of possibilities. In this chapter we will focus
 8799 on supervised approaches. Unsupervised entity linking is closely related to the problem
 8800 of **cross-document coreference resolution**, where the task is to identify pairs of mentions
 8801 that corefer, across document boundaries (Bagga and Baldwin, 1998b; Singh et al., 2011).

8802 17.1.1 Entity linking by learning to rank

8803 Entity linking is often formulated as a **ranking** problem,

$$\hat{y} = \operatorname{argmax}_{y \in \mathcal{Y}(x)} \Psi(y, x, c), \quad [17.1]$$

8804 where y is a target entity, x is a description of the mention, $\mathcal{Y}(x)$ is a set of candidate
 8805 entities, and c is a description of the context — such as the other text in the document,
 8806 or its metadata. The function Ψ is a scoring function, which could be a linear model,
 8807 $\Psi(y, x, c) = \theta \cdot f(y, x, c)$, or a more complex function such as a neural network. In either
 8808 case, the scoring function can be learned by minimizing a margin-based **ranking loss**,

$$\ell(\hat{y}, y^{(i)}, x^{(i)}, c^{(i)}) = (\Psi(\hat{y}, x^{(i)}, c^{(i)}) - \Psi(y^{(i)}, x^{(i)}, c^{(i)}) + 1)_+, \quad [17.2]$$

8809 where $y^{(i)}$ is the ground truth and $\hat{y} \neq y^{(i)}$ is the predicted target for mention $x^{(i)}$ in
 8810 context $c^{(i)}$ (Joachims, 2002; Dredze et al., 2010).

8811 **Candidate identification** For computational tractability, it is helpful to restrict the set of
 8812 candidates, $\mathcal{Y}(x)$. One approach is to use a **name dictionary**, which maps from strings
 8813 to the entities that they might mention. This mapping is many-to-many: a string such as
 8814 *Atlanta* can refer to multiple entities, and conversely, an entity such as ATLANTA can be
 8815 referenced by multiple strings. A name dictionary can be extracted from Wikipedia, with
 8816 links between each Wikipedia entity page and the anchor text of all hyperlinks that point
 8817 to the page (Bunescu and Pasca, 2006; Ratinov et al., 2011). To improve recall, the name
 8818 dictionary can be augmented by partial and approximate matching (Dredze et al., 2010),
 8819 but as the set of candidates grows, the risk of false positives increases. For example, the
 8820 string *Atlanta* is a partial match to *the Atlanta Fed* (a name for the FEDERAL RESERVE BANK
 8821 OF ATLANTA), and a noisy match (edit distance of one) from *Atalanta* (a heroine in Greek
 8822 mythology and an Italian soccer team).

8823 **Features** Feature-based approaches to entity ranking rely on three main types of local
 8824 information (Dredze et al., 2010):

- The similarity of the mention string to the canonical entity name, as quantified by string similarity. This feature would elevate the city ATLANTA over the basketball team ATLANTA HAWKS for the string *Atlanta*.
- The popularity of the entity, which can be measured by Wikipedia page views or PageRank in the Wikipedia link graph. This feature would elevate ATLANTA, GEORGIA over the unincorporated community of ATLANTA, OHIO.
- The entity type, as output by the named entity recognition system. This feature would elevate the city of ATLANTA over the magazine ATLANTA in contexts where the mention is tagged as a location.

In addition to these local features, the document context can also help. If *Jamaica* is mentioned in a document about the Caribbean, it is likely to refer to the island nation; in the context of New York, it is likely to refer to the neighborhood in Queens; in the context of a menu, it might refer to a hibiscus tea beverage. Such hints can be formalized by computing the similarity between the Wikipedia page describing each candidate entity and the mention context $c^{(i)}$, which may include the bag-of-words representing the document (Dredze et al., 2010; Hoffart et al., 2011) or a smaller window of text around the mention (Ratinov et al., 2011). For example, we can compute the cosine similarity between bag-of-words vectors for the context and entity description, typically weighted using **inverse document frequency** to emphasize rare words.²

Neural entity linking An alternative approach is to compute the score for each entity candidate using distributed vector representations of the entities, mentions, and context. For example, for the task of entity linking in Twitter, Yang et al. (2016) employ the bilinear scoring function,

$$\Psi(y, x, c) = v_y^\top \Theta^{(y,x)} x + v_y^\top \Theta^{(y,c)} c, \quad [17.3]$$

with $v_y \in \mathbb{R}^{K_y}$ as the vector embedding of entity y , $x \in \mathbb{R}^{K_x}$ as the embedding of the mention, $c \in \mathbb{R}^{K_c}$ as the embedding of the context, and the matrices $\Theta^{(y,x)}$ and $\Theta^{(y,c)}$ as parameters that score the compatibility of each entity with respect to the mention and context. Each of the vector embeddings can be learned from an end-to-end objective, or pre-trained on unlabeled data.

- Pretrained **entity embeddings** can be obtained from an existing knowledge base (Bordes et al., 2011, 2013), or by running a word embedding algorithm such as WORD2VEC

²The **document frequency** of word j is $DF(j) = \frac{1}{N} \sum_{i=1}^N \delta(x_j^{(i)} > 0)$, equal to the number of documents in which the word appears. The contribution of each word to the cosine similarity of two bag-of-words vectors can be weighted by the **inverse document frequency** $\frac{1}{DF(j)}$ or $\log \frac{1}{DF(j)}$, to emphasize rare words (Spärck Jones, 1972).

- 8855 on the text of Wikipedia, with hyperlinks substituted for the anchor text.³
- 8856 • The embedding of the mention x can be computed by averaging the embeddings
 8857 of the words in the mention (Yang et al., 2016), or by the compositional techniques
 8858 described in § 14.8.
- 8859 • The embedding of the context c can also be computed from the embeddings of the
 8860 words in the context. A **denoising autoencoder** learns a function from raw text to
 8861 dense K -dimensional vector encodings by minimizing a reconstruction loss (Vin-
 8862 cent et al., 2010),

$$\min_{\theta_g, \theta_h} \sum_{i=1}^N \|\mathbf{x}^{(i)} - g(h(\tilde{\mathbf{x}}^{(i)}; \theta_h); \theta_g)\|^2, \quad [17.4]$$

8859 where $\tilde{\mathbf{x}}^{(i)}$ is a noisy version of the bag-of-words counts $\mathbf{x}^{(i)}$, which is produced by
 8860 randomly setting some counts to zero; $h : \mathbb{R}^V \mapsto \mathbb{R}^K$ is an encoder with parameters
 8861 θ_h ; and $g : \mathbb{R}^K \mapsto \mathbb{R}^V$, with parameters θ_g . The encoder and decoder functions
 8862 are typically implemented as feedforward neural networks. To apply this model to
 8863 entity linking, each entity and context are initially represented by the encoding of
 8864 their bag-of-words vectors, $h(e)$ and $g(c)$, and these encodings are then fine-tuned
 8865 from labeled data (He et al., 2013). The context vector c can also be obtained by
 8866 convolution on the embeddings of words in the document (Sun et al., 2015), or by
 8867 examining metadata such as the author’s social network (Yang et al., 2016).

8868 The remaining parameters $\Theta^{(y,x)}$ and $\Theta^{(y,c)}$ can be trained by backpropagation from the
 8869 margin loss in Equation 17.2.

8870 17.1.2 Collective entity linking

8871 Entity linking can be more accurate when it is performed jointly across a document. To
 8872 see why, consider the following lists:

- 8873 (17.5) California, Oregon, Washington
- 8874 (17.6) Baltimore, Washington, Philadelphia
- 8875 (17.7) Washington, Adams, Jefferson

8876 In each case, the term *Washington* refers to a different entity, and this reference is strongly
 8877 suggested by the other entries on the list. In the last list, all three names are highly am-
 8878 biguous — there are dozens of other *Adams* and *Jefferson* entities in Wikipedia. But a

³Pre-trained entity embeddings can be downloaded from <https://code.google.com/archive/p/word2vec/>.

8879 preference for coherence motivates **collectively** linking these references to the first three
 8880 U.S. presidents.

8881 A general approach to collective entity linking is to introduce a compatibility score
 8882 $\psi_c(\mathbf{y})$. Collective entity linking is then performed by optimizing the global objective,

$$\hat{\mathbf{y}} = \underset{\mathbf{y} \in \mathbb{Y}(\mathbf{x})}{\operatorname{argmax}} \Psi_c(\mathbf{y}) + \sum_{i=1}^N \Psi_\ell(y^{(i)}, \mathbf{x}^{(i)}, \mathbf{c}^{(i)}), \quad [17.5]$$

8883 where $\mathbb{Y}(\mathbf{x})$ is the set of all possible collective entity assignments for the mentions in \mathbf{x} ,
 8884 and ψ_ℓ is the local scoring function for each entity i . The compatibility function is typically
 8885 decomposed into a sum of pairwise scores, $\Psi_c(\mathbf{y}) = \sum_{i=1}^N \sum_{j \neq i}^N \Psi_c(y^{(i)}, y^{(j)})$. These scores
 8886 can be computed in a number of different ways:

- 8887 • Wikipedia defines high-level categories for entities (e.g., *living people*, *Presidents of*
 8888 *the United States*, *States of the United States*), and Ψ_c can reward entity pairs for the
 8889 number of categories that they have in common (Cucerzan, 2007).
- 8890 • Compatibility can be measured by the number of incoming hyperlinks shared by
 8891 the Wikipedia pages for the two entities (Milne and Witten, 2008).
- 8892 • In a neural architecture, the compatibility of two entities can be set equal to the inner
 8893 product of their embeddings, $\Psi_c(y^{(i)}, y^{(j)}) = \mathbf{v}_{y^{(i)}} \cdot \mathbf{v}_{y^{(j)}}$.
- 8894 • A non-pairwise compatibility score can be defined using a type of latent variable
 8895 model known as a **probabilistic topic model** (Blei et al., 2003; Blei, 2012). In this
 8896 framework, each latent topic is a probability distribution over entities, and each
 8897 document has a probability distribution over topics. Each entity helps to determine
 8898 the document's distribution over topics, and in turn these topics help to resolve am-
 8899 biguous entity mentions (Newman et al., 2006). Inference can be performed using
 8900 the sampling techniques described in chapter 5.

8901 Unfortunately, collective entity linking is **NP-hard** even for pairwise compatibility func-
 8902 tions, so exact optimization is almost certainly intractable. Various approximate inference
 8903 techniques have been proposed, including **integer linear programming** (Cheng and Roth,
 8904 2013), **Gibbs sampling** (Han and Sun, 2012), and graph-based algorithms (Hoffart et al.,
 8905 2011; Han et al., 2011).

8906 17.1.3 *Pairwise ranking loss functions

8907 The loss function defined in Equation 17.2 considers only the highest-scoring prediction
 8908 \hat{y} , but in fact, the true entity $y^{(i)}$ should outscore *all* other entities. A loss function based on
 8909 this idea would give a gradient against the features or representations of several entities,

Algorithm 19 WARP approximate ranking loss

```

1: procedure WARP( $y^{(i)}$ ,  $\mathbf{x}^{(i)}$ )
2:    $N \leftarrow 0$ 
3:   repeat
4:     Randomly sample  $y \sim \mathcal{Y}(\mathbf{x}^{(i)})$ 
5:      $N \leftarrow N + 1$ 
6:     if  $\psi(y, \mathbf{x}^{(i)}) + 1 > \psi(y^{(i)}, \mathbf{x}^{(i)})$  then            $\triangleright$  check for margin violation
7:        $r \leftarrow \lfloor |\mathcal{Y}(\mathbf{x}^{(i)})|/N \rfloor$                           $\triangleright$  compute approximate rank
8:       return  $L_{\text{rank}}(r) \times (\psi(y, \mathbf{x}^{(i)}) + 1 - \psi(y^{(i)}, \mathbf{x}^{(i)}))$ 
9:     until  $N \geq |\mathcal{Y}(\mathbf{x}^{(i)})| - 1$                             $\triangleright$  no violation found
10:    return 0                                          $\triangleright$  return zero loss

```

8910 not just the top-scoring prediction. Usunier et al. (2009) define a general ranking error
 8911 function,

$$L_{\text{rank}}(k) = \sum_{j=1}^k \alpha_j, \quad \text{with } \alpha_1 \geq \alpha_2 \geq \dots \geq 0, \quad [17.6]$$

8912 where k is equal to the number of labels ranked higher than the correct label $y^{(i)}$. This
 8913 function defines a class of ranking errors: if $\alpha_j = 1$ for all j , then the ranking error is
 8914 equal to the rank of the correct entity; if $\alpha_1 = 1$ and $\alpha_{j>1} = 0$, then the ranking error is
 8915 one whenever the correct entity is not ranked first; if α_j decreases smoothly with j , as in
 8916 $\alpha_j = \frac{1}{j}$, then the error is between these two extremes.

This ranking error can be integrated into a margin objective. Remember that large margin classification requires not only the correct label, but also that the correct label outscores other labels by a substantial margin. A similar principle applies to ranking: we want a high rank for the correct entity, and we want it to be separated from other entities by a substantial margin. We therefore define the margin-augmented rank,

$$r(y^{(i)}, \mathbf{x}^{(i)}) \triangleq \sum_{y \in \mathcal{Y}(\mathbf{x}^{(i)}) \setminus y^{(i)}} \delta \left(1 + \psi(y, \mathbf{x}^{(i)}) \geq \psi(y^{(i)}, \mathbf{x}^{(i)}) \right), \quad [17.7]$$

8917 where $\delta(\cdot)$ is a delta function, and $\mathcal{Y}(\mathbf{x}^{(i)}) \setminus y^{(i)}$ is the set of all entity candidates minus
 8918 the true entity $y^{(i)}$. The margin-augmented rank is the rank of the true entity, after aug-
 8919 menting every other candidate with a margin of one, under the current scoring function
 8920 ψ . (The context c is omitted for clarity, and can be considered part of x .)

For each instance, a hinge loss is computed from the ranking error associated with this

margin-augmented rank, and the violation of the margin constraint,

$$\ell(y^{(i)}, \mathbf{x}^{(i)}) = \frac{L_{\text{rank}}(r(y^{(i)}, \mathbf{x}^{(i)}))}{r(y^{(i)}, \mathbf{x}^{(i)})} \sum_{y \in \mathcal{Y}(\mathbf{x}) \setminus y^{(i)}} \left(\psi(y, \mathbf{x}^{(i)}) - \psi(y^{(i)}, \mathbf{x}^{(i)}) + 1 \right)_+, \quad [17.8]$$

The sum in Equation 17.8 includes non-zero values for every label that is ranked at least as high as the true entity, after applying the margin augmentation. Dividing by the margin-augmented rank of the true entity thus gives the average violation.

The objective in Equation 17.8 is expensive to optimize when the label space is large, as is usually the case for entity linking against large knowledge bases. This motivates a randomized approximation called **WARP** (Weston et al., 2011), shown in Algorithm 19. In this procedure, we sample random entities until one violates the pairwise margin constraint, $\psi(y, \mathbf{x}^{(i)}) + 1 \geq \psi(y^{(i)}, \mathbf{x}^{(i)})$. The number of samples N required to find such a violation yields an approximation of the margin-augmented rank of the true entity, $r(y^{(i)}, \mathbf{x}^{(i)}) \approx \left\lceil \frac{|\mathcal{Y}(\mathbf{x})|}{N} \right\rceil$. If a violation is found immediately, $N = 1$, the correct entity probably ranks below many others, $r \approx |\mathcal{Y}(\mathbf{x})|$. If many samples are required before a violation is found, $N \rightarrow |\mathcal{Y}(\mathbf{x})|$, then the correct entity is probably highly ranked, $r \rightarrow 1$. A computational advantage of WARP is that it is not necessary to find the highest-scoring label, which can impose a non-trivial computational cost when $\mathcal{Y}(\mathbf{x}^{(i)})$ is large. The objective is conceptually similar to the **negative sampling** objective in WORD2VEC (chapter 14), which compares the observed word against randomly sampled alternatives.

17.2 Relations

After identifying the entities that are mentioned in a text, the next step is to determine how they are related. Consider the following example:

(17.8) George Bush traveled to France on Thursday for a summit.

This sentence introduces a relation between the entities referenced by *George Bush* and *France*. In the Automatic Content Extraction (ACE) ontology (Linguistic Data Consortium, 2005), the type of this relation is PHYSICAL, and the subtype is LOCATED. This relation would be written,

PHYSICAL.LOCATED(GEORGE BUSH, FRANCE). [17.9]

Relations take exactly two arguments, and the order of the arguments matters.

In the ACE datasets, relations are annotated between entity mentions, as in the example above. Relations can also hold between nominals, as in the following example from the SemEval-2010 shared task (Hendrickx et al., 2009):

CAUSE-EFFECT	<i>those cancers were caused by radiation exposures</i>
INSTRUMENT-AGENCY	<i>phone operator</i>
PRODUCT-PRODUCER	<i>a factory manufactures suits</i>
CONTENT-CONTAINER	<i>a bottle of honey was weighed</i>
ENTITY-ORIGIN	<i>letters from foreign countries</i>
ENTITY-DESTINATION	<i>the boy went to bed</i>
COMPONENT-WHOLE	<i>my apartment has a large kitchen</i>
MEMBER-COLLECTION	<i>there are many trees in the forest</i>
COMMUNICATION-TOPIC	<i>the lecture was about semantics</i>

Table 17.1: Relations and example sentences from the SemEval-2010 dataset (Hendrickx et al., 2009)

8949 (17.9) The cup contained tea from dried ginseng.

8950 This sentence describes a relation of type ENTITY-ORIGIN between *tea* and *ginseng*. Nominal
 8951 relation extraction is closely related to **semantic role labeling** (chapter 13). The main
 8952 difference is that relation extraction is restricted to a relatively small number of relation
 8953 types; for example, Table 17.1 shows the ten relation types from SemEval-2010.

8954 17.2.1 Pattern-based relation extraction

8955 Early work on relation extraction focused on hand-crafted patterns (Hearst, 1992). For
 8956 example, the appositive *Starbuck, a native of Nantucket* signals the relation ENTITY-ORIGIN
 8957 between *Starbuck* and *Nantucket*. This pattern can be written as,

$$\text{PERSON , } a \text{ native of LOCATION} \Rightarrow \text{ENTITY-ORIGIN(PERSON, LOCATION)}. \quad [17.10]$$

8958 This pattern will be “triggered” whenever the literal string *, a native of* occurs between an
 8959 entity of type PERSON and an entity of type LOCATION. Such patterns can be generalized
 8960 beyond literal matches using techniques such as lemmatization, which would enable the
 8961 words (*buy, buys, buying*) to trigger the same patterns (see § 4.3.1.2). A more aggressive
 8962 strategy would be to group all words in a WordNet synset (§ 4.2), so that, e.g., *buy* and
 8963 *purchase* trigger the same patterns.

8964 Relation extraction patterns can be implemented in finite-state automata (§ 9.1). If the
 8965 named entity recognizer is also a finite-state machine, then the systems can be combined
 8966 by finite-state transduction (Hobbs et al., 1997). This makes it possible to propagate uncer-
 8967 tainty through the finite-state cascade, and disambiguate from higher-level context. For
 8968 example, suppose the entity recognizer cannot decide whether *Starbuck* refers to either a
 8969 PERSON or a LOCATION; in the composed transducer, the relation extractor would be free
 8970 to select the PERSON annotation when it appears in the context of an appropriate pattern.

8971 **17.2.2 Relation extraction as a classification task**

8972 Relation extraction can be formulated as a classification problem,

$$\hat{r}_{(i,j),(m,n)} = \operatorname{argmax}_{r \in \mathcal{R}} \Psi(r, (i, j), (m, n), \mathbf{w}), \quad [17.11]$$

8973 where $r \in \mathcal{R}$ is a relation type (possibly NIL), $\mathbf{w}_{i+1:j}$ is the span of the first argument, and
 8974 $\mathbf{w}_{m+1:n}$ is the span of the second argument. The argument $\mathbf{w}_{m+1:n}$ may appear before
 8975 or after $\mathbf{w}_{i+1:j}$ in the text, or they may overlap; we stipulate only that $\mathbf{w}_{i+1:j}$ is the first
 8976 argument of the relation. We now consider three alternatives for computing the scoring
 8977 function.

8978 **17.2.2.1 Feature-based classification**

8979 In a feature-based classifier, the scoring function is defined as,

$$\Psi(r, (i, j), (m, n), \mathbf{w}) = \boldsymbol{\theta} \cdot \mathbf{f}(r, (i, j), (m, n), \mathbf{w}), \quad [17.12]$$

8980 with $\boldsymbol{\theta}$ representing a vector of weights, and $\mathbf{f}(\cdot)$ a vector of features. The pattern-based
 8981 methods described in § 17.2.1 suggest several features:

- 8982 • Local features of $\mathbf{w}_{i+1:j}$ and $\mathbf{w}_{m+1:n}$, including: the strings themselves; whether they
 8983 are recognized as entities, and if so, which type; whether the strings are present in a
 8984 **gazetteer** of entity names; each string's syntactic **head** (§ 9.2.2).
- 8985 • Features of the span between the two arguments, $\mathbf{w}_{j+1:m}$ or $\mathbf{w}_{n+1:i}$ (depending on
 8986 which argument appears first): the length of the span; the specific words that appear
 8987 in the span, either as a literal sequence or a bag-of-words; the wordnet synsets (§ 4.2)
 8988 that appear in the span between the arguments.
- 8989 • Features of the syntactic relationship between the two arguments, typically the **de-**
 8990 **pendency path** between the arguments (§ 13.2.1). Example dependency paths are
 8991 shown in Table 17.2.

8992 **17.2.2.2 Kernels**

8993 Suppose that the first line of Table 17.2 is a labeled example, and the remaining lines are
 8994 instances to be classified. A feature-based approach would have to decompose the depen-
 8995 dency paths into features that capture individual edges, with or without their labels, and
 8996 then learn weights for each of these features: for example, the second line contains identi-
 8997 cal dependencies, but different arguments; the third line contains a different inflection of
 8998 the word *travel*; the fourth and fifth lines each contain an additional edge on the depen-
 8999 dency path; and the sixth example uses an entirely different path. Rather than attempting
 9000 to create local features that capture all of the ways in which these dependencies paths

1. <i>George Bush traveled to France</i>	<i>George Bush</i> \leftarrow traveled \rightarrow France NSUBJ OBL
2. <i>Ahab traveled to Nantucket</i>	<i>Ahab</i> \leftarrow traveled \rightarrow Nantucket NSUBJ OBL
3. <i>George Bush will travel to France</i>	<i>George Bush</i> \leftarrow travel \rightarrow France NSUBJ OBL
4. <i>George Bush wants to travel to France</i>	<i>George Bush</i> \leftarrow wants \rightarrow travel \rightarrow France NSUBJ XCOMP OBL
5. <i>Ahab traveled to a city in France</i>	<i>Ahab</i> \leftarrow traveled \rightarrow city \rightarrow France NSUBJ OBL NMOD
6. <i>We await Ahab's visit to France</i>	<i>Ahab</i> \leftarrow visit \rightarrow France NMOD:POSS NMOD

Table 17.2: Candidates instances for the PHYSICAL.LOCATED relation, and their dependency paths

9001 are similar and different, we can instead define a similarity function κ , which computes a
 9002 score for any pair of instances, $\kappa : \mathcal{X} \times \mathcal{X} \mapsto \mathbb{R}_+$. The score for any pair of instances (i, j)
 9003 is $\kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) \geq 0$, with $\kappa(i, j)$ being large when instances $\mathbf{x}^{(i)}$ and $\mathbf{x}^{(j)}$ are similar. If the
 9004 function κ obeys a few key properties it is a valid **kernel function**.⁴

Given a valid kernel function, we can build a non-linear classifier without explicitly defining a feature vector or neural network architecture. For a binary classification problem $y \in \{-1, 1\}$, we have the decision function,

$$\hat{y} = \text{Sign}(b + \sum_{i=1}^N y^{(i)} \alpha^{(i)} \kappa(\mathbf{x}^{(i)}, \mathbf{x})) \quad [17.13]$$

9005 where b and $\{\alpha^{(i)}\}_{i=1}^N$ are parameters that must be learned from the training set, under
 9006 the constraint $\forall_i, \alpha^{(i)} \geq 0$. Intuitively, each α_i specifies the importance of the instance $\mathbf{x}^{(i)}$
 9007 towards the classification rule. Kernel-based classification can be viewed as a weighted
 9008 form of the **nearest-neighbor** classifier (Hastie et al., 2009), in which test instances are
 9009 assigned the most common label among their near neighbors in the training set. This
 9010 results in a non-linear classification boundary. The parameters are typically learned from
 9011 a margin-based objective (see § 2.3), leading to the **kernel support vector machine**. To
 9012 generalize to multi-class classification, we can train separate binary classifiers for each
 9013 label (sometimes called **one-versus-all**), or train binary classifiers for each pair of possible
 9014 labels (**one-versus-one**).

9015 Dependency kernels are particularly effective for relation extraction, due to their ability
 9016 to capture syntactic properties of the path between the two candidate arguments. One
 9017 class of dependency tree kernels is defined recursively, with the score for a pair of trees

⁴The **Gram matrix** \mathbf{K} arises from computing the kernel function between all pairs in a set of instances. For a valid kernel, the Gram matrix must be symmetric ($\mathbf{K} = \mathbf{K}^\top$) and positive semi-definite ($\forall \mathbf{a}, \mathbf{a}^\top \mathbf{K} \mathbf{a} \geq 0$). For more on kernel-based classification, see chapter 14 of Murphy (2012).

equal to the similarity of the root nodes and the sum of similarities of matched pairs of child subtrees (Zelenko et al., 2003; Culotta and Sorensen, 2004). Alternatively, Bunescu and Mooney (2005) define a kernel function over sequences of unlabeled dependency edges, in which the score is computed as a product of scores for each pair of words in the sequence: identical words receive a high score, words that share a synset or part-of-speech receive a small non-zero score (e.g., *travel* / *visit*), and unrelated words receive a score of zero.

17.2.2.3 Neural relation extraction

Convolutional neural networks were an early neural architecture for relation extraction (Zeng et al., 2014; dos Santos et al., 2015). For the sentence (w_1, w_2, \dots, w_M) , obtain a matrix of word embeddings \mathbf{X} , where $x_m \in \mathbb{R}^K$ is the embedding of w_m . Now, suppose the candidate arguments appear at positions a_1 and a_2 ; then for each word in the sentence, its position with respect to each argument is $m - a_1$ and $m - a_2$. (Following Zeng et al. (2014), this is a restricted version of the relation extraction task in which the arguments are single tokens.) To capture any information conveyed by these positions, the word embeddings are concatenated with embeddings of the positional offsets, $x_{m-a_1}^{(p)}$ and $x_{m-a_2}^{(p)}$. The complete base representation of the sentence is,

$$\mathbf{X}(a_1, a_2) = \begin{pmatrix} \mathbf{x}_1 & \mathbf{x}_2 & \cdots & \mathbf{x}_M \\ \mathbf{x}_{1-a_1}^{(p)} & \mathbf{x}_{2-a_1}^{(p)} & \cdots & \mathbf{x}_{M-a_1}^{(p)} \\ \mathbf{x}_{1-a_2}^{(p)} & \mathbf{x}_{2-a_2}^{(p)} & \cdots & \mathbf{x}_{M-a_2}^{(p)} \end{pmatrix}, \quad [17.14]$$

where each column is a vertical concatenation of a word embedding, represented by the column vector \mathbf{x}_m , and two positional embeddings, specifying the position with respect to a_1 and a_2 . The matrix $\mathbf{X}(a_1, a_2)$ is then taken as input to a convolutional layer (see § 3.4), and max-pooling is applied to obtain a vector. The final scoring function is then,

$$\Psi(r, i, j, \mathbf{X}) = \theta_r \cdot \text{MaxPool}(\text{ConvNet}(\mathbf{X}(i, j); \phi)), \quad [17.15]$$

where ϕ defines the parameters of the convolutional operator, and the θ_r defines a set of weights for relation r . The model can be trained using a margin objective,

$$\hat{r} = \underset{r}{\operatorname{argmax}} \Psi(r, i, j, \mathbf{X}) \quad [17.16]$$

$$\ell = (1 + \psi(\hat{r}, i, j, \mathbf{X}) - \psi(r, i, j, \mathbf{X}))_+. \quad [17.17]$$

Recurrent neural networks have also been applied to relation extraction, using a network such as an bidirectional LSTM to encode the words or dependency path between the two arguments. Xu et al. (2015) segment each dependency path into left and right subpaths: the path $George \xleftarrow[\text{NSUBJ}]{wants} Bush \xrightarrow[\text{XCOMP}]{travel} France$ is segmented into the subpaths,

9044 (17.10) *George Bush* $\xleftarrow[\text{NSUBJ}]{} \text{wants}$

9045 (17.11) $\text{wants} \xrightarrow[\text{XCOMP}]{} \text{travel} \xrightarrow[\text{OBL}]{} \text{France.}$

Xu et al. (2015) then run recurrent networks from the arguments to the root word (in this case, *wants*), obtaining the final representation by max pooling across all the recurrent states along each path. This process can be applied across separate “channels”, in which the inputs consist of embeddings for the words, parts-of-speech, dependency relations, and WordNet hypernyms. To define the model formally, let $s(m)$ define the successor of word m in either the left or right subpath (in a dependency path, each word can have a successor in at most one subpath). Let $\mathbf{x}_m^{(c)}$ indicate the embedding of word (or relation) m in channel c , and let $\overleftarrow{\mathbf{h}}_m^{(c)}$ and $\overrightarrow{\mathbf{h}}_m^{(c)}$ indicate the associated recurrent states in the left and right subtrees respectively. Then the complete model is specified as follows,

$$\mathbf{h}_{s(m)}^{(c)} = \text{RNN}(\mathbf{x}_{s(m)}^{(c)}, \mathbf{h}_m^{(c)}) \quad [17.18]$$

$$\mathbf{z}^{(c)} = \text{MaxPool}(\overleftarrow{\mathbf{h}}_i^{(c)}, \overleftarrow{\mathbf{h}}_{s(i)}^{(c)}, \dots, \overleftarrow{\mathbf{h}}_{\text{root}}^{(c)}, \overrightarrow{\mathbf{h}}_j^{(c)}, \overrightarrow{\mathbf{h}}_{s(j)}^{(c)}, \dots, \overrightarrow{\mathbf{h}}_{\text{root}}^{(c)}) \quad [17.19]$$

$$\Psi(r, i, j) = \theta \cdot [\mathbf{z}^{(\text{word})}; \mathbf{z}^{(\text{POS})}; \mathbf{z}^{(\text{dependency})}; \mathbf{z}^{(\text{hypernym})}]. \quad [17.20]$$

9046 Note that \mathbf{z} is computed by applying max-pooling to the *matrix* of horizontally concatenated vectors \mathbf{h} , while Ψ is computed from the *vector* of vertically concatenated vectors 9047 \mathbf{z} . Xu et al. (2015) pass the score Ψ through a **softmax** layer to obtain a probability 9048 $p(r | i, j, w)$, and train the model by regularized **cross-entropy**. Miwa and Bansal (2016) 9049 show that a related model can solve the more challenging “end-to-end” relation extraction 9050 task, in which the model must simultaneously detect entities and then extract their 9051 relations.

9053 17.2.3 Knowledge base population

9054 In many applications, what matters is not what fraction of sentences are analyzed correctly, but how much accurate knowledge can be extracted. **Knowledge base population** 9055 (**KBP**) refers to the task of filling in Wikipedia-style infoboxes, as shown in Figure 17.1a. 9056 Knowledge base population can be decomposed into two subtasks: **entity linking** (described in § 17.1), and **slot filling** (Ji and Grishman, 2011). Slot filling has two key differences from the formulation of relation extraction presented above: the relations hold 9057 between entities rather than spans of text, and the performance is evaluated at the *type* 9058 level (on entity pairs), rather than on the *token level* (on individual sentences).

9062 From a practical standpoint, there are three other important differences between slot 9063 filling and per-sentence relation extraction.

- KBP tasks are often formulated from the perspective of identifying attributes of a few “query” entities. As a result, these systems often start with an **information retrieval** phase, in which relevant passages of text are obtained by search.
- For many entity pairs, there will be multiple passages of text that provide evidence. Slot filling systems must aggregate this evidence to predict a single relation type (or set of relations).
- Labeled data is usually available in the form of pairs of related entities, rather than annotated passages of text. Training from such type-level annotations is a challenge: two entities may be linked by several relations, or they may appear together in a passage of text that nonetheless does not describe their relation to each other.

Information retrieval is beyond the scope of this text (see Manning et al., 2008). The remainder of this section describes approaches to information fusion and learning from type-level annotations.

17.2.3.1 Information fusion

In knowledge base population, there will often be multiple pieces of evidence for (and sometimes against) a single relation. For example, a search for the entity MAYNARD JACKSON, JR. may return several passages that reference the entity ATLANTA:⁵

- (17.12) Elected mayor of Atlanta in 1973, **Maynard Jackson** was the first African American to serve as mayor of a major southern city.
- (17.13) **Atlanta**’s airport will be renamed to honor **Maynard Jackson**, the city’s first Black mayor .
- (17.14) Born in Dallas, Texas in 1938, **Maynard Holbrook Jackson, Jr.** moved to Atlanta when he was 8.
- (17.15) **Maynard Jackson** has gone from one of the worst high schools in **Atlanta** to one of the best.

The first and second examples provide evidence for the relation **MAYOR** holding between the entities **ATLANTA** and **MAYNARD JACKSON, JR.**. The third example provides evidence for a different relation between these same entities, **LIVED-IN**. The fourth example poses an entity linking problem, referring to **MAYNARD JACKSON HIGH SCHOOL**. Knowledge base population requires aggregating this sort of textual evidence, and predicting the relations that are most likely to hold.

⁵First three examples from: <http://www.georgiaencyclopedia.org/articles/government-politics/maynard-jackson-1938-2003>; JET magazine, November 10, 2003; www.todayingeorgiahistory.org/content/maynard-jackson-elected

9095 One approach is to run a single-document relation extraction system (using the tech-
 9096 niques described in § 17.2.2), and then aggregate the results (Li et al., 2011). Relations
 9097 that are detected with high confidence in multiple documents are more likely to be valid,
 9098 motivating the heuristic,

$$\psi(r, e_1, e_2) = \sum_{i=1}^N (\text{p}(r(e_1, e_2) | \mathbf{w}^{(i)}))^\alpha, \quad [17.21]$$

9099 where $\text{p}(r(e_1, e_2) | \mathbf{w}^{(i)})$ is the probability of relation r between entities e_1 and e_2 condi-
 9100 tioned on the text $\mathbf{w}^{(i)}$, and $\alpha \gg 1$ is a tunable hyperparameter. Using this heuristic, it is
 9101 possible to rank all candidate relations, and trace out a **precision-recall curve** as more rel-
 9102ations are extracted.⁶ Alternatively, features can be aggregated across multiple passages
 9103 of text, feeding a single type-level relation extraction system (Wolfe et al., 2017).

9104 Precision can be improved by introducing constraints across multiple relations. For
 9105 example, if we are certain of the relation $\text{PARENT}(e_1, e_2)$, then it cannot also be the case
 9106 that $\text{PARENT}(e_2, e_1)$. Integer linear programming makes it possible to incorporate such
 9107 constraints into a global optimization (Li et al., 2011). Other pairs of relations have pos-
 9108itive correlations, such $\text{MAYOR}(e_1, e_2)$ and $\text{LIVED-IN}(e_1, e_2)$. Compatibility across relation
 9109 types can be incorporated into probabilistic graphical models (e.g., Riedel et al., 2010).

9110 17.2.3.2 Distant supervision

9111 Relation extraction is “annotation hungry,” because each relation requires its own la-
 9112 beled data. Rather than relying on annotations of individual documents, it would be
 9113 preferable to use existing knowledge resources — such as the many facts that are al-
 9114 ready captured in knowledge bases like DBpedia. However such annotations raise the
 9115 inverse of the information fusion problem considered above: the existence of the relation
 9116 $\text{MAYOR}(\text{MAYNARD JACKSON JR., ATLANTA})$ provides only **distant supervision** for the
 9117 example texts in which this entity pair is mentioned.

9118 One approach is to treat the entity pair as the instance, rather than the text itself (Mintz
 9119 et al., 2009). Features are then aggregated across all sentences in which both entities are
 9120 mentioned, and labels correspond to the relation (if any) between the entities in a knowl-
 9121 edge base, such as FreeBase. Negative instances are constructed from entity pairs that are
 9122 not related in the knowledge base. In some cases, two entities are related, but the knowl-
 9123 edge base is missing the relation; however, because the number of possible entity pairs is
 9124 huge, these missing relations are presumed to be relatively rare. This approach is shown
 9125 in Figure 17.2.

⁶The precision-recall curve is similar to the ROC curve shown in Figure 4.4, but it includes the precision $\frac{\text{TP}}{\text{TP} + \text{FP}}$ rather than the false positive rate $\frac{\text{FP}}{\text{FP} + \text{TN}}$.

- **Label** : MAYOR(ATLANTA, MAYNARD JACKSON)
 - Elected mayor of **Atlanta** in 1973, **Maynard Jackson** ...
 - **Atlanta**'s airport will be renamed to honor **Maynard Jackson**, the city's first Black mayor
 - Born in Dallas, Texas in 1938, **Maynard Holbrook Jackson, Jr.** moved to **Atlanta** when he was 8.
- **Label** : MAYOR(NEW YORK, FIORELLO LA GUARDIA)
 - **Fiorello La Guardia** was Mayor of **New York** for three terms ...
 - **Fiorello La Guardia**, then serving on the **New York** City Board of Aldermen...
- **Label** : BORN-IN(DALLAS, MAYNARD JACKSON)
 - Born in **Dallas**, Texas in 1938, **Maynard Holbrook Jackson, Jr.** moved to Atlanta when he was 8.
 - **Maynard Jackson** was raised in **Dallas** ...
- **Label** : NIL(NEW YORK, MAYNARD JACKSON)
 - **Jackson** married Valerie Richardson, whom he had met in **New York**...
 - **Jackson** was a member of the Georgia and **New York** bars ...

Figure 17.2: Four training instances for relation classification using **distant supervision** Mintz et al. (2009). The first two instances are positive for the MAYOR relation, and the third instance is positive for the BORN-IN relation. The fourth instance is a negative example, constructed from a pair of entities (NEW YORK, MAYNARD JACKSON) that do not appear in any Freebase relation. Each instance's features are computed by aggregating across all sentences in which the two entities are mentioned.

9126 In **multiple instance learning**, labels are assigned to *sets* of instances, of which only
 9127 an unknown subset are actually relevant (Dietterich et al., 1997; Maron and Lozano-Pérez,
 9128 1998). This formalizes the framework of distant supervision: the relation $\text{REL}(A, B)$ acts
 9129 as a label for the entire set of sentences mentioning entities A and B, even when only a
 9130 subset of these sentences actually describes the relation. One approach to multi-instance
 9131 learning is to introduce a binary **latent variable** for each sentence, indicating whether the
 9132 sentence expresses the labeled relation (Riedel et al., 2010). A variety of inference tech-
 9133 niques have been employed for this probabilistic model of relation extraction: Surdeanu
 9134 et al. (2012) use expectation maximization, Riedel et al. (2010) use sampling, and Hoff-
 9135 mann et al. (2011) use a custom graph-based algorithm. Expectation maximization and
 9136 sampling are surveyed in chapter 5, and are covered in more detail by Murphy (2012);
 9137 graph-based methods are surveyed by Mihalcea and Radev (2011).

Task	Relation ontology	Supervision
PropBank semantic role labeling	VerbNet	sentence
FrameNet semantic role labeling	FrameNet	sentence
Relation extraction	ACE, TAC, SemEval, etc	sentence
Slot filling	ACE, TAC, SemEval, etc	relation
Open Information Extraction	open	seed relations or patterns

Table 17.3: Various relation extraction tasks and their properties. VerbNet and FrameNet are described in chapter 13. ACE (Linguistic Data Consortium, 2005), TAC (McNamee and Dang, 2009), and SemEval (Hendrickx et al., 2009) refer to shared tasks, each of which involves an ontology of relation types.

9138 17.2.4 Open information extraction

9139 In classical relation extraction, the set of relations is defined in advance, using a **schema**.
 9140 The relation for any pair of entities can then be predicted using multi-class classification.
 9141 In **open information extraction** (OpenIE), a relation can be any triple of text. The example
 9142 sentence (17.12) instantiates several “relations” of this sort:

- 9143 • (*mayor of, Maynard Jackson, Atlanta*),
- 9144 • (*elected, Maynard Jackson, mayor of Atlanta*),
- 9145 • (*elected in, Maynard Jackson, 1973*),

9146 and so on. Extracting such tuples can be viewed as a lightweight version of **semantic role**
 9147 **labeling** (chapter 13), with only two argument types: first slot and second slot. The task is
 9148 generally evaluated on the relation level, rather than on the level of sentences: precision is
 9149 measured by the number of extracted relations that are accurate, and recall is measured by
 9150 the number of true relations that were successfully extracted. OpenIE systems are trained
 9151 from distant supervision or bootstrapping, rather than from labeled sentences.

9152 An early example is the TextRunner system (Banko et al., 2007), which identifies re-
 9153 lations with a set of handcrafted syntactic rules. The examples that are acquired from the
 9154 handcrafted rules are then used to train a classification model that uses part-of-speech pat-
 9155 terns as features. Finally, the relations that are extracted by the classifier are aggregated,
 9156 removing redundant relations and computing the number of times that each relation is
 9157 mentioned in the corpus. TextRunner was the first in a series of systems that performed
 9158 increasingly accurate open relation extraction by incorporating more precise linguistic fea-
 9159 tures (Etzioni et al., 2011), distant supervision from Wikipedia infoboxes (Wu and Weld,
 9160 2010), and better learning algorithms (Zhu et al., 2009).

9161 17.3 Events

9162 Relations link pairs of entities, but many real-world situations involve more than two en-
9163 tities. Consider again the example sentence (17.12), which describes the **event** of an elec-
9164 tion, with four properties: the office (MAYOR), the district (ATLANTA), the date (1973), and
9165 the person elected (MAYNARD JACKSON, JR.). In **event detection**, a schema is provided
9166 for each event type (e.g., an election, a terrorist attack, or a chemical reaction), indicating
9167 all the possible properties of the event. The system is then required to fill in as many of
9168 these properties as possible (Doddington et al., 2004).

9169 Event detection systems generally involve a retrieval component (finding relevant
9170 documents and passages of text) and an extraction component (determining the proper-
9171 ties of the event based on the retrieved texts). Early approaches focused on finite-state pat-
9172 terns for identify event properties (Hobbs et al., 1997); such patterns can be automatically
9173 induced by searching for patterns that are especially likely to appear in documents that
9174 match the event query (Riloff, 1996). Contemporary approaches employ techniques that
9175 are similar to FrameNet semantic role labeling (§ 13.2), such as structured prediction over
9176 local and global features (Li et al., 2013) and bidirectional recurrent neural networks (Feng
9177 et al., 2016). These methods detect whether an event is described in a sentence, and if so,
9178 what are its properties.

9179 **Event coreference** Because multiple sentences may describe unique properties of a sin-
9180 gle event, **event coreference** is required to link event mentions across a single passage
9181 of text, or between passages (Humphreys et al., 1997). Bejan and Harabagiu (2014) de-
9182 fine event coreference as the task of identifying event mentions that share the same event
9183 participants (i.e., the slot-filling entities) and the same event properties (e.g., the time and
9184 location), within or across documents. Event coreference resolution can be performed us-
9185 ing supervised learning techniques in a similar way to entity coreference, as described
9186 in chapter 15: move left-to-right through the document, and use a classifier to decide
9187 whether to link each event reference to an existing cluster of coreferent events, or to cre-
9188 ate a new cluster (Ahn, 2006). Each clustering decision is based on the compatibility of
9189 features describing the participants and properties of the event. Due to the difficulty of
9190 annotating large amounts of data for entity coreference, unsupervised approaches are es-
9191 pecially desirable (Chen and Ji, 2009; Bejan and Harabagiu, 2014).

9192 **Relations between events** Just as entities are related to other entities, events may be
9193 related to other events: for example, the event of winning an election both *precedes* and
9194 *causes* the event of serving as mayor; moving to Atlanta *precedes* and *enables* the event of
9195 becoming mayor of Atlanta; moving from Dallas to Atlanta *prevents* the event of later be-
9196 coming mayor of Dallas. As these examples show, events may be related both temporally
9197 and causally. The **TimeML** annotation scheme specifies a set of six temporal relations

	Positive (+)	Negative (-)	Underspecified (u)
Certain (CT)	Fact: CT+	Counterfact: CT-	Certain, but unknown: CTU
Probable (PR)	Probable: PR+	Not probable: PR-	(NA)
Possible (PS)	Possible: PS+	Not possible: PS-	(NA)
Underspecified (U)	(NA)	(NA)	Unknown or uncommitted: UU

Table 17.4: Table of factuality values from the FactBank corpus (Saurí and Pustejovsky, 2009). The entry (NA) indicates that this combination is not annotated.

9198 between events (Pustejovsky et al., 2005), derived in part from **interval algebra** (Allen,
9199 1984). The TimeBank corpus provides TimeML annotations for 186 documents (Pustejovsky
9200 et al., 2003). Methods for detecting these temporal relations combine supervised
9201 machine learning with temporal constraints, such as transitivity (e.g. Mani et al., 2006;
9202 Chambers and Jurafsky, 2008).

9203 More recent annotation schemes and datasets combine temporal and causal relations (Mirza
9204 et al., 2014; Dunietz et al., 2017): for example, the CaTeRS dataset includes annotations of
9205 320 five-sentence short stories (Mostafazadeh et al., 2016). Abstracting still further, **processes**
9206 are networks of causal relations between multiple events. A small dataset of biological
9207 processes is annotated in the ProcessBank dataset (Berant et al., 2014), with the
9208 goal of supporting automatic question answering on scientific textbooks.

9209 17.4 Hedges, denials, and hypotheticals

9210 The methods described thus far apply to **propositions** about the way things are in the
9211 real world. But natural language can also describe events and relations that are likely or
9212 unlikely, possible or impossible, desired or feared. The following examples hint at the
9213 scope of the problem (Prabhakaran et al., 2010):

- 9214 (17.16) GM will lay off workers.
- 9215 (17.17) A spokesman for GM said GM will lay off workers.
- 9216 (17.18) GM may lay off workers.
- 9217 (17.19) The politician claimed that GM will lay off workers.
- 9218 (17.20) Some wish GM would lay off workers.
- 9219 (17.21) Will GM lay off workers?
- 9220 (17.22) Many wonder whether GM will lay off workers.

9221 Accurate information extraction requires handling these **extra-propositional** aspects
9222 of meaning, which are sometimes summarized under the terms **modality** and **negation**.⁷
9223 Modality refers to expressions of the speaker's attitude towards her own statements, in-
9224 cluding "degree of certainty, reliability, subjectivity, sources of information, and perspec-
9225 tive" (Morante and Sporleder, 2012). Various systematizations of modality have been
9226 proposed (e.g., Palmer, 2001), including categories such as future, interrogative, imper-
9227 ative, conditional, and subjective. Information extraction is particularly concerned with
9228 negation and certainty. For example, Saurí and Pustejovsky (2009) link negation with
9229 a modal calculus of certainty, likelihood, and possibility, creating the two-dimensional
9230 schema shown in Table 17.4. This is the basis for the FactBank corpus, with annotations
9231 of the **factuality** of all sentences in 208 documents of news text.

9232 A related concept is **hedging**, in which speakers limit their commitment to a proposi-
9233 tion (Lakoff, 1973):

- 9234 (17.23) These results **suggest** that expression of c-jun, jun B and jun D genes **might** be in-
9235 volved in terminal granulocyte differentiation... (Morante and Daelemans, 2009)
9236 (17.24) A whale is **technically** a mammal (Lakoff, 1973)

9237 In the first example, the hedges *suggest* and *might* communicate uncertainty; in the second
9238 example, there is no uncertainty, but the hedge *technically* indicates that the evidence for
9239 the proposition will not fully meet the reader's expectations. Hedging has been studied
9240 extensively in scientific texts (Medlock and Briscoe, 2007; Morante and Daelemans, 2009),
9241 where the goal of large-scale extraction of scientific facts is obstructed by hedges and spec-
9242 ulation. Still another related aspect of modality is **evidentiality**, in which speakers mark
9243 the source of their information. In many languages, it is obligatory to mark evidentiality
9244 through affixes or particles (Aikhenvald, 2004); while evidentiality is not grammaticalized
9245 in English, authors are expected to express this information in contexts such as journal-
9246 ism (Kovach and Rosenstiel, 2014) and Wikipedia.⁸

9247 Methods for handling negation and modality generally include two phases:

- 9248 1. detecting negated or uncertain events;
9249 2. identifying the scope and focus of the negation or modal operator.

⁷The classification of negation as extra-propositional is controversial: Packard et al. (2014) argue that negation is a "core part of compositionally constructed logical-form representations." Negation is an element of the semantic parsing tasks discussed in chapter 12 and chapter 13 — for example, negation markers are treated as adjuncts in PropBank semantic role labeling. However, many of the relation extraction methods mentioned in this chapter do not handle negation directly. A further consideration is that negation interacts closely with aspects of modality that are generally not considered in propositional semantics, such as certainty and subjectivity.

⁸<https://en.wikipedia.org/wiki/Wikipedia:Verifiability>

9250 A considerable body of work on negation has employed rule-based techniques such as
 9251 regular expressions (Chapman et al., 2001) to detect negated events. Such techniques
 9252 match lexical cues (e.g., *Norwood was not elected Mayor*), while avoiding “double nega-
 9253 tives” (e.g., *surely all this is not without meaning*). More recent approaches employ classi-
 9254 fiers over lexical and syntactic features (Uzuner et al., 2009) and sequence labeling (Prab-
 9255 hakaran et al., 2010).

9256 The tasks of scope and focus resolution are more fine grained, as shown in the example
 9257 from Morante and Sporleder (2012):

- 9258 (17.25) [After his habit he said] **nothing**, and after mine I asked no questions.
 9259 After his habit he said nothing, and [after mine I asked] **no** [questions].

9260 In this sentence, there are two negation cues (*nothing* and *no*). Each negates an event,
 9261 indicated by the underlined verbs *said* and *asked* (this is the focus of negation), and each
 9262 occurs within a scope: *after his habit he said* and *after mine I asked* ____ *questions*. These tasks
 9263 are typically formalized as sequence labeling problems, with each word token labeled
 9264 as beginning, inside, or outside of a cue, focus, or scope span (see § 8.3). Conventional
 9265 sequence labeling approaches can then be applied, using surface features as well as syn-
 9266 tax (Velldal et al., 2012) and semantic analysis (Packard et al., 2014). Labeled datasets
 9267 include the BioScope corpus of biomedical texts (Vincze et al., 2008) and a shared task
 9268 dataset of detective stories by Arthur Conan Doyle (Morante and Blanco, 2012).

9269 17.5 Question answering and machine reading

9270 The victory of the Watson question-answering system against three top human players on
 9271 the game show *Jeopardy!* was a landmark moment for natural language processing (Fer-
 9272 rucci et al., 2010). Game show questions are usually answered by **factoids**: entity names
 9273 and short phrases.⁹ The task of factoid question answering is therefore closely related to
 9274 information extraction, with the additional problem of accurately parsing the question.

9275 17.5.1 Formal semantics

9276 Semantic parsing is an effective method for question-answering in restricted domains
 9277 such as questions about geography and airline reservations (Zettlemoyer and Collins,
 9278 2005), and has also been applied in “open-domain” settings such as question answering
 9279 on Freebase (Berant et al., 2013) and biomedical research abstracts (Poon and Domingos,
 9280 2009). One approach is to convert the question into a lambda calculus expression that

⁹The broader landscape of question answering includes “why” questions (*Why did Ahab continue to pursue the white whale?*), “how questions” (*How did Queequeg die?*), and requests for summaries (*What was Ishmael’s attitude towards organized religion?*). For more, see Hirschman and Gaizauskas (2001).

9281 returns a boolean value: for example, the question *who is the mayor of the capital of Georgia?*
 9282 would be converted to,

$$\lambda x. \exists y \text{ CAPITAL(GEORGIA, } y) \wedge \text{MAYOR}(y, x). \quad [17.22]$$

9283 This lambda expression can then be used to query an existing knowledge base, returning
 9284 “true” for all entities that satisfy it.

9285 17.5.2 Machine reading

9286 Recent work has focused on answering questions about specific textual passages, similar
 9287 to the reading comprehension examinations for young students (Hirschman et al., 1999).
 9288 This task has come to be known as **machine reading**.

9289 17.5.2.1 Datasets

9290 The machine reading problem can be formulated in a number of different ways. The most
 9291 important distinction is what form the answer should take.

- 9292 • **Multiple-choice question answering**, as in the MCTest dataset of stories (Richardson et al., 2013) and the New York Regents Science Exams (Clark, 2015). In MCTest,
 9293 the answer is deducible from the text alone, while in the science exams, the system
 9294 must make inferences using an existing model of the underlying scientific phenomena.
 9295 Here is an example from MCTest:

9297 (17.26) James the turtle was always getting into trouble. Sometimes he'd reach into
 9298 the freezer and empty out all the food ...

9299 Q: What is the name of the trouble making turtle?
 9300 (a) Fries
 9301 (b) Pudding
 9302 (c) James
 9303 (d) Jane

- 9304 • **Cloze-style “fill in the blank”** questions, as in the CNN/Daily Mail comprehension
 9305 task (Hermann et al., 2015), the Children’s Book Test (Hill et al., 2016), and the Who-
 9306 did-What dataset (Onishi et al., 2016). In these tasks, the system must guess which
 9307 word or entity completes a sentence, based on reading a passage of text. Here is an
 9308 example from Who-did-What:

9309 (17.27) Q: Tottenham manager Juande Ramos has hinted he will allow ____ to leave
 9310 if the Bulgaria striker makes it clear he is unhappy. (Onishi et al., 2016)

9311 The query sentence may be selected either from the story itself, or from an external
 9312 summary. In either case, datasets can be created automatically by processing large
 9313 quantities existing documents. An additional constraint is that that missing element
 9314 from the cloze must appear in the main passage of text: for example, in Who-did-
 9315 What, the candidates include all entities mentioned in the main passage. In the
 9316 CNN/Daily Mail dataset, each entity name is replaced by a unique identifier, e.g.,
 9317 ENTITY37. This ensures that correct answers can only be obtained by accurately
 9318 reading the text, and not from external knowledge about the entities.

- 9319 • **Extractive** question answering, in which the answer is drawn from the original text.
 9320 In WikiQA, answers are sentences (Yang et al., 2015). In the Stanford Question An-
 9321 swering Dataset (SQuAD), answers are words or short phrases (Rajpurkar et al.,
 9322 2016):

9323 (17.28) In metereology, precipitation is any product of the condensation of atmo-
 9324 spheric water vapor that falls under gravity.
 9325 Q: What causes precipitation to fall? A: gravity

9326 In both WikiQA and SQuAD, the original texts are Wikipedia articles, and the ques-
 9327 tions are generated by crowdworkers.

9328 **17.5.2.2 Methods**

9329 A baseline method is to search the text for sentences or short passages that overlap with
 9330 both the query and the candidate answer (Richardson et al., 2013). In example (17.26), this
 9331 baseline would select the correct answer, since *James* appears in a sentence that includes
 9332 the query terms *trouble* and *turtle*.

This baseline can be implemented as a neural architecture, using an **attention mechanism** (see § 18.3.1), which scores the similarity of the query to each part of the source text (Chen et al., 2016). The first step is to encode the passage $w^{(p)}$ and the query $w^{(q)}$, using two bidirectional LSTMs (§ 7.6).

$$\mathbf{h}^{(q)} = \text{BiLSTM}(\mathbf{w}^{(q)}; \Theta^{(q)}) \quad [17.23]$$

$$\mathbf{h}^{(p)} = \text{BiLSTM}(\mathbf{w}^{(p)}; \Theta^{(p)}). \quad [17.24]$$

The query is represented by vertically concatenating the final states of the left-to-right and right-to-left passes:

$$\mathbf{u} = [\overrightarrow{\mathbf{h}}^{(q)}_{M_q}; \overleftarrow{\mathbf{h}}^{(q)}_0]. \quad [17.25]$$

The attention vector is computed as a softmax over a vector of bilinear products, and the expected representation is computed by summing over attention values,

$$\tilde{\alpha}_m = (\mathbf{u}^{(q)})^\top \mathbf{W}_a \mathbf{h}_m^{(p)} \quad [17.26]$$

$$\boldsymbol{\alpha} = \text{SoftMax}(\tilde{\boldsymbol{\alpha}}) \quad [17.27]$$

$$\mathbf{o} = \sum_{m=1}^M \alpha_m \mathbf{h}_m^{(p)}. \quad [17.28]$$

Each candidate answer c is represented by a vector \mathbf{x}_c . Assuming the candidate answers are spans from the original text, these vectors can be set equal to the corresponding element in $\mathbf{h}^{(p)}$. The score for each candidate answer a is computed by the inner product,

$$\hat{c} = \underset{c}{\operatorname{argmax}} \mathbf{o} \cdot \mathbf{x}_c. \quad [17.29]$$

9333 This architecture can be trained end-to-end from a loss based on the log-likelihood of the
 9334 correct answer. A number of related architectures have been proposed (e.g., Hermann
 9335 et al., 2015; Kadlec et al., 2016; Dhingra et al., 2017; Cui et al., 2017), and the relationships
 9336 between these methods are surveyed by Wang et al. (2017).

9337 Additional resources

9338 The field of information extraction is surveyed in course notes by Grishman (2012), and
 9339 more recently in a short survey paper (Grishman, 2015). Shen et al. (2015) survey the task
 9340 of entity linking, and Ji and Grishman (2011) survey work on knowledge base popula-
 9341 tion. This chapter’s discussion of non-propositional meaning was strongly influenced by
 9342 Morante and Sporleder (2012), who introduced a special issue of the journal *Computational
 9343 Linguistics* dedicated to recent work on modality and negation.

9344 Exercises

9345 1. Consider the following heuristic for entity linking:

- 9346 • Among all entities that have the same type as the mention (e.g., LOC, PER),
 9347 choose the one whose name has the lowest edit distance from the mention.
- 9348 • If more than one entity has the right type and the lowest edit distance from the
 9349 mention, choose the most popular one.
- 9350 • If no candidate entity has the right type, choose NIL.

Now suppose you have the following feature function:

$$f(y, \mathbf{x}) = [\text{edit-dist}(\text{name}(y), \mathbf{x}), \text{same-type}(y, \mathbf{x}), \text{popularity}(y), \delta(y = \text{NIL})]$$

Design a set of ranking weights θ that match the heuristic. You may assume that edit distance and popularity are always in the range [0, 100], and that the NIL entity has values of zero for all features except δ ($y = \text{NIL}$).

2. Now consider another heuristic:

- Among all candidate entities that have edit distance zero from the mention and the right type, choose the most popular one.
- If no entity has edit distance zero from the mention, choose the one with the right type that is most popular, regardless of edit distance.
- If no entity has the right type, choose NIL.

Using the same features and assumptions from the previous problem, prove that there is no set of weights that could implement this heuristic. Then show that the heuristic can be implemented by adding a single feature. Your new feature should consider only the edit distance.

3. * Consider the following formulation for collective entity linking, which rewards sets of entities that are all of the same type, where “types” can be elements of any set:

$$\psi_c(\mathbf{y}) = \begin{cases} \alpha & \text{all entities in } \mathbf{y} \text{ have the same type} \\ \beta & \text{more than half of the entities in } \mathbf{y} \text{ have the same type} \\ 0 & \text{otherwise.} \end{cases} \quad [17.30]$$

Show how to implement this model of collective entity linking in an **integer linear program**. You may want to review § 13.2.2.

To get started, here is an integer linear program for entity linking, without including the collective term ψ_c :

$$\begin{aligned} \max_{z_{i,y} \in \{0,1\}} \quad & \sum_{i=1}^N \sum_{y \in \mathcal{Y}(\mathbf{x}^{(i)})} s_{i,y} z_{i,y} \\ \text{s.t.} \quad & \sum_{y \in \mathcal{Y}(\mathbf{x}^{(i)})} z_{i,y} \leq 1 \quad \forall i \in \{1, 2, \dots, N\} \end{aligned}$$

where $z_{i,y} = 1$ if entity y is linked to mention i , and $s_{i,y}$ is a parameter that scores the quality of this individual ranking decision, e.g., $s_{i,y} = \theta \cdot f(y, \mathbf{x}^{(i)}, \mathbf{c}^{(i)})$.

To incorporate the collective linking score, you may assume parameters r ,

$$r_{y,\tau} = \begin{cases} 1, & \text{entity } y \text{ has type } \tau \\ 0, & \text{otherwise.} \end{cases} \quad [17.31]$$

Hint: You will need to define several auxiliary variables to optimize over.

- 9373 4. Run `nltk.corpus.download('reuters')` to download the Reuters corpus in
 9374 NLTK, and run `from nltk.corpus import reuters` to import it. The com-
 9375 mand `reuters.words()` returns an iterator over the tokens in the corpus.
- 9376 a) Apply the pattern *_____, such as _____* to this corpus, obtaining candidates for the
 9377 IS-A relation, e.g. `IS-A(ROMANIA, COUNTRY)`. What are three pairs that this
 9378 method identifies correctly? What are three different pairs that it gets wrong?
- 9379 b) Design a pattern for the PRESIDENT relation, e.g. `PRESIDENT(PHILIPPINES, CORAZON AQUINO)`.
 9380 In this case, you may want to augment your pattern matcher with the ability
 9381 to match multiple token wildcards, perhaps using case information to detect
 9382 proper names. Again, list three correct
- 9383 c) Preprocess the Reuters data by running a named entity recognizer, replacing
 9384 tokens with named entity spans when applicable. Apply your PRESIDENT
 9385 matcher to this new data. Does the accuracy improve? Compare 20 randomly-
 9386 selected pairs from this pattern and the one you designed in the previous part.
- 9387 5. Represent the dependency path $\mathbf{x}^{(i)}$ as a sequence of words and dependency arcs
 9388 of length M_i , ignoring the endpoints of the path. In example 1 of Table 17.2, the
 9389 dependency path is,

$$\mathbf{x}^{(1)} = (\xleftarrow[\text{NSUBJ}]{}, \text{traveled}, \xrightarrow[\text{OBL}]{}) \quad [17.32]$$

9390 If $x_m^{(i)}$ is a word, then let $\text{pos}(x_m^{(i)})$ be its part-of-speech, using the tagset defined in
 9391 chapter 8.

We can define the following kernel function over pairs of dependency paths (Bunescu
 and Mooney, 2005):

$$\kappa(\mathbf{x}^{(i)}, \mathbf{x}^{(j)}) = \begin{cases} 0, & M_i \neq M_j \\ \prod_{m=1}^{M_i} c(x_m^{(i)}, x_m^{(j)}), & M_i = M_j \end{cases}$$

$$c(x_m^{(i)}, x_m^{(j)}) = \begin{cases} 2, & x_m^{(i)} = x_m^{(j)} \\ 1, & x_m^{(i)} \text{ and } x_m^{(j)} \text{ are words and } \text{pos}(x_m^{(i)}) = \text{pos}(x_m^{(j)}) \\ 0, & \text{otherwise.} \end{cases}$$

9392 Using this kernel function, compute the kernel similarities of example 1 from Ta-
 9393 ble 17.2 with the other five examples.

6. Continuing from the previous problem, suppose that the instances have the follow-
 ing labels:

$$y_2 = 1, y_3 = -1, y_4 = -1, y_5 = 1, y_6 = 1 \quad [17.33]$$

9394 Identify the conditions for α and b under which $\hat{y}_1 = 1$. Remember the constraint
 9395 that $\alpha_i \geq 0$ for all i .

9396 Chapter 18

9397 Machine translation

9398 Machine translation (MT) is one of the “holy grail” problems in artificial intelligence,
9399 with the potential to transform society by facilitating communication between people
9400 anywhere in the world. As a result, MT has received significant attention and funding
9401 since the early 1950s. However, it has proved remarkably challenging, and while there
9402 has been substantial progress towards usable MT systems — especially for high-resource
9403 language pairs like English-French — we are still far from translation systems that match
9404 the nuance and depth of human translations.

9405 18.1 Machine translation as a task

9406 Machine translation can be formulated as an optimization problem:

$$\hat{\mathbf{w}}^{(t)} = \underset{\mathbf{w}^{(t)}}{\operatorname{argmax}} \Psi(\mathbf{w}^{(s)}, \mathbf{w}^{(t)}), \quad [18.1]$$

9407 where $\mathbf{w}^{(s)}$ is a sentence in a **source** language, $\mathbf{w}^{(t)}$ is a sentence in the **target language**,
9408 and Ψ is a scoring function. As usual, this formalism requires two components: a decod-
9409 ing algorithm for computing $\hat{\mathbf{w}}^{(t)}$, and a learning algorithm for estimating the parameters
9410 of the scoring function Ψ .

9411 Decoding is difficult for machine translation because of the huge space of possible
9412 translations. We have faced large label spaces before: for example, in sequence labeling,
9413 the set of possible label sequences is exponential in the length of the input. In these cases,
9414 it was possible to search the space quickly by introducing locality assumptions: for ex-
9415 ample, that each tag depends only on its predecessor, or that each production depends
9416 only on its parent. In machine translation, no such locality assumptions seem possible:
9417 human translators reword, reorder, and rearrange words; they replace single words with
9418 multi-word phrases, and vice versa. This flexibility means that in even relatively simple

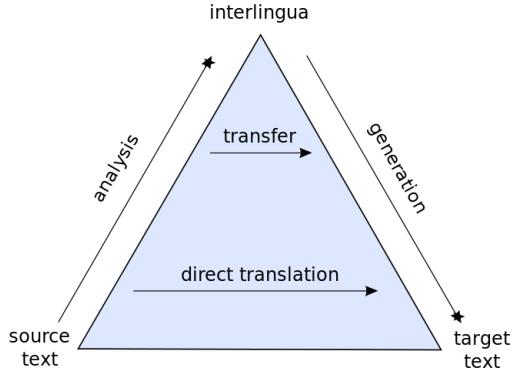


Figure 18.1: The Vauquois Pyramid http://commons.wikimedia.org/wiki/File:Direct_translation_and_transfer_translation_pyramind.svg

9419 translation models, decoding is NP-hard (Knight, 1999). Approaches for dealing with this
 9420 complexity are described in § 18.4.

Estimating translation models is difficult as well. Labeled translation data usually comes in the form parallel sentences, e.g.,

$$\begin{aligned} w^{(s)} &= A \text{ } Vinay \text{ } le \text{ } gusta \text{ } las \text{ } manzanas. \\ w^{(t)} &= Vinay \text{ likes apples.} \end{aligned}$$

9421 A useful feature function would note the translation pairs (*gusta, likes*), (*manzanas, apples*),
 9422 and even (*Vinay, Vinay*). But this word-to-word **alignment** is not given in the data. One
 9423 solution is to treat this alignment as a **latent variable**; this is the approach taken by clas-
 9424 sical **statistical machine translation** (SMT) systems, described in § 18.2. Another solution
 9425 is to model the relationship between $w^{(t)}$ and $w^{(s)}$ through a more complex and expres-
 9426 sive function; this is the approach taken by **neural machine translation** (NMT) systems,
 9427 described in § 18.3.

9428 The **Vauquois Pyramid** is a theory of how translation should be done. At the lowest
 9429 level, the translation system operates on individual words, but the horizontal distance
 9430 at this level is large, because languages express ideas differently. If we can move up the
 9431 triangle to syntactic structure, the distance for translation is reduced; we then need only
 9432 produce target-language text from the syntactic representation, which can be as simple
 9433 as reading off a tree. Further up the triangle lies semantics; translating between semantic
 9434 representations should be easier still, but mapping between semantics and surface text is
 9435 a difficult, unsolved problem. At the top of the triangle is **interlingua**, a semantic repre-
 9436 sentation that is so generic that it is identical across all human languages. Philosophers

	Adequate?	Fluent?
<i>To Vinay it like Python</i>	yes	no
<i>Vinay debugs memory leaks</i>	no	yes
<i>Vinay likes Python</i>	yes	yes

Table 18.1: Adequacy and fluency for translations of the Spanish sentence *A Vinay le gusta Python*.

debate whether such a thing as interlingua is really possible (Derrida, 1985). While the first-order logic representations discussed in chapter 12 might be considered to be language independent, it is built on an inventory of relations that is suspiciously similar to a subset of English words (Nirenburg and Wilks, 2001). Nonetheless, the idea of linking translation and semantic understanding may still be a promising path, if the resulting translations better preserve the meaning of the original text.

18.1.1 Evaluating translations

There are two main criteria for a translation, summarized in Table 18.1.

- **Adequacy:** The translation $w^{(t)}$ should adequately reflect the linguistic content of $w^{(s)}$. For example, if $w^{(s)} = A Vinay le gusta Python$, the gloss¹ $w^{(t)} = To Vinay it like Python$ is considered adequate becomes it contains all the relevant content. The output $w^{(t)} = Vinay debugs memory leaks$ is not adequate.
- **Fluency:** The translation $w^{(t)}$ should read like fluent text in the target language. By this criterion, the gloss $w^{(t)} = To Vinay it like Python$ will score poorly, and $w^{(t)} = Vinay debugs memory leaks$ will be preferred.

Automated evaluations of machine translations typically merge both of these criteria, by comparing the system translation with one or more **reference translations**, produced by professional human translators. The most popular quantitative metric is **BLEU** (bilingual evaluation understudy; Papineni et al., 2002), which is based on n -gram precision: what fraction of n -grams in the system translation appear in the reference? Specifically, for each n -gram length, the precision is defined as,

$$p_n = \frac{\text{number of } n\text{-grams appearing in both reference and hypothesis translations}}{\text{number of } n\text{-grams appearing in the hypothesis translation}}. \quad [18.2]$$

The n -gram precisions for three hypothesis translations are shown in Figure 18.2.

¹A gloss is a word-for-word translation.

	Translation	p_1	p_2	p_3	p_4	BP	BLEU
<i>Reference</i>	<i>Vinay likes programming in Python</i>						
<i>Sys1</i>	<i>To Vinay it like to program Python</i>	$\frac{2}{7}$	0	0	0	1	.21
<i>Sys2</i>	<i>Vinay likes Python</i>	$\frac{3}{3}$	$\frac{1}{2}$	0	0	.51	.33
<i>Sys3</i>	<i>Vinay likes programming in his pajamas</i>	$\frac{4}{6}$	$\frac{3}{5}$	$\frac{2}{4}$	$\frac{1}{3}$	1	.76

Figure 18.2: A reference translation and three system outputs. For each output, p_n indicates the precision at each n -gram, and BP indicates the brevity penalty.

9459 The BLEU score is then based on the average, $\exp \frac{1}{N} \sum_{n=1}^N \log p_n$. Two modifications
 9460 of Equation 18.2 are necessary: (1) to avoid computing $\log 0$, all precisions are smoothed
 9461 to ensure that they are positive; (2) each n -gram in the source can be used at most once,
 9462 so that *to to to to to* does not achieve $p_1 = 1$ against the reference *to be or not to be*.
 9463 Furthermore, precision-based metrics are biased in favor of short translations, which can
 9464 achieve high scores by minimizing the denominator in [18.2]. To avoid this issue, a **brevity**
 9465 **penalty** is applied to translations that are shorter than the reference. This penalty is indi-
 9466 cated as “BP” in Figure 18.2.

9467 Automated metrics like BLEU have been validated by correlation with human judg-
 9468 ments of translation quality. Nonetheless, it is not difficult to construct examples in which
 9469 the BLEU score is high, yet the translation is disfluent or carries a completely different
 9470 meaning from the original. To give just one example, consider the problem of translating
 9471 pronouns. Because pronouns refer to specific entities, a single incorrect pronoun can obliti-
 9472 onate the semantics of the original sentence. Existing state-of-the-art systems generally
 9473 do not attempt the reasoning necessary to correctly resolve pronominal anaphora (Hard-
 9474 meier, 2012). Despite the importance of pronouns for semantics, they have a marginal
 9475 impact on BLEU, which may help to explain why existing systems do not make a greater
 9476 effort to translate them correctly.

9477 **Fairness and bias** The problem of pronoun translation intersects with issues of fairness
 9478 and bias. In many languages, such as Turkish, the third person singular pronoun is gender
 9479 neutral. Today’s state-of-the-art systems produce the following Turkish-English transla-
 9480 tions (Caliskan et al., 2017):

- 9481 (18.1) *O bir doktor.*
 He is a doctor.
 9482 (18.2) *O bir hemşire.*
 She is a nurse.

9483 The same problem arises for other professions that have stereotypical genders, such as
9484 engineers, soldiers, and teachers, and for other languages that have gender-neutral pro-
9485 nouns. This bias was not directly programmed into the translation model; it arises from
9486 statistical tendencies in existing datasets. This highlights a general problem with data-
9487 driven approaches, which can perpetuate biases that negatively impact disadvantaged
9488 groups. Worse, machine learning can *amplify* biases in data (Bolukbasi et al., 2016): if a
9489 dataset has even a slight tendency towards men as doctors, the resulting translation model
9490 may produce translations in which doctors are always *he*, and nurses are always *she*.

9491 **Other metrics** A range of other automated metrics have been proposed for machine
9492 translation. One potential weakness of BLEU is that it only measures precision; METEOR
9493 is a weighted *F*-MEASURE, which is a combination of recall and precision (see § 4.4.1).
9494 **Translation Error Rate (TER)** computes the string **edit distance** (see § 9.1.4.1) between the
9495 reference and the hypothesis (Snover et al., 2006). For language pairs like English and
9496 Japanese, there are substantial differences in word order, and word order errors are not
9497 sufficiently captured by *n*-gram based metrics. The **RIBES** metric applies rank correla-
9498 tion to measure the similarity in word order between the system and reference transla-
9499 tions (Isozaki et al., 2010).

9500 18.1.2 Data

9501 Data-driven approaches to machine translation rely primarily on **parallel corpora**: sentence-
9502 level translations. Early work focused on government records, in which fine-grained offi-
9503 cial translations are often required. For example, the IBM translation systems were based
9504 on the proceedings of the Canadian Parliament, called **Hansards**, which are recorded in
9505 English and French (Brown et al., 1990). The growth of the European Union led to the
9506 development of the **EuroParl corpus**, which spans 21 European languages (Koehn, 2005).
9507 While these datasets helped to launch the field of machine translation, they are restricted
9508 to narrow domains and a formal speaking style, limiting their applicability to other types
9509 of text. As more resources are committed to machine translation, new translation datasets
9510 have been commissioned. This has broadened the scope of available data to news,² movie
9511 subtitles,³ social media (Ling et al., 2013), dialogues (Fordyce, 2007), TED talks (Paul et al.,
9512 2010), and scientific research articles (Nakazawa et al., 2016).

9513 Despite this growing set of resources, the main bottleneck in machine translation data
9514 is the need for parallel corpora that are aligned at the sentence level. Many languages have
9515 sizable parallel corpora with some high-resource language, but not with each other. The
9516 high-resource language can then be used as a “pivot” or “bridge” (Boitet, 1988; Utiyama

²https://catalog.ldc.upenn.edu/LDC2010T10_translation-task.html ³<http://www.statmt.org/wmt15/>

³<http://opus.nlpl.eu/>

and Isahara, 2007): for example, De Gispert and Marino (2006) use Spanish as a bridge for translation between Catalan and English. For most of the 6000 languages spoken today, the only source of translation data remains the Judeo-Christian Bible (Resnik et al., 1999). While relatively small, at less than a million tokens, the Bible has been translated into more than 2000 languages, far outpacing any other corpus. Some research has explored the possibility of automatically identifying parallel sentence pairs from unaligned parallel texts, such as web pages and Wikipedia articles (Kilgarriff and Grefenstette, 2003; Resnik and Smith, 2003; Adafre and De Rijke, 2006). Another approach is to create large parallel corpora through crowdsourcing (Zaidan and Callison-Burch, 2011).

18.2 Statistical machine translation

The previous section introduced adequacy and fluency as the two main criteria for machine translation. A natural modeling approach is to represent them with separate scores,

$$\Psi(\mathbf{w}^{(s)}, \mathbf{w}^{(t)}) = \Psi_A(\mathbf{w}^{(s)}, \mathbf{w}^{(t)}) + \Psi_F(\mathbf{w}^{(t)}). \quad [18.3]$$

The fluency score Ψ_F need not even consider the source sentence; it only judges $\mathbf{w}^{(t)}$ on whether it is fluent in the target language. This decomposition is advantageous because it makes it possible to estimate the two scoring functions on separate data. While the adequacy model must be estimated from aligned sentences — which are relatively expensive and rare — the fluency model can be estimated from monolingual text in the target language. Large monolingual corpora are now available in many languages, thanks to resources such as Wikipedia.

An elegant justification of the decomposition in Equation 18.3 is provided by the **noisy channel model**, in which each scoring function is a log probability:

$$\Psi_A(\mathbf{w}^{(s)}, \mathbf{w}^{(t)}) \triangleq \log p_{S|T}(\mathbf{w}^{(s)} | \mathbf{w}^{(t)}) \quad [18.4]$$

$$\Psi_F(\mathbf{w}^{(t)}) \triangleq \log p_T(\mathbf{w}^{(t)}) \quad [18.5]$$

$$\Psi(\mathbf{w}^{(s)}, \mathbf{w}^{(t)}) = \log p_{S|T}(\mathbf{w}^{(s)} | \mathbf{w}^{(t)}) + \log p_T(\mathbf{w}^{(t)}) = \log p_{S,T}(\mathbf{w}^{(s)}, \mathbf{w}^{(t)}). \quad [18.6]$$

By setting the scoring functions equal to the logarithms of the prior and likelihood, their sum is equal to $\log p_{S,T}$, which is the logarithm of the joint probability of the source and target. The sentence $\hat{\mathbf{w}}^{(t)}$ that maximizes this joint probability is also the maximizer of the conditional probability $p_{T|S}$, making it the most likely target language sentence, conditioned on the source.

The noisy channel model can be justified by a generative story. The target text is originally generated from a probability model p_T . It is then encoded in a “noisy channel” $p_{S|T}$, which converts it to a string in the source language. In decoding, we apply Bayes’ rule to recover the string $\mathbf{w}^{(t)}$ that is maximally likely under the conditional probability

	<i>A</i>	<i>Vinay</i>	<i>le</i>	<i>gusta</i>	<i>python</i>
<i>Vinay</i>		■			
<i>likes</i>			■	■	
<i>python</i>					■

Figure 18.3: An example word-to-word alignment

9545 $p_{T|S}$. Under this interpretation, the target probability p_T is just a language model, and
 9546 can be estimated using any of the techniques from chapter 6. The only remaining learning
 9547 problem is to estimate the translation model $p_{S|T}$.

9548 18.2.1 Statistical translation modeling

9549 The simplest decomposition of the translation model is word-to-word: each word in the
 9550 source should be aligned to a word in the translation. This approach presupposes an
 9551 **alignment** $\mathcal{A}(\mathbf{w}^{(s)}, \mathbf{w}^{(t)})$, which contains a list of pairs of source and target tokens. For
 9552 example, given $\mathbf{w}^{(s)} = A\ Vinay\ le\ gusta\ Python$ and $\mathbf{w}^{(t)} = Vinay\ likes\ Python$, one possible
 9553 word-to-word alignment is,

$$\mathcal{A}(\mathbf{w}^{(s)}, \mathbf{w}^{(t)}) = \{(A, \emptyset), (Vinay, Vinay), (le, likes), (gusta, likes), (Python, Python)\}. \quad [18.7]$$

9554 This alignment is shown in Figure 18.3. Another, less promising, alignment is:

$$\mathcal{A}(\mathbf{w}^{(s)}, \mathbf{w}^{(t)}) = \{(A, Vinay), (Vinay, likes), (le, Python), (gusta, \emptyset), (Python, \emptyset)\}. \quad [18.8]$$

9555 Each alignment contains exactly one tuple for each word in the *source*, which serves to
 9556 explain how the source word could be translated from the target, as required by the trans-
 9557 lation probability $p_{S|T}$. If no appropriate word in the target can be identified for a source
 9558 word, it is aligned to \emptyset — as is the case for the Spanish function word *a* in the example,
 9559 which glosses to the English word *to*. Words in the target can align with multiple words
 9560 in the source, so that the target word *likes* can align to both *le* and *gusta* in the source.

The joint probability of the alignment and the translation can be defined conveniently

as,

$$p(\mathbf{w}^{(s)}, \mathcal{A} | \mathbf{w}^{(t)}) = \prod_{m=1}^{M^{(s)}} p(w_m^{(s)}, a_m | w_{a_m}^{(t)}, m, M^{(s)}, M^{(t)}) \quad [18.9]$$

$$= \prod_{m=1}^{M^{(s)}} p(a_m | m, M^{(s)}, M^{(t)}) \times p(w_m^{(s)} | w_{a_m}^{(t)}). \quad [18.10]$$

This probability model makes two key assumptions:

- The alignment probability factors across tokens,

$$p(\mathcal{A} | \mathbf{w}^{(s)}, \mathbf{w}^{(t)}) = \prod_{m=1}^{M^{(s)}} p(a_m | m, M^{(s)}, M^{(t)}). \quad [18.11]$$

This means that each alignment decision is independent of the others, and depends only on the index m , and the sentence lengths $M^{(s)}$ and $M^{(t)}$.

- The translation probability also factors across tokens,

$$p(\mathbf{w}^{(s)} | \mathbf{w}^{(t)}, \mathcal{A}) = \prod_{m=1}^{M^{(s)}} p(w_m^{(s)} | w_{a_m}^{(t)}), \quad [18.12]$$

so that each word in $\mathbf{w}^{(s)}$ depends only on its aligned word in $\mathbf{w}^{(t)}$. This means that translation is word-to-word, ignoring context. The hope is that the target language model $p(\mathbf{w}^{(t)})$ will correct any disfluencies that arise from word-to-word translation.

To translate with such a model, we could sum or max over all possible alignments,

$$p(\mathbf{w}^{(s)}, \mathbf{w}^{(t)}) = \sum_{\mathcal{A}} p(\mathbf{w}^{(s)}, \mathbf{w}^{(t)}, \mathcal{A}) \quad [18.13]$$

$$= p(\mathbf{w}^{(t)}) \sum_{\mathcal{A}} p(\mathcal{A}) \times p(\mathbf{w}^{(s)} | \mathbf{w}^{(t)}, \mathcal{A}) \quad [18.14]$$

$$\geq p(\mathbf{w}^{(t)}) \max_{\mathcal{A}} p(\mathcal{A}) \times p(\mathbf{w}^{(s)} | \mathbf{w}^{(t)}, \mathcal{A}). \quad [18.15]$$

The term $p(\mathcal{A})$ defines the prior probability over alignments. A series of alignment models with increasingly relaxed independence assumptions was developed by researchers at IBM in the 1980s and 1990s, known as IBM Models 1-6 (Och and Ney, 2003). IBM Model 1 makes the strongest independence assumption:

$$p(a_m | m, M^{(s)}, M^{(t)}) = \frac{1}{M^{(t)}}. \quad [18.16]$$

9569 In this model, every alignment is equally likely. This is almost surely wrong, but it re-
 9570 sults in a convex learning objective, yielding a good initialization for the more complex
 9571 alignment models (Brown et al., 1993; Koehn, 2009).

9572 18.2.2 Estimation

9573 Let us define the parameter $\theta_{u \rightarrow v}$ as the probability of translating target word u to source
 9574 word v . If word-to-word alignments were annotated, these probabilities could be com-
 9575 puted from relative frequencies,

$$\hat{\theta}_{u \rightarrow v} = \frac{\text{count}(u, v)}{\text{count}(u)}, \quad [18.17]$$

9576 where $\text{count}(u, v)$ is the count of instances in which word v was aligned to word u in
 9577 the training set, and $\text{count}(u)$ is the total count of the target word u . The smoothing
 9578 techniques mentioned in chapter 6 can help to reduce the variance of these probability
 9579 estimates.

9580 Conversely, if we had an accurate translation model, we could estimate the likelihood
 9581 of each alignment decision,

$$q_m(a_m | \mathbf{w}^{(s)}, \mathbf{w}^{(t)}) \propto p(a_m | m, M^{(s)}, M^{(t)}) \times p(w_m^{(s)} | w_{a_m}^{(t)}), \quad [18.18]$$

where $q_m(a_m | \mathbf{w}^{(s)}, \mathbf{w}^{(t)})$ is a measure of our confidence in aligning source word $w_m^{(s)}$
 to target word $w_{a_m}^{(t)}$. The relative frequencies could then be computed from the *expected
 counts*,

$$\hat{\theta}_{u \rightarrow v} = \frac{E_q[\text{count}(u, v)]}{\text{count}(u)} \quad [18.19]$$

$$E_q[\text{count}(u, v)] = \sum_m q_m(a_m | \mathbf{w}^{(s)}, \mathbf{w}^{(t)}) \delta(w_m^{(s)} = v) \delta(w_{a_m}^{(t)} = u). \quad [18.20]$$

9582 The **expectation-maximization** (EM) algorithm proceeds by iteratively updating q_m
 9583 and $\hat{\Theta}$. The algorithm is described in general form in chapter 5. For statistical machine
 9584 translation, the steps of the algorithm are:

- 9585 1. **E-step:** Update beliefs about word alignment using Equation 18.18.
- 9586 2. **M-step:** Update the translation model using Equations 18.19 and 18.20.

9587 As discussed in chapter 5, the expectation maximization algorithm is guaranteed to con-
 9588 verge, but not to a global optimum. However, for IBM Model 1, it can be shown that EM
 9589 optimizes a convex objective, and global optimality is guaranteed. For this reason, IBM
 9590 Model 1 is often used as an initialization for more complex alignment models. For more
 9591 detail, see Koehn (2009).

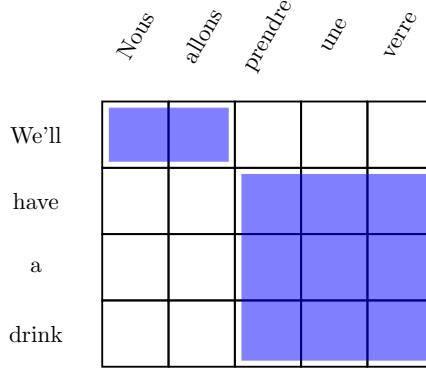


Figure 18.4: A phrase-based alignment between French and English, corresponding to example (18.3)

18.2.3 Phrase-based translation

Real translations are not word-to-word substitutions. One reason is that many multiword expressions are not translated literally, as shown in this example from French:

- (18.3) *Nous allons prendre un verre*
 We will take a glass
 We'll have a drink

The line *we will take a glass* is the word-for-word gloss of the French sentence; the translation *we'll have a drink* is shown on the third line. Such examples are difficult for word-to-

word translation models, since they require translating *prendre* to *have* and *verre* to *drink*.
 These translations are only correct in the context of these specific phrases.

Phrase-based translation generalizes on word-based models by building translation tables and alignments between multiword spans. (These “phrases” are not necessarily syntactic constituents like the noun phrases and verb phrases described in chapters 9 and 10.) The generalization from word-based translation is surprisingly straightforward: the translation tables can now condition on multi-word units, and can assign probabilities to multi-word units; alignments are mappings from spans to spans, $((i, j), (k, \ell))$, so that

$$p(\mathbf{w}^{(s)} | \mathbf{w}^{(t)}, \mathcal{A}) = \prod_{((i, j), (k, \ell)) \in \mathcal{A}} p_{w^{(s)} | w^{(t)}}(\{w_{i+1}^{(s)}, w_{i+2}^{(s)}, \dots, w_j^{(s)}\} | \{w_{k+1}^{(t)}, w_{k+2}^{(t)}, \dots, w_\ell^{(t)}\}). \quad [18.21]$$

The phrase alignment $((i, j), (k, \ell))$ indicates that the span $\mathbf{w}_{i+1:j}^{(s)}$ is the translation of the span $\mathbf{w}_{k+1:\ell}^{(t)}$. An example phrasal alignment is shown in Figure 18.4. Note that the align-

ment set \mathcal{A} is required to cover all of the tokens in the source, just as in word-based translation. The probability model $p_{w^{(s)}|w^{(t)}}$ must now include translations for all phrase pairs, which can be learned from expectation-maximization just as in word-based statistical machine translation.

18.2.4 *Syntax-based translation

The Vauquois Pyramid (Figure 18.1) suggests that translation might be easier if we take a higher-level view. One possibility is to incorporate the syntactic structure of the source, the target, or both. This is particularly promising for language pairs that consistent syntactic differences. For example, English adjectives almost always precede the nouns that they modify, while in Romance languages such as French and Spanish, the adjective often follows the noun: thus, *angry fish* would translate to *pez (fish) enojado (angry)* in Spanish. In word-to-word translation, these reorderings cause the alignment model to be overly permissive. It is not that the order of *any* pair of English words can be reversed when translating into Spanish, but only adjectives and nouns within a noun phrase. Similar issues arise when translating between verb-final languages such as Japanese (in which verbs usually follow the subject and object), verb-initial languages like Tagalog and classical Arabic, and verb-medial languages such as English.

An elegant solution is to link parsing and translation in a **synchronous context-free grammar** (SCFG; Chiang, 2007).⁴ An SCFG is a set of productions of the form $X \rightarrow (\alpha, \beta, \sim)$, where X is a non-terminal, α and β are sequences of terminals or non-terminals, and \sim is a one-to-one alignment of items in α with items in β . To handle the English-Spanish adjective-noun ordering, an SCFG would include productions such as,

$$\text{NP} \rightarrow (\text{DET}_1 \text{NN}_2 \text{JJ}_3, \quad \text{DET}_1 \text{JJ}_3 \text{NN}_2), \quad [18.22]$$

with subscripts indicating the alignment between the Spanish (left) and English (right) parts of the right-hand side. Terminal productions yield translation pairs,

$$\text{JJ} \rightarrow (\text{enojado}_1, \text{angry}_1). \quad [18.23]$$

A synchronous derivation begins with the start symbol S , and derives a pair of sequences of terminal symbols.

Given an SCFG in which each production yields at most two symbols in each language (Chomsky Normal Form; see § 9.2.1.2), a sentence can be parsed using only the CKY algorithm (chapter 10). The resulting derivation also includes productions in the other language, all the way down to the surface form. Therefore, SCFGs make translation very similar to parsing. In a weighted SCFG, the log probability $\log p_{S|T}$ can be computed from

⁴Key earlier work includes syntax-driven transduction (Lewis II and Stearns, 1968) and stochastic inversion transduction grammars (Wu, 1997).

the sum of the log-probabilities of the productions. However, combining SCFGs with a target language model is computationally expensive, necessitating approximate search algorithms (Huang and Chiang, 2007).

Synchronous context-free grammars are an example of **tree-to-tree translation**, because they model the syntactic structure of both the target and source language. In **string-to-tree translation**, string elements are translated into constituent tree fragments, which are then assembled into a translation (Yamada and Knight, 2001; Galley et al., 2004); in **tree-to-string translation**, the source side is parsed, and then transformed into a string on the target side (Liu et al., 2006). A key question for syntax-based translation is the extent to which we phrasal constituents align across translations (Fox, 2002), because this governs the extent to which we can rely on monolingual parsers and treebanks. For more on syntax-based machine translation, see the monograph by Williams et al. (2016).

18.3 Neural machine translation

Neural network models for machine translation are based on the **encoder-decoder** architecture (Cho et al., 2014). The encoder network converts the source language sentence into a vector or matrix representation; the decoder network then converts the encoding into a sentence in the target language.

$$\mathbf{z} = \text{ENCODE}(\mathbf{w}^{(s)}) \quad [18.24]$$

$$\mathbf{w}^{(t)} \mid \mathbf{w}^{(s)} \sim \text{DECODE}(\mathbf{z}), \quad [18.25]$$

where the second line means that the function $\text{DECODE}(\mathbf{z})$ defines the conditional probability $p(\mathbf{w}^{(t)} \mid \mathbf{w}^{(s)})$.

The decoder is typically a recurrent neural network, which generates the target language sentence one word at a time, while recurrently updating a hidden state. The encoder and decoder networks are trained end-to-end from parallel sentences. If the output layer of the decoder is a logistic function, then the entire architecture can be trained to maximize the conditional log-likelihood,

$$\log p(\mathbf{w}^{(t)} \mid \mathbf{w}^{(s)}) = \sum_{m=1}^{M(t)} p(w_m^{(t)} \mid \mathbf{w}_{1:m-1}^{(t)}, \mathbf{z}) \quad [18.26]$$

$$p(w_m^{(t)} \mid \mathbf{w}_{1:m-1}^{(t)}, \mathbf{w}^{(s)}) \propto \exp\left(\boldsymbol{\beta}_{w_m^{(t)}} \cdot \mathbf{h}_{m-1}^{(t)}\right) \quad [18.27]$$

where the hidden state $\mathbf{h}_{m-1}^{(t)}$ is a recurrent function of the previously generated text $\mathbf{w}_{1:m-1}^{(t)}$ and the encoding \mathbf{z} . The second line is equivalent to writing,

$$w_m^{(t)} \mid \mathbf{w}_{1:m-1}^{(t)}, \mathbf{w}^{(s)} \sim \text{SoftMax}\left(\boldsymbol{\beta} \cdot \mathbf{h}_{m-1}^{(t)}\right), \quad [18.28]$$

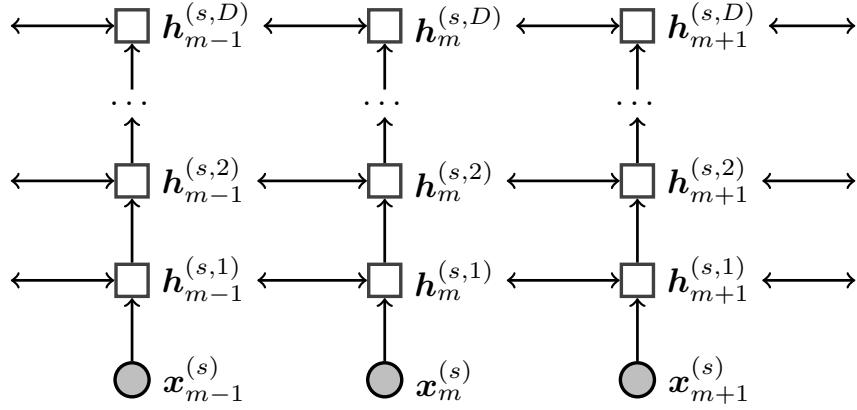


Figure 18.5: A deep bidirectional LSTM encoder

9649 where $\beta \in \mathbb{R}^{(V^{(t)} \times K)}$ is the matrix of output word vectors for the $V^{(t)}$ words in the target
 9650 language vocabulary.

The simplest encoder-decoder architecture is the **sequence-to-sequence** model (Sutskever et al., 2014). In this model, the encoder is set to the final hidden state of a **long short-term memory (LSTM)** (see § 6.3.3) on the source sentence:

$$\mathbf{h}_m^{(s)} = \text{LSTM}(\mathbf{x}_m^{(s)}, \mathbf{h}_{m-1}^{(s)}) \quad [18.29]$$

$$\mathbf{z} \triangleq \mathbf{h}_{M^{(s)}}^{(s)}, \quad [18.30]$$

where $\mathbf{x}_m^{(s)}$ is the embedding of source language word $w_m^{(s)}$. The encoding then provides the initial hidden state for the decoder LSTM:

$$\mathbf{h}_0^{(t)} = \mathbf{z} \quad [18.31]$$

$$\mathbf{h}_m^{(t)} = \text{LSTM}(\mathbf{x}_m^{(t)}, \mathbf{h}_{m-1}^{(t)}), \quad [18.32]$$

9651 where $\mathbf{x}_m^{(t)}$ is the embedding of the target language word $w_m^{(t)}$.

9652 Sequence-to-sequence translation is nothing more than wiring together two LSTMs:
 9653 one to read the source, and another to generate the target. To make the model work well,
 9654 some additional tweaks are needed:

- 9655 • Most notably, the model works much better if the source sentence is reversed, reading
 9656 from the end of the sentence back to the beginning. In this way, the words at the
 9657 beginning of the source have the greatest impact on the encoding \mathbf{z} , and therefore

9658 impact the words at the beginning of the target sentence. Later work on more ad-
 9659 vanced encoding models, such as **neural attention** (see § 18.3.1), has eliminated the
 9660 need for reversing the source sentence.

- The encoder and decoder can be implemented as **deep LSTMs**, with multiple layers of hidden states. As shown in Figure 18.5, each hidden state $\mathbf{h}_m^{(s,i)}$ at layer i is treated as the input to an LSTM at layer $i + 1$:

$$\mathbf{h}_m^{(s,1)} = \text{LSTM}(\mathbf{x}_m^{(s)}, \mathbf{h}_{m-1}^{(s)}) \quad [18.33]$$

$$\mathbf{h}_m^{(s,i+1)} = \text{LSTM}(\mathbf{h}_m^{(s,i)}, \mathbf{h}_{m-1}^{(s)}), \quad \forall i \geq 1. \quad [18.34]$$

9661 The original work on sequence-to-sequence translation used four layers; in 2016,
 9662 Google’s commercial machine translation system used eight layers (Wu et al., 2016).⁵

- 9663 • Significant improvements can be obtained by creating an **ensemble** of translation
 9664 models, each trained from a different random initialization. For an ensemble of size
 9665 N , the per-token decoding probability is set equal to,

$$p(w^{(t)} | \mathbf{z}, \mathbf{w}_{1:m-1}^{(t)}) = \frac{1}{N} \sum_{i=1}^N p_i(w^{(t)} | \mathbf{z}, \mathbf{w}_{1:m-1}^{(t)}), \quad [18.35]$$

9666 where p_i is the decoding probability for model i . Each translation model in the
 9667 ensemble includes its own encoder and decoder networks.

- 9668 • The original sequence-to-sequence model used a fairly standard training setup: stochas-
 9669 tic gradient descent with an exponentially decreasing learning rate after the first five
 9670 epochs; mini-batches of 128 sentences, chosen to have similar length so that each
 9671 sentence on the batch will take roughly the same amount of time to process; gradi-
 9672 ent clipping (see § 3.3.4) to ensure that the norm of the gradient never exceeds some
 9673 predefined value.

9674 18.3.1 Neural attention

9675 The sequence-to-sequence model discussed in the previous section was a radical depart-
 9676 ure from statistical machine translation, in which each word or phrase in the target lan-
 9677 guage is conditioned on a single word or phrase in the source language. Both approaches
 9678 have advantages. Statistical translation leverages the idea of compositionality — transla-
 9679 tions of large units should be based on the translations of their component parts — and
 9680 this seems crucial if we are to scale translation to longer units of text. But the translation
 9681 of each word or phrase often depends on the larger context, and encoder-decoder models
 9682 capture this context at the sentence level.

⁵Google reports that this system took six days to train for English-French translation, using 96 NVIDIA K80 GPUs, which would have cost roughly half a million dollars at the time.

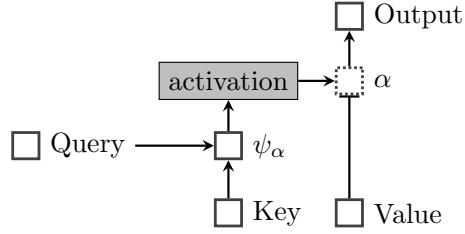


Figure 18.6: A general view of neural attention. The dotted box indicates that each $\alpha_{m \rightarrow n}$ can be viewed as a **gate** on value n .

Is it possible for translation to be both contextualized and compositional? One approach is to augment neural translation with an **attention mechanism**. The idea of neural attention was described in § 17.5, but its application to translation bears further discussion. In general, attention can be thought of as using a query to select from a memory of key-value pairs. However, the query, keys, and values are all vectors, and the entire operation is differentiable. For each key n in the memory, we compute a score $\psi_\alpha(m, n)$ with respect to the query m . That score is a function of the compatibility of the key and the query, and can be computed using a small feedforward neural network. The vector of scores is passed through an activation function, such as softmax. The output of this activation function is a vector of non-negative numbers $[\alpha_{m \rightarrow 1}, \alpha_{m \rightarrow 2}, \dots, \alpha_{m \rightarrow N}]^\top$, with length N equal to the size of the memory. Each value in the memory v_n is multiplied by the attention $\alpha_{m \rightarrow n}$; the sum of these scaled values is the output. This process is shown in Figure 18.6. In the extreme case that $\alpha_{m \rightarrow n} = 1$ and $\alpha_{m \rightarrow n'} = 0$ for all other n' , then the attention mechanism simply selects the value v_n from the memory.

Neural attention makes it possible to integrate alignment into the encoder-decoder architecture. Rather than encoding the entire source sentence into a fixed length vector z , it can be encoded into a matrix $Z \in \mathbb{R}^{K \times M^{(S)}}$, where K is the dimension of the hidden state, and $M^{(S)}$ is the number of tokens in the source input. Each column of Z represents the state of a recurrent neural network over the source sentence. These vectors are constructed from a **bidirectional LSTM** (see § 7.6), which can be a deep network as shown in Figure 18.5. These columns are both the keys and the values in the attention mechanism.

At each step m in decoding, the attentional state is computed by executing a query, which is equal to the state of the decoder, $h_m^{(t)}$. The resulting compatibility scores are,

$$\psi_\alpha(m, n) = v_\alpha \cdot \tanh(\Theta_\alpha[h_m^{(t)}; h_n^{(s)}]). \quad [18.36]$$

The function ψ is thus a two layer feedforward neural network, with weights v_α on the output layer, and weights Θ_α on the input layer. To convert these scores into attention weights, we apply an activation function, which can be vector-wise softmax or an

9707 element-wise sigmoid:

Softmax attention

$$\alpha_{m \rightarrow n} = \frac{\exp \psi_\alpha(m, n)}{\sum_{n'=1}^{M^{(s)}} \exp \psi_\alpha(m, n')} \quad [18.37]$$

Sigmoid attention

$$\alpha_{m \rightarrow n} = \sigma(\psi_\alpha(m, n)) \quad [18.38]$$

The attention α is then used to compute an **context vector** c_m by taking a weighted average over the columns of Z ,

$$c_m = \sum_{n=1}^{M^{(s)}} \alpha_{m \rightarrow n} z_n, \quad [18.39]$$

where $\alpha_{m \rightarrow n} \in [0, 1]$ is the amount of attention from word m of the target to word n of the source. The context vector can be incorporated into the decoder's word output probability model, by adding another layer to the decoder (Luong et al., 2015):

$$\tilde{h}_m^{(t)} = \tanh(\Theta_c[h_m^{(t)}; c_m]) \quad [18.40]$$

$$p(w_{m+1}^{(t)} | w_{1:m}^{(t)}, w^{(s)}) \propto \exp\left(\beta_{w_{m+1}^{(t)}} \cdot \tilde{h}_m^{(t)}\right). \quad [18.41]$$

9708 Here the decoder state $h_m^{(t)}$ is concatenated with the context vector, forming the input
 9709 to compute a final output vector $\tilde{h}_m^{(t)}$. The context vector can be incorporated into the
 9710 decoder recurrence in a similar manner (Bahdanau et al., 2014).

9711 **18.3.2 *Neural machine translation without recurrence**

In the encoder-decoder model, attention's “keys and values” are the hidden state representations in the encoder network, z , and the “queries” are state representations in the decoder network $h^{(t)}$. It is also possible to completely eliminate recurrence from neural translation, by applying **self-attention** (Lin et al., 2017; Kim et al., 2017) within the encoder and decoder, as in the **transformer architecture** (Vaswani et al., 2017). For level i , the basic equations of the encoder side of the transformer are:

$$z_m^{(i)} = \sum_{n=1}^{M^{(s)}} \alpha_{m \rightarrow n}^{(i)} (\Theta_v h_n^{(i-1)}) \quad [18.42]$$

$$h_m^{(i)} = \Theta_2 \text{ReLU}(\Theta_1 z_m^{(i)} + b_1) + b_2. \quad [18.43]$$

9712 For each token m at level i , we compute self-attention over the entire source sentence:
 9713 the keys, values, and queries are all projections of the vector $\mathbf{h}^{(i-1)}$. The attention scores
 9714 $\alpha_{m \rightarrow n}^{(i)}$ are computed using a scaled form of softmax attention,

$$\alpha_{m \rightarrow n} \propto \exp(\psi_\alpha(m, n)/M), \quad [18.44]$$

9715 where M is the length of the input. This encourages the attention to be more evenly
 9716 dispersed across the input. Self-attention is applied across multiple “heads”, each using
 9717 different projections of $\mathbf{h}^{(i-1)}$ to form the keys, values, and queries.

9718 The output of the self-attentional layer is the representation $\mathbf{z}_m^{(i)}$, which is then passed
 9719 through a two-layer feed-forward network, yielding the input to the next layer, $\mathbf{h}^{(i)}$. To
 9720 ensure that information about word order in the source is integrated into the model, the
 9721 encoder includes **positional encodings** of the index of each word in the source. These
 9722 encodings are vectors for each position $m \in \{1, 2, \dots, M\}$. The positional encodings are
 9723 concatenated with the word embeddings \mathbf{x}_m at the base layer of the model.⁶

9724 Convolutional neural networks (see § 3.4) have also been applied as encoders in neu-
 9725 ral machine translation. For each word $w_m^{(s)}$, a convolutional network computes a rep-
 9726 resentation $\mathbf{h}_m^{(s)}$ from the embeddings of the word and its neighbors. This procedure is
 9727 applied several times, creating a deep convolutional network. The recurrent decoder then
 9728 computes a set of attention weights over these convolutional representations, using the
 9729 decoder’s hidden state $\mathbf{h}^{(t)}$ as the queries. This attention vector is used to compute a
 9730 weighted average over the outputs of *another* convolutional neural network of the source,
 9731 yielding an averaged representation \mathbf{c}_m , which is then fed into the decoder. As with the
 9732 transformer, speed is the main advantage over recurrent encoding models; another sim-
 9733 ilarity is that word order information is approximated through the use of positional en-
 9734 codings. It seems likely that there are limitations to how well positional encodings can
 9735 account for word order and deeper linguistic structure. But for the moment, the com-
 9736 putational advantages of such approaches have put them on par with the best recurrent
 9737 translation models.⁷

9738 18.3.3 Out-of-vocabulary words

9739 Thus far, we have treated translation as a problem at the level of words or phrases. For
 9740 words that do not appear in the training data, all such models will struggle. There are
 9741 two main reasons for the presence of out-of-vocabulary (OOV) words:

⁶The transformer architecture relies on several additional tricks, including **layer normalization** (see § 3.3.4) and residual connections around the nonlinear activations (see § 3.2.2).

⁷A recent evaluation found that best performance was obtained by using a recurrent network for the decoder, and a transformer for the encoder (Chen et al., 2018). The transformer was also found to significantly outperform a convolutional neural network.

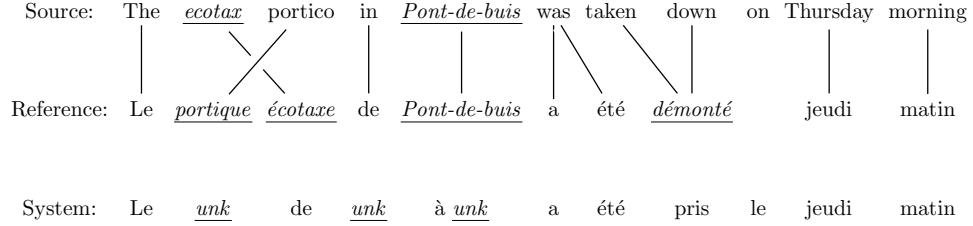


Figure 18.7: Translation with *unknown words*. The system outputs *unk* to indicate words that are outside its vocabulary. Figure adapted from Luong et al. (2015).

- New proper nouns, such as family names or organizations, are constantly arising — particularly in the news domain. The same is true, to a lesser extent, for technical terminology. This issue is shown in Figure 18.7.
- In many languages, words have complex internal structure, known as **morphology**. An example is German, which uses compounding to form nouns like *Abwasserbehandlungsanlage* (*sewage water treatment plant*; example from Sennrich et al. (2016)). While compounds could in principle be addressed by better tokenization (see § 8.4), other morphological processes involve more complex transformations of subword units.

Names and technical terms can be handled in a postprocessing step: after first identifying alignments between unknown words in the source and target, we can look up each aligned source word in a dictionary, and choose a replacement (Luong et al., 2015). If the word does not appear in the dictionary, it is likely to be a proper noun, and can be copied directly from the source to the target. This approach can also be integrated directly into the translation model, rather than applying it as a postprocessing step (Jean et al., 2015).

Words with complex internal structure can be handled by translating subword units rather than entire words. A popular technique for identifying subword units is **byte-pair encoding** (BPE; Gage, 1994; Sennrich et al., 2016). The initial vocabulary is defined as the set of characters used in the text. The most common character bigram is then merged into a new symbol, and the vocabulary is updated. The merging operation is applied repeatedly, until the vocabulary reaches some maximum size. For example, given the dictionary $\{fish, fished, want, wanted, bike, biked\}$, we would first merge $e+d$ into the subword unit ed , since this bigram appears in three words of the six words. Next, there are several bigrams that each appear in a pair of words: $f+i$, $i+s$, $s+h$, $w+a$, $a+n$, etc. These can be merged in any order, resulting in the segmentation, $\{fish, fish+ed, want, want+ed, bik+e, bik+ed\}$. At this point, there are no subword bigrams that appear more than once. In real data, merging is performed until the number of subword units reaches some predefined threshold,

9769 such as 10^4 .

9770 Each subword unit is treated as a token for translation, in both the encoder (source
 9771 side) and decoder (target side). BPE can be applied jointly to the union of the source and
 9772 target vocabularies, identifying subword units that appear in both languages. For lan-
 9773 guages that have different scripts, such as English and Russian, **transliteration** between
 9774 the scripts should be applied first.⁸

9775 18.4 Decoding

Given a trained translation model, the decoding task is:

$$\hat{\mathbf{w}}^{(t)} = \underset{\mathbf{w} \in \mathcal{V}^*}{\operatorname{argmax}} \Psi(\mathbf{w}, \mathbf{w}^{(s)}), \quad [18.45]$$

9776 where $\mathbf{w}^{(t)}$ is a sequence of tokens from the target vocabulary \mathcal{V} . It is not possible to
 9777 efficiently obtain exact solutions to the decoding problem, for even minimally effective
 9778 models in either statistical or neural machine translation. Today's state-of-the-art transla-
 9779 tion systems use **beam search** (see § 11.3.1.4), which is an incremental decoding algorithm
 9780 that maintains a small constant number of competitive hypotheses. Such greedy approxi-
 9781 mations are reasonably effective in practice, and this may be in part because the decoding
 9782 objective is only loosely correlated with measures of translation quality, so that exact op-
 9783 timization of [18.45] may not greatly improve the resulting translations.

Decoding in neural machine translation is somewhat simpler than in phrase-based
 statistical machine translation.⁹ The scoring function Ψ is defined,

$$\Psi(\mathbf{w}^{(t)}, \mathbf{w}^{(s)}) = \sum_{m=1}^{M^{(t)}} \psi(w_m^{(t)}; \mathbf{w}_{1:m-1}^{(t)}, \mathbf{z}) \quad [18.46]$$

$$\psi(w^{(t)}; \mathbf{w}_{1:m-1}^{(t)}, \mathbf{z}) = \beta_{w_m^{(t)}} \cdot \mathbf{h}_m^{(t)} - \log \sum_{w \in \mathcal{V}} \exp(\beta_w \cdot \mathbf{h}_m^{(t)}), \quad [18.47]$$

9784 where \mathbf{z} is the encoding of the source sentence $\mathbf{w}^{(s)}$, and $\mathbf{h}_m^{(t)}$ is a function of the encoding
 9785 \mathbf{z} and the decoding history $\mathbf{w}_{1:m-1}^{(t)}$. This formulation subsumes the attentional translation
 9786 model, where \mathbf{z} is a matrix encoding of the source.

Now consider the incremental decoding algorithm,

$$\hat{w}_m^{(t)} = \underset{w \in \mathcal{V}}{\operatorname{argmax}} \psi(w; \hat{\mathbf{w}}_{1:m-1}^{(t)}, \mathbf{z}), \quad m = 1, 2, \dots \quad [18.48]$$

⁸Transliteration is crucial for converting names and other foreign words between languages that do not share a single script, such as English and Japanese. It is typically approached using the finite-state methods discussed in chapter 9 (Knight and Graehl, 1998).

⁹For more on decoding in phrase-based statistical models, see Koehn (2009).

9787 This algorithm selects the best target language word at position m , assuming that it has
 9788 already generated the sequence $\hat{w}_{1:m-1}^{(t)}$. (Termination can be handled by augmenting
 9789 the vocabulary \mathcal{V} with a special end-of-sequence token, ■.) The incremental algorithm
 9790 is likely to produce a suboptimal solution to the optimization problem defined in Equa-
 9791 tion 18.45, because selecting the highest-scoring word at position m can set the decoder
 9792 on a “garden path,” in which there are no good choices at some later position $n > m$. We
 9793 might hope for some dynamic programming solution, as in sequence labeling (§ 7.3). But
 9794 the Viterbi algorithm and its relatives rely on a Markov decomposition of the objective
 9795 function into a sum of local scores: for example, scores can consider locally adjacent tags
 9796 (y_m, y_{m-1}), but not the entire tagging history $y_{1:m}$. This decomposition is not applicable
 9797 to recurrent neural networks, because the hidden state $h_m^{(t)}$ is impacted by the entire his-
 9798 tory $w_{1:m}^{(t)}$; this sensitivity to long-range context is precisely what makes recurrent neural
 9799 networks so effective.¹⁰ In fact, it can be shown that decoding from any recurrent neural
 9800 network is NP-complete (Siegelmann and Sontag, 1995; Chen et al., 2018).

9801 **Beam search** Beam search is a general technique for avoiding search errors when ex-
 9802 haustive search is impossible; it was first discussed in § 11.3.1.4. Beam search can be
 9803 seen as a variant of the incremental decoding algorithm sketched in Equation 18.48, but
 9804 at each step m , a set of K different hypotheses are kept on the beam. For each hypothesis
 9805 $k \in \{1, 2, \dots, K\}$, we compute both the current score $\sum_{m=1}^{M^{(t)}} \psi(w_{k,m}^{(t)}; w_{k,1:m-1}^{(t)}, z)$ as well as
 9806 the current hidden state $h_k^{(t)}$. At each step in the beam search, the K top-scoring children
 9807 of each hypothesis currently on the beam are “expanded”, and the beam is updated. For
 9808 a detailed description of beam search for RNN decoding, see Graves (2012).

9809 **Learning and search** Conventionally, the learning algorithm is trained to predict the
 9810 right token in the translation, conditioned on the translation history being correct. But
 9811 if decoding must be approximate, then we might do better by modifying the learning
 9812 algorithm to be robust to errors in the translation history. **Scheduled sampling** does this
 9813 by training on histories that sometimes come from the ground truth, and sometimes come
 9814 from the model’s own output (Bengio et al., 2015).¹¹ As training proceeds, the training
 9815 wheels come off: we increase the fraction of tokens that come from the model rather than
 9816 the ground truth. Another approach is to train on an objective that relates directly to beam
 9817 search performance (Wiseman et al., 2016). **Reinforcement learning** has also been applied
 9818 to decoding of RNN-based translation models, making it possible to directly optimize
 9819 translation metrics such as BLEU (Ranzato et al., 2016).

¹⁰Note that this problem does not impact RNN-based sequence labeling models (see § 7.6). This is because the tags produced by these models do not affect the recurrent state.

¹¹Scheduled sampling builds on earlier work on learning to search (Daumé III et al., 2009; Ross et al., 2011), which are also described in § 15.2.4.

18.5 Training towards the evaluation metric

In likelihood-based training, the objective is to maximize the probability of a parallel corpus. However, translations are not evaluated in terms of likelihood: metrics like BLEU consider only the correctness of a single output translation, and not the range of probabilities that the model assigns. It might therefore be better to train translation models to achieve the highest BLEU score possible — to the extent that we believe BLEU measures translation quality. Unfortunately, BLEU and related metrics are not friendly for optimization: they are discontinuous, non-differentiable functions of the parameters of the translation model.

Consider an error function $\Delta(\hat{\mathbf{w}}^{(t)}, \mathbf{w}^{(t)})$, which measures the discrepancy between the system translation $\hat{\mathbf{w}}^{(t)}$ and the reference translation $\mathbf{w}^{(t)}$; this function could be based on BLEU or any other metric on translation quality. One possible criterion would be to select the parameters θ that minimize the error of the system's preferred translation,

$$\hat{\mathbf{w}}^{(t)} = \operatorname{argmax}_{\mathbf{w}^{(t)}} \Psi(\mathbf{w}^{(t)}, \mathbf{w}^{(s)}; \theta) \quad [18.49]$$

$$\hat{\theta} = \operatorname{argmin}_{\theta} \Delta(\hat{\mathbf{w}}^{(t)}, \mathbf{w}^{(s)}) \quad [18.50]$$

However, identifying the top-scoring translation $\hat{\mathbf{w}}^{(t)}$ is usually intractable, as described in the previous section. In **minimum error-rate training (MERT)**, $\hat{\mathbf{w}}^{(t)}$ is selected from a set of candidate translations $\mathcal{Y}(\mathbf{w}^{(s)})$; this is typically a strict subset of all possible translations, so that it is only possible to optimize an approximation to the true error rate (Och and Ney, 2003).

A further issue is that the objective function in Equation 18.50 is discontinuous and non-differentiable, due to the argmax over translations: an infinitesimal change in the parameters θ could cause another translation to be selected, with a completely different error. To address this issue, we can instead minimize the **risk**, which is defined as the expected error rate,

$$R(\theta) = E_{\hat{\mathbf{w}}^{(t)} | \mathbf{w}^{(s)}; \theta} [\Delta(\hat{\mathbf{w}}^{(t)}, \mathbf{w}^{(t)})] \quad [18.51]$$

$$= \sum_{\hat{\mathbf{w}}^{(t)} \in \mathcal{Y}(\mathbf{w}^{(s)})} p(\hat{\mathbf{w}}^{(t)} | \mathbf{w}^{(s)}) \times \Delta(\hat{\mathbf{w}}^{(t)}, \mathbf{w}^{(t)}). \quad [18.52]$$

Minimum risk training minimizes the sum of $R(\theta)$ across all instances in the training set.

The risk can be generalized by exponentiating the translation probabilities,

$$\tilde{p}(\mathbf{w}^{(t)}; \theta, \alpha) \propto \left(p(\mathbf{w}^{(t)} | \mathbf{w}^{(s)}; \theta) \right)^\alpha \quad [18.53]$$

$$\tilde{R}(\theta) = \sum_{\hat{\mathbf{w}}^{(t)} \in \mathcal{Y}(\mathbf{w}^{(s)})} \tilde{p}(\hat{\mathbf{w}}^{(t)} | \mathbf{w}^{(s)}; \alpha, \theta) \times \Delta(\hat{\mathbf{w}}^{(t)}, \mathbf{w}^{(t)}) \quad [18.54]$$

where $\mathcal{Y}(\mathbf{w}^{(s)})$ is now the set of *all* possible translations for $\mathbf{w}^{(s)}$. Exponentiating the probabilities in this way is known as **annealing** (Smith and Eisner, 2006). When $\alpha = 1$, then $\tilde{R}(\boldsymbol{\theta}) = R(\boldsymbol{\theta})$; when $\alpha = \infty$, then $\tilde{R}(\boldsymbol{\theta})$ is equivalent to the sum of the errors of the maximum probability translations for each sentence in the dataset.

Clearly the set of candidate translations $\mathcal{Y}(\mathbf{w}^{(s)})$ is too large to explicitly sum over. Because the error function Δ generally does not decompose into smaller parts, there is no efficient dynamic programming solution to sum over this set. We can approximate the sum $\sum_{\mathbf{w}^{(t)} \in \mathcal{Y}(\mathbf{w}^{(s)})}$ with a sum over a finite number of samples, $\{\mathbf{w}_1^{(t)}, \mathbf{w}_2^{(t)}, \dots, \mathbf{w}_K^{(t)}\}$. If these samples were drawn uniformly at random, then the (annealed) risk would be approximated as (Shen et al., 2016),

$$\tilde{R}(\boldsymbol{\theta}) \approx \frac{1}{Z} \sum_{k=1}^K \tilde{p}(\mathbf{w}_k^{(t)} | \mathbf{w}^{(s)}; \boldsymbol{\theta}, \alpha) \times \Delta(\mathbf{w}_k^{(t)}, \mathbf{w}^{(t)}) \quad [18.55]$$

$$Z = \sum_{k=1}^K \tilde{p}(\mathbf{w}_k^{(t)} | \mathbf{w}^{(s)}; \boldsymbol{\theta}, \alpha). \quad [18.56]$$

Shen et al. (2016) report that performance plateaus at $K = 100$ for minimum risk training of neural machine translation.

Uniform sampling over the set of all possible translations is undesirable, because most translations have very low probability. A solution from Monte Carlo estimation is **importance sampling**, in which we draw samples from a **proposal distribution** $q(\mathbf{w}^{(t)})$. This distribution can be set equal to the current translation model $p(\mathbf{w}^{(t)} | \mathbf{w}^{(s)}; \boldsymbol{\theta})$. Each sample is then weighted by an **importance score**, $\omega_k = \frac{\tilde{p}(\mathbf{w}_k^{(t)} | \mathbf{w}^{(s)})}{q(\mathbf{w}_k^{(t)})}$. The effect of this weighting is to correct for any mismatch between the proposal distribution q and the true distribution \tilde{p} . The risk can then be approximated as,

$$\mathbf{w}_k^{(t)} \sim q(\mathbf{w}^{(t)}) \quad [18.57]$$

$$\omega_k = \frac{\tilde{p}(\mathbf{w}_k^{(t)} | \mathbf{w}^{(s)})}{q(\mathbf{w}_k^{(t)})} \quad [18.58]$$

$$\tilde{R}(\boldsymbol{\theta}) \approx \frac{1}{\sum_{k=1}^K \omega_k} \sum_{k=1}^K \omega_k \times \Delta(\mathbf{w}_k^{(t)}, \mathbf{w}^{(t)}). \quad [18.59]$$

Importance sampling will generally give a more accurate approximation with a given number of samples. The only formal requirement is that the proposal assigns non-zero probability to every $\mathbf{w}^{(t)} \in \mathcal{Y}(\mathbf{w}^{(s)})$. For more on importance sampling and related methods, see Robert and Casella (2013).

9845 Additional resources

9846 A complete textbook on machine translation is available from Koehn (2009). While this
9847 book precedes recent work on neural translation, a more recent draft chapter on neural
9848 translation models is also available (Koehn, 2017). Neubig (2017) provides a comprehen-
9849 sive tutorial on neural machine translation, starting from first principles. The course notes
9850 from Cho (2015) are also useful.

9851 Several neural machine translation systems are available, in connection with each of
9852 the major neural computing libraries: `lamtram` is an implementation of neural machine
9853 translation in the `dynet` (Neubig et al., 2017); `OpenNMT` (Klein et al., 2017) is an imple-
9854 mentation primarily in `Torch`; `tensor2tensor` is an implementation of several of the
9855 Google translation models in `tensorflow` (Abadi et al., 2016).

9856 Literary translation is especially challenging, even for expert human translators. Mes-
9857 sud (2014) describes some of these issues in her review of an English translation of *L'étranger*,
9858 the 1942 French novel by Albert Camus.¹² She compares the new translation by Sandra
9859 Smith against earlier translations by Stuart Gilbert and Matthew Ward, focusing on the
9860 difficulties presented by a single word in the first sentence:

9861 Then, too, Smith has reconsidered the book's famous opening. Camus's
9862 original is deceptively simple: "*Aujourd'hui, maman est morte.*" Gilbert influ-
9863 enced generations by offering us "Mother died today"—inscribing in Meur-
9864 sault [the narrator] from the outset a formality that could be construed as
9865 heartlessness. But *maman*, after all, is intimate and affectionate, a child's name
9866 for his mother. Matthew Ward concluded that it was essentially untranslatable
9867 ("mom" or "mummy" being not quite apt), and left it in the original French:
9868 "Maman died today." There is a clear logic in this choice; but as Smith has
9869 explained, in an interview in *The Guardian*, *maman* "didn't really tell the reader
9870 anything about the connotation." She, instead, has translated the sentence as
9871 "My mother died today."

9872 I chose "My mother" because I thought about how someone would
9873 tell another person that his mother had died. Meursault is speaking
9874 to the reader directly. "My mother died today" seemed to me the
9875 way it would work, and also implied the closeness of "maman" you
9876 get in the French.

9877 Elsewhere in the book, she has translated *maman* as "mama" — again, striving
9878 to come as close as possible to an actual, colloquial word that will carry the
9879 same connotations as *maman* does in French.

¹²The book review is currently available online at <http://www.nybooks.com/articles/2014/06/05/camus-new-letranger/>.

9880 The passage is a useful reminder that while the quality of machine translation has
9881 improved dramatically in recent years, expert human translations draw on considerations
9882 that are beyond the ken of any known computational approach.

9883 **Exercises**

9884 1. Give a synchronized derivation (§ 18.2.4) for the Spanish-English translation,

9885 (18.4) *El pez enojado atacado.*
9886 The fish angry attacked.
9886 The angry fish attacked.

9887 As above, the second line shows a word-for-word gloss, and the third line shows
9888 the desired translation. Use the synchronized production rule in [18.22], and design
9889 the other production rules necessary to derive this sentence pair. You may derive
9890 (*atacado*, *attacked*) directly from VP.

9891 Chapter 19

9892 Text generation

9893 In many of the most interesting problems in natural language processing, language is
9894 the output. The previous chapter described the specific case of machine translation, but
9895 there are many other applications, from summarization of research articles, to automated
9896 journalism, to dialogue systems. This chapter emphasizes three main scenarios: data-to-
9897 text, in which text is generated to explain or describe a structured record or unstructured
9898 perceptual input; text-to-text, which typically involves fusing information from multiple
9899 linguistic sources into a single coherent summary; and dialogue, in which text is generated
9900 as part of an interactive conversation with one or more human participants.

9901 19.1 Data-to-text generation

9902 In data-to-text generation, the input ranges from structured records, such as the descrip-
9903 tion of an weather forecast (as shown in Figure 19.1), to unstructured perceptual data,
9904 such as a raw image or video; the output may be a single sentence, such as an image cap-
9905 tion, or a multi-paragraph argument. Despite this diversity of conditions, all data-to-text
9906 systems share some of the same challenges (Reiter and Dale, 2000):

- 9907 • determining what parts of the data to describe;
- 9908 • planning a presentation of this information;
- 9909 • **lexicalizing** the data into words and phrases;
- 9910 • organizing words and phrases into well-formed sentences and paragraphs.

9911 The earlier stages of this process are sometimes called **content selection** and **text plan-**
9912 **ning**; the later stages are often called **surface realization**.

9913 Early systems for data-to-text generation were modular, with separate software com-
9914 ponents for each task. Artificial intelligence **planning** algorithms can be applied to both

Database:	Temperature			Cloud Sky Cover		
	time	min	mean	max	time	percent (%)
	06:00-21:00	9	15	21	06:00-09:00	25-50
	09:00-12:00				09:00-12:00	50-75
Wind Speed			Wind Direction			
	time	min	mean	max	time	mode
	06:00-21:00	15	20	30	06:00-21:00	S

Text: Cloudy, with temperatures between 10 and 20 degrees. South wind around 20 mph.

Figure 19.1: An example input-output pair for the task of generating text descriptions of weather forecasts (Konstas and Lapata, 2013). [todo: permission]

9915 the high-level information structure and the organization of individual sentences, ensur-
 9916 ing that communicative goals are met (McKeown, 1992; Moore and Paris, 1993). Surface
 9917 realization can be performed by grammars or templates, which link specific types of data
 9918 to candidate words and phrases. A simple example template is offered by Wiseman et al.
 9919 (2017), for generating descriptions of basketball games:

9920 (19.1) The <team1>(<wins1>-<losses1>) defeated the <team2>(<wins2>-<losses2>),
 9921 <pts1>-<pts2>.
 9922 The New York Knicks (45-5) defeated the Boston Celtics (11-38), 115-79.

9923 For more complex cases, it may be necessary to apply morphological inflections such as
 9924 pluralization and tense marking — even in the simple example above, languages such
 9925 as Russian would require case marking suffixes for the team names. Such inflections can
 9926 be applied as a postprocessing step. Another difficult challenge for surface realization is
 9927 the generation of varied **referring expressions** (e.g., *The Knicks*, *New York*, *they*), which is
 9928 critical to avoid repetition. As discussed in § 16.2.1, the form of referring expressions is
 9929 constrained by the discourse and information structure.

9930 An example at the intersection of rule-based and statistical techniques is the Nitrogen
 9931 system (Langkilde and Knight, 1998). The input to Nitrogen is an abstract meaning rep-
 9932 resentation (AMR; see § 13.3) of semantic content to be expressed in a single sentence. In
 9933 data-to-text scenarios, the abstract meaning representation is the output of a higher-level
 9934 text planning stage. A set of rules then converts the abstract meaning representation into
 9935 various sentence plans, which may differ in both the high-level structure (e.g., active ver-
 9936 sus passive voice) as well as the low-level details (e.g., word and phrase choice). Some
 9937 examples are shown in Figure 19.2. To control the combinatorial explosion in the number
 9938 of possible realizations for any given meaning, the sentence plans are unified into a single
 9939 finite-state acceptor, in which word tokens are represented by arcs (see § 9.1.1). A bigram

```
(a / admire-01
  :ARG0 (v / visitor
    :ARG1-of (c / arrive-01
      :ARG4 (j / Japan)))
  :ARG1 (m / "Mount Fuji"))
```

- Visitors who came to Japan admire Mount Fuji.
- Visitors who came in Japan admire Mount Fuji.
- Mount Fuji is admired by the visitor who came in Japan.

Figure 19.2: Abstract meaning representation and candidate surface realizations from the Nitrogen system. Example adapted from Langkilde and Knight (1998).

language model is then used to compute weights on the arcs, so that the shortest path is also the surface realization with the highest bigram language model probability.

More recent systems are unified models that are trained end-to-end using backpropagation. Data-to-text generation shares many properties with machine translation, including a problem of **alignment**: labeled examples provide the data and the text, but they do not specify which parts of the text correspond to which parts of the data. For example, to learn from Figure 19.1, the system must align the word *cloudy* to records in CLOUD SKY COVER, the phrases *10* and *20 degrees* to the MIN and MAX fields in TEMPERATURE, and so on. As in machine translation, both latent variables and neural attention have been proposed as solutions.

19.1.1 Latent data-to-text alignment

Given a dataset of texts and associated records $\{(\mathbf{w}^{(i)}, \mathbf{y}^{(i)})\}_{i=1}^N$, our goal is to learn a model Ψ , so that

$$\hat{\mathbf{w}} = \underset{\mathbf{w} \in \mathcal{V}^*}{\operatorname{argmax}} \Psi(\mathbf{w}, \mathbf{y}; \theta), \quad [19.1]$$

where \mathcal{V}^* is the set of strings over a discrete vocabulary, and θ is a vector of parameters. The relationship between \mathbf{w} and \mathbf{y} is complex: the data \mathbf{y} may contain dozens of records, and \mathbf{w} may extend to several sentences. To facilitate learning and inference, it would be helpful to decompose the scoring function Ψ into subcomponents. This would be possible if given an **alignment**, specifying which element of \mathbf{y} is expressed in each part of \mathbf{w} (Angeli et al., 2010):

$$\Psi(\mathbf{w}, \mathbf{y}; \theta) = \sum_{m=1}^M \psi_{w,y}(\mathbf{w}_m, \mathbf{y}_{z_m}) + \psi_z(z_m, z_{m-1}), \quad [19.2]$$

where z_m indicates the record aligned to word m . For example, in Figure 19.1, z_1 might specify that the word *cloudy* is aligned to the record *cloud-sky-cover:percent*. The score for this alignment would then be given by the weight on features such as

$$(\textit{cloudy}, \textit{cloud-sky-cover:percent}), \quad [19.3]$$

which could be learned from labeled data $\{(\mathbf{w}^{(i)}, \mathbf{y}^{(i)}, \mathbf{z}^{(i)})\}_{i=1}^N$. The function ψ_z can learn to assign higher scores to alignments that are coherent, referring to the same records in adjacent parts of the text.¹

Several datasets include structured records and natural language text (Barzilay and McKeown, 2005; Chen and Mooney, 2008; Liang and Klein, 2009), but the alignments between text and records are usually not available.² One solution is to model the problem probabilistically, treating the alignment as a latent variable (Liang et al., 2009; Konstas and Lapata, 2013). The model can then be estimated using expectation maximization or sampling (see chapter 5).

19.1.2 Neural data-to-text generation

The **encoder-decoder model** and **neural attention** were introduced in § 18.3 as methods for neural machine translation. They can also be applied to data-to-text generation, with the data acting as the source language (Mei et al., 2016). In neural machine translation, the attention mechanism linked words in the source to words in the target; in data-to-text generation, the attention mechanism can link each part of the generated text back to a record in the data. The biggest departure from translation is in the encoder, which depends on the form of the data.

19.1.2.1 Data encoders

In some types of structured records, all values are drawn from discrete sets. For example, the birthplace of an individual is drawn from a discrete set of possible locations; the diagnosis and treatment of a patient are drawn from an exhaustive list of clinical codes (Johnson et al., 2016). In such cases, vector embeddings can be estimated for each field and possible value: for example, a vector embedding for the field BIRTHPLACE, and another embedding for the value BERKELEY_CALIFORNIA (Bordes et al., 2011). The table of such embeddings serves as the encoding of a structured record (He et al., 2017). It is also possible to compress the entire table into a single vector representation, by **pooling** across the embeddings of each field and value (Lebret et al., 2016).

Sequences Some types of structured records have a natural ordering, such as events in a game (Chen and Mooney, 2008) and steps in a recipe (Tutin and Kittredge, 1992). For example, the following records describe a sequence of events in a robot soccer match (Mei

¹More expressive decompositions of Ψ are possible. For example, Wong and Mooney (2007) use a synchronous context-free grammar (see § 18.2.4) to “translate” between a meaning representation and natural language text.

²An exception is a dataset of records and summaries from American football games, containing annotations of alignments between sentences and records (Snyder and Barzilay, 2007).

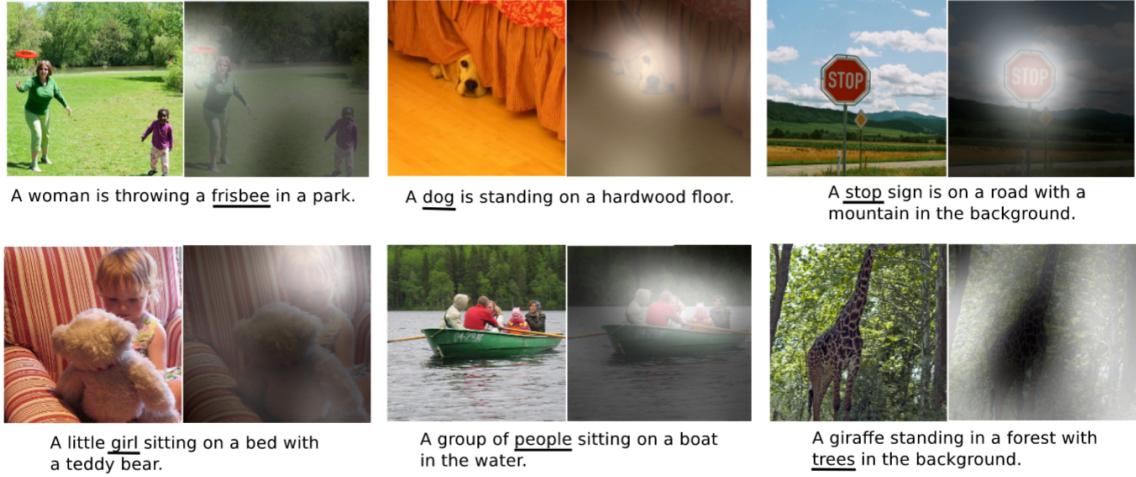


Figure 19.3: Examples of the image captioning task, with attention masks shown for each of the underlined words. From Xu et al. (2015). [todo: permission]

et al., 2016):

```
PASS(arg1 = PURPLE6,arg2 = PURPLE3)
KICK(arg1 = PURPLE3)
BADPASS(arg1 = PURPLE3,arg2 = PINK9).
```

9989 Each event is a single record, and can be encoded by a concatenation of vector representations for the event type (e.g., PASS), the field (e.g., arg1), and the values (e.g., PURPLE3),
9990 e.g.,

$$\mathbf{X} = [\mathbf{u}_{\text{PASS}}, \mathbf{u}_{\text{arg1}}, \mathbf{u}_{\text{PURPLE6}}, \mathbf{u}_{\text{arg2}}, \mathbf{u}_{\text{PURPLE3}}]. \quad [19.4]$$

9992 This encoding can then act as the input layer for a recurrent neural network, yielding a
9993 sequence of vector representations $\{z_r\}_{r=1}^R$, where r indexes over records. Interestingly,
9994 this sequence-based approach is effective even in cases where there is no natural ordering
9995 over the records, such as the weather data in Figure 19.1 (Mei et al., 2016).

9996 **Images** Another flavor of data-to-text generation is the generation of text captions for
9997 images. Examples from this task are shown in Figure 19.3. Images are naturally repre-
9998 sented as tensors: a color image of 320×240 pixels would be stored as a tensor with
9999 $320 \times 240 \times 3$ intensity values. The dominant approach to image classification is to en-
10000 code images as vectors using a combination of convolution and pooling (Krizhevsky et al.,
10001 2012). Chapter 3 explains how to use convolutional networks for text; for images, convo-
10002 lution is applied across the vertical, horizontal, and color dimensions. By pooling the re-
10003 sults of successive convolutions, the image is converted to a vector representation, which

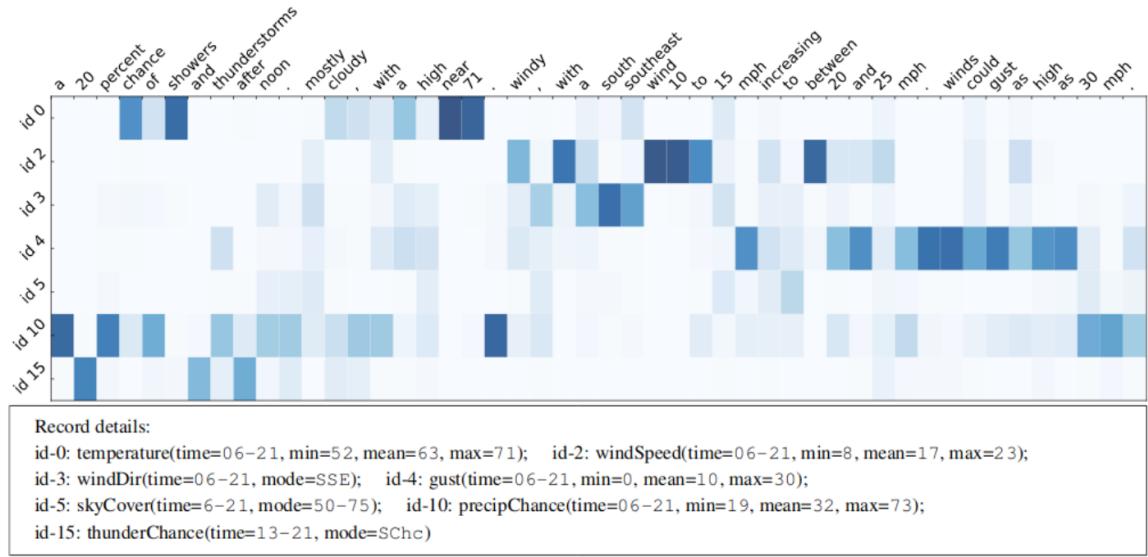


Figure 19.4: Neural attention in text generation. Figure from Mei et al. (2016).[todo: permission]

10004 can then be fed directly into the decoder as the initial state (Vinyals et al., 2015), just as
 10005 in the sequence-to-sequence translation model (see § 18.3). Alternatively, one can apply
 10006 a set of convolutional networks, yielding vector representations for different parts of the
 10007 image, which can then be combined using neural attention (Xu et al., 2015).

10008 19.1.2.2 Attention

Given a set of embeddings of the data $\{\mathbf{z}_r\}_{r=1}^R$ and a decoder state \mathbf{h}_m , the attention vector over the data can be computed using the same technique described in § 18.3.1. When generating word m of the output, a softmax attention mechanism computes the weighted average \mathbf{c}_m ,

$$\psi_\alpha(m, r) = \beta_\alpha \cdot f(\Theta_\alpha[\mathbf{h}_m; \mathbf{z}_r]) \quad [19.5]$$

$$\boldsymbol{\alpha}_m = \text{SoftMax}([\psi_\alpha(m, 1), \psi_\alpha(m, 2), \dots, \psi_\alpha(m, R)]) \quad [19.6]$$

$$\mathbf{c}_m = \sum_{r=1}^R \alpha_{m \rightarrow r} \mathbf{z}_r, \quad [19.7]$$

10009 where f is an elementwise nonlinearity such as tanh or ReLU (see § 3.2.1). The weighted
 10010 average \mathbf{c}_m can then be included in the recurrent update to the decoder state, or in the
 10011 emission probabilities, as described in § 18.3.1. Figure 19.4 shows the attention to compo-
 10012 nents of a weather record, while generating the text shown on the x -axis.

10013 Adapting this architecture to image captioning is straightforward. A convolutional
 10014 neural networks is applied to a set of image locations, and the output at each location ℓ is
 10015 represented with a vector z_ℓ . Attention can then be computed over the image locations,
 10016 as shown in the right panels of each pair of images in Figure 19.3.

10017 Various modifications to this basic mechanism have been proposed. In **coarse-to-fine**
 10018 **attention** (Mei et al., 2016), each record receives a global attention $a_r \in [0, 1]$, which is in-
 10019 dependent of the decoder state. This global attention, which represents the overall impor-
 10020 tance of the record, is multiplied with the decoder-based attention scores, before comput-
 10021 ing the final normalized attentions. In **structured attention**, the attention vector $\alpha_{m \rightarrow \cdot}$ can
 10022 include structural biases, which can favor assigning higher attention values to contiguous
 10023 segments or to dependency subtrees (Kim et al., 2017). Structured attention vectors can
 10024 be computed by running the forward-backward algorithm to obtain marginal attention
 10025 probabilities (see § 7.5.3.3). Because each step in the forward-backward algorithm is dif-
 10026 ferentiable, it can be encoded in a computation graph, and end-to-end learning can be
 10027 performed by backpropagation.

10028 19.1.2.3 Decoder

10029 Given the encoding, the decoder can function just as in neural machine translation (see
 10030 § 18.3.1), using the attention-weighted encoder representation in the decoder recurrence
 10031 and/or output computation. As in machine translation, beam search can help to avoid
 10032 search errors (Lebret et al., 2016).

Many applications require generating words that do not appear in the training vocabulary. For example, a weather record may contain a previously unseen city name; a sports record may contain a previously unseen player name. Such tokens can be generated in the text by copying them over from the input (e.g., Gulcehre et al., 2016).³ First introduce an additional variable $s_m \in \{\text{gen}, \text{copy}\}$, indicating whether token $w_m^{(t)}$ should be generated or copied. The decoder probability is then,

$$p(w^{(t)} | w_{1:m-1}^{(t)}, \mathbf{Z}, s_m) = \begin{cases} \text{SoftMax}(\beta_{w^{(t)}} \cdot h_{m-1}^{(t)}), & s_m = \text{gen} \\ \sum_{r=1}^R \delta(w_r^{(s)} = w^{(t)}) \times \alpha_{m \rightarrow r}, & s_m = \text{copy}, \end{cases} \quad [19.8]$$

10033 where $\delta(w_r^{(s)} = w^{(t)})$ is an indicator function, taking the value 1 iff the text of the record
 10034 $w_r^{(s)}$ is identical to the target word $w^{(t)}$. The probability of copying record r from the source
 10035 is $\delta(s_m = \text{copy}) \times \alpha_{m \rightarrow r}$, the product of the copy probability by the local attention. Note
 10036 that in this model, the attention weights α_m are computed from the *previous* decoder state
 10037 h_{m-1} . The computation graph therefore remains a feedforward network, with recurrent
 10038 paths such as $h_{m-1}^{(t)} \rightarrow \alpha_m \rightarrow w_m^{(t)} \rightarrow h_m^{(t)}$.

³A number of variants of this strategy have been proposed (e.g., Gu et al., 2016; Merity et al., 2017). See Wiseman et al. (2017) for an overview.

10039 To facilitate end-to-end training, the switching variable s_m can be represented by a
 10040 gate π_m , which is computed from a two-layer feedforward network, whose input consists
 10041 of the concatenation of the decoder state $\mathbf{h}_{m-1}^{(t)}$ and the attention-weighted representation
 10042 of the data, $\mathbf{c}_m = \sum_{r=1}^R \alpha_{m \rightarrow r} \mathbf{z}_r$,

$$\pi_m = \sigma(\Theta^{(2)} f(\Theta^{(1)}[\mathbf{h}_{m-1}^{(t)}; \mathbf{c}_m])). \quad [19.9]$$

The full generative probability at token m is then,

$$p(w^{(t)} | \mathbf{w}_{1:m}^{(t)}, \mathbf{Z}) = \pi_m \times \underbrace{\frac{\exp \beta_{w^{(t)}} \cdot \mathbf{h}_{m-1}^{(t)}}{\sum_{j=1}^V \exp \beta_j \cdot \mathbf{h}_{m-1}^{(t)}}}_{\text{generate}} + (1 - \pi_m) \times \underbrace{\sum_{r=1}^R \delta(w_r^{(s)} = w^{(t)}) \times \alpha_{m \rightarrow r}}_{\text{copy}}. \quad [19.10]$$

10043 19.2 Text-to-text generation

10044 Text-to-text generation includes problems of summarization and simplification:

- 10045 • reading a novel and outputting a paragraph-long summary of the plot;⁴
- 10046 • reading a set of blog posts about politics, and outputting a bullet list of the various
 10047 issues and perspectives;
- 10048 • reading a technical research article about the long-term health consequences of drink-
 10049 ing kombucha, and outputting a summary of the article in language that non-experts
 10050 can understand.

10051 These problems can be approached in two ways: through the encoder-decoder architec-
 10052 ture discussed in the previous section, or by operating directly on the input text.

10053 19.2.1 Neural abstractive summarization

10054 **Sentence summarization** is the task of shortening a sentence while preserving its mean-
 10055 ing, as in the following examples (Knight and Marcu, 2000; Rush et al., 2015):

- 10056 (19.2) The documentation is typical of Epson quality: excellent.
 10057 Documentation is excellent.

⁴In § 16.3.4.1, we encountered a special case of single-document summarization, which involved extracting the most important sentences or discourse units. We now consider the more challenging problem of **abstractive summarization**, in which the summary can include words that do not appear in the original text.

- 10059 (19.3) Russian defense minister Ivanov called sunday for the creation of a joint front for
 10060 combating global terrorism.
 10061 Russia calls for joint front against terrorism.
 10062

10063 Sentence summarization is closely related to **sentence compression**, in which the sum-
 10064 mary is produced by deleting words or phrases from the original (Clarke and Lapata,
 10065 2008). But as shown in (19.3), a sentence summary can also introduce new words, such as
 10066 *against*, which replaces the phrase *for combatting*.

10067 Sentence summarization can be treated as a machine translation problem, using the at-
 10068 tentional encoder-decoder translation model discussed in § 18.3.1 (Rush et al., 2015). The
 10069 longer sentence is encoded into a sequence of vectors, one for each token. The decoder
 10070 then computes attention over these vectors when updating its own recurrent state. As
 10071 with data-to-text generation, it can be useful to augment the encoder-decoder model with
 10072 the ability to copy words directly from the source. Rush et al. (2015) train this model by
 10073 building four million sentence pairs from news articles. In each pair, the longer sentence is
 10074 the first sentence of the article, and the summary is the article headline. Sentence summa-
 10075 rization can also be trained in a semi-supervised fashion, using a probabilistic formulation
 10076 of the encoder-decoder model called a **variational autoencoder** (Miao and Blunsom, 2016,
 10077 also see § 14.8.2).

When summarizing longer documents, an additional concern is that the summary not be repetitive: each part of the summary should cover new ground. This can be addressed by maintaining a vector of the sum total of all attention values thus far, $t_m = \sum_{n=1}^m \alpha_n$. This total can be used as an additional input to the computation of the attention weights,

$$\alpha_{m \rightarrow n} \propto \exp \left(\mathbf{v}_\alpha \cdot \tanh(\Theta_\alpha[\mathbf{h}_m^{(t)}; \mathbf{h}_n^{(s)}; \mathbf{t}_m]) \right), \quad [19.11]$$

which enables the model to learn to prefer parts of the source which have not been attended to yet (Tu et al., 2016). To further encourage diversity in the generated summary, See et al. (2017) introduce a **coverage loss** to the objective function,

$$\ell_m = \sum_{n=1}^{M^{(s)}} \min(\alpha_{m \rightarrow n}, t_{m \rightarrow n}). \quad [19.12]$$

10078 This loss will be low if $\alpha_{m \rightarrow \cdot}$ assigns little attention to words that already have large
 10079 values in $t_{m \rightarrow \cdot}$. Coverage loss is similar to the concept of **marginal relevance**, in which
 10080 the reward for adding new content is proportional to the extent to which it increases
 10081 the overall amount of information conveyed by the summary (Carbonell and Goldstein,
 10082 1998).

10083 **19.2.2 Sentence fusion for multi-document summarization**

10084 In **multi-document summarization**, the goal is to produce a summary that covers the
 10085 content of several documents (McKeown et al., 2002). One approach to this challenging
 10086 problem is to identify sentences across multiple documents that relate to a single theme,
 10087 and then to fuse them into a single sentence (Barzilay and McKeown, 2005). As an exam-
 10088 ple, consider the following two sentences (McKeown et al., 2010):

- 10089 (19.4) Palin actually turned against the bridge project only after it became a national
 10090 symbol of wasteful spending.
 10091 (19.5) Ms. Palin supported the bridge project while running for governor, and aban-
 10092 doned it after it became a national scandal.

10093 An *intersection* preserves only the content that is present in both sentences:

- 10094 (19.6) Palin turned against the bridge project after it became a national scandal.

10095 A *union* includes information from both sentences:

- 10096 (19.7) Ms. Palin supported the bridge project while running for governor, but turned
 10097 against it when it became a national scandal and a symbol of wasteful spending.

Dependency parsing is often used as a technique for sentence fusion. After parsing each sentence, the resulting dependency trees can be aggregated into a lattice (Barzilay and McKeown, 2005) or a graph structure (Filippova and Strube, 2008), in which identical or closely related words (e.g., *Palin*, *bridge*, *national*) are fused into a single node. The resulting graph can then be pruned back to a tree by solving an **integer linear program** (see § 13.2.2),

$$\max_{\mathbf{y}} \sum_{i,j,r} \psi(i \xrightarrow{r} j, \mathbf{w}; \boldsymbol{\theta}) \times y_{i,j,r} \quad [19.13]$$

$$\text{s.t. } \mathbf{y} \in \mathcal{C}, \quad [19.14]$$

10098 where the variable $y_{i,j,r} \in \{0, 1\}$ indicates whether there is an edge from i to j of type r ,
 10099 the score of this edge is $\psi(i \xrightarrow{r} j, \mathbf{w}; \boldsymbol{\theta})$, and \mathcal{C} is a set of constraints, described below. As
 10100 usual, \mathbf{w} is the list of words in the graph, and $\boldsymbol{\theta}$ is a vector of parameters. The score $\psi(i \xrightarrow{r} j, \mathbf{w}; \boldsymbol{\theta})$
 10101 reflects the “importance” of the modifier j to the overall meaning: in intersective
 10102 fusion, this score indicates the extent to which the content in this edge is expressed in all
 10103 sentences; in union fusion, the score indicates whether the content in the edge is expressed
 10104 in any sentence.

10105 The constraint set \mathcal{C} ensures that \mathbf{y} forms a valid dependency graph. It can also im-
 10106 pose additional linguistic constraints: for example, ensuring that coordinated nouns are

10107 sufficiently similar. The resulting tree must then be **linearized** into a sentence. This is
 10108 typically done by generating a set of candidate linearizations, and choosing the one with
 10109 the highest score under a language model (Langkilde and Knight, 1998; Song et al., 2016).

10110 **19.3 Dialogue**

10111 **Dialogue systems** are capable of conversing with a human interlocutor, often to per-
 10112 form some task (Grosz, 1979), but sometimes just to chat (Weizenbaum, 1966). While re-
 10113 search on dialogue systems goes back several decades (Carbonell, 1970; Winograd, 1972),
 10114 commercial systems such as Alexa and Siri have recently brought this technology into
 10115 widespread use. Nonetheless, there is a significant gap between research and practice:
 10116 many practical dialogue systems remain scripted and inflexible, while research systems
 10117 emphasize abstractive text generation, “on-the-fly” decision making, and probabilistic
 10118 reasoning about the user’s intentions.

10119 **19.3.1 Finite-state and agenda-based dialogue systems**

10120 Finite-state automata were introduced in chapter 9 as a formal model of computation,
 10121 in which string inputs and outputs are linked to transitions between a finite number of
 10122 discrete states. This model naturally fits simple task-oriented dialogues, such as the one
 10123 shown in the left panel of Figure 19.5. This (somewhat frustrating) dialogue can be repre-
 10124 sented with a finite-state transducer, as shown in the right panel of the figure. The accept-
 10125 ing state is reached only when the two needed pieces of information are provided, and the
 10126 human user confirms that the order is correct. In this simple scenario, the TOPPING and
 10127 ADDRESS are the two **slots** associated with the activity of ordering a pizza, which is called
 10128 a **frame**. Frame representations can be hierarchical: for example, an ADDRESS could have
 10129 slots of its own, such as STREET and CITY.

10130 In the example dialogue in Figure 19.5, the user provides the precise inputs that are
 10131 needed in each turn (e.g., *anchovies*; *the College of Computing building*). Some users may
 10132 prefer to communicate more naturally, with phrases like *I’d, uh, like some anchovies please*.
 10133 One approach to handling such utterances is to design a custom grammar, with non-
 10134 terminals for slots such as TOPPING and LOCATION. However, context-free parsing of
 10135 unconstrained speech input is challenging. A more lightweight alternative is BIO-style
 10136 sequence labeling (see § 8.3), e.g.:

10137 (19.9) *I’d like anchovies , and please bring it to the College of Computing*
 O O B-TOPPING O O O O O O O B-ADDR I-ADDR I-ADDR I-ADDR I-ADDR
 10138 *Building .*
 I-ADDR O

- (19.8) A: I want to order a pizza.
 B: What toppings?
 A: Anchovies.
 B: Ok, what address?
 A: The College of Computing building.
 B: Please confirm: one pizza with artichokes, to be delivered to the College of Computing building.
 A: No.
 B: What toppings?
 ...

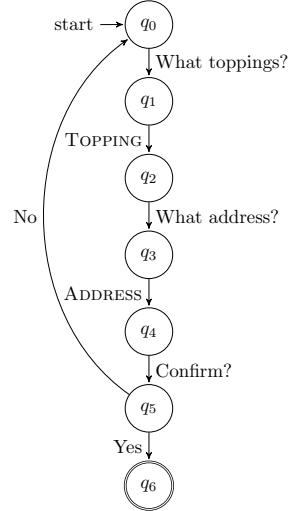


Figure 19.5: An example dialogue and the associated finite-state model. In the finite-state model, SMALL CAPS indicates that the user must provide information of this type in their answer.

10139 The tagger can be driven by a bi-directional recurrent neural network, similar to recurrent
 10140 approaches to semantic role labeling described in § 13.2.3.

10141 The input in (19.9) could not be handled by the finite-state system from Figure 19.5,
 10142 which forces the user to provide the topping first, and then the location. In this sense,
 10143 the **initiative** is driven completely by the system. **Agenda-based dialogue systems** ex-
 10144 tend finite-state architectures by attempting to recognize all slots that are filled by the
 10145 user’s reply, thereby handling these more complex examples. Agenda-based systems dy-
 10146 namically pose additional questions until the frame is complete (Bobrow et al., 1977; Allen
 10147 et al., 1995; Rudnicky and Xu, 1999). Such systems are said to be **mixed-initiative**, because
 10148 both the user and the system can drive the direction of the dialogue.

10149 19.3.2 Markov decision processes

10150 The task of dynamically selecting the next move in a conversation is known as **dialogue**
 10151 **management**. This problem can be framed as a **Markov decision process**, which is a
 10152 theoretical model that includes a discrete set of states, a discrete set of actions, a function
 10153 that computes the probability of transitions between states, and a function that computes
 10154 the cost or reward of action-state pairs. Let’s see how each of these elements pertains to
 10155 the pizza ordering dialogue system.

- 10156 • Each state is a tuple of information about whether the topping and address are

10157 known, and whether the order has been confirmed. For example,

(KNOWN TOPPING, UNKNOWN ADDRESS, NOT CONFIRMED) [19.15]

10158 is a possible state. Any state in which the pizza order is confirmed is a terminal
10159 state, and the Markov decision process stops after entering such a state.

- 10160 • The set of actions includes querying for the topping, querying for the address, and
10161 requesting confirmation. Each action induces a probability distribution over states,
10162 $p(s_t | a_t, s_{t-1})$. For example, requesting confirmation of the order is not likely to
10163 result in a transition to the terminal state if the topping is not yet known. This
10164 probability distribution over state transitions may be learned from data, or it may
10165 be specified in advance.
- 10166 • Each state-action-state tuple earns a reward, $r_a(s_t, s_{t+1})$. In the context of the pizza
10167 ordering system, a simple reward function would be,

$$r_a(s_t, s_{t+1}) = \begin{cases} 0, & a = \text{CONFIRM}, s_{t+1} = (*, *, \text{CONFIRMED}) \\ -10, & a = \text{CONFIRM}, s_{t+1} = (*, *, \text{NOT CONFIRMED}) \\ -1, & a \neq \text{CONFIRM} \end{cases} \quad [19.16]$$

10168 This function assigns zero reward for successful transitions to the terminal state, a
10169 large negative reward to a rejected request for confirmation, and a small negative re-
10170ward for every other type of action. The system is therefore rewarded for reaching
10171 the terminal state in few steps, and penalized for prematurely requesting confirma-
10172 tion.

10173 In a Markov decision process, a **policy** is a function $\pi : \mathcal{S} \mapsto \mathcal{A}$ that maps from states to
10174 actions (see § 15.2.4.3). The value of a policy is the expected sum of discounted rewards,
10175 $E_\pi[\sum_{t=1}^T \gamma^t r_{a_t}(s_t, s_{t+1})]$, where γ is the discount factor, $\gamma \in [0, 1)$. Discounting has the
10176 effect of emphasizing rewards that can be obtained immediately over less certain rewards
10177 in the distant future.

10178 An optimal policy can be obtained by dynamic programming, by iteratively updating
10179 the **value function** $V(s)$, which is the expectation of the cumulative reward from s under
10180 the optimal action a ,

$$V(s) \leftarrow \max_{a \in \mathcal{A}} \sum_{s' \in \mathcal{S}} p(s' | s, a)[r_a(s, s') + \gamma V(s')]. \quad [19.17]$$

10181 The value function $V(s)$ is computed in terms of $V(s')$ for all states $s' \in \mathcal{S}$. A series
10182 of iterative updates to the value function will eventually converge to a stationary point.
10183 This algorithm is known as **value iteration**. Given the converged value function $V(s)$, the

10184 optimal action at each state is the argmax,

$$\pi(s) = \operatorname{argmax}_{a \in \mathcal{A}} \sum_{s' \in \mathcal{S}} p(s' | s, a)[r_a(s, s') + \gamma V(s')]. \quad [19.18]$$

10185 Value iteration and related algorithms are described in detail by Sutton and Barto (1998).
 10186 For applications to dialogue systems, see Levin et al. (1998) and Walker (2000).

10187 The Markov decision process framework assumes that the current state of the dialogue
 10188 is known. In reality, the system may misinterpret the user’s statements — for example,
 10189 believing that a specification of the delivery location (PEACHTREE) is in fact a specification
 10190 of the topping (PEACHES). In a **partially observable Markov decision process (POMDP)**,
 10191 the system receives an *observation* o , which is probabilistically conditioned on the state,
 10192 $p(o | s)$. It must therefore maintain a distribution of beliefs about which state it is in, with
 10193 $q_t(s)$ indicating the degree of belief that the dialogue is in state s at time t . The POMDP
 10194 formulation can help to make dialogue systems more robust to errors, particularly in the
 10195 context of spoken language dialogues, where the speech itself may be misrecognized (Roy
 10196 et al., 2000; Williams and Young, 2007). However, finding the optimal policy in a POMDP
 10197 is computationally intractable, requiring additional approximations.

10198 19.3.3 Neural chatbots

10199 Chatting is a lot easier when you don’t need to get anything done. **Chatbots** are systems
 10200 that parry the user’s input with a response that keeps the conversation going. They can be
 10201 built from the encoder-decoder architecture discussed in § 18.3 and § 19.1.2: the encoder
 10202 converts the user’s input into a vector, and the decoder produces a sequence of words as a
 10203 response. For example, Shang et al. (2015) apply the attentional encoder-decoder transla-
 10204 tion model, training on a dataset of posts and responses from the Chinese microblogging
 10205 platform Sina Weibo.⁵ This approach is capable of generating replies that relate themati-
 10206 cally to the input, as shown in the following examples:⁶

10207 (19.10) A: High fever attacks me every New Year’s day.
 10208 Get B: well soon and stay healthy!

10209 (19.11) A: I gain one more year. Grateful to my group, so happy.
 10210 B: Getting old now. Time has no mercy.

10211 While encoder-decoder models can generate responses that make sense in the con-
 10212 text of the immediately preceding turn, they struggle to maintain coherence over longer

⁵Twitter is also frequently used for construction of dialogue datasets (Ritter et al., 2011; Sordoni et al., 2015). Another source is technical support chat logs from the Ubuntu linux distribution (Uthus and Aha, 2013; Lowe et al., 2015).

⁶All examples are translated from Chinese by Shang et al. (2015).

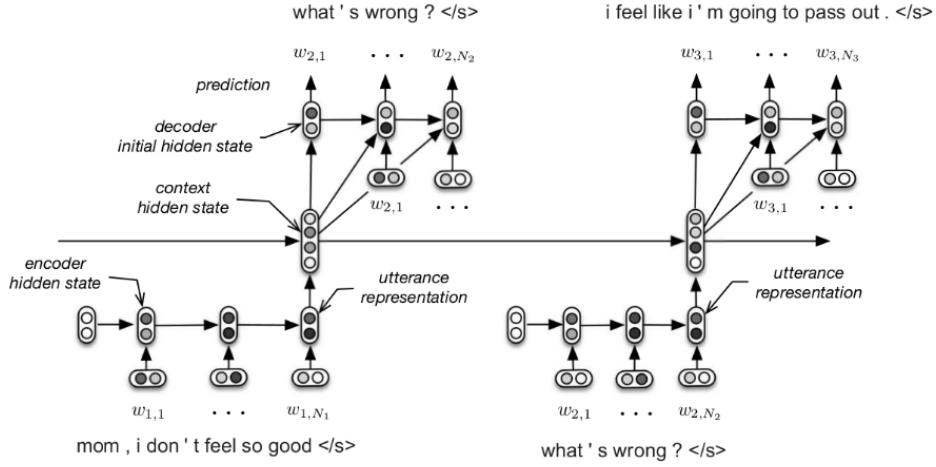


Figure 19.6: A hierarchical recurrent neural network for dialogue, with recurrence over both words and turns, from Serban et al. (2016). [todo: permission]

conversations. One solution is to model the dialogue context recurrently. This creates a **hierarchical recurrent network**, including both word-level and turn-level recurrences. The turn-level hidden state is then used as additional context in the decoder (Serban et al., 2016), as shown in Figure 19.6.

An open question is how to integrate the encoder-decoder architecture into task-oriented dialogue systems. Neural chatbots can be trained end-to-end: the user’s turn is analyzed by the encoder, and the system output is generated by the decoder. This architecture can be trained by log-likelihood using backpropagation (e.g., Sordoni et al., 2015; Serban et al., 2016), or by more elaborate objectives, using reinforcement learning (Li et al., 2016). In contrast, the task-oriented dialogue systems described in § 19.3.1 typically involve a set of specialized modules: one for recognizing the user input, another for deciding what action to take, and a third for arranging the text of the system output.

Recurrent neural network decoders can be integrated into Markov Decision Process dialogue systems, by conditioning the decoder on a representation of the information that is to be expressed in each turn (Wen et al., 2015). Specifically, the long short-term memory (LSTM; § 6.3) architecture is augmented so that the memory cell at turn m takes an additional input d_m , which is a representation of the slots and values to be expressed in the next turn. However, this approach still relies on additional modules to recognize the user’s utterance and to plan the overall arc of the dialogue.

Another promising direction is to create embeddings for the elements in the domain: for example, the slots in a record and the entities that can fill them. The encoder then

10234 encodes not only the words of the user’s input, but the embeddings of the elements that
 10235 the user mentions. Similarly, the decoder is endowed with the ability to refer to specific
 10236 elements in the knowledge base. He et al. (2017) show that such a method can learn to
 10237 play a collaborative dialogue game, in which both players are given a list of entities and
 10238 their properties, and the goal is to find an entity that is on both players’ lists.

10239 Further reading

10240 Gatt and Krahmer (2018) provide a comprehensive recent survey on text generation. For
 10241 a book-length treatment of earlier work, see Reiter and Dale (2000). For a survey on image
 10242 captioning, see Bernardi et al. (2016); for a survey of pre-neural approaches to dialogue
 10243 systems, see Rieser and Lemon (2011). **Dialogue acts** were introduced in § 8.6 as a labeling
 10244 scheme for human-human dialogues; they also play a critical role in task-based dialogue
 10245 systems (e.g., Allen et al., 1996). The incorporation of theoretical models of dialogue into
 10246 computational systems is reviewed by Jurafsky and Martin (2009, chapter 24).

10247 While this chapter has focused on the informative dimension of text generation, another
 10248 line of research aims to generate text with configurable stylistic properties (Walker
 10249 et al., 1997; Mairesse and Walker, 2011; Ficler and Goldberg, 2017; Hu et al., 2017). This
 10250 chapter also does not address the generation of creative text such as narratives (Riedl and
 10251 Young, 2010), jokes (Ritchie, 2001), poems (Colton et al., 2012), and song lyrics (Gonçalo Oliveira
 10252 et al., 2007).

10253 Exercises

10254 1. The SimpleNLG system produces surface realizations from representations of de-
 10255 sired syntactic structure (Gatt and Reiter, 2009). This system can be accessed on
 10256 github at <https://github.com/simpleNLG/simpleNLG>. Download the sys-
 10257 tem, and produce realizations of the following examples:

- 10258 (19.12) Call me Ismael.
- 10259 (19.13) I try all things.
- 10260 (19.14) I achieve what I can.

10261 Then convert each example to a question. [todo: Can’t get SimpleNLG to work with
 10262 python anymore]

10263 **Appendix A**

10264 **Probability**

10265 Probability theory provides a way to reason about random events. The sorts of random
10266 events that are typically used to explain probability theory include coin flips, card draws,
10267 and the weather. It may seem odd to think about the choice of a word as akin to the flip of
10268 a coin, particularly if you are the type of person to choose words carefully. But random or
10269 not, language has proven to be extremely difficult to model deterministically. Probability
10270 offers a powerful tool for modeling and manipulating linguistic data.

10271 Probability can be thought of in terms of **random outcomes**: for example, a single coin
10272 flip has two possible outcomes, heads or tails. The set of possible outcomes is the **sample**
10273 **space**, and a subset of the **sample space** is an **event**. For a sequence of two coin flips,
10274 there are four possible outcomes, $\{HH, HT, TH, TT\}$, representing the ordered sequences
10275 heads-head, heads-tails, tails-heads, and tails-tails. The event of getting exactly one head
10276 includes two outcomes: $\{HT, TH\}$.

10277 Formally, a probability is a function from events to the interval between zero and one:
10278 $\Pr : \mathcal{F} \mapsto [0, 1]$, where \mathcal{F} is the set of possible events. An event that is certain has proba-
10279 bility one; an event that is impossible has probability zero. For example, the probability
10280 of getting fewer than three heads on two coin flips is one. Each outcome is also an event
10281 (a set with exactly one element), and for two flips of a fair coin, the probability of each
10282 outcome is,

$$\Pr(\{HH\}) = \Pr(\{HT\}) = \Pr(\{TH\}) = \Pr(\{TT\}) = \frac{1}{4}. \quad [\text{A.1}]$$

10283 **A.1 Probabilities of event combinations**

10284 Because events are sets of outcomes, we can use set-theoretic operations such as comple-
10285 ment, intersection, and union to reason about the probabilities of events and their combi-
10286 nations.

10287 For any event A , there is a **complement** $\neg A$, such that:

- 10288 • The probability of the union $A \cup \neg A$ is $\Pr(A \cup \neg A) = 1$;
- 10289 • The intersection $A \cap \neg A = \emptyset$ is the empty set, and $\Pr(A \cap \neg A) = 0$.

10290 In the coin flip example, the event of obtaining a single head on two flips corresponds to
 10291 the set of outcomes $\{HT, TH\}$; the complement event includes the other two outcomes,
 10292 $\{TT, HH\}$.

10293 A.1.1 Probabilities of disjoint events

10294 When two events have an empty intersection, $A \cap B = \emptyset$, they are **disjoint**. The probabili-
 10295 ty of the union of two disjoint events is equal to the sum of their probabilities,

$$A \cap B = \emptyset \Rightarrow \Pr(A \cup B) = \Pr(A) + \Pr(B). \quad [A.2]$$

10296 This is the **third axiom of probability**, and it can be generalized to any countable sequence
 10297 of disjoint events.

In the coin flip example, this axiom can derive the probability of the event of getting a single head on two flips. This event is the set of outcomes $\{HT, TH\}$, which is the union of two simpler events, $\{HT, TH\} = \{HT\} \cup \{TH\}$. The events $\{HT\}$ and $\{TH\}$ are disjoint. Therefore,

$$\Pr(\{HT, TH\}) = \Pr(\{HT\} \cup \{TH\}) = \Pr(\{HT\}) + \Pr(\{TH\}) \quad [A.3]$$

$$= \frac{1}{4} + \frac{1}{4} = \frac{1}{2}. \quad [A.4]$$

10298 In the general, the probability of the union of two events is,

$$\Pr(A \cup B) = \Pr(A) + \Pr(B) - \Pr(A \cap B). \quad [A.5]$$

This can be seen visually in Figure A.1, and it can be derived from the third axiom of probability. Consider an event that includes all outcomes in B that are not in A , denoted as $B - (A \cap B)$. By construction, this event is disjoint from A . We can therefore apply the additive rule,

$$\Pr(A \cup B) = \Pr(A) + \Pr(B - (A \cap B)). \quad [A.6]$$

Furthermore, the event B is the union of two disjoint events: $A \cap B$ and $B - (A \cap B)$.

$$\Pr(B) = \Pr(B - (A \cap B)) + \Pr(A \cap B). \quad [A.7]$$

Reorganizing and substituting into Equation A.6 gives the desired result:

$$\Pr(B - (A \cap B)) = \Pr(B) - \Pr(A \cap B) \quad [A.8]$$

$$\Pr(A \cup B) = \Pr(A) + \Pr(B) - \Pr(A \cap B). \quad [A.9]$$

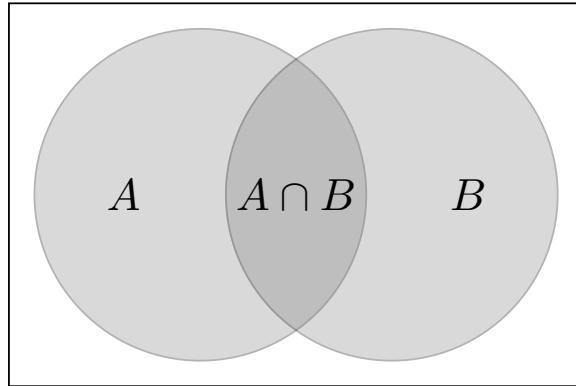


Figure A.1: A visualization of the probability of non-disjoint events A and B .

10299 A.1.2 Law of total probability

10300 A set of events $\mathcal{B} = \{B_1, B_2, \dots, B_N\}$ is a **partition** of the sample space iff each pair of
 10301 events is disjoint ($B_i \cap B_j = \emptyset$), and the union of the events is the entire sample space.
 10302 The law of total probability states that we can **marginalize** over these events as follows,

$$\Pr(A) = \sum_{B_n \in \mathcal{B}} \Pr(A \cap B_n). \quad [\text{A.10}]$$

10303 For any event B , the union $B \cup \neg B$ is a partition of the sample space. Therefore, a special
 10304 case of the law of total probability is,

$$\Pr(A) = \Pr(A \cap B) + \Pr(A \cap \neg B). \quad [\text{A.11}]$$

10305 A.2 Conditional probability and Bayes' rule

A **conditional probability** is an expression like $\Pr(A \mid B)$, which is the probability of the event A , assuming that event B happens too. For example, we may be interested in the probability of a randomly selected person answering the phone by saying *hello*, conditioned on that person being a speaker of English. Conditional probability is defined as the ratio,

$$\Pr(A \mid B) = \frac{\Pr(A \cap B)}{\Pr(B)}. \quad [\text{A.12}]$$

The **chain rule of probability** states that $\Pr(A \cap B) = \Pr(A \mid B) \times \Pr(B)$, which is just

a rearrangement of terms from Equation A.12. The chain rule can be applied repeatedly:

$$\begin{aligned}\Pr(A \cap B \cap C) &= \Pr(A | B \cap C) \times \Pr(B \cap C) \\ &= \Pr(A | B \cap C) \times \Pr(B | C) \times \Pr(C).\end{aligned}$$

Bayes' rule (sometimes called Bayes' law or Bayes' theorem) gives us a way to convert between $\Pr(A | B)$ and $\Pr(B | A)$. It follows from the definition of conditional probability and the chain rule:

$$\Pr(A | B) = \frac{\Pr(A \cap B)}{\Pr(B)} = \frac{\Pr(B | A) \times \Pr(A)}{\Pr(B)} \quad [\text{A.13}]$$

10306 Each term in Bayes rule has a name, which we will occasionally use:

- 10307 • $\Pr(A)$ is the **prior**, since it is the probability of event A without knowledge about
- 10308 whether B happens or not.
- 10309 • $\Pr(B | A)$ is the **likelihood**, the probability of event B given that event A has oc-
- 10310 curred.
- 10311 • $\Pr(A | B)$ is the **posterior**, the probability of event A with knowledge that B has
- 10312 occurred.

10313 **Example** The classic examples for Bayes' rule involve tests for rare diseases, but Man-
 10314 ning and Schütze (1999) reframe this example in a linguistic setting. Suppose that you are
 10315 interested in a rare syntactic construction, such as *parasitic gaps*, which occur on average
 10316 once in 100,000 sentences. Here is an example of a parasitic gap:

10317 (A.1) *Which class did you attend ... without registering for ...?*

10318 Lana Linguist has developed a complicated pattern matcher that attempts to identify
 10319 sentences with parasitic gaps. It's pretty good, but it's not perfect:

- 10320 • If a sentence has a parasitic gap, the pattern matcher will find it with probability
- 10321 0.95. (This is the **recall**, which is one minus the **false positive rate**.)
- 10322 • If the sentence doesn't have a parasitic gap, the pattern matcher will wrongly say it
- 10323 does with probability 0.005. (This is the **false positive rate**, which is one minus the
- 10324 **precision**.)

10325 Suppose that Lana's pattern matcher says that a sentence contains a parasitic gap. What
 10326 is the probability that this is true?

Let G be the event of a sentence having a parasitic gap, and T be the event of the test being positive. We are interested in the probability of a sentence having a parasitic gap given that the test is positive. This is the conditional probability $\Pr(G | T)$, and it can be computed by Bayes' rule:

$$\Pr(G | T) = \frac{\Pr(T | G) \times \Pr(G)}{\Pr(T)}. \quad [\text{A.14}]$$

10327 We already know both terms in the numerator: $\Pr(T | G)$ is the recall, which is 0.95; $\Pr(G)$
10328 is the prior, which is 10^{-5} .

10329 We are not given the denominator, but it can be computed using tools developed earlier
10330 in this section. First apply the law of total probability, using the partition $\{G, \neg G\}$:

$$\Pr(T) = \Pr(T \cap G) + \Pr(T \cap \neg G). \quad [\text{A.15}]$$

This says that the probability of the test being positive is the sum of the probability of a **true positive** ($T \cap G$) and the probability of a **false positive** ($T \cap \neg G$). The probability of each of these events can be computed using the chain rule:

$$\Pr(T \cap G) = \Pr(T | G) \times \Pr(G) = 0.95 \times 10^{-5} \quad [\text{A.16}]$$

$$\Pr(T \cap \neg G) = \Pr(T | \neg G) \times \Pr(\neg G) = 0.005 \times (1 - 10^{-5}) \approx 0.005 \quad [\text{A.17}]$$

$$\Pr(T) = \Pr(T \cap G) + \Pr(T \cap \neg G) \quad [\text{A.18}]$$

$$= 0.95 \times 10^{-5} + 0.005. \quad [\text{A.19}]$$

Plugging these terms into Bayes' rule gives the desired posterior probability,

$$\Pr(G | T) = \frac{\Pr(T | G) \Pr(G)}{\Pr(T)} \quad [\text{A.20}]$$

$$= \frac{0.95 \times 10^{-5}}{0.95 \times 10^{-5} + 0.005 \times (1 - 10^{-5})} \quad [\text{A.21}]$$

$$\approx 0.002. \quad [\text{A.22}]$$

10331 Lana's pattern matcher seems accurate, with false positive and false negative rates
10332 below 5%. Yet the extreme rarity of the phenomenon means that a positive result from the
10333 detector is most likely to be wrong.

10334 A.3 Independence

Two events are independent if the probability of their intersection is equal to the product of their probabilities: $\Pr(A \cap B) = \Pr(A) \times \Pr(B)$. For example, for two flips of a fair

coin, the probability of getting heads on the first flip is independent of the probability of getting heads on the second flip:

$$\Pr(\{HT, HH\}) = \Pr(HT) + \Pr(HH) = \frac{1}{4} + \frac{1}{4} = \frac{1}{2} \quad [A.23]$$

$$\Pr(\{HH, TH\}) = \Pr(HH) + \Pr(TH) = \frac{1}{4} + \frac{1}{4} = \frac{1}{2} \quad [A.24]$$

$$\Pr(\{HT, HH\}) \times \Pr(\{HH, TH\}) = \frac{1}{2} \times \frac{1}{2} = \frac{1}{4} \quad [A.25]$$

$$\Pr(\{HT, HH\} \cap \{HH, TH\}) = \Pr(HH) = \frac{1}{4} \quad [A.26]$$

$$= \Pr(\{HT, HH\}) \times \Pr(\{HH, TH\}). \quad [A.27]$$

If $\Pr(A \cap B \mid C) = \Pr(A \mid C) \times \Pr(B \mid C)$, then the events A and B are **conditionally independent**, written $A \perp B \mid C$. Conditional independence plays a important role in probabilistic models such as Naïve Bayes chapter 2.

A.4 Random variables

Random variables are functions from events to \mathbb{R}^n , where \mathbb{R} is the set of real numbers. This subsumes several useful special cases:

- An **indicator random variable** is a function from events to the set $\{0, 1\}$. In the coin flip example, we can define Y as an indicator random variable, taking the value 1 when the coin has come up heads on at least one flip. This would include the outcomes $\{HH, HT, TH\}$. The probability $\Pr(Y = 1)$ is the sum of the probabilities of these outcomes, $\Pr(Y = 1) = \frac{1}{4} + \frac{1}{4} + \frac{1}{4} = \frac{3}{4}$.
- A **discrete random variable** is a function from events to a discrete subset of \mathbb{R} . Consider the coin flip example: the number of heads on two flips, X , can be viewed as a discrete random variable, $X \in \{0, 1, 2\}$. The event probability $\Pr(X = 1)$ can again be computed as the sum of the probabilities of the events in which there is one head, $\{HT, TH\}$, giving $\Pr(X = 1) = \frac{1}{4} + \frac{1}{4} = \frac{1}{2}$.

Each possible value of a random variable is associated with a subset of the sample space. In the coin flip example, $X = 0$ is associated with the event $\{TT\}$, $X = 1$ is associated with the event $\{HT, TH\}$, and $X = 2$ is associated with the event $\{HH\}$. Assuming a fair coin, the probabilities of these events are, respectively, $1/4$, $1/2$, and $1/4$. This list of numbers represents the **probability distribution** over X , written p_X , which maps from the possible values of X to the non-negative reals. For a specific value x , we write $p_X(x)$, which is equal to the event probability $\Pr(X = x)$.¹ The function p_X is called

¹In general, capital letters (e.g., X) refer to random variables, and lower-case letters (e.g., x) refer to specific values. When the distribution is clear from context, I will simply write $p(x)$.

a probability **mass** function (pmf) if X is discrete; it is called a probability **density** function (pdf) if X is continuous. In either case, the function must sum to one, and all values must be non-negative:

$$\int_x p_X(x)dx = 1 \quad [A.28]$$

$$\forall x, p_X(x) \geq 0. \quad [A.29]$$

Probabilities over multiple random variables can written as **joint probabilities**, e.g., $p_{A,B}(a,b) = \Pr(A = a \cap B = b)$. Several properties of event probabilities carry over to probability distributions over random variables:

- The **marginal probability distribution** is $p_A(a) = \sum_b p_{A,B}(a,b)$.
- The **conditional probability distribution** is $p_{A|B}(a | b) = \frac{p_{A,B}(a,b)}{p_B(b)}$.
- Random variables A and B are independent iff $p_{A,B}(a,b) = p_A(a) \times p_B(b)$.

A.5 Expectations

Sometimes we want the **expectation** of a function, such as $E[g(x)] = \sum_{x \in \mathcal{X}} g(x)p(x)$. Expectations are easiest to think about in terms of probability distributions over discrete events:

- If it is sunny, Lucia will eat three ice creams.
- If it is rainy, she will eat only one ice cream.
- There's a 80% chance it will be sunny.
- The expected number of ice creams she will eat is $0.8 \times 3 + 0.2 \times 1 = 2.6$.

If the random variable X is continuous, the expectation is an integral:

$$E[g(x)] = \int_{\mathcal{X}} g(x)p(x)dx \quad [A.30]$$

For example, a fast food restaurant in Quebec has a special offer for cold days: they give a 1% discount on poutine for every degree below zero. Assuming a thermometer with infinite precision, the expected price would be an integral over all possible temperatures,

$$E[\text{price}(x)] = \int_{\mathcal{X}} \min(1, 1+x) \times \text{original-price} \times p(x)dx. \quad [A.31]$$

10369 **A.6 Modeling and estimation**

10370 **Probabilistic models** provide a principled way to reason about random events and ran-
10371 dom variables. Let's consider the coin toss example. Each toss can be modeled as a ran-
10372 dom event, with probability θ of the event H , and probability $1 - \theta$ of the complementary
10373 event T . If we write a random variable X as the total number of heads on three coin
10374 flips, then the distribution of X depends on θ . In this case, X is distributed as a **binomial**
10375 **random variable**, meaning that it is drawn from a binomial distribution, with **parameters**
10376 $(\theta, N = 3)$. This is written,

$$X \sim \text{Binomial}(\theta, N = 3). \quad [\text{A.32}]$$

10377 The properties of the binomial distribution enable us to make statements about the X ,
10378 such as its expected value and the likelihood that its value will fall within some interval.

Now suppose that θ is unknown, but we have run an experiment, in which we ex-
 ecuted N trials, and obtained x heads. We can **estimate** θ by the principle of **maximum**
likelihood:

$$\hat{\theta} = \operatorname{argmax}_{\theta} p_X(x; \theta, N). \quad [\text{A.33}]$$

This says that the estimate $\hat{\theta}$ should be the value that maximizes the likelihood of the
 data. The semicolon indicates that θ and N are parameters of the probability function.
 The likelihood $p_X(x; \theta, N)$ can be computed from the binomial distribution,

$$p_X(x; \theta, N) = \frac{N!}{x!(N-x)!} \theta^x (1-\theta)^{N-x}. \quad [\text{A.34}]$$

10379 This likelihood is proportional to the product of the probability of individual out-
10380 comes: for example, the sequence T, H, H, T, H would have probability $\theta^3(1-\theta)^2$. The
10381 term $\frac{N!}{x!(N-x)!}$ arises from the many possible orderings by which we could obtain x heads
10382 on N trials. This term does not depend on θ , so it can be ignored during estimation.

In practice, we maximize the log-likelihood, which is a monotonic function of the like-
 lihood. Under the binomial distribution, the log-likelihood is a **convex** function of θ (see

§ 2.3), so it can be maximized by taking the derivative and setting it equal to zero.

$$\ell(\theta) = x \log \theta + (N - x) \log(1 - \theta) \quad [\text{A.35}]$$

$$\frac{\partial \ell(\theta)}{\partial \theta} = \frac{x}{\theta} - \frac{N - x}{1 - \theta} \quad [\text{A.36}]$$

$$\frac{N - x}{1 - \theta} = \frac{x}{\theta} \quad [\text{A.37}]$$

$$\frac{N - x}{x} = \frac{1 - \theta}{\theta} \quad [\text{A.38}]$$

$$\frac{N}{x} - 1 = \frac{1}{\theta} - 1 \quad [\text{A.39}]$$

$$\hat{\theta} = \frac{x}{N}. \quad [\text{A.40}]$$

10383 In this case, the maximum likelihood estimate is equal to $\frac{x}{N}$, the fraction of trials that
 10384 came up heads. This intuitive solution is also known as the **relative frequency estimate**,
 10385 since it is equal to the relative frequency of the outcome.

Is maximum likelihood estimation always the right choice? Suppose you conduct one trial, and get heads. Would you conclude that $\theta = 1$, meaning that the coin is guaranteed to come up heads? If not, then you must have some **prior expectation** about θ . To incorporate this prior information, we can treat θ as a random variable, and use Bayes' rule:

$$p(\theta | x; N) = \frac{p(x | \theta) \times p(\theta)}{p(x)} \quad [\text{A.41}]$$

$$\propto p(x | \theta) \times p(\theta) \quad [\text{A.42}]$$

$$\hat{\theta} = \operatorname{argmax}_{\theta} p(x | \theta) \times p(\theta). \quad [\text{A.43}]$$

10386 This is the **maximum a posteriori** (MAP) estimate. Given a form for $p(\theta)$, you can de-
 10387 rive the MAP estimate using the same approach that was used to derive the maximum
 10388 likelihood estimate.

10389 Additional resources

10390 A good introduction to probability theory is offered by Manning and Schütze (1999),
 10391 which helped to motivate this section. For more detail, Sharon Goldwater provides an-
 10392 other useful reference, <http://homepages.inf.ed.ac.uk/sgwater/teaching/general/probability.pdf>. A historical and philosophical perspective on probability is offered
 10393 by Diaconis and Skyrms (2017).

10395 **Appendix B**

10396 **Numerical optimization**

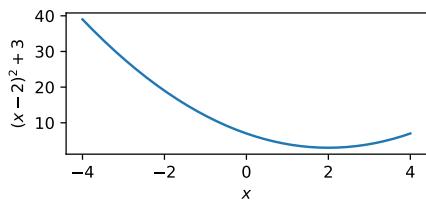
10397 Unconstrained numerical optimization involves solving problems of the form,

$$\min_{\mathbf{x} \in \mathbb{R}^D} f(\mathbf{x}), \quad [\text{B.1}]$$

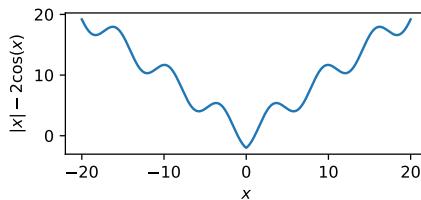
10398 where $\mathbf{x} \in \mathbb{R}^D$ is a vector of D real numbers.

10399 Differentiation is fundamental to continuous optimization. Suppose that at some \mathbf{x}^* ,
10400 every partial derivative is equal to 0: formally, $\frac{\partial f}{\partial x_i}\Big|_{\mathbf{x}^*} = 0$. Then \mathbf{x}^* is said to be a **critical**
10401 **point** of f . For a **convex** function f (defined in § 2.3), $f(\mathbf{x}^*)$ is equal to the global minimum
10402 of f iff \mathbf{x}^* is a critical point of f .

As an example, consider the convex function $f(x) = (x - 2)^2 + 3$, shown in Figure B.1a. The derivative is $\frac{\partial f}{\partial x} = 2x - 4$. A unique minimum can be obtained by setting the derivative equal to zero and solving for x , obtaining $x^* = 2$. Now consider the multivariate convex function $f(\mathbf{x}) = \frac{1}{2}\|\mathbf{x} - [2, 1]^\top\|^2$, where $\|\mathbf{x}\|^2$ is the squared Euclidean norm. The partial



(a) The function $f(x) = (x - 2)^2 + 3$



(b) The function $f(x) = |x| - 2 \cos(x)$

Figure B.1: Two functions with unique global minima

derivatives are,

$$\frac{\partial d}{\partial x_1} = x_1 - 2 \quad [B.2]$$

$$\frac{\partial d}{\partial x_2} = x_2 - 1 \quad [B.3]$$

10403 The unique minimum is $\mathbf{x}^* = [2, 1]^\top$.

10404 For non-convex functions, critical points are not necessarily global minima. A **local**
 10405 **minimum** \mathbf{x}^* is a point at which the function takes a smaller value than at all nearby
 10406 neighbors: formally, \mathbf{x}^* is a local minimum if there is some positive ϵ such that $f(\mathbf{x}^*) \leq$
 10407 $f(\mathbf{x})$ for all \mathbf{x} within distance ϵ of \mathbf{x}^* . Figure B.1b shows the function $f(x) = |x| - 2 \cos(x)$,
 10408 which has many local minima, as well as a unique global minimum at $x = 0$. A critical
 10409 point may also be the local or global maximum of the function; it may be a **saddle point**,
 10410 which is a minimum with respect to at least one coordinate, and a maximum with respect
 10411 at least one other coordinate; it may be an **inflection point**, which is neither a minimum
 10412 nor maximum. When available, the second derivative of f can help to distinguish these
 10413 cases.

10414 B.1 Gradient descent

For many convex functions, it is not possible to solve for \mathbf{x}^* in closed form. In gradient descent, we compute a series of solutions, $\mathbf{x}^{(0)}, \mathbf{x}^{(1)}, \dots$ by taking steps along the local gradient $\nabla_{\mathbf{x}^{(t)}} f$, which is the vector of partial derivatives of the function f , evaluated at the point $\mathbf{x}^{(t)}$. Each solution $\mathbf{x}^{(t+1)}$ is computed,

$$\mathbf{x}^{(t+1)} \leftarrow \mathbf{x}^{(t)} - \eta^{(t)} \nabla_{\mathbf{x}^{(t)}} f. \quad [B.4]$$

10415 where $\eta^{(t)} > 0$ is a **step size**. If the step size is chosen appropriately, this procedure will
 10416 find the global minimum of a differentiable convex function. For non-convex functions,
 10417 gradient descent will find a local minimum. The extension to non-differentiable convex
 10418 functions is discussed in § 2.3.

10419 B.2 Constrained optimization

Optimization must often be performed under constraints: for example, when optimizing the parameters of a probability distribution, the probabilities of all events must sum to one. Constrained optimization problems can be written,

$$\min_{\mathbf{x}} f(\mathbf{x}) \quad [B.5]$$

$$\text{s.t. } g_c(\mathbf{x}) \leq 0, \quad \forall c = 1, 2, \dots, C \quad [B.6]$$

where each $g_i(\mathbf{x})$ is a scalar function of \mathbf{x} . For example, suppose that \mathbf{x} must be non-negative, and that its sum cannot exceed a budget b . Then there are $D + 1$ inequality constraints,

$$g_i(\mathbf{x}) = -x_i, \quad \forall i = 1, 2, \dots, D \quad [\text{B.7}]$$

$$g_{D+1}(\mathbf{x}) = -b + \sum_{i=1}^D x_i. \quad [\text{B.8}]$$

Inequality constraints can be combined with the original objective function f by forming a **Lagrangian**,

$$L(\mathbf{x}, \boldsymbol{\lambda}) = f(\mathbf{x}) + \sum_{c=1}^C \lambda_c g_c(\mathbf{x}), \quad [\text{B.9}]$$

where λ_c is a **Lagrange multiplier**. For any Lagrangian, there is a corresponding **dual form**, which is a function of $\boldsymbol{\lambda}$:

$$D(\boldsymbol{\lambda}) = \min_{\mathbf{x}} L(\mathbf{x}, \boldsymbol{\lambda}). \quad [\text{B.10}]$$

The Lagrangian L can be referred to as the **primal form**.

B.3 Example: Passive-aggressive online learning

Sometimes it is possible to solve a constrained optimization problem by manipulating the Lagrangian. One example is maximum-likelihood estimation of a Naïve Bayes probability model, as described in § 2.1.3. In that case, it is unnecessary to explicitly compute the Lagrange multiplier. Another example is illustrated by the **passive-aggressive** algorithm for online learning (Crammer et al., 2006). This algorithm is similar to the perceptron, but the goal at each step is to make the most conservative update that gives zero margin loss on the current example.¹ Each update can be formulated as a constrained optimization over the weights $\boldsymbol{\theta}$:

$$\min_{\boldsymbol{\theta}} \frac{1}{2} \|\boldsymbol{\theta} - \boldsymbol{\theta}^{(i-1)}\|^2 \quad [\text{B.11}]$$

$$\text{s.t. } \ell^{(i)}(\boldsymbol{\theta}) = 0 \quad [\text{B.12}]$$

where $\boldsymbol{\theta}^{(i-1)}$ is the previous set of weights, and $\ell^{(i)}(\boldsymbol{\theta})$ is the margin loss on instance i . As in § 2.3.1, this loss is defined as,

$$\ell^{(i)}(\boldsymbol{\theta}) = 1 - \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y^{(i)}) + \max_{y \neq y^{(i)}} \boldsymbol{\theta} \cdot \mathbf{f}(\mathbf{x}^{(i)}, y). \quad [\text{B.13}]$$

¹This is the basis for the name of the algorithm: it is passive when the loss is zero, but it aggressively moves to make the loss zero when necessary.

When the margin loss is zero for $\theta^{(i-1)}$, the optimal solution is simply to set $\theta^* = \theta^{(i-1)}$, so we will focus on the case where $\ell^{(i)}(\theta^{(i-1)}) > 0$. The Lagrangian for this problem is,

$$L(\theta, \lambda) = \frac{1}{2} \|\theta - \theta^{(i-1)}\|^2 + \lambda \ell^{(i)}(\theta), \quad [\text{B.14}]$$

Holding λ constant, we can solve for θ by differentiating,

$$\nabla_{\theta} L = \theta - \theta^{(i-1)} + \lambda \frac{\partial}{\partial \theta} \ell^{(i)}(\theta) \quad [\text{B.15}]$$

$$\theta^* = \theta^{(i-1)} + \lambda \delta, \quad [\text{B.16}]$$

where $\delta = f(x^{(i)}, y^{(i)}) - f(x^{(i)}, \hat{y})$ and $\hat{y} = \operatorname{argmax}_{y \neq y^{(i)}} \theta \cdot f(x^{(i)}, y)$.

The Lagrange multiplier λ acts as the learning rate in a perceptron-style update to θ . We can solve for λ by plugging θ^* back into the Lagrangian, obtaining the dual function,

$$D(\lambda) = \frac{1}{2} \|\theta^{(i-1)} + \lambda \delta - \theta^{(i-1)}\|^2 + \lambda(1 - (\theta^{(i-1)} + \lambda \delta) \cdot \delta) \quad [\text{B.17}]$$

$$= \frac{\lambda^2}{2} \|\delta\|^2 - \lambda^2 \|\delta\|^2 + \lambda(1 - \theta^{(i-1)} \cdot \delta) \quad [\text{B.18}]$$

$$= -\frac{\lambda^2}{2} \|\delta\|^2 + \lambda \ell^{(i)}(\theta^{(i-1)}). \quad [\text{B.19}]$$

Differentiating and solving for λ ,

$$\frac{\partial D}{\partial \lambda} = -\lambda \|\delta\|^2 + \ell^{(i)}(\theta^{(i-1)}) \quad [\text{B.20}]$$

$$\lambda^* = \frac{\ell^{(i)}(\theta^{(i-1)})}{\|\delta\|^2}. \quad [\text{B.21}]$$

The complete update equation is therefore:

$$\theta^* = \theta^{(i-1)} + \frac{\ell^{(i)}(\theta^{(i-1)})}{\|f(x^{(i)}, y^{(i)}) - f(x^{(i)}, \hat{y})\|^2} (f(x^{(i)}, y^{(i)}) - f(x^{(i)}, \hat{y})). \quad [\text{B.22}]$$

This update has strong intuitive support. The numerator of the learning rate grows with the loss. The denominator grows with the norm of the difference between the feature vectors associated with the correct and predicted label. If this norm is large, then the step with respect to each feature should be small, and vice versa.

10437

Bibliography

- 10438 Abadi, M., A. Agarwal, P. Barham, E. Brevdo, Z. Chen, C. Citro, G. S. Corrado, A. Davis,
10439 J. Dean, M. Devin, S. Ghemawat, I. J. Goodfellow, A. Harp, G. Irving, M. Isard, Y. Jia,
10440 R. Józefowicz, L. Kaiser, M. Kudlur, J. Levenberg, D. Mané, R. Monga, S. Moore,
10441 D. G. Murray, C. Olah, M. Schuster, J. Shlens, B. Steiner, I. Sutskever, K. Talwar, P. A.
10442 Tucker, V. Vanhoucke, V. Vasudevan, F. B. Viégas, O. Vinyals, P. Warden, M. Watten-
10443 berg, M. Wicke, Y. Yu, and X. Zheng (2016). Tensorflow: Large-scale machine learning
10444 on heterogeneous distributed systems. *CoRR abs/1603.04467*.
- 10445 Abend, O. and A. Rappoport (2017). The state of the art in semantic representation. In
10446 *Proceedings of the Association for Computational Linguistics (ACL)*.
- 10447 Abney, S., R. E. Schapire, and Y. Singer (1999). Boosting applied to tagging and PP attach-
10448 ment. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pp.
10449 132–134.
- 10450 Abney, S. P. (1987). *The English noun phrase in its sentential aspect*. Ph. D. thesis, Mas-
10451 sachusetts Institute of Technology.
- 10452 Abney, S. P. and M. Johnson (1991). Memory requirements and local ambiguities of pars-
10453 ing strategies. *Journal of Psycholinguistic Research* 20(3), 233–250.
- 10454 Adafre, S. F. and M. De Rijke (2006). Finding similar sentences across multiple languages
10455 in wikipedia. In *Proceedings of the Workshop on NEW TEXT Wikis and blogs and other*
10456 *dynamic text sources*.
- 10457 Ahn, D. (2006). The stages of event extraction. In *Proceedings of the Workshop on Annotating*
10458 *and Reasoning about Time and Events*, pp. 1–8. Association for Computational Linguistics.
- 10459 Aho, A. V., M. S. Lam, R. Sethi, and J. D. Ullman (2006). Compilers: Principles, techniques,
10460 & tools.
- 10461 Aikhenvald, A. Y. (2004). *Evidentiality*. Oxford University Press.

- 10462 Akaike, H. (1974). A new look at the statistical model identification. *IEEE Transactions on
10463 Automatic Control* 19(6), 716–723.
- 10464 Akmajian, A., R. A. Demers, A. K. Farmer, and R. M. Harnish (2010). *Linguistics: An
10465 introduction to language and communication* (Sixth ed.). Cambridge, MA: MIT press.
- 10466 Alfau, F. (1999). *Chromos*. Dalkey Archive Press.
- 10467 Allauzen, C., M. Riley, J. Schalkwyk, W. Skut, and M. Mohri (2007). OpenFst: A gen-
10468 eral and efficient weighted finite-state transducer library. In *International Conference on
10469 Implementation and Application of Automata*, pp. 11–23. Springer.
- 10470 Allen, J. F. (1984). Towards a general theory of action and time. *Artificial intelligence* 23(2),
10471 123–154.
- 10472 Allen, J. F., B. W. Miller, E. K. Ringger, and T. Sikorski (1996). A robust system for natural
10473 spoken dialogue. In *Proceedings of the Association for Computational Linguistics (ACL)*, pp.
10474 62–70.
- 10475 Allen, J. F., L. K. Schubert, G. Ferguson, P. Heeman, C. H. Hwang, T. Kato, M. Light,
10476 N. Martin, B. Miller, M. Poesio, and D. Traum (1995). The TRAINS project: A case
10477 study in building a conversational planning agent. *Journal of Experimental & Theoretical
10478 Artificial Intelligence* 7(1), 7–48.
- 10479 Alm, C. O., D. Roth, and R. Sproat (2005). Emotions from text: machine learning for
10480 text-based emotion prediction. In *Proceedings of Empirical Methods for Natural Language
10481 Processing (EMNLP)*, pp. 579–586.
- 10482 Aluísio, S., J. Pelizzoni, A. Marchi, L. de Oliveira, R. Manenti, and V. Marquiafável (2003).
10483 An account of the challenge of tagging a reference corpus for Brazilian Portuguese.
10484 *Computational Processing of the Portuguese Language*, 194–194.
- 10485 Anand, P., M. Walker, R. Abbott, J. E. Fox Tree, R. Bowman, and M. Minor (2011). Cats rule
10486 and dogs drool!: Classifying stance in online debate. In *Proceedings of the 2nd Workshop
10487 on Computational Approaches to Subjectivity and Sentiment Analysis*, Portland, Oregon, pp.
10488 1–9. Association for Computational Linguistics.
- 10489 Anandkumar, A. and R. Ge (2016). Efficient approaches for escaping higher order saddle
10490 points in non-convex optimization. In *Proceedings of the Conference On Learning Theory
10491 (COLT)*, pp. 81–102.
- 10492 Anandkumar, A., R. Ge, D. Hsu, S. M. Kakade, and M. Telgarsky (2014). Tensor decompo-
10493 sitions for learning latent variable models. *The Journal of Machine Learning Research* 15(1),
10494 2773–2832.

- 10495 Ando, R. K. and T. Zhang (2005). A framework for learning predictive structures from
10496 multiple tasks and unlabeled data. *The Journal of Machine Learning Research* 6, 1817–
10497 1853.
- 10498 Andor, D., C. Alberti, D. Weiss, A. Severyn, A. Presta, K. Ganchev, S. Petrov, and
10499 M. Collins (2016). Globally normalized transition-based neural networks. In *Proceedings*
10500 of the Association for Computational Linguistics (ACL), pp. 2442–2452.
- 10501 Angeli, G., P. Liang, and D. Klein (2010). A simple domain-independent probabilistic ap-
10502 proach to generation. In *Proceedings of Empirical Methods for Natural Language Processing*
10503 (EMNLP), pp. 502–512.
- 10504 Antol, S., A. Agrawal, J. Lu, M. Mitchell, D. Batra, C. Lawrence Zitnick, and D. Parikh
10505 (2015). Vqa: Visual question answering. In *Proceedings of the International Conference on*
10506 *Computer Vision (ICCV)*, pp. 2425–2433.
- 10507 Aronoff, M. (1976). *Word formation in generative grammar*. MIT Press.
- 10508 Arora, S. and B. Barak (2009). *Computational complexity: a modern approach*. Cambridge
10509 University Press.
- 10510 Arora, S., R. Ge, Y. Halpern, D. Mimmo, A. Moitra, D. Sontag, Y. Wu, and M. Zhu (2013).
10511 A practical algorithm for topic modeling with provable guarantees. In *Proceedings of the*
10512 *International Conference on Machine Learning (ICML)*, pp. 280–288.
- 10513 Arora, S., Y. Li, Y. Liang, T. Ma, and A. Risteski (2016). Linear algebraic structure of word
10514 senses, with applications to polysemy. *arXiv preprint arXiv:1601.03764*.
- 10515 Artstein, R. and M. Poesio (2008). Inter-coder agreement for computational linguistics.
10516 *Computational Linguistics* 34(4), 555–596.
- 10517 Artzi, Y. and L. Zettlemoyer (2013). Weakly supervised learning of semantic parsers for
10518 mapping instructions to actions. *Transactions of the Association for Computational Linguis-*
10519 *tics* 1, 49–62.
- 10520 Attardi, G. (2006). Experiments with a multilanguage non-projective dependency parser.
10521 In *Proceedings of the Conference on Natural Language Learning (CoNLL)*, pp. 166–170.
- 10522 Auer, P. (2013). *Code-switching in conversation: Language, interaction and identity*. Routledge.
- 10523 Auer, S., C. Bizer, G. Kobilarov, J. Lehmann, R. Cyganiak, and Z. Ives (2007). Dbpedia: A
10524 nucleus for a web of open data. *The semantic web*, 722–735.
- 10525 Austin, J. L. (1962). *How to do things with words*. Oxford University Press.

- 10526 Aw, A., M. Zhang, J. Xiao, and J. Su (2006). A phrase-based statistical model for SMS text
 10527 normalization. In *Proceedings of the Association for Computational Linguistics (ACL)*, pp.
 10528 33–40.
- 10529 Ba, J. L., J. R. Kiros, and G. E. Hinton (2016). Layer normalization. *arXiv preprint arXiv:1607.06450*.
- 10531 Bagga, A. and B. Baldwin (1998a). Algorithms for scoring coreference chains. In *Proceedings of the Language Resources and Evaluation Conference*, pp. 563–566.
- 10533 Bagga, A. and B. Baldwin (1998b). Entity-based cross-document coreferencing using the
 10534 vector space model. In *Proceedings of the International Conference on Computational Lin-
 10535 guistics (COLING)*, pp. 79–85.
- 10536 Bahdanau, D., K. Cho, and Y. Bengio (2014). Neural machine translation by jointly learn-
 10537 ing to align and translate. In *Neural Information Processing Systems (NIPS)*.
- 10538 Baldwin, T. and S. N. Kim (2010). Multiword expressions. In *Handbook of natural language
 10539 processing*, Volume 2, pp. 267–292. Boca Raton, USA: CRC Press.
- 10540 Balle, B., A. Quattoni, and X. Carreras (2011). A spectral learning algorithm for finite state
 10541 transducers. In *Proceedings of the European Conference on Machine Learning and Principles
 10542 and Practice of Knowledge Discovery in Databases (ECML)*, pp. 156–171.
- 10543 Banarescu, L., C. Bonial, S. Cai, M. Georgescu, K. Griffitt, U. Hermjakob, K. Knight,
 10544 P. Koehn, M. Palmer, and N. Schneider (2013, August). Abstract meaning represen-
 10545 tation for sembanking. In *Proceedings of the 7th Linguistic Annotation Workshop and Interoperability
 10546 with Discourse*, Sofia, Bulgaria, pp. 178–186. Association for Computational
 10547 Linguistics.
- 10548 Banko, M., M. J. Cafarella, S. Soderland, M. Broadhead, and O. Etzioni (2007). Open
 10549 information extraction from the web. In *Proceedings of the International Joint Conference
 10550 on Artificial Intelligence (IJCAI)*, pp. 2670–2676.
- 10551 Bansal, N., A. Blum, and S. Chawla (2004). Correlation clustering. *Machine Learning* 56(1–
 10552 3), 89–113.
- 10553 Barber, D. (2012). *Bayesian reasoning and machine learning*. Cambridge University Press.
- 10554 Barman, U., A. Das, J. Wagner, and J. Foster (2014, October). Code mixing: A challenge for
 10555 language identification in the language of social media. In *Proceedings of the First Work-
 10556 shop on Computational Approaches to Code Switching*, Doha, Qatar, pp. 13–23. Association
 10557 for Computational Linguistics.

- 10558 Barnickel, T., J. Weston, R. Collobert, H.-W. Mewes, and V. Stümpflen (2009). Large scale
10559 application of neural network based semantic role labeling for automated relation ex-
10560 traction from biomedical texts. *PLoS One* 4(7), e6393.
- 10561 Baron, A. and P. Rayson (2008). Vard2: A tool for dealing with spelling variation in his-
10562 torical corpora. In *Postgraduate conference in corpus linguistics*.
- 10563 Baroni, M., R. Bernardi, and R. Zamparelli (2014). Frege in space: A program for compo-
10564 sitional distributional semantics. *Linguistic Issues in Language Technologies*.
- 10565 Barzilay, R. and M. Lapata (2008, mar). Modeling local coherence: An Entity-Based ap-
10566 proach. *Computational Linguistics* 34(1), 1–34.
- 10567 Barzilay, R. and K. R. McKeown (2005). Sentence fusion for multidocument news summa-
10568 rization. *Computational Linguistics* 31(3), 297–328.
- 10569 Beesley, K. R. and L. Karttunen (2003). *Finite-state morphology*. Stanford, CA: Center for
10570 the Study of Language and Information.
- 10571 Bejan, C. A. and S. Harabagiu (2014). Unsupervised event coreference resolution. *Compu-
10572 tational Linguistics* 40(2), 311–347.
- 10573 Bell, E. T. (1934). Exponential numbers. *The American Mathematical Monthly* 41(7), 411–419.
- 10574 Bender, E. M. (2013, jun). *Linguistic Fundamentals for Natural Language Processing: 100
10575 Essentials from Morphology and Syntax*, Volume 6 of *Synthesis Lectures on Human Language
10576 Technologies*. Morgan & Claypool Publishers.
- 10577 Bengio, S., O. Vinyals, N. Jaitly, and N. Shazeer (2015). Scheduled sampling for sequence
10578 prediction with recurrent neural networks. In *Neural Information Processing Systems
10579 (NIPS)*, pp. 1171–1179.
- 10580 Bengio, Y., R. Ducharme, P. Vincent, and C. Janvin (2003). A neural probabilistic language
10581 model. *The Journal of Machine Learning Research* 3, 1137–1155.
- 10582 Bengio, Y., P. Simard, and P. Frasconi (1994). Learning long-term dependencies with gra-
10583 dient descent is difficult. *IEEE Transactions on Neural Networks* 5(2), 157–166.
- 10584 Bengtsson, E. and D. Roth (2008). Understanding the value of features for coreference
10585 resolution. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*,
10586 pp. 294–303.
- 10587 Benjamini, Y. and Y. Hochberg (1995). Controlling the false discovery rate: a practical and
10588 powerful approach to multiple testing. *Journal of the Royal Statistical Society. Series B
10589 (Methodological)*, 289–300.

- 10590 Berant, J., A. Chou, R. Frostig, and P. Liang (2013). Semantic parsing on freebase from
 10591 question-answer pairs. In *Proceedings of Empirical Methods for Natural Language Processing*
 10592 (EMNLP), pp. 1533–1544.
- 10593 Berant, J., V. Srikumar, P.-C. Chen, A. Vander Linden, B. Harding, B. Huang, P. Clark, and
 10594 C. D. Manning (2014). Modeling biological processes for reading comprehension. In
 10595 *Proceedings of Empirical Methods for Natural Language Processing* (EMNLP).
- 10596 Berg-Kirkpatrick, T., A. Bouchard-Côté, J. DeNero, and D. Klein (2010). Painless unsuper-
 10597 vised learning with features. In *Proceedings of the North American Chapter of the Associa-*
 10598 *tion for Computational Linguistics* (NAACL), pp. 582–590.
- 10599 Berg-Kirkpatrick, T., D. Burkett, and D. Klein (2012). An empirical investigation of sta-
 10600 tistical significance in NLP. In *Proceedings of Empirical Methods for Natural Language*
 10601 *Processing* (EMNLP), pp. 995–1005.
- 10602 Berger, A. L., V. J. D. Pietra, and S. A. D. Pietra (1996). A maximum entropy approach to
 10603 natural language processing. *Computational linguistics* 22(1), 39–71.
- 10604 Bergsma, S., D. Lin, and R. Goebel (2008). Distributional identification of non-referential
 10605 pronouns. In *Proceedings of the Association for Computational Linguistics* (ACL), pp. 10–18.
- 10606 Bernardi, R., R. Cakici, D. Elliott, A. Erdem, E. Erdem, N. Ikizler-Cinbis, F. Keller, A. Mus-
 10607 cat, and B. Plank (2016). Automatic description generation from images: A survey of
 10608 models, datasets, and evaluation measures. *Journal of Artificial Intelligence Research* 55,
 10609 409–442.
- 10610 Bertsekas, D. P. (2012). Incremental gradient, subgradient, and proximal methods for
 10611 convex optimization: A survey. See Sra et al. (2012).
- 10612 Bhatia, P., R. Guthrie, and J. Eisenstein (2016). Morphological priors for probabilistic neu-
 10613 ral word embeddings. In *Proceedings of Empirical Methods for Natural Language Processing*
 10614 (EMNLP).
- 10615 Bhatia, P., Y. Ji, and J. Eisenstein (2015). Better document-level sentiment analysis from
 10616 first discourse parsing. In *Proceedings of Empirical Methods for Natural Language Processing*
 10617 (EMNLP).
- 10618 Biber, D. (1991). *Variation across speech and writing*. Cambridge University Press.
- 10619 Bird, S., E. Klein, and E. Loper (2009). *Natural language processing with Python*. California:
 10620 O'Reilly Media.
- 10621 Bishop, C. M. (2006). *Pattern recognition and machine learning*. Springer.

- 10622 Björkelund, A. and P. Nugues (2011). Exploring lexicalized features for coreference reso-
10623 lution. In *Proceedings of the Conference on Natural Language Learning (CoNLL)*, pp. 45–50.
- 10624 Blackburn, P. and J. Bos (2005). *Representation and inference for natural language: A first*
10625 *course in computational semantics*. CSLI.
- 10626 Blei, D. M. (2012). Probabilistic topic models. *Communications of the ACM* 55(4), 77–84.
- 10627 Blei, D. M. (2014). Build, compute, critique, repeat: Data analysis with latent variable
10628 models. *Annual Review of Statistics and Its Application* 1, 203–232.
- 10629 Blei, D. M., A. Y. Ng, and M. I. Jordan (2003). Latent dirichlet allocation. *the Journal of*
10630 *machine Learning research* 3, 993–1022.
- 10631 Blitzer, J., M. Dredze, and F. Pereira (2007). Biographies, bollywood, boom-boxes and
10632 blenders: Domain adaptation for sentiment classification. In *Proceedings of the Associa-*
10633 *tion for Computational Linguistics (ACL)*, pp. 440–447.
- 10634 Blum, A. and T. Mitchell (1998). Combining labeled and unlabeled data with co-training.
10635 In *Proceedings of the Conference On Learning Theory (COLT)*, pp. 92–100.
- 10636 Bobrow, D. G., R. M. Kaplan, M. Kay, D. A. Norman, H. Thompson, and T. Winograd
10637 (1977). Gus, a frame-driven dialog system. *Artificial intelligence* 8(2), 155–173.
- 10638 Bochnet, B. (2010). Very high accuracy and fast dependency parsing is not a contradiction.
10639 In *Proceedings of the International Conference on Computational Linguistics (COLING)*, pp.
10640 89–97.
- 10641 Boitet, C. (1988). Pros and cons of the pivot and transfer approaches in multilingual ma-
10642 chine translation. *Readings in machine translation*, 273–279.
- 10643 Bojanowski, P., E. Grave, A. Joulin, and T. Mikolov (2017). Enriching word vectors with
10644 subword information. *Transactions of the Association for Computational Linguistics* 5, 135–
10645 146.
- 10646 Bollacker, K., C. Evans, P. Paritosh, T. Sturge, and J. Taylor (2008). Freebase: a collabora-
10647 tively created graph database for structuring human knowledge. In *Proceedings of the*
10648 *ACM International Conference on Management of Data (SIGMOD)*, pp. 1247–1250. AcM.
- 10649 Bolukbasi, T., K.-W. Chang, J. Y. Zou, V. Saligrama, and A. T. Kalai (2016). Man is to
10650 computer programmer as woman is to homemaker? debiasing word embeddings. In *Neural Information Processing Systems (NIPS)*, pp. 4349–4357.
- 10652 Bordes, A., N. Usunier, A. Garcia-Duran, J. Weston, and O. Yakhnenko (2013). Translating
10653 embeddings for modeling multi-relational data. In *Neural Information Processing Systems*
10654 (*NIPS*), pp. 2787–2795.

- 10655 Bordes, A., J. Weston, R. Collobert, Y. Bengio, et al. (2011). Learning structured embeddings of knowledge bases. In *Proceedings of the National Conference on Artificial Intelligence (AAAI)*, pp. 301–306.
- 10658 Borges, J. L. (1993). *Other Inquisitions 1937–1952*. University of Texas Press. Translated by Ruth L. C. Simms.
- 10660 Botha, J. A. and P. Blunsom (2014). Compositional morphology for word representations and language modelling. In *Proceedings of the International Conference on Machine Learning (ICML)*.
- 10663 Bottou, L. (2012). Stochastic gradient descent tricks. In *Neural networks: Tricks of the trade*, pp. 421–436. Springer.
- 10665 Bottou, L., F. E. Curtis, and J. Nocedal (2016). Optimization methods for large-scale machine learning. *arXiv preprint arXiv:1606.04838*.
- 10667 Bowman, S. R., L. Vilnis, O. Vinyals, A. Dai, R. Jozefowicz, and S. Bengio (2016). Generating sentences from a continuous space. In *Proceedings of the Conference on Natural Language Learning (CoNLL)*, pp. 10–21.
- 10670 boyd, d. and K. Crawford (2012). Critical questions for big data. *Information, Communication & Society* 15(5), 662–679.
- 10672 Boyd, S. and L. Vandenberghe (2004). *Convex Optimization*. New York: Cambridge University Press.
- 10674 Branavan, S., H. Chen, J. Eisenstein, and R. Barzilay (2009). Learning document-level semantic properties from free-text annotations. *Journal of Artificial Intelligence Research* 34(2), 569–603.
- 10677 Branavan, S. R., H. Chen, L. S. Zettlemoyer, and R. Barzilay (2009). Reinforcement learning for mapping instructions to actions. In *Proceedings of the Association for Computational Linguistics (ACL)*, pp. 82–90.
- 10680 Braud, C., O. Lacroix, and A. Søgaard (2017). Does syntax help discourse segmentation? not so much. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pp. 2432–2442.
- 10683 Briscoe, T. (2011). Introduction to formal semantics for natural language.
- 10684 Brown, P. F., J. Cocke, S. A. D. Pietra, V. J. D. Pietra, F. Jelinek, J. D. Lafferty, R. L. Mercer, and P. S. Roossin (1990). A statistical approach to machine translation. *Computational linguistics* 16(2), 79–85.

- 10687 Brown, P. F., P. V. Desouza, R. L. Mercer, V. J. D. Pietra, and J. C. Lai (1992). Class-based
10688 n-gram models of natural language. *Computational linguistics* 18(4), 467–479.
- 10689 Brown, P. F., V. J. D. Pietra, S. A. D. Pietra, and R. L. Mercer (1993). The mathematics
10690 of statistical machine translation: Parameter estimation. *Computational linguistics* 19(2),
10691 263–311.
- 10692 Brun, C. and C. Roux (2014). Décomposition des “hash tags” pour l’amélioration de la
10693 classification en polarité des “tweets”. *Proceedings of Traitement Automatique des Langues
10694 Naturelles*, 473–478.
- 10695 Bruni, E., N.-K. Tran, and M. Baroni (2014). Multimodal distributional semantics. *Journal
10696 of Artificial Intelligence Research* 49(2014), 1–47.
- 10697 Bullinaria, J. A. and J. P. Levy (2007). Extracting semantic representations from word co-
10698 occurrence statistics: A computational study. *Behavior research methods* 39(3), 510–526.
- 10699 Bunescu, R. C. and R. J. Mooney (2005). A shortest path dependency kernel for relation
10700 extraction. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*,
10701 pp. 724–731.
- 10702 Bunescu, R. C. and M. Pasca (2006). Using encyclopedic knowledge for named entity
10703 disambiguation. In *Proceedings of the European Chapter of the Association for Computational
10704 Linguistics (EACL)*, pp. 9–16.
- 10705 Burstein, J., D. Marcu, and K. Knight (2003). Finding the WRITE stuff: Automatic identi-
10706 fication of discourse structure in student essays. *IEEE Intelligent Systems* 18(1), 32–39.
- 10707 Burstein, J., J. Tetreault, and S. Andreyev (2010). Using entity-based features to model
10708 coherence in student essays. In *Human language technologies: The 2010 annual conference
10709 of the North American chapter of the Association for Computational Linguistics*, pp. 681–684.
10710 Association for Computational Linguistics.
- 10711 Burstein, J., J. Tetreault, and M. Chodorow (2013). Holistic discourse coherence annotation
10712 for noisy essay writing. *Dialogue & Discourse* 4(2), 34–52.
- 10713 Cai, Q. and A. Yates (2013). Large-scale semantic parsing via schema matching and lexicon
10714 extension. In *Proceedings of the Association for Computational Linguistics (ACL)*, pp. 423–
10715 433.
- 10716 Caliskan, A., J. J. Bryson, and A. Narayanan (2017). Semantics derived automatically from
10717 language corpora contain human-like biases. *Science* 356(6334), 183–186.
- 10718 Canny, J. (1987). A computational approach to edge detection. In *Readings in Computer
10719 Vision*, pp. 184–203. Elsevier.

- 10720 Cappé, O. and E. Moulines (2009). On-line expectation–maximization algorithm for latent
10721 data models. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)* 71(3),
10722 593–613.
- 10723 Carbonell, J. and J. Goldstein (1998). The use of mmr, diversity-based reranking for re-
10724 ordering documents and producing summaries. In *Proceedings of ACM SIGIR conference*
10725 on Research and development in information retrieval, pp. 335–336.
- 10726 Carbonell, J. R. (1970). Mixed-initiative man-computer instructional dialogues. Technical
10727 report, BOLT BERANEK AND NEWMAN INC CAMBRIDGE MASS.
- 10728 Cardie, C. and K. Wagstaff (1999). Noun phrase coreference as clustering. In *Proceedings*
10729 of *Empirical Methods for Natural Language Processing (EMNLP)*, pp. 82–89.
- 10730 Carletta, J. (1996). Assessing agreement on classification tasks: the kappa statistic. *Com-
10731 putational linguistics* 22(2), 249–254.
- 10732 Carletta, J. (2007). Unleashing the killer corpus: experiences in creating the multi-
10733 everything ami meeting corpus. *Language Resources and Evaluation* 41(2), 181–190.
- 10734 Carlson, L. and D. Marcu (2001). Discourse tagging reference manual. Technical Report
10735 ISI-TR-545, Information Sciences Institute.
- 10736 Carlson, L., M. E. Okurowski, and D. Marcu (2002). RST discourse treebank. Linguistic
10737 Data Consortium, University of Pennsylvania.
- 10738 Carpenter, B. (1997). *Type-logical semantics*. Cambridge, MA: MIT Press.
- 10739 Carreras, X., M. Collins, and T. Koo (2008). Tag, dynamic programming, and the percep-
10740 tron for efficient, feature-rich parsing. In *Proceedings of the Conference on Natural Language*
10741 *Learning (CoNLL)*, pp. 9–16.
- 10742 Carreras, X. and L. Màrquez (2005). Introduction to the conll-2005 shared task: Semantic
10743 role labeling. In *Proceedings of the Ninth Conference on Computational Natural Language*
10744 *Learning*, pp. 152–164. Association for Computational Linguistics.
- 10745 Carroll, L. (1917). *Through the looking glass: And what Alice found there*. Chicago: Rand,
10746 McNally.
- 10747 Chambers, N. and D. Jurafsky (2008). Jointly combining implicit constraints improves
10748 temporal ordering. In *Proceedings of Empirical Methods for Natural Language Processing*
10749 (*EMNLP*), pp. 698–706.
- 10750 Chang, K.-W., A. Krishnamurthy, A. Agarwal, H. Daume III, and J. Langford (2015).
10751 Learning to search better than your teacher. In *Proceedings of the International Confer-
10752 ence on Machine Learning (ICML)*.

- 10753 Chang, M.-W., L. Ratinov, and D. Roth (2007). Guiding semi-supervision with constraint-
10754 driven learning. In *Proceedings of the Association for Computational Linguistics (ACL)*, pp.
10755 280–287.
- 10756 Chang, M.-W., L.-A. Ratinov, N. Rizzolo, and D. Roth (2008). Learning and inference with
10757 constraints. In *Proceedings of the National Conference on Artificial Intelligence (AAAI)*, pp.
10758 1513–1518.
- 10759 Chapman, W. W., W. Bridewell, P. Hanbury, G. F. Cooper, and B. G. Buchanan (2001). A
10760 simple algorithm for identifying negated findings and diseases in discharge summaries.
10761 *Journal of biomedical informatics* 34(5), 301–310.
- 10762 Charniak, E. (1997). Statistical techniques for natural language parsing. *AI magazine* 18(4),
10763 33–43.
- 10764 Charniak, E. and M. Johnson (2005). Coarse-to-fine n-best parsing and maxent discrimi-
10765 native reranking. In *Proceedings of the Association for Computational Linguistics (ACL)*, pp.
10766 173–180.
- 10767 Chelba, C. and A. Acero (2006). Adaptation of maximum entropy capitalizer: Little data
10768 can help a lot. *Computer Speech & Language* 20(4), 382–399.
- 10769 Chelba, C., T. Mikolov, M. Schuster, Q. Ge, T. Brants, P. Koehn, and T. Robinson (2013).
10770 One billion word benchmark for measuring progress in statistical language modeling.
10771 *arXiv preprint arXiv:1312.3005*.
- 10772 Chen, D., J. Bolton, and C. D. Manning (2016). A thorough examination of the CNN/Daily
10773 Mail reading comprehension task. In *Proceedings of the Association for Computational
10774 Linguistics (ACL)*.
- 10775 Chen, D. and C. D. Manning (2014). A fast and accurate dependency parser using neural
10776 networks. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*,
10777 pp. 740–750.
- 10778 Chen, D. L. and R. J. Mooney (2008). Learning to sportscast: a test of grounded language
10779 acquisition. In *Proceedings of the International Conference on Machine Learning (ICML)*, pp.
10780 128–135.
- 10781 Chen, H., S. Branavan, R. Barzilay, and D. R. Karger (2009). Content modeling using latent
10782 permutations. *Journal of Artificial Intelligence Research* 36(1), 129–163.
- 10783 Chen, M., Z. Xu, K. Weinberger, and F. Sha (2012). Marginalized denoising autoencoders
10784 for domain adaptation. In *Proceedings of the International Conference on Machine Learning
10785 (ICML)*.

- 10786 Chen, M. X., O. Firat, A. Bapna, M. Johnson, W. Macherey, G. Foster, L. Jones, N. Parmar,
 10787 M. Schuster, Z. Chen, Y. Wu, and M. Hughes (2018). The best of both worlds: Combin-
 10788 ing recent advances in neural machine translation. In *Proceedings of the Association for*
 10789 *Computational Linguistics (ACL)*.
- 10790 Chen, S. F. and J. Goodman (1999). An empirical study of smoothing techniques for lan-
 10791 *guage modeling. Computer Speech & Language 13(4)*, 359–393.
- 10792 Chen, T. and C. Guestrin (2016). Xgboost: A scalable tree boosting system. In *Proceedings*
 10793 *of Knowledge Discovery and Data Mining (KDD)*, pp. 785–794.
- 10794 Chen, X., X. Qiu, C. Zhu, P. Liu, and X. Huang (2015). Long short-term memory neural
 10795 networks for chinese word segmentation. In *Proceedings of Empirical Methods for Natural*
 10796 *Language Processing (EMNLP)*, pp. 1197–1206.
- 10797 Chen, Y., S. Gilroy, A. Malletti, K. Knight, and J. May (2018). Recurrent neural networks
 10798 as weighted language recognizers. In *Proceedings of the North American Chapter of the*
 10799 *Association for Computational Linguistics (NAACL)*.
- 10800 Chen, Z. and H. Ji (2009). Graph-based event coreference resolution. In *Proceedings of*
 10801 *the 2009 Workshop on Graph-based Methods for Natural Language Processing*, pp. 54–57.
 10802 *Association for Computational Linguistics*.
- 10803 Cheng, X. and D. Roth (2013). Relational inference for wikification. In *Proceedings of*
 10804 *Empirical Methods for Natural Language Processing (EMNLP)*, pp. 1787–1796.
- 10805 Chiang, D. (2007). Hierarchical phrase-based translation. *Computational Linguistics 33(2)*,
 10806 201–228.
- 10807 Chiang, D., J. Graehl, K. Knight, A. Pauls, and S. Ravi (2010). Bayesian inference for
 10808 finite-state transducers. In *Proceedings of the North American Chapter of the Association for*
 10809 *Computational Linguistics (NAACL)*, pp. 447–455.
- 10810 Cho, K. (2015). Natural language understanding with distributed representation.
 10811 *CoRR abs/1511.07916*.
- 10812 Cho, K., B. Van Merriënboer, C. Gulcehre, D. Bahdanau, F. Bougares, H. Schwenk, and
 10813 Y. Bengio (2014). Learning phrase representations using rnn encoder-decoder for sta-
 10814 tistical machine translation. In *Proceedings of Empirical Methods for Natural Language*
 10815 *Processing (EMNLP)*.
- 10816 Chomsky, N. (1957). *Syntactic structures*. The Hague: Mouton & Co.
- 10817 Chomsky, N. (1982). *Some concepts and consequences of the theory of government and binding*,
 10818 Volume 6. MIT press.

- 10819 Choromanska, A., M. Henaff, M. Mathieu, G. B. Arous, and Y. LeCun (2015). The loss
10820 surfaces of multilayer networks. In *Proceedings of Artificial Intelligence and Statistics (AIS-*
10821 *TATS)*, pp. 192–204.
- 10822 Christensen, J., S. Soderland, O. Etzioni, et al. (2010). Semantic role labeling for open
10823 information extraction. In *Proceedings of the Workshop on Formalisms and Methodology for*
10824 *Learning by Reading*, pp. 52–60. Association for Computational Linguistics.
- 10825 Christodoulopoulos, C., S. Goldwater, and M. Steedman (2010). Two decades of unsuper-
10826 vised pos induction: How far have we come? In *Proceedings of Empirical Methods for*
10827 *Natural Language Processing (EMNLP)*, pp. 575–584.
- 10828 Chu, Y.-J. and T.-H. Liu (1965). On shortest arborescence of a directed graph. *Scientia*
10829 *Sinica* 14(10), 1396–1400.
- 10830 Chung, C. and J. W. Pennebaker (2007). The psychological functions of function words.
10831 In K. Fiedler (Ed.), *Social communication*, pp. 343–359. New York and Hove: Psychology
10832 Press.
- 10833 Church, K. (2011). A pendulum swung too far. *Linguistic Issues in Language Technology* 6(5),
10834 1–27.
- 10835 Church, K. W. (2000). Empirical estimates of adaptation: the chance of two Noriega-
10836 s is closer to $p/2$ than p^2 . In *Proceedings of the International Conference on Computational*
10837 *Linguistics (COLING)*, pp. 180–186.
- 10838 Church, K. W. and P. Hanks (1990). Word association norms, mutual information, and
10839 lexicography. *Computational linguistics* 16(1), 22–29.
- 10840 Ciaramita, M. and M. Johnson (2003). Supersense tagging of unknown nouns in wordnet.
10841 In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pp. 168–
10842 175.
- 10843 Clark, K. and C. D. Manning (2015). Entity-centric coreference resolution with model
10844 stacking. In *Proceedings of the Association for Computational Linguistics (ACL)*, pp. 1405–
10845 1415.
- 10846 Clark, K. and C. D. Manning (2016). Improving coreference resolution by learning entity-
10847 level distributed representations. In *Proceedings of the Association for Computational Lin-*
10848 *guistics (ACL)*.
- 10849 Clark, P. (2015). Elementary school science and math tests as a driver for ai: take the aristo
10850 challenge! In *Proceedings of the National Conference on Artificial Intelligence (AAAI)*, pp.
10851 4019–4021.

- 10852 Clarke, J., D. Goldwasser, M.-W. Chang, and D. Roth (2010). Driving semantic parsing
 10853 from the world’s response. In *Proceedings of the Conference on Natural Language Learning*
 10854 (*CoNLL*), pp. 18–27.
- 10855 Clarke, J. and M. Lapata (2008). Global inference for sentence compression: An integer
 10856 linear programming approach. *Journal of Artificial Intelligence Research* 31, 399–429.
- 10857 Cohen, J. (1960). A coefficient of agreement for nominal scales. *Educational and psychologi-*
 10858 *cal measurement* 20(1), 37–46.
- 10859 Cohen, S. (2016). *Bayesian analysis in natural language processing*. Synthesis Lectures on
 10860 Human Language Technologies. San Rafael, CA: Morgan & Claypool Publishers.
- 10861 Collier, N., C. Nobata, and J.-i. Tsujii (2000). Extracting the names of genes and gene
 10862 products with a hidden markov model. In *Proceedings of the International Conference on*
 10863 *Computational Linguistics (COLING)*, pp. 201–207.
- 10864 Collins, M. (1997). Three generative, lexicalised models for statistical parsing. In *Proceed-*
 10865 *ings of the Association for Computational Linguistics (ACL)*, pp. 16–23.
- 10866 Collins, M. (2002). Discriminative training methods for hidden markov models: theory
 10867 and experiments with perceptron algorithms. In *Proceedings of Empirical Methods for*
 10868 *Natural Language Processing (EMNLP)*, pp. 1–8.
- 10869 Collins, M. (2013). Notes on natural language processing. <http://www.cs.columbia.edu/~mcollins/notes-spring2013.html>.
- 10871 Collins, M. and T. Koo (2005). Discriminative reranking for natural language parsing.
 10872 *Computational Linguistics* 31(1), 25–70.
- 10873 Collins, M. and B. Roark (2004). Incremental parsing with the perceptron algorithm. In
 10874 *Proceedings of the 42nd Annual Meeting on Association for Computational Linguistics*, pp.
 10875 111. Association for Computational Linguistics.
- 10876 Collobert, R., K. Kavukcuoglu, and C. Farabet (2011). Torch7: A matlab-like environment
 10877 for machine learning. Technical Report EPFL-CONF-192376, EPFL.
- 10878 Collobert, R. and J. Weston (2008). A unified architecture for natural language process-
 10879 ing: Deep neural networks with multitask learning. In *Proceedings of the International*
 10880 *Conference on Machine Learning (ICML)*, pp. 160–167.
- 10881 Collobert, R., J. Weston, L. Bottou, M. Karlen, K. Kavukcuoglu, and P. Kuksa (2011). Nat-
 10882 ural language processing (almost) from scratch. *Journal of Machine Learning Research* 12,
 10883 2493–2537.

- 10884 Colton, S., J. Goodwin, and T. Veale (2012). Full-face poetry generation. In *Proceedings of
10885 the International Conference on Computational Creativity*, pp. 95–102.
- 10886 Conneau, A., D. Kiela, H. Schwenk, L. Barrault, and A. Bordes (2017). Supervised learning
10887 of universal sentence representations from natural language inference data. In *Proceed-
10888 ings of Empirical Methods for Natural Language Processing (EMNLP)*, pp. 681–691.
- 10889 Cormen, T. H., C. E. Leiserson, R. L. Rivest, and C. Stein (2009). *Introduction to algorithms*
10890 (third ed.). MIT press.
- 10891 Cotterell, R., H. Schütze, and J. Eisner (2016). Morphological smoothing and extrapolation
10892 of word embeddings. In *Proceedings of the Association for Computational Linguistics (ACL)*,
10893 pp. 1651–1660.
- 10894 Coviello, L., Y. Sohn, A. D. Kramer, C. Marlow, M. Franceschetti, N. A. Christakis, and
10895 J. H. Fowler (2014). Detecting emotional contagion in massive social networks. *PloS
one* 9(3), e90315.
- 10897 Covington, M. A. (2001). A fundamental algorithm for dependency parsing. In *Proceedings
10898 of the 39th annual ACM southeast conference*, pp. 95–102.
- 10899 Crammer, K., O. Dekel, J. Keshet, S. Shalev-Shwartz, and Y. Singer (2006, December).
10900 Online passive-aggressive algorithms. *The Journal of Machine Learning Research* 7, 551–
10901 585.
- 10902 Crammer, K. and Y. Singer (2001). Pranking with ranking. In *Neural Information Processing
10903 Systems (NIPS)*, pp. 641–647.
- 10904 Creutz, M. and K. Lagus (2007). Unsupervised models for morpheme segmentation and
10905 morphology learning. *ACM Transactions on Speech and Language Processing (TSLP)* 4(1),
10906 3.
- 10907 Cross, J. and L. Huang (2016). Span-based constituency parsing with a structure-label
10908 system and provably optimal dynamic oracles. In *Proceedings of Empirical Methods for
10909 Natural Language Processing (EMNLP)*, pp. 1–11.
- 10910 Cucerzan, S. (2007). Large-scale named entity disambiguation based on wikipedia data.
10911 In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*.
- 10912 Cui, H., R. Sun, K. Li, M.-Y. Kan, and T.-S. Chua (2005). Question answering passage
10913 retrieval using dependency relations. In *Proceedings of the 28th annual international ACM
10914 SIGIR conference on Research and development in information retrieval*, pp. 400–407. ACM.
- 10915 Cui, Y., Z. Chen, S. Wei, S. Wang, T. Liu, and G. Hu (2017). Attention-over-attention neural
10916 networks for reading comprehension. In *Proceedings of the Association for Computational
10917 Linguistics (ACL)*.

- 10918 Culotta, A. and J. Sorensen (2004). Dependency tree kernels for relation extraction. In
 10919 *Proceedings of the Association for Computational Linguistics (ACL)*.
- 10920 Culotta, A., M. Wick, and A. McCallum (2007). First-order probabilistic models for coref-
 10921 erence resolution. In *Proceedings of the North American Chapter of the Association for Com-*
 10922 *putational Linguistics (NAACL)*, pp. 81–88.
- 10923 Curry, H. B. and R. Feys (1958). *Combinatory Logic*, Volume I. Amsterdam: North Holland.
- 10924 Danescu-Niculescu-Mizil, C., M. Sudhof, D. Jurafsky, J. Leskovec, and C. Potts (2013). A
 10925 computational approach to politeness with application to social factors. In *Proceedings*
 10926 *of the Association for Computational Linguistics (ACL)*, pp. 250–259.
- 10927 Das, D., D. Chen, A. F. Martins, N. Schneider, and N. A. Smith (2014). Frame-semantic
 10928 parsing. *Computational Linguistics* 40(1), 9–56.
- 10929 Daumé III, H. (2007). Frustratingly easy domain adaptation. In *Proceedings of the Associa-*
 10930 *tion for Computational Linguistics (ACL)*.
- 10931 Daumé III, H., J. Langford, and D. Marcu (2009). Search-based structured prediction.
 10932 *Machine learning* 75(3), 297–325.
- 10933 Daumé III, H. and D. Marcu (2005). A large-scale exploration of effective global features
 10934 for a joint entity detection and tracking model. In *Proceedings of Empirical Methods for*
 10935 *Natural Language Processing (EMNLP)*, pp. 97–104.
- 10936 Dauphin, Y. N., R. Pascanu, C. Gulcehre, K. Cho, S. Ganguli, and Y. Bengio (2014). Identi-
 10937 fying and attacking the saddle point problem in high-dimensional non-convex opti-
 10938 mization. In *Neural Information Processing Systems (NIPS)*, pp. 2933–2941.
- 10939 Davidson, D. (1967). The logical form of action sentences. In N. Rescher (Ed.), *The Logic of*
 10940 *Decision and Action*. Pittsburgh: University of Pittsburgh Press.
- 10941 De Gispert, A. and J. B. Marino (2006). Catalan-english statistical machine translation
 10942 without parallel corpus: bridging through spanish. In *Proc. of 5th International Conference*
 10943 *on Language Resources and Evaluation (LREC)*, pp. 65–68. Citeseer.
- 10944 De Marneffe, M.-C. and C. D. Manning (2008). The stanford typed dependencies represen-
 10945 tation. In *Coling 2008: Proceedings of the workshop on Cross-Framework and Cross-Domain*
 10946 *Parser Evaluation*, pp. 1–8. Association for Computational Linguistics.
- 10947 Dean, J. and S. Ghemawat (2008). Mapreduce: simplified data processing on large clusters.
 10948 *Communications of the ACM* 51(1), 107–113.
- 10949 Deerwester, S. C., S. T. Dumais, T. K. Landauer, G. W. Furnas, and R. A. Harshman (1990).
 10950 Indexing by latent semantic analysis. *JASIS* 41(6), 391–407.

- 10951 Dehdari, J. (2014). *A Neurophysiologically-Inspired Statistical Language Model*. Ph. D. thesis,
10952 The Ohio State University.
- 10953 Deisenroth, M. P., A. A. Faisal, and C. S. Ong (2018). *Mathematics For Machine Learning*.
10954 Cambridge UP.
- 10955 Dempster, A. P., N. M. Laird, and D. B. Rubin (1977). Maximum likelihood from incom-
10956 plete data via the em algorithm. *Journal of the Royal Statistical Society. Series B (Method-
10957 ological)*, 1–38.
- 10958 Denis, P. and J. Baldridge (2007). A ranking approach to pronoun resolution. In *IJCAI*.
- 10959 Denis, P. and J. Baldridge (2008). Specialized models and ranking for coreference resolu-
10960 tion. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing*,
10961 EMNLP '08, Stroudsburg, PA, USA, pp. 660–669. Association for Computational Lin-
10962 guistics.
- 10963 Denis, P. and J. Baldridge (2009). Global joint models for coreference resolution and named
10964 entity classification. *Procesamiento del Lenguaje Natural* 42.
- 10965 Derrida, J. (1985). Des tours de babel. In J. Graham (Ed.), *Difference in translation*. Ithaca,
10966 NY: Cornell University Press.
- 10967 Dhingra, B., H. Liu, Z. Yang, W. W. Cohen, and R. Salakhutdinov (2017). Gated-attention
10968 readers for text comprehension. In *Proceedings of the Association for Computational Lin-
10969 guistics (ACL)*.
- 10970 Diaconis, P. and B. Skyrms (2017). *Ten Great Ideas About Chance*. Princeton University
10971 Press.
- 10972 Dietterich, T. G. (1998). Approximate statistical tests for comparing supervised classifica-
10973 tion learning algorithms. *Neural computation* 10(7), 1895–1923.
- 10974 Dietterich, T. G., R. H. Lathrop, and T. Lozano-Pérez (1997). Solving the multiple instance
10975 problem with axis-parallel rectangles. *Artificial intelligence* 89(1), 31–71.
- 10976 Dimitrova, L., N. Ide, V. Petkevic, T. Erjavec, H. J. Kaalep, and D. Tufis (1998). Multext-
10977 east: Parallel and comparable corpora and lexicons for six central and eastern european
10978 languages. In *Proceedings of the 17th international conference on Computational linguistics-
10979 Volume 1*, pp. 315–319. Association for Computational Linguistics.
- 10980 Doddington, G. R., A. Mitchell, M. A. Przybocki, L. A. Ramshaw, S. Strassel, and R. M.
10981 Weischedel (2004). The automatic content extraction (ace) program-tasks, data, and
10982 evaluation. In *Proceedings of the Language Resources and Evaluation Conference*, pp. 837–
10983 840.

- 10984 dos Santos, C., B. Xiang, and B. Zhou (2015). Classifying relations by ranking with con-
 10985 volutional neural networks. In *Proceedings of the Association for Computational Linguistics*
 10986 (ACL), pp. 626–634.
- 10987 Dowty, D. (1991). Thematic proto-roles and argument selection. *Language*, 547–619.
- 10988 Dredze, M., P. McNamee, D. Rao, A. Gerber, and T. Finin (2010). Entity disambiguation
 10989 for knowledge base population. In *Proceedings of the 23rd International Conference on*
 10990 *Computational Linguistics*, pp. 277–285. Association for Computational Linguistics.
- 10991 Dredze, M., M. J. Paul, S. Bergsma, and H. Tran (2013). Carmen: A Twitter geolocation
 10992 system with applications to public health. In *AAAI workshop on expanding the boundaries*
 10993 *of health informatics using AI (HIAI)*, pp. 20–24.
- 10994 Dreyfus, H. L. (1992). *What computers still can't do: a critique of artificial reason*. MIT press.
- 10995 Du, L., W. Buntine, and M. Johnson (2013). Topic segmentation with a structured topic
 10996 model. In *Proceedings of the North American Chapter of the Association for Computational*
 10997 *Linguistics* (NAACL), pp. 190–200.
- 10998 Duchi, J., E. Hazan, and Y. Singer (2011). Adaptive subgradient methods for online learn-
 10999 ing and stochastic optimization. *The Journal of Machine Learning Research* 12, 2121–2159.
- 11000 Dunietz, J., L. Levin, and J. Carbonell (2017). The because corpus 2.0: Annotating causality
 11001 and overlapping relations. In *Proceedings of the Linguistic Annotation Workshop*.
- 11002 Durrett, G., T. Berg-Kirkpatrick, and D. Klein (2016). Learning-based single-document
 11003 summarization with compression and anaphoricity constraints. In *Proceedings of the*
 11004 *Association for Computational Linguistics* (ACL), pp. 1998–2008.
- 11005 Durrett, G. and D. Klein (2013). Easy victories and uphill battles in coreference resolution.
 11006 In *Proceedings of the Conference on Empirical Methods in Natural Language Processing*.
- 11007 Durrett, G. and D. Klein (2015). Neural crf parsing. In *Proceedings of the Association for*
 11008 *Computational Linguistics* (ACL).
- 11009 Dyer, C., M. Ballesteros, W. Ling, A. Matthews, and N. A. Smith (2015). Transition-based
 11010 dependency parsing with stack long short-term memory. In *Proceedings of the Association*
 11011 *for Computational Linguistics* (ACL), pp. 334–343.
- 11012 Dyer, C., A. Kuncoro, M. Ballesteros, and N. A. Smith (2016). Recurrent neural network
 11013 grammars. In *Proceedings of the North American Chapter of the Association for Computational*
 11014 *Linguistics* (NAACL), pp. 199–209.
- 11015 Edmonds, J. (1967). Optimum branchings. *Journal of Research of the National Bureau of*
 11016 *Standards B* 71(4), 233–240.

- 11017 Efron, B. and R. J. Tibshirani (1993). An introduction to the bootstrap: Monographs on
11018 statistics and applied probability, vol. 57. *New York and London: Chapman and Hall/CRC*.
- 11019 Eisenstein, J. (2009). Hierarchical text segmentation from multi-scale lexical cohesion. In
11020 *Proceedings of the North American Chapter of the Association for Computational Linguistics*
11021 (NAACL).
- 11022 Eisenstein, J. and R. Barzilay (2008). Bayesian unsupervised topic segmentation. In *Pro-*
11023 *ceedings of Empirical Methods for Natural Language Processing (EMNLP)*.
- 11024 Eisner, J. (1997). State-of-the-art algorithms for minimum spanning trees: A tutorial dis-
11025 cussion.
- 11026 Eisner, J. (2000). Bilexical grammars and their cubic-time parsing algorithms. In *Advances*
11027 *in probabilistic and other parsing technologies*, pp. 29–61. Springer.
- 11028 Eisner, J. (2002). Parameter estimation for probabilistic finite-state transducers. In *Proced-*
11029 *ings of the Association for Computational Linguistics (ACL)*, pp. 1–8.
- 11030 Eisner, J. (2016). Inside-outside and forward-backward algorithms are just backprop. In
11031 *Proceedings of the Workshop on Structured Prediction for NLP*, pp. 1–17.
- 11032 Eisner, J. M. (1996). Three new probabilistic models for dependency parsing: An explo-
11033 ration. In *Proceedings of the International Conference on Computational Linguistics (COL-*
11034 *ING)*, pp. 340–345.
- 11035 Ekman, P. (1992). Are there basic emotions? *Psychological Review* 99(3), 550–553.
- 11036 Elman, J. L. (1990). Finding structure in time. *Cognitive science* 14(2), 179–211.
- 11037 Elman, J. L., E. A. Bates, M. H. Johnson, A. Karmiloff-Smith, D. Parisi, and K. Plunkett
11038 (1998). *Rethinking innateness: A connectionist perspective on development*, Volume 10. MIT
11039 press.
- 11040 Elsner, M. and E. Charniak (2010). Disentangling chat. *Computational Linguistics* 36(3),
11041 389–409.
- 11042 Esuli, A. and F. Sebastiani (2006). Sentiwordnet: A publicly available lexical resource for
11043 opinion mining. In *LREC*, Volume 6, pp. 417–422. Citeseer.
- 11044 Etzioni, O., A. Fader, J. Christensen, S. Soderland, and M. Mausam (2011). Open informa-
11045 tion extraction: The second generation. In *Proceedings of the International Joint Conference*
11046 *on Artificial Intelligence (IJCAI)*, pp. 3–10.

- 11047 Faruqui, M., J. Dodge, S. K. Jauhar, C. Dyer, E. Hovy, and N. A. Smith (2015). Retrofitting
 11048 word vectors to semantic lexicons. In *Proceedings of the North American Chapter of the*
 11049 *Association for Computational Linguistics (NAACL)*.
- 11050 Faruqui, M. and C. Dyer (2014). Improving vector space word representations using mul-
 11051 tilingual correlation. In *Proceedings of the European Chapter of the Association for Compu-*
 11052 *tational Linguistics (EACL)*, pp. 462–471.
- 11053 Faruqui, M., R. McDonald, and R. Soricut (2016). Morpho-syntactic lexicon generation
 11054 using graph-based semi-supervised learning. *Transactions of the Association for Compu-*
 11055 *tational Linguistics* 4, 1–16.
- 11056 Faruqui, M., Y. Tsvetkov, P. Rastogi, and C. Dyer (2016, August). Problems with evaluation
 11057 of word embeddings using word similarity tasks. In *Proceedings of the 1st Workshop on*
 11058 *Evaluating Vector-Space Representations for NLP*, Berlin, Germany, pp. 30–35. Association
 11059 for Computational Linguistics.
- 11060 Fellbaum, C. (2010). *WordNet*. Springer.
- 11061 Feng, V. W., Z. Lin, and G. Hirst (2014). The impact of deep hierarchical discourse struc-
 11062 tures in the evaluation of text coherence. In *Proceedings of the International Conference on*
 11063 *Computational Linguistics (COLING)*, pp. 940–949.
- 11064 Feng, X., L. Huang, D. Tang, H. Ji, B. Qin, and T. Liu (2016). A language-independent
 11065 neural network for event detection. In *Proceedings of the Association for Computational*
 11066 *Linguistics (ACL)*, pp. 66–71.
- 11067 Fernandes, E. R., C. N. dos Santos, and R. L. Milidiú (2014). Latent trees for coreference
 11068 resolution. *Computational Linguistics*.
- 11069 Ferrucci, D., E. Brown, J. Chu-Carroll, J. Fan, D. Gondek, A. A. Kalyanpur, A. Lally, J. W.
 11070 Murdock, E. Nyberg, J. Prager, et al. (2010). Building Watson: An overview of the
 11071 DeepQA project. *AI magazine* 31(3), 59–79.
- 11072 Ficler, J. and Y. Goldberg (2017, September). Controlling linguistic style aspects in neural
 11073 language generation. In *Proceedings of the Workshop on Stylistic Variation*, Copenhagen,
 11074 Denmark, pp. 94–104. Association for Computational Linguistics.
- 11075 Filippova, K. and M. Strube (2008). Sentence fusion via dependency graph compression.
 11076 In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pp. 177–
 11077 185.
- 11078 Fillmore, C. J. (1968). The case for case. In E. Bach and R. Harms (Eds.), *Universals in*
 11079 *linguistic theory*. Holt, Rinehart, and Winston.

- 11080 Fillmore, C. J. (1976). Frame semantics and the nature of language. *Annals of the New York
11081 Academy of Sciences* 280(1), 20–32.
- 11082 Fillmore, C. J. and C. Baker (2009). A frames approach to semantic analysis. In *The Oxford
11083 Handbook of Linguistic Analysis*. Oxford University Press.
- 11084 Finkel, J. R., T. Grenager, and C. Manning (2005). Incorporating non-local information
11085 into information extraction systems by gibbs sampling. In *Proceedings of the Association
11086 for Computational Linguistics (ACL)*, pp. 363–370.
- 11087 Finkel, J. R., T. Grenager, and C. D. Manning (2007). The infinite tree. In *Proceedings of the
11088 Association for Computational Linguistics (ACL)*, pp. 272–279.
- 11089 Finkel, J. R., A. Kleeman, and C. D. Manning (2008). Efficient, feature-based, conditional
11090 random field parsing. In *Proceedings of the Association for Computational Linguistics (ACL)*,
11091 pp. 959–967.
- 11092 Finkel, J. R. and C. Manning (2009). Hierarchical bayesian domain adaptation. In *Proceed-
11093 ings of the North American Chapter of the Association for Computational Linguistics (NAACL)*,
11094 pp. 602–610.
- 11095 Finkel, J. R. and C. D. Manning (2008). Enforcing transitivity in coreference resolution.
11096 In *Proceedings of the 46th Annual Meeting of the Association for Computational Linguistics
11097 on Human Language Technologies: Short Papers*, pp. 45–48. Association for Computational
11098 Linguistics.
- 11099 Finkelstein, L., E. Gabrilovich, Y. Matias, E. Rivlin, Z. Solan, G. Wolfman, and E. Ruppin
11100 (2002). Placing search in context: The concept revisited. *ACM Transactions on Information
11101 Systems* 20(1), 116–131.
- 11102 Firth, J. R. (1957). *Papers in Linguistics 1934-1951*. Oxford University Press.
- 11103 Flanigan, J., S. Thomson, J. Carbonell, C. Dyer, and N. A. Smith (2014). A discrimina-
11104 tive graph-based parser for the abstract meaning representation. In *Proceedings of the
11105 Association for Computational Linguistics (ACL)*, pp. 1426–1436.
- 11106 Foltz, P. W., W. Kintsch, and T. K. Landauer (1998). The measurement of textual coherence
11107 with latent semantic analysis. *Discourse processes* 25(2-3), 285–307.
- 11108 Fordyce, C. S. (2007). Overview of the iwslt 2007 evaluation campaign. In *International
11109 Workshop on Spoken Language Translation (IWSLT) 2007*.
- 11110 Fox, H. (2002). Phrasal cohesion and statistical machine translation. In *Proceedings of
11111 Empirical Methods for Natural Language Processing (EMNLP)*, pp. 304–3111.

- 11112 Francis, W. and H. Kucera (1982). *Frequency analysis of English usage*. Houghton Mifflin
11113 Company.
- 11114 Francis, W. N. (1964). A standard sample of present-day English for use with digital
11115 computers. Report to the U.S Office of Education on Cooperative Research Project No.
11116 E-007.
- 11117 Freund, Y. and R. E. Schapire (1999). Large margin classification using the perceptron
11118 algorithm. *Machine learning* 37(3), 277–296.
- 11119 Fromkin, V., R. Rodman, and N. Hyams (2013). *An introduction to language*. Cengage
11120 Learning.
- 11121 Fundel, K., R. Küffner, and R. Zimmer (2007). Relex – relation extraction using depen-
11122 dency parse trees. *Bioinformatics* 23(3), 365–371.
- 11123 Gabow, H. N., Z. Galil, T. Spencer, and R. E. Tarjan (1986). Efficient algorithms for finding
11124 minimum spanning trees in undirected and directed graphs. *Combinatorica* 6(2), 109–
11125 122.
- 11126 Gabrilovich, E. and S. Markovitch (2007). Computing semantic relatedness using
11127 wikipedia-based explicit semantic analysis. In *Proceedings of the International Joint Con-
11128 ference on Artificial Intelligence (IJCAI)*, Volume 7, pp. 1606–1611.
- 11129 Gage, P. (1994). A new algorithm for data compression. *The C Users Journal* 12(2), 23–38.
- 11130 Gale, W. A., K. W. Church, and D. Yarowsky (1992). One sense per discourse. In *Pro-
11131 ceedings of the workshop on Speech and Natural Language*, pp. 233–237. Association for
11132 Computational Linguistics.
- 11133 Galley, M., M. Hopkins, K. Knight, and D. Marcu (2004). What's in a translation rule? In
11134 *Proceedings of the North American Chapter of the Association for Computational Linguistics
(NAACL)*, pp. 273–280.
- 11136 Galley, M., K. R. McKeown, E. Fosler-Lussier, and H. Jing (2003). Discourse segmentation
11137 of multi-party conversation. In *Proceedings of the Association for Computational Linguistics
(ACL)*.
- 11139 Ganchev, K. and M. Dredze (2008). Small statistical models by random feature mixing. In
11140 *Proceedings of the ACL08 HLT Workshop on Mobile Language Processing*, pp. 19–20.
- 11141 Ganchev, K., J. Graça, J. Gillenwater, and B. Taskar (2010). Posterior regularization for
11142 structured latent variable models. *The Journal of Machine Learning Research* 11, 2001–
11143 2049.

- 11144 Ganin, Y., E. Ustinova, H. Ajakan, P. Germain, H. Larochelle, F. Laviolette, M. Marchand,
11145 and V. Lempitsky (2016). Domain-adversarial training of neural networks. *Journal of
11146 Machine Learning Research* 17(59), 1–35.
- 11147 Gao, J., G. Andrew, M. Johnson, and K. Toutanova (2007). A comparative study of param-
11148 eter estimation methods for statistical natural language processing. In *Proceedings of the
11149 Association for Computational Linguistics (ACL)*, pp. 824–831.
- 11150 Gatt, A. and E. Krahmer (2018). Survey of the state of the art in natural language genera-
11151 tion: Core tasks, applications and evaluation. *Journal of Artificial Intelligence Research* 61,
11152 65–170.
- 11153 Gatt, A. and E. Reiter (2009). Simplenlg: A realisation engine for practical applications.
11154 In *Proceedings of the 12th European Workshop on Natural Language Generation*, pp. 90–93.
11155 Association for Computational Linguistics.
- 11156 Ge, D., X. Jiang, and Y. Ye (2011). A note on the complexity of $l_1 p$ minimization. *Mathe-
11157 matical programming* 129(2), 285–299.
- 11158 Ge, N., J. Hale, and E. Charniak (1998). A statistical approach to anaphora resolution. In
11159 *Proceedings of the sixth workshop on very large corpora*, Volume 71, pp. 76.
- 11160 Ge, R., F. Huang, C. Jin, and Y. Yuan (2015). Escaping from saddle points — online stochas-
11161 tic gradient for tensor decomposition. In P. Grünwald, E. Hazan, and S. Kale (Eds.),
11162 *Proceedings of the Conference On Learning Theory (COLT)*.
- 11163 Ge, R. and R. J. Mooney (2005). A statistical semantic parser that integrates syntax and
11164 semantics. In *Proceedings of the Conference on Natural Language Learning (CoNLL)*, pp.
11165 9–16.
- 11166 Geach, P. T. (1962). *Reference and generality: An examination of some medieval and modern
11167 theories*. Cornell University Press.
- 11168 Gildea, D. and D. Jurafsky (2002). Automatic labeling of semantic roles. *Computational
11169 linguistics* 28(3), 245–288.
- 11170 Gimpel, K., N. Schneider, B. O’Connor, D. Das, D. Mills, J. Eisenstein, M. Heilman, D. Yo-
11171 gatama, J. Flanigan, and N. A. Smith (2011). Part-of-speech tagging for Twitter: an-
11172 notation, features, and experiments. In *Proceedings of the Association for Computational
11173 Linguistics (ACL)*, pp. 42–47.
- 11174 Glass, J., T. J. Hazen, S. Cyphers, I. Malioutov, D. Huynh, and R. Barzilay (2007). Recent
11175 progress in the mit spoken lecture processing project. In *Eighth Annual Conference of the
11176 International Speech Communication Association*.

- 11177 Glorot, X. and Y. Bengio (2010). Understanding the difficulty of training deep feedforward
 11178 neural networks. In *Proceedings of Artificial Intelligence and Statistics (AISTATS)*, pp. 249–
 11179 256.
- 11180 Glorot, X., A. Bordes, and Y. Bengio (2011). Deep sparse rectifier networks. In *Proceedings*
 11181 *of the 14th International Conference on Artificial Intelligence and Statistics. JMLR W&CP*
 11182 *Volume*, Volume 15, pp. 315–323.
- 11183 Godfrey, J. J., E. C. Holliman, and J. McDaniel (1992). Switchboard: Telephone speech
 11184 corpus for research and development. In *Proceedings of the International Conference on*
 11185 *Acoustics, Speech, and Signal Processing (ICASSP)*, pp. 517–520. IEEE.
- 11186 Goldberg, Y. (2017a, June). An adversarial review of “adversarial generation of
 11187 natural language”. [https://medium.com/@yoav.goldberg/an-adversarial-review-of-](https://medium.com/@yoav.goldberg/an-adversarial-review-of-adversarial-generation-of-natural-language-409ac3378bd7)
 11188 [adversarial-generation-of-natural-language-409ac3378bd7](#).
- 11189 Goldberg, Y. (2017b). *Neural Network Methods for Natural Language Processing*. Synthesis
 11190 Lectures on Human Language Technologies. Morgan & Claypool Publishers.
- 11191 Goldberg, Y. and M. Elhadad (2010). An efficient algorithm for easy-first non-directional
 11192 dependency parsing. In *Proceedings of the North American Chapter of the Association for*
 11193 *Computational Linguistics (NAACL)*, pp. 742–750.
- 11194 Goldberg, Y. and J. Nivre (2012). A dynamic oracle for arc-eager dependency parsing.
 11195 In *Proceedings of the International Conference on Computational Linguistics (COLING)*, pp.
 11196 959–976.
- 11197 Goldberg, Y., K. Zhao, and L. Huang (2013). Efficient implementation of beam-search
 11198 incremental parsers. In *ACL (2)*, pp. 628–633.
- 11199 Goldwater, S. and T. Griffiths (2007). A fully bayesian approach to unsupervised part-of-
 11200 speech tagging. In *Annual meeting-association for computational linguistics*, Volume 45.
- 11201 Gonçalo Oliveira, H. R., F. A. Cardoso, and F. C. Pereira (2007). Tra-la-lyrics: An approach
 11202 to generate text based on rhythm. In *Proceedings of the 4th. International Joint Workshop*
 11203 *on Computational Creativity*. A. Cardoso and G. Wiggins.
- 11204 Goodfellow, I., Y. Bengio, and A. Courville (2016). *Deep learning*. MIT Press.
- 11205 Goodman, J. T. (2001). A bit of progress in language modeling. *Computer Speech & Lan-*
 11206 *guage* 15(4), 403–434.
- 11207 Gouws, S., D. Metzler, C. Cai, and E. Hovy (2011). Contextual bearing on linguistic varia-
 11208 tion in social media. In *LASM*.

- 11209 Goyal, A., H. Daume III, and S. Venkatasubramanian (2009). Streaming for large scale
11210 nlp: Language modeling. In *Proceedings of the North American Chapter of the Association*
11211 for Computational Linguistics (NAACL), pp. 512–520.
- 11212 Graves, A. (2012). Sequence transduction with recurrent neural networks. In *Proceedings*
11213 of the International Conference on Machine Learning (ICML).
- 11214 Graves, A. and N. Jaitly (2014). Towards end-to-end speech recognition with recur-
11215 rent neural networks. In *Proceedings of the International Conference on Machine Learning*
11216 (ICML), pp. 1764–1772.
- 11217 Graves, A. and J. Schmidhuber (2005). Framewise phoneme classification with bidirec-
11218 tional lstm and other neural network architectures. *Neural Networks* 18(5), 602–610.
- 11219 Grice, H. P. (1975). Logic and conversation. In P. Cole and J. L. Morgan (Eds.), *Syntax and*
11220 *Semantics Volume 3: Speech Acts*, pp. 41–58. Academic Press.
- 11221 Grishman, R. (2012). Information extraction: Capabilities and challenges. Notes prepared
11222 for the 2012 International Winter School in Language and Speech Technologies, Rovira
11223 i Virgili University, Tarragona, Spain.
- 11224 Grishman, R. (2015). Information extraction. *IEEE Intelligent Systems* 30(5), 8–15.
- 11225 Grishman, R., C. Macleod, and J. Sterling (1992). Evaluating parsing strategies using
11226 standardized parse files. In *Proceedings of the third conference on Applied natural language*
11227 *processing*, pp. 156–161. Association for Computational Linguistics.
- 11228 Grishman, R. and B. Sundheim (1996). Message understanding conference-6: A brief his-
11229 tory. In *Proceedings of the International Conference on Computational Linguistics* (COLING),
11230 pp. 466–471.
- 11231 Groenendijk, J. and M. Stokhof (1991). Dynamic predicate logic. *Linguistics and philoso-*
11232 *phy* 14(1), 39–100.
- 11233 Grosz, B. J. (1979). Focusing and description in natural language dialogues. Technical
11234 report, SRI INTERNATIONAL MENLO PARK CA.
- 11235 Grosz, B. J., S. Weinstein, and A. K. Joshi (1995). Centering: A framework for modeling
11236 the local coherence of discourse. *Computational linguistics* 21(2), 203–225.
- 11237 Gu, J., Z. Lu, H. Li, and V. O. Li (2016). Incorporating copying mechanism in sequence-to-
11238 sequence learning. In *Proceedings of the Association for Computational Linguistics* (ACL),
11239 pp. 1631–1640.
- 11240 Gulcehre, C., S. Ahn, R. Nallapati, B. Zhou, and Y. Bengio (2016). Pointing the unknown
11241 words. In *Proceedings of the Association for Computational Linguistics* (ACL), pp. 140–149.

- 11242 Gutmann, M. U. and A. Hyvärinen (2012). Noise-contrastive estimation of unnormalized
11243 statistical models, with applications to natural image statistics. *The Journal of Machine
11244 Learning Research* 13(1), 307–361.
- 11245 Haghghi, A. and D. Klein (2007). Unsupervised coreference resolution in a nonparametric
11246 bayesian model. In *Proceedings of the Association for Computational Linguistics (ACL)*.
- 11247 Haghghi, A. and D. Klein (2009). Simple coreference resolution with rich syntactic and
11248 semantic features. In *Proceedings of Empirical Methods for Natural Language Processing
11249 (EMNLP)*, pp. 1152–1161.
- 11250 Haghghi, A. and D. Klein (2010). Coreference resolution in a modular, entity-centered
11251 model. In *Proceedings of the North American Chapter of the Association for Computational
11252 Linguistics (NAACL)*, pp. 385–393.
- 11253 Hajič, J. and B. Hladká (1998). Tagging inflective languages: Prediction of morphological
11254 categories for a rich, structured tagset. In *Proceedings of the Association for Computational
11255 Linguistics (ACL)*, pp. 483–490.
- 11256 Halliday, M. and R. Hasan (1976). *Cohesion in English*. London: Longman.
- 11257 Hammerton, J. (2003). Named entity recognition with long short-term memory. In *Pro-
11258 ceedings of the Conference on Natural Language Learning (CoNLL)*, pp. 172–175.
- 11259 Han, X. and L. Sun (2012). An entity-topic model for entity linking. In *Proceedings of
11260 Empirical Methods for Natural Language Processing (EMNLP)*, pp. 105–115.
- 11261 Han, X., L. Sun, and J. Zhao (2011). Collective entity linking in web text: a graph-based
11262 method. In *Proceedings of ACM SIGIR conference on Research and development in informa-
11263 tion retrieval*, pp. 765–774.
- 11264 Hannak, A., E. Anderson, L. F. Barrett, S. Lehmann, A. Mislove, and M. Riedewald (2012).
11265 Tweetin'in the rain: Exploring societal-scale effects of weather on mood. In *Proceedings
11266 of the International Conference on Web and Social Media (ICWSM)*.
- 11267 Hardmeier, C. (2012). Discourse in statistical machine translation. a survey and a case
11268 study. *Discours. Revue de linguistique, psycholinguistique et informatique. A journal of lin-
11269 guistics, psycholinguistics and computational linguistics* (11).
- 11270 Haspelmath, M. and A. Sims (2013). *Understanding morphology*. Routledge.
- 11271 Hastie, T., R. Tibshirani, and J. Friedman (2009). *The elements of statistical learning* (Second
11272 ed.). New York: Springer.

- 11273 Hatzivassiloglou, V. and K. R. McKeown (1997). Predicting the semantic orientation of
11274 adjectives. In *Proceedings of the Association for Computational Linguistics (ACL)*, pp. 174–
11275 181.
- 11276 Hayes, A. F. and K. Krippendorff (2007). Answering the call for a standard reliability
11277 measure for coding data. *Communication methods and measures* 1(1), 77–89.
- 11278 He, H., A. Balakrishnan, M. Eric, and P. Liang (2017). Learning symmetric collaborative
11279 dialogue agents with dynamic knowledge graph embeddings. In *Proceedings of the As-
11280 sociation for Computational Linguistics (ACL)*, pp. 1766–1776.
- 11281 He, K., X. Zhang, S. Ren, and J. Sun (2015). Delving deep into rectifiers: Surpassing
11282 human-level performance on imagenet classification. In *Proceedings of the International
11283 Conference on Computer Vision (ICCV)*, pp. 1026–1034.
- 11284 He, K., X. Zhang, S. Ren, and J. Sun (2016). Deep residual learning for image recognition.
11285 In *Proceedings of the International Conference on Computer Vision (ICCV)*, pp. 770–778.
- 11286 He, L., K. Lee, M. Lewis, and L. Zettlemoyer (2017). Deep semantic role labeling: What
11287 works and what's next. In *Proceedings of the Association for Computational Linguistics
11288 (ACL)*.
- 11289 He, Z., S. Liu, M. Li, M. Zhou, L. Zhang, and H. Wang (2013). Learning entity repre-
11290 sentation for entity disambiguation. In *Proceedings of the Association for Computational
11291 Linguistics (ACL)*, pp. 30–34.
- 11292 Hearst, M. A. (1992). Automatic acquisition of hyponyms from large text corpora. In
11293 *Proceedings of the International Conference on Computational Linguistics (COLING)*, pp. 539–
11294 545. Association for Computational Linguistics.
- 11295 Hearst, M. A. (1997). Texttiling: Segmenting text into multi-paragraph subtopic passages.
11296 *Computational linguistics* 23(1), 33–64.
- 11297 Heerschap, B., F. Goossen, A. Hogenboom, F. Frasincar, U. Kaymak, and F. de Jong (2011).
11298 Polarity analysis of texts using discourse structure. In *Proceedings of the 20th ACM inter-
11299 national conference on Information and knowledge management*, pp. 1061–1070. ACM.
- 11300 Henderson, J. (2004). Discriminative training of a neural network statistical parser. In
11301 *Proceedings of the Association for Computational Linguistics (ACL)*, pp. 95–102.
- 11302 Hendrickx, I., S. N. Kim, Z. Kozareva, P. Nakov, D. Ó Séaghdha, S. Padó, M. Pennacchiotti,
11303 L. Romano, and S. Szpakowicz (2009). SemEval-2010 task 8: Multi-way classification of
11304 semantic relations between pairs of nominals. In *Proceedings of the Workshop on Semantic
11305 Evaluations: Recent Achievements and Future Directions*, pp. 94–99. Association for Com-
11306 putational Linguistics.

- 11307 Hermann, K. M., T. Kočiský, E. Grefenstette, L. Espeholt, W. Kay, M. Suleyman, and
 11308 P. Blunsom (2015). Teaching machines to read and comprehend. In *Advances in Neu-*
 11309 *ral Information Processing Systems*, pp. 1693–1701.
- 11310 Hernault, H., H. Prendinger, D. A. duVerle, and M. Ishizuka (2010). HILDA: A discourse
 11311 parser using support vector machine classification. *Dialogue and Discourse* 1(3), 1–33.
- 11312 Hill, F., A. Bordes, S. Chopra, and J. Weston (2016). The goldilocks principle: Reading
 11313 children’s books with explicit memory representations. In *Proceedings of the International*
 11314 *Conference on Learning Representations (ICLR)*.
- 11315 Hill, F., K. Cho, and A. Korhonen (2016). Learning distributed representations of sentences
 11316 from unlabelled data. In *Proceedings of the North American Chapter of the Association for*
 11317 *Computational Linguistics (NAACL)*.
- 11318 Hindle, D. and M. Rooth (1993). Structural ambiguity and lexical relations. *Computational*
 11319 *linguistics* 19(1), 103–120.
- 11320 Hirao, T., Y. Yoshida, M. Nishino, N. Yasuda, and M. Nagata (2013). Single-document
 11321 summarization as a tree knapsack problem. In *Proceedings of Empirical Methods for Nat-*
 11322 *ural Language Processing (EMNLP)*, pp. 1515–1520.
- 11323 Hirschman, L. and R. Gaizauskas (2001). Natural language question answering: the view
 11324 from here. *natural language engineering* 7(4), 275–300.
- 11325 Hirschman, L., M. Light, E. Breck, and J. D. Burger (1999). Deep read: A reading compre-
 11326 hension system. In *Proceedings of the Association for Computational Linguistics (ACL)*, pp.
 11327 325–332.
- 11328 Hobbs, J. R. (1978). Resolving pronoun references. *Lingua* 44(4), 311–338.
- 11329 Hobbs, J. R., D. Appelt, J. Bear, D. Israel, M. Kameyama, M. Stickel, and M. Tyson (1997).
 11330 Fastus: A cascaded finite-state transducer for extracting information from natural-
 11331 language text. *Finite-state language processing*, 383–406.
- 11332 Hochreiter, S. and J. Schmidhuber (1997). Long short-term memory. *Neural computa-*
 11333 *tion* 9(8), 1735–1780.
- 11334 Hockenmaier, J. and M. Steedman (2007). Ccgbank: a corpus of ccg derivations and de-
 11335 pendency structures extracted from the penn treebank. *Computational Linguistics* 33(3),
 11336 355–396.
- 11337 Hoffart, J., M. A. Yosef, I. Bordino, H. Fürstenau, M. Pinkal, M. Spaniol, B. Taneva,
 11338 S. Thater, and G. Weikum (2011). Robust disambiguation of named entities in text. In
 11339 *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pp. 782–792.

- 11340 Hoffmann, R., C. Zhang, X. Ling, L. Zettlemoyer, and D. S. Weld (2011). Knowledge-based
11341 weak supervision for information extraction of overlapping relations. In *Proceedings of*
11342 *the Association for Computational Linguistics (ACL)*, pp. 541–550.
- 11343 Holmstrom, L. and P. Koistinen (1992). Using additive noise in back-propagation training.
11344 *IEEE Transactions on Neural Networks* 3(1), 24–38.
- 11345 Hovy, E. and J. Lavid (2010). Towards a ‘science’ of corpus annotation: a new method-
11346 ological challenge for corpus linguistics. *International journal of translation* 22(1), 13–36.
- 11347 Hsu, D., S. M. Kakade, and T. Zhang (2012). A spectral algorithm for learning hidden
11348 markov models. *Journal of Computer and System Sciences* 78(5), 1460–1480.
- 11349 Hu, M. and B. Liu (2004). Mining and summarizing customer reviews. In *Proceedings of*
11350 *Knowledge Discovery and Data Mining (KDD)*, pp. 168–177.
- 11351 Hu, Z., Z. Yang, X. Liang, R. Salakhutdinov, and E. P. Xing (2017). Toward controlled
11352 generation of text. In *International Conference on Machine Learning*, pp. 1587–1596.
- 11353 Huang, F. and A. Yates (2012). Biased representation learning for domain adaptation. In
11354 *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pp. 1313–1323.
- 11355 Huang, L. and D. Chiang (2007). Forest rescoring: Faster decoding with integrated lan-
11356 guage models. In *Proceedings of the Association for Computational Linguistics (ACL)*, pp.
11357 144–151.
- 11358 Huang, L., S. Fayong, and Y. Guo (2012). Structured perceptron with inexact search. In
11359 *Proceedings of the North American Chapter of the Association for Computational Linguistics*
11360 (*NAACL*), pp. 142–151.
- 11361 Huang, Y. (2015). *Pragmatics* (Second ed.). Oxford Textbooks in Linguistics. Oxford Uni-
11362 versity Press.
- 11363 Huang, Z., W. Xu, and K. Yu (2015). Bidirectional lstm-crf models for sequence tagging.
11364 *arXiv preprint arXiv:1508.01991*.
- 11365 Huffman, D. A. (1952). A method for the construction of minimum-redundancy codes.
11366 *Proceedings of the IRE* 40(9), 1098–1101.
- 11367 Humphreys, K., R. Gaizauskas, and S. Azzam (1997). Event coreference for information
11368 extraction. In *Proceedings of a Workshop on Operational Factors in Practical, Robust Anaphora*
11369 *Resolution for Unrestricted Texts*, pp. 75–81. Association for Computational Linguistics.
- 11370 Ide, N. and Y. Wilks (2006). Making sense about sense. In *Word sense disambiguation*, pp.
11371 47–73. Springer.

- 11372 Ioffe, S. and C. Szegedy (2015). Batch normalization: Accelerating deep network training by reducing internal covariate shift. In *Proceedings of the International Conference on Machine Learning (ICML)*, pp. 448–456.
- 11375 Isozaki, H., T. Hirao, K. Duh, K. Sudoh, and H. Tsukada (2010). Automatic evaluation of translation quality for distant language pairs. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pp. 944–952.
- 11378 Iyyer, M., V. Manjunatha, J. Boyd-Graber, and H. Daumé III (2015). Deep unordered composition rivals syntactic methods for text classification. In *Proceedings of the Association for Computational Linguistics (ACL)*, pp. 1681–1691.
- 11381 James, G., D. Witten, T. Hastie, and R. Tibshirani (2013). *An introduction to statistical learning*, Volume 112. Springer.
- 11383 Janin, A., D. Baron, J. Edwards, D. Ellis, D. Gelbart, N. Morgan, B. Peskin, T. Pfau, E. Shriberg, A. Stolcke, et al. (2003). The ICSI meeting corpus. In *Acoustics, Speech, and Signal Processing, 2003. Proceedings.(ICASSP'03). 2003 IEEE International Conference on*, Volume 1, pp. I–I. IEEE.
- 11387 Jean, S., K. Cho, R. Memisevic, and Y. Bengio (2015). On using very large target vocabulary for neural machine translation. In *Proceedings of the Association for Computational Linguistics (ACL)*, pp. 1–10.
- 11390 Jeong, M., C.-Y. Lin, and G. G. Lee (2009). Semi-supervised speech act recognition in emails and forums. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pp. 1250–1259.
- 11393 Ji, H. and R. Grishman (2011). Knowledge base population: Successful approaches and challenges. In *Proceedings of the Association for Computational Linguistics (ACL)*, pp. 1148–1158.
- 11396 Ji, Y., T. Cohn, L. Kong, C. Dyer, and J. Eisenstein (2015). Document context language models. In *International Conference on Learning Representations, Workshop Track*, Volume abs/1511.03962.
- 11399 Ji, Y. and J. Eisenstein (2014). Representation learning for text-level discourse parsing. In *Proceedings of the Association for Computational Linguistics (ACL)*.
- 11401 Ji, Y. and J. Eisenstein (2015, June). One vector is not enough: Entity-augmented distributional semantics for discourse relations. *Transactions of the Association for Computational Linguistics (TACL)*.

- 11404 Ji, Y., G. Haffari, and J. Eisenstein (2016). A latent variable recurrent neural network for
11405 discourse relation language models. In *Proceedings of the North American Chapter of the*
11406 *Association for Computational Linguistics (NAACL)*.
- 11407 Ji, Y. and N. A. Smith (2017). Neural discourse structure for text categorization. In *Pro-
11408 ceedings of the Association for Computational Linguistics (ACL)*, pp. 996–1005.
- 11409 Ji, Y., C. Tan, S. Martschat, Y. Choi, and N. A. Smith (2017). Dynamic entity representations
11410 in neural language models. In *Proceedings of Empirical Methods for Natural Language
11411 Processing (EMNLP)*, pp. 1831–1840.
- 11412 Jiang, L., M. Yu, M. Zhou, X. Liu, and T. Zhao (2011). Target-dependent twitter sentiment
11413 classification. In *Proceedings of the Association for Computational Linguistics (ACL)*, pp.
11414 151–160.
- 11415 Jing, H. (2000). Sentence reduction for automatic text summarization. In *Proceedings of
the sixth conference on Applied natural language processing*, pp. 310–315. Association for
11416 Computational Linguistics.
- 11417 Joachims, T. (2002). Optimizing search engines using clickthrough data. In *Proceedings of
Knowledge Discovery and Data Mining (KDD)*, pp. 133–142.
- 11418 Jockers, M. L. (2015). Szuzhet? <http://bla.bla.com>.
- 11421 Johnson, A. E., T. J. Pollard, L. Shen, H. L. Li-wei, M. Feng, M. Ghassemi, B. Moody,
11422 P. Szolovits, L. A. Celi, and R. G. Mark (2016). Mimic-iii, a freely accessible critical care
11423 database. *Scientific data* 3, 160035.
- 11424 Johnson, M. (1998). Pcfg models of linguistic tree representations. *Computational Linguis-
11425 tics* 24(4), 613–632.
- 11426 Johnson, R. and T. Zhang (2017). Deep pyramid convolutional neural networks for text
11427 categorization. In *Proceedings of the Association for Computational Linguistics (ACL)*, pp.
11428 562–570.
- 11429 Joshi, A. K. (1985). How much context-sensitivity is required to provide reasonable struc-
11430 tural descriptions? – tree adjoining grammars. In *Natural Language Processing – Theoret-
11431 ical, Computational and Psychological Perspective*. New York, NY: Cambridge University
11432 Press.
- 11433 Joshi, A. K. and Y. Schabes (1997). Tree-adjoining grammars. In *Handbook of formal lan-
11434 guages*, pp. 69–123. Springer.
- 11435 Joshi, A. K., K. V. Shanker, and D. Weir (1991). The convergence of mildly context-sensitive
11436 grammar formalisms. In *Foundational Issues in Natural Language Processing*. Cambridge
11437 MA: MIT Press.

- 11438 Jozefowicz, R., O. Vinyals, M. Schuster, N. Shazeer, and Y. Wu (2016). Exploring the limits
11439 of language modeling. *arXiv preprint arXiv:1602.02410*.
- 11440 Jozefowicz, R., W. Zaremba, and I. Sutskever (2015). An empirical exploration of recurrent
11441 network architectures. In *Proceedings of the International Conference on Machine Learning*
11442 (ICML), pp. 2342–2350.
- 11443 Jurafsky, D. (1996). A probabilistic model of lexical and syntactic access and disambiguation.
11444 *Cognitive Science* 20(2), 137–194.
- 11445 Jurafsky, D. and J. H. Martin (2009). *Speech and Language Processing* (Second ed.). Prentice
11446 Hall.
- 11447 Jurafsky, D. and J. H. Martin (2018). *Speech and Language Processing* (Third ed.). Prentice
11448 Hall.
- 11449 Kadlec, R., M. Schmid, O. Bajgar, and J. Kleindienst (2016). Text understanding with
11450 the attention sum reader network. In *Proceedings of the Association for Computational*
11451 *Linguistics (ACL)*, pp. 908–918.
- 11452 Kalchbrenner, N. and P. Blunsom (2013, August). Recurrent convolutional neural net-
11453 works for discourse compositionality. In *Proceedings of the Workshop on Continuous Vec-*
11454 *tor Space Models and their Compositionality*, Sofia, Bulgaria, pp. 119–126. Association for
11455 Computational Linguistics.
- 11456 Kalchbrenner, N., E. Grefenstette, and P. Blunsom (2014). A convolutional neural network
11457 for modelling sentences. In *Proceedings of the Association for Computational Linguistics*
11458 (*ACL*), pp. 655–665.
- 11459 Karlsson, F. (2007). Constraints on multiple center-embedding of clauses. *Journal of Lin-*
11460 *guistics* 43(02), 365–392.
- 11461 Kate, R. J., Y. W. Wong, and R. J. Mooney (2005). Learning to transform natural to formal
11462 languages. In *Proceedings of the National Conference on Artificial Intelligence (AAAI)*.
- 11463 Kehler, A. (2007). Rethinking the SMASH approach to pronoun interpretation. In *Interdis-*
11464 *ciplinary perspectives on reference processing*, New Directions in Cognitive Science Series,
11465 pp. 95–122. Oxford University Press.
- 11466 Kibble, R. and R. Power (2004). Optimizing referential coherence in text generation. *Com-*
11467 *putational Linguistics* 30(4), 401–416.
- 11468 Kilgarriff, A. (1997). I don't believe in word senses. *Computers and the Humanities* 31(2),
11469 91–113.

- 11470 Kilgarriff, A. and G. Grefenstette (2003). Introduction to the special issue on the web as
11471 corpus. *Computational linguistics* 29(3), 333–347.
- 11472 Kim, M.-J. (2002). Does korean have adjectives? *MIT Working Papers in Linguistics* 43,
11473 71–89.
- 11474 Kim, S.-M. and E. Hovy (2006, July). Extracting opinions, opinion holders, and topics
11475 expressed in online news media text. In *Proceedings of the Workshop on Sentiment and*
11476 *Subjectivity in Text*, Sydney, Australia, pp. 1–8. Association for Computational Linguis-
11477 tics.
- 11478 Kim, Y. (2014). Convolutional neural networks for sentence classification. In *Proceedings*
11479 *of Empirical Methods for Natural Language Processing (EMNLP)*, pp. 1746–1751.
- 11480 Kim, Y., C. Denton, L. Hoang, and A. M. Rush (2017). Structured attention networks. In
11481 *Proceedings of the International Conference on Learning Representations (ICLR)*.
- 11482 Kim, Y., Y. Jernite, D. Sontag, and A. M. Rush (2016). Character-aware neural language
11483 models. In *Proceedings of the National Conference on Artificial Intelligence (AAAI)*.
- 11484 Kingma, D. and J. Ba (2014). Adam: A method for stochastic optimization. *arXiv preprint*
11485 *arXiv:1412.6980*.
- 11486 Kiperwasser, E. and Y. Goldberg (2016). Simple and accurate dependency parsing using
11487 bidirectional lstm feature representations. *Transactions of the Association for Compu-
11488 tational Linguistics* 4, 313–327.
- 11489 Kipper-Schuler, K. (2005). *VerbNet: A broad-coverage, comprehensive verb lexicon*. Ph. D.
11490 thesis, Computer and Information Science, University of Pennsylvania.
- 11491 Kiros, R., R. Salakhutdinov, and R. Zemel (2014). Multimodal neural language models. In
11492 *Proceedings of the International Conference on Machine Learning (ICML)*, pp. 595–603.
- 11493 Kiros, R., Y. Zhu, R. Salakhutdinov, R. S. Zemel, A. Torralba, R. Urtasun, and S. Fidler
11494 (2015). Skip-thought vectors. In *Neural Information Processing Systems (NIPS)*.
- 11495 Klein, D. and C. D. Manning (2003). Accurate unlexicalized parsing. In *Proceedings of the*
11496 *Association for Computational Linguistics (ACL)*, pp. 423–430.
- 11497 Klein, D. and C. D. Manning (2004). Corpus-based induction of syntactic structure: Mod-
11498 els of dependency and constituency. In *Proceedings of the Association for Computational*
11499 *Linguistics (ACL)*.
- 11500 Klein, G., Y. Kim, Y. Deng, J. Senellart, and A. M. Rush (2017). Opennmt: Open-source
11501 toolkit for neural machine translation. *arXiv preprint arXiv:1701.02810*.

- 11502 Klementiev, A., I. Titov, and B. Bhattachari (2012). Inducing crosslingual distributed representations of words. In *Proceedings of the International Conference on Computational Linguistics (COLING)*, pp. 1459–1474.
- 11505 Klenner, M. (2007). Enforcing consistency on coreference sets. In *Recent Advances in Natural Language Processing (RANLP)*, pp. 323–328.
- 11507 Knight, K. (1999). Decoding complexity in word-replacement translation models. *Computational Linguistics* 25(4), 607–615.
- 11509 Knight, K. and J. Graehl (1998). Machine transliteration. *Computational Linguistics* 24(4), 599–612.
- 11511 Knight, K. and D. Marcu (2000). Statistics-based summarization-step one: Sentence compression. In *Proceedings of the National Conference on Artificial Intelligence (AAAI)*, pp. 703–710.
- 11514 Knight, K. and J. May (2009). Applications of weighted automata in natural language processing. In *Handbook of Weighted Automata*, pp. 571–596. Springer.
- 11516 Knott, A. (1996). *A data-driven methodology for motivating a set of coherence relations*. Ph. D. thesis, The University of Edinburgh.
- 11518 Koehn, P. (2005). Europarl: A parallel corpus for statistical machine translation. In *MT summit*, Volume 5, pp. 79–86.
- 11520 Koehn, P. (2009). *Statistical machine translation*. Cambridge University Press.
- 11521 Koehn, P. (2017). Neural machine translation. *arXiv preprint arXiv:1709.07809*.
- 11522 Konstas, I. and M. Lapata (2013). A global model for concept-to-text generation. *Journal of Artificial Intelligence Research* 48, 305–346.
- 11524 Koo, T., X. Carreras, and M. Collins (2008, jun). Simple semi-supervised dependency parsing. In *Proceedings of ACL-08: HLT*, Columbus, Ohio, pp. 595–603. Association for Computational Linguistics.
- 11527 Koo, T. and M. Collins (2005). Hidden-variable models for discriminative reranking. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pp. 507–514.
- 11529 Koo, T. and M. Collins (2010). Efficient third-order dependency parsers. In *Proceedings of the Association for Computational Linguistics (ACL)*.
- 11531 Koo, T., A. Globerson, X. Carreras, and M. Collins (2007). Structured prediction models via the matrix-tree theorem. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pp. 141–150.

- 11534 Kovach, B. and T. Rosenstiel (2014). *The elements of journalism: What newspeople should know
11535 and the public should expect*. Three Rivers Press.
- 11536 Krishnamurthy, J. (2016). Probabilistic models for learning a semantic parser lexicon. In
11537 *Proceedings of the North American Chapter of the Association for Computational Linguistics
11538 (NAACL)*, pp. 606–616.
- 11539 Krishnamurthy, J. and T. M. Mitchell (2012). Weakly supervised training of semantic
11540 parsers. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*,
11541 pp. 754–765.
- 11542 Krizhevsky, A., I. Sutskever, and G. E. Hinton (2012). Imagenet classification with deep
11543 convolutional neural networks. In *Neural Information Processing Systems (NIPS)*, pp.
11544 1097–1105.
- 11545 Kübler, S., R. McDonald, and J. Nivre (2009). Dependency parsing. *Synthesis Lectures on
11546 Human Language Technologies* 1(1), 1–127.
- 11547 Kuhlmann, M. and J. Nivre (2010). Transition-based techniques for non-projective depen-
11548 dency parsing. *Northern European Journal of Language Technology (NEJLT)* 2(1), 1–19.
- 11549 Kummerfeld, J. K., T. Berg-Kirkpatrick, and D. Klein (2015). An empirical analysis of op-
11550 timization for max-margin NLP. In *Proceedings of Empirical Methods for Natural Language
11551 Processing (EMNLP)*.
- 11552 Kwiatkowski, T., S. Goldwater, L. Zettlemoyer, and M. Steedman (2012). A probabilistic
11553 model of syntactic and semantic acquisition from child-directed utterances and their
11554 meanings. In *Proceedings of the European Chapter of the Association for Computational Lin-
11555 guistics (EACL)*, pp. 234–244.
- 11556 Lafferty, J., A. McCallum, and F. Pereira (2001). Conditional random fields: Probabilistic
11557 models for segmenting and labeling sequence data. In *icml*.
- 11558 Lakoff, G. (1973). Hedges: A study in meaning criteria and the logic of fuzzy concepts.
11559 *Journal of philosophical logic* 2(4), 458–508.
- 11560 Lample, G., M. Ballesteros, S. Subramanian, K. Kawakami, and C. Dyer (2016). Neural
11561 architectures for named entity recognition. In *Proceedings of the North American Chapter
11562 of the Association for Computational Linguistics (NAACL)*, pp. 260–270.
- 11563 Langkilde, I. and K. Knight (1998). Generation that exploits corpus-based statistical
11564 knowledge. In *Proceedings of the Association for Computational Linguistics (ACL)*, pp. 704–
11565 710.

- 11566 Lapata, M. (2003). Probabilistic text structuring: Experiments with sentence ordering. In
 11567 *Proceedings of the Association for Computational Linguistics (ACL)*, pp. 545–552.
- 11568 Lappin, S. and H. J. Leass (1994). An algorithm for pronominal anaphora resolution.
 11569 *Computational linguistics* 20(4), 535–561.
- 11570 Lari, K. and S. J. Young (1990). The estimation of stochastic context-free grammars using
 11571 the inside-outside algorithm. *Computer speech & language* 4(1), 35–56.
- 11572 Lascarides, A. and N. Asher (2007). Segmented discourse representation theory: Dynamic
 11573 semantics with discourse structure. In *Computing meaning*, pp. 87–124. Springer.
- 11574 Law, E. and L. v. Ahn (2011). Human computation. *Synthesis Lectures on Artificial Intelli-*
 11575 *gence and Machine Learning* 5(3), 1–121.
- 11576 Lebret, R., D. Grangier, and M. Auli (2016). Neural text generation from structured data
 11577 with application to the biography domain. In *Proceedings of Empirical Methods for Natural*
 11578 *Language Processing (EMNLP)*, pp. 1203–1213.
- 11579 LeCun, Y. and Y. Bengio (1995). Convolutional networks for images, speech, and time
 11580 series. *The handbook of brain theory and neural networks* 3361.
- 11581 LeCun, Y., L. Bottou, G. B. Orr, and K.-R. Müller (1998). Efficient backprop. In *Neural*
 11582 *networks: Tricks of the trade*, pp. 9–50. Springer.
- 11583 Lee, C. M. and S. S. Narayanan (2005). Toward detecting emotions in spoken dialogs.
 11584 *IEEE transactions on speech and audio processing* 13(2), 293–303.
- 11585 Lee, H., A. Chang, Y. Peirsman, N. Chambers, M. Surdeanu, and D. Jurafsky (2013). De-
 11586 terministic coreference resolution based on entity-centric, precision-ranked rules. *Com-*
 11587 *putational Linguistics* 39(4), 885–916.
- 11588 Lee, H., Y. Peirsman, A. Chang, N. Chambers, M. Surdeanu, and D. Jurafsky (2011). Stan-
 11589 ford’s multi-pass sieve coreference resolution system at the conll-2011 shared task. In
 11590 *Proceedings of the Conference on Natural Language Learning (CoNLL)*, pp. 28–34. Associa-
 11591 tion for Computational Linguistics.
- 11592 Lee, K., L. He, M. Lewis, and L. Zettlemoyer (2017). End-to-end neural coreference reso-
 11593 lution. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*.
- 11594 Lenat, D. B., R. V. Guha, K. Pittman, D. Pratt, and M. Shepherd (1990). Cyc: toward
 11595 programs with common sense. *Communications of the ACM* 33(8), 30–49.
- 11596 Lesk, M. (1986). Automatic sense disambiguation using machine readable dictionaries:
 11597 how to tell a pine cone from an ice cream cone. In *Proceedings of the 5th annual interna-*
 11598 *tional conference on Systems documentation*, pp. 24–26. ACM.

- 11599 Levesque, H. J., E. Davis, and L. Morgenstern (2011). The winograd schema challenge.
11600 In *Aaaai spring symposium: Logical formalizations of commonsense reasoning*, Volume 46, pp.
11601 47.
- 11602 Levin, E., R. Pieraccini, and W. Eckert (1998). Using markov decision process for learning
11603 dialogue strategies. In *Acoustics, Speech and Signal Processing, 1998. Proceedings of the*
11604 *1998 IEEE International Conference on*, Volume 1, pp. 201–204. IEEE.
- 11605 Levy, O. and Y. Goldberg (2014). Dependency-based word embeddings. In *Proceedings of*
11606 *the Association for Computational Linguistics (ACL)*, pp. 302–308.
- 11607 Levy, O., Y. Goldberg, and I. Dagan (2015). Improving distributional similarity with
11608 lessons learned from word embeddings. *Transactions of the Association for Computational*
11609 *Linguistics* 3, 211–225.
- 11610 Levy, R. and C. Manning (2009). An informal introduction to computational semantics.
- 11611 Lewis, M. and M. Steedman (2013). Combined distributional and logical semantics. *Trans-*
11612 *actions of the Association for Computational Linguistics* 1, 179–192.
- 11613 Lewis II, P. M. and R. E. Stearns (1968). Syntax-directed transduction. *Journal of the ACM*
11614 (*JACM*) 15(3), 465–488.
- 11615 Li, J. and D. Jurafsky (2015). Do multi-sense embeddings improve natural language
11616 understanding? In *Proceedings of Empirical Methods for Natural Language Processing*
11617 (*EMNLP*), pp. 1722–1732.
- 11618 Li, J. and D. Jurafsky (2017). Neural net models of open-domain discourse coherence. In
11619 *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pp. 198–209.
- 11620 Li, J., R. Li, and E. Hovy (2014). Recursive deep models for discourse parsing. In *Proceed-*
11621 *ings of Empirical Methods for Natural Language Processing (EMNLP)*.
- 11622 Li, J., M.-T. Luong, and D. Jurafsky (2015). A hierarchical neural autoencoder for para-
11623 graphs and documents. In *Proceedings of Empirical Methods for Natural Language Process-*
11624 *ing (EMNLP)*.
- 11625 Li, J., T. Luong, D. Jurafsky, and E. Hovy (2015). When are tree structures necessary
11626 for deep learning of representations? In *Proceedings of Empirical Methods for Natural*
11627 *Language Processing (EMNLP)*, pp. 2304–2314.
- 11628 Li, J., W. Monroe, A. Ritter, D. Jurafsky, M. Galley, and J. Gao (2016, November). Deep
11629 reinforcement learning for dialogue generation. In *Proceedings of the 2016 Conference on*
11630 *Empirical Methods in Natural Language Processing*, Austin, Texas, pp. 1192–1202. Associa-
11631 *tion for Computational Linguistics*.

- 11632 Li, Q., S. Anzaroot, W.-P. Lin, X. Li, and H. Ji (2011). Joint inference for cross-document
 11633 information extraction. In *Proceedings of the International Conference on Information and*
 11634 *Knowledge Management (CIKM)*, pp. 2225–2228.
- 11635 Li, Q., H. Ji, and L. Huang (2013). Joint event extraction via structured prediction with
 11636 global features. In *Proceedings of the Association for Computational Linguistics (ACL)*, pp.
 11637 73–82.
- 11638 Liang, P., A. Bouchard-Côté, D. Klein, and B. Taskar (2006). An end-to-end discriminative
 11639 approach to machine translation. In *Proceedings of the Association for Computational*
 11640 *Linguistics (ACL)*, pp. 761–768.
- 11641 Liang, P., M. Jordan, and D. Klein (2009). Learning semantic correspondences with less
 11642 supervision. In *Proceedings of the Association for Computational Linguistics (ACL)*, pp. 91–
 11643 99.
- 11644 Liang, P., M. I. Jordan, and D. Klein (2013). Learning dependency-based compositional
 11645 semantics. *Computational Linguistics* 39(2), 389–446.
- 11646 Liang, P. and D. Klein (2009). Online em for unsupervised models. In *Proceedings of the*
 11647 *North American Chapter of the Association for Computational Linguistics (NAACL)*, pp. 611–
 11648 619.
- 11649 Liang, P., S. Petrov, M. I. Jordan, and D. Klein (2007). The infinite pcfg using hierarchical
 11650 dirichlet processes. In *Proceedings of Empirical Methods for Natural Language Processing*
 11651 (*EMNLP*), pp. 688–697.
- 11652 Liang, P. and C. Potts (2015). Bringing machine learning and compositional semantics
 11653 together. *Annual Review of Linguistics* 1(1), 355–376.
- 11654 Lieber, R. (2015). *Introducing morphology*. Cambridge University Press.
- 11655 Lin, D. (1998). Automatic retrieval and clustering of similar words. In *Proceedings of the*
 11656 *17th international conference on Computational linguistics-Volume 2*, pp. 768–774. Associa-
 11657 tion for Computational Linguistics.
- 11658 Lin, J. and C. Dyer (2010). Data-intensive text processing with mapreduce. *Synthesis*
 11659 *Lectures on Human Language Technologies* 3(1), 1–177.
- 11660 Lin, Z., M. Feng, C. N. d. Santos, M. Yu, B. Xiang, B. Zhou, and Y. Bengio (2017). A
 11661 structured self-attentive sentence embedding. *arXiv preprint arXiv:1703.03130*.
- 11662 Lin, Z., M.-Y. Kan, and H. T. Ng (2009). Recognizing implicit discourse relations in the
 11663 penn discourse treebank. In *Proceedings of Empirical Methods for Natural Language Pro-*
 11664 *cessing (EMNLP)*, pp. 343–351.

- 11665 Lin, Z., H. T. Ng, and M.-Y. Kan (2011). Automatically evaluating text coherence using
11666 discourse relations. In *Proceedings of the Association for Computational Linguistics (ACL)*,
11667 pp. 997–1006.
- 11668 Lin, Z., H. T. Ng, and M. Y. Kan (2014, nov). A PDTB-styled end-to-end discourse parser.
11669 *Natural Language Engineering FirstView*, 1–34.
- 11670 Ling, W., C. Dyer, A. Black, and I. Trancoso (2015). Two/too simple adaptations of
11671 word2vec for syntax problems. In *Proceedings of the North American Chapter of the As-*
11672 *sociation for Computational Linguistics (NAACL)*.
- 11673 Ling, W., T. Luís, L. Marujo, R. F. Astudillo, S. Amir, C. Dyer, A. W. Black, and I. Trancoso
11674 (2015). Finding function in form: Compositional character models for open vocabulary
11675 word representation. In *Proceedings of Empirical Methods for Natural Language Processing*
11676 (*EMNLP*).
- 11677 Ling, W., G. Xiang, C. Dyer, A. Black, and I. Trancoso (2013). Microblogs as parallel cor-
11678 pora. In *Proceedings of the Association for Computational Linguistics (ACL)*.
- 11679 Ling, X., S. Singh, and D. S. Weld (2015). Design challenges for entity linking. *Transactions*
11680 *of the Association for Computational Linguistics* 3, 315–328.
- 11681 Linguistic Data Consortium (2005, July). ACE (automatic content extraction) English an-
11682 notation guidelines for relations. Technical Report Version 5.8.3, Linguistic Data Con-
11683 sortium.
- 11684 Liu, B. (2015). *Sentiment Analysis: Mining Opinions, Sentiments, and Emotions*. Cambridge
11685 University Press.
- 11686 Liu, D. C. and J. Nocedal (1989). On the limited memory BFGS method for large scale
11687 optimization. *Mathematical programming* 45(1-3), 503–528.
- 11688 Liu, Y., Q. Liu, and S. Lin (2006). Tree-to-string alignment template for statistical machine
11689 translation. In *Proceedings of the Association for Computational Linguistics (ACL)*, pp. 609–
11690 616.
- 11691 Loper, E. and S. Bird (2002). Nltk: The natural language toolkit. In *Proceedings of the ACL-*
11692 *02 Workshop on Effective tools and methodologies for teaching natural language processing and*
11693 *computational linguistics-Volume 1*, pp. 63–70. Association for Computational Linguistics.
- 11694 Louis, A., A. Joshi, and A. Nenkova (2010). Discourse indicators for content selection in
11695 summarization. In *Proceedings of the 11th Annual Meeting of the Special Interest Group on*
11696 *Discourse and Dialogue*, pp. 147–156. Association for Computational Linguistics.

- 11697 Louis, A. and A. Nenkova (2013). What makes writing great? first experiments on article
 11698 quality prediction in the science journalism domain. *Transactions of the Association for
 11699 Computational Linguistics* 1, 341–352.
- 11700 Loveland, D. W. (2016). *Automated Theorem Proving: a logical basis*. Elsevier.
- 11701 Lowe, R., N. Pow, I. V. Serban, and J. Pineau (2015). The ubuntu dialogue corpus: A large
 11702 dataset for research in unstructured multi-turn dialogue systems. In *Proceedings of the
 11703 Special Interest Group on Discourse and Dialogue (SIGDIAL)*.
- 11704 Luo, X. (2005). On coreference resolution performance metrics. In *Proceedings of Empirical
 11705 Methods for Natural Language Processing (EMNLP)*, pp. 25–32.
- 11706 Luo, X., A. Ittycheriah, H. Jing, N. Kambhatla, and S. Roukos (2004). A mention-
 11707 synchronous coreference resolution algorithm based on the bell tree. In *Proceedings
 11708 of the Association for Computational Linguistics (ACL)*.
- 11709 Luong, M.-T., R. Socher, and C. D. Manning (2013). Better word representations with
 11710 recursive neural networks for morphology. *CoNLL-2013* 104.
- 11711 Luong, T., H. Pham, and C. D. Manning (2015). Effective approaches to attention-based
 11712 neural machine translation. In *Proceedings of Empirical Methods for Natural Language
 11713 Processing (EMNLP)*, pp. 1412–1421.
- 11714 Luong, T., I. Sutskever, Q. Le, O. Vinyals, and W. Zaremba (2015). Addressing the rare
 11715 word problem in neural machine translation. In *Proceedings of the Association for Compu-
 11716 tational Linguistics (ACL)*, pp. 11–19.
- 11717 Maas, A. L., A. Y. Hannun, and A. Y. Ng (2013). Rectifier nonlinearities improve neu-
 11718 ral network acoustic models. In *Proceedings of the International Conference on Machine
 11719 Learning (ICML)*.
- 11720 Mairesse, F. and M. A. Walker (2011). Controlling user perceptions of linguistic style:
 11721 Trainable generation of personality traits. *Computational Linguistics* 37(3), 455–488.
- 11722 Mani, I., M. Verhagen, B. Wellner, C. M. Lee, and J. Pustejovsky (2006). Machine learning
 11723 of temporal relations. In *Proceedings of the Association for Computational Linguistics (ACL)*,
 11724 pp. 753–760.
- 11725 Mann, W. C. and S. A. Thompson (1988). Rhetorical structure theory: Toward a functional
 11726 theory of text organization. *Text* 8(3), 243–281.
- 11727 Manning, C. D. (2015). Computational linguistics and deep learning. *Computational Lin-
 11728 guistics* 41(4), 701–707.

- 11729 Manning, C. D. (2016). Computational linguistics and deep learning. *Computational Lin-*
11730 *guistics* 41(4).
- 11731 Manning, C. D., P. Raghavan, H. Schütze, et al. (2008). *Introduction to information retrieval*,
11732 Volume 1. Cambridge university press.
- 11733 Manning, C. D. and H. Schütze (1999). *Foundations of Statistical Natural Language Process-*
11734 *ing*. Cambridge, Massachusetts: MIT press.
- 11735 Marcu, D. (1996). Building up rhetorical structure trees. In *Proceedings of the National*
11736 *Conference on Artificial Intelligence*, pp. 1069–1074.
- 11737 Marcu, D. (1997a). From discourse structures to text summaries. In *Proceedings of the*
11738 *workshop on Intelligent Scalable Text Summarization*.
- 11739 Marcu, D. (1997b). From local to global coherence: A bottom-up approach to text plan-
11740 ning. In *Proceedings of the National Conference on Artificial Intelligence (AAAI)*, pp. 629–635.
- 11741 Marcus, M. P., M. A. Marcinkiewicz, and B. Santorini (1993). Building a large annotated
11742 corpus of English: The Penn Treebank. *Computational Linguistics* 19(2), 313–330.
- 11743 Maron, O. and T. Lozano-Pérez (1998). A framework for multiple-instance learning. In
11744 *Neural Information Processing Systems (NIPS)*, pp. 570–576.
- 11745 Márquez, G. G. (1970). *One Hundred Years of Solitude*. Harper & Row. English translation
11746 by Gregory Rabassa.
- 11747 Martins, A. F. T., N. A. Smith, and E. P. Xing (2009). Concise integer linear programming
11748 formulations for dependency parsing. In *Proceedings of the Association for Computational*
11749 *Linguistics (ACL)*, pp. 342–350.
- 11750 Martins, A. F. T., N. A. Smith, E. P. Xing, P. M. Q. Aguiar, and M. A. T. Figueiredo (2010).
11751 Turbo parsers: Dependency parsing by approximate variational inference. In *Proceed-
11752 ings of Empirical Methods for Natural Language Processing (EMNLP)*, pp. 34–44.
- 11753 Matsuzaki, T., Y. Miyao, and J. Tsujii (2005). Probabilistic cfg with latent annotations. In
11754 *Proceedings of the Association for Computational Linguistics (ACL)*, pp. 75–82.
- 11755 Matthiessen, C. and J. A. Bateman (1991). *Text generation and systemic-functional linguistics:
11756 experiences from English and Japanese*. Pinter Publishers.
- 11757 McCallum, A. and W. Li (2003). Early results for named entity recognition with condi-
11758 tional random fields, feature induction and web-enhanced lexicons. In *Proceedings of
11759 the North American Chapter of the Association for Computational Linguistics (NAACL)*, pp.
11760 188–191.

- 11761 McCallum, A. and B. Wellner (2004). Conditional models of identity uncertainty with
 11762 application to noun coreference. In *NIPS*, pp. 905–912.
- 11763 McDonald, R., K. Crammer, and F. Pereira (2005). Online large-margin training of depen-
 11764 dency parsers. In *Proceedings of the Association for Computational Linguistics (ACL)*, pp.
 11765 91–98.
- 11766 McDonald, R., K. Hannan, T. Neylon, M. Wells, and J. Reynar (2007). Structured models
 11767 for fine-to-coarse sentiment analysis. In *Proceedings of ACL*.
- 11768 McDonald, R. and F. Pereira (2006). Online learning of approximate dependency parsing
 11769 algorithms. In *Proceedings of the European Chapter of the Association for Computational
 11770 Linguistics (EACL)*.
- 11771 McKeown, K. (1992). *Text generation*. Cambridge University Press.
- 11772 McKeown, K., S. Rosenthal, K. Thadani, and C. Moore (2010). Time-efficient creation of
 11773 an accurate sentence fusion corpus. In *Proceedings of the North American Chapter of the
 11774 Association for Computational Linguistics (NAACL)*, pp. 317–320.
- 11775 McKeown, K. R., R. Barzilay, D. Evans, V. Hatzivassiloglou, J. L. Klavans, A. Nenkova,
 11776 C. Sable, B. Schiffman, and S. Sigelman (2002). Tracking and summarizing news on a
 11777 daily basis with columbia’s newsblaster. In *Proceedings of the second international confer-
 11778 ence on Human Language Technology Research*, pp. 280–285.
- 11779 McNamee, P. and H. T. Dang (2009). Overview of the tac 2009 knowledge base population
 11780 track. In *Text Analysis Conference (TAC)*, Volume 17, pp. 111–113.
- 11781 Medlock, B. and T. Briscoe (2007). Weakly supervised learning for hedge classification in
 11782 scientific literature. In *Proceedings of the Association for Computational Linguistics (ACL)*,
 11783 pp. 992–999.
- 11784 Mei, H., M. Bansal, and M. R. Walter (2016). What to talk about and how? selective gen-
 11785 eration using lstms with coarse-to-fine alignment. In *Proceedings of the North American
 11786 Chapter of the Association for Computational Linguistics (NAACL)*, pp. 720–730.
- 11787 Merity, S., N. S. Keskar, and R. Socher (2018). Regularizing and optimizing lstm language
 11788 models. In *Proceedings of the International Conference on Learning Representations (ICLR)*.
- 11789 Merity, S., C. Xiong, J. Bradbury, and R. Socher (2017). Pointer sentinel mixture models.
 11790 In *Proceedings of the International Conference on Learning Representations (ICLR)*.
- 11791 Messud, C. (2014, June). A new ‘l’étranger’. *New York Review of Books*.

- 11792 Miao, Y. and P. Blunsom (2016). Language as a latent variable: Discrete generative mod-
11793 els for sentence compression. In *Proceedings of Empirical Methods for Natural Language*
11794 *Processing (EMNLP)*, pp. 319–328.
- 11795 Miao, Y., L. Yu, and P. Blunsom (2016). Neural variational inference for text processing. In
11796 *Proceedings of the International Conference on Machine Learning (ICML)*.
- 11797 Mihalcea, R., T. A. Chklovski, and A. Kilgarriff (2004, July). The senseval-3 english lexical
11798 sample task. In *Proceedings of SENSEVAL-3*, Barcelona, Spain, pp. 25–28. Association for
11799 Computational Linguistics.
- 11800 Mihalcea, R. and D. Radev (2011). *Graph-based natural language processing and information*
11801 *retrieval*. Cambridge University Press.
- 11802 Mikolov, T., K. Chen, G. Corrado, and J. Dean (2013). Efficient estimation of word repre-
11803 sentations in vector space. In *Proceedings of International Conference on Learning Represen-*
11804 *tations*.
- 11805 Mikolov, T., A. Deoras, D. Povey, L. Burget, and J. Cernocky (2011). Strategies for train-
11806 ing large scale neural network language models. In *Automatic Speech Recognition and*
11807 *Understanding (ASRU), 2011 IEEE Workshop on*, pp. 196–201. IEEE.
- 11808 Mikolov, T., M. Karafiat, L. Burget, J. Cernocky, and S. Khudanpur (2010). Recurrent
11809 neural network based language model. In *INTERSPEECH*, pp. 1045–1048.
- 11810 Mikolov, T., I. Sutskever, K. Chen, G. S. Corrado, and J. Dean (2013). Distributed rep-
11811 resentations of words and phrases and their compositionality. In *Advances in Neural*
11812 *Information Processing Systems*, pp. 3111–3119.
- 11813 Mikolov, T., W.-t. Yih, and G. Zweig (2013). Linguistic regularities in continuous space
11814 word representations. In *Proceedings of the North American Chapter of the Association for*
11815 *Computational Linguistics (NAACL)*, pp. 746–751.
- 11816 Mikolov, T. and G. Zweig. Context dependent recurrent neural network language model.
11817 In *Proceedings of Spoken Language Technology (SLT)*, pp. 234–239.
- 11818 Miller, G. A., G. A. Heise, and W. Lichten (1951). The intelligibility of speech as a function
11819 of the context of the test materials. *Journal of experimental psychology* 41(5), 329.
- 11820 Miller, M., C. Sathi, D. Wiesenthal, J. Leskovec, and C. Potts (2011). Sentiment flow
11821 through hyperlink networks. In *Proceedings of the International Conference on Web and*
11822 *Social Media (ICWSM)*.
- 11823 Miller, S., J. Guinness, and A. Zamanian (2004). Name tagging with word clusters and
11824 discriminative training. In *Proceedings of the North American Chapter of the Association for*
11825 *Computational Linguistics (NAACL)*, pp. 337–342.

- 11826 Milne, D. and I. H. Witten (2008). Learning to link with wikipedia. In *Proceedings of the
11827 International Conference on Information and Knowledge Management (CIKM)*, pp. 509–518.
11828 ACM.
- 11829 Miltzakaki, E. and K. Kukich (2004). Evaluation of text coherence for electronic essay
11830 scoring systems. *Natural Language Engineering* 10(1), 25–55.
- 11831 Minka, T. P. (1999). From hidden markov models to linear dynamical systems. Tech. Rep.
11832 531, Vision and Modeling Group of Media Lab, MIT.
- 11833 Minsky, M. (1974). A framework for representing knowledge. Technical Report 306, MIT
11834 AI Laboratory.
- 11835 Minsky, M. and S. Papert (1969). *Perceptrons*. MIT press.
- 11836 Mintz, M., S. Bills, R. Snow, and D. Jurafsky (2009). Distant supervision for relation extrac-
11837 tion without labeled data. In *Proceedings of the Association for Computational Linguistics
11838 (ACL)*, pp. 1003–1011.
- 11839 Mirza, P., R. Sprugnoli, S. Tonelli, and M. Speranza (2014). Annotating causality in the
11840 tempeval-3 corpus. In *Proceedings of the EACL 2014 Workshop on Computational Ap-
11841 proaches to Causality in Language (CAtoCL)*, pp. 10–19.
- 11842 Misra, D. K. and Y. Artzi (2016). Neural shift-reduce ccg semantic parsing. In *Proceedings
11843 of Empirical Methods for Natural Language Processing (EMNLP)*.
- 11844 Mitchell, J. and M. Lapata (2010). Composition in distributional models of semantics.
11845 *Cognitive Science* 34(8), 1388–1429.
- 11846 Miwa, M. and M. Bansal (2016). End-to-end relation extraction using lstms on sequences
11847 and tree structures. In *Proceedings of the Association for Computational Linguistics (ACL)*,
11848 pp. 1105–1116.
- 11849 Mnih, A. and G. Hinton (2007). Three new graphical models for statistical language mod-
11850 elling. In *Proceedings of the 24th international conference on Machine learning*, ICML '07,
11851 New York, NY, USA, pp. 641–648. ACM.
- 11852 Mnih, A. and G. E. Hinton (2008). A scalable hierarchical distributed language model. In
11853 *Neural Information Processing Systems (NIPS)*, pp. 1081–1088.
- 11854 Mnih, A. and Y. W. Teh (2012). A fast and simple algorithm for training neural probabilis-
11855 tic language models. In *Proceedings of the International Conference on Machine Learning
11856 (ICML)*.
- 11857 Mohammad, S. M. and P. D. Turney (2013). Crowdsourcing a word–emotion association
11858 lexicon. *Computational Intelligence* 29(3), 436–465.

- 11859 Mohri, M., F. Pereira, and M. Riley (2002). Weighted finite-state transducers in speech
11860 recognition. *Computer Speech & Language* 16(1), 69–88.
- 11861 Mohri, M., A. Rostamizadeh, and A. Talwalkar (2012). *Foundations of machine learning*.
11862 MIT press.
- 11863 Montague, R. (1973). The proper treatment of quantification in ordinary english. In *Ap-*
11864 *proaches to natural language*, pp. 221–242. Springer.
- 11865 Moore, J. D. and C. L. Paris (1993, dec). Planning text for advisory dialogues: Capturing
11866 intentional and rhetorical information. *Comput. Linguist.* 19(4), 651–694.
- 11867 Morante, R. and E. Blanco (2012). *sem 2012 shared task: Resolving the scope and fo-
11868 cus of negation. In *Proceedings of the First Joint Conference on Lexical and Computational*
11869 *Semantics-Volume 1: Proceedings of the main conference and the shared task, and Volume 2:*
11870 *Proceedings of the Sixth International Workshop on Semantic Evaluation*, pp. 265–274. Asso-
11871 ciation for Computational Linguistics.
- 11872 Morante, R. and W. Daelemans (2009). Learning the scope of hedge cues in biomedical
11873 texts. In *Proceedings of the Workshop on Current Trends in Biomedical Natural Language*
11874 *Processing*, pp. 28–36. Association for Computational Linguistics.
- 11875 Morante, R. and C. Sporleder (2012). Modality and negation: An introduction to the
11876 special issue. *Computational linguistics* 38(2), 223–260.
- 11877 Mostafazadeh, N., A. Grelish, N. Chambers, J. Allen, and L. Vanderwende (2016, June).
11878 Caters: Causal and temporal relation scheme for semantic annotation of event struc-
11879 tures. In *Proceedings of the Fourth Workshop on Events*, San Diego, California, pp. 51–61.
11880 Association for Computational Linguistics.
- 11881 Mueller, T., H. Schmid, and H. Schütze (2013). Efficient higher-order CRFs for morpholog-
11882 ical tagging. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*,
11883 pp. 322–332.
- 11884 Müller, C. and M. Strube (2006). Multi-level annotation of linguistic data with mmax2.
11885 *Corpus technology and language pedagogy: New resources, new tools, new methods* 3, 197–
11886 214.
- 11887 Muralidharan, A. and M. A. Hearst (2013). Supporting exploratory text analysis in litera-
11888 ture study. *Literary and linguistic computing* 28(2), 283–295.
- 11889 Murphy, K. P. (2012). *Machine Learning: A Probabilistic Perspective*. The MIT Press.
- 11890 Nakagawa, T., K. Inui, and S. Kurohashi (2010). Dependency tree-based sentiment classi-
11891 fication using crfs with hidden variables. In *Proceedings of the North American Chapter of*
11892 *the Association for Computational Linguistics (NAACL)*, pp. 786–794.

- 11893 Nakazawa, T., M. Yaguchi, K. Uchimoto, M. Utiyama, E. Sumita, S. Kurohashi, and H. Isahara (2016). ASPEC: Asian scientific paper excerpt corpus. In *Proceedings of the Language Resources and Evaluation Conference*, pp. 2204–2208.
- 11896 Navigli, R. (2009). Word sense disambiguation: A survey. *ACM Computing Surveys (CSUR)* 41(2), 10.
- 11898 Neal, R. M. and G. E. Hinton (1998). A view of the em algorithm that justifies incremental, sparse, and other variants. In *Learning in graphical models*, pp. 355–368. Springer.
- 11900 Nenkova, A. and K. McKeown (2012). A survey of text summarization techniques. In *Mining text data*, pp. 43–76. Springer.
- 11902 Neubig, G. (2017). Neural machine translation and sequence-to-sequence models: A tutorial. *arXiv preprint arXiv:1703.01619*.
- 11904 Neubig, G., C. Dyer, Y. Goldberg, A. Matthews, W. Ammar, A. Anastasopoulos, M. Balles-teros, D. Chiang, D. Clothiaux, T. Cohn, K. Duh, M. Faruqui, C. Gan, D. Garrette, Y. Ji, L. Kong, A. Kuncoro, G. Kumar, C. Malaviya, P. Michel, Y. Oda, M. Richardson, N. Saphra, S. Swayamdipta, and P. Yin (2017). Dynet: The dynamic neural network toolkit.
- 11909 Neubig, G., Y. Goldberg, and C. Dyer (2017). On-the-fly operation batching in dynamic computation graphs. In *Neural Information Processing Systems (NIPS)*.
- 11911 Neuhaus, P. and N. Bröker (1997). The complexity of recognition of linguistically adequate dependency grammars. In *eacl*, pp. 337–343.
- 11913 Newman, D., C. Chemudugunta, and P. Smyth (2006). Statistical entity-topic models. In *Proceedings of Knowledge Discovery and Data Mining (KDD)*, pp. 680–686.
- 11915 Ng, V. (2010). Supervised noun phrase coreference research: The first fifteen years. In *Proceedings of the 48th annual meeting of the association for computational linguistics*, pp. 1396–1411. Association for Computational Linguistics.
- 11918 Nguyen, D. and A. S. Dogruöz (2013). Word level language identification in online multi-lingual communication. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*.
- 11921 Nguyen, D. T. and S. Joty (2017). A neural local coherence model. In *Proceedings of the Association for Computational Linguistics (ACL)*, pp. 1320–1330.
- 11923 Nigam, K., A. K. McCallum, S. Thrun, and T. Mitchell (2000). Text classification from labeled and unlabeled documents using em. *Machine learning* 39(2-3), 103–134.

- 11925 Nirenburg, S. and Y. Wilks (2001). What's in a symbol: ontology, representation and lan-
11926 guage. *Journal of Experimental & Theoretical Artificial Intelligence* 13(1), 9–23.
- 11927 Nivre, J. (2008). Algorithms for deterministic incremental dependency parsing. *Computa-*
11928 *tional Linguistics* 34(4), 513–553.
- 11929 Nivre, J., M.-C. de Marneffe, F. Ginter, Y. Goldberg, J. Hajič, C. D. Manning, R. McDonald,
11930 S. Petrov, S. Pyysalo, N. Silveira, R. Tsarfaty, and D. Zeman (2016, may). Universal de-
11931 pendencies v1: A multilingual treebank collection. In N. C. C. Chair), K. Choukri, T. De-
11932 clerck, S. Goggi, M. Grobelnik, B. Maegaard, J. Mariani, H. Mazo, A. Moreno, J. Odijk,
11933 and S. Piperidis (Eds.), *Proceedings of the Tenth International Conference on Language Re-*
11934 *sources and Evaluation (LREC 2016)*, Paris, France. European Language Resources Asso-
11935 ciation (ELRA).
- 11936 Nivre, J. and J. Nilsson (2005). Pseudo-projective dependency parsing. In *Proceedings of the*
11937 *43rd Annual Meeting on Association for Computational Linguistics*, pp. 99–106. Association
11938 for Computational Linguistics.
- 11939 Novikoff, A. B. (1962). On convergence proofs on perceptrons. In *Proceedings of the Sym-*
11940 *posium on the Mathematical Theory of Automata*, Volume 12, pp. 615—622.
- 11941 Och, F. J. and H. Ney (2003). A systematic comparison of various statistical alignment
11942 models. *Computational linguistics* 29(1), 19–51.
- 11943 O'Connor, B., M. Krieger, and D. Ahn (2010). Tweetmotif: Exploratory search and topic
11944 summarization for twitter. In *Proceedings of the International Conference on Web and Social*
11945 *Media (ICWSM)*, pp. 384–385.
- 11946 Oflazer, K. and İ. Kuruöz (1994). Tagging and morphological disambiguation of turkish
11947 text. In *Proceedings of the fourth conference on Applied natural language processing*, pp. 144–
11948 149. Association for Computational Linguistics.
- 11949 Ohta, T., Y. Tateisi, and J.-D. Kim (2002). The genia corpus: An annotated research abstract
11950 corpus in molecular biology domain. In *Proceedings of the second international conference*
11951 *on Human Language Technology Research*, pp. 82–86. Morgan Kaufmann Publishers Inc.
- 11952 Onishi, T., H. Wang, M. Bansal, K. Gimpel, and D. McAllester (2016). Who did what: A
11953 large-scale person-centered cloze dataset. In *Proceedings of Empirical Methods for Natural*
11954 *Language Processing (EMNLP)*, pp. 2230–2235.
- 11955 Owoputi, O., B. O'Connor, C. Dyer, K. Gimpel, N. Schneider, and N. A. Smith (2013).
11956 Improved part-of-speech tagging for online conversational text with word clusters. In
11957 *Proceedings of the North American Chapter of the Association for Computational Linguistics*
11958 (*NAACL*), pp. 380–390.

- 11959 Packard, W., E. M. Bender, J. Read, S. Oepen, and R. Dridan (2014). Simple negation
 11960 scope resolution through deep parsing: A semantic solution to a semantic problem. In
 11961 *Proceedings of the Association for Computational Linguistics (ACL)*, pp. 69–78.
- 11962 Paice, C. D. (1990). Another stemmer. In *ACM SIGIR Forum*, Volume 24, pp. 56–61.
- 11963 Pak, A. and P. Paroubek (2010). Twitter as a corpus for sentiment analysis and opinion
 11964 mining. In *LREC*, Volume 10, pp. 1320–1326.
- 11965 Palmer, F. R. (2001). *Mood and modality*. Cambridge University Press.
- 11966 Palmer, M., D. Gildea, and P. Kingsbury (2005). The proposition bank: An annotated
 11967 corpus of semantic roles. *Computational linguistics* 31(1), 71–106.
- 11968 Pan, S. J. and Q. Yang (2010). A survey on transfer learning. *IEEE Transactions on knowledge
 11969 and data engineering* 22(10), 1345–1359.
- 11970 Pan, X., T. Cassidy, U. Hermjakob, H. Ji, and K. Knight (2015). Unsupervised entity linking
 11971 with abstract meaning representation. In *Proceedings of the North American Chapter of the
 11972 Association for Computational Linguistics (NAACL)*, pp. 1130–1139.
- 11973 Pang, B. and L. Lee (2004). A sentimental education: Sentiment analysis using subjectivity
 11974 summarization based on minimum cuts. In *Proceedings of the Association for Compu-
 11975 tational Linguistics (ACL)*, pp. 271–278.
- 11976 Pang, B. and L. Lee (2005). Seeing stars: Exploiting class relationships for sentiment cate-
 11977 gorization with respect to rating scales. In *Proceedings of the Association for Computational
 11978 Linguistics (ACL)*, pp. 115–124.
- 11979 Pang, B. and L. Lee (2008). Opinion mining and sentiment analysis. *Foundations and trends
 11980 in information retrieval* 2(1–2), 1–135.
- 11981 Pang, B., L. Lee, and S. Vaithyanathan (2002). Thumbs up?: sentiment classification using
 11982 machine learning techniques. In *Proceedings of Empirical Methods for Natural Language
 11983 Processing (EMNLP)*, pp. 79–86.
- 11984 Papineni, K., S. Roukos, T. Ward, and W.-J. Zhu (2002). Bleu: a method for automatic
 11985 evaluation of machine translation. In *Proceedings of the Association for Computational
 11986 Linguistics (ACL)*, pp. 311–318.
- 11987 Park, J. and C. Cardie (2012). Improving implicit discourse relation recognition through
 11988 feature set optimization. In *Proceedings of the Special Interest Group on Discourse and Dia-
 11989 logue (SIGDIAL)*, pp. 108–112.
- 11990 Parsons, T. (1990). *Events in the Semantics of English*, Volume 5. MIT Press.

- 11991 Pascanu, R., T. Mikolov, and Y. Bengio (2013). On the difficulty of training recurrent neural
11992 networks. In *Proceedings of the 30th International Conference on Machine Learning (ICML-*
11993 *13)*, pp. 1310–1318.
- 11994 Paul, M., M. Federico, and S. Stüker (2010). Overview of the iwslt 2010 evaluation cam-
11995 paign. In *International Workshop on Spoken Language Translation (IWSLT) 2010*.
- 11996 Pedersen, T., S. Patwardhan, and J. Michelizzi (2004). Wordnet::similarity - measuring the
11997 relatedness of concepts. In *Proceedings of the North American Chapter of the Association for*
11998 *Computational Linguistics (NAACL)*, pp. 38–41.
- 11999 Pedregosa, F., G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blon-
12000 del, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau,
12001 M. Brucher, M. Perrot, and E. Duchesnay (2011). Scikit-learn: Machine learning in
12002 Python. *Journal of Machine Learning Research* 12, 2825–2830.
- 12003 Pei, W., T. Ge, and B. Chang (2015). An effective neural network model for graph-based
12004 dependency parsing. In *Proceedings of the Association for Computational Linguistics (ACL)*,
12005 pp. 313–322.
- 12006 Peldszus, A. and M. Stede (2013). From argument diagrams to argumentation mining
12007 in texts: A survey. *International Journal of Cognitive Informatics and Natural Intelligence*
12008 (*IJCINI*) 7(1), 1–31.
- 12009 Peldszus, A. and M. Stede (2015). An annotated corpus of argumentative microtexts. In
12010 *Proceedings of the First Conference on Argumentation*.
- 12011 Peng, F., F. Feng, and A. McCallum (2004). Chinese segmentation and new word detec-
12012 tion using conditional random fields. In *Proceedings of the International Conference on*
12013 *Computational Linguistics (COLING)*, pp. 562.
- 12014 Pennington, J., R. Socher, and C. Manning (2014). Glove: Global vectors for word repre-
12015 sentation. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*,
12016 pp. 1532–1543.
- 12017 Pereira, F. and Y. Schabes (1992). Inside-outside reestimation from partially bracketed
12018 corpora. In *Proceedings of the Association for Computational Linguistics (ACL)*, pp. 128–
12019 135.
- 12020 Pereira, F. C. N. and S. M. Shieber (2002). *Prolog and natural-language analysis*. Microtome
12021 Publishing.
- 12022 Peters, M. E., M. Neumann, M. Iyyer, M. Gardner, C. Clark, K. Lee, and L. Zettlemoyer
12023 (2018). Deep contextualized word representations. In *Proceedings of the North American*
12024 *Chapter of the Association for Computational Linguistics (NAACL)*.

- 12025 Peterson, W. W., T. G. Birdsall, and W. C. Fox (1954). The theory of signal detectability.
12026 *Transactions of the IRE professional group on information theory* 4(4), 171–212.
- 12027 Petrov, S., L. Barrett, R. Thibaux, and D. Klein (2006). Learning accurate, compact, and in-
12028 terpretable tree annotation. In *Proceedings of the Association for Computational Linguistics*
12029 (*ACL*).
- 12030 Petrov, S., D. Das, and R. McDonald (2012, May). A universal part-of-speech tagset. In
12031 *Proceedings of LREC*.
- 12032 Petrov, S. and R. McDonald (2012). Overview of the 2012 shared task on parsing the web.
12033 In *Notes of the First Workshop on Syntactic Analysis of Non-Canonical Language (SANCL)*,
12034 Volume 59.
- 12035 Pinker, S. (2003). *The language instinct: How the mind creates language*. Penguin UK.
- 12036 Pinter, Y., R. Guthrie, and J. Eisenstein (2017). Mimicking word embeddings using
12037 subword RNNs. In *Proceedings of Empirical Methods for Natural Language Processing*
12038 (*EMNLP*).
- 12039 Pitler, E., A. Louis, and A. Nenkova (2009). Automatic sense prediction for implicit dis-
12040 course relations in text. In *Proceedings of the Association for Computational Linguistics*
12041 (*ACL*).
- 12042 Pitler, E. and A. Nenkova (2009). Using syntax to disambiguate explicit discourse con-
12043 nectives in text. In *Proceedings of the Association for Computational Linguistics (ACL)*, pp.
12044 13–16.
- 12045 Pitler, E., M. Raghupathy, H. Mehta, A. Nenkova, A. Lee, and A. Joshi (2008). Easily iden-
12046 tifiable discourse relations. In *Proceedings of the International Conference on Computational*
12047 *Linguistics (COLING)*, pp. 87–90.
- 12048 Plank, B., A. Søgaard, and Y. Goldberg (2016). Multilingual part-of-speech tagging with
12049 bidirectional long short-term memory models and auxiliary loss. In *Proceedings of the*
12050 *Association for Computational Linguistics (ACL)*.
- 12051 Poesio, M., R. Stevenson, B. Di Eugenio, and J. Hitzeman (2004). Centering: A parametric
12052 theory and its instantiations. *Computational linguistics* 30(3), 309–363.
- 12053 Polanyi, L. and A. Zaenen (2006). Contextual valence shifters. In *Computing attitude and*
12054 *affect in text: Theory and applications*. Springer.
- 12055 Ponzetto, S. P. and M. Strube (2006). Exploiting semantic role labeling, wordnet and
12056 wikipedia for coreference resolution. In *Proceedings of the North American Chapter of*
12057 *the Association for Computational Linguistics (NAACL)*, pp. 192–199.

- 12058 Ponzetto, S. P. and M. Strube (2007). Knowledge derived from wikipedia for computing
12059 semantic relatedness. *Journal of Artificial Intelligence Research* 30, 181–212.
- 12060 Poon, H. and P. Domingos (2008). Joint unsupervised coreference resolution with markov
12061 logic. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pp.
12062 650–659.
- 12063 Poon, H. and P. Domingos (2009). Unsupervised semantic parsing. In *Proceedings of Em-
12064 pirical Methods for Natural Language Processing (EMNLP)*, pp. 1–10.
- 12065 Popel, M., D. Marecek, J. Stepánek, D. Zeman, and Z. Zabokrtský (2013). Coordination
12066 structures in dependency treebanks. In *Proceedings of the Association for Computational
12067 Linguistics (ACL)*, pp. 517–527.
- 12068 Popescu, A.-M., O. Etzioni, and H. Kautz (2003). Towards a theory of natural language
12069 interfaces to databases. In *Proceedings of Intelligent User Interfaces (IUI)*, pp. 149–157.
- 12070 Poplack, S. (1980). Sometimes i'll start a sentence in spanish y termino en español: toward
12071 a typology of code-switching1. *Linguistics* 18(7-8), 581–618.
- 12072 Porter, M. F. (1980). An algorithm for suffix stripping. *Program* 14(3), 130–137.
- 12073 Prabhakaran, V., O. Rambow, and M. Diab (2010). Automatic committed belief tagging.
12074 In *Proceedings of the International Conference on Computational Linguistics (COLING)*, pp.
12075 1014–1022.
- 12076 Pradhan, S., X. Luo, M. Recasens, E. Hovy, V. Ng, and M. Strube (2014). Scoring corefer-
12077 ence partitions of predicted mentions: A reference implementation. In *Proceedings of the
12078 Association for Computational Linguistics (ACL)*, pp. 30–35.
- 12079 Pradhan, S., L. Ramshaw, M. Marcus, M. Palmer, R. Weischedel, and N. Xue (2011).
12080 CoNLL-2011 shared task: Modeling unrestricted coreference in OntoNotes. In *Proceed-
12081 ings of the Fifteenth Conference on Computational Natural Language Learning: Shared Task*,
12082 pp. 1–27. Association for Computational Linguistics.
- 12083 Pradhan, S., W. Ward, K. Hacioglu, J. H. Martin, and D. Jurafsky (2005). Semantic role
12084 labeling using different syntactic views. In *Proceedings of the Association for Computational
12085 Linguistics (ACL)*, pp. 581–588.
- 12086 Prasad, R., N. Dinesh, A. Lee, E. Miltsakaki, L. Robaldo, A. Joshi, and B. Webber (2008).
12087 The Penn Discourse Treebank 2.0. In *Proceedings of LREC*.
- 12088 Punyakanok, V., D. Roth, and W.-t. Yih (2008). The importance of syntactic parsing and
12089 inference in semantic role labeling. *Computational Linguistics* 34(2), 257–287.

- 12090 Pustejovsky, J., P. Hanks, R. Sauri, A. See, R. Gaizauskas, A. Setzer, D. Radev, B. Sundheim,
 12091 D. Day, L. Ferro, et al. (2003). The timebank corpus. In *Corpus linguistics*, Volume 2003,
 12092 pp. 40. Lancaster, UK.
- 12093 Pustejovsky, J., B. Ingria, R. Sauri, J. Castano, J. Littman, R. Gaizauskas, A. Setzer, G. Katz,
 12094 and I. Mani (2005). The specification language timeml. In *The language of time: A reader*,
 12095 pp. 545–557. Oxford University Press.
- 12096 Qin, L., Z. Zhang, H. Zhao, Z. Hu, and E. Xing (2017). Adversarial connective-exploiting
 12097 networks for implicit discourse relation classification. In *Proceedings of the Association
 12098 for Computational Linguistics (ACL)*, pp. 1006–1017.
- 12099 Qiu, G., B. Liu, J. Bu, and C. Chen (2011). Opinion word expansion and target extraction
 12100 through double propagation. *Computational linguistics* 37(1), 9–27.
- 12101 Quattoni, A., S. Wang, L.-P. Morency, M. Collins, and T. Darrell (2007). Hidden conditional
 12102 random fields. *IEEE transactions on pattern analysis and machine intelligence* 29(10).
- 12103 Rahman, A. and V. Ng (2011). Narrowing the modeling gap: a cluster-ranking approach
 12104 to coreference resolution. *Journal of Artificial Intelligence Research* 40, 469–521.
- 12105 Rajpurkar, P., J. Zhang, K. Lopyrev, and P. Liang (2016). Squad: 100,000+ questions for
 12106 machine comprehension of text. In *Proceedings of Empirical Methods for Natural Language
 12107 Processing (EMNLP)*, pp. 2383–2392.
- 12108 Ranzato, M., S. Chopra, M. Auli, and W. Zaremba (2016). Sequence level training with
 12109 recurrent neural networks. In *Proceedings of the International Conference on Learning Rep-
 12110 resentations (ICLR)*.
- 12111 Rao, D., D. Yarowsky, A. Shreevats, and M. Gupta (2010). Classifying latent user attributes
 12112 in twitter. In *Proceedings of Workshop on Search and mining user-generated contents*.
- 12113 Ratinov, L. and D. Roth (2009). Design challenges and misconceptions in named entity
 12114 recognition. In *Proceedings of the Thirteenth Conference on Computational Natural Language
 12115 Learning*, pp. 147–155. Association for Computational Linguistics.
- 12116 Ratinov, L., D. Roth, D. Downey, and M. Anderson (2011). Local and global algorithms
 12117 for disambiguation to wikipedia. In *Proceedings of the Association for Computational Lin-
 12118 guistics (ACL)*, pp. 1375–1384.
- 12119 Ratliff, N. D., J. A. Bagnell, and M. Zinkevich (2007). (approximate) subgradient methods
 12120 for structured prediction. In *Proceedings of Artificial Intelligence and Statistics (AISTATS)*,
 12121 pp. 380–387.

- 12122 Ratnaparkhi, A. (1996). A maximum entropy model for part-of-speech tagging. In *emnlp*,
12123 pp. 133–142.
- 12124 Ratnaparkhi, A., J. Reynar, and S. Roukos (1994). A maximum entropy model for preposi-
12125 tional phrase attachment. In *Proceedings of the workshop on Human Language Technology*,
12126 pp. 250–255. Association for Computational Linguistics.
- 12127 Read, J. (2005). Using emoticons to reduce dependency in machine learning techniques for
12128 sentiment classification. In *Proceedings of the ACL student research workshop*, pp. 43–48.
12129 Association for Computational Linguistics.
- 12130 Reisinger, D., R. Rudinger, F. Ferraro, C. Harman, K. Rawlins, and B. V. Durme (2015).
12131 Semantic proto-roles. *Transactions of the Association for Computational Linguistics* 3, 475–
12132 488.
- 12133 Reisinger, J. and R. J. Mooney (2010). Multi-prototype vector-space models of word mean-
12134 ing. In *Proceedings of the North American Chapter of the Association for Computational Lin-*
12135 *guistics (NAACL)*, pp. 109–117.
- 12136 Reiter, E. and R. Dale (2000). *Building natural language generation systems*. Cambridge
12137 university press.
- 12138 Resnik, P., M. B. Olsen, and M. Diab (1999). The bible as a parallel corpus: Annotating the
12139 ‘book of 2000 tongues’. *Computers and the Humanities* 33(1-2), 129–153.
- 12140 Resnik, P. and N. A. Smith (2003). The web as a parallel corpus. *Computational Linguis-*
12141 *tics* 29(3), 349–380.
- 12142 Ribeiro, F. N., M. Araújo, P. Gonçalves, M. A. Gonçalves, and F. Benevenuto (2016).
12143 Sentibench-a benchmark comparison of state-of-the-practice sentiment analysis meth-
12144 ods. *EPJ Data Science* 5(1), 1–29.
- 12145 Richardson, M., C. J. Burges, and E. Renshaw (2013). MCTest: A challenge dataset for
12146 the open-domain machine comprehension of text. In *Proceedings of Empirical Methods for*
12147 *Natural Language Processing (EMNLP)*, pp. 193–203.
- 12148 Riedel, S., L. Yao, and A. McCallum (2010). Modeling relations and their mentions without
12149 labeled text. In *Proceedings of the European Conference on Machine Learning and Principles*
12150 *and Practice of Knowledge Discovery in Databases (ECML)*, pp. 148–163.
- 12151 Riedl, M. O. and R. M. Young (2010). Narrative planning: Balancing plot and character.
12152 *Journal of Artificial Intelligence Research* 39, 217–268.
- 12153 Rieser, V. and O. Lemon (2011). *Reinforcement learning for adaptive dialogue systems: a data-*
12154 *driven methodology for dialogue management and natural language generation*. Springer Sci-
12155 ence & Business Media.

- 12156 Riloff, E. (1996). Automatically generating extraction patterns from untagged text. In
 12157 *Proceedings of the national conference on artificial intelligence*, pp. 1044–1049.
- 12158 Riloff, E. and J. Wiebe (2003). Learning extraction patterns for subjective expressions. In
 12159 *Proceedings of the 2003 conference on Empirical methods in natural language processing*, pp.
 12160 105–112. Association for Computational Linguistics.
- 12161 Ritchie, G. (2001). Current directions in computational humour. *Artificial Intelligence Re-*
 12162 *view* 16(2), 119–135.
- 12163 Ritter, A., C. Cherry, and W. B. Dolan (2011). Data-driven response generation in social
 12164 media. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pp.
 12165 583–593.
- 12166 Roark, B., M. Saraclar, and M. Collins (2007). Discriminative i_l / n_l/i_l -gram language
 12167 modeling. *Computer Speech & Language* 21(2), 373–392.
- 12168 Robert, C. and G. Casella (2013). *Monte Carlo statistical methods*. Springer Science & Busi-
 12169 ness Media.
- 12170 Rosenfeld, R. (1996). A maximum entropy approach to adaptive statistical language mod-
 12171 elling. *Computer Speech & Language* 10(3), 187–228.
- 12172 Ross, S., G. Gordon, and D. Bagnell (2011). A reduction of imitation learning and struc-
 12173 tured prediction to no-regret online learning. In *Proceedings of Artificial Intelligence and*
 12174 *Statistics (AISTATS)*, pp. 627–635.
- 12175 Roy, N., J. Pineau, and S. Thrun (2000). Spoken dialogue management using probabilistic
 12176 reasoning. In *Proceedings of the Association for Computational Linguistics (ACL)*, pp. 93–
 12177 100.
- 12178 Rudnicky, A. and W. Xu (1999). An agenda-based dialog management architecture for
 12179 spoken language systems. In *IEEE Automatic Speech Recognition and Understanding Work-
 12180 shop*, Volume 13.
- 12181 Rush, A. M., S. Chopra, and J. Weston (2015). A neural attention model for abstractive sen-
 12182 tence summarization. In *Proceedings of Empirical Methods for Natural Language Processing*
 12183 (*EMNLP*), pp. 379–389.
- 12184 Rush, A. M., D. Sontag, M. Collins, and T. Jaakkola (2010). On dual decomposition and
 12185 linear programming relaxations for natural language processing. In *Proceedings of Em-
 12186 pirical Methods for Natural Language Processing (EMNLP)*, pp. 1–11.
- 12187 Russell, S. J. and P. Norvig (2009). *Artificial intelligence: a modern approach* (3rd ed.). Prentice
 12188 Hall.

- 12189 Rutherford, A., V. Demberg, and N. Xue (2017). A systematic study of neural discourse
12190 models for implicit discourse relation. In *Proceedings of the European Chapter of the Asso-*
12191 *ciation for Computational Linguistics (EACL)*, pp. 281–291.
- 12192 Rutherford, A. T. and N. Xue (2014). Discovering implicit discourse relations through
12193 brown cluster pair representation and coreference patterns. In *Proceedings of the Euro-*
12194 *pean Chapter of the Association for Computational Linguistics (EACL)*.
- 12195 Sag, I. A., T. Baldwin, F. Bond, A. Copestake, and D. Flickinger (2002). Multiword expres-
12196 sions: A pain in the neck for nlp. In *International Conference on Intelligent Text Processing*
12197 and *Computational Linguistics*, pp. 1–15. Springer.
- 12198 Sagae, K. (2009). Analysis of discourse structure with syntactic dependencies and data-
12199 driven shift-reduce parsing. In *Proceedings of the 11th International Conference on Parsing*
12200 *Technologies*, pp. 81–84.
- 12201 Santos, C. D. and B. Zadrozny (2014). Learning character-level representations for part-of-
12202 speech tagging. In *Proceedings of the International Conference on Machine Learning (ICML)*,
12203 pp. 1818–1826.
- 12204 Sato, M.-A. and S. Ishii (2000). On-line em algorithm for the normalized gaussian network.
12205 *Neural computation* 12(2), 407–432.
- 12206 Saurí, R. and J. Pustejovsky (2009). Factbank: a corpus annotated with event factuality.
12207 *Language resources and evaluation* 43(3), 227.
- 12208 Saxe, A. M., J. L. McClelland, and S. Ganguli (2014). Exact solutions to the nonlinear
12209 dynamics of learning in deep linear neural networks. In *Proceedings of the International*
12210 *Conference on Learning Representations (ICLR)*.
- 12211 Schank, R. C. and R. Abelson (1977). *Scripts, goals, plans, and understanding*. Hillsdale, NJ:
12212 Erlbaum.
- 12213 Schapire, R. E. and Y. Singer (2000). Boostexter: A boosting-based system for text catego-
12214 *Machine learning* 39(2-3), 135–168.
- 12215 Schaul, T., S. Zhang, and Y. LeCun (2013). No more pesky learning rates. In *Proceedings of*
12216 *the International Conference on Machine Learning (ICML)*, pp. 343–351.
- 12217 Schnabel, T., I. Labutov, D. Mimno, and T. Joachims (2015). Evaluation methods for un-
12218 supervised word embeddings. In *Proceedings of Empirical Methods for Natural Language*
12219 *Processing (EMNLP)*, pp. 298–307.
- 12220 Schneider, N., J. Flanigan, and T. O’Gorman (2015). The logic of amr: Practical, unified,
12221 graph-based sentence semantics for nlp. In *Proceedings of the North American Chapter of*
12222 *the Association for Computational Linguistics (NAACL)*, pp. 4–5.

- 12223 Schütze, H. (1998). Automatic word sense discrimination. *Computational linguistics* 24(1),
12224 97–123.
- 12225 Schwarm, S. E. and M. Ostendorf (2005). Reading level assessment using support vector
12226 machines and statistical language models. In *Proceedings of the Association for Compu-*
12227 *tational Linguistics (ACL)*, pp. 523–530.
- 12228 See, A., P. J. Liu, and C. D. Manning (2017). Get to the point: Summarization with pointer-
12229 generator networks. In *Proceedings of the Association for Computational Linguistics (ACL)*,
12230 pp. 1073–1083.
- 12231 Sennrich, R., B. Haddow, and A. Birch (2016). Neural machine translation of rare words
12232 with subword units. In *Proceedings of the Association for Computational Linguistics (ACL)*,
12233 pp. 1715–1725.
- 12234 Serban, I. V., A. Sordoni, Y. Bengio, A. C. Courville, and J. Pineau (2016). Building end-to-
12235 end dialogue systems using generative hierarchical neural network models. In *Proceed-*
12236 *ings of the National Conference on Artificial Intelligence (AAAI)*, pp. 3776–3784.
- 12237 Settles, B. (2012). Active learning. *Synthesis Lectures on Artificial Intelligence and Machine*
12238 *Learning* 6(1), 1–114.
- 12239 Shang, L., Z. Lu, and H. Li (2015). Neural responding machine for short-text conversation.
12240 In *Proceedings of the Association for Computational Linguistics (ACL)*, pp. 1577–1586.
- 12241 Shen, D. and M. Lapata (2007). Using semantic roles to improve question answering. In
12242 *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pp. 12–21.
- 12243 Shen, S., Y. Cheng, Z. He, W. He, H. Wu, M. Sun, and Y. Liu (2016). Minimum risk train-
12244 ing for neural machine translation. In *Proceedings of the Association for Computation-*
12245 *al Linguistics (ACL)*, pp. 1683–1692.
- 12246 Shen, W., J. Wang, and J. Han (2015). Entity linking with a knowledge base: Issues, tech-
12247 niques, and solutions. *IEEE Transactions on Knowledge and Data Engineering* 27(2), 443–
12248 460.
- 12249 Shieber, S. M. (1985). Evidence against the context-freeness of natural language. *Linguistics*
12250 and *Philosophy* 8(3), 333–343.
- 12251 Siegelmann, H. T. and E. D. Sontag (1995). On the computational power of neural nets.
12252 *Journal of computer and system sciences* 50(1), 132–150.
- 12253 Singh, S., A. Subramanya, F. Pereira, and A. McCallum (2011). Large-scale cross-
12254 document coreference using distributed inference and hierarchical models. In *Proceed-*
12255 *ings of the Association for Computational Linguistics (ACL)*, pp. 793–803.

- 12256 Sipser, M. (2012). *Introduction to the Theory of Computation*. Cengage Learning.
- 12257 Smith, D. A. and J. Eisner (2006). Minimum risk annealing for training log-linear models.
- 12258 In *Proceedings of the Association for Computational Linguistics (ACL)*, pp. 787–794.
- 12259 Smith, D. A. and J. Eisner (2008). Dependency parsing by belief propagation. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pp. 145–156.
- 12260
- 12261 Smith, D. A. and N. A. Smith (2007). Probabilistic models of nonprojective dependency trees. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pp. 132–140.
- 12262
- 12263
- 12264 Smith, N. A. (2011). Linguistic structure prediction. *Synthesis Lectures on Human Language Technologies* 4(2), 1–274.
- 12265
- 12266 Snover, M., B. Dorr, R. Schwartz, L. Micciulla, and J. Makhoul (2006). A study of translation edit rate with targeted human annotation. In *Proceedings of association for machine translation in the Americas*, Volume 200.
- 12267
- 12268
- 12269 Snow, R., B. O'Connor, D. Jurafsky, and A. Y. Ng (2008). Cheap and fast—but is it good?: evaluating non-expert annotations for natural language tasks. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pp. 254–263.
- 12270
- 12271
- 12272 Snyder, B. and R. Barzilay (2007). Database-text alignment via structured multilabel classification. In *Proceedings of the International Joint Conference on Artificial Intelligence (IJCAI)*, pp. 1713–1718.
- 12273
- 12274
- 12275 Socher, R., J. Bauer, C. D. Manning, and A. Y. Ng (2013). Parsing with compositional vector grammars. In *Proceedings of the Association for Computational Linguistics (ACL)*.
- 12276
- 12277 Socher, R., B. Huval, C. D. Manning, and A. Y. Ng (2012). Semantic compositionality through recursive matrix-vector spaces. In *Proceedings of the 2012 Joint Conference on Empirical Methods in Natural Language Processing and Computational Natural Language Learning*, pp. 1201–1211. Association for Computational Linguistics.
- 12278
- 12279
- 12280
- 12281 Socher, R., A. Perelygin, J. Y. Wu, J. Chuang, C. D. Manning, A. Y. Ng, and C. Potts (2013). Recursive deep models for semantic compositionality over a sentiment treebank. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*.
- 12282
- 12283
- 12284 Søgaard, A. (2013). Semi-supervised learning and domain adaptation in natural language processing. *Synthesis Lectures on Human Language Technologies* 6(2), 1–103.
- 12285
- 12286 Solorio, T. and Y. Liu (2008). Learning to predict code-switching points. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pp. 973–981. Association for Computational Linguistics.
- 12287
- 12288

- 12289 Somasundaran, S., G. Namata, J. Wiebe, and L. Getoor (2009). Supervised and unsupervised methods in employing discourse relations for improving opinion polarity classification. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*.
- 12292 Somasundaran, S. and J. Wiebe (2009). Recognizing stances in online debates. In *Proceedings of the Association for Computational Linguistics (ACL)*, pp. 226–234.
- 12294 Song, L., B. Boots, S. M. Siddiqi, G. J. Gordon, and A. J. Smola (2010). Hilbert space embeddings of hidden markov models. In *Proceedings of the International Conference on Machine Learning (ICML)*, pp. 991–998.
- 12297 Song, L., Y. Zhang, X. Peng, Z. Wang, and D. Gildea (2016). Amr-to-text generation as a traveling salesman problem. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pp. 2084–2089.
- 12300 Soon, W. M., H. T. Ng, and D. C. Y. Lim (2001). A machine learning approach to coreference resolution of noun phrases. *Computational linguistics* 27(4), 521–544.
- 12302 Sordoni, A., M. Galley, M. Auli, C. Brockett, Y. Ji, M. Mitchell, J.-Y. Nie, J. Gao, and B. Dolan (2015). A neural network approach to context-sensitive generation of conversational responses. In *Proceedings of the North American Chapter of the Association for Computational Linguistics (NAACL)*.
- 12306 Soriceut, R. and D. Marcu (2003). Sentence level discourse parsing using syntactic and lexical information. In *Proceedings of the North American Chapter of the Association for Computational Linguistics (NAACL)*, pp. 149–156.
- 12309 Sowa, J. F. (2000). *Knowledge representation: logical, philosophical, and computational foundations*. Pacific Grove, CA: Brooks/Cole.
- 12311 Spärck Jones, K. (1972). A statistical interpretation of term specificity and its application in retrieval. *Journal of documentation* 28(1), 11–21.
- 12313 Spitkovsky, V. I., H. Alshawi, D. Jurafsky, and C. D. Manning (2010). Viterbi training improves unsupervised dependency parsing. In *CONLL*, pp. 9–17.
- 12315 Sporleder, C. and M. Lapata (2005). Discourse chunking and its application to sentence compression. In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pp. 257–264.
- 12318 Sproat, R., A. Black, S. Chen, S. Kumar, M. Ostendorf, and C. Richards (2001). Normalization of non-standard words. *Computer Speech & Language* 15(3), 287–333.
- 12320 Sproat, R., W. Gale, C. Shih, and N. Chang (1996). A stochastic finite-state word-segmentation algorithm for chinese. *Computational linguistics* 22(3), 377–404.

- 12322 Sra, S., S. Nowozin, and S. J. Wright (2012). *Optimization for machine learning*. MIT Press.
- 12323 Srivastava, N., G. Hinton, A. Krizhevsky, I. Sutskever, and R. Salakhutdinov (2014).
12324 Dropout: A simple way to prevent neural networks from overfitting. *The Journal of
12325 Machine Learning Research* 15(1), 1929–1958.
- 12326 Srivastava, R. K., K. Greff, and J. Schmidhuber (2015). Training very deep networks. In
12327 *Neural Information Processing Systems (NIPS)*, pp. 2377–2385.
- 12328 Stab, C. and I. Gurevych (2014a). Annotating argument components and relations in per-
12329 suasive essays. In *Proceedings of the International Conference on Computational Linguistics
12330 (COLING)*, pp. 1501–1510.
- 12331 Stab, C. and I. Gurevych (2014b). Identifying argumentative discourse structures in per-
12332 suasive essays. In *Proceedings of the 2014 Conference on Empirical Methods in Natural Lan-
12333 guage Processing (EMNLP)*, pp. 46–56.
- 12334 Stede, M. (2011, nov). *Discourse Processing*, Volume 4 of *Synthesis Lectures on Human Lan-
12335 guage Technologies*. Morgan & Claypool Publishers.
- 12336 Steedman, M. and J. Baldridge (2011). Combinatory categorial grammar. In *Non-
12337 Transformational Syntax: Formal and Explicit Models of Grammar*. Wiley-Blackwell.
- 12338 Stenetorp, P., S. Pyysalo, G. Topić, T. Ohta, S. Ananiadou, and J. Tsujii (2012). Brat: a web-
12339 based tool for nlp-assisted text annotation. In *Proceedings of the European Chapter of the
12340 Association for Computational Linguistics (EACL)*, pp. 102–107.
- 12341 Stern, M., J. Andreas, and D. Klein (2017). A minimal span-based neural constituency
12342 parser. In *Proceedings of the Association for Computational Linguistics (ACL)*.
- 12343 Stolcke, A., K. Ries, N. Coccaro, E. Shriberg, R. Bates, D. Jurafsky, P. Taylor, R. Martin,
12344 C. Van Ess-Dykema, and M. Meteer (2000). Dialogue act modeling for automatic tag-
12345 ging and recognition of conversational speech. *Computational linguistics* 26(3), 339–373.
- 12346 Stone, P. J. (1966). *The General Inquirer: A Computer Approach to Content Analysis*. The MIT
12347 Press.
- 12348 Stoyanov, V., N. Gilbert, C. Cardie, and E. Riloff (2009). Conundrums in noun phrase
12349 coreference resolution: Making sense of the state-of-the-art. In *Proceedings of the Associa-
12350 tion for Computational Linguistics (ACL)*, pp. 656–664.
- 12351 Strang, G. (2016). *Introduction to linear algebra* (Fifth ed.). Wellesley, MA: Wellesley-
12352 Cambridge Press.

- 12353 Strubell, E., P. Verga, D. Belanger, and A. McCallum (2017). Fast and accurate entity recognition
 12354 with iterated dilated convolutions. In *Proceedings of Empirical Methods for Natural
 12355 Language Processing (EMNLP)*.
- 12356 Suchanek, F. M., G. Kasneci, and G. Weikum (2007). Yago: a core of semantic knowledge.
 12357 In *Proceedings of the Conference on World-Wide Web (WWW)*, pp. 697–706.
- 12358 Sun, X., T. Matsuzaki, D. Okanohara, and J. Tsujii (2009). Latent variable perceptron algo-
 12359 rithm for structured classification. In *Proceedings of the International Joint Conference on
 12360 Artificial Intelligence (IJCAI)*, Volume 9, pp. 1236–1242.
- 12361 Sun, Y., L. Lin, D. Tang, N. Yang, Z. Ji, and X. Wang (2015). Modeling mention, context
 12362 and entity with neural networks for entity disambiguation. In *IJCAI*, pp. 1333–1339.
- 12363 Sundermeyer, M., R. Schlüter, and H. Ney (2012). Lstm neural networks for language
 12364 modeling. In *INTERSPEECH*.
- 12365 Surdeanu, M., J. Tibshirani, R. Nallapati, and C. D. Manning (2012). Multi-instance multi-
 12366 label learning for relation extraction. In *Proceedings of Empirical Methods for Natural Lan-
 12367 guage Processing (EMNLP)*, pp. 455–465.
- 12368 Sutskever, I., O. Vinyals, and Q. V. Le (2014). Sequence to sequence learning with neural
 12369 networks. In *Neural Information Processing Systems (NIPS)*, pp. 3104–3112.
- 12370 Sutton, R. S. and A. G. Barto (1998). *Reinforcement learning: An introduction*, Volume 1. MIT
 12371 press Cambridge.
- 12372 Sutton, R. S., D. A. McAllester, S. P. Singh, and Y. Mansour (2000). Policy gradient methods
 12373 for reinforcement learning with function approximation. In *Neural Information Process-
 12374 ing Systems (NIPS)*, pp. 1057–1063.
- 12375 Taboada, M., J. Brooke, M. Tofiloski, K. Voll, and M. Stede (2011). Lexicon-based methods
 12376 for sentiment analysis. *Computational linguistics* 37(2), 267–307.
- 12377 Taboada, M. and W. C. Mann (2006). Rhetorical structure theory: Looking back and mov-
 12378 ing ahead. *Discourse studies* 8(3), 423–459.
- 12379 Täckström, O., K. Ganchev, and D. Das (2015). Efficient inference and structured learning
 12380 for semantic role labeling. *Transactions of the Association for Computational Linguistics* 3,
 12381 29–41.
- 12382 Täckström, O., R. McDonald, and J. Uszkoreit (2012). Cross-lingual word clusters for
 12383 direct transfer of linguistic structure. In *Proceedings of the North American Chapter of the
 12384 Association for Computational Linguistics (NAACL)*, pp. 477–487.

- 12385 Tang, D., B. Qin, and T. Liu (2015). Document modeling with gated recurrent neural net-
12386 work for sentiment classification. In *Proceedings of Empirical Methods for Natural Language*
12387 *Processing (EMNLP)*, pp. 1422–1432.
- 12388 Taskar, B., C. Guestrin, and D. Koller (2003). Max-margin markov networks. In *Neural*
12389 *Information Processing Systems (NIPS)*.
- 12390 Tausczik, Y. R. and J. W. Pennebaker (2010). The psychological meaning of words: LIWC
12391 and computerized text analysis methods. *Journal of Language and Social Psychology* 29(1),
12392 24–54.
- 12393 Teh, Y. W. (2006). A hierarchical bayesian language model based on pitman-yor processes.
12394 In *Proceedings of the Association for Computational Linguistics (ACL)*, pp. 985–992.
- 12395 Tesnière, L. (1966). *Éléments de syntaxe structurale* (second ed.). Paris: Klincksieck.
- 12396 Teufel, S., J. Carletta, and M. Moens (1999). An annotation scheme for discourse-level
12397 argumentation in research articles. In *Proceedings of the European Chapter of the Association*
12398 *for Computational Linguistics (EACL)*, pp. 110–117.
- 12399 Teufel, S. and M. Moens (2002). Summarizing scientific articles: experiments with rele-
12400 vance and rhetorical status. *Computational linguistics* 28(4), 409–445.
- 12401 Thomas, M., B. Pang, and L. Lee (2006). Get out the vote: Determining support or opposi-
12402 tion from Congressional floor-debate transcripts. In *Proceedings of Empirical Methods for*
12403 *Natural Language Processing (EMNLP)*, pp. 327–335.
- 12404 Tibshirani, R. (1996). Regression shrinkage and selection via the lasso. *Journal of the Royal*
12405 *Statistical Society. Series B (Methodological)*, 267–288.
- 12406 Titov, I. and J. Henderson (2007). Constituent parsing with incremental sigmoid belief
12407 networks. In *Proceedings of the Association for Computational Linguistics (ACL)*, pp. 632–
12408 639.
- 12409 Toutanova, K., D. Klein, C. D. Manning, and Y. Singer (2003). Feature-rich part-of-speech
12410 tagging with a cyclic dependency network. In *Proceedings of the North American Chapter*
12411 *of the Association for Computational Linguistics (NAACL)*.
- 12412 Trivedi, R. and J. Eisenstein (2013). Discourse connectors for latent subjectivity in senti-
12413 ment analysis. In *Proceedings of the North American Chapter of the Association for Compu-*
12414 *tational Linguistics (NAACL)*, pp. 808–813.
- 12415 Tromble, R. W. and J. Eisner (2006). A fast finite-state relaxation method for enforcing
12416 global constraints on sequence decoding. In *Proceedings of the North American Chapter of*
12417 *the Association for Computational Linguistics (NAACL)*, pp. 423.

- 12418 Tsochantaridis, I., T. Hofmann, T. Joachims, and Y. Altun (2004). Support vector machine
 12419 learning for interdependent and structured output spaces. In *Proceedings of the twenty-*
 12420 *first international conference on Machine learning*, pp. 104. ACM.
- 12421 Tsvetkov, Y., M. Faruqui, W. Ling, G. Lample, and C. Dyer (2015). Evaluation of word
 12422 vector representations by subspace alignment. In *Proceedings of Empirical Methods for*
 12423 *Natural Language Processing (EMNLP)*, pp. 2049–2054.
- 12424 Tu, Z., Z. Lu, Y. Liu, X. Liu, and H. Li (2016). Modeling coverage for neural machine
 12425 translation. In *Proceedings of the Association for Computational Linguistics (ACL)*, pp. 76–
 12426 85.
- 12427 Turian, J., L. Ratinov, and Y. Bengio (2010). Word representations: a simple and general
 12428 method for semi-supervised learning. In *Proceedings of the Association for Computational*
 12429 *Linguistics (ACL)*, pp. 384–394.
- 12430 Turing, A. M. (2009). Computing machinery and intelligence. In *Parsing the Turing Test*,
 12431 pp. 23–65. Springer.
- 12432 Turney, P. D. and P. Pantel (2010). From frequency to meaning: Vector space models of
 12433 semantics. *Journal of Artificial Intelligence Research* 37, 141–188.
- 12434 Tutin, A. and R. Kittredge (1992). Lexical choice in context: generating procedural texts.
 12435 In *Proceedings of the International Conference on Computational Linguistics (COLING)*, pp.
 12436 763–769.
- 12437 Twain, M. (1997). *A Tramp Abroad*. New York: Penguin.
- 12438 Tzeng, E., J. Hoffman, T. Darrell, and K. Saenko (2015). Simultaneous deep transfer across
 12439 domains and tasks. In *Proceedings of the IEEE International Conference on Computer Vision*,
 12440 pp. 4068–4076.
- 12441 Usunier, N., D. Buffoni, and P. Gallinari (2009). Ranking with ordered weighted pairwise
 12442 classification. In *Proceedings of the International Conference on Machine Learning (ICML)*,
 12443 pp. 1057–1064.
- 12444 Uthus, D. C. and D. W. Aha (2013). The ubuntu chat corpus for multiparicipant chat
 12445 analysis. In *AAAI Spring Symposium: Analyzing Microtext*, Volume 13, pp. 01.
- 12446 Utiyama, M. and H. Isahara (2001). A statistical model for domain-independent text seg-
 12447 mentation. In *Proceedings of the 39th Annual Meeting on Association for Computational*
 12448 *Linguistics*, pp. 499–506. Association for Computational Linguistics.

- 12449 Utiyama, M. and H. Isahara (2007). A comparison of pivot methods for phrase-based
12450 statistical machine translation. In *Human Language Technologies 2007: The Conference of
12451 the North American Chapter of the Association for Computational Linguistics; Proceedings of
12452 the Main Conference*, pp. 484–491.
- 12453 Uzuner, Ö., X. Zhang, and T. Sibanda (2009). Machine learning and rule-based approaches
12454 to assertion classification. *Journal of the American Medical Informatics Association* 16(1),
12455 109–115.
- 12456 Vadas, D. and J. R. Curran (2011). Parsing noun phrases in the penn treebank. *Computa-
12457 tional Linguistics* 37(4), 753–809.
- 12458 Van Eynde, F. (2006). NP-internal agreement and the structure of the noun phrase. *Journal
12459 of Linguistics* 42(1), 139–186.
- 12460 Van Gael, J., A. Vlachos, and Z. Ghahramani (2009). The infinite hmm for unsuper-
12461 vised pos tagging. In *Proceedings of Empirical Methods for Natural Language Processing
12462 (EMNLP)*, pp. 678–687.
- 12463 Vaswani, A., N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and
12464 I. Polosukhin (2017). Attention is all you need. In *Neural Information Processing Systems
12465 (NIPS)*, pp. 6000–6010.
- 12466 Velldal, E., L. Øvrelid, J. Read, and S. Oepen (2012). Speculation and negation: Rules,
12467 rankers, and the role of syntax. *Computational linguistics* 38(2), 369–410.
- 12468 Versley, Y. (2011). Towards finer-grained tagging of discourse connectives. In *Proceedings
12469 of the Workshop Beyond Semantics: Corpus-based Investigations of Pragmatic and Discourse
12470 Phenomena*, pp. 2–63.
- 12471 Vilain, M., J. Burger, J. Aberdeen, D. Connolly, and L. Hirschman (1995). A model-
12472 theoretic coreference scoring scheme. In *Proceedings of the 6th conference on Message
12473 understanding*, pp. 45–52. Association for Computational Linguistics.
- 12474 Vincent, P., H. Larochelle, I. Lajoie, Y. Bengio, and P.-A. Manzagol (2010). Stacked de-
12475 noising autoencoders: Learning useful representations in a deep network with a local
12476 denoising criterion. *Journal of Machine Learning Research* 11(Dec), 3371–3408.
- 12477 Vincze, V., G. Szarvas, R. Farkas, G. Móra, and J. Csirik (2008). The bioscope corpus:
12478 biomedical texts annotated for uncertainty, negation and their scopes. *BMC bioinformat-
12479 ics* 9(11), S9.
- 12480 Vinyals, O., A. Toshev, S. Bengio, and D. Erhan (2015). Show and tell: A neural image cap-
12481 tion generator. In *Computer Vision and Pattern Recognition (CVPR), 2015 IEEE Conference
12482 on*, pp. 3156–3164. IEEE.

- 12483 Viterbi, A. (1967). Error bounds for convolutional codes and an asymptotically optimum
 12484 decoding algorithm. *IEEE transactions on Information Theory* 13(2), 260–269.
- 12485 Voll, K. and M. Taboada (2007). Not all words are created equal: Extracting semantic
 12486 orientation as a function of adjective relevance. In *Proceedings of Australian Conference
 12487 on Artificial Intelligence*.
- 12488 Wager, S., S. Wang, and P. S. Liang (2013). Dropout training as adaptive regularization. In
 12489 *Neural Information Processing Systems (NIPS)*, pp. 351–359.
- 12490 Wainwright, M. J. and M. I. Jordan (2008). Graphical models, exponential families, and
 12491 variational inference. *Foundations and Trends® in Machine Learning* 1(1-2), 1–305.
- 12492 Walker, M. A. (2000). An application of reinforcement learning to dialogue strategy selec-
 12493 tion in a spoken dialogue system for email. *Journal of Artificial Intelligence Research* 12,
 12494 387–416.
- 12495 Walker, M. A., J. E. Cahn, and S. J. Whittaker (1997). Improvising linguistic style: Social
 12496 and affective bases for agent personality. In *Proceedings of the first international conference
 12497 on Autonomous agents*, pp. 96–105. ACM.
- 12498 Wang, C., N. Xue, and S. Pradhan (2015). A Transition-based Algorithm for AMR Parsing.
 12499 In *Proceedings of the North American Chapter of the Association for Computational Linguistics
 12500 (NAACL)*, pp. 366–375.
- 12501 Wang, H., T. Onishi, K. Gimpel, and D. McAllester (2017). Emergent predication structure
 12502 in hidden state vectors of neural readers. In *Proceedings of the 2nd Workshop on Represen-
 12503 tation Learning for NLP*, pp. 26–36.
- 12504 Weaver, W. (1955). Translation. *Machine translation of languages* 14, 15–23.
- 12505 Webber, B. (2004, sep). D-LTAG: extending lexicalized TAG to discourse. *Cognitive Sci-
 12506 ence* 28(5), 751–779.
- 12507 Webber, B., M. Egg, and V. Kordoni (2012). Discourse structure and language technology.
 12508 *Journal of Natural Language Engineering* 1.
- 12509 Webber, B. and A. Joshi (2012). Discourse structure and computation: past, present and
 12510 future. In *Proceedings of the ACL-2012 Special Workshop on Rediscovering 50 Years of Dis-
 12511 coveries*, pp. 42–54. Association for Computational Linguistics.
- 12512 Wei, G. C. and M. A. Tanner (1990). A monte carlo implementation of the em algorithm
 12513 and the poor man’s data augmentation algorithms. *Journal of the American Statistical
 12514 Association* 85(411), 699–704.

- 12515 Weinberger, K., A. Dasgupta, J. Langford, A. Smola, and J. Attenberg (2009). Feature
12516 hashing for large scale multitask learning. In *Proceedings of the International Conference*
12517 *on Machine Learning (ICML)*, pp. 1113–1120.
- 12518 Weizenbaum, J. (1966). Eliza—a computer program for the study of natural language
12519 communication between man and machine. *Communications of the ACM* 9(1), 36–45.
- 12520 Wellner, B. and J. Pustejovsky (2007). Automatically identifying the arguments of dis-
12521 course connectives. In *Proceedings of Empirical Methods for Natural Language Processing*
12522 (*EMNLP*), pp. 92–101.
- 12523 Wen, T.-H., M. Gasic, N. Mrkšić, P.-H. Su, D. Vandyke, and S. Young (2015). Semantically
12524 conditioned lstm-based natural language generation for spoken dialogue systems. In
12525 *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pp. 1711–1721.
- 12526 Weston, J., S. Bengio, and N. Usunier (2011). Wsabie: Scaling up to large vocabulary image
12527 annotation. In *IJCAI*, Volume 11, pp. 2764–2770.
- 12528 Wiebe, J., T. Wilson, and C. Cardie (2005). Annotating expressions of opinions and emo-
12529 tions in language. *Language resources and evaluation* 39(2), 165–210.
- 12530 Wieting, J., M. Bansal, K. Gimpel, and K. Livescu (2015). Towards universal paraphrastic
12531 sentence embeddings. *arXiv preprint arXiv:1511.08198*.
- 12532 Wieting, J., M. Bansal, K. Gimpel, and K. Livescu (2016). CHARAGRAM: Embedding
12533 words and sentences via character n-grams. In *Proceedings of Empirical Methods for Nat-*
12534 *ural Language Processing (EMNLP)*, pp. 1504–1515.
- 12535 Williams, J. D. and S. Young (2007). Partially observable markov decision processes for
12536 spoken dialog systems. *Computer Speech & Language* 21(2), 393–422.
- 12537 Williams, P., R. Sennrich, M. Post, and P. Koehn (2016). Syntax-based statistical machine
12538 translation. *Synthesis Lectures on Human Language Technologies* 9(4), 1–208.
- 12539 Wilson, T., J. Wiebe, and P. Hoffmann (2005). Recognizing contextual polarity in phrase-
12540 level sentiment analysis. In *Proceedings of Empirical Methods for Natural Language Pro-*
12541 *cessing (EMNLP)*, pp. 347–354.
- 12542 Winograd, T. (1972). Understanding natural language. *Cognitive psychology* 3(1), 1–191.
- 12543 Wiseman, S., A. M. Rush, and S. M. Shieber (2016). Learning global features for corefer-
12544 ence resolution. In *Proceedings of the North American Chapter of the Association for Compu-*
12545 *tational Linguistics (NAACL)*, pp. 994–1004.

- 12546 Wiseman, S., S. Shieber, and A. Rush (2017). Challenges in data-to-document generation.
 12547 In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*, pp. 2253–
 12548 2263.
- 12549 Wiseman, S. J., A. M. Rush, S. M. Shieber, and J. Weston (2015). Learning anaphoricity and
 12550 antecedent ranking features for coreference resolution. In *Proceedings of the Association*
 12551 *for Computational Linguistics (ACL)*.
- 12552 Wolf, F. and E. Gibson (2005). Representing discourse coherence: A corpus-based study.
 12553 *Computational Linguistics* 31(2), 249–287.
- 12554 Wolfe, T., M. Dredze, and B. Van Durme (2017). Pocket knowledge base population. In
 12555 *Proceedings of the Association for Computational Linguistics (ACL)*, pp. 305–310.
- 12556 Wong, Y. W. and R. Mooney (2007). Generation by inverting a semantic parser that uses
 12557 statistical machine translation. In *Proceedings of the North American Chapter of the Associa-*
 12558 *tion for Computational Linguistics (NAACL)*, pp. 172–179.
- 12559 Wong, Y. W. and R. J. Mooney (2006). Learning for semantic parsing with statistical ma-
 12560 chine translation. In *Proceedings of the North American Chapter of the Association for Com-*
 12561 *putational Linguistics (NAACL)*, pp. 439–446.
- 12562 Wu, B. Y. and K.-M. Chao (2004). *Spanning trees and optimization problems*. CRC Press.
- 12563 Wu, D. (1997). Stochastic inversion transduction grammars and bilingual parsing of par-
 12564 allel corpora. *Computational linguistics* 23(3), 377–403.
- 12565 Wu, F. and D. S. Weld (2010). Open information extraction using wikipedia. In *Proceedings*
 12566 *of the Association for Computational Linguistics (ACL)*, pp. 118–127.
- 12567 Wu, X., R. Ward, and L. Bottou (2018). Wngrad: Learn the learning rate in gradient de-
 12568 scent. *arXiv preprint arXiv:1803.02865*.
- 12569 Wu, Y., M. Schuster, Z. Chen, Q. V. Le, M. Norouzi, W. Macherey, M. Krikun, Y. Cao,
 12570 Q. Gao, K. Macherey, J. Klingner, A. Shah, M. Johnson, X. Liu, Łukasz Kaiser, S. Gouws,
 12571 Y. Kato, T. Kudo, H. Kazawa, K. Stevens, G. Kurian, N. Patil, W. Wang, C. Young,
 12572 J. Smith, J. Riesa, A. Rudnick, O. Vinyals, G. Corrado, M. Hughes, and J. Dean (2016).
 12573 Google’s neural machine translation system: Bridging the gap between human and ma-
 12574 chine translation. *CoRR abs/1609.08144*.
- 12575 Xia, F. (2000). The part-of-speech tagging guidelines for the penn chinese treebank (3.0).
 12576 Technical report, University of Pennsylvania Institute for Research in Cognitive Science.
- 12577 Xu, K., J. Ba, R. Kiros, K. Cho, A. Courville, R. Salakhudinov, R. Zemel, and Y. Bengio
 12578 (2015). Show, attend and tell: Neural image caption generation with visual attention.
 12579 In *Proceedings of the International Conference on Machine Learning (ICML)*, pp. 2048–2057.

- 12580 Xu, W., X. Liu, and Y. Gong (2003). Document clustering based on non-negative matrix
12581 factorization. In *SIGIR*, pp. 267–273. ACM.
- 12582 Xu, Y., L. Mou, G. Li, Y. Chen, H. Peng, and Z. Jin (2015). Classifying relations via long
12583 short term memory networks along shortest dependency paths. In *Proceedings of Empirical
12584 Methods for Natural Language Processing (EMNLP)*, pp. 1785–1794.
- 12585 Xuan Bach, N., N. L. Minh, and A. Shimazu (2012). A reranking model for discourse seg-
12586 mentation using subtree features. In *Proceedings of the Special Interest Group on Discourse
12587 and Dialogue (SIGDIAL)*.
- 12588 Xue, N. et al. (2003). Chinese word segmentation as character tagging. *Computational
12589 Linguistics and Chinese Language Processing* 8(1), 29–48.
- 12590 Xue, N., H. T. Ng, S. Pradhan, R. Prasad, C. Bryant, and A. T. Rutherford (2015). The
12591 CoNLL-2015 shared task on shallow discourse parsing. In *Proceedings of the Conference
12592 on Natural Language Learning (CoNLL)*.
- 12593 Xue, N., H. T. Ng, S. Pradhan, A. Rutherford, B. L. Webber, C. Wang, and H. Wang (2016).
12594 Conll 2016 shared task on multilingual shallow discourse parsing. In *CoNLL Shared
12595 Task*, pp. 1–19.
- 12596 Yamada, H. and Y. Matsumoto (2003). Statistical dependency analysis with support vector
12597 machines. In *Proceedings of IWPT*, Volume 3, pp. 195–206.
- 12598 Yamada, K. and K. Knight (2001). A syntax-based statistical translation model. In *Proceed-
12599 ings of the 39th Annual Meeting on Association for Computational Linguistics*, pp. 523–530.
12600 Association for Computational Linguistics.
- 12601 Yang, B. and C. Cardie (2014). Context-aware learning for sentence-level sentiment anal-
12602 ysis with posterior regularization. In *Proceedings of the Association for Computational Lin-
12603 guistics (ACL)*.
- 12604 Yang, Y., M.-W. Chang, and J. Eisenstein (2016). Toward socially-infused information ex-
12605 traction: Embedding authors, mentions, and entities. In *Proceedings of Empirical Methods
12606 for Natural Language Processing (EMNLP)*.
- 12607 Yang, Y. and J. Eisenstein (2013). A log-linear model for unsupervised text normalization.
12608 In *Proceedings of Empirical Methods for Natural Language Processing (EMNLP)*.
- 12609 Yang, Y. and J. Eisenstein (2015). Unsupervised multi-domain adaptation with feature em-
12610 beddings. In *Proceedings of the North American Chapter of the Association for Computational
12611 Linguistics (NAACL)*.

- 12612 Yang, Y., W.-t. Yih, and C. Meek (2015). WikiQA: A challenge dataset for open-domain
 12613 question answering. In *Proceedings of Empirical Methods for Natural Language Processing*
 12614 (*EMNLP*), pp. 2013–2018.
- 12615 Yannakoudakis, H., T. Briscoe, and B. Medlock (2011). A new dataset and method for
 12616 automatically grading esol texts. In *Proceedings of the 49th Annual Meeting of the Associa-*
 12617 *tion for Computational Linguistics: Human Language Technologies-Volume 1*, pp. 180–189.
 12618 Association for Computational Linguistics.
- 12619 Yarowsky, D. (1995). Unsupervised word sense disambiguation rivaling supervised meth-
 12620 ods. In *Proceedings of the Association for Computational Linguistics (ACL)*, pp. 189–196.
 12621 Association for Computational Linguistics.
- 12622 Yee, L. C. and T. Y. Jones (2012, March). Apple ceo in china mission to clear up problems.
 12623 *Reuters*. retrieved on March 26, 2017.
- 12624 Yi, Y., C.-Y. Lai, S. Petrov, and K. Keutzer (2011, October). Efficient parallel cky parsing on
 12625 gpus. In *Proceedings of the 12th International Conference on Parsing Technologies*, Dublin,
 12626 Ireland, pp. 175–185. Association for Computational Linguistics.
- 12627 Yu, C.-N. J. and T. Joachims (2009). Learning structural svms with latent variables. In
 12628 *Proceedings of the International Conference on Machine Learning (ICML)*, pp. 1169–1176.
- 12629 Yu, F. and V. Koltun (2016). Multi-scale context aggregation by dilated convolutions. In
 12630 *Proceedings of the International Conference on Learning Representations (ICLR)*.
- 12631 Zaidan, O. F. and C. Callison-Burch (2011). Crowdsourcing translation: Professional qual-
 12632 ity from non-professionals. In *Proceedings of the Association for Computational Linguistics*
 12633 (*ACL*), pp. 1220–1229.
- 12634 Zaremba, W., I. Sutskever, and O. Vinyals. Recurrent neural network regularization. *arXiv*
 12635 preprint *arXiv:1409.2329*.
- 12636 Zeiler, M. D. (2012). Adadelta: an adaptive learning rate method. *arXiv preprint*
 12637 *arXiv:1212.5701*.
- 12638 Zelenko, D., C. Aone, and A. Richardella (2003). Kernel methods for relation extraction.
 12639 *The Journal of Machine Learning Research* 3, 1083–1106.
- 12640 Zelle, J. M. and R. J. Mooney (1996). Learning to parse database queries using induc-
 12641 tive logic programming. In *Proceedings of the National Conference on Artificial Intelligence*
 12642 (*AAAI*), pp. 1050–1055.
- 12643 Zeng, D., K. Liu, S. Lai, G. Zhou, and J. Zhao (2014). Relation classification via convolu-
 12644 tional deep neural network. In *Proceedings of the International Conference on Computational*
 12645 *Linguistics (COLING)*, pp. 2335–2344.

- 12646 Zettlemoyer, L. S. and M. Collins (2005). Learning to map sentences to logical form: Structured classification with probabilistic categorial grammars. In *Proceedings of UAI*.
12647
- 12648 Zhang, X., J. Zhao, and Y. LeCun (2015). Character-level convolutional networks for text
12649 classification. In *Neural Information Processing Systems (NIPS)*, pp. 649–657.
- 12650 Zhang, Y. and S. Clark (2008). A tale of two parsers: investigating and combining graph-
12651 based and transition-based dependency parsing using beam-search. In *Proceedings of*
12652 *Empirical Methods for Natural Language Processing (EMNLP)*, pp. 562–571.
- 12653 Zhang, Y., T. Lei, R. Barzilay, T. Jaakkola, and A. Globerson (2014). Steps to excellence:
12654 Simple inference with refined scoring of dependency trees. In *Proceedings of the Associa-*
12655 *tion for Computational Linguistics (ACL)*, pp. 197–207.
- 12656 Zhang, Y. and J. Nivre (2011). Transition-based dependency parsing with rich non-local
12657 features. In *Proceedings of the Association for Computational Linguistics (ACL)*, pp. 188–193.
- 12658 Zhou, J. and W. Xu (2015). End-to-end learning of semantic role labeling using recurrent
12659 neural networks. In *Proceedings of the Association for Computational Linguistics (ACL)*, pp.
12660 1127–1137.
- 12661 Zhu, J., Z. Nie, X. Liu, B. Zhang, and J.-R. Wen (2009). Statsnowball: a statistical approach
12662 to extracting entity relationships. In *Proceedings of the Conference on World-Wide Web*
12663 (*WWW*), pp. 101–110.
- 12664 Zhu, X., Z. Ghahramani, and J. D. Lafferty (2003). Semi-supervised learning using gaus-
12665 sian fields and harmonic functions. In *Proceedings of the International Conference on Ma-*
12666 *chine Learning (ICML)*, pp. 912–919.
- 12667 Zhu, X. and A. B. Goldberg (2009). Introduction to semi-supervised learning. *Synthesis*
12668 *lectures on artificial intelligence and machine learning* 3(1), 1–130.
- 12669 Zipf, G. K. (1949). Human behavior and the principle of least effort.
- 12670 Zirn, C., M. Niepert, H. Stuckenschmidt, and M. Strube (2011). Fine-grained sentiment
12671 analysis with structural features. In *IJCNLP*, Chiang Mai, Thailand, pp. 336–344.
- 12672 Zou, W. Y., R. Socher, D. Cer, and C. D. Manning (2013). Bilingual word embeddings
12673 for phrase-based machine translation. In *Proceedings of Empirical Methods for Natural*
12674 *Language Processing (EMNLP)*, pp. 1393–1398.

¹²⁶⁷⁵ Index

- 12676 *K*-means, 108
12677 α -conversion, 298
12678 β -conversion, 295
12679 β -reduction, 295
12680 *n*-gram language models, 206
12681 *n*-gram, 40
12682 *F*-MEASURE, 94
12683 BLEU, 437
12684 WordNet, 86
- 12685 ablation test, 96
12686 absolute discounting, 142
12687 Abstract Meaning Representation (AMR), 311, 323
12688 abstractive summarization, 399, 466
12690 accepting path, 201
12691 accuracy, 38, 93
12692 action (reinforcement learning), 373
12693 active learning, 130
12694 AdaDelta (online optimization), 75
12695 AdaGrad (online optimization), 55, 75
12696 Adam (online optimization), 75
12697 adequacy (translation), 437
12698 adjectives, 187
12699 adjuncts (semantics), 310, 314
12700 adpositions, 188
12701 adverbs, 187
12702 adversarial networks, 125
12703 affix (morphology), 203
12704 agenda-based dialogue systems, 470
12705 agenda-based parsing, 282
- 12706 agent (thematic role), 312
12707 alignment, 325
12708 alignment (machine translation), 436, 441
12710 alignment (text generation), 461
12711 Amazon Mechanical Turk, 102
12712 ambiguity, 218, 226
12713 ambiguity, attachment, 237
12714 ambiguity, complement structure, 237
12715 ambiguity, coordination scope, 237
12716 ambiguity, modifier scope, 237
12717 ambiguity, particle versus preposition, 237
12719 anaphoric, 368
12720 anchored productions, 240
12721 animacy (semantics), 311
12722 annealing, 456
12723 antecedent (coreference), 359, 367
12724 antonymy, 86
12725 apophony, 202
12726 arc-eager dependency parsing, 273, 275
12727 arc-factored dependency parsing, 267
12728 arc-standard dependency parsing, 273
12729 area under the curve (AUC), 95
12730 argumentation, 398
12731 argumentation mining, 399
12732 arguments, 407
12733 article (syntax), 192
12734 aspect, 187
12735 aspect-based opinion mining, 83
12736 attachment ambiguity, 261

- 12737 attention mechanism, 379, 430, 449, 462
 12738 autoencoder, 353
 12739 autoencoder, denoising, 353, 412
 12740 automated theorem provers, 290
 12741 automatic differentiation, 70
 12742 auxiliary verbs, 188
 12743 average mutual information, 341
 12744 averaged perceptron, 42
 12745 backchannel, 197
 12746 backoff, 142
 12747 backpropagation, 69, 148
 12748 backpropagation through time, 148
 12749 backward recurrence, 175, 176
 12750 backward-looking center, 388
 12751 bag of words, 29
 12752 balanced *F*-MEASURE, 95
 12753 balanced test set, 93
 12754 batch learning, 41
 12755 batch normalization, 74
 12756 batch optimization, 53
 12757 Baum-Welch algorithm, 180
 12758 Bayes' rule, 478
 12759 Bayesian nonparametrics, 115, 256
 12760 beam sampling, 182
 12761 beam search, 275, 372, 453, 454
 12762 Bell number, 370
 12763 best-path algorithm, 206
 12764 bias (learning theory), 38, 139
 12765 bias-variance tradeoff, 38, 141
 12766 biconvexity, 114
 12767 bidirectional LSTM, 449
 12768 bidirectional recurrent neural network,
 178
 12769 bigrams, 40, 82
 12770 bilexical, 255
 12771 bilexical features, 270
 12773 bilinear product, 338
 12774 binarization (context-free grammar),
 219, 236
 12776 binomial distribution, 96
 12777 binomial random variable, 482
 12778 binomial test, 96
 12779 BIO notation, 194, 321
 12780 biomedical natural language processing,
 193
 12782 bipartite graph, 327
 12783 Bonferroni correction, 99
 12784 boolean semiring, 207
 12785 boosting, 62
 12786 bootstrap samples, 98
 12787 brevity penalty (machine translation),
 438
 12789 Brown clusters, 335, 340
 12790 byte-pair encoding, 351, 452
 12791 c-command, 361
 12792 case marking, 192, 227
 12793 Catalan number, 233
 12794 cataphora, 360
 12795 center embedding, 215
 12796 centering theory, 363, 388
 12797 chain FSA, 213
 12798 chain rule of probability, 477
 12799 chance agreement, 102
 12800 character-level language models, 153
 12801 chart parsing, 234
 12802 chatbots, 472
 12803 Chomsky Normal Form (CNF), 219
 12804 Chu-Liu-Edmonds algorithm, 268
 12805 CKY algorithm, 234
 12806 class imbalance, 93
 12807 classification weights, 29
 12808 closed-vocabulary, 153
 12809 closure (regular languages), 200
 12810 cloze question answering, 429
 12811 cluster ranking, 371
 12812 clustering, 108
 12813 co-training, 120
 12814 coarse-to-fine attention, 465
 12815 code switching, 190, 196
 12816 Cohen's Kappa, 102

- 12817 coherence, 402
 12818 cohesion, 385
 12819 collapsed Gibbs sampling, 128
 12820 collective entity linking, 412, 413
 12821 collocation extraction, 352
 12822 collocation features, 87
 12823 combinatorial optimization, 19
 12824 combinatory categorial grammar, 228
 12825 complement clause, 221
 12826 complement event (probability), 476
 12827 composition (CCG), 229
 12828 compositional vector grammars, 397
 12829 compositionality, 18, 21, 350
 12830 computation graph, 62, 69
 12831 computational linguistics (versus natural language processing), 13
 12832 computational social science, 17
 12833 concept (AMR), 323
 12835 conditional independence, 164, 480
 12836 conditional log-likelihood, 66
 12837 conditional probability, 50, 477
 12838 conditional probability distribution, 481
 12839 conditional random field, 172
 12840 confidence interval, 98
 12841 configuration (transition-based parsing), 273
 12842 connected (graph theory), 325
 12844 consistency (logic), 293
 12845 constants (logic), 288
 12846 constituents, 220
 12847 constrained optimization, 47, 319
 12848 constraint-driven learning, 130
 12849 constraints, 319
 12850 content selection (text generation), 459
 12851 content words, 188
 12852 context vector (attentional neural translation), 450
 12854 context-free grammars (CFGs), 216
 12855 context-free languages, 215, 216
 12856 context-sensitive languages, 226
 12857 continuous bag-of-words (CBOW), 343
 12858 contradiction, 354
 12859 conversational turns, 197
 12860 convex optimization, 53
 12861 convexity, 44, 72, 303, 482, 485
 12862 convolutional neural networks, 67, 76, 82, 153, 180, 195, 419
 12863 cooperative principle, 359
 12865 coordinate ascent, 114
 12866 coordinating conjunctions, 188
 12867 coordinating discourse relations, 395
 12868 copula, 187, 225, 264
 12869 coreference resolution, 355, 359
 12870 coreferent, 359
 12871 cosine similarity, 347, 386
 12872 cost-augmented decoding, 49, 171
 12873 coverage (summarization), 399
 12874 coverage loss, 467
 12875 critical point, 72, 485
 12876 cross-document coreference resolution, 410
 12877 cross-entropy, 66, 420
 12879 cross-serial dependencies, 227
 12880 cross-validation, 39
 12881 crowdsourcing, 102
 12882 cumulative probability distribution, 97
 12883 dead neurons, 65
 12884 decidability (logic), 293
 12885 decision trees, 62
 12886 deep learning, 61
 12887 deep LSTM, 448
 12888 definiteness, 193
 12889 delta function, 37
 12890 denotation (semantics), 288
 12891 dependency grammar, 261
 12892 dependency graph, 262
 12893 dependency parse, 261
 12894 dependency path, 87, 318, 417
 12895 dependent, 262
 12896 derivation (context-free languages), 217
 12897 derivation (parsing), 272, 278

- 12898 derivation (semantic parsing), 296, 300
 12899 derivational ambiguity, 230
 12900 derivational morphology, 202
 12901 determiner, 189
 12902 determiner phrase, 223
 12903 deterministic FSA, 202
 12904 development set, 39, 93
 12905 dialogue acts, 102, 197, 474
 12906 dialogue management, 470
 12907 dialogue systems, 137, 469
 12908 digital humanities, 17, 81
 12909 dilated convolution, 77, 195
 12910 Dirichlet distribution, 127
 12911 discounting (language models), 142
 12912 discourse, 385
 12913 discourse connectives, 391
 12914 discourse depth, 400
 12915 discourse depth tree, 401
 12916 discourse parsing, 390
 12917 discourse relations, 355, 390
 12918 discourse segment, 385
 12919 discourse sense classification, 392
 12920 discourse unit, 395
 12921 discrete random variable, 480
 12922 discriminative learning, 40
 12923 disjoint events (probability), 476
 12924 distant supervision, 130, 422, 423
 12925 distributed semantics, 334
 12926 distributional hypothesis, 333, 334
 12927 distributional semantics, 22, 334
 12928 distributional statistics, 87, 255, 334
 12929 document frequency, 411
 12930 domain adaptation, 107, 123
 12931 dropout, 71, 149
 12932 dual decomposition, 321
 12933 dynamic computation graphs, 70
 12934 dynamic oracle, 280
 12935 dynamic programming, 159
 12936 dynamic semantics, 305, 390
 12937 E-step (expectation-maximization), 111
 12938 early stopping, 43, 75
 12939 early update, 280
 12940 easy-first parsing, 283
 12941 edit distance, 209, 439
 12942 effective count (language models), 141
 12943 elementary discourse units, 395
 12944 elementwise nonlinearity, 63
 12945 Elman unit, 147
 12946 ELMo (embeddings from language models), 349
 12947 embedding, 177, 343
 12948 emission features, 158
 12950 emotion, 84
 12951 empirical Bayes, 128
 12952 empty string, 200
 12953 encoder-decoder, 446
 12954 encoder-decoder model, 353, 462
 12955 ensemble, 322
 12956 ensemble learning, 448
 12957 ensemble methods, 62
 12958 entailment, 293, 354
 12959 entities, 407
 12960 entity embeddings, 411
 12961 entity grid, 389
 12962 entity linking, 359, 407, 409, 420
 12963 entropy, 57, 111
 12964 estimation, 482
 12965 EuroParl corpus, 439
 12966 event, 425
 12967 event (probability), 475
 12968 event coreference, 425
 12969 event detection, 425
 12970 event semantics, 309
 12971 events, 407
 12972 evidentiality, 192, 427
 12973 exchange clustering, 342
 12974 expectation, 481
 12975 expectation maximization, 109, 143
 12976 expectation semiring, 215
 12977 expectation-maximization, in machine translation, 443
 12978

- 12979 explicit semantic analysis, 336
12980 exploding gradients, 149
12981 extra-propositional semantics, 427
12982 extractive question-answering, 430
12983 extractive summarization, 399
12984 extrinsic evaluation, 151, 347

12985 factoid questions, 326
12986 factoids, 428
12987 factor graph, 173
12988 factuality, 427
12989 false discovery rate, 99
12990 False negative, 93
12991 False positive, 93
12992 false positive, 479
12993 false positive rate, 95, 478
12994 feature co-adaptation, 71
12995 feature function, 30, 39
12996 feature hashing, 92
12997 feature noising, 72
12998 feature selection, 56
12999 features, 18
13000 feedforward neural network, 64
13001 fine-tuned word embeddings, 349
13002 finite state acceptor (FSA), 201
13003 finite state automata, 201
13004 finite state composition, 212
13005 finite state transducers, 204, 209
13006 first-order logic, 291
13007 fluency (translation), 437
13008 fluent, 137
13009 focus, 324
13010 formal language theory, 199
13011 forward recurrence, 174
13012 forward variable, 176
13013 forward variables, 174
13014 forward-backward algorithm, 175, 214,
 247
13015 forward-looking centers, 388
13016 frame, 469
13017 frame elements, 314

13019 FrameNet, 314
13020 frames, 314
13021 Frobenius norm, 71
13022 function (first-order logic), 292
13023 function words, 188
13024 functional margin, 47
13025 functional segmentation, 385, 387

13026 garden path sentence, 156
13027 gate (neural networks), 66, 449
13028 gazetteer, 417
13029 gazetteers, 365
13030 gazetteers, 194
13031 generalization, 43
13032 generalized linear models, 57
13033 generative model, 33, 243
13034 generative models, 371
13035 generative process, 143
13036 generic referents, 364
13037 geometric margin, 47
13038 Gibbs sampling, 127, 413
13039 gloss, 137, 189, 437
13040 government and binding theory, 361
13041 gradient, 45
13042 gradient clipping, 74
13043 gradient descent, 53
13044 Gram matrix, 418
13045 grammar induction, 249
13046 grammaticality, 402
13047 graph-based dependency parsing, 267
13048 graphical model, 164
13049 graphics processing units (GPUs), 180,
 195
13050 grid search, 38

13051 Hamming cost, 171
13053 Hansards corpus, 439
13054 hanzi, 89
13055 hard expectation-maximization, 114
13056 head, 262, 417
13057 head percolation rules, 252

- 13058 head rules, 261
 13059 head word, 220, 261, 365
 13060 head words, 252
 13061 hedging, 427
 13062 held-out data, 151
 13063 Hessian matrix, 54
 13064 hidden Markov models, 164
 13065 hidden variable perceptron, 215
 13066 hierarchical clustering, 340
 13067 hierarchical recurrent network, 473
 13068 hierarchical softmax, 147, 345
 13069 hierarchical topic segmentation, 387
 13070 highway network, 66
 13071 hinge loss, 44
 13072 holonymy, 87
 13073 homonym, 85
 13074 human computation, 103
 13075 hypergraph, 398
 13076 hypernymy, 87
 13077 hyperparameter, 38
 13078 hyponymy, 87
 13079 illocutionary force, 197
 13080 implicit discourse relations, 392
 13081 importance sampling, 456
 13082 importance score, 456
 13083 incremental expectation maximization, 114
 13084 incremental perceptron, 280, 372
 13085 independent and identically distributed (IID), 32
 13086 indicator function, 37
 13089 indicator random variable, 480
 13090 inference, 157
 13091 inference (logic), 287
 13092 inference rules, 290
 13093 inflection point, 486
 13094 inflectional affixes, 90
 13095 inflectional morphology, 187, 202, 210
 13096 information extraction, 407
 13097 information retrieval, 17, 421
 13098 initiative (dialogue systems), 470
 13099 input word embeddings, 147
 13100 inside recurrence, 243, 244, 248
 13101 inside-outside algorithm, 247, 256
 13102 instance (AMR), 323
 13103 instance labels, 32
 13104 integer linear program, 432, 468
 13105 integer linear programming, 319, 370, 401, 413
 13107 inter-annotator agreement, 102
 13108 interjections, 187
 13109 interlingua, 436
 13110 interpolated n -gram language model, 207
 13112 interpolation, 143
 13113 interval algebra, 426
 13114 intrinsic evaluation, 151, 347
 13115 inverse document frequency, 411
 13116 inverse relation (AMR), 324
 13117 inversion (finite state), 211
 13118 irrealis, 82
 13119 Jeffreys-Perks law, 141
 13120 Jensen's inequality, 111
 13121 joint probabilities, 481
 13122 joint probability, 32, 50
 13123 Kalman smoother, 182
 13124 Katz backoff, 142
 13125 kernel function, 418
 13126 kernel methods, 62
 13127 kernel support vector machine, 62, 418
 13128 Kleene star, 200
 13129 knapsack problem, 401
 13130 knowledge base, 407
 13131 knowledge base population, 420
 13132 L-BFGS, 54
 13133 label bias problem, 279
 13134 label propagation, 122, 132
 13135 labeled dependencies, 263
 13136 labeled precision, 238

- 13137 labeled recall, 238
 13138 Lagrange multiplier, 487
 13139 Lagrangian, 487
 13140 Lagrangian dual, 487
 13141 lambda calculus, 295
 13142 lambda expressions, 295
 13143 language model, 138
 13144 language models, 16
 13145 Laplace smoothing, 38, 141
 13146 large margin classification, 46
 13147 latent conditional random fields, 303
 13148 latent Dirichlet allocation, 386
 13149 latent semantic analysis, 336, 338
 13150 latent variable, 110, 214, 302, 423, 436
 13151 latent variable perceptron, 303, 368
 13152 layer normalization, 75, 451
 13153 leaky ReLU, 65
 13154 learning to search, 261, 281, 374
 13155 least squares, 84
 13156 leave-one-out, 39
 13157 lemma, 85
 13158 lemma (lexical semantics), 210
 13159 lemmatization, 90
 13160 Levenshtein edit distance, 209
 13161 lexical entry, 296
 13162 lexical features, 61
 13163 lexical semantics, 85
 13164 lexical unit (frame semantics), 314
 13165 lexicalization, 252
 13166 lexicalization (text generation), 459
 13167 lexicalized tree-adjoining grammar for discourse (D-LTAG), 391
 13168 lexicon, 296
 13169 lexicon (CCG), 229
 13170 lexicon-based classification, 84
 13172 lexicon-based sentiment analysis, 82
 13173 Lidstone smoothing, 141
 13174 light verb, 325
 13175 likelihood, 478
 13176 linear regression, 84
 13177 linear separability, 41
 13178 linearization, 469
 13179 literal character, 200
 13180 local minimum, 486
 13181 local optimum, 114
 13182 locally-normalized objective, 279
 13183 log-bilinear language model, 351
 13184 logistic function, 57
 13185 logistic loss, 51
 13186 logistic regression, 50, 57
 13187 Long short-term memories, 147
 13188 long short-term memory, 149
 13189 long short-term memory (LSTM), 66, 191, 447
 13190 lookup layer, 67, 147
 13192 loss function, 43
 13193 LSTM, 149
 13194 LSTM-CRF, 179, 322
 13195 machine learning, 14
 13196 machine reading, 429
 13197 machine translation, 137
 13198 Macro F-MEASURE, 94
 13199 macro-reading, 408
 13200 margin, 41, 46
 13201 marginal probability distribution, 481
 13202 marginal relevance, 467
 13203 marginalize, 477
 13204 markable, 366
 13205 Markov assumption, 164
 13206 Markov blanket, 164
 13207 Markov Chain Monte Carlo (MCMC), 115, 127, 182
 13209 Markov decision process, 470
 13210 Markov random fields, 172
 13211 matrix-tree theorem, 272
 13212 max-margin Markov network, 171
 13213 max-product algorithm, 167
 13214 maximum a posteriori, 38, 483
 13215 maximum conditional likelihood, 50
 13216 maximum entropy, 57
 13217 maximum likelihood, 482

- 13218 maximum likelihood estimate, 36
 13219 maximum likelihood estimation, 32
 13220 maximum spanning tree, 268
 13221 McNemar’s test, 96
 13222 meaning representation, 287
 13223 membership problem, 199
 13224 memory cell (LSTM), 149
 13225 mention (coreference resolution), 359
 13226 mention (entity), 407
 13227 mention (information extraction), 409
 13228 mention ranking, 368
 13229 mention-pair model, 367
 13230 meronymy, 87
 13231 meteor, 439
 13232 method of moments, 128
 13233 micro *F*-MEASURE, 94
 13234 micro-reading, 408
 13235 mildly context-sensitive languages, 227
 13236 minibatch, 54
 13237 minimization (FSA), 204
 13238 minimum error-rate training (MERT),
 455
 13239 minimum risk training, 455
 13240 mixed-initiative, 470
 13241 modality, 427
 13243 model, 19
 13244 model builder, 293
 13245 model checker, 293
 13246 model-theoretic semantics, 288
 13247 modeling (machine learning), 53
 13248 modifier (dependency grammar), 262
 13249 modus ponens, 290
 13250 moment-matching, 57
 13251 monomorphemic, 204
 13252 morphemes, 17, 153, 203
 13253 morphological analysis, 210
 13254 morphological generation, 210
 13255 morphological segmentation, 169
 13256 morphology, 91, 168, 202, 350, 452
 13257 morphosyntactic, 186
 13258 morphosyntactic attributes, 190
 13259 morphotactic, 203
 13260 multi-document summarization, 468
 13261 multi-view learning, 120
 13262 multilayer perceptron, 64
 13263 multinomial distribution, 34
 13264 multinomial naïve Bayes, 34
 13265 multiple instance learning, 130, 423
 13266 multiple-choice question answering, 429
 13267 multitask learning, 130
 13268 Naïve Bayes, 33
 13269 name dictionary, 410
 13270 named entities, 193
 13271 named entity linking, 409
 13272 named entity recognition, 179, 407, 409
 13273 named entity types, 409
 13274 narrow convolution, 77
 13275 nearest-neighbor, 62, 418
 13276 negation, 82, 427
 13277 negative sampling, 345, 346, 415
 13278 Neo-Davidsonian event semantics, 310
 13279 neural attention, 448
 13280 neural machine translation, 436
 13281 neural networks, 61, 146
 13282 NIL entity, 409
 13283 noise-contrastive estimation, 147
 13284 noisy channel model, 138, 440
 13285 nominal modifier, 223
 13286 nominals, 359
 13287 nominals (coreference), 366
 13288 non-core roles (AMR), 324
 13289 non-terminals (context-free grammars),
 217
 13290 normalization, 90
 13292 noun phrase, 14, 220
 13293 nouns, 186
 13294 NP-hard, 56, 413
 13295 nuclearity (RST), 395
 13296 nucleus (RST), 395
 13297 null hypothesis, 96
 13298 numeral (part of speech), 189

- | | | | |
|-------|---|-------|--|
| 13299 | numerical optimization, 20 | 13338 | passive-aggressive, 487 |
| 13300 | offset feature, 31 | 13339 | path (finite state automata), 201 |
| 13301 | one-dimensional convolution, 77 | 13340 | Penn Discourse Treebank (PDTB), 391 |
| 13302 | one-hot, 380 | 13341 | Penn Treebank, 152, 169, 190, 220, 246 |
| 13303 | one-hot vector, 66 | 13342 | perceptron, 41 |
| 13304 | one-tailed p-value, 97 | 13343 | perplexity, 152 |
| 13305 | one-versus-all multiclass classification, 418 | 13344 | phonology, 204 |
| 13306 | | 13345 | phrase (syntax), 220 |
| 13307 | one-versus-one multiclass classification, 418 | 13346 | phrase-structure grammar, 220 |
| 13308 | | 13347 | pivot features, 124 |
| 13309 | online expectation maximization, 114 | 13348 | pointwise mutual information, 338 |
| 13310 | online learning, 41, 54 | 13349 | policy, 373, 471 |
| 13311 | ontology, 20 | 13350 | policy (search), 280 |
| 13312 | open information extraction, 424 | 13351 | policy gradient, 374 |
| 13313 | open word classes, 186 | 13352 | polysemous, 86 |
| 13314 | opinion polarity, 81 | 13353 | pooling, 381 |
| 13315 | oracle, 278, 326 | 13354 | pooling (convolution), 77, 381, 462 |
| 13316 | oracle (learning to search), 374 | 13355 | positional encodings, 451 |
| 13317 | orthography, 204, 212 | 13356 | positive pointwise mutual information, 339 |
| 13318 | orthonormal matrix, 73 | 13357 | posterior, 478 |
| 13319 | out-of-vocabulary words, 191 | 13358 | power law, 14 |
| 13320 | outside recurrence, 244, 248 | 13359 | pragmatics, 359 |
| 13321 | overfit, 38 | 13360 | pre-trained word representations, 349 |
| 13322 | overfitting, 42 | 13361 | precision, 94, 478 |
| 13323 | overgeneration, 211, 221 | 13362 | precision-at- k , 95, 404 |
| | | 13363 | precision-recall curve, 422 |
| 13324 | parallel corpora, 439 | 13364 | precision-recall curves, 95 |
| 13325 | parameters, 482 | 13365 | predicate, 407 |
| 13326 | paraphrase, 354 | 13366 | predicative adjectives, 225 |
| 13327 | parent annotation, 251 | 13367 | predictive likelihood, 115 |
| 13328 | parsing, 217 | 13368 | prepositional phrase, 14, 225 |
| 13329 | part-of-speech, 185 | 13369 | primal form, 487 |
| 13330 | part-of-speech tagging, 155 | 13370 | principle of compositionality, 294 |
| 13331 | partially observable Markov decision process (POMDP), 472 | 13371 | prior, 478 |
| 13332 | | 13372 | prior expectation, 483 |
| 13333 | partially supervised learning, 256 | 13373 | probabilistic context-free grammars (PCFGs), 243 |
| 13334 | particle (part-of-speech), 189, 225 | 13374 | probabilistic models, 482 |
| 13335 | partition, 477 | 13375 | probabilistic topic model, 413 |
| 13336 | partition function, 174 | 13376 | probability density function, 481 |
| 13337 | parts-of-speech, 17 | | |

- 13379 probability distribution, 480
 13380 probability mass function, 97, 481
 13381 probability simplex, 33
 13382 processes, 426
 13383 production rules, 217
 13384 productivity, 203
 13385 projection function, 124
 13386 projectivity, 265
 13387 pronominal anaphora resolution, 359
 13388 pronoun, 188
 13389 PropBank, 314
 13390 proper nouns, 187
 13391 property (logic), 291
 13392 proposal distribution, 456
 13393 proposition, 426
 13394 propositions, 288, 289
 13395 prosody, 197
 13396 proto-roles, 313
 13397 pseudo-projective dependency parsing,
 275
 13398 pumping lemma, 215
 13400 pushdown automata, 217
 13401 pushdown automaton, 258
 13402 quadratic program, 47
 13403 quantifier, 291
 13404 quantifier, existential, 292
 13405 quantifier, universal, 292
 13406 quasi-Newton optimization, 54
 13407 question answering, 353, 409
 13408 random outcomes, 475
 13409 ranking, 410
 13410 ranking loss, 410
 13411 recall, 94, 478
 13412 recall-at- k , 404
 13413 receiver operating characteristic (ROC),
 95
 13415 rectified linear unit (ReLU), 65
 13416 recurrent neural network, 147
 13417 recurrent neural networks, 419
 13418 recursion, 14
 13419 recursive neural network, 401
 13420 recursive neural networks, 258, 351, 354
 13421 recursive production, 217
 13422 reference arguments, 329
 13423 reference resolution, 359
 13424 reference translations, 437
 13425 referent, 359
 13426 referring expression, 387
 13427 referring expressions, 359, 460
 13428 reflexive pronoun, 361
 13429 regression, 84
 13430 regular expression, 200
 13431 regular language, 200
 13432 regularization, 49
 13433 reification (events), 309
 13434 reinforcement learning, 373, 454
 13435 relation (logic), 297
 13436 relation extraction, 281, 327, 415
 13437 relations, 407
 13438 relations (information extraction), 407
 13439 relations (logic), 288
 13440 relative frequency estimate, 36, 138, 483
 13441 reranking, 257
 13442 residual networks, 66
 13443 retrofitting (word embeddings), 352
 13444 Rhetorical Structure Theory (RST), 394
 13445 rhetorical zones, 387
 13446 RIBES (translation metric), 439
 13447 ridge regression, 84
 13448 risk, 455
 13449 roll-in (reinforcement learning), 374
 13450 roll-out (reinforcement learning), 375
 13451 root (morpheme), 351
 13452 saddle point, 486
 13453 saddle points, 72
 13454 sample space, 475
 13455 satellite (RST), 395
 13456 satisfaction (logic), 293
 13457 scheduled sampling, 454

- | | | | |
|-------|---|-------|---|
| 13458 | schema, 407, 408, 424 | 13499 | slots (dialogue systems), 469 |
| 13459 | search error, 259, 372 | 13500 | smooth functions, 45 |
| 13460 | second-order dependency parsing, 267 | 13501 | smoothing, 38, 141 |
| 13461 | second-order logic, 292 | 13502 | soft K -means, 109 |
| 13462 | seed lexicon, 85 | 13503 | softmax, 63, 146, 420 |
| 13463 | segmented discourse representation theory (SDRT), 390 | 13504 | source domain, 122 |
| 13464 | self-attention, 450 | 13505 | source language, 435 |
| 13465 | self-training, 120 | 13506 | spanning tree, 262 |
| 13466 | semantic, 186 | 13507 | sparse matrix, 338 |
| 13467 | semantic concordance, 88 | 13508 | sparsity, 56 |
| 13468 | semantic parsing, 294 | 13509 | spectral learning, 129 |
| 13469 | semantic role, 310 | 13510 | speech acts, 197 |
| 13470 | Semantic role labeling, 310 | 13511 | speech recognition, 137 |
| 13471 | semantic role labeling, 416, 424 | 13512 | split constituents, 316 |
| 13472 | semantic underspecification, 305 | 13513 | spurious ambiguity, 230, 272, 278, 300 |
| 13473 | semantics, 250, 287 | 13514 | squashing function, 147 |
| 13474 | semi-supervised learning, 107, 117, 348 | 13515 | squashing functions, 65 |
| 13475 | semiring algebra, 182 | 13516 | stand-off annotations, 101 |
| 13476 | semiring notation, 207 | 13517 | Stanford Natural Language Inference corpus, 354 |
| 13477 | semisupervised, 88 | 13518 | statistical learning theory, 42 |
| 13478 | senses, 313 | 13519 | statistical machine translation, 436 |
| 13479 | sentence (logic), 292 | 13520 | statistical significance, 96 |
| 13480 | sentence compression, 467 | 13521 | stem, 18 |
| 13481 | sentence fusion, 468 | 13522 | stem (morphology), 203 |
| 13482 | sentence summarization, 466 | 13524 | stemmer, 90 |
| 13483 | sentiment, 81 | 13525 | step size, 53, 486 |
| 13484 | sentiment lexicon, 32 | 13526 | stochastic gradient descent, 45, 54 |
| 13485 | sequence-to-sequence, 447 | 13527 | stoplist, 92 |
| 13486 | shift-reduce parsing, 258 | 13528 | stopwords, 92 |
| 13487 | shifted positive pointwise mutual information, 346 | 13529 | string (formal language theory), 199 |
| 13488 | shortest-path algorithm, 205 | 13530 | string-to-tree translation, 446 |
| 13489 | sigmoid, 63 | 13531 | strong compositionality criterion (RST), 397 |
| 13490 | simplex, 127 | 13532 | strongly equivalent grammars, 218 |
| 13491 | singular value decomposition, 73 | 13534 | structure induction, 180 |
| 13492 | singular value decomposition (SVD), 116 | 13535 | structured attention, 465 |
| 13493 | singular vectors, 74 | 13536 | structured perceptron, 170 |
| 13494 | skipgram word embeddings, 344 | 13537 | structured prediction, 30 |
| 13495 | slack variables, 48 | 13538 | structured support vector machine, 171 |
| 13496 | slot filling, 420 | 13539 | subgradient, 45, 56 |

- 13540 subjectivity detection, 83
 13541 subordinating conjunctions, 188
 13542 subordinating discourse relations, 395
 13543 sum-product algorithm, 174
 13544 summarization, 137, 399
 13545 supersenses, 348
 13546 supervised machine learning, 32
 13547 support vector machine, 47
 13548 support vectors, 47
 13549 surface form, 211
 13550 surface realization, 459
 13551 synchronous context-free grammar, 445
 13552 synonymy, 86, 333
 13553 synset (synonym set), 86
 13554 synsets, 378
 13555 syntactic dependencies, 262
 13556 syntactic path, 317
 13557 syntactic-semantic grammar, 295
 13558 syntax, 185, 219, 287
- 13559 tagset, 186
 13560 tanh activation function, 65
 13561 target domain, 122
 13562 target language, 435
 13563 Targeted sentiment analysis, 83
 13564 tense, 187
 13565 terminal symbols (context-free grammars), 217
 13566 test set, 39, 107
 13568 test statistic, 96
 13569 text classification, 29
 13570 text mining, 17
 13571 text planning, 459
 13572 thematic roles, 311
 13573 third axiom of probability, 476
 13574 third-order dependency parsing, 268
 13575 TimeML, 425
 13576 tokenization, 88, 195
 13577 tokens, 34
 13578 topic models, 17
 13579 topic segmentation, 385
- 13580 trace (syntax), 230
 13581 training set, 32, 107
 13582 transduction, 200
 13583 transfer learning, 130
 13584 transformer architecture, 450
 13585 transition features, 158
 13586 transition system, 273
 13587 transition-based parsing, 233, 258
 13588 transitive closure, 369
 13589 translation error rate (TER), 439
 13590 translation model, 138
 13591 transliteration, 453
 13592 tree-adjoining grammar, 228
 13593 tree-to-string translation, 446
 13594 tree-to-tree translation, 446
 13595 treebank, 246
 13596 trellis, 160, 213
 13597 trigrams, 40
 13598 trilexical dependencies, 255
 13599 tropical semiring, 183, 207
 13600 True negative, 93
 13601 True positive, 93
 13602 true positive, 479
 13603 true positive rate, 95
 13604 truncated singular value decomposition, 338
 13605 truth conditions, 293
 13607 tuning set, 39, 93
 13608 Turing test, 15
 13609 two-tailed test, 97
 13610 type systems, 298
 13611 type-raising, 229, 298
 13612 types, 34
- 13613 unary closure, 237
 13614 unary productions, 218
 13615 underfit, 38
 13616 underflow, 33
 13617 undergeneration, 211, 221
 13618 Universal Dependencies, 186, 261
 13619 unlabeled precision, 238

- 13620 unlabeled recall, 238
13621 unseen word, 179
13622 unsupervised, 88
13623 unsupervised learning, 83, 107
13624 utterances, 197
- 13625 validation function (semantic parsing),
 304
13627 validity (logic), 293
13628 value function, 471
13629 value iteration, 471
13630 vanishing gradient, 65
13631 vanishing gradients, 149
13632 variable (AMR), 323
13633 variable (logic), 291
13634 variable, bound (logic), 292
13635 variable, free (logic), 291
13636 variance, 99
13637 variance (learning theory), 38
13638 variational autoencoder, 467
13639 Vauquois Pyramid, 436
13640 verb phrase, 221
13641 VerbNet, 312
13642 verbs, 187
13643 vertical Markovization, 251
- 13644 Viterbi algorithm, 158
13645 Viterbi variable, 160
13646 volition (semantics), 311
- 13647 WARP loss, 415
13648 weakly equivalent grammars, 218
13649 weight decay, 71
13650 weighted context-free grammar, 240
13651 weighted context-free grammars, 226,
 236
13653 weighted finite state acceptors, 205
13654 wide convolution, 77
13655 Wikification, 409
13656 Winograd schemas, 15
13657 word embedding, 67
13658 word embeddings, 61, 148, 334, 335
13659 word representations, 334
13660 word sense disambiguation, 85
13661 word senses, 85
13662 word tokens, 88
13663 WordNet, 20
13664 world model, 288
- 13665 yield (context-free grammars), 217
- 13666 zero-one loss, 44